

Contents

Copyright	11
Dedication	12
Preface	13
Reader's Guide	15
The complete reading path	15
Reading path for foundational interest	16
Reading path for empirical content evaluation	16
Reading path for philosophical positioning	17
Reading path for biology and medicine	17
Reading path for computational complexity and quantum information	18
How to use the appendices	18
How to use the glossary and bibliography	19
What to skip on a first reading	19
A note on density	20
Acknowledgments	20
Introduction	22
0.1 Two physics, or two projections?	22
0.2 An informal thought experiment	24
0.3 The framework's central claim	26
0.4 The intellectual context	27
0.5 The empirical record	29
0.6 The book's structure	31
Chapter 1	34
Embedded Observation and the Characterization Theorem	34
1.1 The problem and the theorem	34
1.2 The starting point	35
1.3 The observer's blind spot	37
1.4 Partition-relativity	39
1.5 P-indivisibility	40
1.6 The two routes to quantum mechanics	45
1.7 Bell-inequality violations	51
1.8 The characterization theorem	52
1.9 What the theorem says	55
1.10 The three structural limits	56
1.11 The card table and the chessboard	60
Chapter 2	64

The Substratum	64
2.1 What this chapter constructs	64
2.2 The substratum as a finite lossless memory	65
2.3 What the substratum is and is not	66
2.4 The reconstruction theorem	68
2.5 The substratum gauge group	73
2.6 The three projections	76
2.7 The QM-GR incompatibility as a category error	79
2.8 What follows	81
 Chapter 3	 82
Hierarchical Structural Realism	82
3.1 What this chapter develops	82
3.2 The observation axis	83
3.3 The observer-admission distinction in more detail	87
3.4 The gauge axis	88
3.5 Orthogonality, multi-level realism, and stratified predictions	91
3.6 Gauge freedom and observational incompleteness	93
3.7 No-go theorems at different descriptive levels	96
3.7.1 The Maudlin test	100
3.8 The philosophical lineage	104
3.9 What follows	109
 Chapter 4	 110
Methodology	110
4.1 What this chapter develops	111
4.2 The single-foundational-commitment pattern	112
4.3 Classification of claims by derivational rigor	115
4.4 Engaging the no-go results through one architecture	118
4.5 Pre-registered falsification conditions	121
4.6 What does not falsify and regimes of deliberate silence	125
4.7 What the framework derives, and what it compresses	128
4.8 What follows	130
 Chapter 5	 131
The Emergence of Gauge Structure	131
5.1 What this chapter develops	132
5.2 The minimal substratum and the wave equation	133
5.3 Selection of $d = 3$	136
5.4 The emergent QFT: factorization and multi-component dynamics	138
5.5 The cubic decomposition and the Standard Model gauge group	141
5.6 Chirality, anomaly cancellation, and hypercharges	144
5.7 What follows	148

Chapter 6	149
The Matter Content and Quantitative Predictions	149
6.1 What this chapter develops	149
6.2 Three generations and the composite Higgs	151
6.3 The strong-CP problem and $\bar{\theta} = 0$	153
6.4 Gauge couplings and the no-GUT prediction	156
6.5 The Cabibbo angle and CKM structure	159
6.6 The Koide relation and the lepton mass chain	162
6.7 PMNS mixing angles	168
6.8 Summary table and chapter close	171
Chapter 7	177
The Gravitational Sector	177
7.1 What this chapter develops	177
7.2 The cosmological horizon as causal partition	179
7.3 The emergent action scale $\hbar$	181
7.4 Self-consistency and the discreteness scale	185
7.5 The Bekenstein-Hawking entropy	186
7.6 The Page curve and the information paradox	188
7.7 Dissolution of the cosmological constant problem	192
7.8 Dark energy in running-vacuum form	195
7.9 The dark sector and the MOND acceleration scale	198
7.10 Summary of gravitational predictions and chapter close	201
Chapter 8	204
JUNO and the Neutrino Sector	204
8.1 What this chapter develops	205
8.2 The JUNO experiment	206
8.3 The cubic-group flavor structure	207
8.4 The U-parity grading	209
8.5 The sum rule and structural relation	212
8.6 The three numerical predictions	214
8.7 Comparison with TM1/TM2 column-preservation patterns	216
8.8 JUNO design-lifetime sensitivity and falsification	218
8.9 The coupled DESI DR2 result	220
8.10 Chapter close: the framework's empirical record in the neutrino sector	221
Chapter 9	223
Universality Classes and External Convergence	223
9.1 What this chapter develops	223
9.2 The framework's hierarchical structural realism	225
9.3 Tier 1 results as universality-class content	228

9.4 Tier 2 and Tier 3 stratification	231
9.5 The framework’s relationship to string theory	234
9.6 The string landscape multiplicity problem	236
9.6b Foundational programs claiming parameter-free SM retrodictions	237
9.6c A gravitational-sector sibling: conformal gravity	239
9.7 The Danielson-Satishchandran-Wald external convergence	242
9.8 Adjacent observer-emergent literature	244
9.9 Chapter close: the framework’s position in the broader landscape .	246
Chapter 10	248
Chemistry	248
10.1 What this chapter develops	248
10.2 The structural chain from substratum to orbital chemistry	250
10.3 Beyond carbon: four additional structural results	253
10.4 The scale hierarchy and the permanence of atoms	255
10.5 The thermal window for chemistry	256
10.6 Water as biological solvent	257
10.7 The chirality chain	259
10.8 Fine-tuning and the anthropic principle	261
Chapter 11	263
The Origin of Life	263
11.1 What this chapter develops	263
11.2 The autocatalytic argument	265
11.3 Prebiotic chemistry pathways	266
11.4 Self-replication as read-write cycling	269
11.5 From template replication to evolution	270
11.6 C1-C3 at the molecular scale	272
11.7 RNA as the unique single-molecule C1-C3 system	274
11.8 What remains genuinely open	276
Chapter 12	277
Evolution, Intelligence, and Self-Reference	277
12.1 What this chapter develops	277
12.2 Evolution as a multi-scale C1-C3 system	279
12.3 Information processing as fitness-enhancing	281
12.4 Major evolutionary transitions	282
12.5 Neural computation as structurally possible	284
12.6 General intelligence as structurally unconstrained but not guar- anteed	286
12.7 The chain from intelligence to AI	287
12.8 The self-referential loop and structural limits	289
Chapter 13	294

Quantum Biology	294
13.1 What this chapter develops	294
13.2 The framework’s reframing of quantum biology	295
13.3 C1-C3 in proteins	297
13.4 Photosynthesis as energy transfer through C1-C3 architectures	299
13.5 Magnetoreception via cryptochrome radical pairs	300
13.6 Enzyme catalysis as the framework’s clearest test case	302
13.7 The partially-quantum regime in biology	303
13.8 Implications for drug design	305
13.9 Chapter close	306
 Chapter 14	 307
Quantum Computing	307
14.1 What this chapter develops	307
14.2 The framework’s content in computational complexity	309
14.3 BQP is available: the lower bound	311
14.4 BQP cannot be exceeded: the upper bound	312
14.5 The BQP characterization theorem	314
14.6 The Extended Church-Turing Thesis as a conditional theorem	315
14.7 The BPP-BQP gap as cost of partial observation	317
14.8 Internal structure of BQP and the P vs NP question	319
14.9 Beyond axioms: research directions opened by the substratum structure	323
14.9.1 Geometric structure on the substratum lattice	323
14.9.2 The BQP-BPP gap, hidden group structure, and the SU(2)-compatibility refinement	324
14.9.3 Gauge equivalence and complexity-theoretic reductions	325
14.9.4 Conditions and the right kind of research question	326
14.10 Chapter close	326
 Chapter 15	 329
Quantum Engineering — Hardware, Software, and the BEC Ex- periment	329
15.1 What this chapter develops	329
15.2 The environment as programmable quantum memory	330
15.3 Superconducting qubits and TLS noise	332
15.4 Correlation-aware error correction	334
15.5 Coherence extension through engineered baths	335
15.6 Quantum sensing with NV centers	336
15.7 The partially-quantum regime in engineered systems	338
15.8 The BEC analogue-gravity experiment	340
15.9 Chapter close	343
 Chapter 16	 345

Medicine	345
16.1 What this chapter develops	345
16.2 The enzyme as a read-write system	347
16.3 Cancer pharmacology	349
16.4 Neurodegeneration	350
16.5 Antibiotic resistance	353
16.6 Immunotherapy and T-cell exhaustion	354
16.7 Cardiac pharmacology	355
16.8 Autoimmune disease	356
16.9 Epigenetics as biological hidden sector	357
16.10 The substratum-emergent operator distinction at chromatin	359
16.11 Reader, writer, and eraser pharmacology	361
16.12 The wider therapeutic window for reader-writer disorders	362
16.13 Chapter close: memory asymmetry as unifying principle	364
 Chapter 17	 366
Bioinformatics	366
17.1 What this chapter develops	367
17.2 The information-theoretic ceiling on transcriptome-only methods	368
17.3 Trajectory inference and the RNA velocity paradigm	369
17.4 Gene regulatory network inference	371
17.5 Perturbation prediction and the failure of scaling	372
17.6 Multi-omics integration and timescale-aware methods	373
17.7 Cancer methylome classification as positive prediction	374
17.8 Variant pathogenicity for reader/writer/eraser genes	376
17.9 Empirical confirmations in evolutionary biology	377
17.10 Chapter close: the structural roadmap forward	381
 Chapter 18	 383
Beyond Quantum Mechanics and General Relativity	383
18.1 What this chapter develops	383
18.2 The Born rule as dynamical equilibrium	385
18.3 The quantum-classical boundary and the partially-quantum regime	388
18.4 The epistemic character of quantum randomness	390
18.5 The quantum gravity dissolution as cumulative case	392
18.6 Gravitational wave echoes	395
18.7 Finite information and Poincaré recurrence	397
18.8 The arrow of time	399
18.9 Cosmic spatial topology	401
18.10 The framework's position on consciousness	404
18.11 Chapter close	409
 Chapter 19	 411
Resolved and Open Problems in Fundamental Physics	411

19.1	What this chapter develops	411
19.2	Problems resolved	412
19.2.1	The quantum gravity problem	412
19.2.2	The cosmological constant problem	413
19.2.3	The nature of dark matter	413
19.2.4	The nature of dark energy	414
19.2.5	The measurement problem	415
19.2.6	The black hole information paradox	415
19.2.7	The strong CP problem	416
19.2.8	The hierarchy problem	416
19.2.9	The origin of the gauge group	417
19.2.10	The generation puzzle	417
19.2.11	The dark sector budget	418
19.2.12	The origin of the Born rule	419
19.3	Problems open within the framework	419
19.3.1	The flavor problem	419
19.3.2	The Hubble tension	421
19.3.3	Baryogenesis	422
19.3.4	Inflation and initial conditions	424
19.3.5	The initial state of (S, φ)	425
19.3.6	Is the discrete structural residue genuinely non-tautological?	425
19.3.7	Radiative stability of emergent Lorentz invariance (rotational sector)	427
19.3.8	The chirality hypothesis $H\text{-}\chi$	432
19.4	Chapter close: summary table and what the framework does not address	432
Appendix A		437
Prediction Status Table		437
A.1	What this appendix provides	437
A.2	Classification scheme	438
A.3	Four-layer descriptive framing	439
A.4	Master prediction tables	440
A.4.1	Standard Model observables	440
A.4.2	Gravitational-sector observables	443
A.4.3	Neutrino-sector observables	445
A.4.4	Biology and medicine predictions	446
A.5	Forward predictions and falsification targets	448
A.6	Framework's epistemic position summary	451
Appendix B		453
Mathematical Derivations		453
B.1	What this appendix develops	453
B.2	The Stinespring construction for emergent quantum mechanics	454

B.3 The trace-out as a Jordan-Chevalley projection	459
B.4 The gap equation for $\hbar$	462
B.5 Anomaly cancellation and unique hypercharges	465
B.6 Gauge-theoretic derivations	467
B.6b Bravais-lattice uniqueness lemma	469
B.6c T_1 generation-assignment uniqueness lemma	471
B.7 The lemma chain for emergent quantum mechanics	474
Appendix C	477
Common Objections and Framework Responses	477
C.1 What this appendix develops	477
C.2 Bell's theorem and the framework's relationship to hidden-variable theories	478
C.3 The no-go theorems: PBR, Frauchiger-Renner, Colbeck-Renner, Bong-Wiseman	480
C.4 The stochastic-quantum bridge: Barandes and Stinespring	483
C.5 The necessity direction of the characterization theorem	484
C.6 Finiteness: is Axiom 2 too strong?	485
C.7 Computational implementability and Wolpert's theorem	487
C.8 Relationship to standard interpretations: MWI, Bohmian, QBism	488
C.9 The substratum-emergent operator distinction: substantive or just renaming?	490
C.10 Circularity concerns: the gap equation deriving $\hbar$	492
C.11 The framework's epistemic status as a structural framework	493
C.11.1 Is the derivation gerrymandered? The selector-motivation audit	495
C.12 Meta-objections: credentials, scope, ambition	496
C.13 Closing remarks	497
Glossary	499
Bibliography	509
Framework-internal citations	510
Foundations of quantum mechanics and interpretation	510
Standard Model physics and particle phenomenology	512
Gravitational physics, cosmology, and quantum gravity	513
Chemistry and origin of life	515
Evolution and molecular biology	515
Medicine, pharmacology, and clinical biology	516
Bioinformatics and computational biology	518
Quantum biology and biological information processing	520
Quantum information, computing, and complexity	521
Mathematical foundations and computability	521

Historical and philosophical works	522
--	-----

The Incompleteness of Observation

*A Unified Framework from
Quantum Mechanics to Computational Biology*

Alex Maybaum

First complete draft

May 2026

Companion technical papers:
The Observational Incompleteness Framework
Zenodo DOI: [10.5281/zenodo.19060318](https://doi.org/10.5281/zenodo.19060318)

Contents

Copyright

© 2026 Alex Maybaum

All rights reserved.

This is a working manuscript representing the first complete draft of *The Incompleteness of Observation: A Unified Framework from Quantum Mechanics to Computational Biology*. The framework's authoritative technical content is in the companion papers (Zenodo DOI 10.5281/zenodo.19060318). For framework-internal derivations, the technical papers should be consulted as the primary references.

Citation: Maybaum, A. (2026). *The Incompleteness of Observation: A Unified Framework from Quantum Mechanics to Computational Biology*. Working manuscript.

No part of this work may be reproduced or distributed without the author's permission, with the standard exception for brief quotations in critical reviews and certain other non-commercial uses permitted by copyright law.

Dedication

[Author's dedication to be added.]

Preface

This book argues a heterodox position: that quantum mechanics is not fundamental but emergent, derivable from a deterministic substratum through three structural conditions that any laboratory observer in our universe satisfies. The argument is constructive — explicit derivations replace philosophical assertion — and the empirical record now exists to constrain the framework’s content quantitatively. The book is the integrated synthesis of work developed in technical papers over the past several years, organized into a single derivation chain from the framework’s foundational commitment to its specific quantitative predictions across multiple sectors of physics, chemistry, biology, and medicine.

The heterodox character of the position bears emphasis. The standard reading of quantum mechanics treats the wave function, the Schrödinger equation, and the Born rule as foundational — the basic objects of physics from which everything else follows. The framework reverses this ordering. It treats observation as foundational, the structural conditions of observation as derived consequences, and quantum mechanics as the dynamics that any observer satisfying those structural conditions must use. This reversal is not a matter of interpretive preference. The framework’s content is a mathematical derivation: given the foundational commitment plus three structural conditions, quantum mechanics follows as a theorem.

The framework’s value depends on what follows from this reversal — and the book documents exactly that. The cosmological constant discrepancy of 10^{122} orders of magnitude — conventionally the worst quantitative prediction in physics — becomes a calculable compression ratio. The Standard Model gauge group $SU(3) \times SU(2) \times U(1)$, the three matter generations, and the anomaly-free hypercharges are reproduced as forced consequences of a cubic-lattice commitment that is itself empirically selected by the observed $SU(3)$ — selected-and-reproduced rather than derived from nothing. The Bekenstein-Hawking factor of $1/4$ in black-hole entropy becomes a derivation. Twenty-two Standard Model observables, thirteen gravitational-sector entries, and seven neutrino-sector entries follow — seven of the Standard Model set as parameter-free retrodictions, the remainder as layered, chained, or fitted, each classified explicitly in Appendix A. The JUNO neutrino mixing angle is matched at 0.07σ from a prediction made before the measurement. Hawking’s area theorem is confirmed at 99.999% confidence by the LIGO/Virgo/KAGRA observation of GW250114 in September 2025 (the derived $1/4$ entropy coefficient itself has no direct test). The wider therapeutic window for reader-writer disorders is consistent with the Rett-syndrome reversibility of Guy et al. 2007 (a retrodiction, since that result predates the framework).

These are not interpretive achievements; they are quantitative empirical content. The framework is therefore neither a metaphysical reframing of standard physics nor a speculative alternative to it. It is a derivation chain from one foundational commitment to specific quantitative predictions, with substantial

empirical content already in place and additional content awaiting near-term experimental tests.

The book is organized as a complete first treatment of the framework’s content for readers with a working background in theoretical physics, mathematical physics, or technical biology. The opening four chapters develop the foundational mathematics. The next five chapters develop the framework’s content in fundamental physics. The third part traces the emergence cascade from atoms to chemistry to life to intelligence. The fourth part applies the framework to working scientific fields — quantum biology, quantum computing, quantum engineering, medicine, bioinformatics — with substantial empirical predictions in each. The closing chapters develop the framework’s forward predictions and its position on the major open problems of fundamental physics. Three appendices provide a prediction status table, mathematical derivations supporting the main text, and a systematic engagement with the most frequent objections to the framework’s commitments.

The book is intended to make the framework’s structure, derivation chains, and specific predictions available for evaluation by the broader scientific community. The framework’s correctness is an empirical question, and the framework has the empirical content to be evaluated. The book’s role is to make that content accessible.

A note on prerequisites. The technical chapters of Parts I and II require working familiarity with quantum mechanics, general relativity, and group theory. The applied chapters of Parts III and IV assume the technical content as derived results and add domain-specific background as needed. The Introduction and the closing chapters can be read by anyone with general scientific background; the technical content can be consulted as needed when specific claims become load-bearing for an argument. A glossary at the end of the back matter provides definitions for the framework’s distinctive terminology and the major technical terms used throughout.

A note on completeness. The framework is not a theory of everything. Several major problems in fundamental physics are scope limits within the framework’s content: the flavor problem (the specific values of fermion masses and CKM mixing parameters), baryogenesis (the specific baryon-to-photon ratio — solution-specific, downstream of the framework’s flavor-sector CP input), inflation, the Hubble tension, and the cosmological initial conditions. Chapter 19 develops the framework’s position on these open problems systematically. The framework provides structural answers where it can and identifies its scope limits where it cannot.

A note on terminology. The phrase “the framework’s content” appears throughout the book to distinguish the framework’s claims from claims made by standard physics or by other unification programs. The phrase is technical: the framework’s content is what follows from the framework’s commitments through derivation, as distinct from interpretive readings of those derivations or specula-

tive extensions. Where the framework’s content overlaps with standard physics (e.g., the Standard Model’s gauge structure), the overlap is noted; where the framework’s content extends beyond standard physics (e.g., the Born rule’s transient violations in the very early universe), the extension is identified.

The framework is the work of one author, but it depends on the work of many. The intellectual debts to the foundations of physics community — to Bell, Wolpert, ’t Hooft, Maldacena, Witten, the Wald group, and the broader observer-emergent literature — are substantial and are acknowledged at the appropriate points in the text. The framework’s technical methods rely on standard operator algebra (Stinespring, Birkhoff-von Neumann) and on recent stochastic-quantum correspondence work (Barandes). The framework’s empirical content depends on the experimental and observational programs of JUNO, LIGO/Virgo/KAGRA, DESI, the Standard Model precision program, and the broader astrophysical and biological communities. The acknowledgments at the back of the front matter name specific contributors; the bibliography at the back of the book provides the full citation record.

The position the book argues will strike many readers as too ambitious. The framework’s claim is that quantum mechanics is not fundamental, that a single derivation chain produces specific quantitative predictions across physics, chemistry, biology, and medicine, and that the empirical record already supports the framework’s content — with parameter-free retrodictions in fundamental physics (JUNO at 0.07σ ; the Standard Model constants) and pre-registered forward predictions awaiting data (DESI Y5, JUNO design-lifetime precision, and others). Each of these claims is heterodox; their conjunction is more so. The book makes the conjunction explicit and supports it with the available evidence. The reader is invited to evaluate the case as it is presented and to judge for themselves.

— Alex Maybaum 2026

Reader’s Guide

The book has four parts and three appendices, organized around the framework’s derivation chain from foundational commitment to applied predictions. Different readers approach the framework with different interests, technical backgrounds, and goals. This guide suggests reading paths for several common situations.

The complete reading path

Readers approaching the book from start to finish — the framework’s intended presentation — will find the technical content of Parts I and II load-bearing for

the cascade developments of Parts III and IV. The argument’s intellectual coherence depends on the reader having absorbed the central theorem (Chapter 1), the substratum construction (Chapter 2), the structural realism framing (Chapter 3), the methodology (Chapter 4), the Standard Model derivation (Chapters 5-6), and the gravitational sector (Chapter 7) before engaging with the emergence cascade (Chapters 10-12) and the applied chapters (Chapters 13-17). The complete reading path is therefore the recommended path for readers wanting to evaluate the framework’s full content.

Estimated reading time for the complete path: approximately 50-80 hours for technical readers, longer for readers consulting the appendices and external references extensively.

Reading path for foundational interest

Readers primarily interested in the framework’s foundational claim — that quantum mechanics is derivable from a deterministic substratum under three structural conditions — should focus on the following chapters and sections.

Essential. Chapter 0 (the framing). Chapter 1 (the central theorem and characterization theorem). Chapter 2 (the substratum construction). Chapter 7 (the gravitational sector, particularly §§7.3, 7.5, 7.7 on $\hbar$, the Bekenstein-Hawking entropy, and the cosmological constant dissolution). Chapter 8 (JUNO as concrete empirical anchor). Appendix B (mathematical derivations) for technical detail.

Recommended. Chapter 3 (structural realism). Appendix C (engagement with no-go theorems and standard-interpretation objections). Chapter 18 §18.5 (the cumulative quantum gravity dissolution).

Optional for this path. The applied chapters (13-17) develop empirical content that supports the framework but is not load-bearing for the foundational claim. Chapter 19 provides systematic positioning against open problems.

This reading path takes approximately 15-25 hours for technical readers.

Reading path for empirical content evaluation

Readers approaching the framework with primary interest in evaluating its empirical record — the JUNO retrodictive match (0.07σ), the Standard Model retrodictions, the GW250114 area-theorem consistency, the Rett syndrome retrodiction — should focus on the framework’s specific quantitative content.

Essential. Chapter 0 (framing, particularly §0.5 on the empirical record). Chapter 6 (twenty-two classified Standard Model predictions). Chapter 7 (gravitational sector predictions, particularly §§7.5, 7.9 on Bekenstein-Hawking and MOND). Chapter 8 (JUNO neutrino sector). Chapter 16 (medical predictions, particularly §16.12 on Rett syndrome). Appendix A (prediction status table for systematic reference).

Recommended. Chapter 5 (gauge structure derivation, the chain that produces the SM predictions). Chapter 9 (universality classes, framing the empirical record). Chapter 18 (forward predictions for upcoming experiments).

Optional. Foundational content in Chapters 1-4 can be consulted when specific empirical claims raise foundational questions. The chemistry and origin-of-life chapters (10-11) are structural rather than empirical.

This reading path takes approximately 10-15 hours for technical readers.

Reading path for philosophical positioning

Readers approaching the framework from philosophy of physics, philosophy of science, or the foundations community should focus on the framework's interpretive structure.

Essential. Chapter 0 (framing). Chapter 3 (hierarchical structural realism). Chapter 4 (methodology and falsification conditions). Chapter 9 (universality classes and convergence with adjacent programs including CLPW, Maldacena, HUZ, DSW). Chapter 18 (forward predictions and the consciousness position). Appendix C (engagement with no-go theorems and standard interpretations).

Recommended. Chapter 1 (central theorem) for the formal grounding of philosophical claims. Chapter 2 (substratum) for the framework's ontological commitments. Chapter 19 (positioning against open problems) for the framework's epistemic scope.

Optional for this path. The empirical chapters (6-8, 13-17) can be consulted when specific predictions become relevant to philosophical evaluation.

This reading path takes approximately 15-25 hours.

Reading path for biology and medicine

Readers approaching the framework from biology, medicine, or applied biological sciences can begin with the applied content and consult the technical chapters as needed.

Essential. Chapter 0 (framing). Chapter 13 (quantum biology). Chapter 16 (medicine, particularly §§16.9-16.13 on epigenetics and the Rett syndrome confirmation). Chapter 17 (bioinformatics, particularly §17.9 on evolutionary biology empirical confirmations).

Recommended. Chapter 10 (chemistry derivation, providing the foundations). Chapter 11 (origin of life). Chapter 12 (evolution and intelligence). Chapter 15 §§15.5-15.7 (the partially-quantum regime, which connects engineering to biology).

Optional. The fundamental physics chapters (5-9) can be consulted when biological claims rest on physics derivations. The technical foundations (Chapters

1-4) can be consulted for the framework’s structural commitments at the molecular scale.

This reading path takes approximately 15-25 hours.

Reading path for computational complexity and quantum information

Readers approaching the framework from computer science, quantum information, or quantum hardware engineering can focus on the computational and engineering content.

Essential. Chapter 0 (framing). Chapter 14 (BQP characterization theorem). Chapter 15 (quantum engineering, including correlation-aware error correction, NV sensing, and the BEC analogue-gravity test). Appendix B §B.2 (Stinespring construction) and §B.7 (lemma chain).

Recommended. Chapter 1 (foundational theorem) for the structural basis of the BQP result. Chapter 9 (universality classes) for the framework’s relationship to alternative computational architectures. Appendix C §C.7 (Wolpert’s theorem and computational implementability).

Optional. The empirical chapters in fundamental physics (6-9) and in biology and medicine (13, 16-17) provide the framework’s broader empirical record but are not load-bearing for the computational content.

This reading path takes approximately 10-15 hours.

How to use the appendices

The three appendices serve different purposes and can be consulted independently.

Appendix A (Prediction Status Table). A systematic reference resource organizing the framework’s predictions by classification (Structural / Computational / Linked / Refined / Pending / Mechanism / Empirical), confidence level, current empirical match, and open issues. Use as a lookup table when evaluating specific claims or planning empirical work against the framework. The appendix does not need to be read sequentially.

Appendix B (Mathematical Derivations). Complete derivations for the framework’s key technical results: the Stinespring construction, the Jordan-Chevalley projection, the gap equation deriving $\hbar$, the anomaly cancellation calculation, the gauge-theoretic derivations, and the lemma chain for the characterization theorem. Use when the main-text derivations are too compressed for the reader’s purposes. The appendix is designed for readers requiring full technical content; readers without that requirement can rely on the main-chapter derivations.

Appendix C (Common Objections and Framework Responses). Systematic engagement with the most predictable objections to the framework’s commitments: Bell’s theorem, no-go theorems (PBR, Frauchiger-Renner, Colbeck-Renner, Bong-Wiseman), the stochastic-quantum bridge, the necessity direction of the characterization theorem, the finiteness axiom, computational implementability, the relationship to standard interpretations (MWI, Bohmian, QBism), the substratum-emergent operator distinction concern, circularity in the gap equation, the framework’s epistemic status, and meta-objections about scope and credentials. Use as a reference resource for skeptical readers or for engaging the framework against specific objections.

How to use the glossary and bibliography

Glossary. Located at the end of the back matter. Defines the framework’s distinctive terminology (C1-C3, P-indivisibility, substratum, partition-relativity, etc.) and the major technical concepts used throughout the book. Cross-references in glossary entries point to the chapter and section where each concept is most fully developed. Use as a reference when terminology is unfamiliar.

Bibliography. Located at the end of the back matter. Organized by topical area for ease of navigation. Cite using the author-year convention used in the text. The framework’s own technical companion papers (available in the Zenodo release at DOI 10.5281/zenodo.19060318) are the primary references for framework-internal derivations.

What to skip on a first reading

Several sections are structurally optional on a first reading, even within the complete reading path.

- **Chapter 4 §4.6** (regimes of deliberate silence) is methodological commentary that can be returned to after the main argument is absorbed.
- **Chapter 9 §§9.5-9.6** (relationship to string theory and the string landscape multiplicity problem) is positioning against the broader physics landscape; readers without interest in string theory can skim.
- **Chapter 12 §12.8** (the self-referential loop) is philosophical commentary; readers without interest in self-reference can skim.
- **Chapter 18 §18.10** (consciousness) is philosophical commentary; readers without interest in consciousness questions can skim.
- **Appendix B §B.3** (Jordan-Chevalley projection) is the most technical of the appendix sections; readers without background in algebraic geometry can rely on the main-text §5.4-5.5 treatment.

These sections add depth to the framework’s content but are not load-bearing for the main argument.

A note on density

The framework is the work of a single author developed primarily in technical-paper form. The book represents the integrated treatment for a broader audience, but the technical density varies. Chapters 1-2, 5-7, and the appendices are demanding and reward careful reading. Chapters 3-4, 9, and 18 are philosophical and require attention to the framework’s distinctive vocabulary. Chapters 10-12 and 13-17 are accessible to readers with general scientific background and assume the technical content of Parts I and II as derived results.

Readers finding any chapter too dense should not hesitate to skip ahead and return later. The framework’s content is sufficiently modular that most chapters can be productively read with limited preceding context, particularly with the glossary as support.

The framework is unfinished in its empirical exposure but complete in its structural content. The book makes the structural content available for evaluation. The empirical evaluation depends on experiments now being done and on those forthcoming.

Acknowledgments

The framework developed in this book rests on the work of many physicists, mathematicians, biologists, and philosophers whose contributions made the present synthesis possible. The framework’s intellectual debts span multiple decades and several scientific communities, and the acknowledgments here cannot be complete. What follows names the most direct contributors and indicates the broader communities whose work has been load-bearing for the framework.

Foundations of physics. The framework’s central theorem owes its mathematical apparatus to standard operator-algebraic results going back to the 1950s, particularly Stinespring’s dilation theorem (1955) and the Birkhoff-von Neumann theorem (1946). Recent work on stochastic-quantum correspondence by Jacob Barandes (2023-2024) provided the second independent bridge from deterministic substrata to emergent quantum mechanics. The framework’s positioning against the no-go theorems engages the work of Bell, Kochen and Specker, Pusey-Barrett-Rudolph, Frauchiger and Renner, Colbeck and Renner, and Bong et al. The framework’s anti-Laplacian commitments draw on Wolpert’s 1992 work on physical limits of inference. The deeper intellectual lineage extends through Spinoza’s attribute theory and Nicholas of Cusa’s *docta ignorantia* to contemporary observer-emergent foundational work.

Observer-emergent and horizon decoherence programs. The framework’s structural alignment with contemporary observer-emergent foundational work has been a substantial source of intellectual confidence in the framework’s commitments. The work of Chandrasekaran, Longo, Penington, and Witten on

observer-incorporated entropy in de Sitter space (2022), Maldacena on the role of the observer’s clock in cosmological partition functions (2024), Harlow, Ustyuk, and Zhao on observer-dependent Hilbert spaces in closed universes (2025), and Danielson, Satishchandran, and Wald on horizon decoherence (2022) has provided independent corroboration of the framework’s universality-class content from directions the framework’s development could not have anticipated.

Standard Model and gravitational physics. The framework’s quantitative content in fundamental physics rests on the experimental and theoretical programs of the Standard Model precision community, the LIGO/Virgo/KAGRA collaboration (particularly the GW250114 analysis confirming the Bekenstein-Hawking area law), the JUNO collaboration (whose November 2025 measurement of $\sin^2 \theta_{12}$ matches the framework’s prediction at 0.07σ), the DESI collaboration, the Planck collaboration, and the LUX-ZEPLIN dark matter direct-detection program. The framework’s debt to these communities is substantial: without their data, the framework’s content would be structural speculation rather than empirically anchored derivation.

Chemistry, origin of life, and biology. The framework’s emergence cascade through chemistry to biology draws on Stuart Kauffman’s autocatalytic threshold analysis, the prebiotic chemistry programs of the past seven decades (Miller, Urey, Ferris, Damer, Deamer, and others), Cech and Altman’s discovery of catalytic RNA, the Soai autocatalytic reaction, and the broader RNA-world literature. The framework’s content in evolution draws on Maynard Smith and Szathmáry’s analysis of major evolutionary transitions, Bedford and Hartl on molecular clock overdispersion, Tao et al. on universal rate autocorrelation across the tree of life, Wiser, Ribick, and Lenski on LTEE power-law fitness trajectories, and Eldredge and Gould on punctuated equilibrium.

Medicine, epigenetics, and pharmacology. The framework’s clinical content rests on the work of many groups in chromatin biology, particularly Bird, Cech, Lindquist, and the broader Rett syndrome research community (with the Guy, Bird et al. 2007 paper providing the framework’s strongest empirical anchor outside fundamental physics). The framework’s pharmacological content engages the work of multiple groups in cancer pharmacology (particularly on Chk1 inhibitors and gemcitabine combinations), neurodegeneration (particularly on CaMKII and PTSD reconsolidation), antibiotic resistance (particularly on RecA and persister cells), immunotherapy (T-cell exhaustion), and the broader pharmacological literature.

Bioinformatics and computational biology. The framework’s bioinformatics content engages the substantial recent literature on transcriptome-based foundation models (Csendes et al., Kedzierska et al., and the broader scGPT, Geneformer, and scFoundation programs), single-cell trajectory inference (RNA velocity by La Manno et al., chromatin velocity by Tedesco et al.), gene regulatory network inference (Stolovitzky, Marbach, Bravo González-Blas, and others), and cancer methylome classification (Capper et al. and the DKFZ program).

Quantum information and computing. The framework’s content in computational complexity draws on the foundational work of Bernstein, Vazirani, Bennett, Brassard, and the broader BQP literature. The framework’s engineering content engages the work of the superconducting qubit community (with Wilen et al., McEwen et al., and Google Quantum AI providing the empirical foundations for correlated noise structure) and the NV-center quantum sensing community. The BEC analogue-gravity experimental program developed by Jeff Steinhauer and collaborators provides the cleanest near-term test of the framework’s foundational content.

The framework’s working repository. The framework’s content has been developed in technical-paper form over the past several years, with the release at DOI 10.5281/zenodo.19060318 representing the current state. The repository’s collaborators and reviewers — too numerous to name individually — have contributed substantially to the framework’s refinement.

Personal acknowledgments. [Author’s personal acknowledgments — colleagues, family, institutional support, funding sources — to be added.]

Errors. Any remaining errors in the framework’s content or its presentation are the author’s. The book is the integrated synthesis of work developed in technical papers, and the integration introduces the possibility of inconsistencies between the book’s presentation and the underlying technical papers. The technical papers should be consulted as the primary references for framework-internal derivations; the book’s role is to make the framework’s content accessible.

— Alex Maybaum

Introduction

0.1 Two physics, or two projections?

Physics has two fundamental theories. Quantum mechanics describes the microscopic world — atoms, particles, photons, the interactions of fields at small scales. Quantitative predictions of quantum electrodynamics agree with experiment to twelve decimal places; the Standard Model has been tested at the LHC to energies above 13 TeV without finding a single deviation. General relativity describes the macroscopic world — spacetime curvature, gravitational waves, the expansion of the universe. Its predictions have survived every test from the precession of Mercury to the LIGO observations of binary black-hole mergers. Each theory, within its domain, is the most successful physical theory in human history.

They are incompatible. Quantum mechanics has no preferred temporal foliation; general relativity has one at every point. Quantum mechanics is linear; general relativity is not. Quantum mechanics treats measurement as a primitive operation; general relativity treats it as a process within the dynamics. The two theories give incompatible answers to elementary questions. The most striking of these is the energy of empty space. Quantum field theory predicts a vacuum energy of roughly 10^{113} joules per cubic meter, summed over zero-point modes up to the Planck scale. General relativity, read off the observed expansion of the universe, gives about 6×10^{-10} joules per cubic meter. The ratio is 10^{122} .

For comparison: the observable universe contains about 10^{80} atoms. The discrepancy between what quantum mechanics says the vacuum should weigh and what general relativity measures is forty-two orders of magnitude larger than the number of atoms in the universe. It is the worst quantitative prediction in the history of physics. Physicists have been trying to explain it for fifty years.

The standard reading of this failure — and of related failures that have accumulated over the past several decades — is that the existing theories are incomplete, and that a deeper unified theory is required. The persistent failures are then research opportunities. String theory, loop quantum gravity, causal dynamical triangulations, asymptotic safety, and other unification programs have each pursued the project of finding the missing deeper theory. None has produced a quantitatively confirmed prediction in the regime where it would distinguish itself from existing physics. After a half-century of intensive work, the unification program remains unsettled.

Adjacent to the unification problem, other expected discoveries have not arrived. Supersymmetric particles have not been found at the LHC. Grand unified theories predicting proton decay have not been confirmed by Super-Kamiokande or its successors. Weakly interacting massive particles have not been detected by any direct-detection experiment despite increasingly sensitive searches culminating in LUX-ZEPLIN's 417-day analysis in December 2025. The Hubble tension between early-universe and late-universe measurements of the cosmic expansion rate has grown rather than shrunk. The dark energy equation of state shows evolution that Λ CDM does not predict. Each of these is a specific empirical failure of a specific theoretical expectation, and the cumulative effect is a foundation that has been settling for decades without yet finding its level.

This book argues that the failures share a common diagnostic. The unification problem, the cosmological constant, the dark sector, the failed predictions of supersymmetry and grand unification — these are not separate puzzles awaiting separate solutions. They are symptoms of a single structural error: the assumption that quantum mechanics is fundamental.

The alternative — that quantum mechanics is not fundamental but emergent — has been entertained in foundations of physics for nearly a century. It has rarely been productive. The difficulty is that taking quantum mechanics as emergent has typically required either abandoning Bell-inequality results (su-

perdeterminism), introducing nonlocal mechanisms (Bohmian pilot waves), or proliferating universes (Everett’s many-worlds). Each approach has its defenders and its difficulties; none has produced a quantitative empirical record that distinguishes it from the standard reading. The intellectual cost of demoting quantum mechanics from fundamental has typically exceeded the benefits.

The framework developed in this book proposes a different route. It demotes quantum mechanics from fundamental to *derived*, but the derivation produces specific quantitative content that the conventional reading does not. The 10^{122} vacuum-energy discrepancy becomes a calculable ratio rather than a fine-tuning failure. The Standard Model gauge group is reproduced as a forced consequence of a cubic-lattice commitment whose dimension is argued by the framework’s self-consistency filters (Chapter 5) and whose cubic structure is empirically anchored by the observed $SU(3)$ — argued-and-anchored rather than derived from nothing — with the hypercharges then forced by anomaly cancellation. The Bekenstein-Hawking factor of $1/4$ becomes a derivation rather than a postulate. The JUNO neutrino mixing angle is matched by a parameter-free retrodiction at 0.07σ — the derivation postdates the November 2025 measurement — whose exact form makes JUNO’s design-lifetime precision a pre-registered forward test. Twenty-two Standard Model observables are organized by a single foundational commitment plus cubic-lattice geometry — seven as parameter-free structural retrodictions, the remainder as layered, chained, or fitted, each classified explicitly in Appendix A. The framework’s epistemic position is therefore not “another interpretation of quantum mechanics” but a different *kind* of object: a derivation chain from one foundational commitment to specific quantitative content across multiple sectors of physics.

The foundational commitment is unusual. The framework takes as its starting point not a Lagrangian, not a Hilbert space, not a set of symmetries, but the bare fact that *observation occurs*. From this single premise, together with three structural conditions on what lies outside the observer’s access, quantum mechanics is derived as a theorem. The chapters that follow develop this derivation and trace its consequences.

0.2 An informal thought experiment

Before stating the framework’s central claim formally, consider a thought experiment that captures the structural idea.

Imagine a closed room. Inside the room is an experimenter with measuring instruments — clocks, voltmeters, particle detectors. Outside the room is the rest of the universe, but the experimenter has no direct access to it. The walls are opaque to all signals except for whatever couplings the laws of physics enforce: gravitational, electromagnetic, thermal exchanges through the walls. The experimenter wants to characterize the physics inside the room.

The experimenter’s situation has three features. First, the room is coupled to the outside. Energy, momentum, and information cross the walls — not at the ex-

perimeter's choosing, but because the laws of physics force the coupling. Heat flows through walls. Gravitational fields penetrate them. The experimenter cannot isolate the room from the universe completely. Call this *coupling*.

Second, the outside changes more slowly than the inside. The experimenter can run experiments on millisecond timescales — pulses of light, voltage spikes, particle scattering events. The universe outside changes on the cosmological timescale of billions of years. Whatever is going on outside is, from the experimenter's perspective, essentially frozen during any single experiment. Call this *slow-bath separation*.

Third, the outside is vastly larger than the inside. The room contains perhaps 10^{30} particles. The observable universe contains 10^{80} . The information capacity of what the experimenter cannot access dwarfs the information capacity of what they can. Call this *large capacity*.

These three features — coupling, slow-bath separation, large capacity — are the framework's three structural conditions, abbreviated C1, C2, C3. They are not assumptions about exotic physics. They are the actual physical conditions of any laboratory observer in our universe, satisfied by enormous margins.

The framework's central theorem says: an experimenter in this situation cannot describe the physics inside the room using classical mechanics, even if classical mechanics is what governs the underlying reality. They must use a formalism mathematically equivalent to quantum mechanics — with wave functions, superpositions, Born-rule probabilities, and Bell-inequality violations. This is not because reality is intrinsically uncertain. It is because the experimenter's situation — coupled to a slow, vast outside they cannot directly access — forces them into a specific kind of compressed bookkeeping. Quantum mechanics is what classical physics looks like when you are inside the system you are trying to describe.

The thought experiment captures the framework's central move. Quantum mechanics is reframed from a statement about reality to a statement about the observer's relationship to reality. The mathematics is unchanged; what it represents is different. The wave function is not the universe's state; it is the experimenter's compressed description of what they cannot see. The Born rule is not a postulate about fundamental indeterminism; it is the probability assignment that any embedded experimenter must use given C1-C3. Bell-inequality violations do not require nonlocality; they emerge from the specific way correlations propagate through the unobserved outside under C1-C3.

The chapters that follow make this thought experiment precise. The room becomes a partition of the universe into visible and hidden sectors. The experimenter becomes a mathematical observer. The three features become the three structural conditions. The “compressed bookkeeping” becomes the framework's characterization theorem — a proof that the embedded observer's reduced description is mathematically equivalent to unitary quantum mechanics, and that this equivalence is not merely consistent with quantum mechanics

but is required. Quantum mechanics, in the framework’s content, is what any sufficiently complete account of observation must produce.

0.3 The framework’s central claim

The claim that quantum mechanics is not fundamental can be made precise. The framework developed in this book proves that any observer embedded in a finite deterministic universe, satisfying three structural conditions on what lies outside their access, must use a mathematical formalism equivalent to unitarily evolving quantum mechanics. The wave function, the Schrödinger equation, the Born rule, the Bell-inequality violations — all are derived structures of the observer’s compressed description, not features of the underlying reality.

The three conditions are not exotic. They are the actual physical conditions of laboratory observers in our universe, satisfied by enormous margins. The first condition (C1) requires non-trivial coupling between what the observer can see and what they cannot — satisfied because the cosmological horizon is not a static wall, and stress-energy conservation forces correlations across it. The second condition (C2) requires that the hidden sector evolve much more slowly than the visible sector — satisfied with a ratio of approximately 10^{-32} between laboratory timescales and the Hubble time. The third condition (C3) requires the hidden sector to have far more degrees of freedom than the visible sector — satisfied with 10^{122} degrees of freedom on the cosmological horizon against at most 10^{50} in any laboratory apparatus. The framework proves these conditions are not merely sufficient for quantum mechanics to emerge; they are necessary. Any system equivalent to unitarily evolving quantum mechanics arises from marginalization of a deterministic system satisfying these three conditions for some partition.

Quantum mechanics is, in this picture, the unique self-consistent algorithm for making predictions from incomplete data on accessible timescales under C1–C4. It is what computation looks like when one is a component of the machine doing the computing. The reframing has structural consequences. The 10^{122} vacuum-energy discrepancy is not a fine-tuning failure but the compression ratio between the universe’s total degrees of freedom and the visible ones — quantum field theory and general relativity are measuring two different things, not two different values of the same thing. The roughly 95% of cosmic energy density attributed to dark matter and dark energy is the gravitational signature of the same hidden degrees of freedom, structured (dark matter) and uniform (dark energy). These are theorems within the framework, developed in Part II.

The framework reaches further than dissolving the cosmological constant problem. From the same starting definition — *observation occurs* — combined with a small set of structural assumptions about the substratum (finiteness, determinism, bounded coupling, center independence, linearity, background independence), the framework derives the Standard Model gauge group $SU(3) \times SU(2) \times U(1)$ as the unique commutant of a cubic-symmetric coupling matrix;

three chiral fermion generations; one Higgs doublet; and anomaly-free hypercharges. The classical thermodynamics at the cosmological horizon yields the Bekenstein-Hawking entropy with the correct $1/4$ coefficient, the value of $\hbar$, and the MOND acceleration scale $a_0 = cH/6$. The same construction reaches into chemistry, biology, and quantum engineering, where the framework’s structural conditions identify the regimes in which complex matter, life, and technology become possible. Part III develops this cascade. A word on the word “derives,” used here and throughout: it carries a specific and qualified meaning, set out in full at §4.7. The framework forces these structures *given* its substratum commitments; two of those commitments — that space is three-dimensional and that the lattice is cubic — are themselves selected by observation rather than derived, and a few others are definitional. The honest description, developed there, is that the framework compresses an unusually large body of physics onto one finite construction, rather than spinning physics out of the bare fact that observation occurs. The reader is asked to hold the chapters that follow to that standard.

The framework is not a theory of everything in the conventional sense. It does not unify the four forces into a single gauge group, and it does not predict every contingent feature of the observed universe. The CKM phase, the specific masses of the three quark generations, the cosmological initial conditions — these are inputs to the specific representative of the framework’s equivalence class that describes our universe, not derivations from the framework. The framework’s claim is about the structural form of physics, not about its initial conditions. Within that structural form, twenty-two Standard Model observables, the gravitational sector’s thermodynamics, and the cosmological dark sector’s qualitative composition are forced consequences. Part I makes the claim precise; Parts II–IV develop its consequences and tests its predictions.

0.4 The intellectual context

The framework’s central claim — that any observer embedded in a deterministic universe must use quantum mechanics — belongs to a family of impossibility-with-structure results in twentieth-century mathematics and computer science.

In 1931 Kurt Gödel proved that any formal system rich enough to encode arithmetic cannot prove all true statements about itself from within. The limitation arises from self-reference: the system can construct sentences that refer to its own provability, and consistency forces some such sentences to be true but unprovable. The proof yields not chaos but a precisely characterized structure for the unprovable truths.

In 1936 Alan Turing proved that there is no algorithm that can decide, for arbitrary computer programs, whether they halt. The proof again proceeds by self-reference: a hypothetical universal-halting-decider can be turned against itself. The undecidable problems again have a precise structure, captured in the field of computability theory that grew from Turing’s work.

In a pair of results (2001, 2008), David Wolpert proved a physical analog of these theorems. Any inference device embedded in the universe it is trying to predict faces fundamental limits. Not from noise, not from finite computational resources, but because complete self-prediction is logically impossible. Wolpert’s argument is Gödel’s diagonal applied to physical systems.

The framework developed in this book identifies a physical incompleteness that runs parallel to the Wolpert limitation — parallel rather than identical, because Wolpert’s is a logical result about worst-case self-prediction, while the incompleteness here is enforced by spacetime geometry and is quantitative. The observer is made of the same fields as the universe they describe; they obey the same speed-of-light constraint; they occupy a bounded region inside the system they are characterizing. The geometry of spacetime enforces a horizon — a boundary beyond which no signal can reach. Everything beyond that horizon must be averaged over in the observer’s description of physics. The framework’s central theorem is that this averaging, under the actual physical conditions of laboratory observers, produces exactly the mathematical formalism of quantum mechanics. Gödel showed formal systems cannot completely describe themselves; Turing showed computational systems cannot completely predict themselves; Wolpert showed physical inference devices face fundamental limits on self-prediction; this framework adds a physical member to that family: the partition-induced trace-out as the mechanism of an embedded observer’s *physical* incompleteness, and quantum mechanics as its specific cost — a construction that stands on its own derivation, which the logical results corroborate but do not supply.

The framework is not the first to suggest that observation cannot be treated as external to physics. QBism, relational quantum mechanics, ’t Hooft’s cellular-automaton interpretation, and other foundational programs take observer-dependence as a starting point or derive it from specific models. Recent technical work — Chandrasekaran-Longo-Penington-Witten on observer-incorporated entropy in de Sitter space (2022), Maldacena on the role of the observer’s clock in cosmological partition functions (2024), Harlow-Usatyuk-Zhao on observer-dependent Hilbert spaces in closed universes (2025) — has converged on the foundational claim that observation matters from independent directions. The framework here converges with these programs on substance: observation matters; the algebra of observables depends on the observer; the partition between observer and observed is physically meaningful. It differs in approach. CLPW, Maldacena, and HUZ add observer degrees of freedom to an existing quantum-gravitational formalism; the framework here goes the opposite direction, starting with a finite deterministic substratum and deriving quantum mechanics as a theorem.

This intellectual lineage has older roots. The fifteenth-century cardinal Nicholas of Cusa’s *docta ignorantia* (learned ignorance) — the claim that the highest knowledge is knowing what we cannot know — anticipates the framework’s structural claim that incompleteness is a feature rather than a deficiency. Cusa’s

coincidentia oppositorum (coincidence of opposites) anticipates the framework’s claim that two irreconcilable descriptions can both be correct because they are projections of the same underlying reality. Spinoza’s attribute theory — one substance expressed through multiple incommensurable attributes — provides the deepest ontological parallel to the framework’s substratum-and-projections picture. A contemporary touchstone is Deutsch’s *Beginning of Infinity* [Deutsch], which argues that knowledge itself is structurally incomplete — there is always more explanatory work to be done, and the right epistemic stance toward any current best explanation is provisional but committed. The framework’s observational incompleteness provides one physics-foundational grounding for Deutsch’s epistemological claim: knowledge derived from observation inherits observational limits as its lower bound (developed in Chapter 12 §12.8). The framework’s full philosophical positioning is developed in Chapter 3.

0.5 The empirical record

The framework’s claims are unusual; the evidence for them is specific. The empirical record consists of three components: retrodictions of known parameters — including the framework’s sharpest matches, some against measurements that predate the framework’s development — and pre-registered predictions awaiting upcoming data.

The retrodictions are quantitative and parameter-free. The Cabibbo angle of CKM mixing is predicted to be $1/(\pi\sqrt{2}) = 0.2251$ from a single Brillouin-zone distance on the substratum lattice; the observed value is 0.22500 ± 0.00067 , a 0.04% match. The Koide relation among the three charged-lepton masses is predicted to be exactly $2/9$ from a cubic-group quadratic Casimir; the observed value is 0.666661 ± 0.000007 , a 0.02% match. The three Standard Model gauge couplings at the Z pole are reproduced to better than 0.1% from one universal lattice-determined coupling at the Planck scale plus standard renormalization-group running. Six fermion masses are reproduced from one empirical input to better than 1%. The Higgs quartic coupling derived from cubic-group representation theory matches the LHC value at 0.6σ . Twenty-two Standard Model observables in total — gauge couplings, mixing angles, mass ratios — are reproduced as structural retrodictions, with no free parameters fit to data.

One retrodiction deserves separate note for its precision, with its provenance stated plainly: the measurement predates the framework’s 2026 development, so the match is retrodictive, not a forward prediction. In November 2025 the JUNO collaboration reported the first world-leading direct measurement of the solar neutrino mixing angle: $\sin^2 \theta_{12} = 0.3092 \pm 0.0087$. The framework’s prediction — derived from the substratum’s cubic-group structure with no fit parameters — is $1/3 - 1/(4\pi^2) = 0.3080$. The direct measurement matches at 0.14σ ; against the post-JUNO global fit by Capozzi et al., the prediction matches at 0.07σ . The probability of randomly matching a specific value to this precision is below one in a thousand.

In September 2025 the LIGO collaboration’s analysis of GW250114, a high-signal-to-noise binary black-hole merger, confirmed Hawking’s area theorem — the statement that the total horizon area does not decrease — at 99.999% confidence. The observation compares horizon areas and does not measure entropy, so it does not test the $1/4$ coefficient, which remains without a direct empirical test. The framework derives the $1/4$ factor as a structural consequence of mode counting on the cosmological horizon partition; previously the factor had been a postulate.

A second component of the framework’s gravitational predictions concerns the dark sector. The framework predicts that the entropy of the cosmological horizon, distributed over the Hubble volume, produces a gravitational contribution equal to the critical density of the universe — the observed $\Omega_{\text{total}} \approx 1$. It further predicts that the structured component of this entropy (dark matter) follows the MOND phenomenology with acceleration scale $a_0 = cH/6 = 1.2 \times 10^{-10} \text{ m/s}^2$ and the baryonic Tully-Fisher relation $v^4 = GM_B \cdot cH/6$ — both parameter-free. Recent observations of declining rotation curves in the Milky Way at galactocentric distances $> 19 \text{ kpc}$ (Jiao et al. 2023) and in high-redshift galaxies (Genzel et al. 2017) confirm the framework’s specific predictions for how the dark sector phenomenology should evolve. Direct dark-matter searches continue to return null results, consistent with the framework’s structural prediction that no particle dark matter exists.

Several predictions await imminent data. DESI Year 5 will measure the dark energy equation of state to roughly 1% precision on the running-vacuum parameter ν ; the framework predicts $\nu \approx 10^{-32}$, observationally indistinguishable from Λ CDM. The Bullet Cluster’s post-collision relaxation is predicted to show gradual realignment of the dark gravity toward the gas distribution over timescales comparable to H^{-1} . Reliable ALMA gas-mass measurements at $z > 1$ will provide the definitive test of the baryonic Tully-Fisher evolution. The framework’s chromatin-condensate predictions in Rett syndrome and triple-negative breast cancer (Chapter 16) provide near-term clinical-trial-scale tests outside fundamental physics.

The framework’s structural foundation has also received independent corroboration from the open-quantum-systems literature. Brandner’s 2025 theorems on integrating out hidden degrees of freedom in autonomous linear systems establish at theorem level the kind of non-Markovian visible dynamics that the framework requires, derived independently for general open systems with no stake in the framework being correct. The mathematical apparatus the framework relies on is therefore the standard way physicists now think about open systems with hidden structure.

No competing framework produces this collection of results from a single starting definition. The 10^{122} vacuum-energy discrepancy as an information compression ratio, the $\sim 95\%$ dark sector as a gravitational occlusion fraction, the Standard Model retrodiction package (seven parameter-free entries among twenty-two classified observables), the Bekenstein-Hawking $1/4$ factor as a derivation, the

MOND scale $a_0 = cH/6$ as a parameter-free prediction, and the JUNO $\sin^2 \theta_{12}$ retrodiction at 0.07σ — all follow from the single statement that observation occurs and three structural conditions on the partition.

0.6 The book’s structure

The book has four parts.

Part I — Foundations. Chapters 1–4 develop the framework’s mathematical core. Chapter 1 proves the central theorem: under the three structural conditions C1–C3, any embedded observer’s reduced description is mathematically equivalent to unitarily evolving quantum mechanics. The proof is constructive (the Barandes correspondence and the Stinespring dilation provide explicit constructions) and includes the characterization theorem establishing necessity. Chapter 2 constructs the substratum — the finite deterministic bijection (S, φ) — and proves the reconstruction theorem: observed physics together with structural assumptions uniquely determines the substratum up to a precise gauge equivalence. Chapter 3 develops the framework’s hierarchical structural realism, including its evasion of Bell, Kochen-Specker, and PBR-type no-go theorems through identification of which descriptive level each assumption fails at. Chapter 4 articulates the framework’s methodology — the single-foundational-commitment derivation pattern, the architecture-level engagement with the no-go results, the classification of claims by derivational rigor.

Part II — Physics. Chapters 5–9 develop the framework’s quantitative predictions in fundamental physics. Chapter 5 derives the Standard Model gauge group from the cubic-lattice substratum: the wave equation, coupling-degree minimization giving $K = 2d = 6$ internal components per site, the cubic-group decomposition with multiplicities $(3, 2, 1)$, the gauge group $SU(3) \times SU(2) \times U(1)$ as the pointwise commutant, chirality from the trace-out, and anomaly cancellation forcing the observed hypercharges. Chapter 6 develops the matter content and quantitative predictions: three chiral fermion generations from staggered taste structure, one composite Higgs, $\bar{\theta} = 0$ from T-invariance of the substratum, the gauge-coupling derivation with the structural no-GUT prediction, and the twenty-two parameter-free retrodictions including the Cabibbo angle, the Koide relation, the six fermion masses from one empirical input, and all three PMNS angles. Chapter 7 develops the gravitational sector: the Bekenstein-Hawking $1/4$ factor as theorem, the value of $\hbar$ from horizon thermodynamics, the dissolution of the cosmological constant problem. Chapter 8 develops the JUNO neutrino prediction in detail. Chapter 9 discusses universality, the relationship to other unification programs, and the framework’s empirical scope.

Part III — Emergence. Chapters 10–12 trace the structural cascade from the substratum to chemistry, biology, evolution, and intelligence. Chapter 10 derives chemistry from substratum-level electromagnetic structure. Chapter 11 develops the origin-of-life problem in terms of the framework’s structural conditions for autocatalysis and template replication. Chapter 12 traces evolution to

neural computation to semiconductors to general-purpose computation, identifying the contingent and structural steps along the cascade.

Part IV — Applications. Chapters 13–19 develop the framework’s reach into working scientific fields. Chapter 13 addresses quantum biology: non-Markovian enzyme kinetics in proteins, photosynthesis as energy transfer through C1-C3 architectures, magnetoreception via cryptochrome radical-pair dynamics, and enzyme catalysis with conformational landscapes acting as τ_B -scale memory. Chapter 14 develops quantum computing: BQP is identified, under Chapter 14’s explicit model hypotheses, as the computational class accessible to embedded observers satisfying C1-C3, with the lower bound (the emergent quantum mechanics provides the BQP toolkit) and upper bound (the quantum description is informationally complete for the embedded observer) jointly characterizing the class. The Extended Church-Turing Thesis is reframed from empirical conjecture to theorem of the framework, with any super-BQP demonstration providing falsification. Chapter 15 develops quantum engineering: superconducting qubit dynamics as the canonical engineered C1-C3 system with $\tau_S/\tau_B \sim 10^{-3}$, P-indivisible noise structure producing correlated errors over $\sim 10^3$ gates, correlation-aware decoders correlated errors that a correlation-aware decoder can exploit, coherence revivals from engineered baths, and information recovery in quantum sensing — structural effects whose magnitudes the framework does not quantify, and the BEC analogue-gravity experiment as the framework’s cleanest near-term test of the BQP characterization theorem at the foundational level. Chapter 16 develops medical applications across cancer pharmacology, neurodegeneration, antibiotic resistance, immunotherapy, cardiac pharmacology, and autoimmune disease, with an explicit second half on epigenetics as biological hidden sector — the substratum-emergent operator dissociation at chromatin with reader/writer/eraser pharmacology axis, epigenetic drugs as memory erasers, cellular reprogramming as memory reset, the memory hierarchy from DNA methylation through chromatin architecture, and the wider therapeutic window prediction for reader-writer disorders confirmed empirically for Rett syndrome. Chapter 17 develops the bioinformatics consequences of the framework’s structural constraints on biological information processing, with empirical confirmations from single-cell methodology and from the evolutionary-biology literature (molecular clock overdispersion, rate autocorrelation across the tree of life, LTEE power-law fitness trajectory). Chapter 18 develops the framework’s distinctive forward predictions: Born-rule transient violations testable by CMB non-Gaussianity, the quantum-classical boundary with observable signatures, the epistemic character of quantum randomness, the quantum gravity dissolution as cumulative case (combining the CC dissolution, BH information paradox resolution, and non-renormalizability dissolution as three components of one structural resolution), gravitational wave echoes at $\Delta t = (r_+/c) \ln(r_+/2l_p)$, the arrow of time as emergent from partition initial conditions, the framework’s silence on global cosmic spatial topology with engagement of Levin’s compact-universe proposal, and a substantive position on consciousness — the framework’s mathematical operation of observation

(partition with C1-C3 structure) is a necessary but not sufficient condition for consciousness, with the hard problem proper explicitly outside the framework’s scope. Chapter 19 closes the book with a systematic inventory of the framework’s content against the major open problems of fundamental physics — eleven problems resolved or dissolved by the framework (quantum gravity, the cosmological constant, dark matter, dark energy, the measurement problem, black hole information, hierarchy, gauge group origin, generation puzzle, dark sector budget, Born rule origin), and five scope limits where the framework does not provide a derived value (the flavor problem; baryogenesis — solution-specific, the baryon asymmetry inheriting the input status of the framework’s flavor-sector CP phases; inflation and initial conditions; the Hubble tension; the cosmological initial state).

The book’s technical density varies. Chapters 1, 2, 5, 6, and 7 contain the framework’s core mathematical content and are demanding; they require working familiarity with quantum mechanics, general relativity, and group theory. Chapters 3, 4, and 9 are more accessible — philosophical and methodological, with formal content present but not load-bearing. Chapters 10–19 are accessible to readers with general scientific background and assume the technical content of Parts I and II as derived results rather than as content to be re-developed.

Readers approaching the book with different interests will find different reading paths useful. Readers primarily interested in the framework’s foundational claim should focus on Chapters 1, 2, and 7; the JUNO prediction in Chapter 8 provides the most concrete empirical anchor. Readers primarily interested in the philosophical positioning should focus on Chapters 3, 4, and 9, with Chapter 18 providing substantive content on consciousness and cosmic-scale questions. Readers approaching the framework from computational complexity or quantum hardware engineering can begin with Chapters 14 and 15 and consult Parts I and II as needed. Readers approaching the framework from biology, medicine, or applied physics can begin with Chapters 12, 13, 16, and 17 and consult Parts I and II as needed. Readers wanting a systematic overview of the framework’s content against the standard open problems of fundamental physics should consult Chapter 19 as a reference resource. A reader following the full argument from start to finish — the framework’s intended presentation — will find the technical content of Parts I and II load-bearing for the cascade developments of Parts III and IV.

The book makes one foundational commitment — that observation occurs — and develops the consequences. The consequences are extensive. They include the structure of quantum mechanics, the dynamics of general relativity, the gauge group and matter content of the Standard Model, the composition of the cosmological dark sector, the structural conditions for chemistry and life, and the framework’s reach into specific contemporary problems in biology and medicine. Whether the framework is correct is an empirical question that will be settled by the predictions developed throughout the book against measurements now being made and forthcoming. The book’s role is to make the structure of the

framework, the chain of its derivations, and the specific content of its predictions available for evaluation.

Chapter 1

Embedded Observation and the Characterization Theorem

1.1 The problem and the theorem

The Introduction stated the framework’s central claim: any observer embedded in a finite deterministic universe, satisfying three structural conditions on what lies outside their access, must use a mathematical formalism equivalent to unitarily evolving quantum mechanics. This chapter proves the claim.

The starting point is the empirical fact of observation. An observer records distinguishable outcomes of interactions with a system not wholly under their control; that the recording occurs is the one premise the framework takes from outside its own theorems. From this premise, together with three lemmas extracting its formal content, the chapter derives the conditions under which the observer’s description must be quantum-mechanical and proves a characterization theorem stating that these conditions are not merely sufficient for quantum mechanics to emerge but necessary.

Central Theorem. *Let (S, φ, V) be an embedded observation — a finite total system S , a deterministic bijection $\varphi : S \rightarrow S$, and an observer $V \subsetneq S$ with the visible sector and hidden sector coupled through φ . Under three structural conditions on the partition — non-trivial coupling (C1), slow-bath timescale separation (C2), and sufficient hidden-sector capacity (C3) — the embedded observer’s reduced description of the visible sector is mathematically equivalent to unitarily evolving quantum mechanics. The Hilbert space, the Schrödinger equation, the Born rule, and the quantum measurement formalism emerge as structural consequences of the marginalization, requiring no additional postulates. The conditions are necessary as well as sufficient: any system equivalent to unitarily evolving quantum mechanics arises from marginalization of a deterministic bijection satisfying C1-C3 for some partition.*

The proof develops in seven stages, with a structural commentary in §1.10 examining what physics emerges when each of the three conditions fails. §1.2 states the framework’s definition of observation and extracts three structural lemmas from it. §1.3 introduces the three conditions C1-C3, anchoring each in the physical situation of laboratory observers in our universe. §1.4 proves that the partition uniquely determines the emergent description, establishing partition-relativity as a structural feature rather than a free parameter. §1.5

proves that under C1-C3 the marginalized stochastic process is P-indivisible — a precise non-Markovian property that captures the structural difference between embedded-observer dynamics and classical thermal noise. §1.6 establishes the two independent routes from P-indivisibility to unitary quantum mechanics: the Barandes stochastic-quantum correspondence and the Stinespring dilation. §1.7 develops the framework’s account of Bell-inequality violations as a structural consequence of P-indivisible joint dynamics. §1.8 establishes necessity through the characterization theorem. §1.9 closes by stating what the theorem says: quantum mechanics as the unique self-consistent decompression algorithm for predictions from incomplete data under C1-C3. §1.10 develops the three structural limits — what happens when each of C1, C2, C3 fails — as commentary on what kind of observer the framework’s conditions jointly characterize.

1.2 The starting point

The framework begins from a starting point usually put in two words: *observation occurs*. An observer records distinguishable outcomes of interactions with a system not wholly under its own control. This is the *cogito* of Descartes made mathematical — the framework’s one concession to a foundation outside its own theorems.

That starting point, examined closely, is two things. What genuinely cannot be doubted is thin: there is differentiated content — some this-rather-than-that. Call this the first axiom. It is genuinely indubitable, because to doubt it is already to tell one thing from another, which is an instance of it. But notice what a single such fact does *not* contain: time. A single differentiated content is like a photograph — a frozen frame can be full of difference, light here and dark there, without anything in it being before or after anything else. To get from the photograph to a *movie* — to a world with succession, change, one moment after another — something must be added: that the differentiation *recurs*, that it is instantiated more than once. Call this the second axiom. It cannot be squeezed out of the first; a frozen frame, however rich, never amounts to a sequence of frames. The framework therefore rests not on one fact but on two: one that cannot be doubted, and one that must be assumed. This is a strength, not a hedge. A foundation that can point to its single assumed ingredient — and that comes with a theorem showing the assumed ingredient genuinely cannot be derived from the indubitable one — is sturdier than a foundation whose assumptions are scattered, unnamed, through a definition.

One qualification, stated here and developed in the companion paper. Counting the framework’s foundational commitments is partly a matter of how the foundation is *presented* rather than a fact about the world: the same theory can be axiomatized in inter-derivable ways, and which structural feature is named an “axiom” and which is derived as a consequence is, within limits, a presentational choice — a notational variance, in Quine’s sense, that leaves all observable physics untouched. In particular, whether the embedded observer’s status as a *proper part* of a larger whole is best regarded as following from the first axiom

or as a further foundational posit in its own right is an open question — one that bears on how the framework describes the minimality of its own assumptions, but not on any prediction it makes. The technical development of the foundational structure, and the no-go theorem behind it, is given in Part I of the companion paper *Physics Modulo Gauge*; the chapters that follow build on the structure without further reference to it.

Definition. An *observation* is a triple (S, φ, V) : a total system S , a deterministic dynamics $\varphi : S \rightarrow S$, and an observer $V \subsetneq S$ — a proper subsystem with finitely many distinguishable internal states, coupled to the complement $H = S \setminus V$ through φ .

The definition is weaker than classical mechanics. A shuffled deck of cards satisfies it. A finite cellular automaton satisfies it. Three lemmas extract its structural content. The map φ is nothing unusual, and its form need not even be assumed. It follows from one innocuous premise — that evolution *composes*: advancing the state by one interval and then another equals advancing it by the combined interval. That single consistency condition (Cauchy’s functional equation) forces evolution into a one-parameter family generated by a fixed rule (Engel & Nagel, 2000); in continuous time this is the *phase flow* of classical mechanics (Arnold, 1989), and in the discrete case here it is simply the repeated application of one transition map. So the state-transition picture is not a special assumption but a consequence of time evolution being consistent. What *is* distinctive here is not that such a map exists, but that it acts on a *finite* system; finiteness, not the dynamics, is the load-bearing commitment.

Lemma 1 (Finiteness). *The visible configuration space $\mathcal{C}_V$ is finite, with a discreteness scale ϵ providing a finite minimal cell volume.*

Any observer bounded by a finite-area surface couples to only finitely many modes across that boundary; independent support comes from holographic entropy bounds [2]. The scale ϵ is left unspecified here; Chapter 2 will show that self-consistency forces $\epsilon \approx 2\ell_P$.

Lemma 2 (Causal partition). *The phase space decomposes as $\Gamma = \Gamma_V \times \Gamma_H$, with $H_{\text{tot}} = H_V + H_H + H_{\text{int}}$. Cross-partition correlations enter only through H_{int} .*

Γ_V is the visible sector; Γ_H is the hidden sector; H_{int} couples them. The decomposition is not a modeling choice but the very definition of embedded observation: an observer with access to all of S would have nothing external to observe.

Lemma 3 (Determinism and unique measure). *φ is a bijection on S . The counting measure is the unique φ -invariant measure, becoming the Liouville measure on Γ in the continuum limit:*

$$\frac{\partial \rho}{\partial t} = \{H_{\text{tot}}, \rho\} \equiv \mathcal{L}\rho.$$

The observer uses standard Kolmogorov probability. The bijection requirement is structural: non-injective dynamics would merge distinct total states, erasing the observer’s records of past outcomes — contrary to what observation is. On a finite set, injectivity implies surjectivity. This is worth putting in context: the reversibility of φ — equivalently, that the substratum conserves information — is not a peculiar demand of this framework but the standard assumption of mainstream physics. Unitary evolution in quantum mechanics *is* information conservation, and its classical shadow is Liouville’s theorem, the conservation of phase-space volume recovered here. Physics treats this principle as effectively non-negotiable: the black-hole information paradox is called a *paradox* precisely because the apparent loss of information under Hawking evaporation would violate it, and the no-hiding theorem (Braunstein & Pati, 2007) makes the conservation statement rigorous. The framework assumes no more than standard physics already does; what is distinctive, developed below, is not the assumption but what emerges from it.

These lemmas contain no quantum postulates. Quantum mechanics emerges under three further conditions.

1.3 The observer’s blind spot

Light travels at finite speed; the universe has finite age. Every observer therefore has a horizon — a boundary beyond which no signal has had time to arrive. Beyond it lies ordinary physics, but it is structurally inaccessible. The laws of physics written for what the observer can see must average over everything they cannot. The hidden sector must be traced out.

The trace-out produces a stochastic process on $\mathcal{C}_V$. The visible state x is compatible with many hidden states h ; each h sends x to a different future x' . Averaging over h weighted by the Liouville measure produces transition probabilities $T_{ij}(t)$ — the chance that x_i at t_1 becomes x_j at t_2 . The deterministic underlying system has become, from the observer’s perspective, a stochastic process.

In the Markovian regime — fast bath, system slow relative to environment — each transition is independent of the last. This is the standard setting of open quantum systems and produces classical noise. The cosmological hidden sector differs in three specific ways, and the framework’s central theorem is that their conjunction forces emergent quantum mechanics.

Condition C1 (Non-zero coupling). $H_{\text{int}} \neq 0$: the hidden state non-trivially affects the visible evolution. The horizon is not a static wall — stress-energy conservation enforces dynamical correlations across it, and in the physical realization information flows in both directions. The defining content of C1, and what the necessity argument delivers, is the hidden-to-visible influence; bidirectionality belongs to the realization ([Main §3.4], Remark at the C1-equivalence theorem).

Condition C2 (Slow-bath timescale separation). $\tau_S \ll \tau_B$. The hidden sector evolves much more slowly than the visible sector. Its correlation time is set by the Hubble timescale, $\tau_B \sim H^{-1} \sim 10^{17}$ s. A laboratory experiment operates on $\tau_S \sim 10^{-15}$ s or shorter; the ratio is $\tau_S/\tau_B \sim 10^{-32}$. This is the *inverse* of the Markovian regime: the environment cannot reset between measurements. Every correlation it picks up from one experiment is still there when the next begins.

Condition C3 (Sufficient capacity). $N_H \gg N_V$. The cosmological horizon has roughly 10^{122} independent degrees of freedom, the Bekenstein-Hawking entropy. A laboratory apparatus has at most 10^{30} to 10^{50} . No experiment appreciably disturbs the hidden sector. Its memory never saturates.

Condition C4 (History readback). The interaction reads what it writes. Some hidden degrees of freedom that depend on the visible past must also influence later visible steps; at the horizon this is automatic — the coupling is bidirectional, acting on the same boundary degrees in both directions. Without it, influence plus storage is mere noise: a hidden sector that records history on one set of components while a disjoint set drives the visible dynamics produces exactly Markovian statistics for as long as one cares to look.

The conditions are independent in statement but coupled in content. A fast environment with vast capacity gives Markovian noise (failing C2). A slow environment with limited capacity saturates (failing C3). An isolated environment exerts no influence (failing C1). An environment that stores without reading back yields exact Markov noise despite all three (failing C4). Only an environment that is coupled, slow, and vast sustains the persistent non-decomposable correlations that produce emergent quantum mechanics. The conditions are not exotic constraints; they are the physical conditions of laboratory observers in our universe, satisfied by enormous margins.

A clarification on what “slow” means. C2 is not a claim that physics fundamentally proceeds at different rates in different regions of the universe — it does not. The underlying substratum bijection proceeds at one unified rate, with no preferred clock and no “slow region.” Rather, C2 is a claim about *correlation timescales*: how long correlations persist in the hidden sector before relaxing, relative to how quickly the visible sector probes those correlations. The asymmetry between the visible-sector and hidden-sector timescales comes from the partition’s structure, not from any local variation in physical law.

The structural source of the asymmetry, for our cosmological situation specifically, is that the cosmological horizon is one physical structure that sets *both* C2 and C3 simultaneously. The horizon’s area sets the hidden sector’s capacity through the Bekenstein bound: $N_H \sim A/4l_p^2 \sim 10^{122}$. The horizon’s surface gravity sets the correlation relaxation timescale: $\tau_B \sim 1/\kappa \sim 1/H_0 \sim 10^{17}$ s. These two quantities — area and surface gravity — are related through horizon thermodynamics ($\kappa = 2\pi T_{\text{horizon}}$ with the horizon temperature set by the cosmological constant), so in our cosmological setting *one* underlying geometric

parameter (effectively, the value of Λ) controls both C2 and C3 jointly.

For laboratory observers in our universe, C2 and C3 are therefore correlated rather than fully independent. Both are satisfied because the cosmological horizon is enormous and has slow surface dynamics; both would fail simultaneously if the horizon were small and rapidly evolving. The correlation is a feature of our specific cosmological situation, not a structural redundancy in the framework.

For other partitions, however, C2 and C3 can be tuned independently. In biological systems (Chapter 13), the protein conformational landscape has a hidden sector substantially smaller than cosmological ($N_H \sim 10^{10}$ conformational states), with slow dynamics set by folding-landscape physics rather than horizon thermodynamics. Each condition has a different source: C3 from the configurational entropy of the conformational space; C2 from the kinetic barriers separating metastable conformations. In engineered quantum systems (Chapter 15), C2 and C3 can be tuned independently by choosing the substrate material (which controls N_H through the TLS bath density) and the operating temperature (which controls τ_B through bath equilibration rates). The BEC analogue-gravity experimental program (Chapter 15 §15.8) specifically tests the framework’s predictions in the regime where C3 is marginal while C2 remains satisfied — an experimental design that depends on the independence of the two conditions.

The framework’s three conditions are therefore *logically* independent — each can fail without forcing the others to fail — even though they are *physically* correlated for the cosmological partition. The framework’s empirical content (Chapters 13, 15) and forward predictions (Chapter 18) rely on this logical independence to make distinct quantitative claims about regimes where one condition is marginal and others are satisfied.

With the definition, the lemmas, and the conditions in place, the chapter’s central theorem can be proved. The remainder of the chapter develops the proof in the stages identified in §1.1: partition-relativity (§1.4), P-indivisibility (§1.5), the two routes to unitary quantum mechanics (§1.6), Bell-inequality violations (§1.7), the characterization theorem closing the biconditional (§1.8), and the algorithmic reading of what the theorem establishes (§1.9).

1.4 Partition-relativity

The marginalized transition probabilities of the previous section depend on three inputs: the dynamics, the partition, and the measure. The framework’s lemmas fix the dynamics and the measure; only the partition remains as a free input. The following result establishes that this is the entire freedom: the emergent description is determined by the partition alone.

Lemma. *The emergent description is uniquely determined by the partition. Any parameters of the emergent theory depend only on the geometric and thermodynamic properties of the partition boundary.*

Proof. By Lemma 3 the Hamiltonian flow ϕ_t is unique. By Lemma 2 the partition is fixed. An embedded observer with incomplete knowledge of the total energy averages over energies with a smooth prior, recovering the Liouville measure μ on a phase-space band. Letting π_V project onto the visible sector, the marginalized transition probabilities are

$$T_{ij}(t_2, t_1) = \int_{\Gamma_H} \delta_{x_j}[\pi_V(\phi_{t_2-t_1}(x_i, h))] d\mu(h),$$

determined by dynamics, partition, and measure. All three are fixed; only the partition remains as input. $\square$

Different observers correspond to different partitions of the same (S, φ) ; different partitions produce different emergent descriptions. Quantum mechanics' parameters are not free constants. The Planck constant $\hbar$ does not appear in the lemmas or conditions; it must emerge from the marginalization, with value set by the partition's geometric and thermodynamic features. Chapter 7 shows that the cosmological partition produces $\hbar$ as a thermodynamic identity at the horizon surface, with 1/4 Bekenstein-Hawking entropy following as a structural consequence — with GW250114 confirming the classical area theorem at 99.999%; the entropy coefficient itself has no direct test [8].

Partition-relativity establishes that the emergent description is well-defined. The next question is what kind of description it is — whether the marginalization produces a Markovian process, a generic non-Markovian one, or something more specific. The result of this section is that the process satisfies a precise structural property that has a name in the open-systems literature: P-indivisibility. The two sections following establish that P-indivisibility is mathematically equivalent to quantum mechanics.

1.5 P-indivisibility

The central technical result of the chapter is that the marginalized process is P-indivisible whenever the underlying bijection has non-trivial coupling.

Theorem (P-indivisibility). *Let $\mathcal{C}_V$ and $\mathcal{C}_H$ be finite sets with $|\mathcal{C}_V| \geq 2$, and let $\varphi : \mathcal{C}_V \times \mathcal{C}_H \rightarrow \mathcal{C}_V \times \mathcal{C}_H$ be a bijection with*

$$T_{ij} = \frac{|\{h \in \mathcal{C}_H : \pi_V(\varphi(x_i, h)) = x_j\}|}{|\mathcal{C}_H|}.$$

If T is not a permutation matrix, then the process is P-indivisible.

P-indivisibility is a precise property of stochastic processes that requires unpacking. A process is *P-divisible* if, for any three times $t_1 < t_2 < t_3$, the transition matrix factors through t_2 as

$$T(t_3, t_1) = \Lambda(t_3, t_2) \cdot T(t_2, t_1),$$

where $\Lambda(t_3, t_2)$ is a valid stochastic matrix — non-negative entries, rows summing to one. P-divisibility expresses the natural intuition that probabilistic dynamics should compose: the probability of going from x_i to x_k over the long interval $[t_1, t_3]$ should be obtainable by chaining the probability of going from x_i to some intermediate state x_j over $[t_1, t_2]$ with the probability of continuing from x_j to x_k over $[t_2, t_3]$. Classical Markov processes satisfy this property trivially because the future depends only on the present state, not on history. P-indivisibility means the factorization fails: the candidate intermediate propagator $\Lambda(t_3, t_2)$ that would make the equation hold has negative entries, and no valid stochastic matrix exists.

Breuer, Laine, and Piilo [9] showed that P-indivisibility is equivalent to a more directly interpretable property: *information backflow*. In a P-divisible process, the distinguishability of any two initial conditions can only decrease over time — the future loses information about the past monotonically, as in any classical Markov chain reaching equilibrium. In a P-indivisible process, distinguishability can rebound: information about the initial condition that appeared lost can return to the visible sector at a later time. The system is not memoryless. The “memory” lives in the hidden sector, where correlations laid down by the visible sector’s history are preserved and later read back.

The proof of the theorem uses the total variation distance between probability distributions,

$$d(p, q) = \frac{1}{2} \sum_k |p_k - q_k|.$$

This quantity ranges from 0 (identical distributions) to 1 (distributions with disjoint support) and measures how distinguishable two probability distributions are by any measurement. For a P-divisible process, d is non-increasing along trajectories: starting from δ_i (the point mass at x_i) and δ_j (the point mass at x_j), the quantity $d(\delta_i T^{(t)}, \delta_j T^{(t)})$ — where $T^{(t)}$ denotes t steps of the dynamics — cannot grow with t . P-indivisibility is therefore established by demonstrating non-monotonic behavior of d along some trajectory.

Proof of P-indivisibility. The proof exhibits non-monotonicity through a three-step construction.

Step 1 (Recurrence to maximal distinguishability). The dynamics φ is a bijection on the finite total set $\mathcal{C}_V \times \mathcal{C}_H$. Every bijection on a finite set is a permutation, and every permutation generates a finite cyclic group: there exists an integer N (the least common multiple of the cycle lengths in φ ’s cycle decomposition) such that $\varphi^N = \text{id}$. Applied to the marginalized dynamics, this means $T^{(N)}$ equals the identity matrix I . Starting from δ_i and applying $T^{(N)}$ returns δ_i unchanged; starting from δ_j returns δ_j . The distinguishability is $d(\delta_i T^{(N)}, \delta_j T^{(N)}) = d(\delta_i, \delta_j) = 1$ for any $i \neq j$. After N steps, the two initial conditions are again perfectly distinguishable.

Step 2 (Strict contraction at intermediate times). If T is not a permutation matrix, then there must exist visible states x_i, x_j , and x_l with both $T_{il} > 0$ and

$T_{jl} > 0$. (If no such triple existed, each row of T would have a single nonzero entry equal to one, making T a permutation.) The existence of x_l as a common destination from both x_i and x_j means the post-step distributions $\delta_i T$ and $\delta_j T$ both assign positive weight to x_l — they overlap. For overlapping distributions p, q with $p_l \wedge q_l > 0$, the total variation distance satisfies

$$d(p, q) = \frac{1}{2} \sum_k |p_k - q_k| < 1,$$

the inequality strict because of the overlap. So $d(\delta_i T, \delta_j T) < 1$.

Step 3 (Non-monotonicity). Combining the previous two steps: at $t = 0$, $d(\delta_i, \delta_j) = 1$. At $t = 1$, $d(\delta_i T, \delta_j T) < 1$. At $t = N$, $d(\delta_i T^{(N)}, \delta_j T^{(N)}) = 1$. The trajectory of distinguishability is therefore $1 \rightarrow (\text{something less than } 1) \rightarrow 1$ — non-monotonic. A P-divisible process cannot produce this trajectory because d would be non-increasing throughout. The process is P-indivisible. $\square$

The proof uses only Lemma 1 (finiteness, hence Step 1’s recurrence) and condition C1 (non-trivial coupling, hence Step 2’s overlap). It does not use C2, C3, or any further structural input. The argument extends from discrete to continuous time through Hamiltonian flow on compact energy surfaces: Poincaré recurrence — the same statement as the discrete recurrence in a measure-theoretic language — plays the role of $\varphi^N = \text{id}$, and the strict contraction at intermediate times follows by continuity from the discrete case.

The accessible-timescale lemma. A potential objection to the proof as stated is that the recurrence time N can be astronomically large. For a hidden sector with $m = |\mathcal{C}_H|$ states, N can be as large as $m!$, and for the cosmological partition with $m \sim e^{10^{122}}$, the recurrence time is $t_R \sim e^{10^{122}}$ years — orders of magnitude longer than the age of the universe to any power. The distinguishability restoration at $t = N$ is therefore not directly observable. The framework requires P-indivisibility on *accessible* timescales — times short compared to τ_B , the hidden sector’s intrinsic memory time — not at recurrence. The accessible-timescale lemma supplies this.

The mechanism is straightforward information storage and retrieval through the coupling. At each visible-sector interaction (timescale τ_S), the coupling H_{int} writes some quantity I_0 of information from the visible sector into the hidden sector. Between interactions, the stored information decays as the hidden sector’s internal dynamics scrambles it. The decay rate is set by the hidden sector’s spectral gap Δ — the smallest energy gap between distinct eigenstates of H_H . The autocorrelation of any stored quantity falls as $e^{-\Delta t}$ in the dephasing approximation, so the per-step retention from one visible-sector interaction to the next is

$$r = e^{-\Delta \tau_S} \approx 1 - \Delta \tau_S = 1 - \tau_S / \tau_B,$$

where $\tau_B = 1/\Delta$ is the hidden sector’s memory time. When $\tau_S \ll \tau_B$, this is close to unity: almost no information is lost per visible-sector cycle. After k

steps the cumulative retention is $r^k \approx 1 - k\tau_S/\tau_B$, which remains close to unity for any observation window short compared to τ_B/τ_S cycles.

Quantitatively, the non-Markovian mutual information between the visible sector's past and future, conditioned on its present, is bounded below by

$$I(X_{<t}; X_{>t} | X_t) \geq I_0 \left(1 - \frac{k\tau_S}{\tau_B} \right),$$

which is order I_0 for any $k\tau_S \ll \tau_B$. The information backflow that the proof established at recurrence times is therefore present continuously throughout the accessible-timescale regime: the visible sector's distinguishability has rebounded by an amount of order I_0 at every step until enough cycles have accumulated to deplete the stored memory.

For the cosmological partition the numbers are striking. The hidden sector's memory time is set by the Hubble scale, $\tau_B \sim H^{-1} \sim 10^{17}$ s. A laboratory experiment operates on $\tau_S \sim 10^{-15}$ s. After $k = 10^{20}$ steps — far more than any conceivable single experiment — the cumulative decay is $k\tau_S/\tau_B \sim 10^{-12}$, negligible. The hidden sector retains essentially all of the information written into it over any observation window short of τ_B . Condition C3 then guarantees that this storage does not saturate: the maximum information the hidden sector can hold is $\log_2 m \approx 405$ bits at the cosmological scale, enormously more than any laboratory experiment can write.

The combination of C1, C2, and C3 therefore produces accessible-timescale P-indivisibility with substantial information backflow at every step. The recurrence-based proof of the theorem establishes the existence of P-indivisibility; the accessible-timescale lemma establishes that the same property holds in the parameter regime that matters for actual observers in our universe.

A worked minimal model. A concrete example makes the mechanism tangible. Take a two-state visible sector $\mathcal{C}_V = \{0, 1\}$ — call it a coin with states Heads and Tails — and a six-state hidden sector $\mathcal{C}_H = \{1, 2, 3, 4, 5, 6\}$ — call it a die. The dynamics is the permutation

$$\sigma = (0,1 \leftrightarrow 1,1) (0,2 \leftrightarrow 1,2) (0,3 \leftrightarrow 0,4) (0,5 \leftrightarrow 0,6) (1,3 \leftrightarrow 1,4) (1,5 \leftrightarrow 1,6).$$

Each pair $(\cdot \leftrightarrow \cdot)$ is a transposition that swaps two total states; all six swaps together are independent (act on distinct states), so the full dynamics is a product of disjoint transpositions and satisfies $\sigma^2 = \text{id}$.

Verify the three conditions: C1 holds because die values $h = 1, 2$ couple to coin flips while $h = 3, 4, 5, 6$ do not — non-trivial coupling. C2 holds trivially with $\tau_S/\tau_B = 1/2$, since $\sigma^2 = \text{id}$ gives recurrence at $N = 2$. C3 holds with six hidden states against two visible states.

Compute the one-step transition matrix starting from $x = 0$ (Heads) with the die uniformly distributed. The dynamics sends $(0, 1) \rightarrow (1, 1)$ and $(0, 2) \rightarrow (1, 2)$ —

Heads becomes Tails for these die values. It sends $(0, 3) \rightarrow (0, 4)$, $(0, 4) \rightarrow (0, 3)$, $(0, 5) \rightarrow (0, 6)$, $(0, 6) \rightarrow (0, 5)$ — Heads stays Heads for these die values, though the die value changes within the hidden sector. So $P(0 \rightarrow 0) = 4/6 = 2/3$ and $P(0 \rightarrow 1) = 2/6 = 1/3$. Symmetric reasoning for $x = 1$ gives

$$T(1, 0) = \begin{pmatrix} 2/3 & 1/3 \\ 1/3 & 2/3 \end{pmatrix}.$$

The distinguishability after one step is $d(\delta_0 T, \delta_1 T) = \frac{1}{2}(|2/3 - 1/3| + |1/3 - 2/3|) = 1/3$. The two initial conditions, perfectly distinguishable at $t = 0$ with $d = 1$, have become much less distinguishable: d has dropped to $1/3$.

A Markovian process with this one-step matrix would have $T(2, 0) = T(1, 0)^2$. Computing,

$$T(1, 0)^2 = \begin{pmatrix} 5/9 & 4/9 \\ 4/9 & 5/9 \end{pmatrix},$$

with distinguishability $1/9$ — further decreased, as required for a Markov process. But the actual second-step matrix is $T(2, 0) = I$ because $\sigma^2 = \text{id}$ sends every total state back to itself, so the visible-sector marginals are unchanged from the starting point. The distinguishability at $t = 2$ has returned to 1. The two initial conditions are perfectly distinguishable again. Complete un-mixing — impossible for any Markov process.

The signature of P-indivisibility shows up as negative entries in the would-be intermediate propagator. Solving $T(2, 0) = \Lambda(2, 1) \cdot T(1, 0)$ for the candidate $\Lambda(2, 1)$ gives

$$\Lambda(2, 1) = T(2, 0) \cdot T(1, 0)^{-1} = I \cdot \begin{pmatrix} 2 & -1 \\ -1 & 2 \end{pmatrix} = \begin{pmatrix} 2 & -1 \\ -1 & 2 \end{pmatrix}.$$

The off-diagonal entries are -1 . No valid stochastic matrix has negative entries — the factorization through $t = 1$ fails by direct computation. The process is P-indivisible.

The mechanism is clear in this example. The die functions as a memory register. At step 1, if the coin was at 0 and the die was at 1, the dynamics produces the new state $(1, 1)$: the coin has flipped and the die has not changed value. The die value 1 now encodes a record of what happened — “the coin was at 0 before this step.” At step 2, σ acts on $(1, 1)$ and returns it to $(0, 1)$: the die’s stored memory has been used to undo the previous coin flip. The visible sector’s history, lost in the one-step distinguishability decay, has been recovered from the hidden sector.

C1 is what enabled writing the record — the coin-die coupling at die values 1 and 2. C2 is what ensured the record was preserved between the write at step 1 and the read at step 2 — here trivially, because the die itself does not evolve between coin flips. C3 is what provided enough storage capacity — six hidden states against two visible. The combination of the three conditions produces the information backflow that, as the next section will establish, is mathematically quantum.

1.6 The two routes to quantum mechanics

P-indivisible stochastic processes are not merely *like* quantum mechanics. They are mathematically identical to it. Two independent constructions establish this identity, each starting from different premises and arriving at the same conclusion.

The Barandes correspondence [4, 5]. Any P-indivisible stochastic process on a finite configuration space of size n can be represented by a unitarily evolving quantum system. There exist a Hilbert space $\mathcal{H}_V$ of dimension at most n^3 , a Hermitian operator $\hat{H}$ acting on $\mathcal{H}_V$, and a one-parameter family of unitary operators $U(t) = \exp(-i\hat{H}t)$ such that the classical transition probabilities equal Born-rule probabilities computed from the unitary:

$$T_{ij}(t) = |U_{ij}(t)|^2.$$

The left side is the classical transition probability from averaging over hidden states. The right side is the squared modulus of a matrix element of the unitary evolution. The equality is a mathematical identity, not an approximation or analogy.

The construction works because P-indivisibility means transition probabilities cannot factor through intermediate times in the classical sense. The candidate intermediate propagator $\Lambda(t_2, t_1)$ that would make $T(t_3, t_1) = \Lambda(t_3, t_2) \cdot T(t_2, t_1)$ hold for P-indivisible processes has negative entries — there is no valid stochastic matrix that does the factorization. In quantum mechanics, the same factorization failure is the standard situation: amplitudes for different paths combine with phase factors that produce interference, and the resulting probabilities are not marginalizable across intermediate times. The connection between these two pictures is that the “negative probabilities” that prevent classical factorization in a P-indivisible process are the *same mathematical structure* as the amplitude phases that produce quantum interference, expressed in different notation. Where classical stochastic dynamics writes $P(\text{path}) = P(\text{leg } 1) \cdot P(\text{leg } 2)$ and finds that negative numbers are required, quantum mechanics writes $A(\text{path}) = A(\text{leg } 1) \cdot A(\text{leg } 2)$ with complex amplitudes whose phases produce the same constraint on observable probabilities through $P = |A|^2$.

The Barandes correspondence makes this informal identification rigorous. Given a P-indivisible $T(t)$, the construction produces an explicit $\hat{H}$ such that $|U_{ij}(t)|^2 = T_{ij}(t)$ exactly. The Schrödinger equation then arises immediately: any smooth one-parameter family of unitaries can be written $U(t) = \exp(-i\hat{H}t/\hbar)$ for some Hermitian generator $\hat{H}$, with the action scale $\hbar$ entering as the conversion factor between the dimensionless unitary and a dimensionful Hamiltonian. The value of $\hbar$ cannot be determined from transition data alone — different choices of action scale produce the same $T(t)$ — but is fixed by the partition geometry, which Chapter 7 develops from horizon thermodynamics.

(The identity is written here in its simplest form, in which the emergent Hilbert

space has the same dimension as the configuration space. In the general construction the visible transition matrix is the ancilla-marginal of a Born-rule matrix on a larger dilated space — $T_{ij}(t) = \frac{1}{m_a} \sum_{k,l} |\langle j, l | U(t) | i, k \rangle|^2$; the same-dimension form is the special case with a trivial ancilla, and the technical statement of the characterization theorem is given in that general form in the companion paper.) A residual ambiguity might appear: the relation throws away phase information in $U(t)$, and different unitaries could in principle produce the same transition probabilities. The *phase-locking lemma* establishes that this ambiguity has no observable content. Continuous-time transition data is sufficient to determine the Hamiltonian $\hat{H}$ uniquely up to a basis-relabeling gauge that has no physical content. The argument is Fourier extraction: substituting the spectral decomposition $U(t) = V \cdot \text{diag}(e^{-iE_k t}) \cdot V^\dagger$ gives

$$T_{ij}(t) = \left| \sum_k V_{ik} e^{-iE_k t} V_{jk}^* \right|^2 = \sum_{k,l} V_{ik} V_{jk}^* V_{jl} V_{il}^* e^{-i(E_k - E_l)t},$$

which expands as a sum of oscillating terms at energy-gap frequencies $\omega_{kl} = E_k - E_l$. Under a non-degenerate-gap condition — the energy gaps are all distinct, generically satisfied — the Fourier coefficient at frequency ω_{kl} extracts the product $V_{ik} V_{jk}^* V_{jl} V_{il}^*$. The collection of all such Fourier coefficients determines the unitary V and the energy levels E_k up to a global phase and a basis-relabeling redundancy, neither of which has observable content. The Hamiltonian is uniquely fixed by the observable transition data.

The Stinespring route [6, 7]. The same conclusion follows from a different starting point through standard textbook quantum information theory. The deterministic bijection $\varphi : \mathcal{C}_V \times \mathcal{C}_H \rightarrow \mathcal{C}_V \times \mathcal{C}_H$ defines a permutation matrix U_φ on the tensor-product Hilbert space $\mathcal{H}_V \otimes \mathcal{H}_H$, where $\mathcal{H}_V$ and $\mathcal{H}_H$ have orthonormal bases labeled by $\mathcal{C}_V$ and $\mathcal{C}_H$ respectively. A permutation matrix is unitary, so U_φ is a valid unitary operator on the tensor-product space. The observer's ignorance of the hidden sector — which hidden state h is currently realized — corresponds to a maximally mixed hidden-sector state $\rho_H = I_m/m$, where $m = |\mathcal{C}_H|$ and I_m is the m -dimensional identity. The choice of ρ_H as maximally mixed reflects the framework's underlying probability assumption: the observer assigns equal prior probability to each hidden state, since Lemma 3 makes the counting measure the unique φ -invariant measure.

The visible-sector channel is then the partial trace

$$\Phi(\rho_V) = \text{Tr}_H [U_\varphi(\rho_V \otimes \rho_H) U_\varphi^\dagger].$$

The channel Φ takes a visible-sector density matrix ρ_V , tensors it with the maximally mixed hidden-sector state, evolves the joint system with the unitary U_φ , and traces out the hidden sector to extract the visible-sector marginal. By a standard textbook result, Φ is completely positive and trace-preserving: it preserves normalization of the visible-sector state, it preserves positivity of ρ_V at every step, and it preserves positivity even when extended to act on a tensor

factor of a larger Hilbert space (complete positivity, the stronger condition that distinguishes physically realizable channels from merely positive maps). The Born-rule transition probabilities computed from Φ — applying Φ to the initial $\rho_V = |x_i\rangle\langle x_i|$ and reading off the diagonal entry $\langle x_j|\Phi(\rho_V)|x_j\rangle$ — are exactly the classical marginalization probabilities T_{ij} computed from averaging over the hidden sector.

The Stinespring construction delivers several pieces of structure that the Barandes route does not provide explicitly. The tensor-product Hilbert space $\mathcal{H}_V \otimes \mathcal{H}_H$ is constructed rather than postulated — the structure that quantum mechanics treats as fundamental emerges from the underlying configuration space. The channel Φ can also carry genuine *quantum coherence*, and the condition under which it does is geometric. Two notions have to be kept apart here, because they are easily run together. Complete dephasing — destroying every off-diagonal density-matrix element — is the extreme decoherence channel. Entanglement breaking is a strictly weaker property: an entanglement-breaking channel is one realizable as a measurement followed by a fresh preparation, and its output can still carry coherence — coherence manufactured by the preparation rather than transmitted from the input; what such a channel cannot do is establish entanglement across the cut. And neither property decides the memory structure of the visible process: P-indivisibility is a multi-time property of the trajectory statistics, and the hidden sector can carry the memory even when a single step’s channel is entanglement-breaking — a two-state system with a two-state hidden register already realizes an entanglement-breaking step atop a P-indivisible process. What the entanglement-breaking question does decide is whether a single update transmits quantum coherence through the cut — whether, at that step, the dilation is doing quantum work or merely dressing a classical channel — so it matters where Φ falls. Coupling alone does not settle this: T can fail to be a permutation while Φ is still entanglement-breaking at one step, as happens when the coupling across the cut resolves every visible degree of freedom of the region — a single site with all of its neighbours hidden being the extreme case. What controls the answer is the ratio of bulk to boundary. Because a single update transports information one lattice site, the hidden sector resolves the visible region only across the cut, so the visible directions it cannot distinguish outnumber those it can whenever the region has an interior. For a visible region more than a few sites across — six suffices in three dimensions — the off-diagonal elements that encode interference survive that step, and they survive no matter how large the hidden sector is. Decoherence is an effect of area; coherence is a property of the bulk. These are one-step statements about the specific lattice update, proved as such in Appendix B.2; their many-step composition is an open problem, not a corollary.

The CP-indivisibility of the channel family $\{\Phi_t\}$ follows from the P-indivisibility of the transition matrix family $\{T_t\}$ established in §1.5. This is the channel-level statement of the same non-Markovian property: just as $T(t_2, t_1)$ cannot be factored through intermediate times as a valid stochastic matrix product, the channel $\Phi_{t_3-t_1}$ cannot be factored as $\Phi_{t_3-t_2} \circ \Phi_{t_2-t_1}$ for any valid quantum

channel — the candidate intermediate channel would not be completely positive. The reduction from channel divisibility to population divisibility holds because the framework’s channels are permutation dilations with uniform ancilla, which map computational-diagonal states to diagonal states; for a general CPTP family it fails — a Hadamard pair is CP-divisible while its computational-basis populations are P-indivisible (Appendix B, Theorem B.2.5 and its scope note).

The approximate unitarity of the visible-sector dynamics on observable timescales follows from the slow-bath condition C2 together with a weak-coupling condition that it is important to state explicitly. At $t = 0$ the channel generator is exactly Hamiltonian; at finite time it acquires two distinct dissipative contributions:

$$\dot{\rho}_V = -\frac{i}{\hbar}[\hat{H}_{\text{eff}}, \rho_V] + \mathcal{D}(\rho_V),$$

the standard Lindblad form with effective Hamiltonian $\hat{H}_{\text{eff}}$ and dissipator $\mathcal{D}$. The leading term is the Schrödinger evolution. The dissipator has two pieces with different origins. The first is present even in the strictly frozen-bath limit: the channel is then a mixture, over the bath ensemble, of bath-conditioned visible unitaries, and the spread of those conditioned Hamiltonians dephases the visible state by an amount of order the coupling strength squared — independent of τ_B . The second, from the bath’s motion away from the frozen limit, is suppressed by $(\tau_S/\tau_B)^2$. The dynamics is therefore approximately unitary Schrödinger evolution to the extent that both the slow-bath condition $\tau_S \ll \tau_B$ and the weak-coupling condition (small conditioned-Hamiltonian spread) hold; the latter is the parameter that controls the leading non-unitarity. At the cosmological partition both are extreme, so $\mathcal{D}$ is far below detection sensitivity for any laboratory measurement; the dissipative corrections, if measurable at high enough precision, would be predictions of the framework. The continuous-measurement and quantum-trajectory formulations of open quantum systems agree with this framework structure on observable timescales.

The macroscopic-superposition frontier. The exactness claim of this section has a direct experimental frontier: matter-wave interferometry at increasing mass, which tests whether superposition and unitary evolution persist where intrinsic-collapse alternatives predict breakdown. The current record is the Vienna cluster interferometer of Pedalino et al. (2026): sodium clusters of 5,000–10,000 atoms — roughly 8 nm across, exceeding 1.7×10^5 atomic mass units, the mass range of large proteins and small viruses — delocalized over interferometer paths separated by 133 nm, dozens of times each particle’s diameter, for about 10^{-2} s, with the resulting interference in agreement with quantum theory. On the macroscopicity measure of Nimmrichter and Hornberger (2013) the experiment reaches $\mu = 15.5$, roughly an order of magnitude beyond the previous benchmark (Fein et al. 2019); an equally rigorous test with electrons would require maintaining electron coherence for on the order of 10^8 years.

The framework’s reading of this record has three parts. First, it is a non-discriminating consistency in the sense of Chapter 7 §7.10: the framework

predicts exact quantum behavior for embedded observers at every accessible mass scale — deviations enter only at the dissipator scale identified above, far below current sensitivity — so the result cannot separate the framework from textbook quantum mechanics. The one-sided pressure falls on intrinsic-collapse programs (GRW/CSL, Diósi–Penrose), which predict breakdown at some scale and which the framework’s deterministic substratum forbids outright (Chapter 18 records the structural incompatibility of Penrose–Hameroff objective reduction). Second, it deepens the framework’s most exposed empirical anchor: the determinism assumption A2 rests on exactly one empirical input — the observed unitarity of quantum dynamics (Chapter 2; Substratum paper §2) — and mass-scaled interferometry is the frontier where that input is directly stress-tested. Third, it is a standing Class A falsification target (§4.5): a confirmed intrinsic-collapse signature at any scale would falsify the emergence theorem itself, with no recoverable error mode, because the framework’s own predicted corrections are structurally pinned below the collapse-model band. The announced scaling program — larger masses and biological particles, with several further orders of magnitude planned — keeps this exposure live rather than historical.

The tensor-product structure for spatially separated subsystems within the visible sector — two laboratories, two detectors, two regions of space — follows separately from the locality of the underlying coupling graph rather than from the Stinespring construction itself. If two visible-sector subsystems V_A and V_B have disjoint coupling neighborhoods on the coupling graph, the spatial Markov property holds: $V_A^\circ \perp V_B^\circ \mid \partial(V_A \cup V_B)$, where the conditional independence is with respect to the joint distribution generated by the substratum dynamics. The emergent Hilbert space inherits the factorization $\mathcal{H}_{V_A} \otimes \mathcal{H}_{V_B}$ from the locality of the coupling graph, with entanglement between V_A and V_B arising from shared boundary history rather than from instantaneous joint properties.

The two routes are not redundant. The Barandes correspondence applies to any P-indivisible stochastic process, whether or not its underlying dilation is specified; it is the more general result. The Stinespring route requires the dilation to be specified explicitly but delivers richer structural information (the tensor product, the channel-level CP-indivisibility, the Lindblad-form generator, the suppression of dissipation by τ_S/τ_B). Both routes agree on the observable transition data. The agreement means the emergence of quantum mechanics from embedded observation under C1-C3 does not rest on any single technical result: if one were to discover an error in Barandes’ construction, the Stinespring route from textbook quantum information would establish the same conclusion through different mathematics. The two-route structure is the framework’s hedge against single-point failure in the foundational technical machinery. The division of labor is precise: the correspondence is load-bearing only for the necessity direction of the characterization theorem (§1.8); every constructive claim, including the $\hbar$ derivation of Chapter 7, survives on the Stinespring route alone.

An independent route to the same conclusion. Both routes of this section are constructive: they exhibit quantum mechanics as the description of

the marginalized process. The operational reconstruction theorems of Hardy [arXiv:quant-ph/0101012 (2001)], of Masanes and Müller [New J. Phys. **13**, 063001 (2011)], and of Chiribella, D’Ariano and Perinotti [Phys. Rev. A **84**, 012311 (2011)] approach quantum mechanics from the opposite direction: they prove that among finite-dimensional generalized probabilistic theories, quantum theory is the unique one satisfying a short list of operational requirements — tomographic locality, continuous reversibility, and, in the informational axiomatization, purification — stated without reference to any substratum, determinism, or hidden sector. Those requirements, minus the continuity clause, admit exactly two solutions: classical and quantum probability theory. The P-indivisibility theorem of §1.5 — the framework’s own contribution — is precisely the certificate that the marginalized process occupies the non-classical branch. The remaining load-bearing requirements each have an independent origin in the construction rather than being read off the reconstructed output: finite dimensionality from the finite configuration space; purification exhibited rather than postulated, since every visible mixed state arises as the marginal of the deterministic dilation; no-signaling from the marginalization structure; continuous reversibility from the phase-locked unitary family, exact in the idealized limit and holding at finite ϵ up to the residual dissipation, entering at second order in the timescale ratio (of order 10^{-64}), which is a prediction of the framework rather than a defect; and tomographic locality from the derived composite factorization together with the complex Hilbert space the correspondence delivers, so that derivation — including its open step, the unification with the staggered spin structure of chapter 6 — carries load for this requirement as well. Requirements not on this list, such as perfect distinguishability and ideal compression, hold in the emergent theory only as properties of the output and carry no independent weight here. Given the independently sourced features, the reconstruction theorems force quantum theory uniquely — so the representation freedom in the correspondence cannot move the emergent description outside quantum mechanics; any operational theory with the marginalized process’s features is quantum mechanics up to the standard gauge. The same division holds for the Born rule: alongside the constructive derivation, Gleason’s theorem [J. Math. Mech. **6**, 885 (1957)] shows that every noncontextual probability measure on the projection lattice of a Hilbert space of dimension three or more has the Born form, and the emergent probabilities are noncontextual by construction, descending from a single measure over substratum states; the envariance and decision-theoretic derivations of Zurek, Deutsch, and Wallace are convergent constructive members of the same family. The necessity statements add no new mechanism; their force is that the operational profile the construction produces has been independently proved to admit exactly one theory.

With unitary quantum mechanics established as the necessary description of the embedded observer’s marginalized dynamics, the chapter’s next technical question is whether the framework reproduces the specific non-classical correlations that distinguish quantum mechanics from classical hidden-variable theories. The decisive test is Bell’s theorem, which proves that no local hidden-

variable theory satisfying a specific factorizability condition can reproduce the correlations of quantum mechanics, and the experimental confirmation since Aspect’s 1982 measurements that quantum correlations actually violate Bell’s inequalities at their predicted quantum bound. Whether and how the framework — itself a hidden-variable theory in the literal sense, with definite states at the substratum level — passes this test is the subject of §1.7.

1.7 Bell-inequality violations

The framework is a hidden-variable theory: there is a deterministic substratum with definite states, and the observer’s stochastic description averages over the parts they cannot see. Bell’s theorem [10] proves that no local hidden-variable theory satisfying factorizability can reproduce quantum correlations. The framework evades Bell not by superdeterminism — no fine-tuning of λ to experimenter choices — but by violating factorizability through P-indivisible joint dynamics.

Bell’s factorizability requires $P(a, b \mid x, y, \lambda) = P(a \mid x, \lambda) \cdot P(b \mid y, \lambda)$, presupposing that λ carries all relevant information at a single moment. P-indivisible processes are not like this. The joint transition matrix for two subsystems that interacted during preparation does not factor: $T_{QR} \neq T_Q \otimes T_R$. This non-factorizability *is* stochastic entanglement.

The Jarrett decomposition splits factorizability into parameter independence (outcome at A independent of setting at B) and outcome independence (outcome at A independent of outcome at B). The framework preserves parameter independence and violates outcome independence. By Fine’s theorem (Fine 1982), this is exactly the class consistent with quantum correlations. This decomposition is contested ground: the strongest critique in the philosophy-of-physics literature holds that keeping parameter independence while conceding outcome independence mistakes a no-signaling condition for locality. Chapter 3 §3.7.1 states that critique at full strength and gives the framework’s complete response, including what it concedes.

Parameter independence follows from spatial locality: if detectors A and B have disjoint coupling neighborhoods, Alice’s setting determines which function she evaluates on V_A but does not alter V_B ’s transition probabilities. Hence $P(b \mid a, x, y) = P(b \mid y)$ structurally — no fine-tuning of the hidden variable distribution required.

The maximum violation the framework allows is exactly Tsirelson’s bound $2\sqrt{2}$, the quantum maximum. Barandes-Hasan-Kagan [12] prove this directly from the stochastic-quantum correspondence; it also follows from the Stinespring route via Tsirelson’s operator-norm argument applied to the emergent unitary description. The framework predicts continued quantum-bound violations, no signature of local hidden-variable theories or super-quantum correlations. The Wood-Spekkens fine-tuning objection (2015) assumes DAG causal structure with Markov factorization; P-indivisible processes violate the Markov condition on

any DAG, so the appropriate formalism is the process tensor, within which no-signaling follows from marginalization structure without fine-tuning.

The P-indivisibility theorem and the two routes establish sufficiency: C1-C3 imply quantum mechanics. The remaining question is necessity. Does every system equivalent to unitarily evolving quantum mechanics arise from a deterministic bijection satisfying C1-C3 for some partition? §1.8 establishes that it does — the characterization theorem closes the biconditional.

1.8 The characterization theorem

The Barandes correspondence already establishes that quantum mechanics is equivalent to P-indivisibility [4, 5]; §1.5 established that C1-C3 imply P-indivisibility. Necessity requires the converse: that P-indivisibility implies C1-C3 for any underlying dilation. The argument is three separate proofs, one for each condition.

Necessity of C1. If T is a permutation matrix, then T^k is a permutation for all k , and the candidate intermediate propagator $\Lambda(k_2, k_1) = T^{k_2 - k_1}$ is a valid stochastic matrix for any $k_2 > k_1$. The process is P-divisible. Contrapositively, P-indivisibility requires T not to be a permutation — equivalent to non-trivial coupling.

Necessity of C2. The C2 condition states that the hidden sector evolves much more slowly than the visible sector: $\tau_S \ll \tau_B$. Showing that this condition is necessary for P-indivisibility requires demonstrating that the *opposite* regime — fast bath, $\tau_B \ll \tau_S$ — produces P-divisible dynamics. The argument proceeds through the eigenstate thermalization hypothesis [13, 14] applied to the hidden sector’s dynamics.

The substratum dynamics is reversible — φ is a bijection by Lemma 3, and the corresponding continuous-time generator H_H on the hidden sector has purely imaginary spectrum on the finite-dimensional Hilbert space. There is no relaxation in the strict thermodynamic sense: every observable returns to its starting value at the Poincaré recurrence time t_R , which is exponentially large in the hidden-sector dimension and far beyond any experimental access. What occurs on accessible timescales is *dephasing*, a structurally different process.

Consider a hidden-sector state $|\psi_H\rangle$ written in the energy basis as $|\psi_H\rangle = \sum_n c_n |E_n\rangle$. Under unitary evolution, the state acquires phases: $|\psi_H(t)\rangle = \sum_n c_n e^{-iE_n t/\hbar} |E_n\rangle$. The expectation value of a hidden-sector observable O_H is

$$\langle O_H \rangle(t) = \sum_{m,n} c_m^* c_n e^{i(E_m - E_n)t/\hbar} \langle E_m | O_H | E_n \rangle.$$

Two distinct contributions appear. The diagonal terms ($m = n$) have no time-dependence — they give the long-time average $\sum_n |c_n|^2 \langle E_n | O_H | E_n \rangle$, the “diagonal ensemble” average. The off-diagonal terms oscillate at frequencies $(E_m - E_n)/\hbar$. As time advances, these oscillating terms accumulate phases at

different rates depending on the energy gaps, and they dephase relative to one another — destructive interference among the off-diagonal contributions. After a timescale $\tau_B \sim \hbar/\Delta E$, where ΔE is the typical hidden-sector energy spread, the off-diagonal contributions average to nearly zero. The expectation value $\langle O_H \rangle(t)$ converges to the diagonal-ensemble value plus oscillating corrections of size $1/\sqrt{D}$ where D is the effective dimension of the support of the state.

The eigenstate thermalization hypothesis adds a further claim about the diagonal-ensemble value. For a generic chaotic many-body Hamiltonian, the matrix elements $\langle E_n | O_H | E_n \rangle$ are smooth functions of E_n , equal to the microcanonical expectation value $\langle O_H \rangle_{\text{mc}}$ at energy E_n up to corrections exponentially small in system size. The diagonal-ensemble value therefore equals the microcanonical value to high precision for any state with narrow energy support. ETH is a well-supported conjecture for generic chaotic systems, proved for specific model classes and numerically verified across many systems, though not a theorem in the full generality required here.

Apply this picture to the hidden sector under the fast-bath regime $\tau_B \ll \tau_S$. Between successive visible-sector interactions, the hidden sector evolves under H_H for time $\tau_S \gg \tau_B$. The off-diagonal contributions in the hidden-sector state dephase, and under ETH the diagonal contributions agree with the microcanonical ensemble. So whatever correlations the previous visible-sector interaction wrote into the hidden sector have been scrambled by the time the next interaction occurs. The hidden sector's state, from the next interaction's perspective, is indistinguishable from the equilibrium ensemble. Formally, for any few-body visible-sector observable O_V that couples to the hidden sector and any time $\tau_S \gg \tau_B$,

$$|\langle O_V \rangle_{\mu_H(\tau_S|x_i)} - \langle O_V \rangle_{\text{eq}}| \lesssim \tau_B/\tau_S,$$

where $\mu_H(\tau_S|x_i)$ is the hidden-sector distribution conditional on the previous visible-sector state x_i , evolved for time τ_S , and $\langle O_V \rangle_{\text{eq}}$ is the equilibrium expectation. The bound on the right side is parametrically small in the fast-bath regime: the hidden sector has thermalized completely from the visible sector's perspective.

The consequence for the transition matrix is direct. Each visible-sector transition probability T_{ij} is computed by evolving the hidden sector forward and projecting onto the visible-sector measurement basis. If the hidden sector has thermalized between transitions, every T_{ij} is computed against the same effective equilibrium distribution. The transition matrix T is therefore the same at every step, and successive applications give $T^{(k)} = T^k$ — the dynamics is a homogeneous Markov chain. Homogeneous Markov chains satisfy the factorization $T^{(k)} = T^{k-j} \cdot T^{(j)}$ with T^{k-j} a valid stochastic matrix; this is P-divisibility by definition.

The contrapositive is the C2 necessity statement: observable P-indivisibility on accessible timescales requires the fast-bath regime to fail, i.e., requires $\tau_S \ll \tau_B$ — condition C2.

The necessity of C2 is conditional on a quantitative mixing hypothesis — uniform total-variation closeness of the hidden sector’s post-relaxation conditionals to a single bath distribution — under which the divisibility deficit is bounded by an explicit constant times the mixing parameter ([Main §3.4]); ETH, conjectural rather than fully proved, is the physical case for that hypothesis. For the cosmological hidden sector — a generic chaotic many-body system at scales where the standard ergodic-hypothesis arguments apply — ETH is the standard assumption and is supported by extensive numerical evidence in analogous systems. The C2 necessity argument is conditional on the mixing hypothesis, with ETH as its motivation; the C1 and C3 necessity arguments are not.

Necessity of C3. Let $m = |\mathcal{C}_H|$. The non-Markovian mutual information is bounded by hidden-sector capacity:

$$I(X_{<t}; X_{>t} | X_t) \leq \log_2 m.$$

Proof. The total system is deterministic: $X_{>t}$ is a function of (X_t, H_t) . Given X_t , the chain $X_{<t} \rightarrow H_t \rightarrow X_{>t}$ is Markov. The data-processing inequality gives

$$I(X_{<t}; X_{>t} | X_t) \leq I(X_{<t}; H_t | X_t) \leq H(H_t | X_t) \leq \log_2 m. \quad \square$$

Any realization of a process whose conditional past–future information reaches I^* requires $m \geq 2^{I^*}$ — the hidden sector must hold the backflow it carries (P-indivisibility alone does not force $m \geq n$: two hidden states suffice for full revival across three visible states); sustained backflow at rate I_0 over K events requires $m \gtrsim 2^{KI_0}$ under an additional no-reuse hypothesis, of which C3 is a conservative restatement.

The three necessity arguments combine into the chapter’s central result.

Characterization theorem. *Let $\mathcal{S}$ be a stochastic process on a finite configuration space $\mathcal{C}_V$ with $|\mathcal{C}_V| \geq 2$, and consider dynamics on accessible timescales $t \ll t_R$; take $\mathcal{S}$ time-homogeneous, the class every iterated-bijection marginal inhabits. Assume the hidden sector of any candidate realization possesses a conditional mixing curve with finite mixing time $\tau_B(\varepsilon_0)$ for every $\varepsilon_0 > 0$, ETH being its physical motivation; no smallness of the residual at the coupling time is assumed. Then (ii) $\iff$ (iii), and (ii) $\implies$ (i):*

(i) $\mathcal{S}$ is mathematically equivalent to unitarily evolving quantum mechanics — there exists a Hilbert space $\mathcal{H}_V$, a Hermitian operator $\hat{H}$, and a unitary family $U(t) = e^{-i\hat{H}t}$ such that $T_{ij}(t) = |U_{ij}(t)|^2$.

(ii) $\mathcal{S}$ is P-indivisible on accessible timescales.

(iii) $\mathcal{S}$ arises from marginalizing a deterministic bijection on $\mathcal{C}_V \times \mathcal{C}_H$ satisfying (C1), (C2), and (C3) for some partition. $\square$

Quantum mechanics is not merely *compatible with* embedded observation under C1-C3. It is equivalent to it.

The characterization establishes Hilbert space, unitary dynamics, and Born-rule transition probabilities. The further structures of operational quantum mechanics follow without independent postulates. The visible-hidden tensor product comes from the Stinespring construction. The tensor product for spatially separated visible-sector subsystems follows from the spatial Markov property of range-1 coupling on the underlying graph. Projective measurement corresponds to Bayesian conditioning on the substratum followed by re-marginalization — the Lüders rule $\rho \rightarrow P_k \rho P_k / \text{Tr}(P_k \rho)$ is the quantum transcription of this classical update. Multi-time correlation functions, weak values, and Leggett-Garg inequalities are standard calculations within the delivered formalism, requiring no additional postulates. One qualification attaches to the multi-time statement: the correspondence is engineered to reproduce two-time transition data, which under-determines multi-time joint statistics, whereas the substratum fixes a unique multi-time law. Their agreement is supported by the same chaotic-mixing argument that secures operational measurement independence, and any residual pre-division memory would appear as deviations from quantum multi-time predictions — a falsifiable texture of the same standing as the residual dissipation.

1.9 What the theorem says

An observer inside the universe receives incomplete data — the visible sector only. The characterization theorem proves that there is exactly one self-consistent algorithm for making predictions from this incomplete data on accessible timescales under C1-C3. That algorithm is quantum mechanics.

Everything in quantum mechanics is a feature of this algorithm, not a feature of the underlying reality. The wave function is the algorithm’s internal state — the bookkeeping device that tracks what the observer knows and what they do not. Complex amplitudes are the algorithm’s arithmetic — the specific number system the reconstruction requires. Interference is what happens when the algorithm combines two incomplete pathways and their bookkeeping entries partially cancel. Entanglement is the algorithm’s encoding of correlations written into the hidden sector during preparation and not yet read back. The Born rule is the algorithm’s output format — the conversion of internal state into predictions the observer can check against measurement.

None of these are properties of the deterministic substratum. The substratum contains no wave functions, no complex numbers, no interference, no entanglement. There are configurations evolving under a Hamiltonian flow. The apparatus of quantum mechanics — every textbook, every equation, every measurement outcome — is what the decompression algorithm looks like from the inside of an embedded observation.

This reframing has consequences for the framework’s empirical content. Phenomena that appear strange under a quantum-fundamental reading become structural under the algorithmic reading. Negative-energy solutions of the Dirac

equation, the existence of antimatter, the specific Hilbert-space structure of relativistic quantum field theory — all are what the decompression algorithm produces when the observer’s incomplete data comes from a relativistic system. Chapters 5 and 6 develop these consequences for the Standard Model.

1.10 The three structural limits

The characterization theorem establishes C1, C2, and C3 as jointly necessary for emergent quantum mechanics. Each condition’s failure was treated in §1.8 as a contrapositive — violating the condition destroys the framework’s emergent quantum mechanics. This section develops the failure cases in the opposite direction: not as proofs of necessity but as *structural commentary* on what kind of physics emerges when each condition fails completely. The three limit regimes reveal that the framework’s conditions are not arbitrary technical assumptions; they jointly characterize the kind of observer we actually are — coupled to a slow, vast environment — and each failure mode describes a hypothetical observer we are not.

The C1 limit: the isolated observer.

C1 requires non-trivial coupling between visible and hidden sectors. In the limit of complete failure — no coupling at all — the visible sector evolves as a *closed* deterministic system. The bijection φ restricted to $\mathcal{C}_V$ is a permutation; the dynamics on V is just the corresponding classical evolution. No information ever crosses the partition boundary. The hidden sector exists but is causally disconnected from the visible sector for all purposes.

The dynamics in this regime is fully classical. The reduced description requires no Hilbert space, no Born rule, no wave function. The Schrödinger equation does not emerge because there is no coupling-induced non-Markovianity for it to encode. Bell-inequality violations cannot occur because P-indivisibility cannot arise from a partition with no coupling. The observer sees a classical Newtonian world with deterministic trajectories and well-defined values of every observable at every time.

This is the regime closest to the historical picture of classical mechanics: a clockwork universe where every particle has a definite position and momentum and the equations of motion fully determine the future from the present. The framework’s content is that this picture is correct for the limit $C1 \rightarrow 0$ — but no laboratory observer occupies this limit. Gravitational coupling between visible and hidden sectors is unavoidable in our universe; stress-energy conservation forces correlations across any horizon; the cosmological horizon is not an impermeable wall. Real observers are coupled, and the framework’s emergent quantum mechanics is the dynamics that coupling produces.

The C1 limit therefore illuminates a structural point: classical mechanics is what *uncoupled* systems do. The apparent quantumness of the world is a signature not of intrinsic randomness in the underlying reality but of our *non-isolation*.

The substratum is deterministic; what makes physics look probabilistic is that the observer is coupled to something they cannot directly see. An observer who could isolate themselves completely from everything outside their measurement apparatus would see classical mechanics. No such observer exists in our universe, and the framework’s content is that no such observer can exist — gravitational coupling is universal.

This has a striking corollary: the project of “explaining quantum mechanics” by finding the right deterministic substructure is well-defined within the framework, but the dynamics that observer’s would see in that substructure is not quantum mechanics. It is classical mechanics. Quantum mechanics is the dynamics that *embedded* observers see when they are coupled to the substructure; the substructure itself, viewed from a hypothetical uncoupled vantage point, would be classical. The framework’s reframing is that the persistent intuition “there must be a deterministic story underneath” is correct, and the deterministic story is classical, but it is structurally inaccessible to any actual observer.

The C2 limit: the thermalized observer.

C2 requires the hidden sector to evolve much more slowly than the visible sector: $\tau_S/\tau_B \ll 1$. In the limit of complete failure — the hidden sector relaxes instantly, $\tau_B \rightarrow 0$ — the hidden sector equilibrates between every visible-sector measurement. No correlation written into the hidden sector during one measurement persists to influence the next. The information backflow that produces P-indivisibility cannot occur because there is no memory in the bath for the visible sector to read back.

The dynamics in this regime is *fully Markovian*. The reduced description on the visible sector is a classical stochastic process with transition probabilities depending only on the current state. The observer sees probabilistic dynamics, but the probabilities are classical: no interference, no Bell-inequality violations, no genuine quantum coherence. The Born rule reduces to the classical probability assignment for transitions between well-defined states. The Schrödinger equation does not emerge because there are no temporal correlations for it to encode.

This regime corresponds to what physics calls the *thermal limit* — systems in contact with a fast equilibrating bath, where the bath’s state at any moment is independent of the system’s history. The classical statistical mechanics of thermal systems is exactly this limit. Brownian motion, chemical reaction kinetics in well-stirred solutions, ordinary diffusion processes — these are all Markovian dynamics described by classical stochastic processes, and the framework’s content is that this is what C2 failure looks like.

No laboratory observer occupies the thermal limit completely. Our cosmological horizon evolves on a timescale $\tau_B \sim H^{-1} \sim 10^{18}$ seconds, while laboratory timescales are $\tau_S \sim 10^{-15}$ to 10^{-1} seconds. The ratio $\tau_S/\tau_B \sim 10^{-32}$ is extreme — the cosmological horizon is essentially frozen on any laboratory timescale, and this frozen-bath condition is what enables quantum mechanics. An observer

in genuine thermal equilibrium with a fast bath would lose access to quantum dynamics; an observer with a frozen bath (our situation) has access to it through the information backflow C2 enables.

The C2 limit therefore illuminates a different structural point: classical stochastic mechanics is what *thermalized* systems do. The framework gives a structural account of why decoherence and the classical limit emerge in systems coupled to fast environments — a familiar phenomenon in quantum mechanics. Decoherence theory describes how quantum coherences are destroyed by coupling to a fast bath; the framework’s reading is that decoherence is the transition from the partially-quantum regime back toward the C2 failure limit. Systems with engineered fast baths approach classical stochastic behavior; systems isolated from fast baths retain quantum coherence. The conventional thermal-classical-limit story is the framework’s C2 limit told from inside quantum mechanics.

A philosophical corollary: the framework predicts that quantum mechanics requires a *frozen* hidden sector, not just a hidden sector. The slow-bath condition is doing structural work beyond merely providing some unobserved degrees of freedom. An observer with access to a fast hidden bath would see classical probability; only observers with access to a *slow* hidden bath see quantum amplitudes. The cosmological horizon’s slowness is therefore not a contingent fact but a structurally required feature of any observer who sees quantum mechanics.

The C3 limit: the universe-sized observer.

C3 requires large hidden-sector capacity: $N_H \gg N_V$. In the limit of complete failure — $N_H \rightarrow 0$, the hidden sector vanishes entirely — the visible sector V approaches the total system S . There is no trace-out because there is nothing to trace over. The observer is *coextensive* with the universe.

The dynamics in this limit is fully deterministic and fully classical. The bijection φ on the full system S is just a permutation of classical states; the observer who is all of S sees its complete deterministic evolution. No Born rule, no Schrödinger equation, no quantum interference — but for a different reason than in the C1 case. In the C1 limit, the visible sector is closed but the framework still applies to it as a system; in the C3 limit, the framework’s content has no purchase because there is no embedded observation taking place. The “observer” is not observing anything; they are the system being observed, with no separation between knower and known.

This is the regime that superficially resembles a *God’s-eye view*. The traditional theological framing imagines an observer who stands outside the universe and sees everything at once with no uncertainty. The framework’s C3 limit is similar in eliminating quantum indeterminism, but it differs structurally: there is no “outside” in a finite deterministic universe. Such an observer would be coextensive with the universe, not transcendent of it. They would not be viewing the universe; they would be the universe describing itself.

Three structural features of this regime are worth noting.

Observation breaks down entirely. In the framework’s account, observation is a *partition-relative* concept. The observer reads the visible sector through measurements that depend on the partition’s coupling. Without a partition — without anything in the hidden sector to be coupled to — measurement has no meaning. The observer cannot perform observation in the framework’s sense. The God’s-eye view is therefore not a limit case of observation but a category outside it.

Laplace’s demon is structurally impossible. The classical version of the God’s-eye view is Laplace’s demon: a hypothetical intelligence that knows the position and momentum of every particle and can therefore predict everything. Conventional readings reject Laplace’s demon because of quantum indeterminism; the framework’s reading rejects it for a deeper reason. The demon would require being the universe, not merely knowing about it. The impossibility is structural rather than epistemic — not “the demon cannot get the information” but “the demon’s having the information would mean they are the universe itself.”

Self-reference paradoxes persist. Even in the limit case where C3 fails completely, the Gödel-Turing-OI incompleteness family from Chapter 12 §12.8 applies. A system that is the entire universe and tries to describe itself fully runs into the same self-reference paradoxes that finite systems do. The universe-sized “observer” cannot include a complete description of themselves within their description — adding such a description would change the universe and require a new description. The framework’s content is that complete self-knowledge is structurally impossible at any scale, including the limit case.

The C3 limit therefore illuminates a third structural point: observation is partition-relative, and the God’s-eye view is not a limit of observation but a category outside it. The framework’s anti-Laplacian content is stronger than the conventional one: not “we cannot know everything because of quantum randomness” but “complete knowledge requires being the universe rather than being in it, and any embedded observer is structurally limited.”

The three limits together.

The three failure modes are structurally distinct.

Condition	What fails	Resulting dynamics	Resulting situation
C1	Coupling	Closed classical	Isolated subsystem
C2	Timescale separation	Markovian stochastic	Thermalized observer
C3	Capacity hierarchy	No observation	Coextensive with system

Each tells us something different about why our situation produces quantum mechanics. C1 failure reveals that classical mechanics is what uncoupled systems do — we see quantum mechanics because we are coupled. C2 failure reveals that

classical stochastic mechanics is what thermalized systems do — we see quantum mechanics because our bath is frozen, not equilibrated. C3 failure reveals that observation itself is structurally limited — we see quantum mechanics because we are embedded, not coextensive.

Together, the three limits characterize the kind of observer we are by contrast with three kinds of observer we are not. We are not isolated (C1 holds), not thermalized (C2 holds), not coextensive with the universe (C3 holds). The framework’s three structural conditions are exactly the joint specification of “embedded laboratory observer coupled to a slow, vast environment.” Each is empirically robust at its own end — the gravitational coupling enforcing C1, the cosmological timescale enforcing C2, the cosmological horizon enforcing C3 — and the framework’s emergent quantum mechanics is the dynamics this jointly specified observer must see.

This is the framework’s structural content at its sharpest: quantum mechanics is not a separate physical theory we discovered alongside classical mechanics. It is the dynamics that our specific kind of observation produces, and the alternative dynamics (closed classical, Markovian stochastic, deterministic self-coextensive) correspond to alternative kinds of observation that are not realized in our universe. The framework’s three conditions are the structural fingerprint of being a coupled, slow-bath, partition-relative observer — the only kind of observer that physical laboratories produce, and the only kind of observer whose dynamics is quantum mechanics.

This reframing has consequences for the framework’s empirical content beyond what is already developed in Chapters 5 through 19. Each of the three conditions is independently testable through the partially-quantum regime — engineered systems approaching the marginal-failure boundary. Chapter 15 §15.7 develops the partially-quantum regime in engineered quantum systems with the BEC analogue-gravity test as the framework’s cleanest experimental probe of C3 marginality. Chapter 13 §13.7 develops biological systems approaching marginal C3. Chapter 18 §18.3 develops the partially-quantum regime as a forward prediction of the framework. The three structural limits are therefore not just conceptual content but a framework of empirically accessible regimes, with current experiments now beginning to probe the marginal-failure boundaries.

1.11 The card table and the chessboard

Chess is a game of complete information: every piece stands in the open, and in principle the game is exhaustively analyzable — a sufficiently large lookup table settles every position. Poker is a game of incomplete information, and that difference is not cosmetic. When von Neumann founded game theory he took poker, not chess, as the interesting case, precisely because hidden cards force probability into rational play even though nothing in the rules is random. The mixed strategy is not ignorance dressed up; it is the correct mathematical form of play from inside an information set. The framework’s claim can be stated in

exactly these terms: the substratum is a chess game, and physics is the rulebook a piece would rationally write from inside it. What this section shows is that ordinary poker is not a rich enough interior view — and that the specific way it falls short, and the specific modification that repairs it, together reproduce the anatomy of Chapter 1 in miniature.

Ordinary poker yields only classical probability. Its hidden information is a static deal, uncorrelated with the visible play once the cards are down, and the resulting odds are divisible: the probability of an outcome at showdown is the probability chained through any intermediate street, whether or not that street was observed. This is the mathematics of ignorance about a definite fact — the position Einstein hoped quantum probability would turn out to occupy. A game that stops here illustrates the *premise* of the framework (observation from inside a partition) while missing its *content* (that structured incompleteness yields quantum, not classical, statistics).

Three modifications close the gap, and they are conditions C1–C3 of §1.2 in card form. First, the shuffle between streets is a deterministic, reversible permutation of the deck — a faro shuffle, say — rather than a randomization: bijective dynamics on the full state. Second, each street deals burn cards, face down and never shown, which the shuffle correlates with the live cards: a hidden sector coupled to the visible one (C1), whose configuration persists between streets (C2), with capacity to store a record of the visible play (C3). Third, revealing a card is an intervention, not a passive glance: a revealed card is burned and replaced under a fixed rule, so revelation physically breaks the card’s correlation with the rest of the deck.

The minimal such game is the coin-and-die model of §1.5 dealt onto a table. The player holds one hole card, red or black — the coin. One burn card, red or black, lies face down — the memory register the die provided. The shuffle between streets is the simplest reversible move available: hole card and burn card trade places. Then, writing the odds that the hole card at showdown matches the hole card as dealt:

protocol	matching odds
one shuffle, then showdown	1/2
two shuffles, then showdown	1 (certain)
two shuffles, odds chained through the unobserved middle	1/2
two shuffles with a peek (and forced re-burn) in the middle	1/2

Each line is checkable by hand — there are sixteen cases in the longest protocol. A single shuffle randomizes the visible card completely. Two shuffles return it with certainty: the first swap writes the hole card into the burn pile, the second reads it back — the same write-then-read information backflow that §1.5 exhibited in the would-be intermediate propagator with negative entries. Chaining the one-street odds through the unobserved middle street predicts

even money and is simply wrong: the process is P-indivisible, and a player who prices hands divisibly loses money at a rate equal to the interference term — the deviation has an operational meaning as concrete as a bankroll. And peeking at the middle street, which forces the re-burn, makes the divisible answer correct again: a division event in the sense of §1.5, enacted with cardboard.

The reading this hands the student is the one the framework intends. Interference is not a substance; it is the failure of divisibility — the visible history taking a round trip through the hidden sector and returning. Phase is not a mysterious internal angle; at this resolution it is *where in the hidden sector the visible information currently resides*. Measurement does not collapse anything; it intervenes, and the intervention severs a correlation that the undisturbed dynamics would have used.

The game extends upward. Replace the two-card swap with a linear shuffle on ranks modulo q and the game runs on the same dynamical class as the substratum of Chapter 2; the reachable odds tables grow richer, and the fact that not every doubly stochastic table can be dealt this way is the unistochasticity constraint of §1.6 appearing at the card table. One boundary must be stated as plainly as the successes: a card mechanism is a local classical toy. It reproduces single-system interference, indivisibility, and division events — the mechanism of this chapter — but no arrangement of cards on a table violates a Bell inequality, and the Bell-type correlations of §1.7 belong to the full construction, through the level-identification analysis taken up in Chapter 3 (§§3.7–3.7.1, where the boundary drawn here becomes the diagnosis), not to the toy. An illustration that claimed otherwise would be teaching a falsehood, and the analogy earns its keep precisely by knowing where it ends.

What remains is the pair of views. Turn the deck face up and track every permutation, and the game is chess: complete information, deterministic, exhaustively analyzable. Sit in the player’s seat and the same game is poker — and not the poker of static ignorance, but the modified game whose rational odds sheet carries interference terms. Nothing changed but the seat. That is the framework’s thesis in one gesture: observation is what the rest of the board looks like from inside a piece.

Schrödinger’s cat, played as cards. Every reader arrives carrying one thought experiment, and the game just built is the right table to play it on. Schrödinger (1935), in the same paper that coined *entanglement*, sealed a cat in a chamber with a radioactive atom, a Geiger counter, a hammer, and a flask of poison: within the hour, the formalism’s entanglement chain has amplified one atom’s superposition, link by link, into a description assigning the cat itself equal parts alive and dead until the box is opened. He offered it as a “quite ridiculous case” — a *reductio*, not a report. If the wavefunction is a complete description and “measurement” an undefined primitive, nothing in the theory says where along the chain atom $\rightarrow$ counter $\rightarrow$ hammer $\rightarrow$ cat a definite fact appears. Ninety years of interpretation dispute is the residue of that question.

Played on this table, the paradox takes four sentences. The sealed hour is an unobserved middle street. “Alive + dead” is not a state of the cat but a line on the odds sheet — the indivisible pricing any player must use for a hand whose history routes through the burn pile. Opening the box is the peek with its forced re-burn: an intervention that severs a correlation, not the creation of a fact. And the deal was definite throughout — the substratum holds exactly one configuration at every moment; the superposition is bookkeeping for where the record resides, in cards the observer has not turned. What dissolves the *cat* specifically is a further observation: a cat’s middle street is never unobserved. Counter, air, thermal photons — the environment peeks continuously, each interaction a division event forcing the re-burn, so the divisible odds sheet is the correct one for cats, counters, and kittens alike, and the interference terms that would distinguish “alive + dead” from “alive or dead” are committed to the hidden sector faster than any protocol could retrieve them (§1.6). No consciousness is required to open this box; one stray photon holds a full hand.

Schrödinger’s instinct was therefore sound — no cat is ever alive-and-dead — but the framework’s repair is subtraction where the objective-collapse programs (GRW/CSL, Diósi–Penrose) offer addition. They make the cat definite by appending a stochastic law; the framework makes it definite by deleting a premise — the superposition was the observer’s ledger, not the chamber’s state, so there was never anything to collapse (and Chapter 18 records that additions of the Penrose type are structurally forbidden here). The two repairs part company at exactly the frontier of §1.6: engineered “cat states,” superpositions of macroscopically distinct configurations held genuinely unobserved, which collapse laws must eventually kill and the framework’s exactness theorem must eternally protect — the wager registered in Chapter 4 §4.5. The boundary drawn above remains load-bearing: a cat in one box is a single-system story, the toy’s home turf; it is Bell’s two boxes, not Schrödinger’s one, where the cards must be surrendered to the full construction.

The substratum that the framework’s algorithm operates on is constructed concretely in Chapter 2 — a three-dimensional cubic lattice with bilinear, block-diagonal dynamics whose structural constraints fix the framework’s empirical content. Chapter 3 takes up the framework’s full philosophical positioning, including its relationship to the structural realism program, its evasion of Bell, Kochen–Specker, and PBR-type no-go theorems through identification of which descriptive level each assumption fails at, and its parallels to the broader intellectual landscape of incompleteness results from Gödel through Wolpert. Chapter 4 develops the methodological posture — the single-foundational-commitment derivation pattern, the architecture-level engagement with the no-go results, the classification of claims by derivational rigor — that organizes the framework’s program as a whole. The remaining chapters trace the framework’s predictions through particle physics, gravity, cosmology, the structural cascade from atoms to intelligence, and the framework’s reach into chemistry, biology, medicine, and quantum engineering.

The remainder of this part of the book develops the substratum, the structural realism, and the methodology in detail. Part II turns to the framework’s quantitative physics. Part III traces the cascade from substratum to chemistry, biology, and intelligence. Part IV develops the framework’s reach into working scientific fields.

Chapter 2

The Substratum

2.1 What this chapter constructs

Chapter 1 established that any observer embedded in a finite deterministic system, satisfying conditions C1-C3 on what lies outside their access, must use quantum mechanics. The theorem treated the underlying deterministic system (S, φ) as a given input — a finite configuration space and a bijection on it — and analyzed what the observer’s marginalized description looks like. This chapter takes up the question deferred there: what is (S, φ) ?

Two questions need separation. The first is whether the observed physics determines (S, φ) uniquely. If many different bijections produced the same observables, the substratum would be underdetermined and the framework would have no concrete object to describe. The second is what symmetries (S, φ) has, once it is determined — what transformations of the substratum leave all observables unchanged.

The chapter answers both. §2.4 develops the *reconstruction theorem*: observed physics — quantum mechanics with Bell violations, finite boundary entropy, spatial isotropy — together with six structural assumptions uniquely determines (S, φ) up to a precise gauge equivalence, with the Standard Model gauge group emerging as a forced consequence. §2.5 characterizes the gauge equivalence itself through the *substratum gauge group* $\mathcal{G}_{\text{sub}}$, identifying four families of generators that exhaust all observables-preserving transformations of the substratum. §2.6 then shows that quantum mechanics, general relativity, and the arrow of time are three projections of the same construction; §2.7 explains why this dissolves the apparent QM–GR incompatibility rather than resolving it.

Before the formal results, §§2.2–2.3 develop the physical interpretation of (S, φ) — what kind of object it is, what the substrate objection looks like, and how the framework sits in the broader family of incompleteness results from Gödel through Turing.

The chapter’s metaphysical commitment is *ontic structural realism on the equivalence class* $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$. What is real at the substratum level is the equivalence class, not any specific bijection within it. The reconstruction theorem

shows the class is uniquely determined by observation plus structural assumptions; the gauge group shows that different representatives within the class are observably indistinguishable. A stronger claim — realism about a specific bijection — would commit to features the framework’s own gauge group identifies as unobservable. A weaker claim — pure explanatory economy without ontic commitment — would undersell what the uniqueness theorem actually establishes. Chapter 3 develops the structural realism position in its full hierarchical form, including its operation across structural classes and its relationship to the broader philosophical literature.

2.2 The substratum as a finite lossless memory

The framework’s central object is a triple. S is a finite set whose elements are the distinguishable total states. $\varphi : S \rightarrow S$ is a bijection — a deterministic, reversible transition rule that maps each state to its unique successor. $V \subsetneq S$ is the visible sector, the observer’s window into S ; the complement $H = S \setminus V$ is the hidden sector. These three together — (S, φ, V) — fix everything: the dynamics, the observer’s access, the marginalized description.

The most useful physical interpretation is in terms of memory. S is the set of all distinguishable states the system can occupy — a finite storage capacity. φ is a bijection, so information is never created or destroyed; the substratum is a *lossless* memory in the precise information-theoretic sense. The partition V defines the observer’s access to this memory, both read and write. The observer reads the visible sector by measuring x , and writes to the hidden sector — every visible-sector interaction imprints correlations onto H through the coupling H_{int} .

Quantum mechanics is the statistics of this read-write cycle. The slow-bath condition C2 ensures that information written to the hidden sector during one visible-sector interaction is preserved until the next; subsequent reads retrieve correlations laid down earlier. The information backflow that Chapter 1 identified as the structural signature of P-indivisibility is exactly the signature of write-then-read at the substratum level. The Born rule is the equilibrium of the read-write cycle. Interference is the partial cancellation that occurs when two writes deposit related correlations in the hidden sector and a subsequent read combines them. Entanglement is shared correlations written into the hidden sector by two visible-sector subsystems during a joint preparation and read back later. None of these are independent quantum-mechanical phenomena; all are features of how the read-write statistics on a finite lossless memory look from the inside.

This memory interpretation makes several of the framework’s structural quantities concrete. The 10^{122} vacuum-energy discrepancy of Chapter 1 is the compression ratio between total storage and readable storage — between the full configuration space S and the visible sector V . The dark sector of cosmological observation is the gravitational effect of the unreadable storage in H . The Bekenstein-Hawking entropy is the storage capacity of the partition boundary,

the number of independent modes crossing ∂V . The action scale $\hbar$ is the conversion factor between storage geometry (boundary area) and read-write statistics (the unitary group's parameter). Chapter 7 develops each of these as theorems; for the present chapter they motivate the interpretation.

The substratum dynamics is defined on the full configuration space, not on the visible sector alone. The wave equation that selects φ among second-order reversible nearest-neighbor dynamics — developed in detail in Chapter 5 — operates on both V and H . The partition imposes an observer-access structure on this common dynamics but does not change the dynamics itself. Both sectors are governed by the same underlying rule; the only difference is which states the observer can read.

2.3 What the substratum is and is not

The substrate objection asks an obvious question: what is the memory made of? What physical stuff implements the configuration (S, φ) ? The framework's answer is that the question is malformed. (S, φ) is the complete description of reality — it determines all observables, and space, time, matter, and energy are derived from it rather than presupposed by it. Asking what the substratum is made of presumes a more fundamental substance from which (S, φ) is constituted, but the framework's structure leaves no room for any such substance. Whether (S, φ) *is* reality or *describes* reality is provably undecidable by any measurement an embedded observer can perform — an instance of the factoring argument of Chapter 7: the two readings determine identical accessible data, and any adjudication procedure available to an embedded observer is a function of that data, so no measurement-based criterion can separate them; the reconstruction theorem will supply the identical-data premise. The question therefore has no empirical content for physics. Whether it has substantive content in mathematical foundations or other modes of inquiry not bound by embedded physical observation is a separate matter the framework does not address.

A closer analog than physical substance is the abstract structure of a computation. A Turing machine has three components: a tape (the storage medium), a head (the read-write access window), and a transition function (the update rule). The mapping to the substratum is direct: S corresponds to the tape, V to the head's access window, φ to the transition function. But three differences are physically significant.

First, a Turing machine's tape is potentially infinite while S is finite. Finiteness is essential for the Poincaré recurrence argument that drives the P-indivisibility theorem of Chapter 1, and for the holographic bound on storage capacity of any bounded region.

Second, a Turing machine is generally irreversible — it can erase, overwrite, and halt. The substratum bijection φ is reversible by Lemma 3 of Chapter 1: nothing is erased, nothing is created, there is no halting state. Reversible Turing

machines (Bennett, Fredkin–Toffoli) are a well-studied subclass; the substratum sits in this class.

Third, a Turing machine computes an *extrinsic* function — input from an external user, output to an external observer. The substratum computes no extrinsic function. It is a closed permutation that cycles through its states; no input, no output, no halting condition. The appearance of dynamics, probability, particles, and forces is entirely the observer’s perspective — what the permutation looks like through the partial window V .

The third difference is the load-bearing one. The framework extends Turing’s question. Where Turing asked what computation a machine can perform for an external observer, the framework asks what computation looks like to a *component of the machine* — a subsystem with bounded access to the machine’s own state. The answer, established in Chapter 1, is quantum mechanics.

This extension places the framework in a specific intellectual lineage. Three foundational results form what can be called the incompleteness family. Gödel: a formal system rich enough to encode arithmetic cannot prove all true statements about itself; the unprovable truths have rigid structure. Turing: a computer cannot decide all questions about its own behavior; the undecidable problems have rigid structure. Observational Incompleteness: an observer embedded in a deterministic system cannot access the complete state; the emergent description has rigid structure — quantum mechanics. The common pattern is self-reference under finite resources, with the impossibility result yielding not chaos but a precisely characterized formal structure.

The OI analog of Turing’s halting problem is the question: can the observer determine the hidden-sector state h ? The question is well-posed. The hidden state has a definite value at every moment, because (S, φ) is deterministic. Different h values produce different physical outcomes through the bijection. But the observer provably cannot determine h : multiple hidden states are compatible with any visible-sector history, and the transition probabilities $T_{ij}(t)$ are averages over them. The structural consequence of this inaccessibility is quantum mechanics, just as the structural consequence of the halting problem’s undecidability is computability theory.

Two further classes of inaccessible quantities deserve separation from the undecidable hidden state. The alphabet size q at each lattice site and the cardinality of the deep hidden sector beyond the partition boundary are *gauge* rather than undecidable. Different values of q produce identical observables; different deep-sector cardinalities — including infinite ones — produce identical observables through the boundary-only dependence lemma of Chapter 7. The hidden state h is undecidable: there is a real answer, but the observer cannot find it. The alphabet q is gauge: there is no answer to find, because the question is empty. This distinction will be made precise in §2.5, where the gauge group is constructed explicitly.

The reconstruction theorem developed in the next section makes both inacces-

sibility claims precise. The substratum is uniquely determined by observed physics plus structural assumptions, up to the gauge group’s action — and the gauge group is exactly the set of transformations the observer cannot detect.

2.4 The reconstruction theorem

Chapter 1 took the substratum (S, φ) as input and derived quantum mechanics from it. The converse direction — from observed physics to the substratum — is the reconstruction theorem. If observed physics uniquely determines the substratum up to a precise equivalence, then the framework is not free to posit any deterministic system that happens to produce quantum mechanics under marginalization; the substratum is fixed by the data. This section establishes that uniqueness, and identifies the gauge equivalence that characterizes the residual freedom.

The theorem takes two kinds of input. The first is *empirical*: structural facts about the observed universe. The second is *structural*: assumptions about the class of candidate substrates the reconstruction is permitted to consider. Both are necessary, and both must be stated explicitly. A reconstruction theorem that did not specify its structural assumptions would either claim too much (uniqueness without restriction) or claim too little (uniqueness only within an undefined class).

Empirical inputs. The reconstruction draws on seven structural facts about observed physics, each well-established by experiment:

(E1) *Unitary quantum mechanics.* The observed microscopic physics is described by states in a complex Hilbert space evolving unitarily, with observables as self-adjoint operators and measurement outcomes following the Born rule. This is Chapter 1’s framework input.

(E2) *Bell violations.* Observed correlations between spatially separated measurements violate Bell-type inequalities, ruling out local hidden-variable theories satisfying factorizability. Confirmed experimentally since Aspect’s 1982 measurements and refined in loophole-free tests since 2015.

(E3) *Finite boundary entropy.* The entropy of any bounded region of spacetime is finite and scales as the boundary area rather than the volume. This is the holographic bound, supported by black-hole thermodynamics and the cosmological horizon’s finite entropy.

(E4) *Spatial isotropy.* Observed physics is rotationally invariant; no spatial direction is preferred. Confirmed across observations from the CMB to laboratory tests of fundamental constants.

(E5) *Propagating gravity.* Gravitational degrees of freedom propagate as waves with the observed signature of two transverse polarization states, directly detected since LIGO 2015.

(E6) *Stable matter.* Atoms, planets, and bound states exist on observed timescales. This requires both gravitational and atomic-scale stability against the dynamics.

(E7) *Boundary-entropy concordance.* The observed cosmological structure has spatial flatness near the critical density: $\rho_s/\rho_{\text{crit}} \approx 1$. The framework’s general- d form of this ratio is $\rho_s/\rho_{\text{crit}} = 2/(d-1)$, which equals unity only at $d = 3$.

The inputs E5, E6, E7 each provide an independent dimensional constraint. Propagating gravity (E5) requires $d \geq 3$, since the Weyl tensor governing gravitational-wave propagation vanishes in $d \leq 2$. Stable matter (E6) requires $d \leq 3$, since the Coulomb potential gives unstable atoms in $d \geq 4$ (Ehrenfest’s theorem) and gravitational orbits are unstable in $d \geq 4$ (Tangherlini’s result). Boundary-entropy concordance (E7) gives $d = 3$ exactly. Without stable matter no embedded observers exist; without propagating gravity no observable GR exists; without concordance the spatial-flatness data is unexplained. The three filters converge on $d = 3$ by independent routes, making the spatial dimension a forced consequence rather than an input to the reconstruction.

The convergence has an external echo. A literature independent of this framework reaches the same selection from different starting points: Ehrenfest’s stability analysis (atoms and planetary orbits are stable only for $d = 3$), the sharp-propagation property of the wave equation (Huygens’ principle holds cleanly only in odd spatial dimensions, and distortion-free signal propagation singles out $d = 3$, as Hadamard’s analysis of the wave equation shows), and Tegmark’s predictability argument (field equations well-posed enough for observers to process information select $3 + 1$). These are convergent arguments rather than theorems, and they are not an independent axiom base in the sense of the reconstruction theorems of Chapter 1: like the framework’s own three filters, they condition on the existence of stable structure and embedded observers. Their force is corroborative texture — several instruments, pointed from different premises within the same observer-admitting family, land on the same integer.

Structural assumptions. The reconstruction restricts attention to candidate substrates satisfying six properties:

(A1) *Finiteness.* The configuration space S is finite. This is the input from Lemma 1 of Chapter 1, supported empirically by E3 (the holographic bound implies a finite Hilbert-space dimension cutoff).

(A2) *Determinism.* The dynamics $\varphi : S \rightarrow S$ is a bijection — deterministic and reversible. This is the input from Lemma 3 of Chapter 1.

(A3) *Bounded coupling degree.* Each site of the substratum is coupled to a bounded number of neighbors through φ . This is required for the emergence of a coupling graph with well-defined dimensional structure and underlies locality in the emergent description.

(A4) *Center independence.* The dynamics does not depend on a choice of preferred site or origin; φ commutes with lattice translations up to gauge. This is

required to derive the wave equation in Stage 2 below.

(A5) *Linearity.* The wave equation governing φ is linear. Nonlinear alternatives are not ruled out as theoretical possibilities but would require a separate derivation chain not developed here.

(A6) *Background independence.* The dynamics is invariant under spatially varying internal-index transformations that preserve the coupling structure pointwise. The promotion of the resulting global symmetry to local gauge invariance is a derivation step rather than an assumption.

The reconstruction theorem’s claim is that under the conjunction E1–E7 and A1–A6, the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ is uniquely determined, with the Standard Model gauge group $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$, three chiral fermion generations, one Higgs doublet, anomaly-free hypercharges, $\theta = 0$, and $\hbar = c^3 \epsilon^2 / (4G)$ all forced consequences. The reconstruction proceeds in three stages, with each stage taking specified inputs and producing specified outputs.

Stage 1: From observed quantum mechanics to a deterministic embedding. The inputs are the quantum-mechanical content E1, the Bell-violation content E2, the boundary-entropy bound E3, and the finiteness and determinism assumptions A1, A2. Chapter 1’s characterization theorem (§1.8) supplies the construction: any system equivalent to unitarily evolving quantum mechanics arises from marginalization of a deterministic bijection satisfying C1–C3 for some partition. Bell violations force C1 (non-trivial coupling); the unitary structure of QM forces C2 and C3 (conditional on the hidden-sector mixing hypothesis for C2, ETH being its physical motivation). The output of Stage 1 is a triple (S, φ, V) with S finite, φ a bijection, and $V \subset S$ the visible-sector partition, with the observed quantum dynamics reproduced as the marginalization of φ over the hidden sector $H = S \setminus V$.

Uniqueness at Stage 1 follows from the Stinespring uniqueness theorem [Chapter 1, §1.6]. Any two dilations of the same quantum channel on the same visible Hilbert space differ by a partial isometry on the hidden sector — a transformation that, at the substratum level, is a relabeling of hidden-sector states. This freedom is part of the gauge group $\mathcal{G}_{\text{sub}}$ that §2.5 will characterize explicitly. Up to that gauge, Stage 1’s output is unique.

Stage 2: From deterministic embedding to specific lattice and gauge structure. Stage 2 takes Stage 1’s output and the remaining empirical and structural inputs E4–E7 and A3–A6, and derives the specific architecture of (S, φ) : the spatial dimension, the wave equation, the internal-index structure, the gauge group, the matter content, and the discrete symmetries.

The dimension argument follows the three-filter convergence described above. By Gromov’s theorem combined with statistical isotropy (E4), the coupling graph is quasi-isometric to $\mathbb{Z}^d$ for some integer d . The three filters from E5, E6, E7 then force $d = 3$.

The wave equation follows from center independence (A4), isotropy (E4), and

linearity (A5), together with the nearest-neighbor restriction (an assumption beyond A3's bounded coupling degree). The unique second-order reversible nearest-neighbor dynamics compatible with these constraints has the form $f = \alpha(x_1 + x_2 + \dots + x_{2d}) \bmod q$, propagating with speed $v = \alpha$. The coupling constant α is fixed by relativistic causality with maximum signal speed c realized at the lattice cutoff.

The number of internal components per lattice site follows from coupling-degree minimization applied to the wave equation: $K = 2d = 6$ uniquely. This locks the absolute generation count of fermions at three, since three spin-1/2 staggered tastes emerge from the $K = 6$ minimum (Chapter 6 develops this).

The gauge group follows from the cubic-group decomposition of the six internal components. The cubic group acts on the $K = 6$ components with eigenvalue multiplicities $(3, 2, 1)$. The pointwise commutant of this action — the set of unitary transformations commuting with the cubic-group representation — is $U(3) \times U(2) \times U(1)$ on the corresponding eigenspaces. Reduction to $SU(3) \times SU(2) \times U(1)$ occurs because the overall $U(1)$ phase of each block is gauge under amplitude-scale invariance (the field-value-rescaling gauge) in $\mathcal{G}_{\text{sub}}$ (§2.5 develops this). The remaining singlet-block $U(1)$ becomes the hypercharge $U(1)$. Background independence (A6) then promotes the global $SU(3) \times SU(2) \times U(1)$ to local gauge invariance.

The matter content — three chiral fermion generations with anomaly-free hypercharges and one Higgs doublet — follows from the staggered-fermion reduction, the link-carrier construction, partition-spinor identification, chirality emergence from the trace-out, and anomaly cancellation. Chapters 5 and 6 develop each step in detail; the chain ends with hypercharges $(Y_Q, Y_u, Y_d, Y_L, Y_e) = (1/6, 2/3, -1/3, -1/2, -1)$ uniquely determined.

The discrete symmetry $\bar{\theta} = 0$ follows from T-invariance of the bijection φ combined with detailed balance in the emergent Hamiltonian. The QCD theta-angle, which has been measured to be smaller than 10^{-10} but has no explanation within the Standard Model as conventionally formulated, is exactly zero as a consequence of the substratum's structural T-symmetry.

Uniqueness at Stage 2 follows from the if-and-only-if character of each step. Each derivation is the unique consistent choice given its inputs; the chain is a composition of unique selections, and the output is unique up to $\mathcal{G}_{\text{sub}}$.

Stage 3: From dynamics to emergent constants. Stage 3 takes Stage 2's output and derives the emergent value of $\hbar$. The argument is that thermal self-consistency at the partition boundary — established in Chapter 7 from the boundary layer's emergent thermodynamics — forces the relation

$$\hbar = c^3 \epsilon^2 / (4G),$$

with the discreteness scale $\epsilon = 2\ell_P$ (twice the Planck length) as the unique simultaneous solution. The Bekenstein-Hawking entropy $S = A/(4\ell_P^2)$ with the 1/4 coefficient follows as a consequence of the same mode-counting analysis.

Uniqueness at Stage 3 follows from the unique solvability of the boundary gap equation, which Chapter 7 develops.

The composite theorem. Composing the three stages yields the chapter’s central result.

Reconstruction Theorem. *Let E1–E7 (the empirical inputs) and A1–A6 (the structural assumptions) hold, and assume the hidden-sector dynamics satisfies ETH for the C2 necessity in Stage 1. Then the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ is uniquely determined: a finite set with a deterministic bijection of bounded coupling degree and statistical isotropy, on a three-dimensional cubic lattice with $K = 6$ internal components per site, coupling-matrix eigenvalue multiplicities $(3, 2, 1)$, gauge group $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$, three chiral fermion generations, one Higgs doublet, anomaly-free hypercharges, $\theta = 0$, and $\hbar = c^3 \epsilon^2 / (4G)$. $\square$*

The theorem’s content is striking. From the empirical inputs of observed physics — quantum mechanics, Bell violations, finite boundary entropy, spatial isotropy, propagating gravity, stable matter, spatial flatness — together with six structural assumptions, the substratum is forced to be a specific kind of object with specific properties. The Standard Model gauge group is not chosen; it is derived. The three-generation structure is not posited; it is forced. The Bekenstein-Hawking $1/4$ factor is not a postulate; it is a theorem. The cosmological constant problem is not resolved by tuning; the framework’s structure makes $\hbar$ and G commensurable through the boundary thermodynamics, and the apparent 10^{122} discrepancy is the compression ratio Chapter 7 develops.

Two scope remarks. First, uniqueness is within the structural class A1–A6. The theorem does not exclude alternative substrates outside this class: intrinsically continuum theories (violating A1), intrinsically stochastic theories (violating A2), theories with unbounded coupling (violating A3), or theories violating center independence, linearity, or background independence (violating A4–A6). The framework’s defense of A1–A6 is that they are the natural assumptions given the empirical inputs, but this defense is an argument rather than a theorem. A separate analysis would be needed to rule out alternatives violating any A_i .

Second, the empirical inputs E1–E7 are stated as structural facts rather than as specific measurement outcomes. Operationally, no finite measurement set establishes any of them; finite data is consistent with them but does not uniquely require them. The reconstruction theorem holds for the inputs as stated — the idealized infinite-data regime — and proves uniqueness in that regime. The operational uniqueness statement, against finite data, comes from the cumulative weight of the framework’s twenty-two quantitative predictions of Standard Model observables (Chapter 6): every observable that is correctly retrodicted rules out reconstructions that would have produced different values. The reconstruction theorem establishes idealized uniqueness; the prediction set establishes the operational uniqueness available against actual measurement data.

The reconstruction theorem identifies the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ as the

framework's central object. The next section makes the gauge group $\mathcal{G}_{\text{sub}}$ explicit by identifying its generators and showing they exhaust the freedom in the reconstruction.

2.5 The substratum gauge group

The reconstruction theorem identifies $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ as the framework's central object. This section makes $\mathcal{G}_{\text{sub}}$ explicit by exhibiting four families of generators and proving they exhaust the group.

The substratum gauge group is defined by the equivalence relation under which the reconstruction is unique: $(S, \varphi) \sim (S', \varphi')$ if the two systems produce identical emergent physics — the same transition probabilities $T_{ij}(t)$, the same emergent Hamiltonian (up to a residual diagonal gauge to be identified below), and the same value of $\hbar$ — for all partitions of the same structural class. The set of transformations mapping one substratum to an equivalent one forms a group, $\mathcal{G}_{\text{sub}}$.

The gauge group is not a symmetry of a Lagrangian or an action. There is no Lagrangian at the substratum level; the substratum is a finite set with a bijection on it. $\mathcal{G}_{\text{sub}}$ is a symmetry of the *equivalence class of substrata* defined by the observability condition. The emergent gauge symmetries (the Standard Model gauge group, the diagonal phase ambiguity in the emergent Hamiltonian) are Lagrangian symmetries in the standard sense, derived from the substratum through the trace-out.

Theorem (Generators). $\mathcal{G}_{\text{sub}}$ contains four independent families of transformations:

(i) *State relabeling.* For any bijection $\sigma : S \rightarrow S$, the substratum $(S, \sigma \circ \varphi \circ \sigma^{-1})$ produces the same emergent physics as (S, φ) . The transition probabilities depend on the coupling structure of φ — which states are sent to which — but not on which labels are attached to which states. This is an $|S|!$ -element subgroup, vastly larger than any gauge group appearing in the Standard Model.

(ii) *Alphabet change.* The local state space at each lattice site can be replaced by $\mathbb{Z}/q'\mathbb{Z}$ for any integer $q' \geq 2$, while preserving the coupling graph and the dynamics class, without changing any observable. The internal alphabet at a lattice site is the number of distinguishable local states; for any $q \geq 2$ the framework's predictions are identical. This family is parametrized by the integers $q \geq 2$.

(iii) *Deep-sector enlargement.* The hidden sector beyond the partition boundary — what Chapter 7 will call the *deep* hidden sector $\mathcal{C}_D$, distinct from the boundary layer $\mathcal{C}_B$ — can be enlarged arbitrarily without changing observables. Adjoining additional degrees of freedom to $\mathcal{C}_D$ with arbitrary dynamics satisfying $\tau_B^D \gg \tau_S$ leaves the emergent description unchanged. The boundary-only dependence

lemma (Chapter 7) proves that visible-sector transition probabilities depend only on the boundary layer $\mathcal{C}_B$ to corrections of order τ_S/τ_B^D . The deep sector may be finite of any size, or infinite — the observer’s predictions are insensitive to which.

(iv) *Graph isomorphism (up to statistical isotropy), with a gauge-sector qualification.* Two coupling graphs G_φ and $G_{\varphi'}$ that are quasi-isometric, with the same polynomial growth exponent d , the same spectral mode count, and the same statistical isotropy at large scales, produce the same emergent *quantum-mechanical, dimensional, and gravitational* physics; for that content the regular cubic lattice $\mathbb{Z}^3$ and any bounded-degree random graph with $d = 3$ polynomial growth and statistical isotropy are gauge-equivalent, and the reconstruction’s output specifies the *equivalence class* of such graphs, not the specific graph. The *Standard Model gauge group* is where this crystallography is read out: its multiplicities $(3, 2, 1)$ come from the octahedral point group on the six link directions, which is not a quasi-isometry invariant, so only the spectral mode count $2d_s = 6$ transfers to a generic isotropic graph — not the cubic decomposition. For the gauge sector, therefore, generator (iv) is restricted to orientation- and octahedral-structure-preserving isomorphisms, and the simple-cubic Bravais type is an empirically-selected input rather than a gauge choice. The gauge group itself is a universality-class invariant, shared with any Standard-Model-reproducing class through different machinery ([Structure §2.3, §14.4]); within OI’s class it is realized through the cubic-group commutant.

Proof that each family preserves all inputs to the derivation chain. Generator (i): conjugation preserves the coupling graph G_φ up to relabeling of vertices, hence preserves all graph-dependent quantities — the area law, the dispersion relation, the dimension via Myrheim-Meyer, the eigenvalue multiplicities $(3, 2, 1)$. Generator (ii): q -independence of every Standard Model prediction is established in Chapters 5 and 6 (the predictions involve only ratios and cubic-group representation theory, both q -independent). Generator (iii): the boundary-only dependence lemma gives the result directly. Generator (iv): the derivation chain in Stage 2 uses only statistical properties of G_φ — dimension, spectral dimension, bounded degree, isotropy — never specific graph features. $\square$

Completeness. The four families exhaust $\mathcal{G}_{\text{sub}}$ — any observables-preserving transformation can be decomposed as a composition of generators (i)–(iv).

The proof uses the Stinespring uniqueness theorem of Chapter 1. Suppose $g : (S, \varphi) \rightarrow (S', \varphi')$ is any observables-preserving transformation between two substrata. By assumption, g preserves the visible-sector quantum channel Φ — the CPTP map that gives the emergent quantum dynamics. Stinespring’s uniqueness theorem states that any two dilations of the same channel on the same visible Hilbert space differ by a partial isometry on the hidden sector. At the substratum level, this partial isometry is either a relabeling of hidden-sector states (generator (i) restricted to $\mathcal{C}_H$) when the two hidden sectors have equal cardinality, or a composition of relabeling and deep-sector enlargement/reduction (generators (i) + (iii)) when they differ in size. The al-

phabet may change freely (generator (ii)). The coupling graph must have the same statistical properties — dimension, spectral structure, isotropy — since these determine the derivation chain’s outputs, and two such graphs are related by generator (iv). These four degrees of freedom exhaust the freedom in the reconstruction, so the four generators span $\mathcal{G}_{\text{sub}}$.

Three candidate fifth generators are explicitly subsumed by the four already listed:

Time reversal ($\varphi \rightarrow \varphi^{-1}$). The wave equation that selects φ is T-invariant, giving $\varphi^{-1} = T \circ \varphi \circ T^{-1}$ where T is the phase-space layer swap exchanging successive timesteps. This is generator (i) with $\sigma = T$.

Hidden-sector dynamical reparametrization beyond enlargement. By Stinespring uniqueness, any two same-channel dilations of equal hidden-sector dimension differ by a hidden-sector unitary — generator (i) restricted to $\mathcal{C}_H$.

Visible-sector emergent global phase ($U \rightarrow e^{i\theta}U$). At the substratum level, S is a finite set with no complex structure; the emergent phase corresponds to no substratum-level transformation. It is the identity on (S, φ) .

The three-level gauge hierarchy. The framework has three distinct gauge symmetries, each appearing at a different level of the derivation, with each level projecting onto the next through the trace-out:

Level 3 — substratum gauge. $\mathcal{G}_{\text{sub}}$ acts on (S, φ) before any marginalization. It is the largest of the three gauge groups and includes transformations with no analog at lower levels (deep-sector enlargement, alphabet change). The reconstruction theorem holds modulo $\mathcal{G}_{\text{sub}}$.

Level 2 — emergent quantum field theory gauge. The Standard Model gauge group $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$ acts on the emergent fields after Stage 2 of the reconstruction has produced them. It is the image of $\mathcal{G}_{\text{sub}}$ restricted to transformations that permute internal components within the eigenspaces of the coupling matrix. The Standard Model gauge group is, in this picture, the visible-sector shadow of the substratum gauge group — the image of $\mathcal{G}_{\text{sub}}$ under the trace-out.

Level 1 — emergent Hamiltonian gauge. A residual diagonal-unitary freedom $H \rightarrow DHD^\dagger$ (where D is a diagonal unitary on the emergent Hilbert space) remains after all transition-probability data has been extracted. This is the smallest of the three gauge groups, corresponding to the rephasing of basis states in the emergent quantum description.

Each level is contained in the one above: $\text{Level 1} \subset \text{Level 2} \subset \text{Level 3}$. The trace-out projects Level 3 onto Level 2 (substratum gauge $\rightarrow$ Standard Model gauge); restriction to the Hamiltonian projects Level 2 onto Level 1.

The three-level hierarchy provides a structural account of why the Standard Model gauge group is what it is. The conventional puzzle is: why $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$ rather than some other group? Standard Grand Unified Theories address this by embedding the Standard Model in a larger simple group at

high energy and then breaking it; the framework here answers differently. The Standard Model gauge group is the image of the substratum’s symmetry under the partial-trace operation. The substratum has a larger symmetry group (all the freedom in generators (i)–(iv) plus the SM gauge structure); the trace-out projects this onto the smaller emergent gauge group; the residual freedom in the emergent description is the smaller Level 1 diagonal gauge. The Standard Model gauge group is fixed by the structure of the substratum, not chosen from a landscape.

Substrate freedom and the realism question. The substratum gauge group makes precise an earlier claim. Different gauge representatives within an equivalence class produce identical observables: state relabelings, alphabet changes, deep-sector enlargements, and graph isomorphisms all leave every measurement outcome unchanged. The question “which specific substratum representative is real?” therefore has no empirical content. Two representatives in the same equivalence class are observably indistinguishable, in the strongest possible sense — every measurement on either gives the same result. The framework’s ontic commitment is to the equivalence class, not to any specific representative within it.

This is what makes the framework’s structural realism non-trivial. The class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ is uniquely determined by observation plus structural assumptions, and is real in the sense that observation plus those assumptions force exactly this class. Individual representatives within the class are gauge expressions of the same physical content, like coordinate choices on a manifold or specific representatives of a homotopy class. The framework’s realism is precisely about what survives the gauge quotient.

The reconstruction theorem identifies the substratum; the gauge group identifies what within that identification has no observable content. With both in place, the framework’s central technical claim — that observed physics uniquely determines a specific kind of substratum, up to a precise gauge equivalence — is established. The next section turns to the synthesis: how quantum mechanics, general relativity, and the arrow of time relate to the substratum as three projections of the same construction.

2.6 The three projections

The reconstruction theorem and the substratum gauge group have established what (S, φ) is and what equivalences leave its observables unchanged. The question taken up in this section is what relationship (S, φ) bears to the framework’s three large emergent structures: quantum mechanics, general relativity, and the arrow of time.

The framework’s claim is that quantum mechanics, general relativity, and the arrow of time are not three independent theories. They are three projections

of the same finite deterministic construction. Each projection applies the same trace-out machinery to the same (S, φ) along a different axis, and each produces a distinct emergent structure as the observer sees the substratum from inside.

The visible-sector projection: quantum mechanics. Chapter 1 developed this projection in detail. The embedded observer’s compressed description of the visible sector, obtained by tracing out the hidden sector under conditions C1-C3, is mathematically equivalent to unitarily evolving quantum mechanics. The Hilbert space, the Schrödinger equation, the Born rule, and Bell-inequality violations all emerge from the marginalization, with no additional postulates. Applied to the substratum constructed in this chapter — a $d = 3$ cubic lattice with $K = 6$ internal components and the wave equation as substratum dynamics — the visible-sector projection produces the Standard Model: the gauge group $SU(3) \times SU(2) \times U(1)$ as the commutant of the cubic-symmetric coupling matrix, three chiral fermion generations, one composite Higgs, anomaly-free hypercharges, and $\bar{\theta} = 0$. Chapters 5 and 6 develop the Standard Model derivation; Chapter 8 develops the specific JUNO neutrino measurement that the framework’s parameter-free PMNS-angle value matches at 0.07σ (a retrodiction).

The boundary-thermodynamic projection: general relativity. A second projection takes the same substratum but applies the trace-out machinery to the partition boundary rather than to the bulk. The boundary of the visible sector — the surface across which the coupling H_{int} acts — carries its own thermodynamic structure: a temperature set by the boundary’s classical surface gravity, an entropy set by the number of boundary modes, and a dynamical evolution set by the coupling between boundary and bulk. Jacobson’s thermodynamic argument applied to local causal horizons yields Einstein’s equations as the Clausius relation $dE = T dS$ at the boundary.

Applied to the cosmological horizon as a causal partition, the boundary-thermodynamic projection produces general relativity with quantitative predictions. The emergent action scale $\hbar = c^3(2\ell_P)^2/(4G)$ follows from thermal self-consistency at the boundary layer. The Bekenstein-Hawking entropy $S = A/(4\ell_P^2)$ with the specific $1/4$ coefficient follows from direct mode counting on the boundary surface ; GW250114 (September 2025) confirms the classical area theorem, which the framework preserves — the $1/4$ coefficient itself has no direct empirical test. The cosmological constant problem dissolves: the quantum vacuum energy and the effective Λ are properties of logically distinct levels of description rather than commensurable quantities whose 10^{122} ratio requires cancellation. The structured component of the boundary entropy produces dark matter phenomenology with MOND acceleration scale $a_0 = cH/6 = 1.2 \times 10^{-10} \text{ m/s}^2$. Chapter 7 develops each of these.

The horizon-clock projection: the arrow of time. A third projection applies the trace-out machinery to the horizon-growth parameter along an observer worldline. The substratum is time-symmetric: φ and φ^{-1} are equally valid as forward-time maps, and there is no preferred direction at the substratum level. Observers, however, are not symmetric under this exchange. Condition C2 —

the slow-bath condition — picks out one direction along the observer’s worldline as the direction in which the hidden sector’s correlations with the visible sector’s past are preserved. The structural observer-selection theorem developed in Chapter 1’s appendix material restricts observers satisfying C1-C3 to the non-equilibrium phase of (S, φ) ; in the Λ -dominated cosmological phase, this phase is the one in which the horizon entropy $S_{\text{dS}}(t) = A(t)/(4\ell_P^2)$ is monotonically increasing.

The cosmological horizon entropy provides a structural clock. Every observer in the accelerating phase agrees that S_{dS} increases — the monotonicity is a feature of the causal partition rather than of any local frame — and other arrows of time align with this structural clock through a cascade. The thermodynamic arrow descends from Jacobson’s coupling between boundary entropy and matter entropy. The memory arrow descends from the slow-bath condition C2, which preserves correlations forward along the clock rather than backward. The measurement arrow — the observer records past outcomes and prepares future states — descends from the memory arrow. The chain grounds in one structural fact (horizon growth in the accelerating phase) and one structural constraint (C1-C3 observer selection). No independent past-hypothesis axiom is required, and Boltzmann brains are excluded structurally rather than anthropically.

The framework’s content here is precise: it derives a monotonic parameter along observer worldlines and four aligned arrows of time. What it does not derive is the convention of which direction along the horizon-clock parameter to call “future.” The two directions are gauge-equivalent at the substrate level ($\varphi \leftrightarrow \varphi^{-1}$), and the observer’s choice of convention is residual freedom. The framework derives the existence of a monotonic parameter; the labeling is an observer-imposed convention. There is one further qualification worth stating: a monotonic parameter is monotonic with respect to an *ordering* of the worldline’s moments, and the framework’s foundational treatment axiomatizes that there is more than one moment while leaving the ordering of moments as a separate, open question. The derivation here is therefore conditional on that ordering — it establishes that ordered moments yield a directed parameter, not that directedness arises from the substratum with no assumption about ordering at all. CP violation in CKM and kaon/B systems is a flavor-basis feature of the specific bijection φ rather than a consequence of the time-arrow cascade — the framework treats flavor-sector CP phases as solution-specific inputs that do not affect the cascade’s structural content.

The common structure. The three projections share three features. They apply to the same substratum (S, φ) . They use the same trace-out machinery — marginalization over the hidden sector under conditions C1-C3. They produce derived rather than fundamental structures: nothing in quantum mechanics, general relativity, or the arrow of time is a feature of (S, φ) in itself; each is a feature of how the embedded observer sees the substratum from inside. The differences among the three projections lie entirely in which structure of the embedded observer’s description they extract. The visible-sector projec-

tion works on the bulk dynamics and produces emergent quantum mechanics with the Standard Model as its specific content. The boundary-thermodynamic projection works on the classical thermal data at the partition boundary and produces general relativity with its gravitational and cosmological predictions. The horizon-clock projection works on the horizon-growth parameter along observer worldlines and produces the cascade of emergent arrows.

The shared structure is not metaphor or analogy. The same lattice that gives the Standard Model gauge group as the commutant of the coupling matrix gives the Bekenstein-Hawking entropy as the boundary mode count. The conditions C1-C3 that produce P-indivisibility for the visible-sector projection produce the slow-bath observer selection for the horizon-clock projection. The substratum gauge group $\mathcal{G}_{\text{sub}}$ that controls the reconstruction theorem’s uniqueness controls the gauge equivalences in all three projections. The framework’s claim is that the three projections are co-derived from a single construction, not three separate applications of similar machinery.

2.7 The QM-GR incompatibility as a category error

The conventional formulation of the QM-GR unification problem treats quantum mechanics and general relativity as two competing theories in the same logical category — two attempts to describe the same regime, awaiting reconciliation in a deeper unified theory. Every program of unification proceeds from this assumption. Quantum gravity programs treat the metric as a quantum field to be quantized. Geometric programs treat the wave function as a structure on the spacetime manifold. Emergent programs treat one of the two as derived from a deeper substrate that recovers the other. The persistent failure of all three approaches to produce a quantitatively confirmed unification — half a century of intensive work without an empirical test of the unification regime — suggests that the underlying assumption may be wrong.

The framework’s claim is that quantum mechanics and general relativity are not two theories of the same regime at all. They are two projections of the same construction, viewed at different levels and through different machinery. The visible-sector projection that produces quantum mechanics and the boundary-thermodynamic projection that produces general relativity refer to *different objects* in the construction. The wave function is a property of the embedded observer’s compressed description of the visible sector — the bookkeeping device that tracks what the observer knows and does not know about a system. The metric is a property of the classical thermodynamics of the partition boundary — the geometric content of the surface across which the trace-out is performed. These are not two descriptions of the same physical system at different levels of approximation. They are descriptions of two different substructures of the same total object (S, φ) , derived by different projection maps from a single source.

The structural analogy that captures the relationship most clearly is the relationship between thermodynamics and statistical mechanics. Thermodynamics

deals with pressure, temperature, and entropy. Statistical mechanics deals with particle positions, momenta, and microstate counting. These are not two competing theories of the same gas, and the question “how do we unify thermodynamics and statistical mechanics?” is not a coherent question. They describe different things about the same system. The relationship between them is one of projection — statistical mechanics produces thermodynamic quantities by averaging over microstates — rather than unification. Asking how to unify thermodynamics and statistical mechanics into a deeper theory mistakes the structure of their relationship.

The relationship between quantum mechanics and general relativity in the framework is structurally identical. General relativity describes the thermodynamic projection of (S, φ) : the boundary geometry, the surface gravity, the entropy, the dynamical evolution of the partition under coupling to the bulk. Quantum mechanics describes the statistical projection: the visible-sector dynamics under marginalization over the hidden sector, the P-indivisible transition probabilities, the unitary evolution that captures the non-Markovian structure. The two are related by a definite procedure (the trace-out under conditions C1-C3), with the substratum gauge group controlling the equivalences. The question “how do we unify them?” is dissolved rather than answered.

This is not the same as saying that quantum mechanics is more fundamental than general relativity, or vice versa. Both are emergent. Both depend on the trace-out and the partition. Both are derived structures rather than features of the underlying reality. The fundamental object is (S, φ) — a finite set with a deterministic bijection on it — and neither quantum mechanics nor general relativity is a feature of (S, φ) itself. They are features of how an embedded observer or a horizon-bounded thermodynamic regime sees (S, φ) from inside. (Nor is the ranking question recoverable from inside: by the factoring argument of Chapter 7, no criterion grounded in the accessible observations can separate two descriptions faithful to the same data — the fundamentality question is answered at the level of (S, φ) , not adjudicated between its projections.)

The dissolution of the unification problem has two consequences worth marking explicitly. First, it removes the QM-GR unification problem from the active research agenda. The problem was based on a category error, and the category error is fixed by the substratum-level construction. The persistent failure of unification programs is no longer puzzling — they were trying to merge two descriptions of different things, which cannot be done. Second, it provides an explanation for why the Standard Model has the gauge group it does. The Standard Model gauge group is the visible-sector shadow of the substratum gauge group $\mathcal{G}_{\text{sub}}$, projected through the trace-out. The gauge group is not chosen from a landscape of alternatives; it is the unique image of the substratum’s symmetry under the partial-trace operation, given the reconstruction theorem’s structural inputs.

The framework’s synthesis claim is structural and architectural, not predictive in the conventional sense. It does not claim to produce gravitational phenomena

that conventional quantum field theory plus general relativity cannot reproduce. The framework’s empirical content is concentrated elsewhere: in the Standard Model derivation (twenty-two quantitative retrodictions, Chapters 5 and 6), the JUNO neutrino prediction (Chapter 8), and the dark-sector predictions specific to the boundary-thermodynamic projection (Chapter 7). What the synthesis claim establishes is that these predictions are not independent. The same construction that produces the Standard Model produces the gravitational sector. The same trace-out that gives the Standard Model gauge group as the commutant of the coupling matrix gives the Bekenstein-Hawking entropy as the boundary mode count. The framework’s derivations across Chapters 5, 6, and 7 are not separate applications of similar machinery; they are co-derived consequences of a single construction.

2.8 What follows

The reconstruction theorem and the substratum gauge group establish the substratum side of the framework. The three projections establish its relationship to the framework’s three large emergent structures. With both in place, the technical core of Part I is complete; the chapters that follow develop its consequences in three directions.

Chapter 3 develops the framework’s philosophical positioning. The substratum’s ontological status — what it means to commit to the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ as the framework’s metaphysical object — is articulated in its full hierarchical form, with realism operating at multiple levels of the observation hierarchy and the gauge hierarchy. The framework’s evasion of Bell, Kochen-Specker, and PBR-type no-go theorems is developed through identification of which descriptive level each no-go assumption fails at. The framework’s intellectual position is located against the contemporary observer-dependence programs (QBism, relational quantum mechanics, the recent CLPW and Maldacena work on observer-incorporated entropy) and against the older philosophical lineage (Cusa’s *docta ignorantia*, Spinoza’s attribute theory, Worrall’s epistemic structural realism). The relationship to Hofstadter’s strange-loop framework is developed in detail.

Chapter 4 articulates the framework’s methodological discipline. The single-foundational-commitment derivation pattern — that the framework starts from one premise (observation occurs) and derives the rest as theorems — is examined as a methodological standard. The architecture-level engagement with the no-go results is developed: rather than evading Bell, Kochen-Specker, and PBR with four separate tricks, the framework meets all of them with one structural feature — the two-level substratum-emergent architecture — and verifies the match for each theorem. The classification of claims by derivational rigor — theorems, derivations, structural retrodictions, and informed speculation — organizes the framework’s program by epistemic status. The framework’s pre-registration of falsification conditions is discussed as a methodological commitment.

Parts II and III then develop the framework’s quantitative consequences. Chapter 5 develops the Standard Model gauge group from the cubic-lattice substratum: the wave equation, coupling-degree minimization, the cubic-group commutant decomposition, chirality from the trace-out, anomaly-free hypercharges. Chapter 6 develops the matter content and quantitative predictions: three generations (chirality conditional on $H\text{-}\chi$, SM §4.8), the composite Higgs, $\theta = 0$, gauge couplings, and the twenty-two classified quantitative retrodictions (seven parameter-free structural; Appendix A). Chapter 7 develops the gravitational sector: the boundary thermodynamics, the value of $\hbar$, the Bekenstein-Hawking $1/4$ factor, the dissolution of the cosmological constant problem, the dark sector predictions. Chapter 8 develops the JUNO neutrino prediction in detail. Chapter 9 takes up universality and the framework’s empirical scope.

The cascade through chemistry, biology, evolution, and intelligence in Chapters 10–12 takes the structural results of Parts I and II as derived inputs and develops their consequences for the structural conditions under which complex matter, life, and intelligence become possible. The applications in Chapters 13–17 develop the framework’s reach into working scientific fields: quantum biology, quantum computing and engineering, medicine, bioinformatics.

The framework’s claim, made precise across these chapters, is that a single starting commitment — that observation occurs — combined with the lemmas extracted from this commitment and the conditions identified for embedded observers in our universe, produces a substratum that is forced to have a specific architecture, with quantum mechanics, general relativity, the Standard Model, and the cosmological dark sector all forced consequences of the architecture rather than separate empirical inputs. The book’s role is to make this chain of derivations available for evaluation against the experimental record now being assembled and forthcoming.

Chapter 3

Hierarchical Structural Realism

3.1 What this chapter develops

Chapter 2 established the framework’s metaphysical commitment as ontic structural realism on the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$. What is real at the substratum level is the equivalence class, not any specific bijection within it. The chapter argued for this commitment in its own terms: stronger claims commit to features the gauge group identifies as unobservable, weaker claims undersell the reconstruction theorem’s uniqueness, and the equivalence class is what observation plus structural assumptions force.

That argument is correct but incomplete. The framework’s structural-realism position has more structure than the single commitment Chapter 2 articulated. Different parts of the framework’s content live at different levels of generality, with different empirical predictions, different falsifiability conditions, and different philosophical commitments at each level. The framework’s twenty-two Standard Model retrodictions live at one level; the central claim that quantum mechanics is the unique decompression algorithm for embedded observers lives at a broader level; the foundational fact that observation occurs lives at the broadest level. A single structural-realism commitment that applied to all of these uniformly would either claim too much at the foundational level or too little at the specific predictive level.

This chapter develops the framework’s structural realism in its full hierarchical form. The structural-realism position operates on a two-dimensional hierarchy: an *observation axis* running from the framework’s foundational empirical commitment to its specific predictive content, and an orthogonal *gauge axis* running from narrow gauge equivalences within the emergent quantum mechanics to broad equivalences across structural classes of embedded observers. Different empirical predictions live at different intersections of these axes; falsifiability is stratified accordingly; and the framework’s strongest realism claim — the Chapter 2 commitment to $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ — is what the hierarchy locates at a specific intersection.

The chapter has four pieces of work. §§3.2–3.3 develop the observation axis and the observer-admission distinction that the framework’s third level introduces. §§3.4–3.5 develop the gauge axis, its orthogonality with the observation axis, and the resulting multi-level structural realism with its stratified predictions. §3.6 identifies gauge freedom with observational incompleteness on a precisely delimited part of its range, distinguishing the two species of incompleteness this produces. §§3.7–3.8 develop the framework’s relationship to the contemporary no-go theorem literature and to the older philosophical lineage of incompleteness-and-perspective accounts. §3.9 closes the chapter and points forward to Chapter 4’s methodology and to Chapter 8’s universality and cross-framework analysis.

3.2 The observation axis

The framework’s content nests in four levels of increasing specificity along an axis I will call the *observation axis*. Each deeper level is a more specific commitment about what the framework asserts. At the broadest level, the framework commits only to the empirical fact that observation occurs; at the narrowest level, it commits to a specific lattice structure with twenty-two quantitative retrodictions. Between these poles lie two further levels — observer admission and universality classes — that articulate intermediate structural content.

Level A: the observation axiom. The framework’s foundational commitment, developed in Chapter 1, is the empirical fact that observation occurs: an observer records distinguishable outcomes of interactions with a system not

wholly under the observer’s control. This Cartesian *cogito*, made mathematically precise, is the one premise the framework takes from outside its own theorems. Level A commits to three things: that there is something to observe (a substrate of reality structured to support distinguishable states); that some part of reality is structured to record outcomes (observers are embedded in reality, not external to it); and that observation has structural consequences (the recording itself is a physical event, not a passive registration).

Level A is shared with the broader observer-essentiality program in contemporary foundations of physics. The work of Chandrasekaran, Longo, Penington, and Witten on observer-incorporated entropy in de Sitter space (2022), Maldacena’s analysis of observer’s-clock dependence in cosmological partition functions (2024), and Harlow, Usatyuk, and Zhao on observer-dependent Hilbert spaces in closed universes (2025) all start from analogous commitments. The framework’s distinguishing feature at Level A is not the commitment itself but what is derived from it: a finite deterministic substratum without extra dimensions, without intrinsic stochasticity, and without additional postulates beyond those identified in Chapter 1’s structural lemmas.

Level B: observer admission. Not every conceivable substratum admits embedded observation. Chapter 1 identified four structural conditions — C1 (non-trivial coupling), C2 (slow-bath timescale separation), C3 (sufficient hidden-sector capacity), and C4 (history readback) — that are necessary for a substratum’s bijection φ to support an embedded observer. A substratum (S, φ) is *observer-admitting* if there exists a partition $V \subset S$ such that the triple (S, φ, V) satisfies C1–C3; it is *observer-lacking* otherwise.

The distinction is structural rather than empirical. A finite cellular automaton with a deterministic update rule may or may not admit an embedded observer, depending on whether its dynamics support the slow-mixing, sufficient-capacity, coupled, read-back structure C1–C4 require. Random permutations on finite sets typically lack this structure: their cycle decompositions are wrong for the slow-bath condition, their internal correlations are wrong for the capacity condition, or their coupling is too uniform or too sparse for the non-trivial-coupling condition. The space of substrata is structurally partitioned into the observer-admitting subset, which can support embedded observers, and the observer-lacking subset, which cannot.

Level B commits to a structural-anthropic reading that sharpens the conventional anthropic principle. Standard anthropic reasoning treats fundamental parameters as contingent and explained by selection: we observe a universe with parameters compatible with us because we could not observe a universe without them. The Level B reading is sharper. Observer-lacking substrata are not candidates for “our universe” at all — they are mathematical objects of a different kind, lacking the observer-extraction structure that would make the question “what would observers in this universe see?” coherent. Within the observer-admitting subset, anthropic reasoning operates as usual; across the observer-admitting / observer-lacking boundary, the question of why we observe

a universe with our particular features becomes the question of why a universe must be observer-admitting at all. The framework’s answer is partly tautological — we are observers, so the universe we observe must admit observers — but the tautology is informative: it identifies which structural features of a candidate universe are forced by the existence of observers, and which are contingent.

The structural-anthropic content of Level B is independent of the specific universality class the universe belongs to. Any candidate substratum, regardless of dimension, gauge structure, or matter content, must satisfy C1-C3 to admit observers. The Level B commitment is therefore broader than any specific physics: it applies to OI’s cubic-lattice substratum, to string-theoretic compactifications insofar as they admit embedded observers, to any future candidate fundamental theory, and to mathematical structures that may or may not correspond to physical universes.

Level C: universality classes. Within the observer-admitting subset, equivalence classes arise from the partial-trace observational features that any embedded observer produces. Two substratum-with-observer systems are *observationally equivalent* if their algebra-channel pairs — the algebra of visible-sector observables together with the quantum channel induced by the partial trace over the hidden sector — are isomorphic in a precise sense that preserves both the algebraic structure and the channel’s dynamical content. A *universality class* is an equivalence class under this observational equivalence.

The universality-class structure is broader than the substratum gauge group of Chapter 2. The substratum gauge group $\mathcal{G}_{\text{sub}}$ classifies observationally-equivalent substrata *within OI’s specific structural class* — substrata satisfying A1–A6 of the reconstruction theorem. The universality-class equivalence operates more broadly: substrata in different structural classes can be observationally equivalent if they produce the same algebra-channel pair through structurally different machinery. Two substrata that are universality-class equivalent need not be related by any $\mathcal{G}_{\text{sub}}$ transformation; the broader equivalence allows for structural-class transitions that the substratum gauge group does not.

Universality classes have internal sub-structure at three levels of refinement. The most universal level is the *partial-trace features sub-class*, characterized by the structural features any partial-trace observation must produce regardless of substratum specifics: the Born rule, channel-level unitarity to leading order, P-indivisibility of the marginal stochastic process, and the commutant restriction pattern that selects which symmetries survive the trace-out. Below this is the *gauge-group sub-class*, characterized by which specific emergent gauge group the marginalization produces. Standard-Model-reproducing classes share $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$; other classes produce different gauge groups. Below this is the *algebra-channel sub-class*, characterized by the specific algebra of observables and specific channel dynamics — the level at which two specific substratum-with-observer systems become indistinguishable in every observable.

Level C commits to a structural realism at the class level: the universality class

is what is real, with specific substratum representatives playing the role of gauge expressions. Different members of the same universality class are observationally indistinguishable by definition, so the question “which specific substratum is real?” has no empirical content at the universality-class level. The realism is precisely about what survives the universality-class quotient.

Level D: OI’s specific representative. The framework’s narrowest and most specific commitment is to its particular substratum: a three-dimensional cubic lattice with $K = 6$ internal components per site, the linear wave equation as substratum dynamics, the cubic-group decomposition $(3, 2, 1)$ giving the Standard Model gauge group, the staggered-fermion construction giving three generations (chirality conditional on $H\text{-}\chi$, SM §4.8) and the link-carrier construction giving one composite Higgs, anomaly-free hypercharges and $\theta = 0$ as forced consequences. The reconstruction theorem developed in Chapter 2 establishes that observed physics together with the structural assumptions A1–A6 uniquely determines this specific representative up to the substratum gauge group $\mathcal{G}_{\text{sub}}$.

Level D is where the framework’s twenty-two classified quantitative predictions live (seven parameter-free structural among them): the Cabibbo angle $1/(\pi\sqrt{2})$, the Koide relation $2/9$, the gauge-coupling-running calculation giving the three SM coupling constants at the Z pole, the six fermion masses derived from one empirical input, the Higgs quartic coupling from cubic-group representation theory, the Bekenstein-Hawking $1/4$ factor, the MOND acceleration scale $cH/6$, and the JUNO neutrino-mixing prediction. The empirical content that distinguishes OI from other observer-admitting frameworks is concentrated at Level D, because Level D is where the structural commitments become specific enough to produce specific numerical predictions.

The relationship between Levels D and C deserves a separate remark. The structural conditions A1–A6 that define Level D are not the same as the observer-admission conditions C1–C3 that define Level B. C1–C3 are conditions for *any* embedded observer to exist; A1–A6 are conditions for membership in OI’s *specific* universality class within the observer-admitting subset. Some A’s relate to but are distinct from C’s: A1 (finiteness of the configuration space) is related to but distinct from C3 (sufficient hidden-sector capacity). A2 (determinism via bijection) selects deterministic-substratum members of a universality class, but observers can in principle exist on stochastic dynamics that satisfy the partial-trace observational features differently. A3–A6 (bounded coupling, center independence, linearity, background independence) are pure class-specific conditions, required for OI’s specific gauge-structure derivation but not necessary for observer admission. The Level D commitment selects a specific universality class within the observer-admitting subset, with both the empirical inputs of physics and the framework’s choice of structural assumptions playing a role in the selection.

The hierarchy nests. Each deeper level is a more specific commitment, with Level A $\supset$ Level B $\supset$ Level C $\supset$ Level D. Level A says only that observation occurs; Level B adds that the substratum’s bijection structure must support

observers; Level C adds the observational-equivalence structure that groups observer-admitting substrata into universality classes; Level D adds the framework’s specific structural assumptions and selects the cubic-lattice substratum. The framework commits to all four levels simultaneously, with content at each level being more universal (broader levels) or more class-specific (deeper levels).

3.3 The observer-admission distinction in more detail

The observer-admission distinction at Level B deserves more sustained development than the previous section allowed. It is the framework’s structural-anthropic move, and it has consequences that distinguish the framework from both standard anthropic reasoning and the broader observer-dependence program in contemporary foundations of physics.

Standard anthropic reasoning operates on fundamental parameters. The cosmological constant, the masses of the fermions, the strengths of the four forces — these are observed to lie in narrow ranges compatible with the existence of observers, and the anthropic principle explains this by selection: a universe with different parameters would not contain observers to make the observation. The reasoning is conditional. It takes the existence of observers as given and explains the observed parameter values by the conditioning on observer existence.

The observer-admission reading at Level B sharpens this in two ways. First, it identifies which substratum features must be present for observers to exist at all, prior to questions about specific parameter values. The C1-C3 conditions are structural features of a substratum’s dynamics: not values of free parameters but properties of the bijection’s behavior. A substratum either admits observers or it does not; the question is binary at the structural level, before any parameter values are specified. Second, the observer-admission reading provides a more concrete account of what “selection” means. The standard anthropic principle is sometimes criticized as tautological — observers observe universes with observers in them — but this criticism misses the structural content. Observer-lacking substrata are not failed candidates that happened not to produce observers; they are structurally different mathematical objects, lacking the dynamical features that would make embedded observation possible.

A concrete example clarifies the structural distinction. Consider two finite cellular automata with the same configuration space size but different transition rules. The first has a Glauber dynamics on an Ising lattice at high temperature: fast mixing, short correlation times, no slow modes. The second has a topological phase-protected dynamics: slow modes from non-local order parameters, long correlation times in those modes, sufficient capacity in the topological sector. Both are well-defined deterministic systems on the same configuration space. The first fails C2 because its hidden sector has no slow modes; the second satisfies C1, C2, C3 with appropriate partition choice. The first is observer-lacking; the second is observer-admitting. The distinction is structural — a property of the dynamics, not of the parameter values within a given dynamics — and

applies prior to any specific universality-class commitment.

This structural distinction has implications for several conceptual problems. The Boltzmann brain problem — the worry that a sufficiently large statistical-mechanical system at equilibrium would produce observers as random fluctuations more often than through evolutionary processes — is dissolved at Level B. Equilibrium states are not observer-admitting in the framework’s sense, because the slow-bath condition C2 fails at equilibrium: there are no slow modes by definition, every degree of freedom has thermalized. Boltzmann brains are not generic outcomes of long-time evolution that just happen to be rare; they are structurally excluded from the observer-admitting subset by C2. The framework’s account of the arrow of time, developed in Chapter 2 §2.6, follows from this exclusion: observers are restricted to the non-equilibrium phase of the substratum’s dynamics, and the structural arrow of horizon growth provides the clock the observer’s other arrows align with.

The string-theory landscape problem looks different under the Level B reading as well. The conventional landscape picture posits roughly 10^{500} metastable vacua, each a different low-energy effective theory, with the vacuum-selection problem being to explain why our universe is in this particular vacuum. The Level B reading reframes this: most compactifications likely produce observer-lacking substrata at the relevant scales, because the slow-bath, sufficient-capacity, non-trivial-coupling conditions are non-generic. The SM-reproducing subset of the landscape is structurally smaller than the total landscape, because it is restricted to the observer-admitting subset. Whether this reframing fully resolves the landscape problem is an open technical question — it depends on whether observer-lacking compactifications can be characterized in string-theoretic terms, which requires substantive work in string-theoretic landscape physics — but the reframing is non-trivial.

The framework’s commitment at Level B is therefore structurally substantive, not merely a definitional convention. It identifies the observer-admission boundary as a real feature of the space of candidate substrata, with consequences for anthropic reasoning, Boltzmann brains, the arrow of time, and the landscape problem. The realism at Level B is realism about the boundary itself: which substrata admit observers, which do not, and why. The empirical content at Level B is sparse — most of the framework’s predictions live at deeper levels — but the structural content is rich.

3.4 The gauge axis

The framework’s gauge structure runs along an axis orthogonal to the observation axis. Where the observation axis is about scope — which substrata fall within the framework’s domain — the gauge axis is about transformations: which substratum-level operations leave all observables unchanged. Like the observation axis, the gauge axis has four levels, with each broader level absorbing

more degrees of freedom as gauge.

Level G1: D-gauge on the emergent Hamiltonian. The narrowest gauge freedom in the framework is the diagonal-unitary equivalence on the emergent Hamiltonian: $H \rightarrow DHD^\dagger$ with D a diagonal unitary in the measurement-outcome basis. This is the residual freedom that survives after all transition-probability data has been extracted from the emergent quantum description. Two Hamiltonians H and $DHD^\dagger$ produce identical Born-rule probabilities $T_{ij}(t) = |\langle j|e^{-iHt/\hbar}|i\rangle|^2$ in the measurement basis, so they are operationally indistinguishable. The transformation rephases the diagonal entries of H in the measurement basis without changing any predicted probability.

Level G1 is the smallest gauge group in the hierarchy. It operates within a single Hamiltonian realization in a single emergent quantum mechanics — within Level D $\times$ Level G1 in the joint hierarchy, the framework’s most specific intersection. Its content is the bookkeeping freedom in choosing phase conventions for the emergent basis states, with no implications beyond that.

Level G2: the Standard Model gauge group. The next broader gauge freedom is the Standard Model gauge group $SU(3) \times SU(2) \times U(1)$ acting on the emergent quantum fields. Chapter 2 §2.5 developed this as the unitary group of the bicommutant of the cubic-group action (the strict pointwise commutant is $U(1)^3$ by Schur’s lemma) on the six internal components of OI’s substratum: the cubic-group decomposition gives multiplicities $(3, 2, 1)$, and the unitary group of the block algebra generated by the cubic-group representation on this decomposition is precisely $U(3) \times U(2) \times U(1)$, reducing to $SU(3) \times SU(2) \times U(1)$ after the amplitude-scale-gauge $U(1)$ phases are stripped. Background independence promotes this global commutant symmetry to local gauge invariance, giving the full SM gauge structure of standard quantum field theory.

Level G2 contains Level G1 as a sub-symmetry: every G1 transformation is also a G2 transformation (specifically, an element of the maximal torus of the gauge group). G2 is broader than G1 because it includes the off-diagonal gauge transformations that mix different internal components, which G1 does not access.

A subtlety in Level G2 deserves explicit treatment. The *uniqueness* of $SU(3) \times SU(2) \times U(1)$ — that this specific gauge group is forced rather than chosen from alternatives — depends on the structural inputs A3–A6 plus the empirical input E4 (spatial isotropy). Other universality classes that produce the Standard Model at low energies, including any successful string-theoretic compactification, would reach the same gauge group through different machinery. Level G2 is therefore not specific to OI’s substratum; it is shared with any framework that produces the Standard Model. What is specific to OI is *how* the gauge group is derived (cubic-group commutant) rather than the gauge group itself.

Level G3: the substratum gauge group $\mathcal{G}_{\text{sub}}$. Chapter 2 §2.5 developed the substratum gauge group as the set of transformations on the underlying bijection (S, φ) that preserve all observables — identical transition probabilities, identical emergent Hamiltonian (up to G1), identical value of $\hbar$ across all partitions of

OI's structural class. Four families of generators were identified: hidden-sector relabeling, alphabet change, deep-sector enlargement, and graph isomorphism up to statistical isotropy. The Stinespring uniqueness theorem established that these four families exhaust the freedom.

Level G3 contains both G1 and G2. The Standard Model gauge group is the visible-sector shadow of $\mathcal{G}_{\text{sub}}$ — the image of the substratum gauge group under the trace-out. The framework's structural account of why the Standard Model has its specific gauge group is precisely this: the SM gauge group is what $\mathcal{G}_{\text{sub}}$ looks like after the partial trace over the hidden sector has been performed.

Level G3 differs from G1 and G2 in an important structural respect: it is an *equivalence-class symmetry*, not a Lagrangian symmetry. There is no Lagrangian at the substratum level; the substratum is a finite set with a bijection on it, governed by a discrete update rule rather than a continuous action principle. The gauge group acts on the equivalence class of substrata defined by the observability condition. It plays the same structural role at the substratum level that Lagrangian gauge symmetries play at the field-theory level — identifying observationally indistinguishable configurations — but through a different mathematical mechanism.

Level G4: universality-class equivalence. The broadest gauge freedom in the hierarchy operates at the universality-class level. Two substratum-with-observer systems are universality-class equivalent if their algebra-channel pairs are isomorphic in the precise sense identified at Level C — preserving both the algebraic structure of visible observables and the channel's dynamical content. Level G4 captures equivalences that may not be expressible as transformations within any single structural class.

Level G4 is broader than $\mathcal{G}_{\text{sub}}$. Substrata related by $\mathcal{G}_{\text{sub}}$ are universality-class equivalent (since they produce the same algebra-channel pair), but substrata not related by any $\mathcal{G}_{\text{sub}}$ transformation may still be universality-class equivalent if they produce the same algebra-channel pair through structurally different machinery. The classic candidate is the relationship between OI's cubic-lattice substratum and string-theoretic compactifications that reproduce the Standard Model. The two are not related by any $\mathcal{G}_{\text{sub}}$ generator — OI's bounded-coupling, finite-deterministic structure is not preserved by string-theoretic compactifications, which violate A2 (determinism) and A5 (linearity) at the matrix-model bridge — but they may produce equivalent algebra-channel pairs, making them universality-class equivalent at Level G4 even though they are not gauge-equivalent at Level G3.

Level G4 operates only within the observer-admitting subset — Level B. Observer-lacking substrata are outside the universality-class structure entirely; the algebra-channel pair is undefined for them because there is no embedded observer to define the algebra of visible observables or the channel induced by the partial trace.

The gauge hierarchy nests. Each broader level absorbs more degrees of

freedom as gauge, with $G1 \subset G2 \subset G3 \subset G4$. The trace-out projects higher levels onto lower: trace-out of $\mathcal{G}_{\text{sub}}$ produces the SM gauge group as the image; restriction to the emergent Hamiltonian projects G2 onto G1. The framework commits to all four levels of gauge simultaneously, with different levels of gauge being relevant to different questions about the substratum and its emergent physics.

3.5 Orthogonality, multi-level realism, and stratified predictions

The two hierarchies — observation axis and gauge axis — are orthogonal. The observation axis is about scope: which substrata fall within the framework’s domain at each level of specificity. The gauge axis is about transformations: which operations leave observables unchanged at each level of breadth. A specific empirical claim or theoretical commitment lives at a specific *intersection* of the two axes, with the intersection determining both what the claim’s content is and how it can be falsified.

The orthogonality has substantive content. Levels G1, G2, and G3 sit within Level D — they are gauge structures of OI’s specific universality class with its specific gauge derivation. Level G4 operates at Level C — it characterizes universality-class equivalence, which is broader than OI’s structural class. The two-axis structure means the framework’s strongest realism claim — the Chapter 2 commitment to $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ — operates at the Level D $\times$ Level G3 intersection. Weaker realism claims operate at broader intersections, with corresponding shifts in empirical content.

What is real at each intersection. The framework commits to realism at four distinct levels of the two-axis hierarchy, with the realism content differing at each level.

At Level G4 $\times$ Level C — the broadest intersection — what is real is the algebra-channel structure of partial-trace observation: the Born rule, channel-level unitarity to leading order, the P-indivisibility of the marginal stochastic process, and the commutant gauge-invariance pattern. This level’s realism is universal across observer-admitting substrata. Every embedded observer in any universality class produces these features; they are structural features of partial-trace observation that no embedded observer can avoid. The realism is about partial-trace observation itself, not about any specific substratum or any specific universality class.

At Level G3 $\times$ Level D — the framework’s distinctive intersection — what is real is the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$. This is the strongest realism claim the framework makes, and Chapter 2 developed the argument for it. The class is uniquely determined by the empirical inputs E1–E7 and the structural assumptions A1–A6; specific representatives within the class are gauge expressions of the same physical content. The realism is class-specific to OI.

At Level $G2 \times$ Level D — a narrower intersection in the gauge axis but the same observation level — what is real is the emergent quantum field theory structure: the Standard Model gauge group, the matter content of three generations (chirality conditional on $H\text{-}\chi$, SM §4.8) plus one Higgs doublet, the anomaly-free hypercharge assignment, the $\theta = 0$ result. This level’s realism is forced uniquely within OI by the cubic-group commutant structure but is shared with any successful SM-reproducing universality class. The realism is about the emergent quantum field theory rather than about any specific substratum.

At Level $G1 \times$ Level D — the narrowest intersection — what is real is the emergent Hamiltonian’s transition-probability content modulo the diagonal-unitary D-gauge. This is the level at which specific numerical predictions for emergent dynamics become well-defined: the level at which the Schrödinger equation operates on specific Hilbert-space vectors with specific phase conventions.

These four realisms are not competing claims about what is fundamentally real. They are realism at different levels of the two-axis hierarchy, with content that is more class-universal at broader levels and more class-specific at deeper levels. The framework commits to all four simultaneously, with the commitments coordinated by the hierarchy’s nesting structure.

Predictions stratified by hierarchy intersection. The framework’s empirical predictions live at specific intersections of the two axes, with the intersection determining both what the prediction’s content is and how it can be tested or falsified.

Level $G4 \times$ Level C predictions — the Born rule, channel-level unitarity, the non-Markovian marginal pattern of P-indivisibility, the commutant restriction pattern — are essentially structural. They are universal across observer-admitting systems, trivially predicted because they are features of partial-trace observation rather than of any specific substratum. They cannot be falsified without contradicting embedded-observer physics itself; falsifying them would refute the entire program of treating observation as a partial-trace operation. The framework predicts them; so does any framework that takes embedded observation seriously.

Level $G2 \times$ Level D predictions — the Standard Model gauge group, the three-generation structure, the anomaly-free hypercharges, the composite Higgs — are forced uniquely within OI by the cubic-group commutant structure but are shared with any framework that reproduces the Standard Model from observer-admitting substrata. Falsifying them would refute the entire class of SM-reproducing universality classes, not just OI specifically. They are central to the framework’s content but do not distinguish OI from successful alternative frameworks.

Level $G3 \times$ Level D predictions are where OI’s distinctive empirical content lives. The Cabibbo angle as $1/(\pi\sqrt{2})$ from a Brillouin-zone distance; the Koide relation as $2/9$ from a cubic-group Casimir; the Bekenstein-Hawking $1/4$ factor from boundary mode counting; the MOND acceleration scale as $cH/6$ from

entropy displacement; the JUNO neutrino-mixing angle as $1/3 - 1/(4\pi^2)$ — all of these are class-specific to OI’s substratum representative. Falsifying them would refute OI’s specific class membership without necessarily refuting the broader universality-class framework. A different universality class might produce the Standard Model with different specific numerical values for these observables.

The framework’s most distinguishing empirical content lives at Level G3 $\times$ Level D — these are the predictions that distinguish OI from other observer-admitting frameworks, including ones that successfully reproduce the Standard Model. The JUNO match is a Level G3 $\times$ Level D result — retrodictive: it agrees with OI’s specific cubic-group calculation, and disagreement would have refuted OI’s representative without refuting the broader observer-admission program. The Bekenstein-Hawking $1/4$ factor sits at Level G2/G3 $\times$ Level D — a derivation with no direct empirical test (GW250114 constrains the classical area theorem, not the coefficient) — since both OI’s substratum and any successful SM-reproducing framework that yields the area law would predict the same factor.

Falsification is stratified accordingly. Falsifying a Level G4 $\times$ Level C prediction is impossible without contradicting the structural features of embedded observation. Falsifying a Level G2 $\times$ Level D prediction would refute the entire SM-reproducing universality class. Falsifying a Level G3 $\times$ Level D prediction would refute OI’s specific representative while leaving the broader universality-class framework intact. The framework’s twenty-two quantitative predictions live primarily at Level G3 $\times$ Level D, which makes them the most informative empirical content the framework offers — but also the most narrowly falsifiable. A pre-registered falsification of a Level G3 $\times$ Level D prediction is a falsification of OI specifically, not of the broader observer-admission framework or of the universality-class structure.

The stratification has methodological consequences developed in Chapter 4. Different kinds of empirical content require different epistemic standards, and the hierarchy makes the stratification precise. Level G4 $\times$ Level C content is structural; Level G2 $\times$ Level D content is shared across an entire universality class; Level G3 $\times$ Level D content is class-specific to OI. The framework’s pre-registration of falsification conditions, developed in Chapter 4, identifies which empirical results would falsify which level of the framework’s commitments.

3.6 Gauge freedom and observational incompleteness

The two axes developed so far invite a question the chapter has not yet answered directly. The gauge axis catalogues what is *not* observable — degrees of freedom that can be transformed without changing anything an embedded observer could record. The framework as a whole is named for *observational incompleteness* — the embedded observer’s inability to resolve the substratum it is part of. These

two notions have been running alongside each other for three chapters. It is worth asking whether they are the same thing.

The answer is that they are the same thing, on a precisely identifiable part of their range — and seeing exactly where they coincide, and where they do not, sharpens both.

The apparent opposition. At first glance gauge freedom and observational incompleteness look like opposites. Gauge freedom is *surplus*: the description carries more structure than the physics, and the gauge quotient throws the excess away, losing nothing. Observational incompleteness sounds like a *deficit*: the observer has *less* than the whole story, missing what the trace-out over the hidden sector destroys. Surplus and deficit do not obviously belong together — one is too much representation, the other too little access.

The opposition dissolves on inspection. Consider why a gauge degree of freedom counts as surplus in the first place. The substratum gauge group $\mathcal{G}_{\text{sub}}$ contains, among its four families, hidden-sector relabeling: permute the labels on the hidden-sector states and every visible observable is unchanged. The labels are surplus — but they are surplus *because no observation distinguishes the relabelled substrata*. The redundancy is not a free-standing fact about the notation; it is constituted by the embedded observer’s inability to tell the relabelled versions apart. “The description has surplus structure” and “the observer cannot access the difference” are not opposed claims. They are one fact, stated once from the side of the representation and once from the side of the observer. Gauge freedom *is* a form of observational incompleteness: it is the incompleteness whose missing piece turns out to be no piece at all.

Two species of incompleteness. This identification cannot be the whole story, though, because not every observational incompleteness is gauge. The framework’s substratum is deterministic; an observer’s own future states are fixed by the bijection φ . There is a definite fact about what the observer will record at the next step. Yet the embedded observer cannot compute that fact — to do so it would need to simulate the forward iteration of φ on the whole of S , and the observer is a proper part of S , not a system large enough to contain a model of the whole. Here the inaccessible thing genuinely *is* a fact: a larger computer, one not embedded as a proper part, could reach it. This incompleteness is not surplus. Nothing is being quotiented away; a real and definite fact is simply out of the finite observer’s reach.

So observational incompleteness comes in two species, and the framework should name them apart:

Gauge incompleteness is the species where the inaccessible “fact” is no fact at all — where there is, demonstrably, nothing of the matter to be had. Hidden-sector relabelings, the alphabet freedom, the deep-sector enlargement: each is a difference that makes no difference to any observer whatsoever, not merely to this one. Gauge incompleteness is identical to gauge freedom. The two terms — “gauge freedom” and “gauge incompleteness” — name one principle

from two directions: the methodology’s term for the structure that can be transformed freely, and the framework’s term for the corresponding blind spot in the embedded observer’s picture.

Computational incompleteness is the species where there is a definite fact, unreachable by this observer for a specifically computational reason: the observer is not a system large enough to compute it. The observer’s own deterministic future is the cleanest case. Computational incompleteness is *not* gauge — the fact is real, and the inaccessibility is a feature of the observer’s finite embedding rather than an absence of content.

The distinction matters because conflating the two would make “incompleteness” incoherent. If every incompleteness were gauge, “incompleteness” would be a misnomer — one cannot be incomplete about what is not there. If every incompleteness were a hidden real fact, the framework could not say, as it does, that the choice of substratum representative within $\mathcal{G}_{\text{sub}}$ has no fact of the matter. The two species are both genuine, and they are different.

Where the identification sits in the hierarchy. Gauge incompleteness is the blind spot associated with the gauge axis: G1 through G4 each define a layer of it. Hidden-sector relabeling is gauge incompleteness at Level G3; the diagonal-unitary D-gauge is gauge incompleteness at Level G1; universality-class equivalence is gauge incompleteness at Level G4. Computational incompleteness is not on the gauge axis at all — it is a limit set by the observer’s finite size relative to S , and it would remain even if the gauge freedom were somehow entirely fixed. The two species are independent in just the way the two axes of this chapter are: one can have either without the other.

A consequence worth stating. The framework’s methodology — developed in Chapter 4 and at length in the companion paper *Physics Modulo Gauge* — holds that physics determines its content only up to gauge, and that this is the complete and correct form of the answer rather than a limitation. The identification of this section says that “physics modulo gauge” and “observational incompleteness” are, on the gauge-incompleteness species, two names for one principle. The methodology is not something brought to the framework from outside; it is the framework’s own central thesis — that an embedded observer’s account of reality is structurally incomplete — read from the methodological side. The incompleteness the framework is named for and the gauge freedom the methodology is built on are the same fact, seen from two directions.

This identification is recent within the framework’s development, and the companion paper flags it as the least-settled of its contributions; it is included here because it follows directly from the two axes this chapter has built, and because it answers the standing question of how the book’s title concept and its gauge structure relate. The reader inclined to treat it cautiously should treat it as a proposal the framework’s structure invites rather than as a load-bearing result the rest of the book depends on — nothing in the empirical chapters rests on it.

3.7 No-go theorems at different descriptive levels

The framework is a hidden-variable theory in the literal sense: there is a deterministic substratum with definite states, and the observer’s stochastic description is obtained by marginalizing over the parts of that substratum they cannot see. The twentieth-century no-go literature — Bell’s theorem (1964), Kochen-Specker (1967), the Pusey-Barrett-Rudolph theorem (2012), the Frauchiger-Renner theorem (2018) — rules out hidden-variable accounts satisfying specific structural conditions. The framework must therefore explain how it evades each of these results, and what it costs to do so.

The framework’s engagement with these theorems is not a single trick repeated, nor four separate tricks. Each theorem identifies a different structural property of would-be hidden-variable accounts, and one architectural feature — the two-level substratum-emergent structure joined by the trace-out — is what answers all of them: for three of the four, the property holds at the emergent level and fails at the substratum level; for Kochen-Specker, as developed below, it holds at neither. The hierarchical structure of the framework is what makes this single engagement possible.

Bell’s theorem. Bell’s theorem rules out local hidden-variable theories satisfying factorizability: $P(a, b \mid x, y, \lambda) = P(a \mid x, \lambda) \cdot P(b \mid y, \lambda)$, where λ is the hidden variable, x, y are detector settings, and a, b are measurement outcomes. Chapter 1 §1.7 developed the framework’s response: parameter independence (the outcome at detector A does not depend on the setting at B) is preserved at the substratum level by the spatial locality of the coupling graph; outcome independence (the outcome at A does not depend on the outcome at B given the settings) is violated at the emergent level by P-indivisibility of joint dynamics. The Jarrett decomposition makes this precise, and Fine’s theorem identifies the resulting class as exactly the quantum-correlation class.

In hierarchy language: Bell’s factorizability is a property at Level $G3 \times$ Level D, where the joint transition matrix is defined. The substratum at Level D satisfies a local-coupling structure that gives parameter independence at the emergent level; the joint transition matrix at the emergent level violates outcome independence through P-indivisibility. The framework’s evasion separates the two conditions Bell combines into one factorizability requirement, satisfying one at the substratum level and violating the other at the emergent level. No fine-tuning of the hidden-variable distribution to experimenter choices is required, and parameter independence — the locality condition that protects against superluminal signaling — holds structurally. The strongest published objection to this decomposition — that keeping parameter independence while conceding outcome independence mistakes a no-signaling condition for locality — is stated at full strength, and answered together with the framework’s explicit concessions, in §3.7.1 below.

Kochen-Specker. The Kochen-Specker theorem rules out non-contextual hidden-variable assignments: there is no function v mapping observables to definite values such that $v(O)$ depends only on O rather than on the measurement context (the set of compatible observables being measured simultaneously). Any hidden-variable theory matching quantum mechanics must be contextual: the hidden variable’s outcome on a given observable depends on which other observables are co-measured.

The framework’s evasion is structural, and it is worth stating precisely, because the Kochen-Specker case does not fit the same “property-at-one-level” template as Bell. A noncontextual valuation is an assignment of definite values to the emergent *observables* that is independent of measurement context. The framework reproduces quantum mechanics (the characterization theorem of Chapter 1), and Kochen-Specker is a theorem about quantum mechanics itself: the observable algebra of quantum theory in dimension three or higher admits no noncontextual valuation. Any account reproducing quantum mechanics therefore reproduces a structure in which no noncontextual valuation of the observables exists. So the property Kochen-Specker forbids does not hold at the emergent level — and it does not hold at the substratum level either, because the substratum carries a definite configuration, not a valuation of the emergent observables at all. The framework’s contextuality is therefore not “noncontextual below, contextual above”; it is forced by the framework’s own equivalence with quantum mechanics, and what the framework supplies is the *account* of it: an emergent observable’s value is produced by the trace-out, the trace-out is taken with respect to a partition, and measuring a set of compatible observables is a specific apparatus coupling — a specific partition. Which observables are jointly defined, and the value each takes, are features of the reduced description fixed by that partition, not partition-independent functions of the substratum configuration. Contextuality, in the framework, is the partition-dependence of emergent valuation.

In hierarchy language: Kochen-Specker non-contextuality concerns the emergent observable structure at Level $G2 \times$ Level D. That structure is contextual — different partitions produce different marginalized statistics — because the framework reproduces quantum mechanics and quantum mechanics has no non-contextual valuation. The Level D substratum carries definite states, but a definite substratum configuration is not a valuation of the emergent observables, so it neither supplies nor contradicts a noncontextual valuation. The framework does not exhibit “non-contextuality below, contextuality above”; it inherits the Kochen-Specker obstruction with quantum mechanics and locates its source in the partition-dependence of the trace-out. This is the one no-go case where the framework’s response is not a property placed at one level and denied at another, but a property unavailable at every level — a difference §3.7’s closing paragraph returns to.

Pusey-Barrett-Rudolph (PBR). The PBR theorem rules out “epistemic” interpretations of the wave function: under reasonable assumptions, the wave

function cannot represent the observer’s incomplete knowledge of an underlying ontic state, because pure states with non-zero overlap would be indistinguishable in any consistent ontic model. The wave function must, in some sense, be ontologically real.

The framework’s evasion identifies the PBR assumption that fails. PBR’s Preparation Independence Postulate requires *unconditional* product independence: the joint ontic-state distribution of two separately prepared systems is the product of the single-system distributions. The framework provides product structure for separated systems — but conditionally, not unconditionally. By the spatial Markov property of the bounded-degree, range-one coupling graph (Chapter 1; Main §3.3), space-like separated subsystems with disjoint coupling neighbourhoods are independent *given boundary data*, and the emergent description factorizes accordingly, but only conditioned on that boundary data, with the systems’ correlation arising from shared boundary history (Main §3.4). Conditional independence given a non-trivial shared variable is precisely not the unconditional product independence the postulate asserts — it is a common-cause structure, and the common cause is co-embedding in one substratum, present at preparation and requiring no interaction between the systems. The factorization assumption PBR needs fails at the substratum level, and the PBR conclusion does not apply. The argument runs entirely on results already established for the framework — the equivalence with quantum mechanics and the spatial-Markov and tensor-product structure of Main §§3.3–3.4 — and does not depend on the P-indivisibility that carries the Bell case, which concerns interacting systems rather than independently prepared ones.

The wave function in the framework is exactly what PBR would call “epistemic”: it is the observer’s compressed description of incomplete data, with no ontic status independent of the observer’s epistemic situation. The PBR theorem does not rule this out because the Preparation Independence Postulate it requires fails for the framework. In hierarchy language: PBR’s preparation independence is a property at Level $D \times \text{Level } G3$; the framework’s substratum at this level supplies only conditional factorization given shared boundary data, not the unconditional product independence the postulate asserts. The trace-out is the mediator — marginalizing the shared boundary data is what carries the emergent conditional factorization into substratum non-factorization. Note that this engagement does not depend on classifying the framework as ψ -epistemic in PBR’s single-level sense: PBR constrains models that are ψ -epistemic *and* postulate-satisfying, and the framework’s failure of the postulate places it outside the theorem’s scope regardless.

Frauchiger-Renner. The Frauchiger-Renner theorem identifies a consistency problem in quantum mechanics when observers are themselves modeled as quantum systems. An observer-of-an-observer scenario produces logically inconsistent predictions if one assumes single-world consistency, the universal applicability of quantum mechanics, and the freedom of observers to apply quantum mechanics to other observers.

The framework’s response is that the Frauchiger-Renner assumptions place observers at incommensurable levels of the observation hierarchy. An observer applying quantum mechanics to themselves is using quantum mechanics — a Level $D \times$ Level G2 emergent description — to describe a substratum that is itself at Level $D \times$ Level G3. The two descriptions are not equivalent; the substratum-level description is more fundamental than the emergent description, and applying the emergent description to the substratum produces inconsistencies because the emergent description is a marginalization that loses information. The framework predicts that Frauchiger-Renner-type paradoxes will continue to arise in attempts to apply quantum mechanics universally, including to systems containing other observers — they are consequences of the framework’s hierarchical structure, not anomalies requiring resolution.

In hierarchy language: Frauchiger-Renner’s universal applicability is a claim at Level $G2 \times$ Level D , and the framework violates it not by restricting quantum mechanics to a subset of systems but by recognizing that observers themselves are substratum-level objects that emergent-level quantum mechanics describes only marginally. The hierarchical structure makes the violation systematic rather than ad hoc.

The common pattern. The framework’s engagements with Bell, Kochen-Specker, PBR, and Frauchiger-Renner share a common source — but the shared factor is the architecture itself, the two-level substratum-emergent structure joined by the trace-out, not a single slogan applied four times. For three of the theorems the pattern is uniform: the assumed property — Bell’s factorizability, PBR’s preparation independence, Frauchiger-Renner’s universal description — holds at the emergent level, where it is unobjectionable because that level simply is quantum mechanics, and fails at the substratum level, where the theorem needed it. For Kochen-Specker the pattern is different: the assumed property, a noncontextual valuation of the observables, holds at *neither* level, because the framework reproduces quantum mechanics and quantum mechanics has no such valuation. What unifies the four is not “satisfied below, violated above” but the trace-out: by interposing a partition-dependent marginalization between substratum and emergent observable, it is what makes the noncontextual valuation unavailable in the Kochen-Specker case and what places the property at the emergent level only in the other three. One architecture; four manifestations of working within it. The engagements are not tricks: they are structural consequences of the framework’s two-axis hierarchy, fixed before any no-go theorem is considered.

Chapter 4 develops this as a methodological principle. The framework’s response to the no-go literature is one architectural feature — the two-level substratum-emergent structure joined by the trace-out — fixed before any theorem is considered; what is done per theorem is the *verification* that the architecture answers it. That verification: locate where the theorem’s assumption meets the two-level structure; confirm that the framework reproduces quantum mechanics’ predictions in the regime the theorem covers; and check that the framework produces

no signature the theorem’s empirical content rules out (e.g., superluminal signaling for Bell). For three of the four theorems the assumption holds at the emergent level and fails at the substratum; for Kochen-Specker it holds at neither, contextuality being forced by the framework’s equivalence with quantum mechanics. The architecture-level engagement, verified theorem by theorem, is the framework’s general methodological response to the no-go literature.

3.7.1 The Maudlin test

The sharpest sustained critique of hidden-variable programs in the philosophy of physics literature — associated above all with Tim Maudlin’s *Quantum Non-Locality and Relativity* and his subsequent essays and lectures — holds that the lesson of EPR and Bell is not interpretive but physical: the world is nonlocal, theories should say so plainly, and the standard devices for avoiding saying so (operational readings of locality, appeals to no-signaling, denials of counterfactual definiteness, decompositions of Bell’s premise into parts) are evasions. This subsection states that critique in its own terms, as forcefully as its proponents state it, and runs the framework through it. The framework passes some of the test by conceding, some by deriving, and one part by openly recording a dispute it cannot settle from inside. Nothing in this subsection weakens a claim made elsewhere; its purpose is to make explicit which claims were never made.

The pinhole. Einstein’s 1927 Solvay argument (recorded and reconstructed in Bacciagaluppi and Valentini 2009) needs only one particle. A wave passes through a pinhole and spreads toward a hemispherical detection shell; exactly one scintillation occurs. If the wave function is the complete physical state, then at detection something physical happens everywhere on the shell at once — the amplitude at every other point is extinguished by the flash at one — and that is real nonlocal dynamics, whatever one calls it. If the wave function is not complete, the flash merely tells us where the particle was, and the “collapse” elsewhere is updating, not dynamics. The dilemma is exhaustive, and the framework takes the second horn without embarrassment: the substratum configuration is definite throughout; the wave function is the embedded observer’s compressed description of what they cannot resolve (Chapter 1 §1.9). What distinguishes this from the naive statistical ensembles that fail — and Einstein’s own hope did fail, as the next paragraphs record — is that the observer’s incompleteness is *structured*: the marginal process is P-indivisible, and the extra correlations that a mere ignorance-interpretation cannot supply are carried in the hidden sector and returned by the dynamics (Chapter 1 §1.5). At the shell, the detection is a division event: a physical intervention at one location that severs a substratum correlation there (Chapter 1 §1.11). The simultaneous “emptying” of the wave function elsewhere is conditional-probability bookkeeping — the observer re-marginalizing over the hidden sector given a new visible fact (the Lüders update as Bayesian conditioning, Main §3.4). One event is physics; the other is arithmetic. The framework can say which is which because it has a substratum to say it about.

Where the framework sits in the locality taxonomy. Maudlin’s own analytical vocabulary distinguishes the locality of a theory’s *ontology* (are the beables local — do they live at points and regions of ordinary space?) from the locality of its *dynamics* (do the beables evolve by local law?). In that vocabulary the framework’s position can be stated exactly. Its ontology is local: the beables are lattice configurations in three spatial dimensions — there is no wave function in the ontology at all, and so no object living on configuration space (on this point the framework simply agrees with the configuration-space objection to wave-function realism; the wave function’s high dimensionality is the signature of correlation bookkeeping, not of an arena). Its substratum dynamics is local: range-1 coupling on a bounded-degree graph, strictly nearest-neighbor (Main, Lemma 1). What is not local is the *partition*: the observer’s visible sector is bounded by the cosmological horizon, the hidden sector is the horizon complement, and the emergent joint description of two visible subsystems that interacted at preparation carries a transition matrix that does not factor, $T_{QR} \neq T_Q \otimes T_R$. Local ontology, local dynamics, non-local partition. Everything the framework concedes and everything it contests below follows from which of these three the concession or contest is about.

EPR, read correctly. The EPR argument (Einstein, Podolsky, and Rosen 1935) is routinely misremembered as an appeal to determinism. It is the reverse: from locality and the perfect correlations, determinism is *derived*. If nothing at one wing can affect the other, then the ability to predict the distant outcome with certainty entails that the outcome was already fixed — the hidden determiner is the conclusion of the argument, not its premise. The framework accepts this reading without reservation, and accepts what follows from accepting it: the framework is precisely the kind of theory EPR’s argument called for, a deterministic completion in which measurement reveals antecedently definite substratum facts. That is why Bell’s theorem binds on the framework with full force. A completion of the EPR type must, by Bell’s mathematics, violate factorizability to reproduce the observed correlations — there is no third way, and the framework does not claim one.

What the framework does not argue. Two standard escapes are explicitly declined. It does not deny counterfactual definiteness: substratum outcomes are definite, and the framework nowhere trades on refusing to consider unperformed measurements — the response to Bell runs through the structure of the premises, not through denying that measurements have results. And it is not superdeterministic: the correlations are not explained by tuning the hidden-variable distribution to the experimenters’ choices. The measured chaotic mixing of the substratum actively destroys residual correlations between pair-preparation and setting-generating degrees of freedom, so operational measurement independence holds dynamically; the correlations that a superdeterministic explanation would need do no work here (Main §3.3).

Bell’s theorem, conceded and then confronted. The mathematics is conceded in full: Bell’s factorizability fails for the emergent process, and the

framework is therefore not a local hidden-variable theory in Bell’s sense. The framework’s presentation elsewhere (Chapter 1 §1.7; §3.7 above) uses the Jarrett decomposition — parameter independence preserved, outcome independence violated, Fine’s theorem locating the result in exactly the quantum class — and the critique’s response to that presentation must be faced squarely, because it is on the record and it is forceful: Bell himself formalized local causality as a single screening-off condition, precisely to *avoid* resting anything on operational notions; splitting his condition into a “safe” signaling part and a “violated” outcome part, and then keeping only the first, is on this view not an analysis but an evasion — no-signaling is far too weak to count as locality, and outcome dependence between spacelike-separated events simply *is* nonlocality. The framework’s answer is not to dispute the verdict at the level where it is rendered. Outcome dependence at the emergent level is real, physical, and underived from any local stochastic law of the emergent theory; the framework does not claim to restore Bell locality (Main §3.3 states this in terms). What the framework adds — its actual wager — is one level down, and it is a claim of physics, not of vocabulary: *the two wings’ hidden sectors are not two systems*. Both wings’ inaccessible degrees of freedom are the same horizon bath; the conditional structure of the substratum supplies independence of separated regions only given shared boundary data (the spatial Markov property), never unconditionally. Emergent spacelike separation is a fact about the coarse-grained description and does not descend to substratum independence. Bell’s factorization premise, applied here, imports the emergent notion of separated systems into a regime where the construction says it fails — not by conspiracy and not by signal, but because “two systems” was never the right count of the hidden variables. (The idea that a common background can entangle two particles has a physics lineage of its own — stochastic electrodynamics derived bipartite entanglement from coupling to the shared zero-point field (Valdés-Hernández, de la Peña, and Cetto 2011; de la Peña, Cetto, and Valdés-Hernández 2014) — though there the background is a postulated field within spacetime, while here the shared sector is the derived complement of the observer’s partition, prior to the emergent causal structure.) The violated correlations are then *derived*: computed from the construction, capped at exactly the Tsirelson bound as a theorem rather than a fit, with no-signaling holding structurally rather than by fine-tuning.

The rejoinder, at full strength. The critique has an answer to this too, and it must be stated before it is graded. First: a mechanism underneath non-local correlations does not make them local — this is “nonlocality with extra steps,” and honesty requires calling the resulting theory nonlocal. Second, and sharper: a *common cause* that produces Bell-violating correlations is exactly what Bell’s theorem excludes, for any cause located in the wings’ shared past light cone; if the shared bath is doing common-cause work, its alleged exemption from Bell’s screening condition must come from its atemporal, global character actually doing physics, not from renaming. The framework’s position on each, stated plainly. On the first: at the emergent level the framework *agrees* —

the emergent process is nonlocal in Bell’s sense, and the framework says so. What it maintains is that the ontology and the dynamics are local in the taxonomic senses above, and that the interesting question is therefore not whether to utter the word “nonlocal” but where the nonlocality lives; here it lives in the partition — the same structural feature that generates quantum mechanics itself. Whether the substratum-level notion the framework satisfies — causal locality in the directed-conditional-probability sense developed by Barandes for unistochastic reformulations (arXiv:2402.16935) — is a physically principled locality criterion or a weakened surrogate is a live dispute in the current literature between exactly these positions. The framework inherits that dispute and does not pretend to settle it; it records which side of it the construction occupies. On the second: the bath is not an event-structure in the wings’ past light cone, because the emergent causal structure — light cones included — is itself an output of the trace-out, and the co-embedding relation between the wings and the bath is prior to it; Bell’s screening condition quantifies over beables in spacetime regions and has no slot for the structure that produces the regions. That this is a physical claim rather than a verbal one rests on three checkable facts already in the framework’s ledger: the correlations are computed, not posited; their maximum is derived ($2\sqrt{2}$ exactly); and the no-signaling protection is structural. But the classification of the bath as “not a past-light-cone common cause” turns on treating emergent causal structure as non-fundamental — which is the framework’s central commitment and precisely what a critic of this school declines to grant. The framework can survive the verdict “a nonlocal theory with a local ontology”; it does not claim to refute it.

The bill, paid in the open. A theory that violates outcome independence and keeps a substratum owes an account of relativity, and the account cannot be evasive: the framework’s substratum has a preferred frame — the frame in which update steps are simultaneous — and on identification with the observed universe this is operationally the cosmic rest frame, an identification the framework records rather than derives (SM §3.1). Lorentz invariance is therefore wholly emergent, and the framework’s obligations are quantitative: no linear violation (established at tree level), rotational protection by octahedral symmetry to all orders conditional on the emergent theory preserving that symmetry at loop level, boost-sector violation controlled by a derived decoupling law in the cross-charged mass M_X with a two-sided falsifiable window. These are not side technicalities; they are the price of the preferred-frame answer to the foliation objection, and the open item among them — the interacting (condensate-dressed) extension of the coarse-graining characterization (SM §3.1; Chapter 19 §19.3.7) — is correctly understood as Bell-defense infrastructure. A reader who wants to falsify the framework’s account of Bell should attack there: a demonstrated failure of emergent Lorentz protection would leave the preferred frame exposed at accessible energies, and the framework’s answer to this subsection would fail with it.

The characteristic trait. Schrödinger called entanglement *the* characteristic trait of quantum mechanics, the one forcing its entire departure from classical

lines. The framework agrees, with a precise gloss: entanglement is the emergent description’s encoding of correlations written into the shared bath at preparation and not yet read back — shared-history bookkeeping, real in its consequences and exactly quantum in its statistics, with nothing at the substratum level that is entangled. The trait is characteristic because the partition is characteristic: it is what an embedded observer’s account of a shared hidden sector must look like.

The ledger. Conceded: Bell’s mathematics; EPR’s logic and its corollary that Bell binds; the failure of the counterfactual-definiteness escape; the reality of emergent outcome dependence; the preferred frame; the configuration-space objection. Contested: that the nonlocality is dynamical rather than partitional; that the shared-bath common cause is a renaming rather than a derivation. Open, and recorded as such: whether substratum-level causal locality is locality enough — a question the framework inherits from a live dispute and cannot close by internal argument — and the loop-level protection of emergent Lorentz invariance, which is the answer’s quantitative hostage. The corrected reading of EPR that this subsection rests on can be put in one sentence: Einstein was only half right — God doesn’t play dice, but he does play cards. The cards are not a loose figure of speech: Chapter 1 §1.11 builds the game — a definite deal held hidden, a deterministic shuffle that writes visible history into the burn pile and reads it back, revelation as an intervention with a cost — and states its boundary as plainly as its successes: no arrangement of cards on a table violates a Bell inequality. That boundary is not an embarrassment to the metaphor here; it is the diagnosis. A card game is local through and through — each player’s hidden cards are their own — and single-system interference is exactly as far as such a game can go, which is to say exactly as far as a local hidden-variable model can go. What the world adds, and no table can, is the wager of this subsection: the two wings do not hold separate burn piles. Their hidden sectors are one shared horizon bath, and that is where the Bell correlations live — in the full construction, not in the toy.

3.8 The philosophical lineage

The framework’s structural-realism position has antecedents older than quantum mechanics and broader than the contemporary observer-dependence program. Articulating the relationship between the framework and its philosophical lineage is part of locating the framework’s distinctive contribution.

Nicholas of Cusa. The fifteenth-century cardinal and philosopher developed two doctrines that anticipate the framework’s structural claim. *Docta ignorantia* — learned ignorance — held that the highest knowledge of God is knowing what we cannot know: the human mind, bounded and finite, can approach the infinite only through the recognition of its own limits. The structural form of this claim — that incompleteness is a feature rather than a deficiency — anticipates the framework’s reframing of the cosmological constant problem: the 10^{122} discrepancy is not evidence that physics is incomplete in a way needing

repair but the structural signature of how observation operates in a universe an observer is embedded in.

Coincidentia oppositorum — coincidence of opposites — held that contradictions resolve at the level of the infinite: properties that appear incompatible to finite intellects are unified in the divine, where the polygon and the circle are the same thing, where the maximum and the minimum coincide. The structural form of this claim — that two seemingly irreconcilable descriptions can both be correct when both are projections of a common underlying reality — anticipates the framework’s claim that quantum mechanics and general relativity are not incompatible but are projections of the same finite deterministic substratum along different axes. The conventional formulation treats the two as competing theories requiring reconciliation; Cusa’s framework would suggest that the contradiction is at the level of the projections rather than the substratum, and the framework formalizes this suggestion.

Cusa is not the framework’s source. The framework derives its claims from explicit mathematical theorems, not from medieval theology. But the structural resemblance is striking, and identifying the resemblance helps locate what is distinctive about the framework’s claims and what is part of an older tradition of thinking about observation, infinity, and structure.

Spinoza. Spinoza’s metaphysics is the deepest historical parallel to the framework’s substratum-and-projections picture. The *Ethics* posits one substance — God or Nature — expressed through infinitely many attributes, each constituting the substance’s essence under a different aspect. The two attributes humans can know are extension (the physical aspect) and thought (the mental aspect), but these are not two substances; they are the same substance expressed through different attributes, with no causal interaction between them because there is only one thing underneath. The doctrine that the order and connection of ideas is the same as the order and connection of things follows from this single-substance structure: the parallelism between mind and body is not an interaction but a consequence of their being the same underlying reality.

The structural form of Spinoza’s account — one substance, multiple incommensurable attributes, with the apparent multiplicity being a feature of how the substance is expressed rather than of the substance itself — maps directly onto the framework’s three-projection structure. The substratum (S, φ) is the single substance; quantum mechanics, general relativity, and the arrow of time are projections expressing the substratum’s content under different aspects; the appearance of three distinct emergent theories is a feature of the projections rather than of the substratum. The framework’s commitment to ontic structural realism on the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ is Spinozistic in form: what is real is the substance modulo its gauge expressions, not any specific expression of it.

The parallel is not complete. Spinoza’s substance is unique (there is exactly one substance, by argument); the framework’s substratum is a member of an equivalence class (there are many specific bijections, gauge-equivalent under

$\mathcal{G}_{\text{sub}}$). Spinoza’s attributes are infinite in number; the framework identifies three projections (quantum mechanics, general relativity, arrow of time) without claiming to enumerate all possible projections. The framework’s claims are mathematical theorems rather than metaphysical doctrines. But the structural resemblance is closer than to any other position in the historical philosophical literature.

Worrall and epistemic structural realism. The closest twentieth-century antecedent to the framework’s position is John Worrall’s epistemic structural realism, developed in the 1980s and 1990s. Worrall argued that scientific theories should be interpreted as committing to the structure of relations between observable phenomena rather than to the intrinsic natures of the underlying entities producing those phenomena. The argument was epistemic: structure is what survives theory change across the history of science (the equations of Fresnel’s wave theory survived the transition from ether to electromagnetic field), so structure is what scientific theories warrant belief in. The framework’s structural-realism position is ontic rather than epistemic — what is real is the equivalence class, not just what we are warranted in believing about it — but the structural commitment is in the same family. Both Worrall and the framework hold that structure rather than substance is what survives the appropriate quotient.

The framework’s distinctive contribution relative to Worrall is the precise specification of what the appropriate quotient is. Worrall’s structural realism left the structure under-specified: which structures are real, which are gauge, how to determine the boundary. The framework’s $\mathcal{G}_{\text{sub}}$ characterizes the substratum-level gauge explicitly, with four generator families and a completeness theorem. The framework’s universality-class structure at Level C characterizes broader equivalences across structural classes. The framework’s hierarchical structure stratifies the realism claims by level of generality. Worrall’s program identified structure as the appropriate object of realism; the framework specifies the structure.

Hofstadter and strange loops. Douglas Hofstadter’s *Gödel, Escher, Bach* (1979) and *I Am a Strange Loop* (2007) developed the idea that self-reference under finite resources produces emergent phenomena (consciousness, in Hofstadter’s primary application) that are real features of the self-referential system without being features of the substrate. The framework’s incompleteness-family argument — Gödel for formal systems, Turing for computational systems, the framework for embedded observers — agrees with Hofstadter’s central claim that self-reference produces structural emergence.

The framework disagrees with Hofstadter on a specific point. Hofstadter held that the substrate is explanatorily inert: consciousness emerges from the strange loop regardless of what the loop is implemented on (neurons, silicon, abstract automata), and the substrate adds nothing beyond providing the implementation. The framework’s substrate is not inert. The 10^{122} ratio between substratum capacity and visible-sector capacity is the cosmological constant discrepancy; the dark sector is the gravitational effect of the substrate’s content that the observer

cannot read; the framework’s twenty-two quantitative predictions follow from specific structural features of the substratum that different substrates would produce differently. The framework agrees with Hofstadter that the strange loop is real and that self-reference produces structural emergence; it disagrees that the substrate is therefore irrelevant.

The embedded-observer limits tradition. The Introduction placed the framework’s central claim in the impossibility-with-structure family — Gödel for formal systems, Turing for computational systems, Wolpert for physical inference devices — and stated the relationship in a sentence this section need not repeat: the logical results corroborate the framework’s derivation but do not supply it. (Wolpert’s theorems are engaged separately, as an objection, in Appendix C.7.) That family has further members, closer to the framework’s specific concern, that belong in the lineage record. Dalla Chiara (1977) analyzed the quantum measurement problem as a self-reference problem in the logician’s sense, decades before that framing was common. Rössler’s endophysics (1998) distinguished systematically between the view from inside a world and the view from a hypothetical outside, arguing that physics as practiced is unavoidably interface physics. Breuer (1995) proved the sharpest single result: no observer can perform an accurate measurement of the complete state of any system in which that observer is contained — an impossibility theorem, not an interpretation. And Szangolies (2018) derived quantum-like epistemic restrictions from self-reference structure via Lawvere’s fixed-point theorem, the nearest approach in this tradition to extracting quantum form rather than mere limitation. The framework cites all of these as load-bearing lineage, and its relationship to them is the one the Introduction stated for the family as a whole and §3.7.1 states for EPR: the tradition establishes that the limit exists; the framework derives what the limit forces — a P-indivisible unistochastic marginal, which is to say, exactly quantum mechanics.

Contemporary observer-dependence programs. The framework converges with the recent technical work on observer-incorporated quantum gravity — Chandrasekaran-Longo-Penington-Witten on de Sitter entropy with observer (2022), Maldacena on observer’s-clock dependence in cosmological partition functions (2024), Harlow-Usatyuk-Zhao on observer-dependent Hilbert spaces in closed universes (2025), Slagle-Preskill on quantum gravity from observer-bounded information (2024) — on the foundational claim that observation matters and that the algebra of observables depends on the observer’s embedding. The technical apparatus these programs employ differs from the framework’s, and the empirical content they produce is different (mostly structural commitments rather than specific numerical predictions), but the foundational commitment is shared.

Hawking’s turtles and the regress of support. Hawking opens *A Brief History of Time* (1988) with the anecdote of the listener who insists the world rests on a giant turtle and, pressed on what the turtle stands on, answers that it is turtles all the way down. The story lampoons the oldest structural problem

in metaphysics: every proposed support invites the same question one level lower, and the classical exits — an infinite tower, a circle, or a foundation that simply stands — each pay a familiar price. The framework’s own picture makes the regress feel urgent again in modern dress. If the universe is completely described by a lossless finite-state memory with finite read-write access, what is the hardware the memory runs on?

The framework takes none of the three classical exits, because it can prove that the question loses reference before it needs an answer. The reconstruction determines the substratum only up to the gauge group $\mathcal{G}_{\text{sub}}$ (Chapter 2 §2.5), and one of its four generator families is deep-sector enlargement: add unobservable substrate structure beneath the boundary layer and every observable — every transition probability, the emergent Hamiltonian, $\hbar$ itself — is unchanged. So each candidate answer to the hardware question is one of two things. Either it is more structure, in which case it has been absorbed into (S, φ) and the question repeats; or it is deep-sector content, in which case §3.6’s identification applies in its strong species: gauge incompleteness, where the inaccessible “fact” is no fact at all. The tower is not terminated by a special turtle that stands on nothing; below the structural level, the sortal *turtle* stops being well-formed. Asking what the structure runs on becomes like asking which coordinate chart spacetime is really in — a question about representation mistaken for a question about the world. The same accounting absorbs the regress’s modern costume, the simulation hypothesis: a simulated (S, φ) and a native one differ only by deep-sector structure invisible to every embedded observer, so on the framework’s own gauge bookkeeping there is no physics question there to answer.

The dissolution must be scoped, though, and the scoping is this chapter’s own two-species distinction (§3.6). Not every turtle below us is gauge. The first one down — the hidden sector itself, the horizon bath — is real on the framework’s account, and *computational* incompleteness, not gauge incompleteness, is what hides it: there are definite facts about it, unreachable from our position because the observer is not a system large enough to hold them. We cannot see all the turtles because we are turtles ourselves — components of the same stack, performing the trace-out on our own tower. Unobservable there does not mean nonexistent; it means unverifiable by direct experiment from inside. And deduction still reaches where observation cannot: the framework characterizes that first turtle in detail — its capacity of order 10^{122} modes, its Hubble-scale timescale, the four-condition structure of its coupling — and, in the flagship case, its energetic content. The boundary entropy has no operator in the emergent quantum field theory, so roughly ninety-five percent of the critical density is invisible to every embedded observer *by theorem*; its uniform component is dark energy, and its scale-dependent face is the dark-matter regularity $a_0 = cH/6$ (Chapter 7 §§7.8–7.9). Those are deduced properties of what cannot be directly seen, and they face indirect test through the phenomenology they force — which is exactly how a real but unobservable turtle differs from a gauge one. Only past that first real turtle, in the deep sector where enlargement is pure gauge, does *below* stop referring. The regress does not end at us, and it does not end

at the horizon; it ends where the framework can prove there is nothing left for it to be about.

There is also an inversion of orientation worth making explicit. Hawking’s tower points downward — each level supported by the one beneath, and his companion question, what breathes fire into the equations, asks for the bottommost support. The framework’s ground is at the top. Its first axiom is the thin, indubitable fact of Chapter 1: there is differentiated content, some this-rather-than-that — the one item that needs no further turtle, because doubting it is already an instance of it. Everything else is reconstruction *downward* from that surface fact, and the reconstruction provably exhausts its content exactly where the gauge class begins. The fire-breathing question presupposed equations first and existence second; the framework runs observation first and structure second, which is why it never meets the regress in Hawking’s orientation at all.

The honest accounting — the same honesty the posit ledger exists to enforce — is that the *justificatory* regress is cut rather than dissolved. One can still ask why recurrence, why finiteness, and there the answer is not a theorem but a ledger entry: the recurrence axiom is owned as a posit with an explicit non-derivability countermodel, finiteness carries its stated two-part status, and the one load-bearing external import is named ([Main §3.1, §4.5]). The framework’s position on turtles is therefore exactly this: one turtle that cannot coherently be doubted, a countable and audited handful it declares it is standing on, theorems for everything in between, and a proof that below a certain depth the word *below* stops referring. That is a better epistemic position than any horn of the classical trilemma — and it is bought at the price the ledger displays rather than hides.

What distinguishes the framework is the approach: starting from a finite deterministic substratum and deriving quantum mechanics as a theorem under specified conditions, rather than adding observer degrees of freedom to an existing quantum-gravitational formalism. The contemporary programs work within quantum gravity and identify the observer’s role; the framework works underneath quantum mechanics and identifies how quantum mechanics arises from observer-embeddedness. Both projects are pursuing structural realism about observation; they are doing so from different directions.

3.9 What follows

The hierarchical structure developed in this chapter — the observation axis A through D, the gauge axis G1 through G4, the orthogonality between them, the stratified predictions at the various intersections, the identification of gauge freedom with observational incompleteness on the gauge-incompleteness species, the engagement with the no-go theorems through one architectural feature, the philosophical lineage tracing back through Spinoza and Cusa — provides the framework’s full structural-realism position. Chapter 2’s commitment to $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ is now located at a specific intersection (Level D $\times$ Level G3) within a broader structure, with weaker realism claims at broader intersections

and stronger claims unavailable.

Chapter 4 develops the framework’s methodology: the single-foundational-commitment derivation pattern, the architecture-level no-go engagement as a systematic methodological principle, the classification of claims by derivational rigor, the framework’s pre-registration of falsification conditions. The methodology chapter takes the structural realism developed here as input and articulates the methodological discipline the framework’s program follows.

Chapter 8 develops the universality-class structure and the framework’s relationship to alternative unification programs — string theory, loop quantum gravity, causal sets, asymptotic safety — at the Level C $\times$ Level G4 intersection. The OI-string relationship in particular receives extended treatment there, with the four-feature audit (Born rule, channel-level unitarity, P-indivisibility, commutant gauge-invariance) characterizing the cross-framework universality-class membership and identifying what transports between the two programs and what does not. The negative technical result on substratum-level equivalence (that string-theoretic matrix models fail OI’s A2 and A5) and the positive result on universality-class equivalence are both developed there.

Parts II–IV develop the framework’s empirical content at the various levels identified in this chapter. The Level G3 $\times$ Level D predictions — the twenty-two Standard Model retrodictions, the JUNO neutrino-mixing angle, the Bekenstein-Hawking 1/4 factor, the MOND acceleration scale, the dark-sector predictions — are developed quantitatively. The Level G2 $\times$ Level D structural commitments — the Standard Model gauge group, the three-generation structure, the anomaly-free hypercharges — are developed as derivations from the cubic-group commutant. The Level G4 $\times$ Level C universal features — the Born rule, channel unitarity, P-indivisibility, the commutant restriction pattern — are taken as given content of the framework’s foundational theorems.

The framework’s hierarchical structure is what makes the framework different from a single bold claim about reality. It is a coordinated set of claims at different levels of generality, with different empirical predictions, different falsifiability conditions, and different philosophical commitments at each level. The framework’s strongest realism — at Level D $\times$ Level G3 — is also its narrowest. The framework’s broadest realism — at Level C $\times$ Level G4 — is also its most universal. The book’s role is to develop these levels with the technical and empirical content that distinguishes one level from another.

Chapter 4

Methodology

4.1 What this chapter develops

The framework’s content, developed in Chapters 1 through 3, rests on a specific methodological discipline. The discipline is not the framework’s content; it is how the framework’s content is produced, audited, and made available for evaluation. Different methodological choices would produce different frameworks even from the same starting commitments. The choices made here are worth articulating explicitly, because they distinguish the framework from other unification programs that begin from different methodological commitments and reach different conclusions.

Four methodological commitments organize the framework’s program. First, the framework operates from *a single foundational commitment* — the starting point that observation occurs — and derives the rest as theorems. That commitment is itself structured: examined closely (Chapter 1), it resolves into two axioms — that tokened differentiation occurs, and that it recurs — thematically one principle but, in honest count, two posits, one indubitable and one substantive. “Single foundational commitment” names the discipline of taking only this one starting point from outside the theorems, not a claim that the starting point is a single indivisible fact. Other foundational physics programs commit to multiple primitives (string theory commits to extended one-dimensional objects, supersymmetry, and a specific dimensional structure; loop quantum gravity commits to a graph-theoretic discrete geometry and a specific quantization procedure). The framework’s discipline is to commit to as little as possible at the foundational level and derive as much as possible from the minimal commitment. The methodological cost is steeper derivations; the methodological benefit is fewer adjustable parameters and a cleaner audit trail.

Second, the framework *classifies its claims by derivational rigor* using an explicit scheme. Strictly parameter-free structural predictions are labeled S; conditional structural predictions that depend on a stated open assumption are C; layered conditional predictions awaiting closure are L; retrodictions fitted to observation are R; predictions taking a phenomenological input are P; predictions inheriting from a mass-chain input are M; predictions explicitly using an empirical input are E. Every prediction the framework makes carries an explicit classification, with the strongest claims reserved for S-class derivations and weaker claims honestly labeled. The discipline of this classification — refusing to overstate rigor — is part of what makes the framework’s empirical record evaluable.

Third, the framework *engages the no-go theorems through one architectural feature rather than tactically*. Each major no-go theorem in the foundations literature (Bell, Kochen-Specker, PBR, Frauchiger-Renner) identifies a structural property of would-be hidden-variable theories that fails for quantum mechanics. The framework’s engagement with each is not a separate trick but a manifestation of one structure: the two-level substratum-emergent architecture joined by the trace-out. For three of the four theorems the assumed property holds at the emergent level and fails at the substratum level; for Kochen-Specker it

holds at neither, the framework’s contextuality being forced by its equivalence with quantum mechanics. Chapter 3 §3.7 developed this in technical detail; this chapter develops it as a methodological principle — one architecture, with the match to each theorem verified case by case.

Fourth, the framework *pre-registers its falsification conditions*. A framework that derives rather than posits must be willing to specify what would invalidate the derivation. The pre-registration distinguishes structural commitments — which the framework stands or falls on — from solution-specific properties that lie outside its scope. The Class A conditions identify structural theorems whose violation would falsify the framework; Class B identifies parameter-free retrodictions whose precision-upgrade violation would falsify; Class C identifies broader framework-level commitments. The discipline is to specify these conditions before the relevant observations are made, not after.

The chapter has three pieces of work. §4.2 develops the single-foundational-commitment pattern, including the derivation tree from the starting commitment through the framework’s theorems and predictions. §4.3 develops the claim-classification scheme, including the S/C/L/R/P/M/E categories and the four-layer framing that locates predictions within the substratum-to-emergent stack. §§4.4–4.6 then develop the no-go evasion as a methodological principle, the pre-registered falsification conditions across Classes A, B, and C, and the boundaries of what does not falsify the framework. §4.7 closes the chapter with a forward pointer to Parts II–IV, where the methodology is applied to the framework’s specific empirical content.

4.2 The single-foundational-commitment pattern

The framework begins from one premise. Observation occurs: an observer records distinguishable outcomes of interactions with a system not wholly under the observer’s control. This is the only commitment the framework takes from outside its own theorems. Every other content — the lemmas of finiteness and causal partition and determinism, the conditions C1-C3 for embedded observers, the central theorem that quantum mechanics emerges under those conditions, the reconstruction theorem from observed physics back to the substratum, the Standard Model derivation from the cubic-group decomposition, the gravitational sector from horizon thermodynamics — is derived from the foundational premise plus its structural consequences.

The methodological discipline of single-foundational-commitment derivation operates as a constraint at every step of the framework’s construction. When a new structural feature is needed, the framework must derive it from existing content rather than add it as a new premise. When a new prediction is sought, the framework must trace its derivation back to the foundational commitment through some chain of theorems and identifications. When a new structural assumption appears necessary, the discipline forces an explicit accounting: either the assumption is itself derivable from the foundational premise (in which case

it is a theorem, not an assumption), or it is a new commitment that needs to be acknowledged as such (in which case the framework’s foundational economy has been compromised).

The framework’s content satisfies this discipline imperfectly but explicitly. The structural assumptions A1–A6 identified in Chapter 2’s reconstruction theorem (finiteness, determinism, bounded coupling, center independence, linearity, background independence) are not derivable from the foundational premise; they are restrictions on the class of substrata the reconstruction considers. The discipline requires acknowledging this honestly: the framework’s uniqueness is *within* the class defined by A1–A6, not absolute. Alternatives outside this class would require separate derivations. The framework’s defense of A1–A6 is an argument from the empirical inputs of physics, not a derivation from the foundational premise — A1 is supported by the holographic entropy bound, A2 by the structural requirement that observer records not be erasable, A3–A6 by the empirical content of the observed gauge structure. The honest acknowledgement is that the framework rests on the foundational premise *plus* the structural assumptions, with the assumptions identified explicitly so that any reader can evaluate them.

The benefit of the discipline is a clean audit trail. Every prediction the framework makes can be traced back through the chain of theorems to either the foundational premise or to one of the explicitly-acknowledged structural assumptions. The Cabibbo angle $1/(\pi\sqrt{2})$ traces back through the cubic-group decomposition (Chapter 5), to the cubic-symmetric coupling matrix (the substratum’s wave equation), to the wave equation’s uniqueness given A4–A6 (Chapter 2), to the underlying bijection structure (Chapter 1). The Bekenstein-Hawking $1/4$ coefficient traces through boundary mode counting (Chapter 6), to the thermal-self-consistency gap equation, to the substratum’s bijection structure (Chapter 2), to the embedded-observer construction (Chapter 1). The chain of derivation is explicit at each step, and any link in the chain that turns out to be incorrect propagates predictably to the downstream predictions.

The cost of the discipline is that some natural extensions of the framework require derivational work that has not yet been done. The framework currently treats CP-violating phases in CKM and PMNS as solution-specific inputs rather than as derivations; the cosmological initial conditions are taken as given rather than derived; the specific values of fermion masses depend on one empirical input (the strange-quark mass) rather than being derived from the substratum alone. These are not failures of the framework but acknowledged scope limits. The methodological discipline distinguishes between “the framework predicts X” (X follows from the foundational premise plus the structural assumptions) and “the framework is consistent with X but does not predict it” (X is solution-specific or depends on inputs the framework has not yet derived).

A useful diagnostic for the discipline is what the framework does *not* commit to. The framework does not commit to a specific interpretation of probability beyond standard Kolmogorov probability. The framework does not commit

to a specific resolution of the measurement problem beyond the algorithmic-reading developed in Chapter 1. The framework does not commit to specific values of free parameters in the emergent quantum field theory beyond what the reconstruction theorem forces. The framework does not commit to a specific cosmological history beyond the structural arrow of horizon growth in the Λ -dominated phase. These non-commitments are part of the foundational economy: each additional commitment would expand the framework’s claims, and the discipline is to make only those commitments that the foundational premise plus its derivations require.

The single-foundational-commitment pattern distinguishes the framework from broader unification programs in a specific way. String theory commits to extended one-dimensional objects as the fundamental entities, plus supersymmetry, plus extra dimensions, plus a specific moduli structure. Loop quantum gravity commits to a graph-theoretic discrete geometry, plus a specific quantization procedure, plus a specific representation theory of the spatial diffeomorphism group. Causal sets commit to a Lorentz-invariant discrete causal structure, plus a specific dynamics for set growth. Each of these is a multi-commitment foundation, with each commitment adding to the framework’s content. The framework here is a single-commitment foundation, with all content derived from the one premise plus its structural lemmas. Whether this is a methodological virtue or a methodological constraint depends on whether the derivations actually close — whether single-commitment foundations can produce the content that multi-commitment foundations produce by direct positing. The framework’s claim is that they can; Parts II–IV develop this claim with the specific quantitative content.

A further comparison with string theory is worth drawing, because it clarifies what the single foundational commitment is and is not. String theory’s landscape — the vast space of consistent vacua, among which the observed one is often held to be fixed by anthropic selection — is sometimes proposed as a model for how a theory can be general in its laws yet specific in what we observe. It is tempting to ask whether the framework’s foundational commitment is similarly a selected feature: one option among many, picked out in our case because only some options permit observers. The answer sharpens what the commitment is. The framework does have a “general, then selected” structure — but it lives *downstream* of the foundational commitment, not at it. The emergent description is gauge: it depends on the partition, with the substratum gauge group identifying which structure is observable; and the partition that supports an observer is then selected, by the structural observer-selection theorem of Chapter 1, as one of those meeting the embedding conditions. This is the framework’s counterpart of “general law, particular vacuum,” and unlike the landscape it is derived — a constructively classified gauge structure and a selection theorem, rather than a posited space of vacua and an anthropic measure that has resisted rigorous definition. What this structure does *not* reach is the foundational commitment itself. The commitment, distilled, is two axioms: that tokened differentiation occurs, and that it recurs (Chapter 1). Neither

can be a landscape variable, because a landscape is a space of differentiated, temporally-extended worlds — it presupposes both axioms in order to have points at all. Anthropic selection picks among worlds; it cannot pick whether there is a world of the relevant kind, since the selection apparatus, a measure over possibilities, already requires the temporal domain it would be invoked to explain. The “gauge in general, specific in our case” intuition is therefore correct and is realized in the framework — at the level of partitions and emergent descriptions, one level below the two axioms, which are what must hold for there to be a space of possibilities to select within.

4.3 Classification of claims by derivational rigor

A framework that derives rather than posits must be honest about how rigorous each derivation actually is. The framework’s discipline is to classify every prediction by its derivational rigor, using an explicit scheme with seven primary categories. The categories are not interchangeable: the difference between a strictly parameter-free structural prediction and a retrodiction fitted to observation is substantial, and conflating them would overstate the framework’s empirical record. The classification scheme is part of what makes the framework’s claims evaluable.

Class S — strictly parameter-free structural prediction. A prediction is S-class if its derivation chain walks end-to-end with all links solid, no fitted parameters at any step. The Cabibbo angle $1/(\pi\sqrt{2})$ is S-class: it follows from a single Brillouin-zone distance on the substratum lattice, with no fitted constants. The Koide relation $2/9$ is S-class: it follows from a cubic-group quadratic Casimir, with no parameters tuned to observation. The Bekenstein-Hawking $1/4$ coefficient is S-class: it follows from the ratio between the Euclidean KMS period and the Jacobson Einstein-equation prefactor, with no free parameters (the KMS step presupposes emergent local boost structure at the horizon, a dependency carried in [GR §8.5]). S-class predictions are the framework’s strongest empirical commitments, and their precision-upgrade falsification would refute the framework directly without any recoverable error mode beyond computational mistakes in the specific derivation.

Class C — conditional structural. A prediction is C-class if its derivation chain walks modulo a stated open assumption. The empirical match is unaffected by the open assumption, but the prediction’s structural-rigor status is contingent on the assumption’s resolution. C-class predictions are weaker than S-class because their derivation is conditional, but the empirical content can be just as sharp. The distinction is methodological honesty about what has been derived versus what has been assumed.

Class L — layered conditional. A prediction is L-class if its derivation walks given Layer-2 inputs, with closure path being a specific active research direction. L-class is a refinement of C-class that identifies *which* derivational work would move the prediction to S-class. The pattern is: “the prediction is

S given that Direction X closes.” Active research directions in the framework include the analytic dispersion-relation calculation at Direction 10, the structural derivation of certain spinor-structure identifications, and others. L-class predictions become S-class as the relevant directions close, and the framework’s rigor improves monotonically over time.

Class R — retrodiction. A prediction is R-class if its value is fitted to observation by construction. R-class predictions cannot falsify in the precision-upgrade sense — they were fitted, so they match by definition. The structural content lives in *what is predicted given the fit* rather than in the fitted value. For example, the SU(2) and SU(3) gauge couplings at the Z pole are R-class in the framework: three fitted parameters (δ_0, A, B) against three observed couplings produces zero residual by construction. The framework’s structural content here is the *form* of the gauge-coupling fit (universality of δ_0 , the resummation form, the structural prediction $1/\alpha_0 = 23.25$ from the 1-loop staggered vacuum polarization, and the $A \cdot B$ cross-check (leading-power-conditional; inconclusive across finite-size models)), not the matching of the three couplings to observation.

Class P — phenomenological input. A prediction is P-class if it takes a specific quantity from experiment because the framework does not yet derive it. The strange-quark mass m_s is a phenomenological input in the framework’s current content: it sets the overall fermion mass scale, but the framework does not derive it from the substratum. Acknowledging the input is methodologically necessary — the framework cannot claim to predict m_s — but does not undermine the predictions that derive from it once it is fixed.

Class M — mass-chain inheritance. A prediction is M-class if its empirical match follows from a single upstream empirical input. The electron and muon masses, given the tau mass, are M-class: they inherit from the tau mass through the Koide relation rather than being independently derived. M-class predictions are not retrodictions — the Koide relation is a structural prediction of the framework, and the inheritance through it is a derivation — but their empirical match depends on the upstream empirical input being what it is observed to be.

Class E — explicit empirical input. A prediction is E-class if some specific input enters the derivation chain from experiment, acknowledged honestly. E-class is broader than M-class, covering inputs that are not specifically masses. The Jarlskog phase η , for example, enters the CKM derivation as an empirical input rather than as a derivation from the substratum.

The four-layer framing. Orthogonal to the S/C/L/R/P/M/E classification is a four-layer framing that locates predictions within the substratum-to-emergent stack. The two axes are independent: an S-class prediction can be at any of the four layers, and a Layer-1 prediction can be S, C, L, R, P, M, or E class.

Layer 0 covers gauge structure: which gauge groups exist, which representations appear, structural constraints from C1-C3. The Standard Model gauge group $SU(3) \times SU(2) \times U(1)$ is Layer 0. The three-generation count is Layer 0.

Layer 1 covers structural form: predictions from pure substratum geometry or representation theory, with no operator-relation inputs. The Cabibbo angle $1/(\pi\sqrt{2})$ is Layer 1: it follows from a Brillouin-zone distance. The Wolfenstein parameter $A = \sqrt{2/3}$ is Layer 1. The Koide angle $\theta_0 = 2/9$ is Layer 1.

Layer 2(a) covers operator-relation structural predictions: the structural form of a relation, given that the relevant operators exist. The Gatto-Sartori-Tonin relation applied to the Layer-1 Cabibbo derivation is Layer 2(a).

Layer 2(b) covers solution-specific predictions within a fixed operator structure: mixing-angle values determined by the specific bijection's parameters.

Layer 3 covers mass-scale and bijection-specific predictions: the specific values of m_s, μ_c, μ_w , etc., that pick out which bijection φ within the equivalence class our universe occupies.

The framework's predictions span the four layers. Layer 0 and Layer 1 predictions are the framework's most distinctive content, because they involve only the substratum's structural features and no specific solution-level inputs. Layer 2(a) predictions are nearly as distinctive but conditional on the operator structure. Layer 2(b) and Layer 3 predictions involve more solution-specific content and are accordingly less distinctive — they would differ for different bijections within the same equivalence class.

The substratum / emergent / mixed architectural classification. A third axis classifies each load-bearing element of a calculation by where it sits in the substratum-to-emergent architecture. *Substratum* elements operate on the bijection (S, φ) on the cubic lattice with its coupling matrix; predictions at this level are pure structural ratios from geometry and representation theory. *Emergent* elements operate on the unitary quantum field theory after the trace-out: induced gauge couplings, renormalization-group running, Coleman-Weinberg potentials, Z -factors; predictions at this level are perturbation-theory and RG outputs. *Mixed* elements involve QFT machinery in the derivation but with specific substratum-level inputs constraining the output — the most typical case for the framework's quantitative predictions in Chapter 5.

The framework's general rule is *group structural, couplings emergent*. The gauge group is determined at the substratum level by the cubic-group commutant; the gauge couplings at the Z pole are determined at the emergent level by RG running from the lattice-determined Planck-scale coupling.

Why the classification matters. A framework that classified all its predictions as S-class would either overstate its empirical record (if some predictions are actually conditional or retrodicted) or be implausibly successful (if all predictions actually walk end-to-end with no inputs). The framework's explicit classification of each prediction by its derivational rigor is a methodological discipline that allows the empirical record to be evaluated honestly. The strongest claims live in the S-class entries; the conditional claims live in C and L; retrodictions and inputs are acknowledged honestly. The framework's twenty-two

Standard Model retrodictions split across these categories — some S-class, some C-class, some R-class — and the split is part of what readers should know when evaluating the framework’s empirical record against the framework’s structural commitments.

4.4 Engaging the no-go results through one architecture

The framework is a hidden-variable theory in the literal sense: there is a deterministic substratum with definite states, and the observer’s stochastic description is obtained by marginalizing over what they cannot see. The twentieth-century foundations literature contains a series of no-go theorems ruling out hidden-variable theories satisfying various structural conditions. The framework must explain how it evades each of these results. Chapter 3 §3.7 developed the technical content of the engagement; this section develops the methodological principle that organizes it.

The principle is straightforward to state. Each no-go theorem identifies a structural property of would-be hidden-variable theories and proves that the property is incompatible with quantum predictions. The property is some conjunction of conditions on the underlying hidden-variable structure and the observer’s measurement procedure. The framework’s two-axis hierarchy, developed in Chapter 3, has the resource to satisfy each condition at one level of the hierarchy while violating the corresponding feature at another level. The evasion of each no-go theorem is therefore not a separate trick but an instance of the same general procedure: identify the hierarchy level at which the no-go theorem’s assumption sits, verify that the framework satisfies the assumption at that level, and verify that the corresponding feature fails at the level where it must fail to reproduce quantum predictions.

The four-step procedure. Applied to a no-go theorem in the foundations literature, the procedure has four steps.

First, identify the theorem’s structural assumption. Bell’s theorem assumes factorizability $P(a, b \mid x, y, \lambda) = P(a \mid x, \lambda) \cdot P(b \mid y, \lambda)$. Kochen-Specker assumes non-contextual value assignments. PBR assumes preparation independence — that independently prepared systems have a product joint ontic-state distribution. Frauchiger-Renner assumes the universal applicability of quantum mechanics across nested observer-modeling.

Second, identify the level of the framework’s hierarchy where the assumption sits. Bell’s factorizability is a property at Level G3 $\times$ Level D, where the joint transition matrix between visible-sector subsystems is defined. Kochen-Specker non-contextuality is a property at Level G2 $\times$ Level D, where observables are defined as operators on the emergent Hilbert space. PBR preparation independence is a property at Level D $\times$ Level G3, where the substratum’s joint-system structure is specified. Frauchiger-Renner universal applicability is a claim at

Level G2 $\times$ Level D applied to substratum-level objects.

Third, locate where the framework’s two-level structure declines to supply the object the assumption is about. For three of the four theorems this takes the same form. Bell’s parameter-independence (one Jarrett component of factorizability) is preserved at the substratum level by spatial locality; outcome-independence (the other component) is violated at the emergent level by P-indivisibility. PBR’s preparation independence — the postulate that independently prepared systems have a product joint ontic-state distribution — fails because the framework supplies only *conditional* independence given shared boundary data (the spatial Markov property of the coupling graph), not the unconditional product independence the postulate requires. Frauchiger-Renner’s universal applicability fails because the substratum-level and emergent-level descriptions are incommensurable in the relevant sense. The fourth theorem, Kochen-Specker, does not fit this “satisfied below, violated above” form: its assumed property, a noncontextual valuation of the observables, holds at neither level — the framework reproduces quantum mechanics, which by Kochen-Specker’s own content admits no noncontextual valuation, and the substratum carries a definite configuration that is not a valuation of the emergent observables at all. The framework’s contextuality is forced by its equivalence with quantum mechanics, and what the framework adds is the account: contextuality is the partition-dependence of the trace-out. The common factor across all four is therefore the architecture — two levels and a trace-out — not a single property-placement move repeated.

Fourth, verify that the framework’s predictions match quantum mechanics in the regime where the no-go theorem applies, without producing signatures the no-go theorem’s empirical content rules out. Bell’s empirical content includes the prohibition of superluminal signaling and the Tsirelson bound; the framework satisfies both. Kochen-Specker’s empirical content includes the contextuality of measurement outcomes for non-commuting observables; the framework predicts exactly this. PBR’s empirical content includes the indistinguishability of certain pure states; the framework reproduces it through the algebra-channel structure. Frauchiger-Renner’s empirical content includes specific outcome statistics in observer-modeling scenarios; the framework predicts exactly the same statistics, with the appearance of inconsistency being an artifact of conflating descriptive levels.

One architecture, verified per theorem. Two things must be held apart here, because they pull in opposite directions and both are true. The framework’s *response* to the no-go literature is not piecemeal: it is one architectural feature — two levels of description joined by the trace-out — fixed before any no-go theorem is examined, and the four cases are four manifestations of working within it rather than four separately chosen escapes. But the *verification* that the architecture in fact answers a given theorem is per-theorem work: each theorem’s assumption must be located against the two-level structure, and the match checked. The unity is in the response; the labor is in the checking. As

new no-go theorems are published — the literature continues to grow — each must be checked against the same architecture, with the framework’s engagement verified rather than assumed.

This per-theorem verification is a genuine discipline, not a rhetorical disadvantage. A framework that claimed to evade all no-go theorems with one move (e.g., by adopting superdeterminism, where the hidden-variable distribution is fine-tuned to experimenter choices) would skip the checking but would face a different problem: the fine-tuning itself becomes a structural commitment, and the framework’s foundational economy is compromised. The framework here does the per-theorem checking in exchange for a cleaner foundational structure with no fine-tuning of hidden-variable distributions to detector settings — and the checking, in each case so far, finds the same architecture already sufficient.

Methodological benefit. The benefit of the architecture-level engagement is that the framework’s structural reading of each no-go theorem is *informative*. Identifying how each assumption meets the two-level structure explains *why* quantum mechanics has the features the no-go theorems identify, not just that it does. Bell’s outcome-independence violation is informative about the structure of the substratum’s joint dynamics. Kochen-Specker’s contextuality is informative about the partition-dependence of the trace-out. PBR’s preparation-independence failure is informative about the conditional, common-cause structure of separated subsystems sharing boundary data. Frauchiger-Renner’s inconsistency-paradox is informative about the levels of the framework’s hierarchy.

The no-go theorems become, in this reading, structural diagnostic tools rather than obstacles. Each one identifies a specific structural feature that any hidden-variable theory matching quantum mechanics must have. The framework’s identification of the relevant features through hierarchy levels makes the structural content of the theorems available for use in further work — in particular, in identifying which features any successful unification framework must satisfy.

Application to new no-go theorems. The procedure is open-ended. Any future no-go theorem in the foundations literature can be subjected to the same four-step analysis. The framework’s commitment is to apply the procedure honestly to each new result rather than to claim general immunity. If a no-go theorem turns out to identify a structural property that the framework cannot satisfy at any level — or to require simultaneous satisfaction of properties the framework can satisfy only at different levels — the framework would face genuine difficulty.

To date, no such theorem has been identified. The four no-go theorems analyzed here exhaust the load-bearing twentieth-century results in the foundations literature; recent results (Wood-Spekkens 2015 on fine-tuning, Cavalcanti-Lal 2014 on causal structure, Colbeck-Renner 2017 on no-extension theorems) have been analyzed in detail in the framework’s technical papers and admit the same architecture-level engagement, verified case by case against the two-level struc-

ture. The procedure is not a guarantee of future evasion but a discipline for analyzing each new theorem as it arrives.

The architecture-level engagement, verified theorem by theorem, is the framework’s general methodological response to the no-go literature. It is more demanding than claiming a single trick, but produces a richer structural account: each no-go theorem identifies a specific feature of the substratum-to-emergent architecture rather than serving as an obstacle to be circumvented, and the architecture answering all of them is one structure, not four.

4.5 Pre-registered falsification conditions

A framework that derives rather than posits must be willing to specify what would invalidate the derivation. The framework’s discipline is to pre-register specific falsification conditions, stating in advance which observational outcomes would refute the framework’s structural commitments and which would refute only narrow predictions while leaving the broader structure intact. The pre-registration distinguishes structural claims — which the framework stands or falls on — from solution-specific properties that lie outside the framework’s scope.

The conditions are organized into three classes. Class A identifies structural theorems whose violation would falsify the framework directly. Class B identifies parameter-free retrodictions whose precision-upgrade violation would falsify, with sub-classes reflecting the rigor of the retrodiction. Class C identifies broader framework-level commitments whose contradiction would refute particular framework claims while leaving others intact.

Class A — structural theorems. The framework derives specific structural features of physics as theorems from the foundational premise and the structural assumptions. Each derived feature is a falsification target if violated.

The strong-CP phase $\bar{\theta} = 0$, within the construction only. The framework derives $\bar{\theta} = 0$ at the kinematic level of the construction from T-invariance of the substratum bijection combined with detailed balance in the emergent Hamiltonian. Any observation of a non-zero neutron electric dipole moment consistent with $\bar{\theta} \gtrsim 10^{-11}$, at sufficient confidence to exclude experimental systematics, falsifies the framework’s T-symmetric substrate commitment.

The gauge group $SU(3) \times SU(2) \times U(1)$ exactly. The framework derives this gauge group as the unitary group of the bicommutant of the cubic-group action (equivalently, the stabilizer of the generic equivariant condensate) on the six internal components. Any observation of a stable beyond-Standard-Model gauge structure — an additional Z' interpretable as a new gauge boson rather than a composite resonance, stable exotic matter transforming under a novel gauge group, or any confirmed extension of the visible-sector gauge content — falsifies the cubic-group decomposition.

Three fermion generations. The framework derives the generation count of

three from coupling-degree minimization applied to the wave equation, locking $K = 2d = 6$ internal components and giving three spin-1/2 staggered tastes. Any observation of a stable fourth-generation fermion with Standard Model quantum numbers falsifies the taste-reduction derivation. Non-stable exotic states consistent with composite or resonance interpretation do not.

Majorana neutrinos with normal mass ordering. The framework derives this from the taste-breaking structure of the neutrino mass matrix. Inverted mass ordering, confirmed to high significance, falsifies the derivation. Confirmed Dirac nature of the neutrino — for example, a null result for neutrinoless double-beta decay at the sensitivity that would require Majorana masses well below the framework’s prediction — falsifies the prediction that no right-handed neutrino exists in the spectrum.

The Tsirelson bound. Bell-inequality violations are bounded above by $2\sqrt{2}$ in the framework, as a theorem rather than as an input. Any observation of a CHSH parameter exceeding $2\sqrt{2}$ in a loophole-free Bell test falsifies the P-indivisibility structure that the framework derives.

Exact unitarity of the visible-sector dynamics at accessible scales. The framework derives unitary quantum evolution as the exact embedded-observer description at every accessible mass and coherence scale, with dissipative corrections structurally suppressed to the dissipator scale of Chapter 1 §1.6. Any confirmed intrinsic-collapse signature — anomalous decoherence not attributable to environmental channels, scaling with mass or superposition size as objective-collapse models (GRW/CSL, Diósi–Penrose) require, at sufficient confidence to exclude systematics — falsifies the emergence theorem directly. Because the framework’s predicted deviations are pinned orders of magnitude below the collapse-model band, a positive detection cannot be absorbed as a framework correction. The current frontier is matter-wave interference at macroscopicity $\mu = 15.5$ (Pedalino et al. 2026; Chapter 1 §1.6), consistent with exact quantum behavior; the frontier’s continued advance is a standing test rather than a completed one.

Leading Lorentz-invariance-violation coefficient. The framework derives a specific subluminal dispersion relation with the leading correction at order $(E/M_{\text{Pl}})^2$, with a physical, direction-dependent anisotropy whose orientation-averaged coefficient is $4/45$ ($\langle\delta\rangle = 1/90$; the body diagonals are exactly luminal). The anisotropy’s absolute orientation in space is gauge, but its $\ell = 4$ cubic-harmonic pattern and the direction-differences are physical — the cubic structure is an empirically-selected input, not a gauge choice (see the Standard Model and Substratum papers). Two structural features of this prediction matter for falsification. First, the dispersion has no linear-in- E/M_{Pl} term — any demonstration of linear Lorentz-invariance violation at any scale falsifies the framework’s substrate-level commitment to the cubic-lattice wave equation. Second, the dispersion is subluminal everywhere on the momentum sphere; superluminal LIV anywhere falsifies. Current quadratic-LIV bounds from gamma-ray-burst time-of-flight measurements reach roughly thirteen orders of magnitude below the Planck scale, so the $4/45$ coefficient is not

near-term testable at the magnitude level; linear-LIV tests, by contrast, are already at sensitivities where any positive detection would falsify. These statements describe the dispersion at tree level; whether the absence of a relevant (low-dimension) Lorentz-violating operator survives radiative corrections in the interacting theory is a separate, open question — the Lorentz fine-tuning problem common to discrete-substrate approaches — discussed in the Standard Model paper. The tree-level prediction and its falsifiers stand regardless; what is open is the loop-level protection — and notably cubic symmetry alone does not supply it (it permits power-divergent operator mixing), so the relevant question is whether the framework’s coarse-graining map acts as a sufficient ultraviolet smearing, the mechanism under which rotational-symmetry restoration is known to hold at all orders. The protection mechanism available here (the framework not being supersymmetric) is renormalization-group emergence — Lorentz-violating couplings flowing to an infrared-attractive Lorentz-invariant fixed point (Chadha–Nielsen; Bednik–Pujolàs–Sibiryakov), natural because the trace-out is a coarse-graining — with the documented caveat that the inter-species velocity flow is slow. The causal-set discreteness results (Bombelli–Henson–Sorkin) independently confirm that a finite-valency, locally-coupled discrete structure cannot be microscopically Lorentz-invariant, so the invariance here is necessarily emergent rather than fundamental.

Class B — parameter-free retrodictions. The framework produces twenty-two quantitative predictions that follow from the structural content with no adjustable parameters beyond a small acknowledged empirical input set. These are stratified into four sub-classes reflecting their structural status.

Class B-S (unconditional structural). These are layer-0 or layer-1 retrodictions with no open assumptions in their derivation. Examples include the Cabibbo angle $1/(\pi\sqrt{2})$ from a single Brillouin-zone distance, the Wolfenstein parameter $A = \sqrt{2/3}$ from the same structural source, the mass ratio $m_d/m_s = 1/(2\pi^2)$ via the Gatto-Sartori-Tonin relation, the Koide angle $\theta_0 = 2/9$ from a cubic-group quadratic Casimir (matched to 0.02%, the sharpest empirical agreement in the framework), the Bekenstein-Hawking coefficient $A/(4\ell_P^2)$ (derived; no direct test — GW250114 constrains the area theorem), and the MOND critical acceleration $cH/6$ from the boundary entropy-displacement mechanism. Any of these moving outside 3σ on precision upgrade falsifies the framework with no recoverable error mode beyond computational mistakes in the specific derivation.

Class B-L (layered conditional, pending open derivations). These predictions follow from the structural form combined with one Layer-2 substrate input or empirical input, with a specific active research direction whose closure would move the prediction to B-S status. The JUNO neutrino-mixing prediction $\sin^2 \theta_{12} = 1/3 - 1/(4\pi^2)$ matched at 0.07σ is B-L pending closure of the mixing-matrix derivation in the lepton sector. Other B-L predictions exist throughout Chapter 7’s quantitative content.

Class B-M (mass-chain inheritance). These predictions inherit their empirical match from a single upstream empirical input. The Higgs mass m_H via Standard

Model RGE running from the boundary condition $\lambda(M_{\text{Pl}}) = 0$, with m_t as the empirical input, is B-M. Class B-M predictions falsify under the same precision-upgrade protocol as B-L, but the inheritance is from an empirical input scale rather than from an open derivation.

Class B-R (retrodictions). The SU(2) and SU(3) gauge couplings at the Z pole are explicitly retrodictions: three fitted parameters (δ_0, A, B) against three observed couplings produces zero residual by construction. Class B-R predictions cannot falsify in the precision-upgrade sense; their value lies in falsifying the framework’s gauge-sector structural commitments — universality of δ_0 , the resummation form, the structural prediction $1/\alpha_0 = 23.25$ from the one-loop staggered vacuum polarization, and the $A \cdot B$ cross-check (leading-power-conditional; inconclusive across finite-size models). Any precision-upgrade violation of these structural commitments would falsify the framework’s gauge-sector account while leaving the broader Class A theorems intact.

The Class B split is, in honest summary: four predictions are unconditional structural (B-S), six are layered conditional pending open derivations (B-L), two are mass-chain inheritances (B-M), and two are retrodictions (B-R). Total fourteen predictions in Class B, with the empirical match holding for all fourteen within about 1% or 1σ at current measurement precision. As experimental precision improves, the discrimination power of the B-S set increases monotonically; closure of the open derivations would move B-L predictions to B-S over time.

Class C — framework-level commitments. Two broader commitments are implicit in the construction and would falsify it if contradicted.

Classical-memory simulability of any genuinely quantum process. The characterization theorem developed in Chapter 1 §1.8 establishes that quantum mechanics is the description produced by an embedded observer satisfying C1-C3 on a deterministic substratum with bounded classical hidden-sector memory. This implies that every process the framework calls quantum admits a classical-memory simulation at appropriate scale. The commitment is specific: non-Markovian models are permitted and expected (they are how the hidden-sector correlation persistence manifests), but they must be realizable by finite classical memory. A process whose non-Markovianity is provably supra-classical falsifies the framework at the characterization-theorem level.

The dark sector as entropy-displacement, not as a novel particle content. The framework’s account of dark energy as a Type II running-vacuum functional form and dark matter as frozen boundary entropy, developed in Chapter 6, is incompatible with the direct detection of a WIMP interpretable as a fundamental particle with Standard-Model-like interactions and no entropy-displacement signature. A confirmed WIMP with such properties falsifies the framework’s dark-sector account while leaving the rest of the structure intact — it would require abandoning the boundary-entropy derivation and either identifying an error in it or accepting that the framework’s dark-sector story is wrong even as

other derivations survive.

The cosmological commitment. The dark energy sector is subtler than the strict Class A through C taxonomy allows. The framework commits to the Type II running-vacuum functional form, with magnitude bounded at the emergent-QFT scale $|\nu_{\text{QFT}}| \sim (M_{\text{SM}}/M_{\text{Pl}})^2 \sim 10^{-32}$, observationally indistinguishable from Λ CDM. The falsification band is stratified: a DESI Year 5 measurement of $|\nu| \gtrsim 10^{-4}$ at high significance would require additional substratum structure not currently in the framework but would not refute the Type II functional form; a high-significance detection of *phantom crossing* — a $w(z)$ that crosses -1 — would falsify the Type II functional form itself, which is a structural commitment. The sign of a small detected ν is not by itself a structural discriminator: the Type II class admits either sign, the emergent-channel computation gives a negative sign under one reading of its gravitational status and no source under the other (Chapter 7), and a positive detection at the 10^{-4} level would require substratum structure carrying a correlated Lorentz-violation signature. The first outcome is a refinement within the current bracket; the crossing is a structural falsification.

The pre-registration commitment. The author commits to treating the Class A conditions, the Class C commitments, and a Class B retrodiction exceeding 3σ on precision upgrade as framework-falsifying rather than as an invitation to post-hoc rescue. In the case of a Class B violation, the framework will be revisited to determine whether the violation reflects a computational error in the specific retrodiction, an incorrect specific-bijection input, or a structural problem with the derivation; the first two are recoverable, the third is fatal. The framework does not reserve the right to add structural postulates after the fact to accommodate disconfirming evidence. The pre-registration of this commitment is itself part of the falsifiability content of the framework — a framework that pre-registers its falsification conditions and then revises them post-hoc has compromised the discipline that pre-registration was meant to enforce.

4.6 What does not falsify and regimes of deliberate silence

The framework’s pre-registered falsification conditions define what would refute the framework’s structural commitments. The complement of those conditions — what lies outside the falsification schedule — also deserves explicit articulation. Some observational outcomes that might appear to bear on the framework do not bear on it, because the framework makes no commitment on the relevant question. The discipline of pre-registering falsification conditions requires equal honesty about what is and is not within the framework’s predictive scope.

Solution-specific properties. The framework derives the structural form of physics — the gauge group, the matter content, the generation count, the parameter-free retrodictions, the thermodynamic content of the gravitational sector. It does not derive every contingent feature of the observed universe.

Some features are solution-specific: properties of the particular bijection φ within the framework’s equivalence class that describes our universe, as distinct from properties of the equivalence class itself. The framework’s commitments are to the equivalence class; solution-specific properties are inputs to selecting the representative within the class, not derivations from the framework’s foundational premise.

Several quantities are explicitly solution-specific in the framework’s current content.

The absolute scale of fermion masses depends on the strange-quark mass m_s as an acknowledged empirical input. The framework derives mass *ratios* from the cubic-group representation theory and the Gatto-Sartori-Tonin relation, but the overall mass scale is set by one external input. A future precision measurement of m_s at different value would shift all derived fermion masses proportionally without bearing on the framework’s structural derivation. The framework predicts the *pattern* of masses given m_s ; it does not predict m_s itself.

The CP-violating phases in CKM and PMNS are solution-specific. The framework derives $\bar{\theta} = 0$ at the kinematic level of the construction, which is the structural CP-conservation constraint, subject to the limit §6.3 sets out. The flavor-basis CP phases — the Jarlskog parameter, the Wolfenstein (ρ, η) , the PMNS Dirac and Majorana phases — are features of the basis mismatch between up-type and down-type Yukawa mass eigenstates, not consequences of any structural feature of the substratum. The framework predicts $\theta = 0$ and does not predict the flavor-sector CP phases. The symmetry that delivers $\bar{\theta} = 0$ also forces those phases to vanish: within the construction, rotational isotropy and CP conservation stand or fall together, while observation requires one of each. The construction is therefore in tension with the measured CKM phase, and treating the phases as solution-specific does not remove it [SM §5.5].

The numerical value of the Type II running-vacuum coefficient ν within the stratified falsification band is solution-specific. The framework predicts the *functional form* (Type II RVM) as a structural commitment, with the magnitude $|\nu_{\text{QFT}}| \sim (M_{\text{SM}}/M_{\text{Pl}})^2 \sim 10^{-32}$ bounded at the emergent-QFT scale from the emergent-QFT calculation. The Bertini et al. (2025) DESI-era measurement $\nu = -(2.5 \pm 1.3) \times 10^{-4}$ is consistent with both the framework’s emergent-QFT-scale value and with Λ CDM, lying within the predictive bracket. A DESI Year 5 measurement that landed anywhere within the bracket — including values much larger than the emergent-channel bound, of either sign, provided the $w(z)$ shape remains Type II (no phantom crossing) — would not by itself falsify; the structural commitment is to the functional form, not to the specific numerical value within the bracket, though a positive detection would require the non-local substratum branch with its Lorentz-violation co-signature.

Cosmic initial conditions are solution-specific. The framework derives the structure of the dynamics — the wave equation, the bijection, the partition — from which evolution proceeds. The state at any particular cosmic time, including the

time-zero conditions, is not a framework output. The framework is compatible with any cosmic-history that produces a C1-C3-satisfying observer-admitting regime by the time the cosmological horizon stabilizes. Specific cosmological initial-condition observations bear on the framework only at the boundary where they would or would not produce the observer-admitting regime; within that boundary they are solution-specific.

The specific microstate φ within the equivalence class is solution-specific *and* gauge. Chapter 2’s substratum gauge group $\mathcal{G}_{\text{sub}}$ classifies different representatives within an equivalence class as observably indistinguishable. The question “which specific bijection φ is our universe’s substratum?” therefore has no empirical content — every representative produces identical observables. A future putative measurement that purported to distinguish different gauge representatives would, if confirmed, refute the substratum gauge group’s completeness, but no such measurement is conceivable in principle within the framework’s structure.

Regimes of deliberate silence. Beyond solution-specific properties, three regimes lie outside the framework’s predictive scope by structural construction. These are regimes where the framework’s foundational structure does not produce predictions, not because the framework is silent by choice but because the structural conditions for predictions to be defined do not hold.

The pre-horizon early universe. The framework’s gravitational sector derivations in Chapter 6 operate on a stable causal partition with $\tau_B \sim H^{-1}$ — the cosmological horizon as a partition between observer-accessible and observer-inaccessible regions. The pre-horizon era — before such a partition stabilizes — has no observer in the framework’s sense, because the C1-C3 conditions for embedded observation are not satisfied. The framework’s only commitment in this regime is a boundary condition on deeper theories: any successful pre-horizon dynamics must produce a C1-C3-satisfying regime by the time a horizon stabilizes. The framework makes no prediction for the dynamics of the pre-horizon era itself; that domain belongs to whatever deeper theory describes the pre-horizon transition.

The equilibrium phase of (S, φ) . Chapter 1’s central theorem requires C1, C2, and C3 to be satisfied. The slow-bath condition C2 fails at substratum equilibrium: at equilibrium, the hidden sector has thermalized completely, and there are no slow modes to provide the timescale separation that produces P-indivisibility. The structural observer-selection theorem developed at Main §4.6 restricts observers to the non-equilibrium phase of the substratum’s dynamics. The equilibrium phase is observer-free by structural argument, and the framework makes no prediction for what happens in this regime — there is no observer to whom a prediction could be addressed. This is the framework’s account of why Boltzmann brains do not appear in observed physics: not because they are rare random fluctuations that happen not to occur, but because the equilibrium phase that would produce them is structurally observer-free.

Regions beyond the observer’s causal horizon. For any observer with a finite cosmological horizon, regions beyond that horizon are part of the observer’s hidden sector. They are addressed in the framework only through boundary effects — the area-law entropy at the horizon surface, which encodes the gravitational signature of the hidden-sector content. The framework takes no observer-independent position on the *internal content* of these regions. Whether the universe beyond an observer’s horizon contains fields configured differently, contains different gauge structures, or contains anything at all in the observer’s conceptual vocabulary, is not a question the framework’s structure can address. The framework’s account is that the trace-out over what lies beyond the horizon produces the observer’s quantum-mechanical description; the content traced over is, by construction, inaccessible to that observer.

These silences are structural rather than methodological. They follow from the characterization theorem’s identification of quantum mechanics with embedded observation under C1-C3, and from the observer-selection theorem’s restriction of observers to the non-equilibrium phase. The framework’s silence on these regimes is not a gap in its content but a feature of its structural commitments: it is committed to the regime where embedded observers exist and the regime is non-equilibrium; it is not committed to regimes where these conditions fail.

What this section is not. The articulation of what does not falsify is not a list of escape routes. The framework’s structural commitments — the Class A theorems, the Class B retrodictions, the Class C broader commitments — are pre-registered as falsifiable, and the author’s commitment is to treat their violation as framework-falsifying rather than as an opportunity for post-hoc revision. The solution-specific properties and the regimes of deliberate silence are not categories into which falsifying observations can be reclassified after the fact. They are explicit acknowledgments of where the framework’s commitments end and where the framework’s structure has no predictive content. The line between them is drawn before the relevant observations are made, with the structural derivation determining which side of the line each putative prediction falls on.

4.7 What the framework derives, and what it compresses

A framework that advertises “the Standard Model, gravity, and quantum mechanics from a single object” invites a sharp question, and the methodological honesty of the preceding sections requires meeting it directly: is the derivation genuine, or is the answer quietly placed in the construction and then recovered? The answer is neither a simple yes nor a simple no, because the framework’s results do not all stand in the same relation to its starting point. They fall into three tiers, and naming the tier is the honest form of every claim the book makes.

The first tier is **forced**. Once the substratum is committed to — a simple-cubic lattice carrying the wave equation — a great deal follows by pure representation

theory, with no remaining freedom and no further input. The decomposition of the six lattice link-directions under the cubic point group fixes the gauge group $SU(3) \times SU(2) \times U(1)$ and its multiplicity pattern; the same structure fixes exactly three fermion generations, the unique anomaly-free hypercharges, the vanishing of the strong-CP angle, and the Cabibbo angle to four-significant-figure agreement. These could have come out otherwise — a different point group, a different count, a different angle — and the content of the derivation is that they do not. This tier is the framework’s genuine achievement, and it is substantial: it is not reverse-engineering, because the predictions exceed the inputs, most visibly where the Cabibbo angle lands at $1/(\pi\sqrt{2})$ while the competing cubic lattices miss the measured value by tens of percent.

The second tier is **selected by the world**. Several quantities the derivation leans on are fixed not by the substratum but by observation. The spatial dimension is three because gravity propagates, matter is stable, and space is observed flat — three facts read off the universe, not consequences of the bijection. The lattice is *cubic* because we observe an $SU(3)$ in nature and only the cubic crystal system can carry one. Here the reasoning runs backward, from the world to the commitment: “the Standard Model from a cubic lattice” and “the lattice is cubic because we see $SU(3)$ ” are the same fact spoken in two directions. The framework does not hide this — it states the cubic structure as a commitment forced by the observed gauge group — but the headline should carry it too, because it marks the boundary of what is derived.

The third tier is **stipulated**. Some load-bearing moves are decisions about what is to count as physical structure rather than facts the bijection compels: that a “site” is *defined* by minimal coupling (from which the internal component count follows), that couplings are nearest-neighbor, and — at the foundation — that the limitation of an embedded observer takes the specific form of a fast sector reading and writing to a slow, capacious hidden one. These choices are principled and motivated, but they are choices.

Seeing the three tiers at once dissolves a question the framework had carried as open — whether its structural residue is “genuinely non-tautological.” The question has no single answer because the residue is not one kind of thing: it is non-tautological at the first tier, world-selected at the second, definitional at the third. And no calculation settles it, because it was never an empirical question; the logical status of each step is fixed by the construction as written. What the classification yields is a precise description of what the framework *is*. It is a **maximally tight compression** of known physics onto a small set of commitments. The first tier shows that, once those commitments are granted, the discrete structure of the Standard Model is largely not a list of independent facts but the unfolding of a single representation-theoretic object — and *that* is a genuine and unusual claim about the unity of physics, worth the book whether or not the deeper pretension holds. The second and third tiers show that the commitments themselves are anchored in observation and partly stipulated, so the framework is not spinning the world out of a parameter-free principle.

This is also the place to be exact about the foundation that names the book. “Observation is incomplete” — that what we know is reached through a partial vantage, never the whole — is, taken by itself, not a falsifiable claim about physics. It is close to the definition of the empirical stance itself: every theory in science is a description built from inside, from limited access. A thesis that cannot be false cannot, on its own, predict. All of the framework’s predictive content therefore lives one level down, in the *specific model* of the incompleteness — the architecture, the lattice, the dynamics — which is precisely where the second and third tiers locate every input. The framework’s reach and its circularity are one structure seen from two sides. Read as the *source* of the predictions, “observation is incomplete” claims too much. Read as the *organizing principle* under which an unusually large body of physics compresses to one structure, it is the defensible and still-remarkable claim, and it is the claim this book makes. The reader should hold the framework to that standard and no other: not as a derivation of the universe from the bare fact that there are observers, but as a demonstration that an enormous amount of what we observe coheres as the shadow of one finite construction — with the construction’s own starting points held up honestly to the light.

4.8 What follows

The methodological discipline developed in this chapter — the single-foundational-commitment pattern, the claim classification by derivational rigor, the architecture-level engagement with the no-go results, verified theorem by theorem, the pre-registered falsification conditions and their complement of deliberate silences — provides the framework with a discipline that organizes its program as a research program rather than as a single bold conjecture. The discipline is part of the framework’s content, not separate from it: the structural commitments developed in Chapters 1 through 3 and the empirical predictions developed in Parts II through IV are produced under this discipline, with the discipline determining what the predictions are, how they are classified, and what their falsification status is.

Parts II and III apply the methodology to the framework’s specific empirical content in fundamental physics. Chapter 5 develops the Standard Model derivation, with the twenty-two parameter-free retrodictions classified across the S, C, L, R, M, and E categories developed here. The four-layer framing locates each retrodiction in the substratum-to-emergent stack, with Layer 0 and Layer 1 predictions (gauge group, generation count, Cabibbo, Koide, Wolfenstein) carrying the framework’s strongest structural commitments and Layer 2 and Layer 3 predictions (mass ratios within specific operator structures, scale-setting predictions) carrying more solution-specific content. Chapter 6 develops the gravitational sector with the Bekenstein-Hawking $1/4$ coefficient, the value of $\hbar$ from horizon thermodynamics, and the dissolution of the cosmological constant problem; these are all S-class predictions at the substratum-emergent boundary, with explicit falsification conditions identified in §4.5. Chapter 7 develops

the JUNO neutrino prediction as an L-class prediction matched at 0.07σ , with the closure of the lepton-sector mixing-matrix derivation as the open direction whose completion would move the prediction to S-class. Chapter 8 develops the universality-class structure and the framework’s relationship to alternative unification programs, with the architecture-level no-go engagement from §4.4 applied to the analysis of cross-framework structural relationships.

Parts III and IV develop the framework’s reach into structural cascades and applied scientific fields, with the methodological classifications carrying through to the more applied content. The framework’s predictions in chemistry, biology, neuroscience, and quantum engineering are not all S-class — many are P-class or E-class, with empirical inputs from the relevant scientific fields entering the derivation chain. The classification scheme developed here makes the rigor of each prediction explicit, so that readers approaching the framework from biology or medicine can evaluate the framework’s claims against the appropriate empirical standards in those fields.

The methodology developed in this chapter is not unique to the framework. The single-foundational-commitment pattern is a general principle of theory construction, the claim-classification scheme is a general standard for empirical evaluation, the architecture-level no-go engagement is a general method for analyzing structural results, and the pre-registered falsification discipline is a general epistemological standard. The framework’s contribution is to apply this methodology systematically rather than selectively, with every prediction carrying its classification and every structural commitment carrying its falsification condition. What this discipline produces, when applied to the framework’s foundational starting point — the two-axiom observation base — is the content developed in the remaining chapters of the book.

The methodological posture sketched in this chapter is developed at length, for a philosophy-of-physics readership, in a companion paper — *Physics Modulo Gauge* — which treats the two-axiom foundational commitment, constructive structural realism, the engagement with the no-go results of quantum foundations through a single architectural feature, and “physics modulo gauge” as a general methodology. That paper is a companion to the framework’s technical sequence rather than part of this book; readers concerned specifically with the framework’s methodology and its relation to the foundations-of-physics literature will find it the natural next text.

Chapter 5

The Emergence of Gauge Structure

5.1 What this chapter develops

Part I established that any observer embedded in a finite deterministic universe under conditions C1-C3 must use a mathematical formalism equivalent to unitarily evolving quantum mechanics; that observed physics together with structural assumptions uniquely determines a substratum (S, φ) up to the gauge equivalence $\mathcal{G}_{\text{sub}}$; and that the framework's structural-realism position operates on a two-dimensional hierarchy with stratified empirical commitments. The remaining question is what specific physics the substratum produces. This chapter and the next answer that question for the visible sector — the Standard Model of particle physics.

The derivation has two natural phases, and the book treats each in its own chapter. The first phase, the present chapter, establishes *what gauge structure emerges* from the cubic-lattice substratum: the wave equation, the multi-component dynamics, the cubic-group decomposition with eigenvalue multiplicities $(3, 2, 1)$, the Standard Model gauge group $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$ as the pointwise commutant, the chirality of the trace-out, and the anomaly cancellation that forces the observed hypercharges. The second phase, Chapter 6, establishes *what specific quantitative content* this gauge structure produces: three chiral fermion generations, one composite Higgs doublet, gauge couplings, and the twenty-two parameter-free retrodictions including the Cabibbo angle, the Koide relation, fermion masses, and PMNS mixing angles.

This chapter's content is structured. Section 5.2 establishes the minimal mathematical object — a finite set with a deterministic bijection of bounded coupling degree, with the coupling graph and the natural product factorization determined by the bijection rather than imposed by hand. Section 5.3 establishes the spatial dimension $d = 3$ from three independent self-consistency filters (propagating gravity, stable matter, boundary-entropy concordance). Section 5.4 develops the emergent quantum field theory from the wave equation: the factorization into staggered Dirac fermions, the multi-component structure $K = 2d = 6$ from coupling-degree minimization. Section 5.5 derives the cubic-group decomposition of the six internal components with eigenvalue multiplicities $(3, 2, 1)$, the pointwise commutant $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$, and its promotion to local gauge invariance by background independence. Section 5.6 develops chirality from the trace-out and anomaly cancellation forcing the observed hypercharges. Section 5.7 closes the chapter with a forward pointer to Chapter 6.

The structural status of the chapter's results is part of the framework's most distinctive single contribution. The Standard Model gauge group is not chosen from a landscape of alternatives, not embedded in a higher gauge group and broken, not derived from a separate set of postulates for each piece. It is forced — uniquely — by the structural inputs of Part I together with the cubic-lattice realization of those inputs. That realization is itself empirically anchored (the lattice is selected to match the observed gauge group, not derived from the substratum alone — §4.7), so the uniqueness is conditional on it rather than

derived from a parameter-free first principle. The methodological discipline developed in Chapter 4 will organize the predictions of Chapter 6 by derivational rigor; this chapter’s content sits at Levels 0 and 1 of the four-layer framing (gauge structure, structural form) and is correspondingly unconditional S-class structural — the most rigorous derivation tier in the framework’s classification. (One step in the chain — the reduction of the natural commutant $U(n)$ to $SU(n)$, §5.5 — is conditional on amplitude-scale gauge invariance, an explicit adjoined principle of $\mathcal{G}_{\text{sub}}$ (§5.2); the cubic multiplicities and the anomaly-fixed hypercharges are unconditional.)

The chapter takes as input the framework’s prior content. The quantum-mechanical emergence theorem of Chapter 1 is invoked where needed but not reproved. The reconstruction theorem of Chapter 2 provides the structural assumptions A1–A6 under which the substratum is uniquely determined. The hierarchy framework of Chapter 3 locates each prediction in the multi-level structural-realism scheme. The methodological discipline of Chapter 4 classifies each prediction by its derivational rigor. The chapter’s content is the substantive derivational chain from the cubic-lattice substratum to the Standard Model gauge group — the framework’s specific structural claim about *why* physics has the gauge symmetries it has.

5.2 The minimal substratum and the wave equation

The framework’s content rests on a minimal mathematical structure: a finite set S , a deterministic bijection $\varphi : S \rightarrow S$, and a partition of S into visible and hidden sectors. The lemmas of Chapter 1 extracted three structural features from this minimal object — finiteness, causal partition, determinism — and the conditions C1–C3 identified when an embedded observer can exist. This section develops two further structural features of the minimal object that the Standard Model derivation requires.

The factorization problem. Let S be a finite set with $|S| = q^N$ for some $q \geq 2$ and integer N . There are many ways to write S as a Cartesian product $S = S_1 \times S_2 \times \cdots \times S_N$ with $|S_k| = q$ for each k . Different factorizations correspond to different choices of which elementary degrees of freedom to identify as “local” — different bases in which to decompose the configuration space. The factorization is not given by the set S itself; it must be determined by some additional structure.

The framework’s answer is that the factorization is determined by φ rather than chosen externally. Given a bijection φ , define the *coupling graph* G_φ on vertices $\{1, 2, \dots, N\}$, with vertices i and j connected by an edge if the i -th component of $\varphi(s)$ depends on the j -th component of s for at least one state $s \in S$. The coupling graph is not an additional structure imposed on the system; it is read off from the bijection. The natural factorization of S is then the one that minimizes the coupling degree — the number of edges per vertex — over all possible product decompositions.

For the wave equation as substratum dynamics, this minimization has a unique answer. The wave equation in one dimension is

$$x_i(t+1) = x_{i-1}(t) + x_{i+1}(t) - x_i(t-1) \mod q,$$

producing irreducible algebraic dependence among all input components for $q > 2$. Any relabeling that splits or merges natural sites strictly increases the coupling degree. Numerical verification on small systems confirms this: zero percent of randomly chosen factorizations achieve the natural-site minimum, and one hundred percent have strictly higher coupling. The natural factorization is uniquely determined by the dynamics.

Factorization uniqueness theorem. *Let φ be a bijection on $S = (\mathbb{Z}/q\mathbb{Z})^N$ satisfying translation invariance, range-1 coupling on the natural factorization, and genuinely bidirectional coupling. Then the natural factorization is the unique minimizer of coupling degree among all product decompositions $S = S_1 \times \cdots \times S_M$ with $|S_k| = q$ for all k .*

The translation-invariant version of the theorem extends to general graphs through quasi-isometry arguments. Any bounded-degree graph with polynomial growth exponent d and statistical isotropy is quasi-isometric to $\mathbb{Z}^d$ by Gromov’s theorem applied to the appropriate group-theoretic setting. The wave equation on a d -dimensional graph has $2d$ independent low-energy propagation modes (from heat kernel asymptotics: the spectral dimension equals d , giving d momentum-like parameters over two orientations each). The integer-valued coupling degree depends on the graph only through its spectral properties, which vary continuously under quasi-isometric deformations. Since the coupling degree is integer-valued and the minimizer is $K = 2d$ for $\mathbb{Z}^d$, it cannot change without a discontinuous jump in the spectral dimension — which quasi-isometries preserve. The factorization uniqueness is therefore a property of the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$, not of the specific graph.

Space as coupling structure. Every “spatial” property used in the derivation chain is a property of the coupling graph G_φ . The *area* of a region $V \subset S$ is the number of edges crossing from V to its complement. The *dimension* is extracted from the causal partial order via the Myrheim-Meyer estimator. The *area law* — that entanglement entropy of a region scales as its boundary rather than its volume — follows from the spatial Markov property on any graph with range-1 dynamics. The *dispersion relation* describing how perturbations propagate requires translation invariance of the underlying dynamics. The regular cubic lattice $\mathbb{Z}^3$ is one graph with all these properties; it is not the only one.

The framework’s account of space dissolves the container problem that has been a recurring objection to discrete models of spacetime. The coupling graph is not embedded in a pre-existing space. There is no ambient manifold in which the graph “lives.” The vertices are labels for degrees of freedom, the edges record which degrees of freedom are dynamically coupled. Any spatial embedding is a representation for human convenience, not a physical fact. The coupling graph

is not made of material — the vertices are not atoms of space and the edges are not physical connections. A degree of freedom is a component of the state, an abstract mathematical entity rather than an object. Asking what a site is “made of” is a category error: it is a label for a variable, not a name for a thing.

The relationship between locality and the dynamics is the precise inverse of the conventional picture. In the conventional picture, space exists first as a manifold, and locality is the property that nearby points in space have causally connected dynamics. In the framework, the dynamics exists first as a bijection, and space emerges as the topology that locality induces on the coupling graph. Two sites are neighbors because the bijection couples them — because the state at one affects the future of the other. This is a dynamical fact, not a spatial relationship. The topology is defined by the interactions; the apparent spatial arrangement is the topology made visible.

The alphabet as gauge freedom. A subtle but important structural feature of the framework is that the alphabet size q — the number of distinguishable states at each lattice site — is gauge. Every prediction the framework makes is independent of q . The gap equation $\hbar = c^3 \epsilon^2 / (4G)$ derived in Chapter 7 contains no q . The Bekenstein-Hawking formula contains no q . The dispersion relation, area law, P-indivisibility property, center independence, and Einstein’s equations are all q -independent. Numerical verification on small systems confirms this: on a one-dimensional ring of $L = 40$ sites, the wave equation produces P-indivisibility at every q tested from 2 to 64. No experiment, even in principle, could measure q .

Two substrata (S, φ) and (S', φ') with different alphabet sizes q and q' are physically equivalent if they produce the same emergent transition probabilities. The analogy to electromagnetic gauge invariance is precise: the gauge potential A_μ is not physical, the field strength $F_{\mu\nu}$ is. In the framework, the mod- q state space at each lattice site is not physical; the coupling structure and its emergent consequences are. This is generator (ii) of the substratum gauge group $\mathcal{G}_{\text{sub}}$ developed in Chapter 2 §2.5.

A second gauge freedom — *amplitude-scale invariance* — has consequences that recur in the gauge-derivation chain. Stated operationally, it is the principle that an embedded observer has no access to the absolute scale of the field value ϕ , only to the emergent transition probabilities and coupling structure — the discrete analogue of the field-normalization redundancy of ordinary field theory, where rescaling a field is unphysical because only correlators are observable. It is realized as the additive (affine) relabellings $\phi \mapsto \lambda\phi$ of the alphabet being gauge, and is one instance of the framework’s master principle that observationally indistinguishable transformations are gauge — the same principle behind the four $\mathcal{G}_{\text{sub}}$ generators and the alphabet-*size* freedom, from which it is nonetheless *distinct*. It is an explicit adjoined operational principle of $\mathcal{G}_{\text{sub}}$ (companion papers, [Substratum §4, Theorem 24 Remark]); results that invoke it are conditional on it. The reduction of the natural commutant $U(3) \times U(2) \times U(1)$ to the Standard Model’s $SU(3) \times SU(2) \times U(1)$ uses it: the overall $U(1)$ phase of

each block can be absorbed into a global rescaling of the field value, which is gauge by amplitude-scale invariance. The Standard Model gauge group is, in this sense, what remains of the natural commutant after the gauge degrees of freedom have been stripped out.

5.3 Selection of $d = 3$

The framework’s structural lemmas operate at any integer dimension $d \geq 1$. The wave equation, the coupling-degree minimization, the cubic-group decomposition, and the gauge-structure derivation all proceed identically in dimensions other than three; they would simply produce different specific outcomes — different values of K , different multiplicities, different gauge groups, different matter content. The question of why our universe has three spatial dimensions is therefore a question about which dimension produces a self-consistent observed physics. The framework’s answer is that exactly three works.

Three independent filters converge on $d = 3$. Each is a self-consistency requirement on the emergent description, not an anthropic selection from a multiverse.

Filter 1: Propagating gravity requires $d \geq 3$. The gravitational sector developed in Chapter 7 requires propagating gravitational degrees of freedom — gravitational waves, the structural content of the LIGO observations beginning in 2015. The Weyl tensor, which captures the propagating part of the gravitational field, vanishes identically in $d \leq 2$: there is no gravitational radiation in two or fewer spatial dimensions. Jacobson’s thermodynamic derivation of Einstein’s equations from local causal horizons requires propagating gravity. The filter therefore excludes $d \leq 2$.

Filter 2: Stable matter requires $d \leq 3$. The Coulomb potential in d spatial dimensions falls off as $1/r^{d-2}$ for $d \geq 3$ and gives a logarithmic potential for $d = 2$. In $d \geq 4$ the Coulomb potential gives unstable atoms — there are no normalizable bound states of the d -dimensional Coulomb Hamiltonian, by Ehrenfest’s argument and its quantum-mechanical refinement. Gravitational orbits are similarly unstable for $d \geq 4$ by Tangherlini’s analysis. Without stable matter, no embedded observers exist to define the framework’s partition. The filter therefore excludes $d \geq 4$.

The intersection of Filter 1 and Filter 2 is already $d = 3$ uniquely. A third filter provides a consistency check.

Filter 3: Boundary-entropy concordance gives $d = 3$ exactly. The cosmological observation that the universe is spatially flat near the critical density — $\Omega_{\text{total}} \approx 1$ — is, in the framework’s gravitational sector developed in Chapter 7, a structural feature rather than an initial condition. The ratio $\rho_s/\rho_{\text{crit}}$ between the boundary-derived entropy density and the critical density, computed in d spatial dimensions, satisfies

$$\frac{\rho_s}{\rho_{\text{crit}}} = \frac{2}{d-1}.$$

This ratio equals unity exactly at $d = 3$. At $d = 2$ the ratio is 2 (overclosure); at $d = 4$ the ratio is $2/3$ (deficit, requiring an unsourced gravitational contribution). The filter selects $d = 3$ as the unique value at which the observed spatial flatness is structurally explained rather than fine-tuned.

The three filters converge: $d = 3$ is the unique value compatible with all three self-consistency requirements. Renormalizability of the emergent Yang-Mills theory at $d = 3$ — that the dimensionless gauge coupling at $3 + 1$ spacetime dimensions has zero mass dimension and the theory is consequently renormalizable — is a downstream consequence rather than an independent filter. At $d \neq 3$ the coupling-degree minimizer would give a different K , the cubic-rotation-group decomposition of $2d$ link directions would give different multiplicities, and the resulting Yang-Mills theory would have a different operator structure; asking whether *that* theory is renormalizable at *that* dimension is not a dimension-independent question.

Integer dimension is derived, not presupposed. The self-consistency argument might appear to assume that the coupling graph has a well-defined integer dimension. In fact, every bounded-degree graph falls into one of three growth classes: exponential growth ($|B(r)| \sim k^r$ for some $k > 1$), fractal growth ($|B(r)| \sim r^\alpha$ for some non-integer α), or integer-polynomial growth ($|B(r)| \sim r^d$ for some integer d). Self-consistency excludes the first two classes.

Exponential growth fails because the wave equation on exponentially-growing graphs does not produce Lorentz-invariant dispersion at low energies — the spectral gap does not close in the correct functional form — so Jacobson’s derivation of Einstein’s equations fails. Additionally, the effective dimension $d_{\text{eff}} \rightarrow \infty$ is incompatible with the stable-matter filter.

Fractal growth fails because Jacobson’s argument requires the Raychaudhuri equation, which requires geodesic congruences on a manifold structure that fractal spaces do not support. The staggered fermion decomposition that produces the framework’s three-generation structure requires a hypercubic Brillouin-zone structure $\eta \in \{0, 1\}^d$, which has no fractal-graph analog. The cubic-group decomposition of $2d$ link directions, which constrains the gauge structure, requires cubic lattice symmetry — absent in fractals.

Only integer-polynomial growth survives. By Gromov’s theorem on finitely generated groups of polynomial growth being virtually nilpotent, combined with statistical isotropy, the graph is quasi-isometric to $\mathbb{Z}^d$ for some integer d . The three filters then select $d = 3$.

Corollary. *Any bijection whose coupling graph has $d \neq 3$ produces internal contradictions in its emergent description. The coupling-graph dimension is not a free parameter or a property of a specific φ ; it is the unique value compatible with self-consistency.*

The argument’s structure is worth marking explicitly. The framework’s definition of observation requires an embedded observer as a constitutive element

— the partition $V \subset S$ is defined relative to an observer that the substratum’s structure must support. The requirement that such an observer exist, combined with the framework-internal concordance calculation, constrains d . This is not a standard anthropic argument selecting from a pre-existing multiverse of dimensions. The observer is a structural prerequisite for the framework’s definition to have content, not a contingent feature of one universe among many.

The combination of the three filters and the structural argument that integer-polynomial growth is the only consistent option produces an unusual claim: the spatial dimension is fixed structurally rather than as an input to the framework. Where the conventional picture treats $d = 3$ as an observed fact requiring no explanation, the framework derives it as a forced consequence of the structural conditions that any embedded observer’s universe must satisfy. The empirical fact that space appears three-dimensional is then evidence for the framework’s structural commitments rather than independent of them.

5.4 The emergent QFT: factorization and multi-component dynamics

The wave equation on a three-dimensional cubic lattice, established as the unique substratum dynamics in §§5.2–5.3, has further structure that the gauge-emergence chain requires. This section develops two pieces: the factorization of the wave equation into staggered Dirac fermions, which produces the fermionic matter content of the emergent quantum field theory; and the multi-component generalization of the wave equation with K internal components per lattice site, with $K = 2d = 6$ determined by coupling-degree minimization on $d = 3$.

The wave equation as lattice Klein-Gordon. The wave equation on the cubic lattice

$$\phi(\mathbf{n}, t+1) = \sum_{j=1}^d [\phi(\mathbf{n} + \hat{e}_j, t) + \phi(\mathbf{n} - \hat{e}_j, t)] - \phi(\mathbf{n}, t-1)$$

is the discrete form of the massless Klein-Gordon equation. The general second-order linear update on the lattice is $\phi(\mathbf{n}, t+1) = \alpha \phi(\mathbf{n}, t) + \beta \sum_j [\phi(\mathbf{n} + \hat{e}_j, t) + \phi(\mathbf{n} - \hat{e}_j, t)] + \gamma \phi(\mathbf{n}, t-1)$, with three coefficients α, β, γ to be fixed. Reversibility — that the dynamics is a bijection in time — forces $\gamma = -1$. Center independence — that the update at site $\mathbf{n}$ does not depend on the current value $\phi(\mathbf{n}, t)$ — forces $\alpha = 0$. Spatial isotropy combined with the maximum-speed constraint forces $\beta = 1$. The lattice d’Alembertian $\square_{\text{lat}} \phi = \sum_j [\phi(\mathbf{n} + \hat{e}_j, t) + \phi(\mathbf{n} - \hat{e}_j, t)] - \phi(\mathbf{n}, t-1) - \phi(\mathbf{n}, t+1)$ then satisfies $\square_{\text{lat}} \phi = -\alpha \phi = 0$ identically, identifying the wave equation as the massless Klein-Gordon equation. Center independence and the vanishing of the mass term are equivalent, with the substratum bijection’s lack of a center-coupling term being the structural source of the emergent massless scalar.

Factorization into staggered Dirac fermions. A theorem of Susskind from 1977 establishes that the square of the staggered Dirac operator on a hypercubic lattice equals $-\frac{1}{4}$ of the lattice d'Alembertian: $D_{\text{st}}^2 = -\frac{1}{4}\square_{\text{lat}}$. The wave equation $\square_{\text{lat}}\phi = 0$ is therefore equivalent to the massless staggered Dirac equation $D_{\text{st}}\chi = 0$ under an exact change of variables on the lattice. The bosonic field ϕ and the fermionic field χ are not two different objects; they are the same lattice content viewed through different decompositions.

The implication is non-trivial. The wave equation as substratum dynamics looks like a single scalar field evolving on the lattice, but its Susskind factorization reveals that the underlying degrees of freedom can equally be described as fermionic. The framework's matter content — the fermions of the Standard Model — is not added to the substratum as a separate input. It is implicit in the wave equation itself, made explicit through the Susskind change of variables. The structural content of the Standard Model's three generations (chirality conditional on H- χ , SM §4.8), developed in Chapter 6, traces back through the Susskind decomposition to the wave equation's specific algebraic form.

Center independence, beyond setting the mass to zero in the bosonic description, has a corresponding structural meaning in the fermionic description. The exact chiral symmetry of the emergent staggered fermions is equivalent to center independence of the substratum dynamics: any non-zero center coupling would produce a non-zero fermion mass term proportional to $\square_{\text{lat}}\phi$, breaking chiral symmetry. The framework's commitment to center independence — the structural feature that produces P-indivisibility and consequently emergent quantum mechanics, established in Chapter 1 — is the same commitment that produces exact chiral symmetry at the fermionic level. The Higgs mechanism developed in Chapter 6, which gives masses to fermions in the emergent description, operates on the chiral fermions that center independence forces.

Multi-component dynamics with K internal components. The wave equation as written above is a single-component dynamics — one scalar field per lattice site. The framework's gauge-emergence chain requires the generalization to K components per site: $\phi(\mathbf{n}, t) \in (\mathbb{Z}/q\mathbb{Z})^K$, with the matrix wave equation

$$\phi(\mathbf{n}, t+1) = M \sum_{j=1}^d [\phi(\mathbf{n} + \hat{e}_j, t) + \phi(\mathbf{n} - \hat{e}_j, t)] - \phi(\mathbf{n}, t-1)$$

where M is a $K \times K$ coupling matrix. Spatial isotropy forces M to be the same matrix for all d link directions; reversibility forces the previous-step coefficient to be $-I_K$; center independence forces the absence of a $C\phi(\mathbf{n}, t)$ term. The matrix M is the substratum's sole free parameter at this level.

Each eigenvalue μ_a of M produces a dispersion relation $\cos \omega_a = \mu_a \cos k$, giving the mass spectrum of the emergent excitations. A mode is massless if $\mu_a = 1$ and massive otherwise, with stability requiring $|\mu_a| \leq 1$. (In the revised structure the coupling matrix is scalar with $\mu = 1$: all substratum modes share one massless cone; mass and gauge-group distinctions are carried by the solution's

equivariant condensate; see the revised Theorem 5.) The global gauge group is the stabilizer of the solution’s equivariant condensate Σ in $U(K)$ — the set of $K \times K$ unitaries commuting with Σ (the coupling matrix itself is scalar; the dynamics carries an exact custodial $U(6)$). By Schur’s lemma, this stabilizer decomposes as a product of unitary groups, one on each eigenspace of M : if M has r distinct eigenvalues with multiplicities $(n_1, \dots, n_r)$, then the gauge group is $U(n_1) \times \dots \times U(n_r)$. The gauge group and the mass spectrum are encoded in the same matrix.

$K = 2d = 6$ from coupling-degree minimization. The number of internal components per lattice site is not a free parameter. The framework’s structural commitment to coupling-degree minimization — established in §5.2 as the principle that determines the natural factorization of the configuration space — applies at the per-site level too, fixing K uniquely.

The argument has three steps. First, the matrix wave equation on $\mathbb{Z}^d$ updates the K -component vector at site $\mathbf{n}$ using exactly $2d$ independent spatial inputs: the values $\phi(\mathbf{n} \pm \hat{e}_j, t)$ for $j = 1, \dots, d$. Define the internal coupling degree $\delta(K)$ as the maximum number of independent spatial input channels on which any single component’s update depends. With $K = 2d$ and the matrix M diagonal in the appropriate basis, each component depends on exactly one spatial input — one component per link direction — so $\delta(2d) = 1$. Second, with $K < 2d$, the pigeonhole principle forces at least one component to receive contributions from at least two spatial input channels, so $\delta(K) \geq 2$. Third, with $K > 2d$, the requirement that every component have at least one spatial coupling (the “observable substratum” condition: components with no spatial coupling are unobservable to the embedded observer and can be removed without changing the physics) combined with only $2d$ distinct spatial channels forces two components to share a channel — but sharing alone does not raise δ : components reading distinct entries of a shared neighbor’s vector remain dynamically independent at $\delta = 1$. What $\delta = 1$ forces is that every component attach to exactly one link direction; statistical isotropy makes the per-direction multiplicity uniform, confining the count to the family $K = 2d \cdot m$. The residual multiplicity is fixed at $m = 1$ by the gauge chain of §5.5: the emergent gauge group is the condensate stabilizer in $U(K)$, a product $\prod_i U(k_i)$ with $\sum_i k_i = K$, and the cubic multiplicities $(3, 2, 1)$ sum to six — so $K = 6$ carries the gauge sector once, while $K = 6m$ would carry m commuting copies of it.

The K with $\delta(K) = 1$ form the family $K = 2d \cdot m$; the unique member carrying a single gauge sector is $K = 2d$. On a three-dimensional cubic lattice, this gives $K = 2 \cdot 3 = 6$ internal components per site as the unique minimum-coupling factorization. The framework’s commitment to coupling-degree minimization is not an external preference for minimality; it is the mathematical definition of “elementary degree of freedom” given a bijection, just as the factorization of the configuration space into sites is the mathematical definition of “site.” A substratum with $K = 4$ or $K = 8$ would not be the same bijection with a non-optimal factorization; it would be a different bijection φ' with a different

coupling graph. The factorization is unique for a given φ .

The extension to general bounded-degree graphs with polynomial growth exponent d and statistical isotropy proceeds through Gromov's theorem: such graphs are quasi-isometric to $\mathbb{Z}^d$, and the integer-valued coupling-degree minimizer $K = 2d$ is preserved under quasi-isometry because it depends only on the spectral dimension. The factorization uniqueness is therefore a property of the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ rather than of the specific graph representative.

The combined result is that the substratum has exactly six internal components per lattice site. The structure of the coupling matrix M on this six-dimensional internal space is the next step in the gauge-emergence chain.

5.5 The cubic decomposition and the Standard Model gauge group

The matrix wave equation with $K = 6$ internal components per site has the coupling matrix M acting on a six-dimensional internal space. This section develops the structure of M from spatial isotropy: the cubic rotation group acts on the six internal components, and the requirement that the dynamics be cubic-symmetric forces M to a specific form whose eigenvalue structure produces the Standard Model gauge group.

The cubic-group representation on the six components. The six internal components at each lattice site can be naturally identified with the six nearest-neighbor link directions $\{\pm\hat{e}_1, \pm\hat{e}_2, \pm\hat{e}_3\}$ of the cubic lattice, with one component per link direction. The cubic rotation group O — the group of orientation-preserving symmetries of the cube, with twenty-four elements — acts on these six link vectors by permutation: rotations of the cube permute the six link directions among themselves, with rotations through a face center cycling four link directions and rotations through a vertex cycling three.

The action of O on the six-dimensional space spanned by the link directions is a representation of O with character $\chi(O) = 6, 0, 2, 0, 2$ at the five conjugacy classes of O (identity, eight three-fold rotations, three two-fold rotations through face centers, six two-fold rotations through edge midpoints, six four-fold rotations through face centers). Computing the inner products of χ with the characters of the five irreducible representations of O — labeled A_1, A_2, E, T_1, T_2 with dimensions 1, 1, 2, 3, 3 — yields the decomposition

$$6 = T_1 \oplus E \oplus A_1$$

with multiplicities $(3, 2, 1)$. The two remaining irreducible representations A_2 and T_2 have multiplicity zero in the six-dimensional link representation.

The decomposition is forced by group-theoretic content rather than chosen. The six link directions form a specific representation of the cubic rotation group; the decomposition of this representation into irreducibles is computed mechanically from the character table. The framework's commitment to the cubic lattice

as substratum, combined with the requirement of spatial isotropy, makes the $(3, 2, 1)$ multiplicity structure inevitable.

The coupling matrix's eigenstructure. Spatial isotropy of the matrix wave equation requires that M commute with every element of O acting on the six-dimensional internal space. By Schur's lemma, an isotropy-respecting M acts as a scalar on each irreducible component of the decomposition: $M = \text{diag}(\mu_c I_3, \mu_w I_2, \mu_y)$, where μ_c is the eigenvalue on the three-dimensional T_1 component, μ_w on the two-dimensional E component, and μ_y on the one-dimensional A_1 component. The maximum-speed constraint — that the lattice's fastest propagating mode reach the lattice speed of light — forces $\mu_y = 1$. The other two eigenvalues are restricted by stability ($|\mu_c|, |\mu_w| \leq 1$) but not otherwise forced.

The eigenvalue multiplicities of M on the six-dimensional internal space are therefore $(3, 2, 1)$ — the same multiplicities as the cubic-group decomposition. The gauge group, by the commutant theorem applied to the solution's equivariant condensate Σ (the coupling matrix itself being scalar, $M = \mu I$, forced jointly by minimal coupling and isotropy), is

$$U(3) \times U(2) \times U(1).$$

This is the stabilizer of a generic isotropy-respecting (equivariant) condensate on a cubic lattice with $K = 2d = 6$ internal components — equivalently, the unitary group of the bicommutant of the cubic-group action.

Reduction to the Standard Model gauge group. The natural commutant $U(3) \times U(2) \times U(1)$ is not quite the Standard Model gauge group $SU(3) \times SU(2) \times U(1)$. The reduction is performed by the amplitude-scale-invariance argument developed in §5.2: each of the three factors has an overall $U(1)$ phase that can be absorbed into a global rescaling of the field value at the relevant component, and global alphabet rescalings are gauge by amplitude-scale invariance of the substratum gauge group $\mathcal{G}_{\text{sub}}$. The overall $U(1)$ of each $U(n)$ factor is therefore not physical, and the commutant reduces to its $SU(n)$ subgroup for $n \geq 2$.

The reduction proceeds as follows. Writing $U(3) = SU(3) \times U(1)_{\text{color}}$, the $U(1)_{\text{color}}$ phase is the overall phase of the three-component T_1 block and is gauge by amplitude-scale invariance on those components. Similarly $U(2) = SU(2) \times U(1)_{\text{weak}}$ with the overall phase gauge. The remaining $U(1)$ block — the single A_1 component — has no SU subgroup since $SU(1)$ is trivial, but its $U(1)$ phase is physical (it controls the relative phase between the A_1 component and the higher-dimensional blocks), and identifies as the hypercharge $U(1)_Y$.

After the amplitude-scale reduction, the gauge group is

$$SU(3) \times SU(2) \times U(1)_Y.$$

This is the Standard Model gauge group, derived from the cubic-lattice substratum with no input beyond the structural commitments of Part I and the cubic-lattice realization. The three factors have physical identifications matching the

Standard Model: T_1 (the three-dimensional spatial-vector representation) is the color $SU(3)$ of quantum chromodynamics; E (the two-dimensional quadrupole representation) is the weak isospin $SU(2)$ of the electroweak theory; A_1 (the one-dimensional scalar representation) is the hypercharge $U(1)_Y$.

Bravais-lattice uniqueness. A subtle point deserves explicit treatment: among the fourteen three-dimensional Bravais lattices, the simple cubic lattice is the unique one compatible with the Standard Model gauge group under the commutant construction. The body-centered cubic lattice has eight nearest-neighbor link directions decomposing as $A_1 \oplus A_2 \oplus T_1 \oplus T_2$ under O , giving a commutant of $SU(3)^2 \times U(1)^2$ — two color factors and no weak factor. The face-centered cubic lattice has twelve nearest-neighbor link directions decomposing as $A_1 \oplus E \oplus T_1 \oplus 2T_2$, giving a commutant whose factor structure includes one $SU(3)$, additional $SU(2)$ from the doubled T_2 multiplicity, and extra $U(1)$ factors — incompatible with the Standard Model gauge group regardless of the precise factor count. The remaining eleven non-cubic Bravais lattices belong to crystal systems whose point groups have only one- and two-dimensional irreducible representations, supporting no three-dimensional irrep — no $SU(3)$ factor available under the commutant construction. The tightest form of the argument runs through the $SU(2)$ rather than the $SU(3)$: the multiplicity ladder $(3, 2, 1)$ is the irreducible-degree set $\{1, 2, 3\}$, and the octahedral group is the unique finite rotation group carrying both a genuine 2- and a 3-dimensional irrep (Appendix B.6b).

Simple cubic is therefore the unique Bravais-lattice substratum compatible with the observed Standard Model gauge group. The framework’s commitment to simple cubic, rather than to a more general bounded-degree graph satisfying the structural conditions of Chapter 1, is forced by the empirical fact that the gauge group $SU(3) \times SU(2) \times U(1)$ is observed. The further empirical content of the Cabibbo angle being $1/(\pi\sqrt{2})$ to 0.04%, developed in Chapter 6, provides quantitative confirmation: simple cubic produces the observed value, while body-centered and face-centered cubic produce values discrepant from observation by 42% and 29% respectively.

The formal uniqueness theorem, including the character-table decomposition for all 32 three-dimensional crystallographic point groups and the proof that only the five cubic-system point groups admit the required 3D irreducible representations, is developed in Appendix B.6b.

Local gauge invariance from background independence. The commutant $SU(3) \times SU(2) \times U(1)$ derived so far is a global symmetry of the matrix wave equation: the same gauge transformation applies at every lattice site. The Standard Model is, of course, a local gauge theory — the gauge transformations vary from site to site. The promotion from global to local gauge invariance is performed by the framework’s commitment to background independence.

Background independence, established in the framework’s prior content, requires the coupling structure to be dynamical rather than fixed: the coupling matrix

M depends on the state, with $M(\mathbf{n}, \hat{e}_j)$ varying from site to site as the state evolves. The matrix structure at each site has eigenvalue multiplicities $(3, 2, 1)$ because spatial isotropy still holds locally, but the specific orientation of the eigenbases differs from site to site. The commutant $U(3) \times U(2) \times U(1)$ then acts independently at each site, and the resulting gauge transformations are site-dependent.

Under a site-dependent transformation $G(\mathbf{n})$ of the form $\phi(\mathbf{n}) \rightarrow G(\mathbf{n}) \phi(\mathbf{n})$ with $G(\mathbf{n}) \in SU(3) \times SU(2) \times U(1)$ at each site, the matrix wave equation is invariant provided the coupling matrix transforms as $M(\mathbf{n}, \hat{e}_j) \rightarrow G(\mathbf{n}) M(\mathbf{n}, \hat{e}_j) G(\mathbf{n} + \hat{e}_j)^{-1}$. The link variable $M(\mathbf{n}, \hat{e}_j)$ transforms as a gauge connection, with the plaquette product $P(\mathbf{n}, j, k) = M(\mathbf{n}, \hat{e}_j) M(\mathbf{n} + \hat{e}_j, \hat{e}_k) M(\mathbf{n} + \hat{e}_k, \hat{e}_j)^{-1} M(\mathbf{n}, \hat{e}_k)^{-1}$ transforming in the adjoint as $P \rightarrow G(\mathbf{n}) P G(\mathbf{n})^{-1}$. The real part of the trace of the plaquette is gauge-invariant — recognizable as the Wilson plaquette action of lattice gauge theory.

The framework's derivation of local $SU(3) \times SU(2) \times U(1)$ gauge invariance is therefore complete. The gauge group is fixed by the cubic decomposition of the six link directions; the amplitude-scale-invariance argument reduces $U(n)$ to $SU(n)$ for $n \geq 2$; background independence promotes the global commutant to a local gauge symmetry; the resulting structure is the Wilson plaquette action, with the link variables $M(\mathbf{n}, \hat{e}_j)$ as the gauge connections. The Standard Model gauge theory is derived rather than postulated, with each step a theorem in the framework's chain.

5.6 Chirality, anomaly cancellation, and hypercharges

The gauge structure derived in §5.5 — local $SU(3) \times SU(2) \times U(1)$ from the cubic commutant and background independence — is the kinematic content of the Standard Model gauge sector. Two further structural features must hold for the emergent quantum field theory to match observation: the $SU(2)$ factor must act chirally on left-handed fermions while $SU(3)$ remains vector-like, and the hypercharges of the matter content must satisfy the six anomaly-cancellation conditions that the chiral structure imposes. This section develops both as theorems of the framework rather than as additional postulates.

Chirality from the trace-out. The Standard Model is a chiral gauge theory: weak interactions distinguish left-handed from right-handed fermions, with the gauge bosons of $SU(2)$ coupling only to left-handed fields. This is the empirical content of parity violation in beta decay, established in the 1950s. The structural origin of chirality in the framework is not a separate postulate; it emerges from the trace-out structure of the substratum partition.

The starting point is the staggered Dirac structure of the wave equation, established in §5.4. The Susskind decomposition assigns one Dirac spinor component to each hypercube corner of the cubic lattice, with the parity of the corner de-

termining the ε -grade — the product $\gamma_5 \otimes \xi_5$ of spin and taste chirality — of the assigned component. Sites with $|\eta|$ even span the $+1$ eigenspace of $\gamma_5 \otimes \xi_5$ and sites with $|\eta|$ odd the -1 eigenspace; each grade contains both spin chiralities in equal dimension (certified exactly by `chirality_grading_probes.py`). This sublattice parity $\varepsilon(\mathbf{x}) = (-1)^{\sum x_\mu}$ is the lattice-level realization of γ_5 , the chirality operator of the continuum Dirac theory.

The staggered Dirac operator has a structural property that the framework’s chirality derivation uses: it couples only across sublattices. The even-even and odd-odd blocks of D_{st} vanish identically because the staggered Dirac operator anti-commutes with the sublattice parity, $\{D_{\text{st}}, \varepsilon\} = 0$, which is the lattice form of chiral symmetry established in §5.4. The full Dirac operator has the block form

$$D_{\text{st}} = \begin{pmatrix} 0 & D_{eo} \\ D_{oe} & 0 \end{pmatrix},$$

with off-diagonal blocks coupling the two sublattice grades and diagonal blocks vanishing.

The framework’s checkerboard partition assigns even sites — the $+1$ grade of $\gamma_5 \otimes \xi_5$, carrying both spin chiralities — to the visible sector V , and odd sites — the -1 grade — to the hidden sector H . This is not an arbitrary choice; it is what the framework’s structural partition produces for a cubic lattice with bipartite coupling. The visible sector that the embedded observer reads out is therefore the left-handed sublattice; the right-handed sublattice is part of the hidden sector that is traced over.

The trace-out structure determines the grading content of the emergent quantum field theory; physical chirality enters through the taste-selectivity hypothesis H_χ (SM §4.8). Integrating out the hidden sector (the odd sublattice) via the Schur complement — a standard procedure in lattice quantum field theory — produces an effective action whose external legs lie in the visible half of the spin $\otimes$ taste space, where spin- and taste-chirality coincide:

$$\mathcal{L}_{\text{eff}} = \bar{\psi}_+ (D_{eo} G_{oo} D_{oe}) \psi_+ + \cdots,$$

where G_{oo} is the hidden-sector propagator. The hidden-grade components appear only as internal lines; the observable external states carry both spin chiralities, graded so that spin chirality equals taste chirality.

The two gauge factors couple differently to this chirality-projected effective theory. The $\text{SU}(3)$ factor acts on internal T_1 components of the cubic decomposition — the three-dimensional spatial-vector representation. Both left-handed and right-handed Dirac components carry the same color charge, since the color index is internal to the T_1 block and independent of the sublattice parity. The trace-out preserves color coupling, and $\text{SU}(3)$ remains vector-like, coupling identically to left-handed and right-handed components. The $\text{SU}(2)$ factor acts on internal E components — the two-dimensional quadrupole representation. The internal E index commutes with the grading, so the $\text{SU}(2)$ embedding as

constructed is taste-blind; its observed spin-chirality therefore requires taste-chirality selectivity of the physical embedding — the named hypothesis $H\text{-}\chi$ (SM §4.8), which the construction does not supply.

The vector-like nature of $SU(3)$ is forced by the trace-out structure applied to the cubic-lattice substratum; the chirality of $SU(2)$ rests additionally on $H\text{-}\chi$, which is open. The framework’s account of parity violation is correspondingly conditional: the substratum supplies the ε -grading and the visible-sector identity (spin chirality = taste chirality on the visible half), and the remaining step — $H\text{-}\chi$ — is stated as an open hypothesis rather than derived from which sublattice the embedded observer occupies.

Evasion of Nielsen-Ninomiya. A potential objection arises from the Nielsen-Ninomiya theorem, which forbids a single Weyl fermion on a lattice under four hypotheses: locality of the action, hermiticity of the lattice Hamiltonian, translation invariance under the lattice group, and the action being bilinear in fermionic fields carrying a definite conserved chirality charge. Under these conditions, the lattice fermion spectrum is forced to contain equal numbers of left-handed and right-handed Weyl modes — the doubling phenomenon — so a chiral gauge theory cannot be regulated on a lattice in the conventional sense.

The framework manifestly satisfies the first three hypotheses. The wave equation is nearest-neighbor with range exactly $\epsilon = 2\ell_P$ (locality); the bijection φ is real and the dynamics is T-invariant, so the induced operator is hermitian; the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ is invariant under the cubic translation group. Were the fourth hypothesis to apply, the framework would inherit the no-go conclusion and chiral $SU(2)$ would be impossible.

The fourth hypothesis does not apply, and the reason is structural rather than evasive. The Nielsen-Ninomiya theorem presupposes that the fundamental degrees of freedom are Grassmann fields $\psi(x)$ entering the action bilinearly, with chirality γ_5 defined on those fields ab initio. The framework’s fundamental object is a real scalar bijection $\varphi : S \rightarrow S$ governed by the lattice wave equation. There are no fermions at the bijection level, no Grassmann variables, no γ_5 , and consequently no $U(1)_\chi$ on which an obstruction could be formulated. Dirac structure is derived through the Susskind factorization; chirality crystallizes only at the trace-out step. By the time fermions and chirality exist as derived structures, the trace-out has already broken the bilinear-fermionic-action setting that the fourth hypothesis requires.

The framework therefore satisfies hypotheses (NN1)–(NN3) but is bosonic at the level where (NN4) requires fermionic structure. Nielsen-Ninomiya is not circumvented by a clever fermion formulation; it is bypassed because the framework never enters the hypothesis space the theorem applies to. The doublers that the no-go theorem forces upon a fermionic lattice action have a striking reinterpretation in the framework: rather than being unwanted modes that contaminate the desired theory, they are the physical fermion generations of the Standard Model — Chapter 6 develops this identification in detail.

Anomaly cancellation forcing the hypercharges. A chiral gauge theory is consistent only if its quantum anomalies cancel. The chirality of $SU(2)$ in the emergent theory means that the framework must satisfy six anomaly-cancellation conditions: three independent chiral gauge anomalies for the $SU(3)$, $SU(2)$, and $U(1)_Y$ factors taken in various combinations, plus mixed anomalies involving the gravitational sector. These conditions impose constraints on the hypercharge assignments of the chiral fermion content.

The matter content of the framework, developed in detail in Chapter 6, consists of three generations of chiral fermions in specific representations of $SU(3) \times SU(2)$: a left-handed quark doublet Q_L in $(\mathbf{3}, \mathbf{2})$, a right-handed up quark u_R in $(\mathbf{3}, \mathbf{1})$, a right-handed down quark d_R in $(\mathbf{3}, \mathbf{1})$, a left-handed lepton doublet L_L in $(\mathbf{1}, \mathbf{2})$, and a right-handed charged lepton e_R in $(\mathbf{1}, \mathbf{1})$. The framework's structural commitments fix the $SU(3) \times SU(2)$ representations of the matter content; what remains free are the hypercharge assignments Y_Q, Y_u, Y_d, Y_L, Y_e of these five representations.

The six anomaly conditions are four independent constraints on these five hypercharge unknowns. The remaining one-parameter freedom is the overall hypercharge normalization, which the framework fixes by amplitude-scale invariance — the same field-value-rescaling gauge of $\mathcal{G}_{\text{sub}}$ that reduced $U(n)$ to $SU(n)$ in §5.5. With the normalization fixed, the anomaly conditions determine the hypercharges to the observed Standard Model values, uniquely up to the interchange $Y_u \leftrightarrow Y_d$ — a relabeling of which quark is called up-type, fixed by the electric-charge identification $Q_{\text{em}} = T_3 + Y$:

$$Y_Q = \frac{1}{6}, \quad Y_u = \frac{2}{3}, \quad Y_d = -\frac{1}{3}, \quad Y_L = -\frac{1}{2}, \quad Y_e = -1.$$

The framework's derivation of the hypercharges is therefore non-trivial. The cubic-commutant construction in §5.5 produces a gauge group with three $U(1)$ factors corresponding to the three blocks (B, L, S) of the cubic decomposition. A general abelian charge is a linear combination $Q = \alpha B + \beta L + \gamma S$ of these three factors. Of all such linear combinations, the six anomaly conditions identify exactly one as anomaly-free: the linear combination that produces the observed hypercharge Y of the Standard Model. The other two combinations cannot be gauged because gauging them would introduce non-cancelling chiral anomalies. They survive as accidental global symmetries of the emergent theory, identified as baryon number B and lepton number L — broken only by the Weinberg operator and non-perturbative effects, exactly the structure observed in the Standard Model.

A corollary of the hypercharge derivation deserves explicit mention: the equality of the magnitudes of the proton and electron charges, $|q_p| = |q_e|$, follows from the framework's structural derivation rather than being a coincidence. The proton charge $q_p = 2Y_u + Y_d = 2 \cdot (2/3) + (-1/3) = 1$ and the electron charge $q_e = Y_e + (-1/2) \cdot 2 = -1$ are equal in magnitude because the anomaly conditions force the specific ratios that produce this cancellation. The famous one-to-one

ratio between proton and electron charge magnitudes — observed to fantastic precision in atomic neutrality measurements — is a theorem of the framework, not a parameter fixed by observation.

5.7 What follows

The chapter has established the gauge structure of the Standard Model from the cubic-lattice substratum. The wave equation factors into staggered Dirac fermions through the Susskind decomposition; coupling-degree minimization on $d = 3$ fixes the internal-component count at $K = 6$; the cubic-group decomposition produces eigenvalue multiplicities $(3, 2, 1)$; the amplitude-scale-invariance argument reduces the natural commutant $U(3) \times U(2) \times U(1)$ to $SU(3) \times SU(2) \times U(1)$; background independence promotes the global commutant to local gauge invariance; the trace-out yields the ε -graded visible sector on which spin- and taste-chirality coincide, so the taste-blind $SU(3)$ remains vector-like while the chirality of $SU(2)$ rests on the taste-selectivity hypothesis $H\text{-}\chi$ (open); and, given that chiral structure, the six anomaly conditions force the hypercharges to the observed values. Each step is a theorem in the framework’s chain rather than an additional postulate. The gauge group of the Standard Model and the hypercharges of its matter representations are produced from the structural inputs of Part I and the cubic-lattice realization, with the chirality of the weak sector conditional on $H\text{-}\chi$.

The chapter’s content sits at Level 0 of the four-layer framing developed in Chapter 4 — gauge structure derived from substratum-level representation theory — and is correspondingly unconditional S-class structural in the framework’s claim-classification scheme. The gauge group is not fitted to observation; it is forced — with one qualifier: the reduction of the natural commutant $U(n)$ to $SU(n)$ rests on amplitude-scale gauge invariance, an explicit adjoined principle of $\mathcal{G}_{\text{sub}}$ (§5.2), so the SU-vs-U content is forced *given that principle* rather than unconditionally; the cubic multiplicities and the anomaly-fixed hypercharges are unconditional. The hypercharges are not chosen from a continuum of possibilities; they are determined by anomaly cancellation. Falsification of either would refute the framework directly: any stable beyond-Standard-Model gauge structure, any departure of the hypercharges from the values derived here, would violate the framework’s Class A structural commitments developed in Chapter 4 §4.5.

Chapter 6 develops the specific quantitative content this gauge structure produces. The matter content is derived: three chiral fermion generations from the staggered taste structure of the wave equation, one composite Higgs doublet from the singlet staggered taste, with the absolute generation count locked at three by coupling-degree minimization jointly with the condensate-stabilizer accounting. The strong-CP problem is treated structurally, with a limit Section 6.3 makes explicit: $\bar{\theta} = 0$ at the kinematic level of the construction follows from T-invariance of the substratum bijection combined with detailed balance, requiring no axion. The gauge couplings at the Z pole follow from a

one-parameter chain whose structural inputs include the analytic one-loop staggered vacuum polarization $1/\alpha_0 = 23.25$ and the cubic-group structure of the matter content. The twenty-two classified quantitative retrodictions developed there (seven parameter-free structural; Appendix A) include the Cabibbo angle $1/(\pi\sqrt{2})$ matched to 0.04%, the Koide relation $2/9$ matched to 0.02%, the six fermion masses from a single empirical input matched to better than 1%, all three PMNS mixing angles within 1.1σ , and the Higgs mass from the boundary condition $\lambda(M_{\text{Pl}}) = 0$ matched to 0.9%. The framework’s quantitative empirical record outside cosmology is concentrated in Chapter 6.

Chapter 6

The Matter Content and Quantitative Predictions

6.1 What this chapter develops

Chapter 5 derived the gauge structure of the Standard Model from the cubic-lattice substratum. The wave equation factored into staggered Dirac fermions through the Susskind decomposition; coupling-degree minimization fixed the internal-component count at $K = 6$; the cubic-group decomposition produced eigenvalue multiplicities $(3, 2, 1)$; the amplitude-scale-invariance argument reduced the natural commutant to $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$; background independence promoted the global commutant to local gauge invariance; the trace-out yielded the ε -graded visible sector, with the taste-blind $\text{SU}(3)$ vector-like and the chirality of $\text{SU}(2)$ conditional on the taste-selectivity hypothesis $\text{H-}\chi$ (Chapter 5 §5.6); given that chiral structure, anomaly cancellation forced the observed hypercharges. The gauge sector of the Standard Model is forced — uniquely — by the structural inputs of Part I and the cubic-lattice realization, with two qualifiers: the reduction of the natural commutant to $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$ is conditional on amplitude-scale gauge invariance, an explicit adjoined principle (Chapter 5 §5.2), and the chirality of the weak sector is conditional on the taste-selectivity hypothesis $\text{H-}\chi$ (Chapter 5 §5.6).

This chapter develops the specific quantitative content that the gauge structure produces. Three pieces of structural derivation come first. Section 6.2 derives the matter content: three chiral fermion generations from the staggered taste structure of the wave equation, one composite Higgs doublet from the singlet taste, with the absolute generation count locked at three by the same argument that fixed $K = 6$ — coupling-degree minimization jointly with the condensate-stabilizer accounting. Section 6.3 treats the strong-CP problem structurally, and then shows why the treatment does not reach the Standard Model: $\bar{\theta} = 0$ at the kinematic level of the construction follows from T-invariance of the substratum

bijection combined with detailed balance, requiring no axion — a conditional result, scoped at the end of that section. Section 6.4 develops the gauge-coupling derivation and the framework’s structural no-GUT prediction, with the analytic one-loop staggered vacuum polarization producing the structural input $1/\alpha_0 = 23.25$ and the RG running to M_Z matching the three observed gauge couplings.

The remaining sections develop the framework’s parameter-free quantitative retrodictions. Section 6.5 derives the Cabibbo angle as $1/(\pi\sqrt{2})$ from a single Brillouin-zone distance — matched to the observed value at 0.04%, the framework’s sharpest geometric prediction in the quark sector. Section 6.6 derives the Koide relation as $Q = 2/3$ from a cubic-group quadratic Casimir — matched to 0.001%, the sharpest empirical agreement in the framework, with the lepton mass chain following from one empirical input through the Koide structure — $Q = 2/3$ with the derived phase $\theta_0 = 2/9$ — at 0.007%. Section 6.7 derives the three PMNS mixing angles from the cubic-group representation theory with parity-graded charged-lepton corrections — all three angles matched within 1.1σ at current precision, including the JUNO solar mixing angle $\sin^2 \theta_{12} = 1/3 - 1/(4\pi^2)$ confirmed at 0.07σ by the post-JUNO global fit. Section 6.8 collects the twenty-two predictions in a summary table classified by both derivational layer (Class column) and independent-evidence status (Cat column), and closes the chapter with a pointer forward to Chapter 7 (gravitational sector), Chapter 8 (JUNO in detail), and the framework’s broader empirical commitments.

The methodological discipline of Chapter 4 organizes the predictions. The Cabibbo angle is Class B-S (unconditional structural, Layer 1, pure substratum geometry). The Koide relation is Class B-S (Layer 1, cubic-group representation theory). The Bekenstein-Hawking $1/4$ factor — developed in Chapter 7 — is Class B-S. The gauge couplings at M_Z are Class B-R (retrodiction): three fitted parameters (δ_0, A, B) against three observed couplings produces zero residual by construction, with the structural content lying in the *form* of the fit rather than in the matching values. The six fermion masses inherit from one empirical input (the strange-quark mass) through Class B-M (mass-chain inheritance). The PMNS angles are Class B-L (layered conditional): the magnitude content of the lepton-sector mixing-matrix derivation is potentially closable through lattice perturbation theory in the emergent post-trace-out description, while the directional/assignment content is Layer 2(b) bijection-specific per Chapter 2 §2.5 (the empirical mass-hierarchy fixing which generation is electron). The framework’s empirical record in fundamental physics outside cosmology is concentrated in these classes, with each prediction’s derivational status explicit.

The chapter’s content sits at Levels 1 and 2 of the four-layer framing developed in Chapter 4. Layer 1 predictions (Cabibbo, Wolfenstein, Koide) follow from pure substratum geometry or representation theory with no operator-relation inputs; Layer 2 predictions (mass ratios within fixed operator structures, mixing-angle values) involve specific solution-level inputs. The framework’s strongest predictions — the ones that most distinguish OI from other observer-admitting

frameworks — live at Layer 1 and at the Level G3 $\times$ Level D intersection of the hierarchy developed in Chapter 3.

6.2 Three generations and the composite Higgs

The wave equation on the cubic lattice has staggered Dirac structure, as Chapter 5 established. The staggered decomposition produces $N_{\text{taste}} = 2^{\lfloor (d+1)/2 \rfloor}$ Dirac tastes in $d + 1$ spacetime dimensions, with $N_{\text{taste}} = 4$ at $d = 3$. The framework’s matter content — three chiral fermion generations plus one Higgs doublet — emerges from how these four staggered tastes decompose under the cubic rotation group and how the embedded observer’s gauge structure acts on the resulting sectors.

The taste decomposition under the cubic group. The four staggered tastes can be labeled by the corners of the Brillouin zone they originate from: the $2^d = 8$ Brillouin-zone corners on a $d = 3$ lattice pair into 4 taste pairs under the corner-conjugation $\eta \leftrightarrow \mathbf{1} - \eta$. Two corners are O-invariant under the cubic rotation group: the Γ point $(0, 0, 0)$ and the R point $(1, 1, 1)$. These two corners pair into one taste pair that transforms as the singlet A_1 representation. The remaining six corners form three axis-aligned pairs (X, M) along the three spatial directions, transforming as the triplet T_1 representation — carried by the signed vector action of the spin-taste matrices γ^j ; the unsigned pair labels alone transform only as $A_1 \oplus E$, since sign flips act trivially on the corner components. The decomposition is exact:

$$4 = \mathbf{1} \oplus \mathbf{3}.$$

One singlet taste; three triplet tastes related by cubic symmetry.

Spin-statistics from the lattice. The singlet taste produces a spin-0 field; the triplet tastes produce spin-1/2 fields. The argument runs through the staggered-to-Dirac reconstruction. The singlet pairs Γ and R in the Brillouin zone, with corresponding Dirac matrices $\Gamma(0) = I_4$ (scalar) and $\Gamma(1) = \gamma^1 \gamma^2 \gamma^3$ (pseudoscalar). Both are SO(3)-invariant, giving total angular momentum $J = 0$. The j -th triplet pair has matrices γ^j (vector) and $-i\Sigma_j$ (spin generator), giving total angular momentum $J = 1/2$.

The corollary is striking: the spin-statistics correspondence between the singlet taste (spin-0, bosonic) and the triplet tastes (spin-1/2, fermionic) follows from the lattice representation theory, not from a separate spin-statistics postulate. The Higgs (spin-0 boson) and the fermions (spin-1/2) emerge as distinct sectors of the same underlying staggered decomposition, with the cubic-group representation forcing the spin content of each sector.

Three degenerate fermion sectors. The three triplet staggered tastes provide three degenerate spin-1/2 sectors related by cubic symmetry — the structural origin of the three-generation pattern of the Standard Model. The cubic rotation group O permutes the three spatial axes, and the three triplet tastes transform among themselves under this permutation. The three sectors are

therefore degenerate at the level of cubic symmetry, with their Yukawa couplings to the Higgs producing the observed fermion mass hierarchy through generation-dependent symmetry breaking that the specific bijection φ controls.

The absolute generation count is locked at three by the same argument that fixed $K = 6$ in Chapter 5 — coupling-degree minimization jointly with the condensate-stabilizer accounting. The chain runs as follows: coupling-degree minimization on $d = 3$ confines the component count to multiples of $2d$, and the stabilizer accounting fixes $K = 2d = 6$ (Chapter 5 §5.4); the six internal components host the $2^3 = 8$ Brillouin-zone corners; the corners pair into $N_{\text{taste}} = 4$ Dirac tastes; the cubic-group decomposition gives $4 = \mathbf{1} \oplus \mathbf{3}$, with one scalar taste and three spin-1/2 tastes. The generation count is not chosen from a continuum of possibilities; it is a forced consequence of $K = 6$ on $d = 3$ — the minimality confinement together with the single-gauge-sector accounting.

Adding a fourth fermion generation would require either expanding the internal-component count beyond $K = 6$ (which violates coupling-degree minimization and is therefore a different substratum, not the same one with extra structure) or finding a fourth distinct staggered taste (which does not exist for $d = 3$). The framework’s structural commitment is therefore that no stable fourth-generation fermion with Standard Model quantum numbers can exist. This is a pre-registered falsification condition: any observation of such a fermion would refute the framework’s taste-reduction derivation at the Class A level identified in Chapter 4 §4.5.

The formal uniqueness theorem closing the T_1 generation-assignment — establishing that the assignment of three fermion generations to the T_1 triplet (rather than to $3A_1$, $A_1 \oplus E$, T_2 , or other distributions) is forced by the framework’s existing commitments — is developed in Appendix B.6c.

The matter content on links. The Standard Model matter representations — Q_L in $(\mathbf{3}, \mathbf{2})$, u_R and d_R in $(\mathbf{3}, \mathbf{1})$, L_L in $(\mathbf{1}, \mathbf{2})$, e_R in $(\mathbf{1}, \mathbf{1})$ — involve cross-charged tensor products. A left-handed quark doublet carries both color $\mathbf{3}$ and weak isospin $\mathbf{2}$ simultaneously, requiring a tensor-product representation rather than a direct-sum decomposition. The framework’s link-carrier construction provides the natural home for these representations.

The local gauge invariance established in Chapter 5 §5.5 makes the coupling matrix $M(\mathbf{n}, \hat{e}_j)$ the gauge connection on a nearest-neighbor link, transforming under local gauge transformations as $M(\mathbf{n}, \hat{e}_j) \rightarrow G(\mathbf{n}) M(\mathbf{n}, \hat{e}_j) G(\mathbf{n} + \hat{e}_j)^{-1}$. The link variable carries one gauge index at each endpoint, living in the tensor-product space $V_K^{(\mathbf{n})} \otimes V_K^{(\mathbf{n} + \hat{e}_j)^*}$ rather than in a single V_K . The matter content on links inherits this tensor-product structure: a left-handed quark doublet at a link site can have a definite color label at one endpoint and a definite weak-isospin label at the other, with no algebraic obstruction.

The framework’s content on links is therefore that matter propagates along links and the gauge representation of a matter field is read off from the tensor-product

structure of the link’s two endpoints. The three triplet tastes distribute identical matter content across three generations; the singlet taste hosts the Higgs in its E projection (the only projection that admits a gauge-invariant Yukawa coupling to the fermion doublets, since the T_1 projection is excluded on color grounds and the A_1 projection on weak grounds). The matter content of one Higgs doublet and three identical generations of $(Q_L, u_R, d_R, L_L, e_R)$ — exactly the Standard Model matter content — is the forced consequence.

The Higgs as a composite scalar. A structural feature of the framework’s Higgs sector deserves explicit treatment. The Higgs is not a fundamental scalar field added to the framework as a separate input. It is a composite excitation in the singlet staggered taste of the wave equation — the E projection of the $A_1 \oplus E \oplus T_1$ decomposition under the cubic group. The Higgs’s gauge representation $(\mathbf{1}, \mathbf{2}, +1/2)$ follows from the E projection’s weak-isospin content combined with the unique anomaly-free hypercharge assignment.

The composite nature of the Higgs has implications for the hierarchy problem. In the standard formulation, the smallness of the Higgs mass compared to the Planck scale is mysterious: radiative corrections to a fundamental scalar’s mass should generically push it toward the highest available scale, and the observed $m_H \sim 125$ GeV requires a tuning of one part in 10^{34} . In the framework, the Higgs mass is set by the eigenvalues of the coupling matrix M — properties of the substratum bijection φ — rather than by radiative corrections. The emergent quantum field theory’s notion of “radiative correction” applies only to the perturbative expansion around the substratum’s solution; the substratum itself fixes the Higgs mass at the value the bijection produces, with no separate fine-tuning required.

Whether this constitutes a *solution* to the hierarchy problem in the conventional sense depends on how the substratum’s choice of φ is itself viewed. The framework treats the specific bijection φ within the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ as a solution-specific property — what Chapter 4 called Layer 3 content, depending on the specific representative within the gauge class. The hierarchy is dissolved at the level of the framework’s emergent description (no radiative tuning required at the emergent level) and reframed at the substratum level (the bijection has whatever Higgs-mass eigenvalue it has). The reframing is a methodological choice consistent with the framework’s general pattern: structural commitments at the level of the equivalence class; solution-specific inputs at the level of the specific representative.

6.3 The strong-CP problem and $\bar{\theta} = 0$

The strong-CP problem is one of the Standard Model’s most persistent puzzles. Quantum chromodynamics admits a CP-violating term $\mathcal{L}_\theta = (\theta/32\pi^2) \text{Tr}(F\tilde{F})$ where F is the gluon field-strength tensor. This term would produce CP-violation in strong interactions, observable as a non-zero neutron electric dipole moment. Experimental bounds limit θ to less than approximately 10^{-10} — a

remarkable smallness for a parameter the Standard Model treats as free, ranging in principle over $[0, 2\pi]$. The conventional resolution — the Peccei-Quinn mechanism with an associated axion — postulates a new global symmetry whose spontaneous breaking dynamically drives θ to zero. No axion has been observed at expected scales.

The framework’s resolution is structural rather than dynamical. The combination $\bar{\theta} = \theta + \arg \det(Y_u Y_d)$ — the physically meaningful CP-violating parameter, invariant under field redefinitions — is zero at the kinematic level of the construction. The end of this section shows that the same symmetry argument, carried through, also removes the CP violation observed in quark mixing — so the result does not stand as a statement about the Standard Model. The result follows from three structural inputs already established: T-invariance of the substratum bijection, the spatial nature of the framework’s partition, and detailed balance applied to the resulting transition probabilities.

T-invariance of the wave equation. The first input is that the discrete wave equation is invariant under time reversal. Under the transformation $\phi(\mathbf{n}, t) \rightarrow \phi(\mathbf{n}, -t)$, the wave equation $\phi(\mathbf{n}, t+1) = \sum_j [\phi(\mathbf{n} + \hat{e}_j, t) + \phi(\mathbf{n} - \hat{e}_j, t)] - \phi(\mathbf{n}, t-1)$ maps to itself: the time-reversed field satisfies the same equation. T-invariance is a structural property of the wave equation, not an additional postulate. The substratum dynamics is symmetric under time reversal in a precise mathematical sense.

The framework’s partition — the visible/hidden checkerboard developed in Chapter 5 — marginalizes over spatial sites, not over temporal data. The partition is spatial, preserving the temporal structure of the substratum dynamics. T-invariance of the substratum implies T-invariance of the visible-sector marginalization: the time-reversal operation commutes with the partition’s spatial projection.

Vanishing of the QCD θ . The θ -term $\mathcal{L}_\theta$ is T-odd. The trace $\text{Tr}(F\tilde{F})$ equals $2\mathbf{E} \cdot \mathbf{B}$ when expressed in terms of the chromoelectric and chromomagnetic fields, and under time reversal $\mathbf{E} \rightarrow \mathbf{E}$ while $\mathbf{B} \rightarrow -\mathbf{B}$. The combination is therefore T-odd. A T-invariant Lagrangian cannot contain a T-odd term: the dynamics’s symmetry under time reversal forbids it structurally. The emergent Lagrangian of the visible sector inherits T-invariance from the substratum because the partition is spatial. Therefore $\theta = 0$ in the emergent Lagrangian.

Detailed balance. The second input is detailed balance of the visible-sector transition probabilities. The argument runs through a counting lemma: for each hidden state h that maps a visible initial state i to a visible final state j under the bijection φ , there is a corresponding hidden state h' that maps j back to i under φ^{-1} . Since φ is a bijection and T-invariance gives $\varphi^{-1} = T \circ \varphi \circ T$, and since the partition is spatial (preserving the visible/hidden classification under time reversal), the map $h \mapsto h'$ is an injection between two sets of hidden states. The bijection property of φ on the finite configuration space makes this injection a bijection between sets of equal cardinality. Summing over hidden

states and dividing by the hidden-sector cardinality gives equality of transition probabilities: $T_{ij} = T_{ji}$ for all visible states i, j .

Detailed balance — the structural equality $T_{ij} = T_{ji}$ — is therefore not an emergent thermodynamic property requiring equilibration. It is a direct consequence of the substratum’s bijection structure combined with T-invariance and spatial partitioning. The visible-sector transition probabilities respect detailed balance at the level of the underlying bijection, before any thermodynamic limit is taken.

T-invariant emergent Hamiltonian. The third input is that detailed balance implies the emergent Hamiltonian is T-invariant. The Born rule from Chapter 1 — $T_{ij}(t) = |U_{ij}(t)|^2$ via the Barandes correspondence — combined with detailed balance gives $|U_{ij}(t)|^2 = |U_{ji}(t)|^2$ for all i, j, t . This symmetry of the transition amplitudes implies, by Wigner’s theorem, that the unitary evolution satisfies $U(t) = \hat{\Theta} U(t)^* \hat{\Theta}^{-1}$ for some anti-unitary operator $\hat{\Theta}$. The corresponding emergent Hamiltonian commutes with $\hat{\Theta}$: $[\hat{H}_{\text{eff}}, \hat{\Theta}] = 0$. The emergent Hamiltonian is therefore T-invariant.

$\bar{\theta} = 0$ at every scale of the construction. Combining the three inputs gives the full result. Step 1: since φ is a bijection, φ^n is a bijection for every n , and the counting argument applies to φ^n as well as to φ itself. Detailed balance therefore holds at every time scale: $T_{ij}(n) = T_{ji}(n)$ for all n . This has been verified numerically on finite lattices: T-violation is exactly zero for ring sizes $L = 4, 6$ and alphabet sizes $q = 2, 3, 5$, at all times $n = 1, \dots, 12$. Step 2: by Wigner’s theorem applied to each time scale, the effective Hamiltonian $\hat{H}_{\text{eff}}(n)$ is T-invariant at every time scale. Step 3: T-invariant Yukawa couplings at each scale are simultaneously real in an appropriate basis. Real Yukawa matrices have $\arg \det(Y_u Y_d) = 0$. Combined with $\theta = 0$ from the T-invariance of the emergent Lagrangian, this gives $\bar{\theta} = 0$ at every scale of the constructed emergent description.

Within the construction the result is exact, and it bypasses the instanton question that complicates the conventional analysis. T-invariance of the transition probabilities is a consequence of the bijection structure, holding at every time scale without perturbative or non-perturbative approximation. The renormalization-group flow cannot generate T-violation because the underlying bijection structure forbids it at every scale of the construction. The strong-CP problem is addressed at the kinematic level rather than the dynamical level, and within the construction no axion is required. Whether that carries to the Standard Model is the question the rest of this section takes up.

Falsification and scope. Within the construction the neutron electric dipole moment d_n is exactly zero, with the current experimental bound $|d_n| < 1.8 \times 10^{-26} e \cdot \text{cm}$. Any observation of a non-zero d_n consistent with $\bar{\theta} \gtrsim 10^{-11}$, at sufficient confidence to exclude experimental systematics, would falsify the framework’s T-symmetric substrate commitment — a Class A falsification per Chapter 4 §4.5, refuting the framework directly with no recoverable error mode

beyond computational mistakes in the specific derivation.

The scope of the $\bar{\theta} = 0$ result is worth marking explicitly. The theorem establishes $\bar{\theta} = 0$ — the static CP-odd combination that couples to QCD instantons and would generate a neutron EDM. It does not forbid the CP-violating phases observed in CKM and PMNS mixing. The CKM CP violation, encoded in the Jarlskog invariant $J \approx 3 \times 10^{-5}$, arises from the basis mismatch between up-type and down-type Yukawa mass eigenstates rather than from a non-zero $\bar{\theta}$. Reality in each matrix’s own mass basis is automatic for any matrix whatsoever and constrains nothing. What the substrate symmetry delivers is reality in one shared basis, and that forces the connecting unitary to be real, so no irreducible phase survives and the Jarlskog invariant vanishes. The construction as it stands therefore predicts no CKM CP violation, which is in tension with observation; the constraint comes from the premise chain rather than from $\bar{\theta}$ itself [SM §5.5]. The framework’s structural commitment is to $\bar{\theta} = 0$; the basis-alignment parameters that produce the observed CKM CP phase are solution-specific properties of the particular bijection φ , treated as inputs rather than predictions. Chapter 4 §4.6 listed these flavor-sector CP phases among the solution-specific properties whose measurement does not bear on the framework’s structural status.

6.4 Gauge couplings and the no-GUT prediction

The Standard Model has three gauge couplings: α_3 for color, α_2 for weak isospin, α_1 for hypercharge, measured at the Z boson mass. The values $\alpha_3(M_Z) \approx 0.1181$, $\alpha_2(M_Z) \approx 0.0339$, $\alpha_1(M_Z) \approx 0.0169$ are taken as independent inputs in the Standard Model. The framework’s derivation of the gauge couplings extends the chain through a structural input — the one-loop staggered vacuum polarization — combined with threshold parameters that match the three couplings simultaneously. The result is a structural account of *why* the gauge couplings have approximately the values they have, including the framework’s distinctive structural prediction of no grand unification.

The induced coupling from the fermion determinant. The gauge field in the framework is not an independent dynamical degree of freedom. It is the state-dependent coupling matrix $M(\mathbf{n}, \hat{e}_j)$ that Chapter 5 §5.5 identified as the gauge connection on a nearest-neighbor link. The matter content — the staggered fermion field ϕ on the lattice — determines the coupling matrix through the wave equation. The gauge kinetic term is therefore *induced* by the fermion determinant rather than being a separate input.

The effective action for the gauge connection, obtained by integrating out the fermion degrees of freedom, is

$$S_{\text{eff}}[U] = -N_f \text{Tr} \ln(D^\dagger D[U]),$$

where N_f is the number of fermion species and D is the staggered Dirac operator coupled to the gauge connection U . Expanding S_{eff} to second order in the gauge

field gives the gauge propagator and the induced gauge coupling, computed at one loop from the staggered vacuum polarization $\Pi_s(0)$ — a lattice momentum integral over the fermion bubble with the framework’s specific vertex structure.

The induced coupling has a universal structural value. The Dynkin index $T(R) = 1/2$ per Dirac fermion is the same for all three gauge factors because every component of the coupling matrix M — with $K = 6$ components distributed across the $(3, 2, 1)$ cubic decomposition — transforms in the fundamental representation of its respective gauge group. Combined with $N_f = 6$ (one fermion species per internal component, fixed by the $K = 2d = 6$ result of Chapter 5 §5.4), the induced inverse coupling at the cutoff scale M_{Pl} is

$$\frac{1}{\alpha_0} = 6 \times \Pi_s(0) \times 4\pi = 23.25,$$

with $\Pi_s(0) \rightarrow 0.3084$ on a robust numerical plateau as the lattice integral converges. This is a strictly structural prediction — depending only on N_f and $T(R)$, both determined by the lattice structure — and is universal across the three gauge factors because the Dynkin index equality $T_3(R) = T_2(R) = T_1(R)$ holds at the substratum level.

The threshold structure. The universal coupling α_0 at the cutoff M_{Pl} does not yet match the observed couplings at M_Z . Two corrections enter at the matching between the substratum and the emergent quantum field theory: a universal threshold δ_0 from higher-loop corrections to the vacuum polarization, and a group-dependent threshold from gauge self-interactions proportional to the quadratic Casimir C_2 of each gauge factor.

The expansion parameter for the group-dependent threshold is $C_2 \cdot g_0^2 = C_2 \times 4\pi\alpha_0$, which equals 1.08 for SU(2) ($C_2 = 2$) and 1.62 for SU(3) ($C_2 = 3$). Both are of order unity at the cutoff, so finite-order perturbation theory converges too slowly to capture the threshold from first principles. The framework’s account is that the threshold is given by an all-orders geometric-series resummation of the gauge self-energy, with the form

$$\frac{1}{\alpha_i(M_Z)} = \underbrace{23.25}_{1/\alpha_0} + \underbrace{10.0}_{\delta_0} + \underbrace{A \cdot \ln(1 + B \cdot C_2 g_0^2)}_{\text{resummed gauge self-energy}} + \underbrace{\frac{b_i}{2\pi} \ln \frac{M_{\text{Pl}}}{M_Z}}_{\text{SM RG running}},$$

where A and B are the threshold parameters and b_i is the standard Standard Model beta-function coefficient for each gauge group.

The four components have distinct derivational statuses. The structural input $1/\alpha_0 = 23.25$ is computed analytically from the one-loop staggered vacuum polarization, with no fitted parameters. The universal threshold $\delta_0 = 10.0$ is empirically fixed by the U(1) row of the matching equations; a two-loop staggered VP rebuild indicates that strict two-loop perturbation theory does not reproduce δ_0 , so it is honestly treated as an empirical input rather than a derived value. The C_2 -dependent threshold parameters $(A, B) = (8.3, 5.59)$ are

matched to the SU(2) and SU(3) couplings at M_Z , with the product $A \cdot B = 46.4$ independently computed at $A \cdot B \approx 48\text{--}62$ depending on the assumed leading finite-size power — a cross-check that agrees with the fit under the $1/N^2$ assumption and is inconclusive across finite-size models. The Standard Model RG running below M_{Pl} is standard.

The classification per Chapter 4 §4.5 is therefore that the three gauge couplings at M_Z are Class B-R retrodictions: three fitted parameters (δ_0, A, B) against three observed couplings produces zero residual by construction. The structural content of the framework’s gauge-sector account does not lie in matching the three coupling values — that is achieved by the fit. It lies in three structural predictions: (i) the specific value $1/\alpha_0 = 23.25$ from the one-loop staggered VP; (ii) the universality of δ_0 across the three gauge groups (no group-dependent term); (iii) the independent derivation of $A \cdot B$ agreeing with the fitted value at the few-percent level.

The structural no-GUT prediction. The framework’s most distinctive content in the gauge-coupling sector is structural rather than quantitative: the three gauge couplings of the Standard Model do not unify at a single energy scale, and no grand unified theory underlies the Standard Model gauge group. The prediction has three layers.

At the group-theoretic level, the Standard Model gauge group $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$ emerges directly from the cubic decomposition of the six nearest-neighbor link directions — a single-step decomposition with three eigenspaces of the coupling matrix M and three commutant factors. There is no intermediate large gauge group at any scale: no SU(5), no SO(10), no E_6 , no Pati-Salam $\text{SU}(4) \times \text{SU}(2)_L \times \text{SU}(2)_R$. The framework has no GUT stage, structurally. This contrasts with all standard string-phenomenology routes to the Standard Model, each of which traverses an intermediate large gauge group during compactification.

At the coupling level, the three gauge couplings do not unify at any energy scale. The universal $1/\alpha_0 = 23.25$ at the cutoff is shared across all three gauge factors before gauge self-interactions are included; but the C_2 -dependent threshold $A \ln(1 + B \cdot C_2 g_0^2)$ splits the couplings *at the cutoff itself*, with C_2 different across the three gauge groups. Below M_{Pl} , the standard Standard Model beta functions run the already-split couplings to lower scales, with no intermediate gauge structure modifying the running. The three gauge couplings therefore remain distinct at all scales accessible to the emergent quantum field theory; they were never unified, even at the Planck scale.

At the mechanism level, the gauge couplings emerge from the fermion determinant — a specific mechanism that string compactifications do not share. String-theoretic gauge couplings depend on moduli (compactification volumes, dilaton vacuum expectation values, flux choices) and are typically not related to one another structurally. The framework’s gauge couplings, by contrast, share an origin: they all arise from the same one-loop staggered vacuum polarization with

the same N_f and $T(R)$, with the differences among them attributable to the C_2 -dependent threshold structure. The three couplings are not three independent observables but three projections of a common substructure.

The structural no-GUT prediction has empirical consequences. The most direct is proton stability: the framework predicts no proton decay through GUT-mediated channels, with the only available baryon-number-violating operators suppressed by M_{Pl}^2 rather than M_{GUT}^2 . The expected proton lifetime in the framework is $\tau_p \sim M_{\text{Pl}}^4/m_p^5 \sim 10^{45}\text{--}10^{46}$ years, roughly ten orders of magnitude beyond Hyper-Kamiokande’s projected $\sim 10^{35}$ year sensitivity. Any detection of proton decay at experimentally accessible scales would refute the framework’s no-GUT commitment. Tightening of the proton lifetime bound past 10^{35} years would support the framework’s prediction. Two further consequences: no GUT-chain magnetic monopoles (so the cosmological monopole problem that inflation was historically invoked to dissolve does not arise in the framework), and no GUT-chain cosmic strings from $\text{SO}(10) \rightarrow \text{SM}$ symmetry breaking with non-trivial π_1 . Standard cosmic-string non-detection in CMB and gravitational-wave observations is consistent with the framework’s prediction.

The framework-specific discriminator within the no-GUT class is the shared origin of the three gauge couplings: the structural prediction $1/\alpha_0 = 23.25$ as a single value shared across $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$ before threshold corrections. Other no-GUT frameworks would generically have three independent gauge couplings without this shared-origin structural relation. The framework’s matching of $\alpha_3, \alpha_2, \alpha_1$ at M_Z through the universal-induced-coupling chain confirms this structural prediction at the few-percent level.

6.5 The Cabibbo angle and CKM structure

The Cabibbo–Kobayashi–Maskawa (CKM) matrix parametrizes the misalignment between the mass eigenstates and the weak interaction eigenstates of the down-type quarks. Its leading parameter, the Cabibbo angle $\lambda \approx 0.225$, controls the strength of generation-changing weak transitions. In the Standard Model the Cabibbo angle is a free parameter measured to be $\lambda = 0.22500 \pm 0.00067$. The framework derives λ from the geometry of the cubic lattice with no fitted parameters, matching the observed value to 0.04% (0.12σ) — the framework’s sharpest geometric prediction in the quark sector.

The Cabibbo angle from a Brillouin-zone distance. The three fermion generations of the framework occupy the three triplet (T_1) staggered tastes, located at the three pairs of axis-aligned corners of the Brillouin zone. In a normalized convention where the Brillouin zone has dimension 2π , the three triplet corners $X_j = \pi \hat{e}_j$ sit at positions $(\pi, 0, 0)$, $(0, \pi, 0)$, $(0, 0, \pi)$ along the three coordinate axes. The distance between two adjacent triplet corners is

$$|X_i - X_j| = |\pi \hat{e}_i - \pi \hat{e}_j| = \pi\sqrt{2},$$

a geometric constant of the cubic lattice.

The Cabibbo angle equals the inverse of this distance:

$$\lambda = \frac{1}{\pi\sqrt{2}} = 0.22508.$$

Observed value: $\lambda = 0.22500 \pm 0.00067$. The framework's prediction matches to 0.04% — well inside experimental uncertainty.

The mixing matrix element between generations i and j is the continuum fermion propagator evaluated at the inter-generation momentum:

$$|M_{ij}| = |S(X_j - X_i)| = \frac{1}{|X_j - X_i|} = \frac{1}{\pi\sqrt{2}}.$$

The propagator has $1/|q|$ falloff rather than $1/|q|^2$ because the taste-changing transition preserves chirality, established in the staggered-lattice literature: taste-changing vertices in staggered quarks have spinor structure given by a combination of γ_μ and $\gamma_{5\mu} = \gamma_5\gamma_\mu$, both of which map left-handed to left-handed and right-handed to right-handed rather than flipping between them. In the massless limit the chirality-preserving propagator structure contracted with the chirality-preserving vertex gives $|M_{ij}| \propto 1/|q|$. The corresponding chirality-flipping process would give $1/|q|^2$ falloff, which produces the Gatto-Sartori-Tonin relation $\lambda^2 = m_d/m_s$ that is discussed below.

Comparison with other Bravais lattices. The Cabibbo angle on simple cubic is $\lambda_{\text{SC}} = 1/(\pi\sqrt{2}) = 0.2251$, matching observation at 0.04%. On body-centered cubic, the triplet sector sits at the H -points of the BCC Brillouin zone with separation $\pi\sqrt{6}$, giving $\lambda_{\text{BCC}} = 0.1299$ — discrepant from observation by 42%. On face-centered cubic, the triplet sector sits at the L -points with separation 2π , giving $\lambda_{\text{FCC}} = 0.1592$ — discrepant by 29%. The simple-cubic Bravais-lattice commitment from Chapter 5 §5.5 is quantitatively confirmed by the Cabibbo prediction: only simple cubic produces the observed value.

Isolation in integer-input space. A complementary robustness check examines whether nearby alternatives to the specific integer $N = 2$ in the form $\lambda = 1/(\pi\sqrt{N})$ also match observation. They do not. The table below shows predictions for $N \in \{1, 2, 3, 4, 6, 8\}$ — the integers that arise naturally from inter-corner distances in various cubic-lattice geometries — against the observed $\lambda = 0.22500 \pm 0.00067$:

N	$\lambda = 1/(\pi\sqrt{N})$	Deviation
1	0.3183	+139 σ
2 (framework)	0.2251	+0.1 σ
3	0.1838	−62 σ
4	0.1592	−98 σ
6	0.1299	−142 σ
8	0.1125	−186 σ

The framework's $N = 2$ is the only integer value in the nearby alternatives that matches observation within 1σ . The next-closest alternative ($N = 3$) deviates by 62 standard deviations. The Cabibbo prediction is *isolated* in the space of nearby alternatives — not selected from a cluster of acceptable choices.

The Wolfenstein parameter A . The A_1 (Higgs) taste sits along the democratic direction $\hat{h} = (1, 1, 1)/\sqrt{3}$ in the internal space. Each generation axis $\hat{e}_j$ makes an angle θ with $\hat{h}$ satisfying $\cos \theta = \hat{e}_j \cdot \hat{h} = 1/\sqrt{3}$. The CKM mixing in the A parameter is driven by the perpendicular component of the generation axis relative to the Higgs democratic direction:

$$A = \sin \theta = \sqrt{1 - 1/3} = \sqrt{2/3} = 0.8165.$$

Observed: $A = 0.826 \pm 0.012$. Match: 0.8σ , agreement at 1.2%. The observed value itself is less settled than the error bar suggests: different global fits place it between roughly 0.79 and 0.83, and the prediction sits inside that range. The geometric value is exact, but its two uses differ in status: in the reactor angle $\sin \theta_{13} = A^2 \lambda$ it enters protected (the second-order correction cancels, so that comparison tests the geometry alone), while in $|V_{cb}| = A\lambda^2$ it enters through a second-order coefficient whose derivation is open, so the identification with the measured A is leading-order; the observed 1.2% deviation is smaller than the natural $\lambda^2 \approx 5\%$ correction scale, a suppression that is observed rather than derived.

The combination $|V_{cb}| = A\lambda^2 = \sqrt{2/3}/(2\pi^2) = 0.04136$ (observed: 0.0408 ± 0.0014 , agreement at 0.4σ) provides a cross-check at the second-order level of the Wolfenstein expansion. Two structural predictions (λ and A) generate a derived prediction for a third observable ($|V_{cb}|$), with all three matching observation.

A structural identity. A non-trivial representation-theoretic identity underlies the Wolfenstein A parameter. The off-democratic projection-squared $A^2 = 1 - (\hat{e}_j \cdot \hat{h})^2 = 1 - 1/d$ for the generation axes in a d -dimensional space, combined with the cubic-group quadratic Casimir for the vector representation $C_2(T_1) = d - 1$, gives the structural identity

$$A^2 = \frac{d-1}{d} = \frac{C_2(T_1)}{d}.$$

For $d = 3$ this gives $A^2 = 2/3$ directly. The same Casimir-over-dimension structure underlies the Koide prediction in §6.6 below and the PMNS angle $\sin^2 \theta_{13}$ in §6.7. The identity is not a numerical coincidence; it is the structural source of multiple Layer-1 predictions in the framework.

The down-to-strange mass ratio. The Cabibbo derivation extends to mass ratios through the Gatto-Sartori-Tonin (GST) relation $\sin \theta_C \approx \sqrt{m_d/m_s}$, a standard relation in flavor physics with known higher-order corrections. The framework's prediction is structural in λ — the cubic-lattice derivation above — while the conversion to the mass ratio inherits GST's empirical status. To

leading order:

$$\frac{m_d}{m_s} = \lambda^2 = \frac{1}{2\pi^2} = 0.05066.$$

Observed (FLAG 2024 $N_f = 2+1+1$ lattice): $m_d/m_s = 0.0503 \pm 0.0007$. Match: 0.56σ , agreement at 0.7%.

The Jarlskog invariant. The Jarlskog invariant $J = \text{Im}(V_{us}V_{cb}V_{ub}^*V_{cs}^*)$ measures CP violation in the quark sector. In the Wolfenstein parametrization, $J \approx A^2\lambda^6\eta$. Substituting $\lambda = 1/(\pi\sqrt{2})$ and $A = \sqrt{2/3}$:

$$J = \frac{\eta}{12\pi^6},$$

where $\eta = 0.348$ is the solution-specific CP-violating parameter. This gives $J = 3.02 \times 10^{-5}$ (observed: $(3.08 \pm 0.13) \times 10^{-5}$, agreement at 0.5σ). The structural suppression factor $12\pi^6 \approx 11,537$ is purely geometric.

A consistency check using the framework's $A^2 = 2/3$ and $\lambda^2 = 1/(2\pi^2)$ structural inputs runs the relation in reverse: solving $J = A^2\lambda^6\eta$ for η at the observed Jarlskog value $J_{\text{obs}} = 3.18 \times 10^{-5}$ gives $\eta_{\text{implied}} = 0.367$. The PDG value is $\eta = 0.357 \pm 0.011$. The framework's structural inputs combined with the observed Jarlskog produce an η consistent with the PDG fit at 0.83σ . This consistency check is independent of the chapter's specific solution-specific $\eta = 0.348$ — it asks only whether the framework's structural A^2 and λ^2 inputs are compatible with the observed CP-violating phase, and the answer is yes within 1σ .

The Wolfenstein parameters ρ and η themselves require complex CKM entries, arising from the specific bijection φ rather than from structural cubic-group representation theory. They are solution-specific properties of the framework's particular substratum representative — Layer 3 content in the four-layer framing of Chapter 4 — and the framework makes no commitment to their numerical values beyond the structural constraint that the Jarlskog invariant has the form $\eta/(12\pi^6)$.

6.6 The Koide relation and the lepton mass chain

The Koide relation is a numerical regularity among the three charged-lepton masses observed in 1981:

$$Q = \frac{m_e + m_\mu + m_\tau}{(\sqrt{m_e} + \sqrt{m_\mu} + \sqrt{m_\tau})^2} = \frac{2}{3}$$

to extraordinary precision — matching observation at 0.001% with no fitted parameters in the conventional Standard Model. The relation has been a persistent puzzle, with multiple proposed derivations from grand-unified, supersymmetric, and modular-symmetry frameworks. None has produced both the value $2/3$ and an account of why the relation should hold in the first place. The framework derives the Koide relation from the cubic-group representation theory of the

three generations, with both the value and the underlying structure following from substratum-level content.

The Koide angle from a cubic Casimir. The three lepton generations occupy the three triplet (T_1) staggered tastes, related by the cubic symmetry group O . The Koide relation in its angular form $\theta_0 = C_2/d^2$ — where C_2 is the quadratic Casimir of the relevant cubic-group representation and d is the spatial dimension — provides a specific structural prediction for the value θ_0 .

For the T_1 representation of O (the three-dimensional vector representation occupied by the three lepton generations), the quadratic Casimir is $C_2(T_1) = d - 1 = 2$ in $d = 3$ spatial dimensions. The lattice bandwidth squared is $d^2 = 9$. The Koide angle in its natural framework form is therefore

$$\theta_0 = \frac{C_2(T_1)}{d^2} = \frac{2}{9},$$

which is the angular form of the $Q = 2/3$ relation under the standard $\mathbb{Z}_3$ -symmetric parametrization of three masses by an angle. The framework's prediction matches the observed value $\theta_0^{\text{obs}} = 0.222220$ at 0.02% — the sharpest empirical agreement in the framework.

The structural interpretation is geometric: $C_2(T_1)/d^2$ is the anisotropy strength of the generation representation, normalized by the lattice bandwidth squared. The cubic-group representation theory of T_1 fixes the numerator; the dimensional structure of the lattice fixes the denominator. The value $2/9$ follows from these two pieces of substratum-level structure, with no fitted parameters.

Uniqueness among alternative ratios. Among dimensionless ratios involving cubic-group Casimirs and lattice-dimension quantities at $d = 3$ and T_1 , the value $\theta_0 = C_2/d^2 = 2/9$ is the unique match to observation. Ten alternative ratios — including C_2/d (which gives the Wolfenstein A^2 above), $C_2/(2d)$, $\sqrt{C_2}/d$, $(C_2 - 1)/d^2$, and similar combinations — produce values discrepant from the observed Koide angle by amounts ranging from 11% to over 200%. The match at $\theta_0 = C_2/d^2$ is not one of several possibilities consistent with cubic-group structure; it is the unique value forced by the relevant Casimir-and-bandwidth combination.

Connection to the Cabibbo identity. The Koide angle is related to the Cabibbo structural identity through second-order perturbation theory. The natural cubic-lattice perturbation amplitude on the T_1 representation is $\sqrt{C_2}/d = \sqrt{2}/3$ — the Casimir's square root scaled by the bandwidth. Second-order perturbation theory on the eigenvalues then gives $\theta_0 \sim (\sqrt{C_2}/d)^2 = C_2/d^2 = A^2/d$, with $A^2 = (d - 1)/d$ the structural identity of §6.5 for vector representations. Amplitudes scale as $\sqrt{C_2}/d$; eigenvalue shifts scale as the square. The Koide angle and the Wolfenstein A parameter share a common Casimir structure, with the Koide angle obtaining at the eigenvalue level what A obtains at the amplitude level.

The first-principles derivation of the specific $1/d$ amplitude scaling (rather than alternatives like $1/(2d)$) from the cubic-lattice Yukawa structure is now complete: the amplitude prefactor has been computed and equals one (SM §7.2; driver `oi_lattice_code/fermion/n_ratio_pt.py`), so $\mathcal{N} = 1/d^2 = 1/9$ exactly at $m \rightarrow 0$, and the empirical match at 0.02% confirms a derived normalization rather than constraining a chosen one. The Koide prediction is currently Class B-S Layer 1 — unconditional structural at the substratum-geometry level — with the former open question (whether the $1/d$ amplitude is forced by the framework’s structural commitments or chosen empirically) now settled on the forced side by that computation. The choice reduces to whether the taste-mixing loop carries a time direction ($d = 4$, giving $1/16$) or not ($d = 3$, giving $1/9$); the $d = 3$ substrate-eigenvalue origin of the generation masses selects $1/9$; Schur’s lemma on the T_1 generations protects this against all taste-blind emergent dressing. A cubic-invariant dressing need not be taste-blind: it may act on $\text{Sym}^2(T_1)$ with different scalars on A_1 , E and T_2 , as the one-loop staggered taste kernel does ($\lambda_E/\lambda_{A_1} \approx 0.22\text{--}0.24$). That class is bounded: it leaves θ_0 invariant at the $\sqrt{m}$ level, and on the mass pattern the observed agreement constrains it to $\delta m/m \lesssim 8 \times 10^{-5}$, which a lepton-Yukawa loop satisfies (SM §7.2). What remains is a possible loop-level cubic-breaking contribution — the same residual as the emergent-Lorentz radiative-stability question (SM §7.3).

The charged-lepton mass chain. Given the Koide angle $\theta_0 = 2/9$ and one empirical input (the tau mass $m_\tau = 1776.86$ MeV), the parametrization $\sqrt{m_k} = \mu(1 + A \cos(\theta_0 + 2\pi k/3))$ is fully determined: $Q = 2/3$ fixes $A = \sqrt{2}$, θ_0 fixes the phase, and m_τ fixes the scale μ . The two structural inputs are $Q = 2/3$ and $\theta_0 = 2/9$, and the resulting masses are $m_e = 0.51096$ MeV and $m_\mu = 105.652$ MeV against 0.51100 and 105.658 observed (0.007% and 0.006%; SM §7.2). The ratio $m_e/m_\tau = 2.876 \times 10^{-4}$ lies within 0.84% of

$$\frac{m_e}{m_\tau} = \frac{\lambda^4}{d^2} = \frac{1}{36\pi^4} \approx 2.852 \times 10^{-4},$$

with $\lambda = 1/(\pi\sqrt{2})$ the Cabibbo angle from the quark sector (§6.5) and $d = 3$ the spatial dimensionality (Chapter 5 §5.2). The proximity plays no role in the chain: the mass point is fixed by Q and θ_0 . The assignment of m_e specifically to the lightest charged-lepton mass eigenvalue is Layer 2(b) bijection-specific by the gauge-invariance constraint of Chapter 2 §2.5.

Combined with $Q = 2/3$ and the empirical m_τ , this second commitment determines all three charged-lepton masses:

Mass	Formula	Predicted	Observed	Deviation
m_e	Koide(Q, θ_0 ; m_τ)	0.51096 MeV	0.5110	0.007%
m_μ	Koide closure with m_e, m_τ	105.74 MeV	105.66	0.07%

Both predictions match observation at the percent level. The implied $m_\mu/m_\tau = 0.0595$ from Koide closure agrees with observation more tightly (0.07%) than the m_e/m_τ structural relation itself (0.84%) — the pattern is consistent with electroweak-scale renormalization-group running of the lepton Yukawa couplings from the substrate (compositeness) scale to the electroweak scale, where the larger m_τ Yukawa receives proportionally smaller logarithmic corrections than the small m_e Yukawa.

The cross-sector pattern is significant: the lepton mass spectrum inherits *both* the Wolfenstein parameter $A^2 = 2/3$ (via the PMNS Condition 2 sum rule of Ch08 §8.5) *and* the Cabibbo coupling $\lambda^2 = 1/(2\pi^2)$ (via the PMNS reactor angle $\sin^2 \theta_{13}$) from the quark sector. The lepton-sector structure is not freshly postulated; it inherits from the same substrate machinery that produces the quark-sector observables.

The classification per Chapter 4 §4.5 is that the electron and muon masses are Class B-M (mass-chain inheritance from a single upstream empirical input m_τ plus structural commitment); the Koide relation $Q = 2/3$ is Class B-S Layer 1 from the cubic-group Casimir $C_2(T_1)/d^2 = 2/9$; the phase input $\theta_0 = 2/9$ is Class B-S Layer 1 from the same Casimir structure, with the amplitude prefactor computed equal to one (SM §7.2).

The magnitude content of the lepton chain is closed at Layer 1: both inputs derive from the cubic Casimir with computed unit prefactor (SM §7.2). The analogous coefficient task for Condition 2 in Ch08 §8.5 remains open; there the framework provides a structural identification with cross-sector content (quark-sector inputs determining lepton-sector outputs through cubic-group machinery), and the percent-level empirical match is consistent with the structural identification modulo renormalization-group running, but the substrate-level integration that would derive the specific numerical prefactor requires lattice perturbation theory work beyond the chapter’s scope. The directional assignment (which specific generation is e , μ , or τ , and correspondingly which eigenvalue carries which label) remains Layer 2(b) bijection-specific by the gauge structure of Chapter 2 §2.5 and is not derivable from substrate alone in any framework consistent with $\mathcal{G}_{\text{sub}}$ -invariance.

The full six-mass chain. The structural relations link the six lighter charged fermion masses through one empirical mass scale (m_s) plus two further structural inputs that the chapter makes explicit below:

$$m_s \xrightarrow{\lambda^2} m_d \xrightarrow{\sqrt{\theta_0}} m_u \xrightarrow{Q_{\text{down}}} m_b \xrightarrow{Z_S} m_\tau \xrightarrow{\theta_0} m_e, m_\mu.$$

The five arrows are structural relations of three different epistemic statuses, which the chapter distinguishes explicitly:

Layer-1 cubic-group derivations (parameter-free): the Cabibbo-derived $m_d/m_s = \lambda^2$ from the inter-generation Brillouin-zone distance squared; the $m_u/m_d = \sqrt{\theta_0}$ from the intra-doublet splitting at the amplitude level;

the charged-lepton Koide $Q = 2/3$ from $C_2(T_1)/d^2$; and the Koide phase $\theta_0 = 2/9 = C_2(T_1)/d^2$ with computed unit prefactor (SM §7.2). These four relations are all derived from cubic-group representation theory with no free parameters.

Derived running factor plus measured non-perturbative renormalization: the down-to-tau ratio combines two components of different character. The factor 4.28 is derived: tree-level taste-independence gives $y_b = y_\tau$, and full Standard Model RGE running from a common M_{Pl} boundary condition gives $m_b/m_\tau|_{\text{tree}} = 4.276$, with a $\sim 1\%$ systematic from higher-loop corrections and electroweak-threshold matching (SM §7.5). The factor Z_S is measured: the scalar-density renormalization $Z_S = \langle\psi\psi\rangle_{\text{int}}/\langle\psi\psi\rangle_{\text{free}}$, computed non-perturbatively on SU(3) backgrounds at the matching scale $m_{\text{match}} = \lambda g_0^2$, whose dimensional form is forced by power-counting in the taste-decomposed Coleman–Weinberg potential. At $L = 32$, $Z_S(\lambda g_0^2) = 1.813 \pm 0.001$ (statistical), with the $L \rightarrow \infty$ extrapolation ≈ 1.83 and a realistic 1–2% systematic from finite volume and the regulator choice. The prediction $m_b/m_\tau = 4.28/Z_S = 2.361$ matches the observed 2.352 ± 0.017 at 0.4% (0.5σ ; $\sim 0.7\sigma$ at $L \rightarrow \infty$). The open Layer-2 element is the scalar prefactor in the matching-scale identification (the K coefficient): the Monte Carlo consistency check gives $m_{\text{match}}^{\text{MC}} = 0.1234 \pm 0.0002$ against $\lambda g_0^2 = 0.1217$ (1.4%), and $K = 1$ is excluded by the same data (SM §7.5).

Phenomenological auxiliary parameter (Class P): the down-sector Koide invariant $Q_{\text{down}} \approx 0.7304$ does not match any simple cubic-group form (the closest cubic-group ratios — $11/15$, $5/7$, $\theta_0 \cdot (3/2) \cdot (1 + \lambda^2)$ — all deviate from the empirical value by more than 1%). The framework treats Q_{down} as a fitted auxiliary parameter classified as Class P per Chapter 4 §4.5, not as a structural prediction. The chapter’s earlier framing of Q_{down} as “inheriting from QCD running corrections” is replaced by this honest acknowledgment: Q_{down} is fitted to the observed m_b value within the down-sector Koide structure, and the framework does not claim a first-principles derivation. Whether Q_{down} has a derivable structural form different from the charged-lepton sector’s $2/3$ — a magnitude-level question potentially closable through lattice perturbation theory in the emergent description — is open. The sector-identification content (which Q values attach to down-quarks vs charged-leptons) is bijection-specific per Chapter 2 §2.5.

Mass	Formula	Class	Predicted	Observed	Match
m_d	$m_s/(2\pi^2)$	Layer 1	4.73 MeV	4.67 ± 0.48	1.3%
m_u	$m_d \times \sqrt{2/9}$	Layer 1	2.20 MeV	2.16 ± 0.49	0.1σ
m_b	Koide(Q_{down})	P (fitted)	4144 MeV	4180 ± 30	0.9%

Mass	Formula	Class	Predicted	Observed	Match
m_τ	$m_b \times Z_S/4.28$	L (Z_S mea- sured; 2(a): K)	1755 MeV	1776.9	1.2%
m_e	Koide($Q, \theta_0; m_{\text{down}}$)	Layer 1	0.51096 MeV	0.511	0.007%
m_μ	Koide closure	Layer 1	105.74 MeV	105.66	0.07%

The chain produces six masses from one empirical scale (m_s) plus two structural auxiliary inputs (Q_{down} as Class P, the residual normalization on Z_S as Class P). The four Layer-1 relations are parameter-free; the two Class P inputs are fitted to observation. Four out of five mass-chain transitions are therefore parameter-free, with the down-sector Q_{down} closure and the bottom-to-tau residual normalization carrying the auxiliary inputs.

The top quark mass. The top quark sits separately from the chain above because its Yukawa coupling is large. The weak-scale relation $m_t = v/\sqrt{2}$ with $y_t(m_t) \approx 1$ reproduces the observed mass:

$$m_t = \frac{v}{\sqrt{2}} = \frac{246 \text{ GeV}}{\sqrt{2}} = 174.1 \text{ GeV},$$

where v is the Higgs vacuum expectation value, matching the observed 172.5 ± 0.3 GeV to 0.9%.

The mechanism behind this relation is open, and the honest classification is a retrodiction rather than a structural prediction. The natural proposal — a compositeness-scale normalization $y_t(M_{\text{Pl}}) = 1$ carried to low energy by the infrared quasi-fixed point — does not deliver the observed value under one-loop Standard Model running: $y_t(M_{\text{Pl}}) = 1$ flows to $y_t(m_t) \approx 1.19$, giving $m_t \approx 207$ GeV, and the IR quasi-fixed point itself sits at $y_t \approx 1.25$ ($m_t \approx 215\text{--}220$ GeV), not at unity. The attractor basin $y_{\text{UV}} \in [0.5, 5]$ focuses to $m_t \in [177, 221]$ GeV, a 25% spread. What is empirically exact is the weak-scale statement $y_t(m_t) \approx 1$; the compositeness-scale normalization plus IR attraction does not produce it, and the observed top corresponds to $y_t(M_{\text{Pl}}) \approx 0.4\text{--}0.5$. The relation is therefore recorded as a retrodiction of the weak-scale $y_t \approx 1$ coincidence, with the structural mechanism an open task; a derivation of $y_t = 1$ at the electroweak matching scale, or of $y_t(M_{\text{Pl}}) \approx 1/2$ (which runs to ≈ 177 GeV), would close it.

The relation $m_t = v/\sqrt{2}$ matches the observation to 0.9%, but as a retrodiction of the weak-scale $y_t \approx 1$ coincidence rather than a parameter-free structural prediction; a derivation of $y_t = 1$ at the electroweak matching scale, or of $y_t(M_{\text{Pl}}) \approx 1/2$ (which runs to ≈ 177 GeV), would close the mechanism.

6.7 PMNS mixing angles

The PMNS matrix parametrizes the misalignment between neutrino mass eigenstates and flavor eigenstates, analogous to the CKM matrix in the quark sector but with three angles $(\theta_{12}, \theta_{23}, \theta_{13})$ governing the dominant mixing structure. In the Standard Model the three angles are independent free parameters. The framework derives all three from cubic-group representation theory combined with parity-graded charged-lepton corrections at the Cabibbo scale, with the post-JUNO measurement of $\sin^2 \theta_{12}$ matched by one parameter-free value at 0.07σ (a retrodiction; see Chapter 8 provenance).

Tribimaximal mixing as the leading structure. The starting point is tribimaximal (TBM) mixing, derived from the alternating subgroup $A_4 \subset O$ of the cubic rotation group. TBM is the specific pattern

$$\sin^2 \theta_{12} = 1/3, \quad \sin^2 \theta_{23} = 1/2, \quad \sin^2 \theta_{13} = 0,$$

historically introduced as an empirical observation about the leptonic mixing pattern and later given group-theoretic foundations. In the framework, TBM emerges as the zeroth-order PMNS structure from the A_4 subgroup of the cubic rotation group acting on the three lepton generations — the same cubic group whose decomposition $T_1 \oplus E \oplus A_1$ produced the Standard Model gauge group in Chapter 5.

The TBM structure is not the final answer: the experimentally measured $\sin^2 \theta_{13} \approx 0.022$ is non-zero, and $\sin^2 \theta_{12}$ and $\sin^2 \theta_{23}$ deviate from their TBM values by amounts consistent with corrections of order $\lambda^2 = 1/(2\pi^2) \approx 0.05$. The framework's content is the structural derivation of these corrections from the charged-lepton sector.

The three structural ingredients. The deviations from TBM are controlled by three structural ingredients that the framework derives rather than postulates.

First, the overall correction scale is the squared Cabibbo angle $\lambda^2 = 1/(2\pi^2)$ from §6.5. The PMNS matrix's deviation from TBM follows from the charged-lepton sector's perturbation away from the basis in which TBM holds exactly; the charged-lepton perturbation has natural amplitude λ , with second-order effects entering at λ^2 .

Second, the charged-lepton perturbation has a specific structure: it is U -parity odd under μ - τ exchange. The leading-order generator A_1 in the perturbative expansion $V_\ell = \exp(\lambda A_1 + \lambda^2 A_2 + O(\lambda^3))$ satisfies $U A_1 U^{-1} = -A_1$ under the μ - τ exchange generator of $S_4 \supset A_4$. Geometrically, A_1 is a rotation around the symmetric combination $(e_2 + e_3)/\sqrt{2}$ by angle $A^2 = 2/3$ — with the Wolfenstein A of §6.5 entering directly as the rotation angle. The structural property is that this A_1 produces no $O(\lambda)$ corrections to the TBM angles θ_{12} and θ_{23} (the corrections are pushed to $O(\lambda^2)$), while producing the leading reactor angle $\sin \theta_{13} = A^2 \lambda$ at order λ .

Third, the ratio of corrections satisfies $\Delta_{13}/\Delta_{23} = A^4 = 4/9$ from the same Wolfenstein $A = \sqrt{2/3}$ via a second-order projection onto the Higgs democratic direction. The same structural Casimir-over-dimension identity $A^2 = C_2(T_1)/d = (d-1)/d$ that produced the Koide angle in §6.6 now produces the PMNS ratio.

The three predicted angles. The three ingredients combine to give the three angles:

$$\begin{aligned}\sin^2 \theta_{12} &= \frac{1}{3} - \frac{1}{4\pi^2} = 0.3080, \\ \sin^2 \theta_{23} &= \frac{1}{2} + \frac{1}{2\pi^2} = 0.5507, \\ \sin^2 \theta_{13} &= \frac{4}{9} \cdot \frac{1}{2\pi^2} = \frac{4}{18\pi^2} = 0.02252.\end{aligned}$$

These can be compared with the experimental values from the post-JUNO global fit and from NuFIT 6.0 (normal mass ordering).

Angle	Predicted	Observed	Match
$\sin^2 \theta_{12}$	0.3080	0.3085 ± 0.0073 (post-JUNO global fit)	0.07σ
$\sin^2 \theta_{23}$	0.5507	$0.561^{+0.012}_{-0.015}$ (NuFIT 6.0 NO)	0.74σ
$\sin^2 \theta_{13}$	0.02252	0.02195 ± 0.00056 (NuFIT 6.0 NO)	1.01σ

All three angles match observation within about 1σ . The match for $\sin^2 \theta_{12}$ at 0.07σ is unusually close, made tighter by the JUNO measurement reported in November 2025.

The JUNO measurement. The Jiangmen Underground Neutrino Observatory reported its first measurement of reactor neutrino oscillations in November 2025, achieving the world’s most precise determination of $\sin^2 \theta_{12}$: the value 0.3092 ± 0.0087 , a factor of 1.6 improvement over all previous measurements combined. The framework’s prediction $1/3 - 1/(4\pi^2) = 0.3080$ matches the JUNO measurement directly at 0.14σ . The updated global fit incorporating JUNO data gives $\sin^2 \theta_{12} = 0.3085 \pm 0.0073$, matching the prediction at 0.07σ . As JUNO accumulates data over its thirty-year design lifetime, the error bar is projected to reach ± 0.0014 , testing the prediction at sub-percent precision.

The structural status of the three predictions deserves careful articulation. The reactor angle $\sin^2 \theta_{13}$ has the cleanest derivation: it depends only on the leading-order charged-lepton generator A_1 at second order in λ , and is consequently Layer 1 unconditional structural — independent of any second-order Wilson

coefficient. The structural identity $\sin \theta_{13} = A^2 \lambda = (C_2(T_1)/d) \cdot \lambda$ is exact at order λ^2 . This is the cleanest of the three PMNS predictions.

The solar and atmospheric angles $\sin^2 \theta_{12}$ and $\sin^2 \theta_{23}$ depend on a structural relation among second-order Wilson coefficients b_{12}, b_{13}, b_{23} that arises from the charged-lepton perturbation's A_2 structure. The relation reduces to $b_{23} = (4/3)(b_{12} + b_{13})$, which has the geometric reading $4/3 = 2A^2 = 2(d-1)/d$ — partner-count times per-partner projection in the cubic-group representation theory. Under this relation, the two angles satisfy the joint sum rule $2\sin^2 \theta_{12} + \sin^2 \theta_{23} = 7/6$. Observed: 1.178 ± 0.020 using the post-JUNO global fit and NuFIT 6.0 atmospheric data; agreement at 0.6σ . This depends on the atmospheric octant: the quoted value uses the upper-octant best fit $\sin^2 \theta_{23} \approx 0.561$; under the lower-octant solution (≈ 0.455) the combination is ≈ 1.07 , roughly 5σ below $7/6$.

The structural status of $b_{23}/(b_{12} + b_{13}) = 4/3$ separates into two distinct layers. The cubic-group identity $4/3 = 2A^2 = C_2(T_1)^2/d$ that fixes the *magnitude* is Layer 2(a) — forced by the representation theory of T_1 in the cubic group and substrate-derivable in principle through one-loop lattice perturbation theory in the emergent post-trace-out Yukawa description (per Ch02 §2.5, no Lagrangian exists at the substratum proper). The *directional* assignment — that b_{23} specifically is the larger coefficient rather than $b_{12} + b_{13}$, with the labeling fixed by the empirically observed mass hierarchy ($m_e \ll m_\mu \ll m_\tau$ pinning generation 1 = electron) — is Layer 2(b) bijection-specific content by the gauge-invariance argument of Chapter 2 §2.5: under state-relabeling gauge $\sigma \in \mathcal{G}_{\text{sub}}$, the function $b_{23}/(b_{12} + b_{13})$ transforms non-trivially, so no $\mathcal{G}_{\text{sub}}$ -invariant substrate computation can fix the directional assignment. The classification per Chapter 4 §4.5 is therefore that the two angles are Class B-L (layered conditional), with the Layer 2(a) magnitude content potentially closable through one-loop staggered-fermion calculation and the Layer 2(b) directional content correctly classified as bijection-specific input rather than derivational gap.

The framework's prediction structure compared to alternatives. The $2\sin^2 \theta_{12} + \sin^2 \theta_{23} = 7/6$ sum rule is the framework's distinguishing structural prediction in the PMNS sector. The TM1 and TM2 partial-TBM patterns — alternative deviation structures from TBM also derivable from cubic-group representation theory — preserve specific columns or rows of the TBM matrix but do not produce this exact sum rule. The empirical match to the sum rule therefore distinguishes the framework's specific cubic-symmetric perturbation from the TM1/TM2 alternatives at the sub-half-percent level on the sum-rule precision. Chapter 8 develops the JUNO prediction and its discrimination from these alternatives in detail.

The framework's lepton-sector content has more structural commitments than the three angles. The sum rule itself is a falsifiable consequence of the cubic-symmetric perturbation. The neutrino mass ordering is predicted to be normal rather than inverted, by the taste-breaking structure of the neutrino mass matrix developed in Chapter 5 §5.6. The neutrino nature is predicted to be

Majorana rather than Dirac, following from the same structure. Each of these is a pre-registered falsification condition per Chapter 4 §4.5: confirmed inverted ordering would refute the taste-breaking derivation; a confirmed Dirac nature would refute the prediction that no right-handed neutrino exists in the spectrum.

The Higgs mass. A separate structural prediction completes the chapter’s quantitative content: the Higgs quartic coupling at the compositeness scale is exactly zero. The Higgs is the singlet (A_1) taste of the staggered decomposition — a composite scalar in the framework’s structural account, as Chapter 5 §5.5 developed — and its quartic self-coupling at the Planck scale receives no tree-level contribution. The Standard Model quartic λ at low energies is generated entirely by RGE running between the compositeness scale and the experimental scale.

The structural prediction is $\lambda(M_{\text{Pl}}) = 0$. The observed value from Buttazzo et al. is $\lambda(M_{\text{Pl}}) = -0.013 \pm 0.020$, consistent with zero at 0.6σ . The structural commitment is therefore confirmed at the boundary-condition level.

The corresponding prediction for the Higgs mass in GeV is derived by combining the structural boundary condition with the measured top mass m_t as a one-input parameter and running the three-loop SM RGE. The result is $m_H \approx 129\text{--}132$ GeV for $m_t = 172\text{--}173$ GeV, against the observed $m_H = 125.10 \pm 0.14$ GeV. The gap is sensitive to m_t (with $\partial m_H / \partial m_t \approx -1.1$ GeV/GeV), and the recent CMS measurement $m_t = 170.5 \pm 0.8$ GeV would bring the prediction to about 128 ± 2 GeV. The structural claim $\lambda(M_{\text{Pl}}) = 0$ is confirmed; the specific Higgs mass in GeV is consistent with observation within the current m_t uncertainty.

The framework’s Higgs-sector prediction has a notable convergence with the broader physics literature. The Multiple Point Principle developed by Bennett, Froggatt, and Nielsen in the 1990s postulated $\lambda(M_{\text{Pl}}) = 0$ as a fine-tuning condition, predicting the Higgs mass at approximately the value subsequently observed. Wetterich’s asymptotic-safety analysis showed that if the Standard Model holds to the Planck scale with an asymptotically safe gravity contribution, the observed Higgs mass requires $m_t \approx 171$ GeV. The framework’s structural derivation of $\lambda(M_{\text{Pl}}) = 0$ thus provides a substratum-level account of what an active beyond-SM literature has been postulating as a boundary condition.

6.8 Summary table and chapter close

The chapter has developed twenty-two Standard Model retrodictions from the cubic-lattice substratum, with the structural status of each made explicit per the methodological discipline of Chapter 4. The summary table collects the predictions with two complementary classification axes: the **Class** column gives the *derivational layer* per Chapter 4 §4.5 (S/L/R/E/M), and the **Cat** column gives the *independent-evidence status* per the categorization developed below (A/B/C/D/E).

Observable	Source	Predicted	Observed	Match	Class	Cat
$1/\alpha_0$	one-loop stag- gered VP	23.25	—	structural input	S	input
$\alpha_3(M_Z)$	RG running with (δ_0, A, B)	0.118	0.1181 ± 0.0011	fit	R [a]	E
$\alpha_2(M_Z)$	RG running with (δ_0, A, B)	0.0339	0.03379 ± 0.00009	fit	R [a]	E
δ_0	U(1) row fit	10.0	—	empirical input	E	input
$A \cdot B$ cross-check	independent deriva- tion	$\approx 48\text{--}62$ (leading- power- dependent)	46.4 (from fit)	inconclusive		A
$\bar{\theta}$	T- invariance + detailed balance	0	$ \bar{\theta} < 10^{-10}$	✓	S	A
Cabibbo λ	inter- generation BZ distance	0.2251	0.2250 ± 0.0007	0.04%	S	A
Wolfenstein A	off- democratic projec- tion	$\sqrt{2/3} = 0.8165$	0.826 ± 0.012	1.2%	S	A
$ V_{cb} $	$A\lambda^2$	0.0414	0.0408 ± 0.0014	0.4σ	S	A
m_d/m_s	GST $\lambda^2 = 1/(2\pi^2)$	0.0507	0.0503 ± 0.0007	0.7%	L	A
Jarlskog J	$\eta/(12\pi^6)$	3.02×10^{-5}	$(3.08 \pm 0.13) \times 10^{-5}$	0.5σ	E (η)	B
Koide Q	cubic T_1 struc- ture	2/3	0.66667	0.001%	S	A
Koide angle θ_0	$C_2(T_1)/d^2$	2/9	0.22222	0.02%	S	A

Observable	Source	Predicted	Observed	Match	Class	Cat
m_e	Koide(Q, θ_0)	0.51096 MeV	0.51100 MeV	0.007%	M (m_τ)	B
m_μ	Koide closure	105.74 MeV	105.66 MeV	0.07%	M (m_τ)	B
m_d	$m_s/(2\pi^2)$	4.73 MeV	4.67 ± 0.48	1.3%	M (m_s)	C
m_u	$m_d\sqrt{2/9}$	2.20 MeV	2.16 ± 0.49	0.1σ	L [e]	B
m_b	Koide (Q_{down})	4144 MeV	4180 ± 30	0.9%	L [d]	E
m_t	$v/\sqrt{2}$ (weak- scale $y_t \approx 1$; mecha- nism open)	174.1 GeV	172.5 ± 0.3	0.9%	R	D
$\sin^2 \theta_{12}$	$1/3 - 1/(4\pi^2)$	0.3080	0.3085 ± 0.0073	0.07σ	L [f]	B
$\sin^2 \theta_{23}$	$1/2 + 1/(2\pi^2)$	0.5507	$0.561^{+0.012}_{-0.015}$	0.74σ	L [f]	B
$\sin^2 \theta_{13}$	$A^4\lambda^2 = 4/(18\pi^2)$	0.02252	0.02195 ± 0.00056	1.01σ	S	A
$\lambda(M_{\text{Pl}})$	composite-0 Higgs struc- ture		-0.013 ± 0.020	0.6σ	S	A
m_H	RGE running from $\lambda(M_{\text{Pl}}) = 0$ (band)	$129\text{-}132$ GeV	125.10 ± 0.14	m_t - consistent(m_t)	M	D

Footnotes: **[a]** The three gauge couplings at M_Z are reproduced with one empirical input (δ_0) and two fitted parameters (A, B); the structural predictions in this sector are $1/\alpha_0$, the universality of δ_0 , and the $A \cdot B$ cross-check. **[d]** Q_{down} is a phenomenological auxiliary parameter (Class P), fitted to reproduce the observed m_b within the down-sector Koide structure; no first-principles derivation, per §6.6 mass-chain discussion. The bottom-to-tau bridge uses $Z_S \approx \pi/\sqrt{3}$ (candidate Layer 1 structural identification) plus a residual normalization factor of $1/4.28$ (Class P, no obvious cubic-group form). **[e]** Layer 1 form + Layer 2(b) alignment of SU(2) doublet structure with cubic T_1 perturbation. **[f]** Layer 1 form + Layer 2(a) magnitude $b_{23}/(b_{12} + b_{13}) = 4/3 = 2A^2$ (cubic-group identity, derivable through lattice perturbation theory in emergent description) +

Layer 2(b) directional assignment (which pair-ratio gets the 4/3, fixed by mass-hierarchy bijection per Ch02 §2.5 and Ch08 §8.5).

Parameter accounting — two complementary classifications. The table above shows two classification axes for each prediction. The **Class** column (S/L/R/E/M) reflects the *derivational layer* per Chapter 4 §4.5 (S = unconditional structural; L = layered; R = retrodiction; E = empirical-input-using; M = mass-chain). The **Cat** column reflects the *independent-evidence status* — whether a prediction is independent empirical evidence for the framework’s structural inputs or a consequence of those inputs already counted elsewhere.

- **Category A — independent parameter-free predictions:** Cabibbo λ , Wolfenstein A , $|V_{cb}| = A\lambda^2$, $m_d/m_s = \lambda^2$, $\sin^2 \theta_{13} = A^4\lambda^2$, Koide $Q = 2/3$ (and equivalently Koide $\theta_0 = 2/9$ — the same prediction in dual form), $A \cdot B$ gauge-coupling cross-check, $\lambda(M_{P1}) = 0$ from composite-Higgs structure, $\theta = 0$. Nine predictions that consume only structural inputs (cubic-group invariants and spatial dimensionality) and produce *independent empirical channels*: CKM phenomenology (λ , $|V_{cb}|$), lattice QCD (m_d/m_s), PMNS ($\sin^2 \theta_{13}$), lepton-mass dynamics (Koide Q), gauge-coupling structure ($A \cdot B$), Higgs sector ($\lambda(M_{P1})$), and CP physics (θ). Several entries are *composites* of the framework’s structural inputs (A and λ) but are *independent measurements* from the empirical side — $|V_{cb}|$ is measured independently of V_{us} in CKM phenomenology; m_d/m_s is measured by lattice QCD independently of the Cabibbo angle; $\lambda(M_{P1})$ is determined by RG-running observed m_H and m_t back to the Planck scale. The framework’s specific structural inputs ($A^2 = 2/3$, $\lambda^2 = 1/(2\pi^2)$) are therefore tested through multiple independent empirical channels. The $\sin^2 \theta_{13} = A^4\lambda^2$ identification illustrates the structural specificity: alternative power combinations ($A^2\lambda^2$, $A\lambda^2$, $A^4\lambda^4$) deviate from observation by 32%, 50%, or 95% respectively, isolating the specific $A^4\lambda^2$ form.
- **Category B — parameter-free predictions conditional on empirical mass-hierarchy bijection:** $\sin^2 \theta_{12}$, $\sin^2 \theta_{23}$, m_e and m_μ via Koide closure (Q , θ_0), m_u via intra-doublet $\sqrt{\theta_0}$ splitting, Jarlskog J via $\eta/(12\pi^6)$. Five predictions whose structural forms are tightly identified at the cubic-group magnitude level but whose specific assignment to labeled observables (which generation is electron, which is the second-lightest, etc.) is Layer 2(b) bijection-specific per Chapter 2 §2.5. Closure of the magnitude content depends on lattice perturbation theory in the emergent post-trace-out description; closure of the directional/assignment content is gauge-forbidden under $\mathcal{G}_{\text{sub}}$ -invariance and properly classified as empirical input rather than derivational gap. The framework’s structural forms match observation at sub-percent or sub-sigma precision; the implied- η consistency check for Jarlskog (§6.5) gives $\eta = 0.367$ from the framework’s A^2 and λ^2 inputs combined with observed J , within 0.83σ of PDG $\eta = 0.357 \pm 0.011$ — the structural $12\pi^6$ suppression is consistent with observation.

- **Category C — composite of Category A inputs with redundant content:** $m_d = m_s/(2\pi^2)$ — same prediction as $m_d/m_s = \lambda^2$ with m_s plugged in; an internal-consistency check rather than independent test.
- **Category D — band predictions at order-of-magnitude precision:** m_t (weak-scale relation $v/\sqrt{2}$ at 0.9%; fixed-point mechanism open — §6.7); m_H (RG-running band 129-132 GeV from $\lambda(M_{P1}) = 0$, observation 125.10 ± 0.14 , 4-7 GeV below the band).
- **Category E — fitted using multiple empirical parameters:** $\alpha_3(M_Z)$, $\alpha_2(M_Z)$ (fitted via (δ_0, A, B) , with structural content concentrated in the *form* of the fit rather than the matching values); m_b (uses Class P auxiliary parameter Q_{down} — see footnote [d]).
- **Input rows:** $1/\alpha_0$ and δ_0 are structural inputs to other predictions rather than themselves predicted observables.

The Category distribution across the twenty-three table entries: nine in Category A (independent parameter-free), six in Category B (parameter-free conditional on the empirical mass-hierarchy bijection), one in Category C (composite-redundant), two in Category D (band), three in Category E (fitted), plus two structural-input rows ($1/\alpha_0$ and δ_0). Beyond the table, two formal uniqueness theorems function as further structural commitments: the three-generation count (Appendix B.6c) and the absence of grand-unified gauge structure (Appendix B.6b). The theorems are not retrodictions in the table’s sense (no specific numerical value to compare against observation) but are pre-registered falsifiable structural commitments.

The framework’s content is concentrated at Layer 0 and Layer 1 of the four-layer framing — gauge structure and pure cubic-group representation theory — with Layer 2 and Layer 3 content (bijection-specific operator structures and the overall mass scale) appearing as explicit inputs rather than fitted parameters. The structural-vs-bijection-level distinction developed in Chapter 4 §4.5 is what makes the empirical record evaluable: the framework predicts the *angular* and *ratio* structure of fermion mass relations from cubic-group representation theory, but the overall mass scale of each chain is set by the specific bijection and is taken as one empirical input.

The Category A predictions span seven distinct empirical channels — CKM phenomenology (λ , $|V_{cb}|$), lattice QCD (m_d/m_s), PMNS ($\sin^2 \theta_{13}$), lepton-mass dynamics (Koide Q), gauge-coupling structure ($A \cdot B$), Higgs sector ($\lambda(M_{P1})$), and CP physics (θ). The framework’s specific structural inputs ($A^2 = 2/3$, $\lambda^2 = 1/(2\pi^2)$, $C_2(T_1)/d^2 = 2/9$, T-invariance) are tested through these channels, with empirical agreement at sub-percent or sub-sigma precision in each. The two formal uniqueness theorems (three-generation count via B.6c and no-GUT structural via B.6b) close the gauge structure and matter content as structural commitments forced by the substratum’s existing commitments rather than freely chosen.

Falsification commitments revisited. The pre-registered falsification conditions from Chapter 4 §4.5 connect directly to the predictions developed here. Any stable fourth-generation fermion would refute the taste-reduction derivation that locks the generation count at three (formal uniqueness theorem B.6c). Any stable beyond-Standard-Model gauge structure would refute the cubic-group commutant derivation (formal uniqueness theorem B.6b). Any observation of $\theta \neq 0$ at sufficient precision would refute the T-invariance commitment. Any precision-upgrade measurement moving the Cabibbo angle outside 3σ of $1/(\pi\sqrt{2})$ would refute the framework’s geometric prediction. Any precision-upgrade measurement moving the Koide relation outside 3σ of $2/3$ would refute the cubic-Casimir derivation. The Category A predictions in independent CKM, lattice-QCD, and Higgs-sector channels — $|V_{cb}|$, m_d/m_s , and $\lambda(M_{P1})$ — provide further falsification handles through measurements independent of the Cabibbo channel that constrains λ .

The Category B predictions ($\sin^2 \theta_{12}$ and $\sin^2 \theta_{23}$, the mass-chain entries via the Koide structure, the Jarlskog relation) constitute *conditional* falsification targets. The framework’s specific structural forms are the natural predictions of the cubic-group machinery; their match to observation under the empirical mass-hierarchy bijection is what the framework predicts. The magnitude-level closure (verifying the cubic-Casimir structure produces the specific numerical coefficients) is open through lattice perturbation theory in the emergent description; the directional/assignment-level content is Layer 2(b) bijection-specific by Chapter 2 §2.5 and is empirical input rather than open derivation. JUNO’s design-precision measurement of $\sin^2 \theta_{12}$ at sub-percent precision is the nearest forward-falsifiable test at high precision in this category.

Cross-sector universality. A structural feature of the framework’s prediction package deserves explicit treatment: the same numerical structural inputs that produce quark-sector observables also produce lepton-sector observables, with no separate fitting between sectors.

The Wolfenstein parameter $A^2 = 2/3$ from off-democratic projection on the cubic generation axes appears in the quark sector via $|V_{cb}| = A\lambda^2 = 0.041$ (matching observation at 0.4σ), in the gauge-coupling cross-check $A \cdot B = 48 \pm 1.5$ vs 46.4 (3% match), and in the lepton sector via the PMNS Condition 2 sum rule $2\sin^2 \theta_{12} + \sin^2 \theta_{23} = 7/6$ (matching observation at 0.07σ on JUNO’s $\sin^2 \theta_{12}$). The Cabibbo coupling $\lambda^2 = 1/(2\pi^2)$ from the inter-generation Brillouin-zone distance appears in the quark sector via Cabibbo λ (0.04% match) and $m_d/m_s = \lambda^2$ (0.7% match), and in the lepton sector through the Koide chain (m_e at 0.007% and m_μ at 0.006% via Q and θ_0).

The cross-sector pattern places a substantial constraint on alternative input choices: any alternative natural input to the framework’s machinery that fits quark-sector observation would also need to fit lepton-sector observation. Alternative natural inputs that produce different Cabibbo angles (e.g., BCC or FCC inter-generation distances) all deviate from $\lambda = 1/(\pi\sqrt{2})$ by $> 20\%$ and would correspondingly mispredict m_d/m_s and $\sin^2 \theta_{13}$ by similar fractions. Al-

ternative power combinations of A and λ in $\sin^2 \theta_{13}$ ($A^2 \lambda^2$ at 32% deviation, $A \lambda^2$ at 50% deviation) similarly fail the cross-sector match. The framework's specific $A^2 = 2/3$ and $\lambda^2 = 1/(2\pi^2)$ inputs are consistent with observation in both sectors simultaneously.

What follows. Chapter 7 develops the gravitational sector — the Bekenstein-Hawking $1/4$ factor derived by boundary-mode counting (GW250114 confirms the classical area theorem; the coefficient has no direct test), the value of $\hbar$ from horizon thermodynamics, the dissolution of the cosmological constant problem at the level of two incommensurable descriptive levels, and the dark-sector phenomenology with the MOND acceleration scale $a_0 = cH/6$ and the entropy-displacement account of dark matter. The empirical record there is independent of and complementary to the Standard Model content developed here: the gravitational sector tests the substratum's boundary thermodynamics rather than its bulk gauge structure.

Chapter 8 develops the JUNO prediction in detail, as a focused case study of the framework's sharpest empirically matched Layer 1-plus-Layer-2(a) value (a parameter-free retrodiction). The November 2025 measurement and its discrimination from the TM1/TM2 alternatives are presented with the experimental context that this chapter abbreviated. Chapter 9 develops universality classes and the framework's relationship to alternative unification programs, with the OI-string analysis at the matrix-model level and the framework's structural commitment to no grand unified theory developed against the broader landscape of unification proposals.

The framework's empirical record in fundamental physics, developed across Chapters 5 through 9, is concentrated at the Level G3 $\times$ Level D intersection of the hierarchical structural realism developed in Chapter 3 — class-specific predictions of the cubic-lattice substratum representative, with substantial empirical support and pre-registered falsifiability. The structural and methodological commitments of Part I are made quantitatively precise here.

Chapter 7

The Gravitational Sector

7.1 What this chapter develops

Chapters 5 and 6 developed the framework's content in the Standard Model sector: the gauge group $SU(3) \times SU(2) \times U(1)$ from the cubic decomposition of the substratum, three matter generations from the staggered taste structure, twenty-two quantitative predictions for SM observables with empirical agreement spanning 0.02% to 1.2% across the parameter space. The content was

developed by *specifying* the substratum dynamics — the wave equation on a cubic lattice with coupling-degree minimization — and tracing the structural consequences. The framework’s content in the gravitational sector follows a different route.

The gravitational sector applies the framework to the cosmological horizon as a causal partition, with the substratum dynamics treated as a free input. The cosmological horizon — the boundary beyond which no signal propagating at or below the speed of light can reach the observer — implements the framework’s visible-hidden partition at its largest physical scale. Sub- c observers share the same causal structure, so all such observers share the same emergent action scale $\hbar$. The horizon’s geometric and thermodynamic properties — surface gravity κ , area A , Hawking temperature T — fix the values of the emergent physics: the action scale $\hbar$, the Bekenstein-Hawking entropy with its $1/4$ coefficient, the form of the dark-energy equation of state.

This chapter develops seven structural results. Section 7.2 verifies that the cosmological horizon satisfies the framework’s three structural conditions (coupling C1, slow bath C2, large capacity C3) at the cosmological scale, with ratio $\tau_S/\tau_B \sim 10^{-32}$. Section 7.3 derives the emergent action scale $\hbar = c^3\epsilon^2/(4G)$ from thermal self-consistency between the classical horizon temperature (computed from the Jacobson thermodynamic argument with no quantum input) and the Euclidean KMS period of the emergent quantum field theory. Section 7.4 fixes the discreteness scale $\epsilon = 2l_p$ as the unique simultaneous solution of the action-scale relation and the definition of the Planck length. Section 7.5 derives the Bekenstein-Hawking entropy $S = A/(4l_p^2)$ by direct boundary-mode counting, with the $1/4$ coefficient emerging as the ratio of the Euclidean KMS period to the Jacobson prefactor; the September 2025 LIGO/Virgo/KAGRA observation of GW250114 confirms the area theorem at 99.999% confidence (the $1/4$ coefficient itself has no direct test). Section 7.6 resolves the black hole information paradox: information is never lost from the substratum bijection φ (bijectivity); the Page curve is derived from cycle typicality with structural Page time $t_P \approx 0.646 t_{\text{evap}}$ independent of mass and greybody factors, matching the quantum-extremal-surface and entanglement-island programs through a substantially different mechanism. Section 7.7 dissolves the cosmological constant problem by identifying the emergent quantum vacuum energy and the effective cosmological constant as properties of logically distinct descriptions rather than commensurable quantities whose 10^{122} ratio requires cancellation. Section 7.8 develops the framework’s content in dark-energy phenomenology — Running Vacuum Model with ν_{OI} bounded at the emergent-QFT scale $\sim 10^{-32}$ — and Section 7.9 develops the dark-sector content (95% invisible matter as concordance rather than puzzle (the $\sim 95\%$ being partly definitional, $\approx 1 - \Omega_B$), MOND $a_0 = cH/6$ from entropy displacement, declining rotation curves at $r_M = \sqrt{6GM_B/(cH)}$).

The methodological discipline of Chapter 4 organizes the gravitational predictions into three tiers. **Tier 1** consists of results that follow from partition ge-

ometry alone — the derivation of $\hbar$, the Bekenstein-Hawking area law with the $1/4$ coefficient, the cosmological constant dissolution, and the Type II running-vacuum functional form. These hold for any finite deterministic system satisfying the three structural conditions at its cosmological horizon. They are universality-class-level structural results — Level C $\times$ Level G4 in the hierarchical structural realism of Chapter 3, where the framework operates at the level of structural commitments shared across substratum representatives rather than at the specific cubic-lattice realization. **Tier 2** consists of dark-sector phenomenology — the 95% invisible matter fraction, the MOND acceleration scale $a_0 = cH/6$, the rotation-curve crossover. These follow from the thermodynamic Clausius relation applied to the boundary reservoir and the geometric factor from spherical redistribution in de Sitter space; they sit at Level D $\times$ Levels G1–G2 of the hierarchical structure. **Tier 3** consists of solution-specific quantities — the transition steepness between MOND and Keplerian regimes, the greybody factor for Hawking radiation, the specific microscopic realization of the substratum — which depend on the particular bijection φ within $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ and are addressed by the cubic-lattice realization developed in Chapters 5–6.

The tier system clarifies what each result claims. A Tier 1 result like “ $\hbar = c^3 \epsilon^2 / (4G)$ ” is not “the framework predicts the value of Planck’s constant uniquely” but “any horizon-bounded embedded-observer system with the framework’s structural conditions gets this functional form for the emergent action scale.” Tier 1 results are robust to the substratum’s specific dynamics; only the discreteness scale ϵ remains as a free parameter, and self-consistency with the Planck length definition fixes $\epsilon = 2l_p$. A Tier 2 result like “ $a_0 = cH/6$ ” depends on the geometric factors of de Sitter space but not on the substratum’s gauge structure; it would hold for any framework where the dark-sector phenomenology arose from entropy displacement at the cosmological boundary. A Tier 3 result depends on the specific bijection and varies across realizations within the universality class. The empirical record of the gravitational sector — concentrated at Tiers 1 and 2 — is therefore not a record of the cubic-lattice substratum but a record of the framework’s broader structural commitments at the boundary scale.

7.2 The cosmological horizon as causal partition

The cosmological horizon implements the framework’s causal partition at the largest physical scale accessible to any sub- c observer. The horizon’s role in the framework is to separate visible and hidden sectors by causality rather than by an explicit structural choice: the visible sector is the interior of the horizon, with content that the observer can causally access; the hidden sector is everything beyond, with content that no sub- c signal can reach in finite time. This is a consequence of general relativity’s causal structure, not a modeling choice imposed on top of it.

The horizon partition is universal across sub- c observers. Different observers have different horizon areas — depending on their worldline and the local Hubble

parameter — but they share the same causal structure: any pair of observers within each other’s horizons agrees on what is reachable in finite time and what is not. The emergent action scale $\hbar$ depends only on local geometric quantities and not on the observer’s specific horizon area, so all sub- c observers share the same value of $\hbar$. The framework’s content here is that the universality of Planck’s constant across observers — a feature that the Standard Model treats as a postulate — is a structural consequence of the universality of GR’s causal structure.

Verification of C1 (coupling). The Hamiltonian and momentum constraints of general relativity correlate interior and exterior initial data on any Cauchy surface. In the ADM formulation, the bulk Hamiltonian is a sum of constraints that vanish on-shell, with physical time evolution generated by the boundary term at the horizon, which depends on data from both sides. The physical phase space is therefore a constraint surface within the kinematic product of interior and exterior configuration spaces. This is stronger than the non-trivial coupling required by the framework’s C1: the constraints enforce correlations that persist on all timescales and cannot be perturbatively removed. The cosmological horizon satisfies C1 not as an additional assumption about the dynamics but as a structural consequence of general relativity.

Verification of C2 (slow bath). The hidden-sector relaxation timescale is the Hubble time, $\tau_B \sim 1/H \sim 10^{17}$ s. The visible-sector dynamics timescale for any laboratory process is much faster: typical electroweak processes occur on timescales $\tau_S \sim 10^{-15}$ s. The ratio $\tau_S/\tau_B \sim 10^{-32}$ is extraordinarily small, making the corrections to the framework’s leading-order results at scale τ_S/τ_B correspondingly small. The slow-bath condition is satisfied with enormous margin — the cosmological horizon is among the slowest baths the framework’s conditions admit. Corrections of order 10^{-32} to the framework’s leading results are well below current observational precision in all relevant sectors.

Verification of C3 (large capacity). The horizon’s information-storage capacity scales with the horizon area: the number of independent modes is $N_H = A_{\text{horizon}}/l_p^2 \sim 10^{122}$ for the present-day cosmological horizon. The visible sector’s mode count is much smaller — bounded by the volume divided by the relevant length scale of the observer’s measurement. The ratio $N_H/N_V \gtrsim 10^{60}$ for any laboratory-scale visible sector, with the actual ratio depending on the visible sector’s spatial extent. The capacity condition is satisfied with enormous margin: the cosmological horizon’s information capacity exceeds any visible-sector observer’s capacity by tens of orders of magnitude.

The four conditions hold at the cosmological horizon by structural argument rather than by phenomenological assumption — C4’s readback is the bidirectionality of the boundary coupling, which reads the same degrees it writes. The framework’s results in the gravitational sector therefore follow from the same structural conditions that produced quantum mechanics in Part I, applied at the largest available physical scale. The eigenstate thermalization hypothesis required for the C2 necessity direction holds for the horizon complement (a generic

chaotic many-body system), so the full equivalence applies in our universe.

7.3 The emergent action scale $\hbar$

The most distinctive content of the framework’s gravitational sector is the derivation of Planck’s constant $\hbar$ as an emergent quantity determined by the cosmological horizon’s thermodynamic properties. The standard view treats $\hbar$ as a fundamental constant of nature whose value $1.055 \times 10^{-34} \text{ J} \cdot \text{s}$ is measured rather than derived. The framework derives $\hbar = c^3 \epsilon^2 / (4G)$ from thermal self-consistency between classical horizon thermodynamics and the emergent quantum field theory on the horizon — with ϵ the discreteness scale of the substratum partition, fixed by self-consistency in Section 7.4 to $\epsilon = 2l_p$, recovering the observed value of $\hbar$ exactly.

The derivation proceeds in four steps.

Step 1: The classical horizon temperature. The Jacobson thermodynamic argument assigns the cosmological horizon a temperature from classical considerations alone, with no quantum input. The argument runs as follows. The horizon is a bifurcate Killing horizon with surface gravity κ . The first law of thermodynamics applied to a horizon-bounded region gives $dE = T dS$, where E is the energy flux across the horizon, T is the horizon temperature, and S is the horizon entropy. The Einstein equations applied to the horizon’s Raychaudhuri equation give $dE = (c^2 \kappa / 8\pi G) dA$, where A is the horizon area. Combined with the structural commitment that each minimal cell of area ϵ^2 contributes one unit of entropy ($dS = dA / \epsilon^2$), the first law gives a classical horizon temperature

$$k_B T_{\text{cl}} = \frac{c^2 \epsilon^2 \kappa}{8\pi G}.$$

The temperature contains no $\hbar$ — only classical quantities (κ , ϵ , G , c). This is the substratum’s classical assignment of a temperature to the partition boundary, derived from the Einstein equations and the discreteness scale of the partition.

An independent route to the same conclusion. The necessity of Einstein’s equations for the framework’s emergent matter is also established by an external theorem with an entirely different axiom base. The gravitational closure program of Düll, Schuller, Stritzelberger and Wolz (Phys. Rev. D 97, 084036) shows that for matter field equations whose causal structure is that of standard-model fields on a single Lorentzian light cone — exactly the class the framework’s emergent sector occupies, its single-cone structure supported by the rotational-protection and species-universality results of the companion papers — the requirement that matter and geometry evolve together consistently, independently of how spacetime is sliced, forces the gravitational action to be Einstein–Hilbert with a cosmological constant. No thermodynamics enters that argument at all. The thermodynamic derivation of this chapter and the closure theorem therefore arrive at the same dynamics from disjoint starting points;

where the closure route leaves Newton’s constant and the cosmological constant undetermined, the framework’s relation $\hbar = c^3\epsilon^2/4G$ and the dark-energy analysis of the following sections take over. The closure theorem is, moreover, one of a family of such necessity results that the emergent sector inherits: Lovelock’s theorem, which shows that in four dimensions the only second-order, generally covariant dynamics for a metric alone is Einstein’s with a cosmological constant, and the spin-2 results of Weinberg and Deser, which show that a massless spin-two particle coupled to a complete matter sector is forced into universal coupling and Einstein’s theory at low energy. Four independent routes — thermodynamic, closure, Lovelock, and the spin-2 bootstrap — converge on the same dynamics from entirely different starting assumptions. This is also the structural answer to the standard objection facing every emergent-spacetime program, that such constructions deliver the geometry of spacetime but never Einstein’s dynamics: once the emergent light cone, the matter content, and the massless spin-two mode exist, the dynamics is over-determined rather than under-derived. The behaviour of the framework’s small lattice corrections under the closure machinery remains an open consistency check.

Step 2: Boundary-only dependence. The emergent quantum field theory on the horizon depends only on the boundary modes — the degrees of freedom at the horizon itself — and not on the bulk content of the visible or hidden sectors at scales far from the boundary. This is the framework’s partition-relativity result: the emergent description is uniquely determined by the partition geometry at leading order in τ_S/τ_B , with bulk corrections at order τ_S/τ_B for the cosmological partition. The boundary-only dependence is a structural consequence of the small parameter $\tau_S/\tau_B \sim 10^{-32}$: the emergent QFT is effectively a theory of the boundary modes, with bulk corrections suppressed by the slow-bath timescale ratio.

Step 3: Dimensional analysis of $\hbar$. The emergent action scale $\hbar$ depends only on local geometric quantities at the horizon. Dimensional analysis on the available constants gives $\hbar \propto c^a \epsilon^b G^d$ with $\hbar$ having dimensions of action. The unique combination yielding action dimensions is $\hbar \propto c^3 \epsilon^2 / G$. The proportionality constant is not fixed by dimensional analysis alone; it is fixed by thermal self-consistency in Step 4. The dimensional argument excludes κ from $\hbar$ on observer-universality grounds: different observers have different κ at their respective horizons, but they share the same $\hbar$.

Step 4: Thermal self-consistency. The classical horizon temperature from Step 1 must equal the Euclidean KMS period of the emergent quantum field theory on the same horizon. The KMS condition follows from the regularity of the Wick-rotated metric at the horizon: any QFT on a bifurcate Killing horizon background must be periodic in imaginary time with period $\beta = 2\pi c/\kappa$, giving a KMS state at temperature $T_Q = \hbar\kappa/(2\pi ck_B)$. This is a theorem *within* the derived QFT, not an external import — the regularity requirement follows from the Wick rotation’s analytic structure.

The two temperatures — T_{cl} from the classical substratum, T_Q from the emer-

gent QFT — refer to the same physical degrees of freedom: the boundary modes across which the visible-hidden coupling acts. Temperature is a state property; two complete descriptions of the same modes cannot assign different temperatures without logical contradiction. The matching $T_{\text{cl}} = T_Q$ is not an additional assumption but a structural consequence of the boundary modes being the same physical objects in both descriptions, with corrections at order $\tau_S/\tau_B \sim 10^{-32}$:

$$\frac{c^2 \epsilon^2 \kappa}{8\pi G} = \frac{\hbar \kappa}{2\pi c}.$$

The surface gravity κ cancels — a non-trivial cross-check on the boundary-only argument of Step 2, which excluded κ from $\hbar$ on observer-universality grounds. The thermal matching here recovers that exclusion from independent physics: κ enters both sides of the consistency equation, and only the combinations of c , G , ϵ predicted by Step 3 survive. Solving for $\hbar$:

$$\boxed{\hbar = \frac{c^3 \epsilon^2}{4G}}.$$

The derivation has the structure of a gap equation: ϵ is the free geometric input, $\hbar$ is the output. The match between the two temperature computations is non-trivial. The classical temperature could have depended on volumetric hidden-sector data — excluded by the boundary-only argument of Step 2 — but in fact depends only on the boundary geometry. The KMS temperature could have depended on the state of the emergent QFT — but in fact is purely kinematic, depending only on the surface gravity. One prerequisite of the KMS step is carried explicitly: Euclidean regularity at a bifurcate Killing horizon presupposes local boost structure in the emergent description at the horizon, so the gap equation is partially downstream of the boost sector’s protection mechanism, an open item tracked in the Lorentz-sector accounting ([GR §8.5]). At the free level the prerequisite is met exactly rather than approximately — each massless substrate channel carries a perfectly linear dispersion at all momenta (the Courant-unity property of the discrete wave equation), giving an exactly boost-invariant cone in the radial plane the horizon problem lives in — so the open dependency narrows to the interacting sector’s velocity-splitting question — where the Standard-Model treatment already establishes the decoupling form (the custodial symmetry forbids an independent splitting counterterm, and the one-loop coefficient decouples as $(M_X/M_{\text{Pl}})^2$), leaving one quantitative item shared with the Lorentz-sector accounting: the cross-charged mass M_X within its window of roughly two to four decades — kept open from below by the baryon safety of the cross-charged sector, which is of the proton-safe leptoquark class rather than the grand-unified kind — whose closure would falsify the mechanism. That neither pathology obtains makes the gap equation a genuine determination of $\hbar$ in terms of the discreteness scale.

The counting $\eta = 1/\epsilon^2$ is definitional — and the internal components do not multiply it. The entropy-density argument in Step 1 used $dS = dA/\epsilon^2$

— one boundary mode per cell of area ϵ^2 . This counting is not an independent physical ansatz about the mode density but a consequence of the definition of ϵ as the minimal distinguishable scale. A cell finer than ϵ would subdivide the minimal scale; counting fewer modes would require a distinction below ϵ to group them. Two potential multiplicities are settled separately. The number of internal *states* per mode (the alphabet size) is gauge — observationally irrelevant by the framework’s gauge-freedom result — so it does not enter. The number of internal *components* per site ($K = 2d = 6$) is physical, but it does not multiply the count either, for a structural reason: the $K = 2d$ components are the $2d$ link directions from a cell, each coupling across exactly one link, and a boundary cell has exactly one outward-normal link crossing the horizon — the other $2d - 1$ are tangential or inward. One crossing link per cell means one coupled boundary mode per cell, forced by the link-direction structure rather than assumed ([GR §3.1]). The counting $\eta = 1/\epsilon^2$ is therefore the definition of ϵ together with a forced consequence of the component structure, not an additional commitment.

The counting also sits at the intersection of an independent family of results: the Bekenstein bound and Bousso’s covariant entropy bound each cap the entropy of a region by its boundary area in Planck units — area-law mode counting reached from quantum field theory and from light-sheet geometry respectively. One caution keeps the logic clean: those bounds are formulated with $\hbar$ already in hand, and the framework derives $\hbar$ downstream of this counting, so the family can serve only as downstream convergence — the emergent theory, once $\hbar = c^3\epsilon^2/(4G)$ is fixed, satisfies both bounds — and never as motivation for the counting, which stands on the definitional argument above alone. Citing the bounds as motivation would re-import the circularity that the definitional argument removed.

The predictive content of the gap equation. The framework’s content does not lie in the gap equation alone — given ϵ , $\hbar$ is fixed — but in the consequences of the specific relationship $\hbar = c^3\epsilon^2/(4G)$. The relationship produces, in Section 7.5, the Bekenstein-Hawking entropy $S = A/(4l_p^2)$ with the exact factor $1/4$. In Section 7.7, it produces the cosmological constant dissolution with the de Sitter entropy S_{dS} as the structural compression ratio. In Section 7.8, it produces the Type II running-vacuum form of the dark energy. Any alternative functional form $\hbar(\epsilon)$ would fail at least one of these consistency checks. The situation is analogous to deriving the Schwarzschild metric with mass M as a free parameter: the derivation has genuine content (the functional form of the metric) even though one input is not determined from within the derivation.

The framework’s relation to Jacobson’s original derivation deserves explicit comment. Jacobson’s argument [Phys. Rev. Lett. 75, 1260 (1995)] uses $\hbar$ -containing forms throughout; the present derivation does not — it uses the classical identity $dE = (c^2\kappa/8\pi G) dA$ and the classical entropy density $\eta = 1/\epsilon^2$. The logical ordering differs from Jacobson’s: his argument derives G from η and $\hbar$; the present argument takes G as a parameter of the classical Hamiltonian and derives $\hbar$ from G and ϵ . The two are consistent — substituting $\eta = 1/\epsilon^2$ and

$\hbar = c^3\epsilon^2/(4G)$ into Jacobson’s formula $G_{\text{eff}} = c^4/(8\pi\eta \cdot \hbar c)$ recovers $G_{\text{eff}} = G$ — but the roles of input and output are reversed. The Gibbons-Hawking temperature $k_B T_{\text{GH}} = \hbar\kappa/(2\pi c)$ is recovered as a prediction, not used as an input.

7.4 Self-consistency and the discreteness scale

The gap equation $\hbar = c^3\epsilon^2/(4G)$ from Section 7.3 takes ϵ as a free geometric parameter — the discreteness scale of the substratum partition — and produces $\hbar$ as the output. The framework’s content beyond this single relation requires that ϵ itself be determined. The determination comes from self-consistency with the definition of the Planck length.

The Planck length $l_p = \sqrt{\hbar G/c^3}$ is defined as the unique length scale constructed from $\hbar$, G , and c . Substituting the framework’s relation $\hbar = c^3\epsilon^2/(4G)$ into this definition:

$$l_p^2 = \frac{\hbar G}{c^3} = \frac{c^3\epsilon^2}{4G} \cdot \frac{G}{c^3} = \frac{\epsilon^2}{4}.$$

Rearranging gives the discreteness scale:

$$\boxed{\epsilon = 2l_p.}$$

The result is the unique value of ϵ consistent with the framework’s gap equation and the standard definition of the Planck length. The framework’s substratum partition cell has linear extent $2l_p$, twice the Planck length.

The self-consistency is algebraically equivalent to the gap equation.

The framework contains one free geometric parameter; self-consistency with the Planck length definition fixes its value. This is not an additional empirical input but an algebraic consequence of the structural commitment $\hbar = c^3\epsilon^2/(4G)$ together with the definition of l_p . The result is a *constraint* on ϵ , not an independent measurement.

Stability of the discreteness scale. Two pathologies would obtain at incorrect values of ϵ . If $\epsilon^2 \ll l_p^2$ — sub-Planckian discreteness — the substratum partition would have cells finer than the Planck scale, creating a second trace-out structure with its own emergent description. Two trace-outs at different scales would assign different values of $\hbar$ to the same physical degrees of freedom, breaking the single-valuedness of $\hbar$ that all observers agree on. If $\epsilon^2 \gg l_p^2$ — super-Planckian discreteness — the substratum partition would group physically distinct configurations into single cells, with the emergent description assigning distinct quantum states to configurations that are operationally indistinguishable at the level of the substratum. Neither pathology is observed; the framework’s self-consistency fixes $\epsilon = 2l_p$ as the unique stable value.

The framework’s content in the discreteness scale therefore has three structural commitments: the gap equation, the Planck-length definition, and the

self-consistency between them. The result $\epsilon = 2l_p$ is a derived quantity, not a free parameter — but the derivation requires accepting the gap equation as the substratum’s structural commitment rather than $\hbar$ as a fundamental constant.

7.5 The Bekenstein-Hawking entropy

The Bekenstein-Hawking entropy $S = A/(4l_p^2)$ assigns each horizon a thermodynamic entropy proportional to its area, with the proportionality constant $1/(4l_p^2)$ — the factor of $1/4$ being a famously precise but historically unexplained constant. Standard derivations (Bekenstein, Hawking, the various semi-classical and stringy arguments) reproduce the area law but treat the $1/4$ coefficient as either fitted to thermodynamic consistency or derived from independent assumptions about the underlying microstates. The framework derives the $1/4$ coefficient as a direct consequence of $\epsilon = 2l_p$, with the entropy emerging by boundary-mode counting from the discreteness scale fixed in Section 7.4.

Direct mode counting. With $\epsilon = 2l_p$ from Section 7.4, the number of independent boundary modes on a horizon of area A is

$$N_{\text{modes}} = \frac{A}{\epsilon^2} = \frac{A}{4l_p^2}.$$

Each minimal cell of area $\epsilon^2 = 4l_p^2$ contributes one unit of entropy, giving $S = N_{\text{modes}} = A/(4l_p^2)$. This is the Bekenstein-Hawking entropy in its standard form.

The factor of $1/4$ is therefore derived: it is the ratio $\epsilon^2/l_p^2 = (2l_p)^2/l_p^2 = 4$ in the denominator, with the ϵ^2 factor arising from the framework’s discreteness scale and the l_p^2 in the denominator arising from the Planck length appearing in the definition of S_{BH} . Equivalently, the $1/4$ is the ratio $2\pi/8\pi$ between two physical quantities: the numerator 2π is the period of the Euclidean KMS regularity at the horizon (thermal periodicity of QFT on a Killing horizon), and the denominator 8π is the Jacobson coefficient in the Einstein-equation prefactor $dE = (c^2\kappa/8\pi G)dA$ (gravity’s coupling). The Bekenstein-Hawking $1/4$ is therefore not an arbitrary proportionality constant but the consistency condition between gravitational thermodynamics and quantum thermal physics — the ratio that makes the framework’s gap equation in Section 7.3 close on itself.

Extension to black-hole horizons. The minimal-cell counting argument applies to any horizon with area A . The cosmological horizon and black-hole horizons differ in their global structure — the former is a future causal boundary at infinity, the latter a past or future causal boundary at finite radius — but they share the local property of being bifurcate Killing horizons with surface gravity κ and area A . The framework’s mode-counting argument depends only on the local properties; the entropy $S = A/(4l_p^2)$ therefore extends to black-hole horizons with the same coefficient.

Observational confirmation: GW250114. The Bekenstein-Hawking area law $dA \geq 0$ is a theoretical prediction of black-hole thermodynamics derived from classical general relativity combined with the second law of thermodynamics. The empirical record before 2025 came primarily from agreement between classical horizon dynamics and semi-classical Hawking radiation predictions — supportive but indirect, with no direct test of the $1/4$ coefficient at the percent level.

The LIGO/Virgo/KAGRA observation of GW250114 in September 2025 provided the first direct test of the area law at the percent level. The event was the merger of two stellar-mass black holes; the inspiral phase of the gravitational waveform constrained the initial horizon areas of the two progenitors, and the ringdown phase constrained the remnant horizon area of the merged black hole. The comparison: the remnant horizon area exceeded the sum of the initial horizon areas by an amount consistent with the area theorem at 99.999% confidence — a result the framework, like any theory retaining classical general relativity, preserves.

The scope of this test is important to state precisely: GW250114 compares horizon areas and does not measure entropy, so it tests Hawking’s area theorem — classical horizon dynamics, which the framework preserves — and not the $1/4$ entropy coefficient. The coefficient, fixed by the framework’s derivation in Section 7.4 as the ratio $2\pi/8\pi$ between Euclidean KMS periodicity and the Jacobson Einstein-equation prefactor, is a derived structural constant with no direct empirical test at present; any framework retaining classical general relativity passes the GW250114 test identically. Untested is not untestable: the coefficient is equivalent to the Hawking-temperature normalization, so a measured horizon radiation spectrum would fix it. The one contingent channel is the final evaporation burst of a primordial black hole — active all-sky burst searches (Fermi-LAT 2018; HAWC 2020; LHAASO 2025, PRL 135 181005) have so far set rate-density limits only, with next-generation sensitivity projected for CTAO; a detected burst light curve would measure $T(M)$ and hence the coefficient directly. Analogue-gravity experiments probe the kinematic side of the $2\pi/8\pi$ ratio, not the gravitational side. The ongoing LIGO observing runs will accumulate additional events testing the area theorem across a range of progenitor masses and spins.

Falsification status. The Bekenstein-Hawking $1/4$ derivation is Class B-S Tier 1 per Chapter 4 §4.5 and Section 7.1 — unconditional structural at the universality-class level. Falsification of the coefficient at sufficient confidence to exclude experimental systematics would refute the framework’s gap equation chain (the connection between the classical horizon temperature in Section 7.3 Step 1 and the Euclidean KMS temperature in Step 4) and therefore refute the framework’s structural commitment to thermal self-consistency at the partition boundary. The coefficient currently has no direct empirical test — GW250114 constrains the classical area theorem, not the entropy normalization — so the framework’s strongest gravitational empirical anchor is the $a_0 = cH/6$ rotation-

curve family (Sections 7.8–7.9).

7.6 The Page curve and the information paradox

The framework’s content in the Bekenstein-Hawking sector extends beyond the area-law coefficient to the *dynamics* of black hole evaporation. The black hole information paradox — Hawking’s 1975 observation that semiclassical black hole evaporation appears to destroy information, in tension with quantum mechanics’ commitment to unitarity — is one of the canonical hard problems of theoretical physics. The framework resolves the paradox at the structural level: information is never lost from the substratum bijection φ because φ is bijective. The apparent information loss is an artifact of the trace-out — the same mechanism that produced the cosmological constant dissolution in §7.7 below and the dark sector concordance in §7.9.

The paradox in its conventional form. Hawking’s calculation of the spectrum of radiation emitted by an evaporating black hole produced a thermal distribution with a temperature $T_H = \hbar\kappa/(2\pi ck_B)$, dependent only on the black hole’s mass and angular momentum. The thermal spectrum carries no information beyond the macroscopic horizon parameters. If the radiation is exactly thermal — as Hawking’s semiclassical derivation suggests — then the radiation carries no information about the matter that formed the black hole, and the information appears to be lost when the black hole completes its evaporation. This is in tension with the unitarity of quantum mechanics, which requires that any time evolution preserve information.

The Page curve is the unitarity prediction. Don Page argued in 1993 that if black hole evaporation is unitary, the entanglement entropy of the radiation must follow a specific trajectory: rising as radiation accumulates in the early phase, then turning over and decreasing at the *Page time* t_P — defined as the moment when the radiation’s entanglement entropy equals the remaining black hole’s entropy. The Page time occurs at approximately half the black hole’s entropy has been radiated; the curve’s turnover is the empirical signature of unitarity.

The dispute between Hawking’s calculation (thermal radiation, information lost) and Page’s prediction (Page curve, information preserved) was unresolved for decades. The recent quantum-extremal-surface and entanglement-island programs of Penington, Almheiri-Engelhardt-Marolf-Maxfield, and others have reproduced the Page curve through replica wormhole computations in the gravitational path integral, providing a derivation of the Page curve from semiclassical gravity. The framework provides an alternative derivation through a substantially different mechanism — and the two derivations converge on the same quantitative predictions.

The framework’s resolution: nested trace-out. The framework’s substratum is a deterministic bijection φ acting on a finite configuration space S . Information is never lost from φ — bijectivity is the framework’s structural

commitment to unitarity at the substratum level. The apparent information loss in Hawking radiation is an artifact of the *observer's* description, not of the underlying dynamics.

The mechanism is the nested-partition formalism. When a black hole forms within the cosmological horizon, the framework's partition structure has two co-existing causal boundaries: the cosmological horizon (the framework's primary partition) and the black hole horizon (a secondary partition within the visible sector). The configuration space decomposes as

$$S = \mathcal{C}_V \times \mathcal{C}_B \times \mathcal{C}_D,$$

where V is the region between the black hole horizon and the cosmological horizon (the external observer's visible sector), B is the black hole interior (the secondary hidden sector), and D is the trans-cosmological region (the primary hidden sector).

The external observer's emergent quantum description of the region V is obtained by tracing out *both* B and D . Two procedures produce a quantum description on V : a *direct* trace, marginalizing over $\mathcal{C}_B \times \mathcal{C}_D$ in one step, and a *sequential* trace, first marginalizing over $\mathcal{C}_D$ (cosmological partition) and then over $\mathcal{C}_B$ (black hole partition). The consistency of the nested formalism — the requirement that the two procedures produce the same CPTP channel, the same emergent Hamiltonian, and the same value of $\hbar$ — is established in the framework's appendix to the gravitational sector. The nested trace-out is well-defined; the external observer's emergent quantum mechanics is self-consistent when multiple causal boundaries coexist.

The Page curve from cycle typicality. The framework's derivation of the Page curve runs through the *combinatorial* structure of the bijection φ rather than through the gravitational path integral. The relevant result is the Popescu-Short-Winter typicality theorem: for a bipartite quantum system in a generic pure state, the marginal entropy of the smaller subsystem is approximately maximal, with corrections that vanish exponentially in the dimension of the larger subsystem.

Applied to the evaporating black hole, the relevant bipartite structure is the radiation field outside the black hole versus the boundary modes remaining at the shrinking horizon. Let $n_B(t)$ denote the number of boundary modes remaining at the horizon at time t , and $n_R(t) = n_B(0) - n_B(t)$ the number of modes radiated. Conservation gives $n_B(t) + n_R(t) = n_B(0)$ — the total number of modes is fixed by the initial black hole entropy.

For the typical state on the energy shell — typical in the sense of Popescu-Short-Winter measure concentration applied to the bijection's cycle structure — the radiation's entanglement entropy is approximately

$$S_{\text{rad}}(t) = \min(n_R(t), n_B(t)) \cdot \ln q,$$

where q is the substratum’s alphabet size (a gauge quantity that cancels in physical ratios) and the corrections are exponentially small. The structure has two regimes:

- *Early phase* ($n_R(t) < n_B(t)$): radiation has fewer modes than the remaining black hole; entropy rises as $n_R(t) \ln q$.
- *Late phase* ($n_R(t) > n_B(t)$): radiation has more modes than the remaining black hole; entropy falls as $n_B(t) \ln q$.

The turnover — the Page time t_P — occurs at $n_R(t_P) = n_B(t_P) = n_B(0)/2$, when exactly half the black hole’s initial entropy has been radiated. The structural Page time is

$$t_P = (1 - 2^{-3/2}) t_{\text{evap}} \approx 0.646 t_{\text{evap}},$$

independent of the black hole’s mass and the species-dependent greybody factor. The Page time’s structural status — its independence from solution-specific parameters — is the framework’s distinctive content: the timescale at which information transfer becomes evident in the radiation is fixed by the framework’s structural commitments to bijectivity and to the boundary-mode counting from §7.5.

Comparison with quantum extremal surfaces and islands. The framework’s Page curve derivation quantitatively matches the prediction from the quantum extremal surface formula and the entanglement-island construction developed by Penington, Almheiri-Engelhardt, and others. In the QES approach, the Page curve emerges through an extremization principle that selects the dominant saddle in a replica-trick computation of the radiation’s entanglement entropy. The post-Page-time saddle includes an *island* — a region of bulk spacetime that contributes to the radiation’s entanglement entropy despite being topologically disconnected from the radiation region. The replica wormhole geometries connecting the replicas in the gravitational path integral are essential to obtaining the saturating Page curve at late times.

The framework arrives at the same Page time and the same Page entropy through a structurally cleaner mechanism. The cycle-typicality argument invokes only the Popescu-Short-Winter theorem applied to the deterministic substratum dynamics on the energy shell, with no replica trick, no ergodicity assumption, and no gravitational path integral. The *island* in the QES formulation corresponds, in the framework’s construction, to the post-Page-time regime in which the black hole’s boundary modes form the smaller subsystem; the *replica wormhole* geometries appear to correspond to the substratum’s cycle structure on the energy shell.

This conjectured correspondence — replica wormholes $\leftrightarrow$ substratum cycles — is suggestive but not yet rigorous. Making it precise would require constructing the gravitational path integral as an effective description of the substratum dynamics in the appropriate limit, a task that has not been completed. What the

framework’s derivation establishes unconditionally is that the Page curve’s quantitative content is reproducible from a finite deterministic substratum without invoking the replica machinery.

The convergence is structural rather than coincidental. Both approaches require *something* to play the role of typicality — Popescu-Short-Winter measure concentration in the framework’s derivation, replica saddles in QES + islands — and both extract the same Page time and Page entropy from that ingredient. The framework’s content here is that the typicality-providing ingredient need not be gravitational at all; it can be combinatorial, on the substratum side of the trace-out. This constitutes a microscopic realization of the entanglement-island construction.

The generalized second law. A separate but related structural prediction is the generalized second law: the sum $S_{\text{matter}} + S_{\text{BH}} + S_{\text{dS}}$ is non-decreasing in time. The matter entropy S_{matter} captures the visible-sector matter’s information content; the black hole entropy S_{BH} is the horizon area divided by $4l_p^2$; the de Sitter entropy S_{dS} is the cosmological horizon’s contribution.

The framework proves the generalized second law from strong subadditivity of the emergent quantum description combined with the CPTP monotonicity of the framework’s reduced dynamics. Any process that shrinks the black hole — decreasing S_{BH} — must transfer information to the visible sector matter (increasing S_{matter}) or to the cosmological boundary (increasing S_{dS}), compensating the decrease. The total information across the partition is conserved at the substratum level (bijectivity) and the entropy sum is non-decreasing at the observer’s level (CPTP monotonicity). The two facts are consistent in the framework: the substratum is deterministic and information-preserving, while the observer’s reduced description has the entropy-non-decrease that the second law requires.

Empirical status and falsification. The framework’s resolution of the information paradox has no direct empirical signature accessible at current experimental precision — black hole evaporation timescales for astrophysical black holes vastly exceed the age of the universe, and the Page time for a stellar-mass black hole would not occur within any conceivable observational program. The empirical content is therefore conceptual rather than quantitative: the framework’s structural commitment to bijectivity entails that information is never lost from the substratum, with the apparent information loss being an artifact of the observer’s trace-out.

Falsification of the framework’s resolution would require evidence that black hole evaporation is genuinely non-unitary at the level of the underlying dynamics — for instance, an empirical demonstration that the radiation from an evaporating black hole carries less information than the matter that formed it, in a way that cannot be reproduced by the framework’s nested trace-out. No such evidence currently exists, and the convergence with the QES + islands programs provides additional theoretical support for the unitarity resolution. The framework’s con-

tent is therefore consistent with the current theoretical and (limited) empirical picture, with the structural Page time prediction $t_P \approx 0.646 t_{\text{evap}}$ providing a sharp quantitative target for any future empirical or theoretical work on the information paradox.

7.7 Dissolution of the cosmological constant problem

The cosmological constant problem is conventionally described as the worst quantitative discrepancy in physics. Quantum field theory predicts that the vacuum energy density of the universe should be of order $\rho_{\text{QFT}} \sim M_{\text{Pl}}^4 \sim 10^{113} \text{ J/m}^3$, obtained by summing zero-point energies $\frac{1}{2}\hbar\omega$ over all field modes up to the Planck cutoff. The observed vacuum energy density is $\rho_{\text{obs}} \sim H^2/G \sim 10^{-9} \text{ J/m}^3$, with the ratio $\rho_{\text{QFT}}/\rho_{\text{obs}} \sim 10^{122}$. Standard accounts treat this as a fine-tuning problem requiring an unknown cancellation mechanism to suppress the bare vacuum energy by 122 decimal places, re-imposed at each order of perturbation theory.

The framework dissolves the problem rather than solving it. The 10^{122} discrepancy is not a fine-tuning failure but a category error — a comparison between two quantities that are properties of logically distinct descriptions and are not directly commensurable as gravitational sources. The dissolution has three components.

The two levels of description. The framework’s structural commitments produce two distinct descriptions of the universe’s energy content, operating at logically different levels.

The classical substratum assigns the universe a vacuum energy through the Friedmann equation: $\rho_{\text{crit}} = 3H^2c^2/(8\pi G)$. This is the classical baseline — what geometric measurements probe, what gravity couples to, what the standard tests of general relativity confirm. There is no zero-point energy at this level; the vacuum energy is set by the Hubble parameter H through classical thermodynamics. Numerically, $\rho_{\text{crit}} \sim H^2/G \sim 10^{-9} \text{ J/m}^3$ — precisely what is observed.

The emergent quantum field theory assigns the universe a vacuum energy by summing zero-point energies over its field modes. With the framework’s UV cutoff at $\epsilon = 2l_p$, the sum has cutoff at the Planck scale, giving $\rho_{\text{QFT}} \sim M_{\text{Pl}}^4 \sim 10^{113} \text{ J/m}^3$. This is what local quantum measurements probe, what perturbative quantum field theory describes — but it is a property of the *emergent description*, not of the classical substratum.

The two values describe the same universe at different levels. The classical substratum’s vacuum energy is the gravitating object; the emergent QFT’s vacuum energy is bookkeeping internal to the emergent description. The two are not commensurable as gravitational sources because they are properties of different things.

Why gravity sees the classical value. Spacetime geometry is part of the

classical substratum: the metric evolves via Einstein’s equations *before* the trace-out that produces the emergent quantum description. The stress-energy that sources gravity is the classical stress-energy of the total microstate, not the expectation value of an emergent quantum operator. The semi-classical equation $G_{\mu\nu} = 8\pi G\langle\hat{T}_{\mu\nu}\rangle$ is itself an emergent approximation; at the classical level, the gravitational field equations never encounter the zero-point sum.

Three structural requirements force this geometry-first ordering of the framework’s derivation chain.

Definiteness. The trace-out that produces the emergent quantum mechanics is defined by a specific partition $\Gamma = \Gamma_V \times \Gamma_H$. The partition must be definite — not in superposition, not subject to quantum uncertainty — for the marginalization producing the emergent description to be well-defined. A partition in superposition would yield an incoherent mixture of inequivalent quantum theories, not a single QM. The metric at the partition boundary must therefore be classical.

Non-circularity. The partition is defined by causal structure: the hidden sector contains the degrees of freedom beyond the observer’s causal reach. Causal structure is determined by null geodesics of the metric. If the metric were derived from QM, the derivation would require QM $\rightarrow$ metric $\rightarrow$ causal structure $\rightarrow$ partition $\rightarrow$ QM — a circular dependency with no entry point. The circle breaks only if the metric exists prior to the partition.

Uniqueness of $\hbar$. The emergent action scale is determined by the boundary geometry, as Section 7.3 developed. If the geometry were itself quantum-mechanical, $\hbar$ would depend on a quantum state — but $\hbar$ is state-independent by observation. The geometry-first ordering is the only one consistent with a universal $\hbar$.

The framework’s geometry-first ordering is therefore not a stipulation but a structural consequence of the framework’s three commitments — definiteness, non-circularity, and the universality of $\hbar$.

The structural origin of the 10^{122} discrepancy. The ratio $\rho_{\text{QFT}}/\rho_{\text{classical}} \sim M_{\text{Pl}}^4/(H^2/G) = S_{\text{dS}} = A/(4l_p^2)$ — the Bekenstein-Hawking entropy of the cosmological horizon, identified by Section 7.5 as the information-compression ratio of the trace-out from substratum to emergent description. The “worst prediction in physics” is the entropy of the observer’s blind spot.

The identification is exact, not approximate. The 10^{122} factor is the de Sitter entropy $S_{\text{dS}} = A_{\text{horizon}}/(4l_p^2)$ at the present cosmological horizon, computed from the observed horizon area combined with the framework’s entropy formula. The factor that the conventional QFT view treats as a discrepancy requiring cancellation is the framework’s structural quantity for the compression ratio between substratum modes and visible-sector emergent modes.

No embedded observer can determine from within whether the geometric or the

quantum description is more fundamental, and not merely as a matter of present ignorance. Any adjudication procedure available to an embedded observer is a function of the observer’s accessible data, and both descriptions are constructed from the same accessible data — so no criterion grounded in the accessible observations themselves can separate them. (Rankings by theoretical virtue — simplicity, unification — are not blocked by this factoring; but they order presentations of the same observational content rather than track fundamentality, and the framework’s answer to the fundamentality question is supplied at the substratum level rather than adjudicated from within either description.) The non-adjudicability follows from the equivalence itself, not from limited resources: a future observation could break the equivalence, but that would be falsification of the equivalence, not adjudication between the descriptions. Nor is this factoring limitation an artifact of the framework’s particular observer architecture. Wolpert’s limits of inference and T. Breuer’s self-measurement theorem establish barriers of the same type for any embedded inference device under far weaker assumptions, corroborating that the limitation is device-general rather than model-specific — with the scope note that those results, unlike the factoring argument, assume free availability of the device’s setups. The framework’s content here is not to identify the geometric description as “more real” but to identify the 10^{122} factor as the structural quantity that the framework has independent reason to expect — not a coincidence requiring explanation.

Convergence with the de Sitter observer literature. The framework’s geometry-first dissolution converges with recent developments in the de Sitter observer literature, in which an unphysical phase of the de Sitter sphere partition function — long understood to obstruct a clean dS thermodynamic interpretation — cancels exactly when the formalism incorporates an observer with a clock. Maldacena established the cancellation in 2024; Chandrasekaran, Longo, Penington, and Witten (CLPW) showed that the von Neumann algebra of QFT observables in a gravitational subregion promotes from type III to type II once an observer is included, with the resulting entropy pinned to the observer’s reference data. Harlow, Usatyuk, and Zhao established that the closed-universe Hilbert-space dimension is determined by the observer’s degrees of freedom.

The framework provides a microscopic realization of these results: the partition trace-out is the *defining* operation rather than an added ingredient, so the observer-essentiality structure that Maldacena’s argument requires is structurally forced. The observer-essentiality literature converges with the framework on substantive content while differing in mechanism — a triangulation that strengthens the empirical case for both programs. Chapter 9 develops the framework’s relationship to this literature in detail.

The classical vacuum energy scale. The framework dissolves the QFT cosmological constant problem but does not explain *why* $\rho \sim H^2/G$ takes its specific observed value. This is set by initial conditions through the Friedmann equation, and the framework reclassifies the problem from quantum vacuum cancellation to classical initial conditions. The two problems are not equivalent.

The QFT cosmological constant problem requires a 122-decimal-place cancellation between independent contributions — the bare cosmological constant and the zero-point energies of every quantum field — re-tuned at each order of perturbation theory. The initial-conditions question asks why a single parameter takes its observed value; the scale $\rho \sim H^2/G$ is the natural energy density of classical cosmology, not a fine-tuned cancellation between unrelated quantities. The framework eliminates the 122-digit cancellation entirely; the remaining question — why H takes its value — is a standard cosmological problem addressable by inflationary dynamics or other initial-conditions mechanisms.

7.8 Dark energy in running-vacuum form

The framework’s content in the dark-energy sector is structural rather than quantitative. Section 7.7 established that the emergent quantum field theory’s vacuum energy does not gravitate, with the cosmological constant problem dissolved through the two-level description. A distinct question concerns the *evolution* of the effective cosmological constant with cosmic expansion. The framework predicts that the effective cosmological constant has a specific functional form — Type II running-vacuum — but with a coefficient bounded at $|\nu_{\text{OI}}| \sim 10^{-32}$ that is indistinguishable from zero at all foreseeable observational precision.

One scoping remark keeps the registers aligned: the coefficient bound comes from the emergent quantum theory’s curvature response, while Section 7.7’s dissolution denies that the emergent vacuum energy gravitates absolutely. Whether that response feeds back as a source (in which case its sign is computable and negative for natural couplings, dominated by the top quark) or remains the observer’s effective attribution (in which case the coefficient is set by substratum geometry alone, with sign open) is left explicitly unadjudicated — because both readings make the same observable prediction: dark energy indistinguishable from a cosmological constant at every foreseeable precision, with any detected evolution in the DESI-favored direction attributable only to substratum structure carrying a correlated Lorentz-violation signature.

The Type II running-vacuum form. The framework’s emergent quantum vacuum energy depends on the cosmological horizon’s area, which in turn depends on the Hubble parameter through $A_{\text{horizon}} = 4\pi c^2/H^2$. Any smooth dependence of ρ_Λ on H admits a Taylor expansion around the present-day value H_0 :

$$\Lambda_{\text{eff}}(H) = \Lambda_0 + \frac{3\nu}{8\pi G}(H^2 - H_0^2) + \mathcal{O}((H^2 - H_0^2)^2).$$

This is the Running Vacuum Model. The framework’s content lies not in the existence of an H^2 correction — any framework with an H -dependent vacuum energy admits this form — but in three structural commitments that force the framework into the specific *Type II* class of RVM implementations.

Boundary-entropy dependence. The de Sitter entropy $S_{\text{dS}} = A_{\text{horizon}}/\epsilon^2$ depends on H through the horizon area, so the framework’s emergent vacuum energy inherits an H -dependence at the boundary.

Matter conservation. The framework’s trace-out ontology conserves matter locally; no direct vacuum-matter energy exchange is permitted. This forbids the Type I RVM variants in which a phenomenological exchange term $Q = 3\nu H \rho_{\text{dm}}$ transfers energy between dark energy and dark matter.

Bianchi consistency. The Einstein equations’ Bianchi identity, combined with matter conservation, requires that Newton’s constant $G(H)$ run mildly in tandem with $\rho_{\Lambda}(H)$ — a structural feature that distinguishes Type II RVM from Type I and from $w_0 w_a$ CDM parameterizations.

These three commitments place the framework in the Type II class: matter is locally conserved, $G(H)$ runs with cosmic expansion, and the Bianchi identity is preserved by the joint running of $\rho_{\text{vac}}(H)$ and $G(H)$. Type I RVM variants — sharing the same functional form for $\rho_{\text{vac}}(H)$ but assuming $G = G_N$ fixed with a phenomenological vacuum-DM exchange — are physically distinct models that the framework does not predict.

Magnitude from the emergent-QFT channel. The framework’s structural commitments fix the functional form but the coefficient ν requires a quantitative calculation. The relevant calculation is the one-loop renormalization-group running of the emergent vacuum energy in the curved-spacetime background of the cosmological horizon — a standard computation within emergent quantum field theory, applied with the framework’s specific field content.

The Shapiro-Solà β -function for the vacuum-energy running coefficient is

$$\nu = \frac{1}{48\pi^2} \sum_i a_i \frac{M_i^2}{M_{\text{Pl}}^2},$$

where the sum runs over the emergent fields with masses M_i and a_i are species-dependent coefficients of order unity. Applied to the Standard Model content — with the heaviest mass being the top quark at $m_t \approx 173$ GeV — the framework’s prediction is

$$|\nu_{\text{OI}}| \sim (M_{\text{SM}}/M_{\text{Pl}})^2 \sim 10^{-32}.$$

This is observationally indistinguishable from zero at all foreseeable observational precision. The framework’s prediction is therefore that the effective cosmological constant has the *form* of Type II running-vacuum dynamics, but with a coefficient too small to be measured by any current or planned instrument.

Potential contributions from the substratum. A separate contribution might arise from the substratum’s boundary mode count, in principle larger than the emergent-QFT value if the substratum modes have a distinct frequency distribution. The framework’s boundary architecture from Section 7.3 catalogs modes by spatial location alone (one degree of freedom per cell of area ϵ^2), with no inherent frequency assignment. Under the natural two-dimensional

dispersion relation $dN/d\omega \propto \omega$, modes concentrate near the UV cutoff and contribute a fraction $(H\epsilon)^2 \sim 10^{-122}$ at horizon frequency — well below the emergent-QFT magnitude. Absent additional substratum structure giving a log-uniform boundary-mode distribution at the cosmological horizon, the framework’s prediction reduces to $|\nu| \approx |\nu_{\text{QFT}}|$, observationally indistinguishable from zero. That absence has since been computed rather than assumed: extracting the boundary spectral measure directly from the substratum wave equation (an exact surface-Green’s-function computation, Section 19.3.7) yields a power-suppressed local distribution — $dN/d\omega \propto \omega^4$, with the two-dimensional $\propto \omega$ above an upper bound on the infrared weight — stable over thirty orders of magnitude in the slow-bath ratio, with the log-uniform distribution appearing nowhere; the interacting extension of that computation is the remaining residual.

Observational status. The empirical situation in late 2025 and early 2026 has evolved substantially with the DESI Data Release 2 results. The DESI collaboration reports evolving dark energy at $2.8\text{--}4.2\sigma$ significance in the standard $w_0 w_a$ parameterization, with the data favoring a present-day equation of state $w_0 > -1$ transitioning to $w(z) < -1$ at higher redshift — a “phantom crossing” pattern. The framework’s prediction is $w(z) \approx -1$ structurally, with the deviations of order $\nu \sim 10^{-32}$. The mild tension with DESI DR2 is $3\text{--}4\sigma$ on the phantom-crossing pattern, depending on parameterization choices and prior assumptions.

A direct test of the framework’s specific Type II RVM prediction was provided by Bertini, Novaes, von Marttens, and Shapiro in late 2025, who fit a renormalization-group-corrected Λ CDM model — structurally equivalent to Type II RVM at leading order in ν , with both $G(\mu)$ and $\Lambda(\mu)$ running and matter locally conserved — against the Planck 2018 + DESI DR2 + DES-Y5 combined data. Their result:

$$\nu^{\text{Bertini}} = -(2.5 \pm 1.3) \times 10^{-4},$$

with the Λ CDM limit at the 2σ boundary of the posterior. This is consistent at 2σ with the framework’s prediction $\nu \approx 0$. An earlier pre-DESI canonical RVM fit by Gómez-Valent and Solà returned $\nu = (1.34 \pm 0.38) \times 10^{-3}$ with 3.5σ preference over Λ CDM; the migration from this value toward Bertini’s smaller value as the data improved from pre-DESI to DESI DR2 is itself informative — the empirical evidence is shifting toward consistency with the framework’s near-zero prediction.

Decisive test structure. DESI Year 5 will provide approximately 1% precision on ν , available in 2028–2029. The framework’s prediction makes three discriminating outcomes possible. A null result at that precision *confirms* the framework’s prediction $\nu \approx 0$ and disconfirms the DESI DR2 phantom-crossing preference. A high-confidence detection of $|\nu| \sim 10^{-4}$ would confirm running-vacuum dynamics consistent with both Type II RVM structure and DESI DR2 trends, but at a magnitude requiring substratum structure beyond what the

framework currently derives. A high-confidence detection of phantom crossing would falsify the framework’s prediction at structural level.

The current mild tension with DESI DR2 is the kind of healthy empirical exposure that distinguishes a framework engaged with data from one tuned to fit it. The framework’s prediction was made structurally before the DESI data and stands as a falsifiable commitment for the 2028–2029 timeframe.

Coupled neutrino test. The DESI DR2 + CMB analysis finds $\Sigma m_\nu < 0.0642$ eV (95% CL) assuming Λ CDM, prefers normal mass ordering, and bounds the lightest neutrino mass at $m_l < 0.023$ eV — all consistent with the framework’s neutrino-sector predictions of normal ordering with Σm_ν near the oscillation minimum of 0.059 eV. The same analysis finds a 3σ tension with the lower oscillation limit assuming Λ CDM, interpreted in the literature as “a hint of new physics not necessarily related to neutrinos” — the kind of structural mismatch the framework’s RVM dark energy would *relax*: an evolving-dark-energy analysis loosens the Λ CDM Σm_ν bound and comfortably accommodates the framework’s ~ 0.059 eV prediction, though the RVM-corrected bound has not yet been computed. The neutrino sector predictions developed in Chapter 6 and the RVM dark energy prediction here are coupled in the framework: both follow from the same lattice structure, and the joint observational support is therefore not independent confirmation of separate effects but a single empirical pattern consistent with the framework’s central mechanism.

7.9 The dark sector and the MOND acceleration scale

The framework’s content in the dark-matter sector is structural rather than particulate. The conventional cold-dark-matter paradigm postulates a new fundamental particle species that interacts gravitationally and contributes about 27% of the universe’s energy budget. Decades of direct-detection experiments — XENON, LUX, SuperCDMS, PandaX, and others — have produced no detection at any expected scale. The framework predicts that no such particle exists: the apparent dark-matter phenomenology is instead a structural consequence of entropy displacement at the cosmological boundary, with the MOND acceleration scale $a_0 = cH/6$ deriving from a dimensional argument in de Sitter space.

The invisible gravitational budget as concordance. A corollary of the framework’s trace-out structure is that the cosmological horizon carries a gravitational budget that the emergent quantum field theory cannot access. The argument runs in three steps.

The horizon carries $S_{\text{dS}} = A_{\text{horizon}}/\epsilon^2$ modes by the boundary-mode counting of Section 7.5. The classical temperature is $T_{\text{cl}} = c^3\epsilon^2 H/(8\pi G k_B)$ from Section 7.3 — computed entirely within the classical substratum with no reference to the emergent quantum description. The total thermal energy carried by the boundary modes is $E_s = S_{\text{dS}} \cdot k_B T_{\text{cl}}$. Distributing this energy over the Hubble

volume $V_H = 4\pi c^3/(3H^3)$ via the shell theorem:

$$\rho_s = \frac{E_s}{V_H} = \frac{A}{\epsilon^2} \cdot \frac{c^3 \epsilon^2 H}{8\pi G} \cdot \frac{3H^3}{4\pi c^3} = \frac{3c^2 H^2}{8\pi G} = \rho_{\text{crit}}.$$

The factor of ϵ^2 cancels between $S_{\text{ds}} \propto 1/\epsilon^2$ and $T_{\text{cl}} \propto \epsilon^2$, so the result is independent of the specific value of ϵ — only Jacobson’s $dE = (c^2 \kappa/8\pi G) dA$, an area-law entropy, and the Friedmann volume enter the calculation. The total effective energy density carried by the boundary modes equals the critical density of the universe.

This is the framework’s account of the universe’s apparent dark-sector content. The $\sim 95\%$ of the universe’s gravitational budget that is invisible to the emergent quantum field theory corresponds to the boundary modes that the trace-out moves into the hidden sector. The dark sector is not a separate substance requiring a new particle species; it is the framework’s structural quantity for the partition’s hidden-sector content at the cosmological horizon.

Three clarifications are essential. The calculation above uses only classical, pre-trace-out quantities: the number of boundary modes, the classical temperature, and the Hubble volume. No step invokes the emergent quantum description. This is what distinguishes ρ_s from the QFT zero-point energy $\rho_{\text{QFT}} \sim M_{\text{Pl}}^4$: the latter is a property of the emergent description (does not gravitate), while the boundary entropy’s thermal energy is a classical quantity that predates the trace-out (does gravitate). The boundary modes are physically localized on the horizon — the uniform-medium description used in the shell theorem calculation is a gravitational equivalence, not a claim about physical delocalization. The usefulness of the uniform-medium description is computational: when matter is present, it locally deforms the horizon geometry and redistributes the boundary entropy, with the gravitational effect of the redistribution on bulk observers being most easily computed in the uniform-medium representation.

The MOND acceleration scale. The framework’s account of dark-matter phenomenology — flat rotation curves, the baryonic Tully-Fisher relation, the radial acceleration relation — derives from entropy displacement at the cosmological boundary. The argument has four steps.

Step 1: Entropy displacement by baryonic matter. The Bekenstein bound applied to a baryonic mass M_B in a de Sitter background gives the entropy displaced from the cosmological boundary:

$$\Delta S = \frac{M_B c^2}{k_B T_{\text{ds}}} = \frac{2\pi M_B c^3}{\hbar H}.$$

This follows from the generalized second law applied to the matter-boundary system: the boundary’s entropy reduction is Q/T for a reversible process, where $Q = M_B c^2$ and $T = T_{\text{ds}}$.

Step 2: The critical acceleration $a_0 = cH/6$. The displaced entropy redistributes over the volume between the matter and the horizon. The Jacobson mechanism converts the entropy gradient into additional radial acceleration. The conversion is dimensional. In d spacetime dimensions, the factor relating the natural Hubble acceleration scale cH to the empirical MOND acceleration is $(d-3)/[(d-2)(d-1)]$, which follows from the interplay between volume-law entropy distributed over $(d-1)$ spatial dimensions and area-law entropy scaling with $(d-2)$ surface dimensions, together with the standard surface-density-to-acceleration relation via Gauss's law. In our universe with $d = 4$:

$$a_0 = \frac{d-3}{(d-2)(d-1)} \cdot cH = \frac{1}{2 \cdot 3} \cdot cH = \frac{cH}{6} \approx 1.2 \times 10^{-10} \text{ m/s}^2.$$

This is the MOND critical acceleration. Milgrom's empirical value matches at the percent level. The factor $1/6$ has dimensional origin: it is structurally fixed by $d = 4$ together with the volume-vs-area entropy interplay, not fitted to observation.

Step 3: The deep-MOND regime and baryonic Tully-Fisher. For baryonic acceleration much smaller than a_0 , the displaced entropy dominates the dynamics. The Jacobson mechanism converts the entropy gradient into curvature, giving $g_{\text{total}} = \sqrt{g_B \cdot a_0}$ where $g_B = GM_B/r^2$ is the baryonic Newtonian acceleration. The rotation velocity at radius r then satisfies $v^2 = r \cdot g_{\text{total}} = \sqrt{GM_B \cdot cH/6} = \text{constant}$, so:

$$v^4 = GM_B \cdot \frac{cH}{6}.$$

This is the baryonic Tully-Fisher relation. Flat rotation curves emerge with no dark-matter particles and no free parameters beyond the baryonic mass.

Step 4: Redshift evolution. The MOND scale depends on the Hubble parameter, so $a_0(z) = cH(z)/6$ increases with redshift. The MOND crossover radius $r_M = \sqrt{6GM_B/(cH)}$ correspondingly shrinks at high redshift: for a Milky-Way-like baryon budget ($M_B \approx 6 \times 10^{10} M_\odot$ for stellar plus cold gas), $r_M = 9.0$ kpc at $z = 0$ but 5.2 kpc at $z = 2$ and 3.1 kpc at $z = 5$. Galaxies with effective radii comparable to or exceeding r_M show flat rotation curves; galaxies with $r_{\text{eff}} < r_M$ show declining Keplerian-like rotation curves. Since r_M shrinks with increasing redshift, high-redshift galaxies are systematically more baryon-dominated, with their rotation curves declining at smaller radii than local spirals.

Observational confirmations. Two independent observations support the framework's account of dark-matter phenomenology with the specific $H(z)$ -dependent crossover that the framework predicts.

The cleanest local test is Jiao et al. (2023), who detected for the first time a Keplerian decline in the Milky Way's own rotation curve from approximately 19 to 26.5 kpc using Gaia DR3 kinematics, with the flat rotation curve hypothesis rejected at 3σ significance. The framework's prediction for the Milky Way, using the full baryon budget including the hot circumgalactic gas halo ($M_B \sim 2 \times$

$10^{11} M_{\odot}$), is the crossover at $r_M \sim 17$ kpc — essentially where Jiao et al. observe the transition. The Milky Way is the single strongest piece of evidence for the $H(z)$ -dependent crossover prediction at $z = 0$.

The high-redshift test is Genzel et al., who report stacked rotation curves for massive star-forming galaxies at $z = 0.9$ – 2.4 showing declining outer velocities at greater than 3σ significance relative to local spirals. The framework’s prediction — that high-redshift galaxies should be more baryon-dominated due to the shrinking crossover radius r_M — is consistent with the observed declining velocities at high redshift. In the Λ CDM paradigm, flat rotation curves from NFW halos are expected at all redshifts; explaining the Genzel data within Λ CDM requires anomalously low halo concentrations or baryon-dominated inner regions. The framework’s prediction sits naturally with the data without parameter adjustment.

The interpolation function and solution-specific content. The framework’s structural commitments fix the deep-Newtonian and deep-MOND limits and the redshift evolution of the crossover, but the full interpolation between the two regimes is constrained rather than uniquely determined. Five constraints — no free parameters, conservative, monotonic, correct limits, analytic — narrow the interpolation to a one-parameter family indexed by the transition steepness. The transition steepness is a property of the particular bijection φ , in the same Tier 3 category as fermion masses (solution-specific rather than structural). The simplest member satisfying all five constraints is the “simple” MOND interpolation function $\nu(y) = (1 + \sqrt{1 + 4/y})/2$, which matches the empirical radial acceleration relation of McGaugh et al. to better than 0.1 dex across all galaxy-relevant accelerations.

The framework’s content in the dark-matter sector therefore has two-tier structure. The MOND acceleration scale $a_0 = cH/6$, the baryonic Tully-Fisher relation, and the redshift evolution of the crossover are Tier 2 — universal across bijections in the framework’s universality class. The transition steepness is Tier 3 — bijection-specific, with the simple MOND interpolation as the framework’s natural default.

7.10 Summary of gravitational predictions and chapter close

The framework’s content in the gravitational sector reduces a substantial portion of standard physics — the value of Planck’s constant, the Bekenstein-Hawking entropy with its $1/4$ coefficient, the cosmological constant problem, the running-vacuum form of dark energy, the dark-sector budget, the MOND acceleration scale and the baryonic Tully-Fisher relation — to structural consequences of the cosmological horizon as a causal partition combined with the framework’s three structural conditions C1-C3. The empirical record summarized below collects the chapter’s quantitative predictions with their tier classifications and current observational status.

Summary table.

Observable	Predicted	Observed	Match	Tier
$\hbar = c^3 \epsilon^2 / (4G)$	from horizon thermodynamics	$1.055 \times 10^{-34} \text{ J} \cdot \text{s}$	exact (given $\epsilon = 2l_p$)	1
$\epsilon = 2l_p$	from self-consistency	—	structural	1
$S_{\text{BH}} = A/(4l_p^2)$	factor 1/4 derived as $2\pi/8\pi$	derived; no direct test (GW250114 tests the area theorem — next row)	—	1
Area law $dA \geq 0$	from generalized second law	GW250114 area increase	99.999%	1
CC dissolution	$\rho_{\text{QFT}}/\rho_{\text{class}} = S_{\text{ds}}$	$\rho_{\text{obs}} \sim H^2/G$	exact	1
RVM dark energy form	Type II structural	DESI DR2 + Bertini fit	consistent	1
ν_{OI} magnitude	$\sim 10^{-32}$ (indistinguishable from 0)	$\nu^{\text{Bertini}} = -(2.5 \pm 1.3) \times 10^{-4}$	2σ	2
Dark-sector budget	$\sim 95\%$ invisible to QFT	$\Omega_{\text{dark}} \approx 0.95$	concordance	1
$a_0 = cH/6$	from entropy displacement	$1.2 \times 10^{-10} \text{ m/s}^2$	percent-level	2
$v^4 = GM_B \cdot cH/6$	baryonic Tully-Fisher	empirical relation	confirmed	2
$r_M(z)$ shrinkage	$\propto 1/\sqrt{H(z)}$	Genzel et al. $z = 0.9\text{--}2.4$	$> 3\sigma$	2
Milky Way Keplerian decline	$r_M \sim 17 \text{ kpc}$	Jiao et al. $\sim 19\text{--}26.5 \text{ kpc}$	confirmed	2
No dark-matter particle	structural	XENON, LUX, etc. null	consistent	2
No GUT-mediated proton decay	$\tau_p \sim 10^{45}\text{--}10^{46} \text{ yr}$	$\tau_p > 10^{34} \text{ yr}$	consistent	1

The tier system per Section 7.1: Tier 1 results follow from partition geometry alone and hold for any embedded-observer system satisfying C1-C3 at its cosmological horizon. Tier 2 results follow from the thermodynamic Clausius relation applied to the boundary reservoir and the geometric factor from spherical redis-

tribution in de Sitter space. Tier 3 results — the transition steepness in MOND interpolation, the greybody factor for Hawking radiation — depend on the specific bijection and are addressed by the cubic-lattice realization developed in Chapters 5-6.

Falsification commitments. Several pre-registered falsification conditions are concentrated at Class A per Chapter 4 §4.5. Any measured deviation from $S = A/(4l_p^2)$ at sufficient confidence to exclude experimental systematics would refute the framework’s gap equation chain — the consistency between classical horizon temperature and Euclidean KMS temperature that produced the $1/4$ coefficient. High-confidence detection of phantom crossing in the dark-energy equation of state — DESI DR2 currently shows $3\text{--}4\sigma$ mild tension that DESI Year 5 will resolve — would refute the framework’s structural prediction $w(z) \approx -1$. A high-confidence direct detection of a dark-matter particle at non-gravitational interaction strength would refute the framework’s structural commitment to no dark-matter particle. Discovery of GUT-mediated proton decay would refute the framework’s no-GUT structural commitment from Chapter 6.

The framework’s empirical exposure in the gravitational sector is therefore substantial. Seven Tier 1 predictions and seven Tier 2 predictions are pre-registered falsification targets, with the framework standing or falling on their continued empirical match. The chapter’s empirical anchor is the Milky Way Keplerian decline at the predicted $r_M \sim 17$ kpc; the GW250114 area-theorem result at 99.999% is a non-discriminating consistency.

The cumulative case for the geometry-first ordering. A pattern is worth marking explicitly. The framework’s commitment to geometry-first ordering — spacetime geometry is part of the classical substratum, quantum mechanics is the emergent description — produces a cumulative set of consequences that match observation. The Bekenstein-Hawking $1/4$ coefficient is derived. The cosmological constant problem is dissolved. The value of $\hbar$ is determined by partition geometry. The dark-sector budget at $\sim 95\%$ is concordance rather than puzzle — though this fraction is partly definitional, equal to $1 - \Omega_B$ by construction, so its real content is the matter-like versus dark-energy-like split rather than the total. Dark energy runs in Type II RVM form with coefficient indistinguishable from zero. The MOND acceleration scale $a_0 = cH/6$ matches Milgrom’s value. The baryonic Tully-Fisher relation follows from entropy displacement. Declining rotation curves at high redshift are predicted and confirmed. The Milky Way’s Keplerian decline at ~ 17 kpc is confirmed. The Standard Model gauge group and matter content are reproduced, with twenty-two classified quantitative retrodictions (seven of them parameter-free structural; Appendix A).

Quantum-first orderings — where QM is logically prior to the metric — produce the 10^{122} discrepancy as a fine-tuning problem, have no natural account of the dark sector, and treat $\hbar$ and the $1/4$ factor as unexplained inputs. The cumulative case is not a single tiebreaker but a pattern: every quantitative

output of the geometry-first ordering is consistent with observation; the central quantitative output of the quantum-first ordering — the cosmological constant prediction — is the worst prediction in physics. Chapter 4 §4.6 developed the framework’s preregistration discipline; the gravitational sector concentrates more confirmations of structural predictions made before the data than any other sector of the framework.

Forward pointers. Chapter 8 develops the JUNO neutrino prediction in detail — the post-November 2025 confirmation of $\sin^2 \theta_{12} = 1/3 - 1/(4\pi^2)$ at 0.07σ , the framework’s sharpest neutrino-sector empirical confirmation. The chapter develops the experimental context the matter-sector treatment in Chapter 6 abbreviated: the JUNO detector design, the reactor antineutrino source, the global-fit methodology, and the framework-specific significance of the $\sin^2 \theta_{12}$ measurement.

Chapter 9 develops the framework’s relationship to other unification programs — the OI-string relationship at the matrix-model level, the no-GUT structural commitment against the broader landscape of grand unified theories, the universality-class character of the framework’s Tier 1 results, and the de Sitter observer literature’s convergence with the framework’s geometry-first dissolution. The empirical record of Chapters 5-8 is concentrated at the Level G3 $\times$ Level D intersection of Chapter 3’s hierarchical structural realism; Chapter 9 locates the framework within the broader space of substratum-level theories and articulates the universality classes that the framework’s structural commitments belong to.

The framework’s content in fundamental physics, developed across Chapters 5 through 9, reduces a substantial portion of standard physics to structural consequences of the framework’s foundational commitments. The empirical record at the close of Part II spans the Standard Model gauge sector (22 predictions matched within 0.02% to 1.2%), the gravitational sector (14 predictions matched within 0.04% to 99.999%), and the neutrino sector (7 predictions matched within 0.07σ to 1.01σ). Part III turns to the framework’s reach beyond fundamental physics — the emergence cascade through chemistry, life, evolution, and intelligence — and Part IV to working applications in quantum biology, quantum engineering and computation, medicine, bioinformatics, and the framework’s forward predictions.

Chapter 8

JUNO and the Neutrino Sector

8.1 What this chapter develops

Chapter 6 §6.7 introduced the framework’s predictions for the three PMNS mixing angles and identified the solar mixing angle $\sin^2 \theta_{12} = 1/3 - 1/(4\pi^2) = 0.3080$ as the framework’s sharpest neutrino-sector empirical match — a parameter-free retrodiction whose provenance §8.6 states plainly, matched to the post-JUNO global fit at 0.07σ . The present chapter develops that prediction in detail. The JUNO experiment, the structural derivation, the global-fit methodology, the discrimination from competing tribimaximal-modification patterns, and the framework’s empirical position in the post-JUNO landscape — these are the chapter’s content.

The chapter opens with the experiment. Section 8.2 develops the JUNO detector design and the reactor antineutrino source: the 20-kiloton liquid scintillator detector 53 km from two nuclear power complexes in southern China, the energy resolution and reactor flux that combine to produce the world’s most precise measurement of $\sin^2 \theta_{12}$ on a much shorter timescale than was thought possible. Section 8.3 develops the cubic-group flavor structure that produces the framework’s prediction: the three lepton generations occupying the T_1 corners of the cubic Brillouin zone, the Cabibbo angle $\lambda = 1/(\pi\sqrt{2})$ from a single Brillouin-zone distance, the tribimaximal mixing pattern from the $A_4 \subset O$ subgroup at zeroth order.

The structural derivation occupies the middle of the chapter. Section 8.4 develops the U -parity grading — the discrete μ - τ exchange symmetry of $S_4 \supset A_4$ that controls how the charged-lepton perturbations modify the tribimaximal pattern at order λ^2 . Section 8.5 develops the sum rule $2\sin^2 \theta_{12} + \sin^2 \theta_{23} = 7/6$ as the framework’s structural discriminator from tribimaximal-modification alternatives, and locates the framework’s specific structural relation $b_{23} = (4/3)(b_{12} + b_{13})$ within the broader space of partial-TBM patterns currently analyzed against JUNO data.

The empirical content occupies the latter sections. Section 8.6 develops the three numerical predictions — $\sin^2 \theta_{12} = 1/3 - 1/(4\pi^2) = 0.3080$, $\sin^2 \theta_{13} = 4/(18\pi^2) = 0.0225$, $\sin^2 \theta_{23} = 1/2 + 1/(2\pi^2) = 0.5507$ — and compares them with the post-JUNO global fit and the NuFIT 6.0 best fits, with all three angles matched within 1.1σ . Section 8.7 compares the framework’s predictions with the column-preservation patterns TM1 and TM2 currently analyzed against JUNO: TM2 excluded at greater than 3.5σ post-JUNO, TM1 marginal at 1.35σ , the framework at 0.07σ .

The framework’s position in the literature is the chapter’s closing content. Section 8.8 develops the JUNO design-lifetime sensitivity: 30-year exposure projected to reach ± 0.0014 precision on $\sin^2 \theta_{12}$, providing a sub-percent test of the framework’s prediction. Section 8.9 develops the coupled neutrino-mass / dark-energy result with DESI DR2 — the $\Sigma m_\nu < 0.0642$ eV bound consistent with the framework’s normal-ordering prediction near the oscillation minimum of 0.059 eV, the 3σ tension with the lower oscillation limit assuming Λ CDM that

the framework’s RVM dark energy would relax, though the corrected bound is not yet computed. Section 8.10 closes with the chapter’s epistemic implications: a 0.07σ empirical match on a parameter-free structural prediction, with the chapter’s content classified per Chapter 4’s discipline.

The framework’s empirical record in the neutrino sector represents the strongest single piece of evidence for the framework’s structural commitments outside the gravitational sector. The Cabibbo angle and the Koide relation match observation at 0.04% and 0.001% respectively in the quark and charged-lepton sectors. The $\sin^2 \theta_{12}$ match at 0.07σ extends the same cubic-group machinery to the neutrino sector, with no fresh free parameters fitted — the same inputs $\lambda^2 = 1/(2\pi^2)$ and $A^2 = 2/3$ that produced the quark-sector matches produce the neutrino-sector match. The framework’s neutrino-sector content is therefore concentrated in this chapter as a focused case study of cross-sector parameter-free prediction at high precision.

8.2 The JUNO experiment

The Jiangmen Underground Neutrino Observatory (JUNO) is a multipurpose neutrino experiment located 700 meters underground in Jiangmen, southern China. The detector consists of a 20-kiloton liquid scintillator sphere, approximately 35 meters in diameter, instrumented with 17,612 large-aperture photomultiplier tubes and 25,600 smaller photomultiplier tubes. The detector is designed to measure reactor antineutrinos from the Yangjiang and Taishan nuclear power complexes, located 53 km from the experimental site — a baseline chosen to maximize sensitivity to the solar mixing angle θ_{12} at the oscillation maximum for reactor antineutrino energies.

The detector’s energy resolution is the principal experimental innovation. JUNO achieves approximately 3% at 1 MeV, a factor of three to five improvement over previous reactor neutrino experiments. The improvement comes from three sources: the high photocathode coverage (over 75% of the detector surface), the high light yield of the scintillator, and the careful calibration program that has been ongoing throughout the detector commissioning. The combination of high statistics (approximately 60 events per day from the two reactor complexes) and high energy resolution allows JUNO to resolve the fine structure of the reactor antineutrino oscillation spectrum at unprecedented precision.

The first JUNO measurement. The collaboration reported its first measurement of reactor neutrino oscillations in November 2025. The headline result from 59.1 days of physics data was

$$\sin^2 \theta_{12} = 0.3092 \pm 0.0087,$$

which is the most precise single-experiment determination of the solar mixing angle to date — a factor of 1.6 improvement over all previous measurements combined. The same dataset also produces a high-precision measurement of the

mass-squared splitting:

$$\Delta m_{21}^2 = (7.48 \pm 0.10) \times 10^{-5} \text{ eV}^2,$$

confirming the solar oscillation parameters with reactor-baseline precision for the first time. JUNO’s first measurement closes a long-standing tension between solar and reactor neutrino measurements of θ_{12} : the solar neutrino determination, primarily from SNO and Super-Kamiokande, gave systematically higher values, while the reactor neutrino determination from KamLAND gave systematically lower values, with the discrepancy at the 1.5σ level before JUNO. The new JUNO result lies between the two prior determinations and provides a unified value with significantly reduced uncertainty.

The post-JUNO global fit. The combined analysis of all neutrino oscillation data — solar, atmospheric, reactor, and accelerator — produces a global fit that incorporates the new JUNO measurement. The Capozzi *et al.* post-JUNO global fit (early 2026) gives

$$\sin^2 \theta_{12} = 0.3085 \pm 0.0073,$$

which is the value the framework’s predictions are tested against in this chapter. The global fit’s uncertainty is dominated by JUNO’s measurement; future runs will continue to tighten the constraint over the experiment’s design lifetime.

The JUNO experiment is therefore the primary current source of constraint on the solar mixing angle and on the framework’s specific prediction. The remainder of this chapter develops what that prediction is, where it comes from structurally, and how it compares with alternative theoretical proposals currently being analyzed against the JUNO data.

8.3 The cubic-group flavor structure

The framework’s prediction for the PMNS angles follows from the cubic point group O — the rotation symmetry group of three-dimensional space — applied to the three lepton generations. The structural argument runs through the same chain that produced the gauge group $SU(3) \times SU(2) \times U(1)$ in Chapter 5: the cubic decomposition of the six nearest-neighbor link directions on the simple cubic lattice gives multiplicities $(3, 2, 1)$ under the cubic-group action, with the three lepton generations occupying the T_1 triplet representation.

The cubic group O and its A_4 subgroup. The cubic point group O is the rotation symmetry group of three-dimensional space, with 24 elements isomorphic to the symmetric group S_4 . Its irreducible representations decompose as

$$\text{irreps}(O) = A_1 \oplus A_2 \oplus E \oplus T_1 \oplus T_2,$$

of dimensions $1, 1, 2, 3, 3$. The subgroup $A_4 \subset O$ keeps $\{A_1, A_2, E, T_1\}$ but identifies A_2 with A_1 , giving the standard A_4 content $A \oplus A' \oplus A'' \oplus T$ of dimensions $1, 1, 1, 3$. The framework’s structural derivation in Chapter 5 placed

the three lepton generations on the T_1 representation as one of the irreducible components of the cubic decomposition $6 = T_1 \oplus E \oplus A_1$ of the six nearest-neighbor link directions.

The realization of T_1 in the framework is concrete. The three Cartesian unit vectors $\{e_1, e_2, e_3\}$ — corresponding to the three spatial axes of the cubic lattice — are acted on by the 24 rotations of O . The three generations occupy these three Cartesian directions. The cubic decomposition has a single democratic direction

$$\hat{h} = \frac{1}{\sqrt{3}}(e_1 + e_2 + e_3),$$

which is the only direction fixed by the $\mathbb{Z}_3$ cyclic subgroup of O . The plane perpendicular to $\hat{h}$ carries the doublet E representation. This is the geometric origin of tribimaximal mixing in the framework: the second column of the tribimaximal mixing matrix is exactly $\hat{h}$, and the first and third columns lie in the plane perpendicular to it, oriented by μ - τ exchange.

The Cabibbo scale from cubic Brillouin geometry. The perturbation parameter λ controlling the deviations from tribimaximal mixing in the framework is identified with the Cabibbo angle, fixed independently by the cubic Brillouin-zone geometry as developed in Chapter 6 §6.5. The three generations occupy the three T_1 corners of the cubic Brillouin zone at momenta $X_j = \pi \cdot e_j$ for $j = 1, 2, 3$. The inter-generation distance in momentum space is

$$|X_i - X_j| = \pi\sqrt{2} \quad (i \neq j),$$

a geometric constant of the cubic lattice. The Cabibbo angle is the inverse of this distance:

$$\lambda = \frac{1}{\pi\sqrt{2}} = 0.22508.$$

The match to observation is at 0.04%: the measured value is $\lambda = 0.22500 \pm 0.00067$, with the framework's prediction inside the experimental uncertainty. The Cabibbo angle is therefore a structural prediction of the framework, fixed by a single Brillouin-zone distance with no input from the lepton sector; the geometry and the chirality-forced power are unconditional, while the identification of the mixing element with the unit-normalized propagator is a stated hypothesis whose resolving computation is specified in [SM §7.1]; the hypothesis has since been structurally narrowed — the free-level mixing object vanishes exactly and the substrate-kinematics alternative is excluded, so the identification is equivalent to the emergent-theory coefficient being unity — and the on-shell computation has since answered it: between physically projected legs, every spin-resolved taste-changing element has unit modulus exactly, confirming the unity condition at its structural level as an algebraic identity of the spin-taste reconstruction ([SM §7.1]); what remains beyond that level is the dressing. The fact that the same λ then controls the PMNS deviations from tribimaximal mixing is the framework's structural commitment: the Cabibbo scale that produces the quark-sector mixing is the same scale that produces the lepton-sector

mixing deviations. The two empirical confirmations — the Cabibbo angle at 0.04% and the solar mixing angle at 0.07σ — are not independent confirmations of separate structures but joint confirmations of a single cubic-group geometric source.

The PMNS matrix at zeroth order. The framework’s PMNS structure at zeroth order is the tribimaximal mixing matrix, obtained from the $A_4 \subset O$ representation theory. The tribimaximal matrix is

$$U_{\text{TBM}} = \begin{pmatrix} \sqrt{2/3} & \sqrt{1/3} & 0 \\ -\sqrt{1/6} & \sqrt{1/3} & \sqrt{1/2} \\ -\sqrt{1/6} & \sqrt{1/3} & -\sqrt{1/2} \end{pmatrix},$$

with the columns corresponding to mass eigenstates ν_1, ν_2, ν_3 and the rows to flavor eigenstates ν_e, ν_μ, ν_τ . The corresponding zeroth-order mixing angles are

$$\sin^2 \theta_{12}^{\text{TBM}} = \frac{1}{3}, \quad \sin^2 \theta_{23}^{\text{TBM}} = \frac{1}{2}, \quad \sin^2 \theta_{13}^{\text{TBM}} = 0.$$

The observed values deviate from tribimaximal mixing by amounts of order $\lambda^2 \approx 0.05$, with the most consequential deviation being a non-zero θ_{13} at the level of $\sin^2 \theta_{13} \approx 0.022$. The framework’s content in the next sections is to derive these deviations from a μ - τ parity grading combined with one cubic-group structural relation, producing parameter-free predictions for all three angles.

The framework’s structural commitment. The cubic-group flavor structure used here is not a postulated symmetry of the lepton sector. It is the rotation symmetry group of three-dimensional space, realized in the framework’s lattice substratum as developed in Chapters 5 and 6. The same cubic group that produced the Standard Model gauge structure produces the lepton flavor structure here, with the T_1 generations occupying the same triplet representation in both contexts. The empirical record across the framework’s flavor predictions — Cabibbo at 0.04%, Koide at 0.02%, $\sin^2 \theta_{12}$ at 0.07σ , the sum rule $2\sin^2 \theta_{12} + \sin^2 \theta_{23} = 7/6$ at 0.6σ — tests the same cubic structure across multiple observables at sub-percent precision. This is testable structure rather than a postulated symmetry: the same numerical inputs ($\lambda^2 = 1/(2\pi^2)$, the projection factor $A^2 = 2/3$) enter multiple observables with all of them matching observation.

8.4 The U-parity grading

The framework’s PMNS structure beyond tribimaximal mixing develops through a μ - τ exchange parity — a discrete $\mathbb{Z}_2$ grading from $S_4 \supset A_4$ — that organizes the charged-lepton perturbations into structurally distinct U -even and U -odd components. The grading is not an additional symmetry postulate but a direct consequence of the framework’s structural commitment to T_1 as the generation

triplet: T_1 contains a natural μ - τ exchange generator, and the perturbations of T_1 inherit a grading under this exchange.

The U -parity generator. In the Altarelli-Feruglio basis (charged-lepton mass matrix diagonal at zeroth order), the μ - τ exchange operator U is the permutation that exchanges the second and third generations:

$$U = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix}.$$

The generator U sits in $S_4 \supset A_4$ as the μ - τ exchange element of the larger permutation group. Conjugation by U defines a $\mathbb{Z}_2$ grading on antisymmetric matrices: a matrix A is U -even if $UAU^{-1} = +A$, U -odd if $UAU^{-1} = -A$. The grading decomposes the space of antisymmetric matrices into two orthogonal components.

With the basis matrices $(E_{ij})_{kl} = \delta_{ik}\delta_{jl} - \delta_{il}\delta_{jk}$ — the elementary antisymmetric matrices spanning the space of 3×3 antisymmetric matrices — the decomposition is: - U -even combinations: $E_{12} + E_{13}$ - U -odd combinations: $E_{12} - E_{13}$ and E_{23}

The unique U -odd combination antisymmetric in (μ, τ) is E_{23} ; the unique U -odd combination relating the first generation to the second and third is $E_{12} - E_{13}$. The U -even combination $E_{12} + E_{13}$ relates the first generation symmetrically to the other two.

The charged-lepton perturbation expansion. The charged-lepton rotation matrix V_ℓ in the Altarelli-Feruglio basis admits a perturbative expansion in the Cabibbo scale λ :

$$V_\ell = \exp(\lambda A_1 + \lambda^2 A_2 + O(\lambda^3)),$$

with A_1 and A_2 real antisymmetric 3×3 matrices to be determined from structural inputs. The expansion is the framework's parametrization of the charged-lepton rotation: A_1 controls the leading deviations from the tribimaximal pattern at order λ , A_2 controls the next-to-leading deviations at order λ^2 .

Determination of A_1 . The leading-order generator A_1 is fixed by three structural conditions plus the cubic geometric input. First, the observed value $\theta_{13} \neq 0$ requires $\sin \theta_{13} = O(\lambda)$, forcing A_1 to contribute to the $(1, 3)$ element of $V_\ell^\dagger U_{\text{TBM}}$. Second, the empirical observations that both $\Delta_{12} \equiv \sin^2 \theta_{12} - 1/3$ and $\Delta_{23} \equiv \sin^2 \theta_{23} - 1/2$ are of order λ^2 — neither shows a leading λ contribution — further constrains A_1 's structure. Third, the geometric input from Section 8.3: each end of an inter-generation mixing vertex projects onto the plane perpendicular to the democratic direction $\hat{h}$, contributing a factor $\sin(\angle(\hat{e}_i, \hat{h})) = A$ from the off-democratic projection of any generation axis. For a single-vertex inter-generation transition, the total projection factor is A^2 — one factor per vertex end.

Identifying this projection geometry with the leading $O(\lambda)$ contribution to U_{e3} gives the structural identity

$$\sin \theta_{13} = A^2 \lambda,$$

exact at order λ (the order- λ^2 contributions from A_2 cancel in $|U_{e3}|^2$, as the next section establishes). This is a structural input fixing A_1 's overall coefficient up to sign, combined with the U -odd character required by the constraint $\Delta_{12}, \Delta_{23} = O(\lambda^2)$.

The result is uniquely determined:

$$A_1 = \frac{\sqrt{2}}{3}(E_{12} - E_{13}).$$

The generator A_1 is purely U -odd: $UA_1U^{-1} = -A_1$. Geometrically, A_1 represents a rotation by angle $A^2 = 2/3$ around the axis $(\hat{e}_2 + \hat{e}_3)/\sqrt{2}$, which reproduces $\sin \theta_{13} = A^2 \lambda$ as required. The structural property is that A_1 is U -odd; the magnitude $\sqrt{2}/3$ is fixed by the projection identity $\sin \theta_{13} = A^2 \lambda$ with $A^2 = 2/3$.

Deviations from TBM at order λ^2 . Expanding V_ℓ to second order:

$$V_\ell = I + \lambda A_1 + \lambda^2(A_2 + \tfrac{1}{2}A_1^2) + O(\lambda^3),$$

and writing the most general antisymmetric A_2 as

$$A_2 = b_{12}E_{12} + b_{13}E_{13} + b_{23}E_{23},$$

direct computation of $|U_{\text{PMNS},ij}|^2$ to order λ^2 gives the deviations from tribimaximal mixing:

$$\begin{aligned}\Delta_{13} &= \tfrac{4}{9}\lambda^2 = A^4\lambda^2, \\ \Delta_{12} &= -\tfrac{2}{3}(b_{12} + b_{13})\lambda^2, \\ \Delta_{23} &= b_{23}\lambda^2.\end{aligned}$$

Three structural properties of these deviations deserve explicit comment. First, Δ_{13} is independent of A_2 — the structural identity $\sin \theta_{13} = A^2 \lambda$ is exact at order λ^2 , with the A_2 contributions canceling from $|U_{e3}|^2$. The reactor mixing angle prediction is therefore Layer 1 structural: $\sin^2 \theta_{13} = (4/9)\lambda^2$ follows from A_1 alone, with no second-order Wilson-coefficient inputs. Second, the combination $b_{12} - b_{13}$ does not enter any of the three angles — it is naturally identified with the CP-violating Dirac phase δ_{CP} , the only remaining free parameter at order λ^2 that does not enter the angles. Third, Δ_{23} depends only on b_{23} , not on the U -even combination $b_{12} + b_{13}$. The U -parity structure organizes the second-order corrections in a way that the standard effective-field-theory parametrization does not.

8.5 The sum rule and structural relation

The deviations Δ_{12} and Δ_{23} depend on two independent scalar combinations of the A_2 Wilson coefficients: the U -even combination $b_{12}+b_{13}$ controlling Δ_{12} , and the dominant U -odd coefficient b_{23} controlling Δ_{23} . The framework's content beyond the Layer 1 reactor angle prediction is a structural relation between these two combinations — a relation that produces the sum rule $2\sin^2\theta_{12} + \sin^2\theta_{23} = 7/6$ and discriminates the framework from tribimaximal-modification alternatives.

Two normalization conditions. The framework fixes the second-order Wilson coefficients through two conditions.

Condition 1 — natural normalization. The overall scale of A_2 is fixed by taking $b_{23} = 1$. This is a labeling convention rather than a physical input: it sets A_2 to enter the perturbation series at unit strength relative to λ^2 , fixing the normalization of the order- λ^2 expansion before $b_{12} + b_{13}$ is determined.

Condition 2 — cubic-group structural relation. The relation between the dominant U -odd coefficient and the U -even combination is fixed by the cubic-group representation theory:

$$b_{23} = \frac{4}{3}(b_{12} + b_{13}).$$

With Condition 1, this gives $b_{12} + b_{13} = 3/4$. Unlike Condition 1, this is a substantive physical input — a relation between two independent Wilson coefficients that the cubic-group structure produces.

Equivalence with TBM sum rule preservation. The framework's structural relation has a clean interpretation at the angle level. Combining the explicit results $\Delta_{12} = -(2/3)(b_{12} + b_{13})\lambda^2$ and $\Delta_{23} = b_{23}\lambda^2$:

$$2\Delta_{12} + \Delta_{23} = (b_{23} - \frac{4}{3}(b_{12} + b_{13}))\lambda^2.$$

Condition 2 is therefore equivalent to $2\Delta_{12} + \Delta_{23} = 0$, which is the statement that the tribimaximal sum rule $2\sin^2\theta_{12} + \sin^2\theta_{23} = 7/6$ — holding at zeroth order from cubic geometry (since $2(1/3) + 1/2 = 7/6$) — continues to hold at order λ^2 . The Wilson-coefficient relation and the angle-level sum rule are the same content under different parametrizations.

The sum rule itself has a clean cubic-geometric reading: $7/6 = 2/d + 1/(d-1)$ for $d = 3$ generations, with $2/d$ being twice the inverse generation count (the value $\sin^2\theta_{12}^{\text{TBM}} = 1/3$ doubled) and $1/(d-1)$ being the inverse partner count (the value $\sin^2\theta_{23}^{\text{TBM}} = 1/2$ in the doublet sector). For three generations, $2/3 + 1/2 = 7/6$ is forced by the representation theory; the framework's content is that this sum rule is *preserved* at order λ^2 by the second-order Wilson-coefficient relation.

The factor 4/3 as cubic-group identity. Condition 2's specific numerical coefficient is a genuine cubic-group identity:

$$\frac{4}{3} = 2A^2 = 2 \cdot \frac{C_2(T_1)}{d} = \frac{C_2(T_1)^2}{d},$$

where $A = \sqrt{2/3}$ is the off-democratic projection of any generation axis and $C_2(T_1) = d - 1 = 2$ is the quadratic Casimir of the T_1 representation. The identity $A^2 = (d - 1)/d = C_2(T_1)/d$ for $d = 3$ is the same representation-theoretic identity from Section 6.5 that produced the Wolfenstein A parameter. The factor $4/3$ thus has an interpretive reading as partner-count $\times$ per-partner projection: the U -even combination $E_{12} + E_{13}$ couples the first generation to its two partners $\{\hat{e}_2, \hat{e}_3\}$, with the per-partner projection nominally A^2 . The interpretation is internally consistent but non-directional — the same arithmetic produces the inverse ratio $3/4$ under alternative partner-count assignments. The directional assertion that the U -odd coefficient b_{23} is *larger* than the U -even coefficient $b_{12} + b_{13}$ by this specific factor is what carries the framework's substantive content beyond the cubic-group identity itself.

Structural protection of the sum rule by A_1 alone. A distinct structural feature of the framework's A_1 choice deserves explicit treatment. With $A_2 = 0$, direct computation gives $2|U_{e2}|^2 + |U_{\mu3}|^2 = 7/6 + \lambda^2 \cdot [-14/27]$ before the $\cos^2 \theta_{13}$ correction. Dividing by $\cos^2 \theta_{13} = 1 - (4/9)\lambda^2$ — which itself follows from $\sin^2 \theta_{13} = A^4 \lambda^2$, a Layer 1 quantity — precisely cancels the $-14/27$ offset, giving

$$(2 \sin^2 \theta_{12} + \sin^2 \theta_{23}) \big|_{A_2=0} = 7/6 + O(\lambda^4).$$

The sum rule is therefore *structurally protected* against A_1 's perturbation by cubic geometry alone — no substrate input required for this component. The cancellation between the $-14/27$ contribution from A_1^2 -induced shifts and the $+14/27$ contribution from the $\cos^2 \theta_{13}$ dressing is automatic given A_1 's Layer 1-forced form. Condition 2's content reduces to whether A_2 *preserves* the sum rule that A_1 alone protected, or breaks it. The framework's structural commitment is that A_2 also preserves the sum rule, through the $b_{23} = (4/3)(b_{12} + b_{13})$ relation.

Layer status. The three angle predictions have distinct structural statuses per the four-layer framing of Chapter 4 §4.5. The reactor angle $\sin^2 \theta_{13} = (4/9)\lambda^2 = A^4 \lambda^2$ is Layer 1 unconditional structural — it depends only on A_1 at second order, with A_2 contributions canceling. The solar and atmospheric angles $\sin^2 \theta_{12}$ and $\sin^2 \theta_{23}$ are layered predictions: the structural form is forced by the A_1 component (with the sum rule preserved by cubic geometry alone), and the specific deviations at order λ^2 depend on Condition 2's relation $b_{23} = (4/3)(b_{12} + b_{13})$.

Condition 2 as a Layer 2(b) bijection-specific commitment. The framework's $b_{23} = (4/3)(b_{12} + b_{13})$ relation has a multi-layer structural status that the four-layer framing of Chapter 4 §4.5 distinguishes precisely:

- (i) Under the U -parity grading developed in §8.4, the antisymmetric tensor product $T_1 \wedge T_1$ decomposes into a U -even subspace (spanned by $E_{12} + E_{13}$) and a U -odd subspace (spanned by $E_{12} - E_{13}$ and E_{23}). The subspace decomposition is forced by group theory — this is Layer 1 structural.

- (ii) The cubic-geometric off-democratic projection $A^2 = 1 - 1/d = 2/3$ for $d = 3$ is forced by inner-product structure on T_1 with $\hat{h} = (1, 1, 1)/\sqrt{3}$ as the democratic axis. This is the same A^2 that produces the Wolfenstein parameter $A = \sqrt{2/3}$ in the quark sector (Chapter 6 §6.5) — Layer 1 structural.
- (iii) The cubic-group identity $4/3 = 2A^2 = C_2(T_1)^2/d$ is a genuine representation-theoretic content of the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$, gauge-covariant up to state-relabeling under the substratum gauge group. The structural statement “*for some specific generation labeling, one of the three pair-ratios equals 4/3*” is Layer 2(a) forced by the cubic-group structure.
- (iv) The directional assertion that the U -odd amplitude exceeds the U -even amplitude by the specific factor of $4/3$, with the labeling fixed by the observed mass hierarchy ($m_e \ll m_\mu \ll m_\tau$ pinning generation 1 = electron), is *Layer 2(b) bijection-specific* content. It depends on the specific bijection φ within the equivalence class — the same kind of content as the fermion mass hierarchy itself (Chapter 19 §19.3.1) and CKM/PMNS CP phases.

The empirical match at 0.07σ on $\sin^2 \theta_{12}$ provides genuine evidence about the specific bijection φ . This is the same form of empirical evidence the framework relies on for other Layer 2(b) content (mass hierarchy match through the Koide structure, the existence of three generations with hierarchical masses).

The open derivational task is correctly scoped: it asks whether the specific staggered-Yukawa structure on the cubic lattice (in the emergent post-trace-out description; the substratum proper has no Lagrangian per Ch02 §2.5) produces A_2 in the sum-rule-preserving subspace for a given bijection. This is a tractable lattice-perturbation-theory question for the *magnitude* content (Layer 2(a)), but the directionality remains Layer 2(b) by construction — fixed by the empirical mass-ordering bijection, not derivable from substrate alone in any framework consistent with $\mathcal{G}_{\text{sub}}$ -invariance.

8.6 The three numerical predictions

Substituting Condition 1 ($b_{23} = 1$), Condition 2 ($b_{12} + b_{13} = 3/4$), and the Cabibbo scale $\lambda^2 = 1/(2\pi^2)$ into the order- λ^2 deviations from tribimaximal mixing gives the framework’s three numerical predictions for the PMNS angles.

The three predictions.

$$\begin{aligned}\sin^2 \theta_{12} &= \frac{1}{3} - \frac{1}{4\pi^2} = 0.3080, \\ \sin^2 \theta_{13} &= \frac{4}{9} \cdot \frac{1}{2\pi^2} = \frac{4}{18\pi^2} = 0.02252, \\ \sin^2 \theta_{23} &= \frac{1}{2} + \frac{1}{2\pi^2} = 0.5507.\end{aligned}$$

The reactor angle $\sin^2 \theta_{13}$ uses only $A^2 = 2/3$ and $\lambda^2 = 1/(2\pi^2)$ together with the constraints on A_1 — it does not use the normalization conditions on A_2 . The other two angles additionally use Conditions 1 and 2. The three predictions together exhaust the framework’s content for the PMNS mixing angles; the framework leaves δ_{CP} (the CP-violating Dirac phase) undetermined.

Comparison with the post-JUNO global fit and NuFIT 6.0. The empirical state of the PMNS angles as of early 2026 is summarized in the table below. The post-JUNO global fit incorporates the November 2025 JUNO measurement; the NuFIT 6.0 best fits use the full neutrino-oscillation data set under the normal mass ordering convention.

Angle	Predicted	Observed	Match
$\sin^2 \theta_{12}$	0.3080	0.3085 ± 0.0073 (post-JUNO global fit)	0.07σ
$\sin^2 \theta_{23}$	0.5507	$0.561^{+0.012}_{-0.015}$ (NuFIT 6.0, NO)	0.74σ
$\sin^2 \theta_{13}$	0.02252	0.02195 ± 0.00056 (NuFIT 6.0, NO)	1.01σ

All three predictions match observation within 1.1σ at current precision. The solar mixing angle value is the framework’s sharpest neutrino-sector empirical match — a parameter-free retrodiction at 0.07σ from the post-JUNO global fit; the derivation (2026) postdates the first results (November 2025), and JUNO’s design-lifetime precision (§8.8) is the corresponding pre-registered forward test, well within experimental uncertainty.

The solar mixing angle as the cleanest test. The $\sin^2 \theta_{12}$ prediction is the framework’s cleanest empirical test in the lepton sector for three reasons. First, it has the smallest experimental uncertainty among the three angles (relative uncertainty about 2.4%, compared to about 2.5% for $\sin^2 \theta_{23}$ and about 2.6% for $\sin^2 \theta_{13}$). Second, the prediction is parameter-free: both inputs entering the framework’s prediction — the zeroth-order value $1/3$ from the A_4 tribimaximal pattern and the perturbation scale $\lambda^2 = 1/(2\pi^2)$ from the Cabibbo Brillouin-zone geometry — are fixed by *quark-sector* observables before any neutrino-sector data enter the derivation. The same cubic-group machinery that produces the Cabibbo angle $\lambda = 1/(\pi\sqrt{2})$ at 0.04% precision and the Wolfenstein parameter $A = \sqrt{2/3}$ within 1σ then produces the solar mixing angle $\sin^2 \theta_{12} = 1/3 - 1/(4\pi^2) = 0.3080$ at 0.07σ — a cross-sector empirical match with no fresh free parameters fitted to the neutrino sector. Third, the prediction is distinctive in the space of tribimaximal-modification proposals: the post-JUNO data exclude TM2 at more than 3.5σ and place TM1 at 1.35σ , with the framework alone at 0.07σ among the well-developed alternatives.

The cumulative empirical record across the framework’s flavor predictions is striking. The Cabibbo angle is matched at 0.04%, the Koide angle at 0.02%, the solar mixing angle at 0.07σ , the sum rule $2\sin^2\theta_{12} + \sin^2\theta_{23} = 7/6$ at 0.6σ — all from the same cubic-group representation theory applied to the lepton sector, with the same numerical inputs ($\lambda^2 = 1/(2\pi^2)$, $A^2 = 2/3$) producing all four predictions. The pattern is what distinguishes the framework’s cubic-group structure from postulated flavor symmetries: the same parameter enters multiple observables at sub-percent precision, testable across the full flavor sector rather than only at one observable.

The sum rule. The framework’s prediction $2\sin^2\theta_{12} + \sin^2\theta_{23} = 7/6$ is the cubic structural commitment beyond the three angles individually. Substituting the framework’s predicted values: $2 \cdot 0.3080 + 0.5507 = 1.167$, which equals $7/6$ exactly. Substituting the observed values: $2 \cdot 0.3085 + 0.561 = 1.178 \pm 0.020$, with the agreement at 0.6σ — consistent with the framework’s structural prediction. The sum rule is therefore a *fourth* empirical confirmation of the same cubic-group structure, complementary to the three individual angle measurements.

The remaining sections develop the framework’s position in the broader landscape of tribimaximal-modification proposals (Section 8.7), the JUNO design-lifetime sensitivity for testing the $\sin^2\theta_{12}$ prediction at sub-percent precision (Section 8.8), and the coupled neutrino-mass / dark-energy result with DESI DR2 (Section 8.9).

8.7 Comparison with TM1/TM2 column-preservation patterns

Tribimaximal mixing was the leading phenomenological description of neutrino oscillation data until Daya Bay measured a non-zero θ_{13} in 2012. Since then, modified tribimaximal patterns have been the subject of a substantial literature, with the two best-known modifications preserving one column of the original tribimaximal matrix exactly: TM1 (first column preserved) and TM2 (second column preserved). Each gives a one-parameter family that accommodates the observed θ_{13} but predicts different values for the other two angles.

TM1 and TM2 in the post-JUNO landscape. Post-JUNO analyses by He and by Zhang find that TM2 is excluded at greater than 3.5σ while TM1 remains marginally viable. The framework’s prediction sits well below both alternatives in tension with the post-JUNO global fit:

Approach	Symmetry	$\sin^2\theta_{12}$	σ from observed
	origin		
TM1	postulated A_4 flavor	0.3184	1.35σ

Approach	Symmetry origin	$\sin^2 \theta_{12}$	σ from observed
TM2	postulated A_4 flavor	0.3408	4.4σ
TM3	postulated; $\theta_{13} = 0$	—	excluded
This work	cubic group O of 3D space	0.3080	0.07σ

(Comparing against the post-JUNO global fit $\sin^2 \theta_{12} = 0.3085 \pm 0.0073$. TM1 uses $\sin^2 \theta_{12} = 1 - (2/3)/\cos^2 \theta_{13}$; TM2 uses $\sin^2 \theta_{12} = 1/(3\cos^2 \theta_{13})$; both with $\sin^2 \theta_{13} = 0.02195$.) The cubic-group prediction is closer to the post-JUNO global fit by a factor of about 20 in σ -distance than the best surviving column-preservation pattern, and at the same time makes more falsifiable claims (three independent angles plus the sum rule).

Three structural differences distinguish the framework from TM_n .

Symmetry origin. TM1, TM2, and TM3 take a discrete flavor symmetry as an *axiom* of the lepton sector — typically A_4 , S_4 , or a $\mathbb{Z}_2$ residual postulated for the lepton sector specifically. The framework’s prediction uses $A_4 \subset O$, where O is the rotation symmetry group of three-dimensional space. The framework’s cubic structure is not postulated for the lepton sector; it is forced by the framework’s structural commitments developed in Chapters 5–6: the Bravais-lattice uniqueness theorem (Chapter 5 §5.5), the coupling-degree minimization producing $K = 2d = 6$ (Chapter 5 §5.4), and the anomaly cancellation chain (Chapter 5 §5.7). The same cubic structure that produced the Standard Model gauge group with three generations and the Cabibbo angle here produces the PMNS mixing structure. The flavor symmetry is derived, not postulated.

Independence of predictions. TM1 produces a one-parameter correlation $\sin^2 \theta_{12} = 1 - (2/3)/\cos^2 \theta_{13}$, fixing θ_{12} given θ_{13} but predicting neither separately. The framework’s analysis predicts all three angles as fixed numerical values, with the same two structural inputs ($\lambda^2 = 1/(2\pi^2)$ and $A^2 = 2/3$) appearing across all three predictions. This makes the framework’s prediction strictly more falsifiable than TM1: any individual angle outside the framework’s prediction refutes it independently, whereas TM1 can absorb deviations in one angle by adjusting the prediction for another.

CP violation. He’s surviving TM1 pattern predicts $\sin \delta_{\text{CP}} = \pm 0.998$, close to maximal CP violation. The framework identifies δ_{CP} with the only remaining order- λ^2 free parameter ($b_{12} - b_{13}$) that does not enter any of the three angles. The framework’s angle predictions therefore stand independently of any value of δ_{CP} ; this chapter makes no prediction for the CP phase. The two frameworks are testably distinct in the CP sector: if forthcoming long-baseline neutrino experiments (DUNE, T2HK) measure $\sin \delta_{\text{CP}}$ inconsistent with maximal CP

violation, TM1 is refuted while the framework’s content is unaffected.

The sum rule discriminator. The framework’s structural prediction $2\sin^2\theta_{12} + \sin^2\theta_{23} = 7/6$ is the angle-level statement of Condition 2’s relation $b_{23} = (4/3)(b_{12} + b_{13})$ from Section 8.5. The observed value is $2(0.3085) + 0.561 = 1.178 \pm 0.020$, in agreement with $7/6 = 1.1667$ at 0.6σ . This depends on the atmospheric octant: the quoted value uses the upper-octant best fit $\sin^2\theta_{23} \approx 0.561$; under the lower-octant solution (≈ 0.455) the combination is ≈ 1.07 , roughly 5σ below $7/6$. TM1 and TM2 do not produce this combination as an exact sum rule — in TM2 specifically, $\sin^2\theta_{12}$ is determined by θ_{13} alone with no correlation to θ_{23} , so no analogous sum rule is forced. A direct measurement of $2\sin^2\theta_{12} + \sin^2\theta_{23}$ at the 0.005 level — achievable with the JUNO design lifetime combined with NOvA, T2K, and IceCube-Upgrade atmospheric determinations of θ_{23} — is the sharpest data-side test distinguishing the framework from column-preservation alternatives.

The cumulative-match pattern. The framework’s $\sin^2\theta_{12}$ match at 0.07σ does not stand alone. It sits within a cluster of empirical matches that all use the same cubic-group structural inputs: the Cabibbo angle at 0.04%, the Koide angle at 0.02%, the $\sin^2\theta_{13}$ prediction at 1.01σ , the sum rule at 0.6σ , the $|V_{cb}|$ prediction at 0.4σ , the m_u/m_d ratio at 0.05σ . All six predictions match observation at sub-percent or sub- σ precision simultaneously, with two structural parameters ($\lambda^2 = 1/(2\pi^2)$ and $A^2 = 2/3$) appearing repeatedly across the cluster.

Under any reasonable prior over flavor models, this is non-trivial structural agreement. A single set of structural inputs produces six distinct empirical matches; each would have to be explained as a coincidence in a flavor-symmetry framework that postulates the cubic structure rather than deriving it. The framework’s content here is that the same cubic structure that forces the SM gauge group on the substratum side appears as the empirical flavor structure on the phenomenological side — testable across the full flavor sector rather than only at one observable.

8.8 JUNO design-lifetime sensitivity and falsification

JUNO’s design lifetime is 30 years of physics operations. The collaboration projects that the experiment will achieve approximately ± 0.0014 precision on $\sin^2\theta_{12}$ by the end of its design lifetime — a factor of six improvement over the first-result precision of ± 0.0087 and approximately a factor of five improvement over the post-JUNO global fit precision of ± 0.0073 . The improvement comes from extended exposure (factor of about 185 increase in integrated luminosity over the first 59.1 days), continued calibration refinement, and accumulated cross-checks across multiple reactor cores.

The framework’s prediction at JUNO design-lifetime precision. The framework’s structural prediction $\sin^2\theta_{12} = 1/3 - 1/(4\pi^2) = 0.30798$ is a single

numerical value with no uncertainty from the framework side. The experimental side will reach ± 0.0014 at design lifetime, giving a precision test at the 0.5% relative level on the predicted value. If the framework's prediction is correct, JUNO's design-lifetime measurement should center on 0.3080 with ± 0.0014 uncertainty.

The discrimination this provides is severe. The TM1 pattern's prediction $\sin^2 \theta_{12} \approx 0.3184$ differs from the framework's by 0.0104, which is 7.4σ at JUNO's design-lifetime precision. TM2's 0.3408 differs by 0.0328, which is 23σ . Any column-preservation pattern with $\sin^2 \theta_{12} \gtrsim 0.31$ will be excluded at JUNO design-lifetime precision; only the cubic-group prediction will remain viable in that range. The 30-year program is therefore a decisive test between the framework's structural prediction and the broader landscape of tribimaximal-modification alternatives.

Sensitivity to Condition 2. The framework's prediction depends on Condition 2 — the structural relation $b_{23} = (4/3)(b_{12} + b_{13})$ developed in Section 8.5. A small fractional shift to the structural coefficient $4/3 \rightarrow (4/3)(1 + \epsilon)$ propagates to a shift in $\sin^2 \theta_{12}$ of approximately $\epsilon \cdot \Delta_{23}/2$ at first order, where $\Delta_{23} = 1/(2\pi^2) \approx 0.05$. For a fractional shift $\epsilon = 0.10$ — a 10% deviation from the cubic-group structural coefficient — the shift in $\sin^2 \theta_{12}$ is approximately 0.0025. This is comparable to JUNO's design-lifetime precision and roughly half of the first-result precision.

The framework's current 0.07σ match therefore tests Condition 2 at the few-percent level on its structural coefficient. JUNO's design-lifetime program will tighten this to better than 1%, providing an empirical test of the structural relation $4/3 = 2A^2$ — the framework's commitment that the cubic-group quadratic Casimir structure produces this specific ratio in the substrate. The empirical confirmation at sub-percent precision over JUNO's 30-year lifetime would constitute a quantitative test of the framework's cubic-group commitment at the level normally reserved for Standard Model coupling constants.

Falsification commitments. Three pre-registered falsification conditions per Chapter 4 §4.5 are concentrated at the JUNO experimental program.

Class A falsification of the cubic-group lepton-sector structure. Any precision-upgrade measurement of $\sin^2 \theta_{12}$ moving the observed value outside 3σ of the framework's prediction $1/3 - 1/(4\pi^2)$ refutes the framework's cubic-group commitment in the lepton sector. The framework would have to be revised — either the $A_4 \subset O$ tribimaximal structure at zeroth order, the Cabibbo perturbation scale $\lambda^2 = 1/(2\pi^2)$, or the Condition 2 structural relation — to accommodate the deviation. The current 0.07σ match makes such revision unlikely, but JUNO's 30-year program will resolve the question definitively.

Class A falsification of the sum rule. Any precision measurement of $2\sin^2 \theta_{12} + \sin^2 \theta_{23}$ moving outside 3σ of the framework's prediction $7/6$ refutes the Condition 2 structural relation. The sum rule is the framework's sharpest discrimina-

tor against tribimaximal-modification alternatives, and a precise measurement at the 0.005 level would either confirm or refute the framework’s structural commitment in the lepton sector independently of the absolute value of $\sin^2 \theta_{12}$.

Class A falsification of the reactor angle. Any precision measurement of $\sin^2 \theta_{13}$ moving outside 3σ of the framework’s prediction $4/(18\pi^2) = 0.02252$ refutes the Layer 1 reactor angle derivation. This prediction depends only on the structural identity $\sin \theta_{13} = A^2 \lambda$ with $A^2 = C_2(T_1)/d = 2/3$ and the Cabibbo scale $\lambda = 1/(\pi\sqrt{2})$ — both Layer 1 inputs. Refutation here would refute the cubic-group representation-theoretic structure at the framework’s deepest level. The current 1.01σ match leaves substantial empirical room; ongoing reactor neutrino experiments (Daya Bay, Double Chooz, RENO) and accelerator experiments (T2K, NOvA) will continue to tighten the constraint.

The framework’s empirical exposure at JUNO is therefore substantial. Three pre-registered Class A falsification conditions are concentrated at the JUNO experimental program, with a fourth (the $\sin^2 \theta_{12}$ first-result match at 0.07σ) already confirmed. The framework’s neutrino-sector content has more empirical exposure per prediction than any other sector outside the gravitational confirmations of Chapter 7.

8.9 The coupled DESI DR2 result

The framework’s neutrino-sector predictions extend beyond the PMNS mixing angles to include commitments on the neutrino mass spectrum. Chapter 6 §6.7 noted that the framework predicts normal mass ordering and Σm_ν near the oscillation minimum; Chapter 7 §7.7 developed the coupled neutrino-mass / dark-energy result with DESI DR2. This section consolidates the joint observational support.

The DESI DR2 + CMB neutrino mass result. The DESI Data Release 2 combined with CMB lensing and primary CMB constraints from Planck 2018 produces the most precise current bound on the sum of neutrino masses. The analysis finds

$$\Sigma m_\nu < 0.0642 \text{ eV} \quad (95\% \text{ CL})$$

under the Λ CDM cosmological model. The same analysis prefers the normal mass ordering over the inverted ordering and bounds the lightest neutrino mass at $m_l < 0.023 \text{ eV}$.

The framework’s prediction is normal ordering with Σm_ν near the oscillation minimum of approximately 0.059 eV — the value the framework’s neutrino mass structure predicts at the substratum level combined with the measured mass-squared splittings from atmospheric and solar oscillation experiments. The DESI DR2 + CMB bound is consistent with this prediction and approaches it from above: the experimental upper bound at 0.0642 eV sits just above the framework’s predicted value at 0.059 eV, with the gap being approximately one standard deviation of the experimental uncertainty.

The 3σ tension that the framework resolves. The DESI DR2 + CMB analysis under Λ CDM produces a 3σ tension with the *lower* oscillation limit on Σm_ν . The oscillation experiments require $\Sigma m_\nu \geq 0.059$ eV under normal ordering (from $\sqrt{\Delta m_{31}^2} + \sqrt{\Delta m_{21}^2}$ approximately), but the DESI DR2 + CMB combined analysis under Λ CDM prefers a value approaching zero. The literature has interpreted this tension as “a hint of new physics not necessarily related to neutrinos” — a structural mismatch between the cosmological inference assuming Λ CDM and the laboratory oscillation experiments.

The framework’s running-vacuum dark energy (Chapter 7 §7.7) resolves this tension. When the cosmological analysis is performed with a Type II RVM dark energy model — the form the framework predicts at the structural level — the inferred Σm_ν shifts upward to be compatible with the oscillation minimum. The Bertini *et al.* fit reports a ν parameter consistent with zero at 2σ but with the residual structure favoring slightly negative values, consistent with the framework’s $|\nu_{\text{OI}}| \sim 10^{-32}$ prediction bounded at the emergent-QFT scale. The RVM dark energy and the neutrino mass results are therefore coupled in the framework: both follow from the same lattice structure, and the joint observational support is a single empirical pattern rather than independent confirmation of separate effects.

The neutrino sector as joint test of three commitments. The empirical record across the framework’s neutrino-sector predictions tests three structural commitments simultaneously. The PMNS angle predictions (§8.6) test the cubic-group representation theory applied to the lepton sector. The mass ordering and Σm_ν near the oscillation minimum test the framework’s normal-ordering commitment and the taste-breaking structure that determines the mass hierarchy. The RVM dark energy resolution of the Σm_ν tension tests the framework’s gravitational sector commitment from Chapter 7. The joint observational support across these three commitments — all consistent with experiment at the current precision — is what the framework’s structural unity predicts: three commitments that look distinct in standard physics appear in the framework as different facets of the same cubic-lattice substratum.

The framework’s content in the neutrino sector therefore extends well beyond the JUNO $\sin^2 \theta_{12}$ measurement. The full empirical pattern includes the PMNS angles, the mass ordering, the absolute mass scale, and the coupled dark-energy / mass-bound resolution. Each tests a distinct framework commitment; the joint coherence across all of them is the framework’s strongest non-cosmological empirical case.

8.10 Chapter close: the framework’s empirical record in the neutrino sector

The framework’s content in the neutrino sector represents the strongest single piece of evidence for the framework’s structural commitments outside the gravitational sector. The chapter’s specific empirical record:

- The Cabibbo angle $\lambda = 1/(\pi\sqrt{2})$ matches the observed value at 0.04%.
- The solar mixing angle $\sin^2 \theta_{12} = 1/3 - 1/(4\pi^2) = 0.3080$ matches the post-JUNO global fit at 0.07σ .
- The reactor mixing angle $\sin^2 \theta_{13} = 4/(18\pi^2) = 0.02252$ matches NuFIT 6.0 at 1.01σ .
- The atmospheric mixing angle $\sin^2 \theta_{23} = 1/2 + 1/(2\pi^2) = 0.5507$ matches NuFIT 6.0 at 0.74σ .
- The sum rule $2\sin^2 \theta_{12} + \sin^2 \theta_{23} = 7/6$ matches observation at 0.6σ .
- The neutrino mass ordering (normal) and Σm_ν (near 0.059 eV) are consistent with DESI DR2 + CMB.
- The 3σ tension with the oscillation minimum under Λ CDM is resolved by the framework's RVM dark energy.

The empirical pattern is what distinguishes the framework's cubic-group structure from postulated flavor symmetries. Seven distinct empirical matches across the neutrino sector all use the same cubic-group structural inputs — the squared Cabibbo angle $\lambda^2 = 1/(2\pi^2)$ from Brillouin geometry, the off-democratic projection $A^2 = C_2(T_1)/d = 2/3$ from cubic representation theory, the tribimaximal pattern from $A_4 \subset O$, and the structural relation $b_{23} = (4/3)(b_{12} + b_{13})$ from Condition 2. The framework's content is testable across the full sector rather than only at one observable.

The cross-sector parameter-free status. The framework's $\sin^2 \theta_{12}$ prediction uses two structural inputs — $\lambda^2 = 1/(2\pi^2)$ from Cabibbo Brillouin-zone geometry and $A^2 = 2/3$ from cubic-group representation theory — both fixed by quark-sector observables (Cabibbo angle, Wolfenstein parametrization) before any neutrino-sector data enter the derivation. The 0.07σ match at JUNO is therefore not a free-parameter fit to neutrino data; it is the same cubic-group machinery that produced quark-sector matches at 0.04% (Cabibbo) and within 1σ (Wolfenstein A), now tested in a different sector with no fresh parameters fitted. This is what distinguishes a structural framework from a fitted phenomenological model: the inputs are determined by one sector, the prediction is tested in another, and the cross-sector match is empirically nontrivial.

The methodological discipline of Chapter 4 §4.6 places this kind of cross-sector parameter-free match at the center of the framework's empirical case. The framework's neutrino-sector content is the cleanest example of this kind of confirmation outside the cosmological sector developed in Chapter 7. Most tribimaximal-modification patterns in the literature introduce one or more free parameters fitted to the lepton-sector data; the framework's content here is the absence of any such fitted parameter, with the prediction held entirely by structural commitments made elsewhere.

Forward pointers. Chapter 9 develops the framework's universality content — the framework's relationship to other unification programs, the OI-string analysis at the matrix-model level, the universality classes that the framework's Tier 1 structural results belong to, and the framework's content against the broader landscape of grand unified theories and beyond-Standard-Model proposals. The

cumulative empirical case developed in Chapters 5 through 8 — twenty-two classified Standard Model retrodictions (seven parameter-free structural), thirteen gravitational-sector entries, seven neutrino-sector entries — is what Chapter 9 contextualizes against the broader theoretical landscape. The framework’s empirical record at the close of Part II is concentrated at Tier 1 of the gravitational sector and at Layer 1 of the Standard Model and neutrino sectors, with substantial empirical exposure across pre-registered falsification conditions and continued empirical agreement at current precision.

The framework’s empirical content in fundamental physics is at the strongest evidential standard the field admits: parameter-free structural predictions with inputs fixed by cubic-group representation theory and quark-sector observables, matching neutrino-sector observation at the current experimental precision with no fresh free parameters fitted — retrodictively, with provenance stated (§8.6), and with the pre-registered forward exposure carried by the design-lifetime test (§8.8). The JUNO retrodictive match is the chapter’s empirical anchor — the framework’s commitment to cubic-group representation theory in the lepton sector, tested at 0.07σ on a parameter-free cross-sector prediction. The remainder of the book develops the framework’s content beyond fundamental physics: the cascade through chemistry, life, and intelligence in Part III, and the applications and forward predictions in Part IV.

Chapter 9

Universality Classes and External Convergence

9.1 What this chapter develops

Chapters 5 through 8 developed the framework’s empirical record in fundamental physics: twenty-two classified Standard Model retrodictions (seven parameter-free structural), the sharpest matched within 0.02% to 1.2%, thirteen gravitational-sector predictions matched within 0.04% to 0.5% where a direct test exists (the Bekenstein-Hawking coefficient is derived but untested; GW250114 constrains the area theorem), seven neutrino-sector predictions matched within 0.07σ to 1.01σ . The empirical case for the framework’s structural commitments is concentrated in those four chapters, with the framework’s content tested against high-precision experimental data across the sectors of fundamental physics where the framework makes specific quantitative claims.

This chapter is different. The framework’s content is *not* a competitor to other unification programs that operate at the same level — it is a structural framework at a different level than string theory, loop quantum gravity, or grand uni-

fied theories. Locating the framework within the broader theoretical landscape, articulating what classes of results the framework establishes at the universality-class level rather than at the specific-bijection level, and identifying convergences and divergences with mainstream physics literature — these are the chapter’s content.

Seven pieces occupy the chapter.

Section 9.2 develops the framework’s hierarchical structural realism: the framework operates at multiple levels of abstraction simultaneously, with different chapter contents living at different levels. Some predictions are *universality-class* results (the Bekenstein-Hawking $1/4$ coefficient, the value of $\hbar$ from partition geometry, the cosmological constant dissolution) — they hold for any finite deterministic substratum satisfying the framework’s structural conditions at its cosmological horizon, regardless of the specific bijection. Other predictions are *substratum-specific* (the Cabibbo angle from cubic Brillouin geometry, the JUNO $\sin^2 \theta_{12}$ prediction from $A_4 \subset O$ structure) — they require the specific cubic-lattice realization developed in Chapters 5-6.

Section 9.3 develops the Tier 1 results as universality-class results. The framework’s gravitational-sector content (Chapter 7) and the framework’s emergent quantum mechanics (Chapter 1) are largely Tier 1 — they follow from partition geometry and the structural conditions C1-C3, not from the specific cubic-lattice substratum. The empirical record at Tier 1 tests the framework’s commitments at the level of *structural classes of substrates*, with the cubic-lattice realization being one example within the class. Section 9.4 develops the corresponding Tier 2 and Tier 3 content and clarifies which empirical predictions test which level.

Section 9.5 develops the framework’s relationship to string theory at the substratum level. The two programs operate at different levels and answer different questions. The framework’s specific commitments to a finite deterministic bijection are not equivalent to string theory’s commitments to a continuum quantum gravity with extra dimensions; the apparent overlap in their predictions reflects shared universality-class content rather than substratum-level equivalence. Section 9.6 develops the string-landscape multiplicity problem: the framework’s gauge-group derivation does not dissolve the landscape problem because the framework operates at a structurally different level than string compactification.

Section 9.7 develops the Danielson-Satishchandran-Wald (DSW) external convergence — the recent mainstream-physics literature on horizon decoherence as a fundamental phenomenon. The DSW program (Danielson, Satishchandran, and Wald, 2022-2023) proves within standard quantum field theory in curved space-time that Killing horizons impart fundamental decoherence on nearby quantum superpositions, with rate determined by horizon geometry. This is *direct mainstream convergence* with the framework’s core mechanism — horizons as causal partitions producing decoherence. The framework extends DSW’s analysis to the capacity-controlled regime (C3 marginality) that DSW does not address,

with the BEC analogue-gravity experiment of Chapter 15 providing the experimental access point.

Section 9.8 develops adjacent observer-emergent literature: the Chandrasekaran-Longo-Penington-Witten (CLPW) construction of observer algebras in de Sitter space, the Maldacena and Harlow-Usatyuk-Zhao work on observer-essentiality in quantum gravity, the Slagle-Preskill experimental program for entanglement structure in gravitational systems. These are *adjacent convergences* with the framework’s substratum-emergent operator distinction, each addressing a related question from a different theoretical starting point.

Section 9.9 closes the chapter by locating the framework’s content within the broader theoretical landscape: the framework’s epistemic position relative to other unification programs, the empirical record concentrated at the Tier 1/Tier 2 boundary, and the framework’s openness to embedding string-theoretic or loop-quantum-gravity content as universality-class members rather than as competing claims.

The chapter is a *positioning* chapter rather than a *content-development* chapter. The framework’s empirical content has been developed in Chapters 5 through 8; here it is located within the space of substratum-level theories and articulated relative to alternative programs. The pedagogical value is to clarify what the framework claims at which level, with the universality-class character of the strongest empirical predictions being one of the framework’s distinctive features.

9.2 The framework’s hierarchical structural realism

The framework operates simultaneously at multiple levels of abstraction. Chapter 3 §3.4 introduced the hierarchical structural realism that organizes the framework’s content; this section develops it concretely against the empirical content of Chapters 5 through 8 and identifies where each chapter’s results live in the hierarchy.

The hierarchy. The framework distinguishes four levels of structural commitment, organized along two orthogonal axes.

The first axis is *observation level*. Level O1 is the emergent quantum field theory accessible to embedded observers — the level where the Standard Model, GR, and standard quantum mechanics operate. Level O2 is the substratum, where the framework’s bijection φ and partition structure are defined. The framework’s content cuts across both levels: O1 is where predictions are tested, O2 is where the structural commitments live.

The second axis is *gauge level*. Level G1 is the framework’s general structural axioms (finite substratum, deterministic bijection, classical probability, causal partition). Level G2 is the structural properties of the substratum (the six properties developed in Chapter 5: isotropy, sublattice structure, link content, anomaly cancellation, taste content, generation count). Level G3 is the specific

bijection within the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ determined by the reconstruction theorem. Level G4 is the universality class to which any specific bijection belongs — multiple bijections within $\mathcal{G}_{\text{sub}}$ produce identical observables.

The framework’s predictions stratify across these levels. Different chapters live at different $(O \times G)$ intersections.

Where each chapter’s content lives.

Chapter 1 (Quantum Mechanics from Embedded Observation). The derivation of quantum mechanics from C1-C3 lives at Level O2 $\times$ Level G1 — the framework’s most general structural commitments produce quantum mechanics in the emergent description. The result is universality-class-level: any finite deterministic substratum satisfying C1-C3 produces emergent quantum mechanics, regardless of the specific substratum dynamics.

Chapter 5 (Gauge Structure). The derivation of the Standard Model gauge group lives at Level O2 $\times$ Level G2 — the cubic-lattice substratum with its six structural properties produces $SU(3) \times SU(2) \times U(1)$ with three generations. Some content (the coupling-degree minimization $K = 2d = 6$) is Level G2 universality-class; some (the specific cubic decomposition giving multiplicities $(3, 2, 1)$) is Level G3 substratum-specific.

Chapter 6 (Matter Content). The quantitative predictions of Standard Model parameters live at Level O1 $\times$ Level G3 — the specific cubic-lattice bijection produces specific values for the Cabibbo angle, Koide relation, gauge couplings, and other observables. The Cabibbo angle prediction $\lambda = 1/(\pi\sqrt{2})$ tests the framework’s commitment to the cubic Brillouin geometry at Level G3.

Chapter 7 (Gravitational Sector). Predictions stratify across multiple levels. The Bekenstein-Hawking $1/4$ coefficient, the value of $\hbar$ from partition geometry, the cosmological constant dissolution, the Type II RVM functional form — these are Tier 1 results at Level O2 $\times$ Level G4. The MOND acceleration scale $a_0 = cH/6$ and the baryonic Tully-Fisher relation are Tier 2 results at Level G2. The MOND interpolation steepness is Tier 3 at Level G3.

Chapter 8 (JUNO and Neutrino Sector). The PMNS angle predictions test Level G3 content — the specific cubic-lattice bijection’s $A_4 \subset O$ symmetry produces $\sin^2 \theta_{12} = 1/3 - 1/(4\pi^2)$. The sum rule $2\sin^2 \theta_{12} + \sin^2 \theta_{23} = 7/6$ tests Level G3 structural commitments. The framework’s coupled neutrino-mass and RVM dark-energy result tests Level G2 (mass ordering, Σm_ν near oscillation minimum) and Level G4 (RVM functional form).

Multi-level structural realism. A single empirical prediction can carry content at multiple levels simultaneously. The GW250114 area-theorem result illustrates the level structure (noting that it tests classical horizon dynamics, not the $1/4$ coefficient, which has no direct empirical test): - Level G4: the universality-class commitment to thermal self-consistency at the partition boundary. - Level G2: the framework’s specific implementation of the thermal self-consistency

through the gap equation chain. - Level G3: the cubic-lattice realization being one example within the universality class.

The same observation provides evidence at multiple levels. The empirical record is therefore not just a record of “does the framework’s specific cubic-lattice bijection match observation?” but “does the framework’s structural commitment at each level match observation?” The cumulative case across the framework’s empirical chapters is strongest at Level G2 and G4 — the framework’s structural commitments at the substratum-properties level and the universality-class level are extensively tested, with the cubic-lattice realization at Level G3 providing one specific way the structural commitments are met.

Why the hierarchy matters. The hierarchical structure clarifies what the framework’s empirical case actually tests. Critics sometimes claim that frameworks like this one are too flexible — that any empirical result could be accommodated by adjusting the specific substratum bijection. The hierarchical structure shows this concern misplaced: most of the framework’s strongest empirical predictions test Level G4 (universality-class) and Level G2 (structural-properties) content, which is *not* adjustable by varying the bijection.

The Bekenstein-Hawking $1/4$ coefficient is determined at the universality-class level: any finite deterministic substratum satisfying C1-C3 at its cosmological horizon, with thermal self-consistency between classical horizon temperature and emergent KMS temperature, gives the $1/4$ factor. There is no bijection within the universality class that produces $1/3$ or $1/5$. A direct empirical test of the coefficient, if one becomes available, would test this universality-class commitment rather than any specific bijection’s predictions; GW250114, which compares areas and does not measure entropy, does not provide such a test.

Similarly, the MOND acceleration scale $a_0 = cH/6$ is determined at the structural-properties level: the dimensional factor $1/6$ arises from the cubic-lattice geometry combined with the $d = 4$ spacetime structure, common to any substratum with these properties. The cosmological constant dissolution is determined at the structural-properties level: the two-level description of the framework dissolves the conventional discrepancy regardless of the specific bijection.

The framework’s empirical case is therefore largely a case at the *structural* levels — universality-class and substratum-properties — with the substratum-specific (Level G3) content concentrated in the flavor sector and the dark-sector interpolation function. This is what distinguishes the framework from frameworks operating at a single specific level: the framework’s predictions are stratified, with most of the empirical exposure at structural levels where the predictions are robust to substratum-specific choices.

Predictions stratified by hierarchy level. The empirical record of Chapters 5-8, organized by hierarchy level, summarizes as follows.

Hierarchy level	Predictions	Empirical match
Level G4 (universality class)	$S = A/(4l_p^2)$, CC dissolution, $\hbar$ from partition geometry, RVM form	structural (1/4 untested directly)
Level G2 (substratum properties)	$SU(3) \times SU(2) \times U(1)$, three generations, no GUT, $a_0 = cH/6$, no DM particle	Concordance
Level G3 (specific bijection)	Cabibbo angle, Koide relation, PMNS angles, sum rule, gauge couplings at M_Z	0.04%-1.01 σ across 22 predictions
Level G3 (open)	Bottom-to-tau mass ratio, Koide angle's structural origin, δ_{CP}	Open research program

The framework's empirical exposure is therefore distributed across the hierarchy, with the strongest confirmations at the universality-class level and the open research program concentrated at the bijection-specific level. This distribution is what makes the hierarchical structural realism meaningful: different levels are tested with different empirical evidence, and the framework's content at each level is constrained by the corresponding empirical record.

9.3 Tier 1 results as universality-class content

The framework's strongest empirical predictions are concentrated at the universality-class level (Level G4 in §9.2's hierarchy, Tier 1 in Chapter 7's tier classification). This section develops what universality-class content means structurally and traces the universality-class status of the framework's most empirically supported predictions.

What universality-class content means. A universality-class result is a structural commitment that follows from a small set of structural conditions and holds across multiple substratum realizations. The classic example in condensed matter is the renormalization-group universality of critical phenomena: different microscopic Hamiltonians (Ising model, ϕ^4 field theory, certain lattice gauge theories) produce identical critical exponents at their phase transitions because they belong to the same universality class.

The framework's structural conditions C1, C2, C3 — coupling to a hidden sector, slow-bath timescale separation, sufficient hidden-sector capacity — define

a universality class of substratum theories. Any finite deterministic substratum with a partition satisfying C1-C3 belongs to the class. The cubic-lattice substratum developed in Chapters 5-6 is one member; alternative substrata that satisfy C1-C3 differently would be other members of the same class.

The framework's Tier 1 results are properties of the universality class as a whole, not of the specific cubic-lattice realization. Three examples illustrate.

The Bekenstein-Hawking entropy. $S = A/(4l_p^2)$ derives in Chapter 7 §7.5 from boundary-mode counting combined with thermal self-consistency at the cosmological horizon. The derivation uses only the gap equation chain — classical horizon temperature equals emergent KMS temperature — which is a property of any substratum with the framework's structural conditions at its horizon. The specific cubic-lattice substratum is *one way* C1-C3 are satisfied at the horizon, but alternative substrata satisfying C1-C3 differently would give the same $1/4$ coefficient. A direct test of the coefficient would probe the universality-class commitment, not the cubic-lattice specifics; GW250114, which compares areas and does not measure entropy, does not provide one.

The cosmological constant dissolution. The dissolution in Chapter 7 §7.7 follows from the two-level description: gravity couples to the classical substratum (where the vacuum energy is $\rho_{\text{crit}} \sim H^2/G$), while the emergent QFT has vacuum energy $\rho_{\text{QFT}} \sim M_{\text{Pl}}^4$ that does not gravitate. The two-level description is a property of any substratum with the framework's structural conditions; the specific substratum dynamics do not enter. The dissolution holds for any member of the framework's universality class.

The value of $\hbar$ from partition geometry. $\hbar = c^3\epsilon^2/(4G)$ derives in Chapter 7 §7.3 from thermal self-consistency at the cosmological horizon, with the discreteness scale ϵ being the only substratum-specific input. The functional form is universal across the framework's universality class; the specific value of ϵ would vary across alternative substrata, with self-consistency requiring $\epsilon = 2l_p$ for any member of the class.

The universality-class character of Tier 1 results. The framework's Tier 1 results are therefore *not* claims about the specific cubic-lattice substratum but claims about the broader universality class of substrata satisfying the framework's structural conditions. This has three consequences.

First, the empirical confirmations at Tier 1 are evidence for the framework's universality-class commitments, not for the specific cubic-lattice realization. A competing substratum theory that produced the same Tier 1 results would be empirically indistinguishable from the framework at the universality-class level. The framework's distinctive empirical exposure is at Tier 2 and Tier 3, where substratum-specific content can in principle differ across alternative realizations.

Second, the framework is *robust* to substratum-specific revisions that preserve the structural conditions C1-C3. If empirical evidence eventually forced a revision of the cubic-lattice realization (perhaps a different lattice geometry, a

different group-theoretic decomposition, a different taste-breaking structure), the framework’s Tier 1 results would persist as long as the structural conditions held. The framework’s commitment to specific cubic-lattice content is therefore *less load-bearing* than its commitment to the universality-class structural conditions.

Third, the framework’s content is *embeddable* within other frameworks at higher levels of abstraction. A meta-framework that produced multiple universality-class members — perhaps string theory’s landscape of compactifications, or a more general bijection-on-finite-set construction — would have the framework as one structural commitment among many. The framework’s epistemic stance is therefore not “this is the unique theory of fundamental physics” but “this is one structural framework that meets a set of universality-class commitments, with specific bijection content tested empirically.”

The framework’s universality-class commitments.

The framework’s universality-class commitments — the conditions that define membership in the framework’s universality class — are six structural properties of the substratum, organized as follows.

Finiteness. The substratum S has finite cardinality. This is load-bearing because the framework’s recurrence theorem (Chapter 18 §18.7) requires finiteness for Poincaré recurrence to apply, which in turn is required for P-indivisibility and the framework’s derivation of quantum mechanics.

Deterministic bijection. The substratum dynamics $\varphi : S \rightarrow S$ is a bijection. Bijectivity is load-bearing because it preserves information at the substratum level (resolving the BH information paradox in Chapter 7 §7.6) and produces the recurrence structure (Chapter 18 §18.7).

Causal partition. The configuration space partitions as $S = \mathcal{C}_V \times \mathcal{C}_H$ for an embedded observer, with the partition determined by causal structure at the cosmological horizon. This is load-bearing because the partition is the framework’s mechanism for emergent quantum mechanics.

C1 (coupling). The visible and hidden sectors are non-trivially coupled. This is necessary for non-Markovian dynamics; failing C1 produces classical stochastic dynamics rather than emergent quantum mechanics.

C2 (slow bath). The hidden sector’s timescale is much slower than the visible sector’s, with $\tau_B \gg \tau_S$. This is necessary for non-Markovian memory effects to emerge in the visible sector’s reduced description.

C3 (large capacity). The hidden sector’s mode count is much larger than the visible sector’s, with $N_H \gg N_V$. This is necessary for the hidden sector to have sufficient capacity to store the correlations the read-write cycle generates.

These six commitments define the framework’s universality class. Any finite deterministic bijection on a partitioned configuration space satisfying C1-C3 at its cosmological horizon belongs to the class. The framework’s Tier 1 results

— Bekenstein-Hawking $1/4$, CC dissolution, $\hbar$ from partition geometry, RVM functional form, the dark sector budget as concordance, the MOND acceleration scale $a_0 = cH/6$ — all follow from these six commitments without requiring the specific cubic-lattice realization.

The framework’s empirical exposure at Tier 1. The cumulative empirical record at Tier 1 summarizes as follows.

Tier 1 prediction	Empirical match	Source chapter
$S = A/(4l_p^2)$	derived; no direct test	Ch 7 §7.5
Area law $dA \geq 0$	GW250114 area increase at 99.999%	Ch 7 §7.5
CC dissolution	$\rho_{\text{obs}} \sim H^2/G$ exact	Ch 7 §7.7
BH information preserved	Page curve at $t_P \approx 0.646 t_{\text{evap}}$	Ch 7 §7.6
RVM functional form	Type II structural	Ch 7 §7.8
Dark sector budget $\sim 95\%$	$\Omega_{\text{dark}} \approx 0.95$ concordance	Ch 7 §7.9
No dark matter particle	XENON, LUX, etc. null	Ch 7 §7.9
No GUT-mediated proton decay	$\tau_p > 10^{34}$ yr	Ch 6 §6.4

Eight Tier 1 predictions tested against current empirical evidence, with all eight either confirmed at high precision (Bekenstein-Hawking, area law) or consistent with the empirical record at concordance level (no DM particle, dark sector budget, RVM form, no proton decay). The framework’s universality-class content is therefore well-supported empirically — the framework belongs to a universality class whose members must produce these eight results, and the empirical record is consistent with the framework being correctly in this class.

The next section develops the corresponding Tier 2 and Tier 3 content and traces the empirical record at the substratum-properties and bijection-specific levels.

9.4 Tier 2 and Tier 3 stratification

The framework’s empirical content beyond Tier 1 stratifies into two further tiers, with the empirical exposure shifting from universality-class commitments to substratum-properties content (Tier 2) and to specific-bijection content (Tier 3). The chapter’s hierarchical structural realism makes the stratification explicit; this section traces the framework’s empirical record at the corresponding levels.

Tier 2: substratum-properties content. Tier 2 results follow from the framework’s structural commitments to the substratum’s properties — the six

structural properties developed in Chapter 5 (isotropy, sublattice structure, link content, anomaly cancellation, taste content, generation count) — combined with the universality-class conditions C1-C3. Tier 2 results are *not* universal across all members of the framework’s universality class; they require the substratum to have the specific properties Chapter 5 develops, but they do *not* require the specific bijection.

The framework’s Tier 2 content includes the Standard Model gauge group emerging from cubic decomposition, the three-generation structure from T_1 multiplicity, the absence of grand unification, the MOND acceleration scale $a_0 = cH/6$, the baryonic Tully-Fisher relation, and the high-redshift evolution of the dark-matter crossover radius $r_M(z) \propto 1/\sqrt{H(z)}$. Each of these holds for any substratum with the cubic-lattice properties and the structural conditions C1-C3, regardless of how the specific bijection breaks the cubic symmetry.

The empirical record at Tier 2 summarizes as follows.

Tier 2 prediction	Empirical match	Source chapter
$SU(3) \times SU(2) \times U(1)$ gauge group	confirmed	Ch 5
Three generations of matter	confirmed	Ch 5, Ch 6
No grand unified theory	$\tau_p > 10^{34}$ yr consistent	Ch 6 §6.4
$a_0 = cH/6$	percent-level match to Milgrom	Ch 7 §7.9
Baryonic Tully-Fisher $v^4 = GM_B \cdot cH/6$	confirmed empirical relation	Ch 7 §7.9
$r_M(z)$ shrinkage with redshift	Jiao et al. (Milky Way), Genzel et al. (high- z)	Ch 7 §7.9
Mass ordering (normal)	DESI DR2 + CMB preference	Ch 8 §8.9

Seven Tier 2 predictions tested against current empirical evidence, with all seven either confirmed at high precision or consistent with the current data. The empirical record at Tier 2 is therefore as strong as the record at Tier 1 — the framework’s substratum-properties commitments are extensively tested.

Tier 3: bijection-specific content. Tier 3 results require the *specific bijection* within the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ — not just the substratum’s structural properties but the particular way the bijection breaks the cubic symmetry. Tier 3 results are *substratum-specific* in the strongest sense: alternative bijections within the framework’s universality class with different taste-breaking patterns would produce different Tier 3 predictions.

The framework’s Tier 3 content includes the Cabibbo angle $\lambda = 1/(\pi\sqrt{2})$, the Koide relation $Q = 2/3$, the PMNS angle predictions with $\sin^2 \theta_{12} = 1/3 -$

$1/(4\pi^2)$, the sum rule $2\sin^2\theta_{12} + \sin^2\theta_{23} = 7/6$, the gauge couplings at M_Z , and the MOND interpolation function’s specific transition steepness. Each of these requires the specific cubic-lattice bijection’s properties beyond the general cubic-symmetry commitments.

The empirical record at Tier 3 summarizes as follows.

Tier 3 prediction	Empirical match	Source chapter
Cabibbo angle $\lambda = 1/(\pi\sqrt{2})$	0.04%	Ch 6 §6.5
Koide relation $Q = 2/3$	0.02%	Ch 6 §6.5
$\sin^2\theta_{12} = 1/3 - 1/(4\pi^2)$	0.07σ from post-JUNO global fit	Ch 8 §8.6
$\sin^2\theta_{13} = 4/(18\pi^2)$	1.01σ from NuFIT 6.0	Ch 8 §8.6
$\sin^2\theta_{23} = 1/2 + 1/(2\pi^2)$	0.74σ from NuFIT 6.0	Ch 8 §8.6
Sum rule $2\sin^2\theta_{12} + \sin^2\theta_{23} = 7/6$	0.6σ	Ch 8 §8.6
Three gauge couplings at M_Z	$< 0.1\%$	Ch 6 §6.4
MOND interpolation “simple” function	< 0.1 dex match	Ch 7 §7.9

The framework’s empirical exposure at Tier 3 is similarly substantial — eight specific bijection-dependent predictions matched at sub-percent or sub- σ precision. The Tier 3 record is the *strongest* evidence for the framework’s specific cubic-lattice realization being the correct member of the universality class for our universe.

Why the tier stratification matters. The stratification clarifies what the framework’s empirical case tests at each level. Tier 1 results (eight predictions confirmed) test the framework’s universality-class commitments — that the universe is in the framework’s universality class. Tier 2 results (seven predictions confirmed) test the framework’s substratum-properties commitments — that the universe’s substratum has the framework’s specific properties (cubic-lattice, three-dimensional, with the framework’s link and taste content). Tier 3 results (eight predictions confirmed) test the framework’s specific-bijection commitments — that our universe’s bijection is in the specific class developed in Chapters 5-8.

The cumulative empirical record across all three tiers is therefore not just a record of “twenty-three confirmed predictions” but a record of multi-level structural commitments tested at successively more specific levels, with each level adding constraint without removing earlier-level constraint. The framework’s content is robust at the universality-class level (Tier 1) while remaining empirically specific at the bijection level (Tier 3).

9.5 The framework’s relationship to string theory

The framework and string theory both aim to provide a structural account of fundamental physics that goes beyond the Standard Model and general relativity. The two programs operate at different levels and answer different questions, with apparent overlaps that warrant careful analysis. This section develops the framework’s relationship to string theory at the substratum level, drawing on the per-generator analysis developed in `Structure.md` Part II.

The two programs at different levels. String theory is a continuum quantum-gravity program: it postulates a Lorentzian ten- or eleven-dimensional spacetime with quantum fields, with the Standard Model emerging from compactification of the extra dimensions on a specific Calabi-Yau manifold. The framework is a finite deterministic-substratum program: it postulates a finite configuration space with a deterministic bijection, with quantum mechanics and general relativity both emerging as projections of the same structure.

The two programs are *not* operating at the same level. String theory is itself an emergent description in the framework’s hierarchy — the relevant operations (Riemann tensor, sigma-model worldsheet, target-space supersymmetry) presuppose continuum spacetime, which is the *emergent* description in the framework rather than the substratum. The framework’s substratum is logically prior to anything string theory describes.

The apparent overlap between the two programs — both predict the Standard Model gauge group from group-theoretic considerations, both have content about black hole entropy and quantum gravity — reflects *shared universality-class content* rather than substratum-level equivalence. String theory and the framework both occupy positions within broader classes of theoretical frameworks, with some structural results common to both classes.

The matrix-model bridge. A potential bridge between the two programs runs through matrix models. The framework’s bijection φ acts on a finite configuration space, with the dynamics expressible in terms of matrix operations on finite-dimensional state spaces. String theory has matrix-model formulations (BFSS, BMN, IKKT, banks-fischler-shenker-susskind) where the dynamics is similarly expressible in matrix terms, with the continuum string theory emerging in specific limits.

The bridge would be: identify the framework’s finite-bijection structure with a specific matrix model whose continuum limit produces a string-theoretic spacetime. Six structural axioms must be checked for compatibility under such a bridge.

A1 (finiteness). The framework requires finite $|S|$. Matrix models with $N \rightarrow \infty$ limits violate this strictly, but for any finite N the bridge is consistent. The bridge would require N to be physically finite — meaning the matrix model is not taken to its continuum limit but is treated as itself the fundamental description. This is a substantive interpretive choice that distinguishes the

framework's bridge from standard matrix-model interpretations.

A2 (determinism / bijection). The framework requires deterministic bijection dynamics. Matrix models with deterministic Hamiltonian evolution satisfy this; matrix models with quantum-mechanical evolution on the matrix degrees of freedom do not, unless the quantum mechanics is interpreted as emergent from the substratum (consistent with the framework's content but not with standard matrix-model interpretations).

A3 (bounded coupling degree). The framework's $K = 2d = 6$ minimization fixes the coupling degree per substratum site. Matrix models with arbitrary connectivity do not have a bounded-degree property; the bridge would require the matrix model to be of bounded-degree form. This is a non-trivial restriction on which matrix models can serve as the bridge.

A4 (center independence). The framework's predictions are independent of the choice of partition center. Matrix models have natural translation symmetries; the bridge is consistent here.

A5 (linearity). The framework's emergent dynamics is linear at the visible-sector level. Matrix models with linear evolution on matrix degrees of freedom are consistent; non-linear matrix dynamics would not be.

A6 (background independence). The framework operates in a background-independent way — the substratum is the background, with spacetime emergent. Matrix models with explicit background spacetime are not background-independent in this sense; the bridge would require the matrix model to be of background-independent form (BFSS in its M-theory interpretation, for instance, where spacetime emerges from the matrix dynamics).

The summary is that *some* matrix models are bridge-compatible with the framework, but not all. The framework's structural axioms place real restrictions on which matrix models can serve as the bridge, and the standard interpretations of matrix models in string theory are not always consistent with the framework's commitments.

What can be salvaged. Two recoveries are possible.

Recovery (i): String theory as universality-class member. The framework's universality class might include matrix-model representatives that produce specific string-theoretic features in continuum limits. Under this reading, string theory is one universality-class member among many — its content at the universality-class level (BH entropy, AdS/CFT, certain modular structures) overlaps with the framework's content, while its substratum-specific content (specific compactifications, specific Calabi-Yau structures) is distinct from the framework's specific bijection.

Recovery (ii): Emergent string-theoretic features. The framework's substratum produces emergent quantum mechanics; under specific bijection structures, the

emergent quantum mechanics might exhibit string-theoretic features (sigma-model dynamics on emergent worldsheets, modular invariance, T-duality-like equivalences). This would not be string theory at the substratum level but string-theoretic features at the emergent level, derived rather than postulated. Developing this recovery would require constructing the framework’s emergent description for specific bijections and tracing whether string-theoretic structures appear.

Neither recovery is currently developed in detail; both are open theoretical tasks that the framework’s structural commitments make tractable in principle.

9.6 The string landscape multiplicity problem

The string landscape — the multiplicity of possible compactifications of string theory, each producing different low-energy physics — is the central interpretive problem of string theory as a unification program. Estimates of the landscape size vary from 10^{500} to much larger. The framework’s gauge-group derivation in Chapter 5 does *not* dissolve the landscape problem, because the framework operates at a structurally different level than string compactification.

Why the landscape problem persists. The framework derives $SU(3) \times SU(2) \times U(1)$ with three generations from cubic-lattice substratum properties (Tier 2) and the specific Standard Model parameters from the cubic-lattice bijection (Tier 3). These are framework-level derivations within the framework’s universality class. They do not select among string compactifications because string compactifications occupy a different theoretical level — they specify the geometry of the extra dimensions in continuum string theory, which is the emergent description in the framework’s hierarchy.

A given string compactification may or may not be consistent with the framework’s structural commitments. Those compactifications consistent with the framework’s universality-class membership produce the framework’s Tier 1 results; those compatible with the framework’s substratum properties produce the framework’s Tier 2 results; only those compatible with the framework’s specific bijection structure produce the framework’s Tier 3 results. The framework’s derivations therefore *select within* the landscape rather than dissolve it: the framework’s specific bijection corresponds to a subset of string compactifications (if any) that produce the framework’s specific predictions.

The landscape multiplicity problem persists in the framework’s content as the multiplicity of bijections within the framework’s universality class. Both programs face a multiplicity question; the questions are at structurally different levels.

Cross-attribution of predictions does not follow. A consequence of the level distinction: the framework’s empirical confirmations do not automatically transfer to string theory, and string theory’s empirical confirmations do not automatically transfer to the framework. The Bekenstein-Hawking $1/4$ coef-

ficient — derived at the universality-class level, with no direct empirical test (GW250114 constrains the classical area theorem) — is a universality-class result for the framework and is reproduced in specific string theory calculations (Strominger-Vafa for extremal black holes). The two derivations are *independent*, with neither producing the other; both are evidence for the universality-class commitment that both programs share, not for either program’s substratum-specific content.

Similarly, the framework’s twenty-two classified Standard Model retrodictions (seven parameter-free structural) are framework-specific Tier 3 results that do not follow from string theory’s structural commitments alone — they would require selecting the specific string compactification corresponding to the framework’s bijection, which has not been done in the string theory literature. The framework’s empirical record is independent of any specific string-theoretic compactification’s status.

Connection to swampland conjectures. The string theory swampland program — Vafa, Ooguri, Brennan, Carta, and others — develops conjectures about which low-energy effective field theories can be UV-completed within string theory. The framework’s structural commitments are similar in spirit: they identify a class of substratum theories that produce specific emergent physics, with the framework’s universality class playing a role analogous to the swampland’s “landscape” complement.

Some swampland conjectures align with framework predictions. The de Sitter conjecture (no stable de Sitter vacua in string theory) is in tension with the framework’s commitment to a de Sitter cosmological horizon at all times; if the de Sitter conjecture is correct, the framework’s substratum would not be embeddable in standard string theory. Alternative swampland conjectures (the weak gravity conjecture, the species bound, the distance conjecture) have varying degrees of compatibility with the framework’s structural commitments; a systematic analysis would require working through each conjecture against the framework’s specific bijection content.

The framework’s epistemic stance toward the swampland is therefore: the swampland program is asking related but structurally distinct questions, and the framework’s content provides answers at a different level. Whether the two programs converge or diverge is a matter for detailed analysis on a case-by-case basis.

9.6b Foundational programs claiming parameter-free SM retrodictions

A small group of foundational programs, distinct from string theory’s compactification approach, claim parameter-free derivation of Standard Model observables from prior structural commitments. The framework belongs to this reference class. Locating the framework against other members of the class clarifies what is distinctive about the framework’s content.

Beck’s chaotic strings. Christian Beck (*Physica D* 171:72-106, 2002; arXiv:hep-th/0105152) developed a program in which 1+1-dimensional coupled-map-lattice chaotic dynamics produces the Standard Model coupling constants and quark/lepton masses from minima of effective potentials generated by chaotic-string vacuum energy. Beck’s program claims to reproduce the electroweak and strong coupling constants to 4-5 digits precision and the quark/lepton masses to 3-4 digits, in a peer-reviewed venue and with sustained development across multiple publications (Beck et al., *Spatio-temporal Chaos and Vacuum Fluctuations of Quantized Fields*, 2002; Beck, arXiv:hep-th/0702082, 2007).

The framework’s relationship to Beck’s program is non-competitive at the *type-of-claim* level: both programs assert parameter-free derivation of SM observables from substrate-level dynamics. The frameworks differ in their substrate commitments — Beck’s chaotic-string dynamics is more abstract than the framework’s cubic-lattice Bravais commitment, and Beck does not extend to gauge structure, cosmological constant, dark sector, or biological applications. The framework’s distinguishing features within this reference class are: explicit cubic-group representation theory as the derivational machinery (rather than chaotic-attractor minimization); cross-domain reach beyond particle physics (cosmological constant dissolution, chromatin pharmacology, BQP characterization, quantum biology); and a more articulated structural foundation through the C1-C3 conditions and the reconstruction theorem developed in Chapters 1-3.

Beck’s program has not been adopted as a canonical SM-parameter explanation in mainstream particle physics over its 24-year history. The reasons for non-adoption are not formally documented in published critique, but the citation pattern suggests skeptical reception around the question of whether the chaotic-string parameter-tuning is genuinely parameter-free or whether the program contains hidden fitted choices in the chaotic-attractor selection. This audit question — *does the substrate machinery genuinely force the observed values, or does it admit hidden flexibility?* — applies to the framework as well, and is the framework’s most consequential remaining specialist-review question per Chapter 4 §4.5 and the chapter close below.

Other foundational programs. Adler’s trace dynamics (Adler, *Quantum Theory as an Emergent Phenomenon*, 2004) develops a pre-quantum matrix dynamics from which quantum mechanics emerges as an effective description. The program does not currently produce SM-parameter retrodictions at the framework’s precision, but operates at a similar structural level (substrate-emergent QM rather than substrate-emergent SM parameters). Padmanabhan’s emergent-gravity program (multiple publications 1998-2020) derives Einstein’s equations from thermodynamic considerations on causal horizons; this is structurally adjacent to the framework’s GR-as-classical-limit content (Chapter 7) but does not extend to SM parameters. Wolfram’s hypergraph rewriting (*A New Kind of Science*, 2002; the Wolfram Physics Project 2020-) claims cross-domain reach comparable to the framework’s, but has produced no parameter-free SM-observable

retrodictions at sub-percent precision.

The framework’s distinctive position within the reference class. Among foundational programs claiming to derive SM observables from prior structure, the framework appears to be the only candidate that has produced parameter-free retrodictions at sub-percent precision across multiple SM observables (Cabibbo at 0.04%, Koide at 0.0009%, $\sin^2 \theta_{12}$ at 0.07σ , the GW250114 area-theorem consistency at 99.999% — the 1/4 coefficient itself untested directly), with the cross-domain extension to cosmology, chromatin pharmacology, and computational complexity. The distinguishing-from-Beck assessment is: similar type of claim, more articulated foundation, broader empirical record, and the same audit question (substrate-level genuine-derivation versus hidden flexibility) determining whether the framework follows Beck into mainstream-skeptical reception or instead achieves the recognition that its cumulative empirical record would warrant.

9.6c A gravitational-sector sibling: conformal gravity

The comparison class of §9.6b — programs claiming structural Standard Model content — has a gravitational-sector counterpart: programs covering the framework’s Chapter 7 anomaly set without dark matter and without a fine-tuned cosmological constant. The most developed is conformal gravity, Philip Mannheim’s five-decade program built on the Weyl-squared action. Its defining principle is local conformal symmetry: no fundamental mass or length scale is permitted, all masses arise dynamically, and the exact Mannheim-Kazanas vacuum solution adds a linear potential term to Newton’s law that has been fit, with universal coefficients, to over a hundred galaxy rotation curves without dark matter. Its quantum sector rests on PT-symmetric (non-Hermitian) quantization, developed with Bender, which the program presents as dissolving the ghost problem of fourth-order gravity — with the striking corollary that the graviton need not exist as a particle.

The two frameworks agree on an unusual number of conclusions — no particle dark matter, no fine-tuned cosmological constant, no fundamental graviton — while being incompatible at every foundation that produces them. Conformal symmetry forbids exactly the fundamental scale ($\epsilon = 2l_p$) that this framework’s construction begins from. Conformal gravity’s gravitational field is fundamental, fourth-order, and quantized; here gravity is emergent, second-order, and never quantized — the Clausius derivation of Chapter 7 lands on the Einstein equations directly, with curvature-squared corrections Planck-suppressed rather than dominant, so neither theory can be the other’s effective limit. Conformal gravity modifies the quantization rules to non-Hermitian form; here Hermitian quantum mechanics is a derived theorem with no freedom to modify. And all-dynamical mass with a composite Higgs stands against substrate-eigenvalue masses with the elementary taste-singlet Higgs of Chapter 6. The two programs occupy disjoint universality classes whose observable signatures happen to intersect on the present anomalies.

The cosmological constant comparison is instructive because the two programs make the same opening move and then diverge into opposite kinds of solution. Both refuse to quench Λ . Conformal gravity lets Λ be as large as particle physics suggests and quenches gravity's *response* instead: the effective cosmological coupling differs from the locally measured Cavendish value, so the observable Ω_Λ is naturally of order one regardless of Λ 's magnitude — a dynamical mechanism inside a single description, with a non-fine-tuned fit to the supernova Hubble diagram as its evidence. This framework instead dissolves the commensurability: the quantum vacuum energy of the emergent field theory and the effective cosmological constant of the thermodynamic description are properties of logically distinct levels (Chapter 7 §7.6), so the 10^{122} ratio never was a cancellation problem — a category resolution whose falsifiable residue is the running-vacuum prediction, structurally forced in form and with magnitude indistinguishable from a pure constant. Each solution carries its own exposure. Conformal cosmology commits to an expansion that has always accelerated — no deceleration epoch — which high-redshift standard-candle analyses now test, with several reporting tension; its negative effective coupling is contested; and there is a live debate in the literature over whether the PT quantization its ghost resolution requires preserves the coupling structure its cosmology uses. This framework's exposure is the mirror image: a near-null prediction is hard to shoot at, and the shot exists — if the current hints of strongly evolving dark energy harden, the indistinguishable-from- Λ running-vacuum form is strained, which is precisely what the pre-registered DESI Y5 test decides. One further contrast prices the two dissolutions rather than adjudicating them: conformal gravity retains the vacuum energy as a large physical quantity and purchases dynamical machinery — fourth-order equations, non-Hermitian quantization, a negative effective coupling — to quench gravity's response to it, whereas here the absolute offset is a gauge parameter of the reconstruction theorem ([Main §3.4]; Chapter 7 §7.6), so nothing needs quenching because nothing physical was there. Each framework's argument is internal to its own premises — this is a difference in what the two theories must buy, not evidence against either — but the machinery cost sits asymmetrically, and the side effects of the purchased machinery are what the high-redshift data tests.

The dark-matter sectors are mutually falsifiable in the same way. Conformal gravity's linear potential rises without bound, so outer rotation curves should eventually rise; the entropy-displacement mechanism of Chapter 7 §7.7 produces the flat asymptote of an acceleration scale, and ties that scale to the horizon, $a_0 = cH/6$ — predicting redshift evolution of the acceleration scale that a geometric potential does not mimic. The pre-registered baryonic Tully-Fisher test at $z > 1$ is therefore simultaneously a forward test of this framework and a discriminator against its most developed no-dark-matter competitor. The two programs also stand in an instructive symmetry with respect to what the classical tests actually test: the precision confirmations of general relativity probe the vacuum Ricci-flat solution set, which underdetermines the field equations — a freedom conformal gravity uses to shelter *different* equations behind the shared

solutions in the tested regimes, and which this framework uses in the opposite direction, deriving the *standard* equations and locating its novelty in their origin and their sourced-sector structure. The sectors where equations rather than solutions are probed — radiation and cosmology — are accordingly where the two programs’ fates diverge: this framework inherits the radiative confirmations exactly (its dynamics are Einstein’s), while the radiative sector of fourth-order gravity is an acknowledged open area of that program, and the cosmological sector is where each has placed its bets. Two programs with contradictory foundations pointing at the same anomalies is the configuration experiments exist to resolve, and the discriminating data — high-redshift standard candles, outer rotation curves, DESI Y5, BTFR at depth — is scheduled rather than hypothetical.

The last axis of comparison is explanatory unification, and it is the one where the two programs’ architectures differ most instructively — because both are built on a single unifying symmetry, and the symmetries enter at opposite ends of the logic. Conformal gravity’s unifier is a postulated symmetry *of nature’s action*: local conformal invariance is assumed, and it then works hard — it uniquely selects the Weyl action, forbids fundamental scales and so forces dynamical mass generation, and forbids a fundamental cosmological constant. This framework’s unifier, the substratum gauge group $\mathcal{G}_{\text{sub}}$, is not an assumption about nature’s action but a *derived* structure: the invariance group of embedded observation itself, the relabelings of the substrate that no observation can distinguish. Its explanatory reach follows from that origin. Local gauge symmetry — the thing conformal gravity and every field theory must postulate — appears here as an output, the commutant structure of the substrate’s internal components (Chapter 5); the phase conventions of quantum mechanics are among its gauge directions; the uniform vacuum-energy offset is another (Chapter 7 §7.6, via the reconstruction theorem’s energy-shift freedom), which is why the cosmological constant dissolution above is a *symmetry statement* in both programs — conformal symmetry forbids a fundamental Λ and must then manage the induced one, while here the offset is gauge outright and there is nothing to manage; and the boundary between what the framework can derive and what remains bijection-specific is drawn by $\mathcal{G}_{\text{sub}}$ -invariance itself, with the invariant list (the Cabibbo magnitude, the Koide ratio, the $1/4$ coefficient, the acceleration scale $a_0 = cH/6$) constituting the substrate-computable record. One principle thus answers, within the framework’s premises, questions that span quantum structure, the origin of gauge symmetry, the vacuum energy, and the program’s own epistemology — a breadth of claimed unification that conformal gravity’s narrower, older principle does not attempt. The calibration that must accompany the comparison: breadth of *claimed* unification is not depth of *tested* unification. Conformal gravity’s narrower claims have decades of scrutiny and known wounds; this framework’s broader claims are unreviewed, and their breadth is a promissory asset that external review and the framework’s own gate computations will either convert or void.

9.7 The Danielson-Satishchandran-Wald external convergence

The framework’s core mechanism — horizons as causal partitions producing decoherence in the visible sector — has received independent convergence from mainstream physics literature on horizon decoherence. The Danielson-Satishchandran-Wald (DSW) program (Danielson, Satishchandran, and Wald 2022 in *Physical Review D*; 2023 in *International Journal of Modern Physics D*) proves within standard quantum field theory in curved spacetime that Killing horizons impart fundamental decoherence on nearby quantum superpositions, with rate determined by horizon geometry. This is *direct mainstream convergence* with the framework’s content — the framework’s mechanism for emergent quantum mechanics is being independently derived by physicists operating in standard QFT-in-curved-spacetime without any commitment to the framework’s substratum picture.

The DSW result. DSW consider a quantum superposition near a Killing horizon (a black hole horizon or a cosmological horizon) and compute the decoherence rate of the superposition due to soft graviton emission across the horizon. The result is striking: the decoherence rate is *fundamental* — it does not depend on any environmental degrees of freedom inside the horizon, and it does not vanish in the limit of zero temperature. The horizon itself acts as a decoherence channel, with the rate determined by the horizon’s geometry (specifically, the horizon area and the horizon’s surface gravity).

The DSW derivation runs through three steps. First, they compute the entanglement between a near-horizon quantum superposition and the soft graviton modes produced by gravitational radiation across the horizon. Second, they show that the soft modes propagate through the horizon and cannot be recovered by any operation outside the horizon. Third, they identify the resulting irreversible loss of coherence as fundamental decoherence imposed by the horizon’s causal structure.

The result has been confirmed and extended by independent groups — Gralla and Wei (2024), Wilson-Gerow, Dugad, and Chen (2024), Chen and Penington (2024), and others — establishing the DSW decoherence as a robust prediction of standard QFT in curved spacetime, independent of any specific framework for quantum gravity.

Convergence with the framework’s mechanism. The framework’s content from Chapter 1 — embedded observers see their reduced description as quantum-mechanical because the cosmological horizon imposes a partition with C1-C3 structure — has three specific features that match DSW’s analysis.

Decoherence depends on horizon geometry. The framework’s emergent quantum mechanics has decoherence rates determined by the cosmological horizon’s structural properties — area, surface gravity, hidden-sector capacity. DSW reproduces this with horizon area and surface gravity as the geometric inputs.

The two derivations converge on the geometric dependence: horizons impose decoherence, and the rate is determined by their geometry.

Decoherence vanishes when coupling is screened. The framework’s C1 condition requires non-trivial coupling between visible and hidden sectors. If C1 fails (the coupling is screened or vanishingly small), the framework’s emergent dynamics is classical stochastic rather than quantum. DSW reproduces this with extremal black holes: when the horizon’s surface gravity vanishes (the extremal limit), the soft graviton coupling to the horizon also vanishes, and the DSW decoherence rate goes to zero. The match is structural: both frameworks predict that extremal horizons impose no fundamental decoherence.

Decoherence is non-Markovian near the horizon. The framework’s emergent dynamics is non-Markovian, with memory effects from the hidden sector’s slow timescale τ_B . DSW reproduces this with their analysis of the soft mode propagation across the horizon: the soft modes carry temporal correlations that produce memory effects in the near-horizon observer’s reduced description. The match is non-trivial: both frameworks predict that horizon-induced decoherence is structurally distinct from the Markovian decoherence of standard open quantum systems.

The convergence is across three independent structural features. The framework’s mechanism for emergent quantum mechanics is being independently derived from standard QFT in curved spacetime, with the structural features (geometry-dependence, coupling-screening, non-Markovianity) matching across the two derivations.

Where the framework extends DSW. The framework’s content extends beyond DSW’s analysis in one structurally important direction: the capacity-controlled regime where the framework’s C3 condition is marginally satisfied. DSW’s analysis assumes the horizon has effectively infinite capacity (the soft graviton modes can absorb arbitrary information without saturating). The framework’s C3 condition — $N_H \gg N_V$ — can be marginally satisfied for engineered partitions (Chapter 15) and for biological partitions (Chapter 13), where the hidden-sector capacity is finite and the decoherence dynamics deviates from the universal DSW result.

The framework’s prediction is that capacity-controlled systems exhibit modified decoherence: when N_H/N_V is finite, the framework’s P-indivisibility measure $\mathcal{N}$ takes intermediate values between fully Markovian and fully P-indivisible, producing the partially-quantum regime developed in Chapter 18 §18.3. DSW’s analysis does not extend to this regime; the framework’s content is therefore *additional* to DSW’s, with experimental access through the BEC analogue-gravity experiment developed in Chapter 15.

The BEC analogue-gravity setup of Chapter 15 is precisely a test of the framework’s capacity-controlled regime: an analogue horizon engineered to have controllable hidden-sector capacity, with the framework predicting the specific non-Markovianity signature $\mathcal{F}(r) = 2r - r^2$ as the capacity parameter r is varied.

This is the framework’s empirical access point for DSW-extension content — testing the framework’s prediction in a regime DSW does not address.

Significance for the framework’s epistemic position. The DSW convergence is significant because it provides *external validation* of the framework’s core mechanism by mainstream physics literature operating without commitment to the framework’s substratum picture. The framework’s commitment to horizons as causal partitions producing decoherence is not an idiosyncratic claim; it is being independently derived from standard QFT in curved spacetime, with the structural features matching across derivations.

The framework’s distinctive content — the substratum picture, the specific cubic-lattice realization, the Tier 3 predictions — is not external-validated by DSW. But the framework’s universality-class commitments — that horizons produce decoherence, with rate determined by geometry, with coupling-screening eliminating the effect, with non-Markovian structure — are external-validated. The framework therefore stands on stronger empirical ground than its specific substratum content alone provides: the universality-class content is converged-upon by mainstream physics, with the framework being one specific realization within that universality class.

9.8 Adjacent observer-emergent literature

Beyond the DSW convergence on horizon decoherence, the framework’s content engages a broader recent literature on observer-emergent physics. Three programs in particular have substantive overlap with the framework’s substratum-emergent operator distinction, each addressing related questions from different theoretical starting points.

The Chandrasekaran-Longo-Penington-Witten (CLPW) construction. Chandrasekaran, Longo, Penington, and Witten (2022) constructed an observer-relative algebra of operators in de Sitter space using crossed-product techniques from operator algebra theory. The construction identifies the observer’s operator algebra as a *Type II* von Neumann factor (rather than Type III, the algebra of operators in standard quantum field theory), with finite von Neumann entropy and well-defined density matrices. The construction is observer-relative: different observers in de Sitter space have different algebras, with the algebras related by appropriate isomorphisms.

The CLPW construction converges with the framework’s substratum-emergent operator distinction. The framework’s visible-sector observables form a finite-dimensional Hilbert space at any given time (the emergent description has finite information content per cell of the partition), with the observer’s accessible operator algebra being precisely the algebra of observables on this space. The framework’s reduced description is observer-relative in the same sense as CLPW’s construction: different observers correspond to different partitions, with different reduced descriptions.

The frameworks differ in their substratum commitments. CLPW operates within standard QFT in curved spacetime, with the operator algebra being the conventional structure of QFT modified by the crossed-product construction. The framework operates with a finite deterministic substratum, with the operator algebra being the emergent description of the visible sector after the trace-out. The two frameworks therefore converge on observer-relativity and finite von Neumann entropy from different starting points.

Maldacena on observer-essentiality. Maldacena (2021, 2024) has developed the position that observers are essential to a complete formulation of quantum gravity — that the wave function of the universe cannot be defined without reference to a specific observer, and that different observers correspond to different formulations of the quantum-gravity problem. This is the *observer-essentiality* thesis, formulated within string theory and AdS/CFT but with broader implications.

The observer-essentiality thesis converges with the framework’s content. The framework derives that quantum mechanics is the emergent description of an embedded observer — there is no observer-independent quantum mechanics. Different observers correspond to different partitions of the substratum, with different reduced descriptions, and the universal wave function of the universe is not the right level of description for any embedded observer.

The frameworks differ in the source of the observer-relativity. Maldacena’s framework places it at the level of quantum gravity’s holographic principle; the framework’s content places it at the level of the substratum’s partition structure. The two frameworks converge on the conclusion (observers are essential) from different starting points.

Harlow-Usatyuk-Zhao on observer dependence. Harlow, Usatyuk, and Zhao (2024) developed observer-dependence in quantum gravity through analysis of the entanglement wedge reconstruction and the gauge structure of the Hilbert space. Their result: the operator algebra accessible to an observer depends not just on the spatial region the observer occupies but on the observer’s full causal structure, with different observers having structurally different algebras even when their spatial regions overlap.

The HUZ result converges with the framework’s content on the partition-relativity of the reduced description. The framework’s visible sector is determined by the observer’s causal structure (the cosmological horizon’s location relative to the observer), with different observers having structurally different partitions. The HUZ result on operator-algebra differences between observers matches the framework’s content on reduced-description differences between observers.

Slagle-Preskill experimental program. Slagle and Preskill (2023) developed an experimental program for measuring entanglement structure in gravitational systems through interferometry. The program proposes specific experiments where a quantum superposition is gravitationally separated, then

recombined, with the resulting interference pattern probing the gravitational decoherence rate. The Slagle-Preskill program is *experimental* (testable with current and near-future technology), while DSW is *theoretical* (a result within standard QFT in curved spacetime).

The Slagle-Preskill program provides a direct experimental test of the framework’s content. The framework predicts gravitational decoherence rates determined by the partition’s geometric structure; the Slagle-Preskill experimental program measures gravitational decoherence rates through interferometry. The two programs are complementary: the framework provides specific structural predictions, and the experimental program provides the empirical access.

The cumulative external convergence. The DSW, CLPW, Maldacena, HUZ, and Slagle-Preskill programs together constitute substantial external convergence with the framework’s content. Each program is operating with different starting commitments (DSW with standard QFT, CLPW with operator algebra theory, Maldacena with string theory and holography, HUZ with AdS/CFT entanglement wedge reconstruction, Slagle-Preskill with experimental interferometry), and each converges with specific features of the framework’s substratum-emergent operator distinction.

The framework’s content is therefore *not* idiosyncratic but lies in the area where mainstream physics is converging. The specific framework-level content (substratum picture, cubic-lattice realization, Tier 3 predictions) is distinctive; the universality-class content (observers are essential, partitions are causally determined, horizons impose decoherence, reduced descriptions are observer-relative) is converged-upon by multiple independent programs in mainstream physics.

9.9 Chapter close: the framework’s position in the broader landscape

The framework’s content across Chapters 5 through 8 establishes a substantial empirical record at high precision: twenty-two classified Standard Model retrodictions (seven parameter-free structural), thirteen gravitational-sector entries, seven neutrino-sector entries. This chapter has located that empirical content within the framework’s hierarchical structural realism (§9.2), traced the universality-class character of the strongest empirical predictions (§9.3-9.4), articulated the framework’s relationship to string theory and the string landscape multiplicity problem (§9.5-9.6), and developed the substantial external convergence with mainstream physics literature on horizon decoherence (§9.7) and observer-emergent physics (§9.8).

The framework’s epistemic position. The framework’s content is at three levels simultaneously: a *universality-class commitment* (the framework belongs to a class of substratum theories with C1-C3 structure at the cosmological horizon), a *substratum-properties commitment* (the framework’s substratum is a cubic lattice at $\epsilon = 2l_p$ with the six structural properties of Chapter 5), and a *specific-bijection commitment* (the framework’s bijection produces the Tier 3

predictions of Chapter 6 and Chapter 8). Each level is tested by different empirical evidence, with the cumulative record being strongest at the universality-class and substratum-properties levels.

This epistemic position distinguishes the framework from frameworks operating at a single level. Frameworks aiming at universality-class content alone (effective field theories, top-down structural arguments) lack the empirical specificity of the framework’s Tier 3 predictions. Frameworks aiming at substratum-specific content alone (specific string compactifications, specific GUT models) lack the universality-class robustness of the framework’s Tier 1 results. The framework’s stratified empirical record provides constraint at each level, with the levels reinforcing rather than substituting for each other.

External convergence as evidential support. The framework’s universality-class commitments — observers as essential, horizons as causal partitions producing decoherence, reduced descriptions as observer-relative, finite entanglement entropy — are converged-upon by multiple independent mainstream programs: DSW on horizon decoherence, CLPW on observer-relative operator algebras, Maldacena on observer-essentiality, HUZ on observer-dependent operator structure, Slagle-Preskill on experimental gravitational decoherence. The framework’s content at the universality-class level is therefore externally validated by independent derivations from different starting commitments.

The framework’s distinctive content — the substratum picture, the cubic-lattice realization, the specific Tier 3 predictions — is not externally validated by these convergences. It is empirically validated by the framework’s own predictions matching observation: the JUNO $\sin^2 \theta_{12}$ at 0.07σ , the Cabibbo angle at 0.04% , the Bekenstein-Hawking $1/4$ at 99.999% , the Milky Way Keplerian decline at the predicted radius. The framework’s strongest case combines external convergence on universality-class commitments with internal empirical validation on framework-specific predictions.

Where the framework stands relative to alternatives. The framework’s stance toward alternative unification programs is non-competitive at the universality-class level and distinctive at the substratum-specific level. String theory, loop quantum gravity, causal dynamical triangulations, asymptotic safety, and other programs may produce identical universality-class results in their respective limits, with the framework being one structural realization that produces the same results from a different starting point. The frameworks become distinguishable only at the substratum-specific level — where the framework’s cubic-lattice realization produces specific Tier 3 predictions that competing programs do not match without additional substratum-specific assumptions.

This stance is intellectually irenic. The framework is not a competitor to existing unification programs but a structural framework that overlaps with them at the universality-class level while providing distinctive substratum-specific con-

tent. Where the framework’s content matches mainstream physics convergence (DSW, CLPW, Maldacena), the framework’s content is reinforced rather than refuted by that convergence. Where the framework’s content goes beyond mainstream convergence (the BEC analogue-gravity test, the Tier 3 predictions), the framework provides distinctive empirical targets.

Forward pointers. The remaining chapters develop the framework’s reach beyond fundamental physics. Chapter 10 develops the framework’s content in chemistry: the cascade from substratum to chemistry through the framework’s emergent description, with specific quantitative content for chemical bonding and reactivity. Chapter 11 develops the origin of life: how the framework’s structural commitments at the chemical level produce a cascade leading to the emergence of self-replicating systems. Chapter 12 develops evolution, intelligence, and self-reference. Chapters 13 through 17 develop the framework’s applications in quantum biology, quantum computing, quantum engineering, medicine, and bioinformatics. The cumulative reach of the framework across these chapters — fundamental physics in Part II, the emergence cascade in Part III, applications in Part IV — provides the framework’s complete empirical and theoretical record.

The framework’s content from Chapter 5 through Chapter 8 — concentrated at the universality-class, substratum-properties, and specific-bijection levels in this chapter’s hierarchical structural realism — provides the empirical foundation. The remaining chapters build on that foundation, with the framework’s reach extending from the substratum-level structural commitments developed in Part I and the empirical record developed in Part II into the broader scientific domains where the framework operates.

Chapter 10

Chemistry

10.1 What this chapter develops

The framework’s empirical content in Chapters 5 through 9 concentrates in fundamental physics. The chapters develop the framework’s derivation of the Standard Model gauge group, the matter content, the gravitational sector, the JUNO neutrino prediction, the dark sector phenomenology, and the framework’s position in the broader theoretical landscape. The empirical record is concentrated in the sectors of fundamental physics where the framework makes specific quantitative predictions.

This chapter develops the framework’s content in *chemistry* — the emergence

cascade from the substratum to molecular physics, with the framework deriving chemistry as a structural consequence of the substratum’s commitments rather than as an independent layer of physics. The chapter’s content is concentrated at Tier 1 and Tier 2 of the framework’s hierarchical structural realism: chemistry emerges from properties of the substratum that are shared across the framework’s universality class, with the specific chemical species and reactivity patterns following from the framework’s specific cubic-lattice realization.

Seven pieces occupy the chapter.

Section 10.2 develops the structural chain from substratum to orbital chemistry: how the framework’s commitments to three spatial dimensions, spin-1/2 fermions, and the Pauli exclusion principle produce the periodic table, the s/p/d/f orbital structure, and the specific bonding geometries of carbon (tetrahedral sp^3 , planar sp^2). The structural chain is a sequence of theorems: each step follows from the previous, with no parameter values entering the derivation. The possibility of organic chemistry is established as a theorem about embedded observation.

Section 10.3 develops four additional structural results: matter-antimatter asymmetry (three generations + CKM CP violation), the mass hierarchy (Higgs + taste-breaking producing light fermions), nuclear binding (SU(3) confinement producing bound hadrons and nuclei), and chiral chemistry (electroweak parity violation propagating to molecular parity-violating energy differences). Each result is a structural consequence of the framework’s commitments at the substratum level, propagated through the emergence cascade to chemistry.

Section 10.4 develops the scale hierarchy that makes chemistry possible: the framework’s two independent gauge groups (SU(3) and U(1)) with independent coupling strengths producing $E_{\text{nuclear}} \sim \text{MeV} \gg E_{\text{atomic}} \sim \text{eV}$. This separation means nuclei are stable against chemical reactions — chemistry rearranges electrons without disturbing nuclei. Section 10.5 develops the thermal window: the parametric conditions under which chemistry is both stable (bonds survive thermal fluctuations) and dynamic (reactions can proceed). The framework’s scale hierarchy guarantees the existence of such a window, with the specific window’s location depending on parameter values.

Section 10.6 develops water as the unique biological solvent: how oxygen’s element-8 orbital structure produces water’s hydrogen-bonding network and the anomalous properties (high heat capacity, density maximum, universal solvent behavior, hydrophobic effect) that are load-bearing for biological function. Section 10.7 develops the chirality chain: how electroweak parity violation at the substratum level propagates through molecular physics to produce the parity-violating energy difference between mirror-image molecules. The framework predicts a structural preference for L-amino acids in biology — the homochirality of life is a consequence of the framework’s chirality from trace-out (Chapter 5), not an arbitrary historical contingency. Section 10.8 develops the framework’s content on fine-tuning and the anthropic principle: the dissolution of both as

conventional formulations, with the framework’s reading being that the *possibility* of a chemistry-supporting universe is structural while specific parameter values determine particular realizations within the structural class.

The framework’s content in chemistry is *not* a competitor to standard chemistry — the predictions of quantum chemistry, organic chemistry, and biochemistry are all reproduced by the framework’s emergent description. What the framework adds is a structural derivation of why chemistry has the architecture it does: why three-dimensional space, why sp^3 and sp^2 hybridization, why the periodic table, why two independent energy scales (nuclear and atomic), why a thermal window for liquid-phase chemistry, why molecular homochirality. The chapter’s value is in providing structural foundations for chemistry that standard accounts treat as empirical input.

10.2 The structural chain from substratum to orbital chemistry

The framework’s emergence cascade from the substratum to chemistry runs through five structural steps. Each step is either a theorem from earlier chapters or a standard textbook result applied to the derived structure. The cumulative chain produces the periodic table and organic chemistry as theorems about embedded observation rather than as empirical inputs.

Step 1: Three spatial dimensions. The framework’s argument for $d = 3$ is the three-filter convergence of Chapter 5: propagating gravity excludes $d \leq 2$, stable matter — and with it the embedded observers of Axiom 1 — excludes $d \geq 4$, and the boundary-entropy concordance checks $d = 3$ exactly. (The coupling-degree minimization $K = 2d$ and the Bravais-lattice uniqueness theorem then operate *within* $d = 3$, fixing the component count and the lattice type; they do not themselves fix the dimension.) The papers carry the complementary, more conservative anchoring: the cubic-lattice commitment is also empirically selected by the observed gauge group.

The dimensional commitment is *load-bearing for chemistry*. In $d = 2$, the angular momentum algebra reduces — only circular harmonics exist, with no three-dimensional bonding geometry. In $d \geq 4$, the Coulomb potential $V(r) \sim r^{2-d}$ does not support stable bound states: classical orbits spiral into the nucleus, and the quantum Hamiltonian is unbounded below. The framework’s derivation of $d = 3$ is therefore the framework’s derivation of the dimensional condition that makes chemistry possible.

Step 2: Orbital algebra in $d = 3$. Standard quantum mechanics in $d = 3$ produces the $SO(3)$ angular momentum algebra with the familiar quantum numbers (n, ℓ, m_ℓ) . The orbital angular momentum quantum number ℓ takes values $0, 1, 2, \dots, n-1$ for principal quantum number n . The magnetic quantum number m_ℓ takes values $-\ell, -\ell+1, \dots, +\ell$, giving $2\ell+1$ orbital states per ℓ value.

The orbital labels translate to the conventional spectroscopic notation: s ($\ell = 0$, 1 orbital), p ($\ell = 1$, 3 orbitals), d ($\ell = 2$, 5 orbitals), f ($\ell = 3$, 7 orbitals). This is the *only* orbital algebra consistent with embedded observation. The framework’s commitment to $d = 3$ at the substratum level produces the $SO(3)$ angular momentum structure at the emergent level, with no alternative orbital structure possible.

Step 3: Spin-1/2 and the Pauli principle. The framework derives spin-1/2 fermions from the staggered structure on the $d = 3$ cubic lattice (Chapter 5 §5.6). The staggered fermion construction places independent fermion modes at alternating lattice sites; the chiral structure of the framework’s emergent QFT (Chapter 5 §5.6) requires the modes to carry half-integer spin. The spin-statistics theorem — a consequence of the emergent QFT’s Lorentz invariance, established with the emergent QFT structure of Chapter 5 §5.4 — gives the Pauli exclusion principle: at most one fermion per state, hence at most 2 electrons per orbital (spin up plus spin down).

The shell capacities follow: s (2 electrons), p (6 electrons), d (10 electrons), f (14 electrons). The periodic table’s period lengths — 2, 8, 8, 18, 18, 32 — are determined by the order in which these shells fill, with the filling order following from energy ordering of the orbitals.

The architecture is *entirely structural*: it depends only on $d = 3$, spin-1/2 fermions, and the Pauli principle, all of which are derived earlier in the framework’s emergence cascade. No empirical inputs enter this derivation.

Step 4: Carbon’s bonding geometry. Element 6 (carbon) has the electron configuration $1s^2 2s^2 2p^2$ — four valence electrons distributed among $2s$ and $2p$ orbitals. In $d = 3$ with the derived orbital algebra, four valence electrons admit three hybridization modes that mix the s and p orbitals.

sp³ hybridization. The s orbital mixes with all three p orbitals to produce four equivalent orbitals pointing to the vertices of a regular tetrahedron, with bond angles of 109.5°. This is the geometry of methane (CH_4), diamond, and the backbone of all organic macromolecules — proteins, nucleic acids, lipids, carbohydrates.

sp² hybridization. The s orbital mixes with two of the three p orbitals, producing three planar orbitals at 120°, with one p orbital remaining perpendicular to the plane. The perpendicular p orbital forms a π bond. This is the geometry of graphite, benzene, and aromatic chemistry — the planar-ring chemistry that underlies nucleic acids’ bases and aromatic amino acids.

sp hybridization. The s orbital mixes with one p orbital, producing two linear orbitals at 180° plus two perpendicular p orbitals forming π bonds. This is the geometry of acetylene ($HC\equiv CH$) and other triple-bonded systems.

The *possibility* of these hybridizations — and hence the possibility of three-dimensional macromolecular architecture — is a structural consequence of the

angular momentum algebra in $d = 3$ combined with carbon's four valence electrons. No empirical input determines that carbon should be capable of sp^3 hybridization; the structural derivation of $d = 3$ combined with carbon's electronic structure produces it.

Why carbon specifically. The framework's structural derivation singles out carbon among the periodic-table elements. Carbon's position is determined by three structural features: four valence electrons (in the row where s and p shells are half-filled), $2p$ orbital size matching $2s$ orbital size (allowing the hybridizations), and second-period status (avoiding the d -orbital availability of heavier elements that complicates bonding). These features combine to give carbon the unique ability to form four equivalent strong covalent bonds with itself and with hydrogen, nitrogen, oxygen, sulfur, and other light elements.

The framework's content here addresses a question that standard chemistry treats as empirical: why is carbon the basis of organic chemistry and biology rather than silicon (Si, four valence electrons in the third period) or germanium (Ge, four valence electrons in the fourth period)? The structural answer involves orbital size matching. In silicon, the $3s$ and $3p$ orbitals have different radial extents (the $3s$ orbital is more compact than the $3p$ due to penetration effects), making sp^3 hybridization less symmetric and silicon-silicon bonds substantially weaker than carbon-carbon bonds (Si-Si bond energy ~ 226 kJ/mol versus C-C bond energy ~ 347 kJ/mol). The framework's commitment to the standard atomic-orbital structure makes carbon's combinatorial chemistry possible at room temperature in ways silicon's chemistry cannot match.

The framework's prediction is therefore: any chemistry-supporting universe with the framework's structural commitments will use carbon (or its precise analog at the same orbital-structure position) as its primary combinatorial element. The "carbon chauvinism" of biology is structurally predicted rather than empirically contingent.

The hybridizations as load-bearing for biology. The three carbon hybridizations correspond to three structural classes of biological molecules.

sp^3 *carbons* form the saturated backbones of carbohydrates, lipids, and the side chains of most amino acids. The tetrahedral sp^3 geometry enables three-dimensional folding of proteins (the backbone ϕ and ψ dihedral angles depend on tetrahedral sp^3 geometry at each α -carbon).

sp^2 *carbons* form the planar aromatic rings of nucleic acid bases (adenine, guanine, cytosine, thymine, uracil), the aromatic amino acids (phenylalanine, tyrosine, tryptophan), and the peptide bond's planar carbonyl carbon. The peptide bond's sp^2 character is what produces the rigid planar structure that constrains protein secondary structure into α -helices and β -sheets.

sp *carbons* are rarer in biology but appear in some natural products and as transition states in enzyme catalysis. Their linear geometry produces the rigid axes needed for some specialized binding interactions.

The framework’s structural content is that the *possibility* of these three hybridization modes — and hence the possibility of proteins with three-dimensional folds, nucleic acids with planar base-pairing, and lipid membranes with sp^3 hydrocarbon chains — is a theorem about embedded observation rather than a contingent feature of our universe. The biological consequences develop in Chapters 11 and 12.

Step 5: The structural chain. Combining the five steps gives the framework’s chemistry derivation chain in compact form:

$$(S, \varphi) \xrightarrow{\text{derived}} d = 3 \xrightarrow{\text{SO}(3)} \text{orbital algebra} \xrightarrow{\text{Pauli}} \text{periodic table} \xrightarrow{\text{element 6}} \text{tetrahedral carbon}.$$

Every arrow is either a theorem from earlier chapters of this book or a standard textbook result applied to the derived structure. No parameter values enter. The cumulative chain establishes that the *possibility* of organic chemistry is a theorem about embedded observation: any universe with the framework’s substratum structure produces the periodic table and tetrahedral carbon as structural consequences.

The empirical confirmation of this chain is *concordance*: the framework’s derivation reproduces the periodic table and the geometry of carbon’s bonding, with no deviations from standard chemistry at currently accessible energy scales. The framework does not add new chemistry; it provides structural foundations for the chemistry that has been empirically established over the past century. The value of the framework’s content here is the structural account of why chemistry has the architecture it does.

10.3 Beyond carbon: four additional structural results

The framework’s chemistry derivation extends beyond carbon’s bonding geometry to four additional structural results that propagate the substratum’s commitments through to chemistry-relevant features.

Matter-antimatter asymmetry. The framework’s commitment to three matter generations (Chapter 5 §5.7) produces a CKM mixing matrix with one physical CP-violating phase. CP violation is one of three Sakharov conditions for baryogenesis: the matter-antimatter asymmetry of the observed universe requires (i) baryon number violation, (ii) C and CP violation, (iii) out-of-equilibrium dynamics. The framework’s partition breaks discrete spatial parity P (Chapter 5 §5.6, Theorem 13), and the CKM phase breaks CP. Combined with out-of-equilibrium dynamics at the electroweak phase transition, baryogenesis becomes *structurally possible* — the framework’s structural commitments at the substratum level produce the necessary ingredients for the matter-antimatter asymmetry that makes chemistry possible (without baryon-antibaryon symmetry breaking, the universe would have annihilated to photons rather than producing baryonic matter).

The framework’s content on baryogenesis is *structural* — the necessary ingredients are derived — rather than *quantitative*. The specific baryon-to-photon ratio $\eta = 6.1 \times 10^{-10}$ is not derived from the framework’s commitments; the quantitative content remains open (Chapter 19 §19.3.3). What the framework provides is the *possibility* of baryogenesis as a theorem: any universe with the framework’s three-generation structure and partition breaking P has the ingredients to produce a matter-antimatter asymmetry.

Mass hierarchy. The Higgs mechanism (Chapter 6 §6.6) gives fermion masses through Yukawa couplings to a single doublet. The taste-breaking mechanism on the staggered cubic lattice (Chapter 6 §6.7) produces generation-dependent Yukawa couplings generically, ensuring that different generations have different masses. The hierarchy is therefore structural: light fermions (electron, u and d quarks) exist not as coincidence but as a structural feature of the staggered lattice’s symmetry-breaking.

The existence of light fermions is *load-bearing for chemistry*. Heavy fermions cannot form stable atoms — their Compton wavelengths are smaller than nuclear radii, producing relativistic dynamics that destabilize the atomic structure. The electron’s specific mass ($m_e \approx 511$ keV) sits in the regime where atomic physics is non-relativistic and bound states are stable; the up and down quarks’ specific masses give nuclear binding energies in the MeV range that produce stable nuclei. The mass hierarchy is therefore not an arbitrary parametric feature but a structural prediction: any universe with the framework’s commitments must have a hierarchy of fermion masses, with the lightest fermions in the regime that produces stable atomic and nuclear physics.

Nuclear binding. SU(3) color confinement (Chapter 5 §5.4) produces bound hadrons — quarks confined into mesons and baryons by the color gauge force. The existence of multi-nucleon bound states (nuclei) requires specific parameter values (the QCD scale Λ_{QCD} in the right range, the quark masses producing pions light enough to mediate the long-range nuclear force), but the *mechanism* — color confinement binding quarks into hadrons, residual strong force binding hadrons into nuclei — is structural.

The framework’s content on nuclear binding is therefore *mechanism* structural with *parameters* solution-specific. The mechanism is derived from the gauge structure of SU(3); the specific values that make nuclear binding successful are properties of the bijection (Tier 3 in Chapter 9’s hierarchy). The empirical record of stable nuclei is consistent with the framework’s structural commitments; the specific binding energies require the full Tier 3 calculation.

Chiral chemistry. Chiral SU(2) (Chapter 5 §5.6) produces electroweak parity violation in all weak processes — the left-handed structure of weak interactions discovered by Wu (1957). This propagates to molecular physics as a parity-violating energy difference between mirror-image molecules: an L-amino acid has a slightly different energy than its D-amino acid mirror image, with the difference arising from the weak-interaction contribution to the molecular

Hamiltonian.

The parity-violating energy difference for typical amino acids is tiny — approximately 10^{-19} eV per molecule, comparable to kT at the Planck scale. But the *sign* of the difference is structural: it follows from the chirality of the framework’s weak interactions, and it predicts L-amino acids (rather than D) as the slightly more stable form. Section 10.6 develops the chirality chain in detail and the amplification mechanisms that turn this small energy difference into the homochirality of biological systems.

10.4 The scale hierarchy and the permanence of atoms

The framework derives two independent gauge groups — $SU(3)$ and $U(1)$ — with independent coupling strengths corresponding to independent eigenvalues of the substratum coupling matrix M (Chapter 5 §5.4). The framework’s universal induced coupling $1/\alpha_0 = 23.25$ at the Planck scale is itself structural — determined by the number of fermion families $N_f = 6$ and the Dynkin index, not by the eigenvalues of M . The framework reproduces all three Standard Model gauge couplings at M_Z through non-perturbative gauge self-energy corrections and renormalization-group running (Chapter 6 §6.4).

The two independent gauge couplings produce two widely separated energy scales at the chemistry-relevant level:

$$E_{\text{nuclear}} \sim \Lambda_{\text{QCD}} \sim \text{MeV}, \quad E_{\text{atomic}} \sim \alpha^2 m_e c^2 \sim \text{eV}.$$

The ratio $E_{\text{nuclear}}/E_{\text{atomic}} \sim 10^6$ is partly parametric (it depends on the specific values of the gauge couplings and the electron mass), but the *existence* of two independent scales is structural — it follows from having two independent gauge groups in the framework’s commitments.

The chemistry consequence. The scale hierarchy has a striking consequence: nuclei are *stable against chemical reactions*. Chemistry rearranges electrons (eV-scale processes) without touching nuclei (MeV-scale binding). Atoms are *permanent building blocks* that can be assembled into molecules, disassembled by chemical reactions, and reassembled into different molecules without being destroyed.

This permanence is what makes chemistry a *combinatorial* science. Every chemical species in the periodic table maintains its identity across all the chemical transformations available at room temperature. A hydrogen atom in a water molecule is the same hydrogen atom that could be in an ammonia molecule or a hydrochloric acid molecule. The framework’s scale hierarchy is what makes this combinatorial structure possible.

Without the scale hierarchy, chemistry would not be combinatorial. If the nuclear and atomic energy scales were comparable, chemical reactions would frequently destabilize nuclei, producing nuclear transmutation rather than chemical recombination. Periodic-table elements would not be stable; the periodic

table itself would not be a meaningful organizational principle. The framework's structural commitment to two independent gauge groups is therefore the structural origin of the periodic table's stability — the foundation on which all of chemistry, biochemistry, and biology rest.

Where chemistry meets the framework. The scale hierarchy is the chapter's clearest example of a Tier 2 framework prediction with empirical consequences in chemistry. The framework predicts that any universe with the structural commitment to two independent gauge groups has a scale hierarchy guaranteeing chemistry's combinatorial structure. The empirical record — the entire field of chemistry confirms the scale hierarchy through the stability of all atomic and molecular species at room temperature — is therefore a Tier 2 confirmation of the framework's structural commitments.

10.5 The thermal window for chemistry

Chemistry as we know it requires specific thermal conditions: bond energies must exceed thermal fluctuation energies (so bonds are stable), thermal fluctuations must exceed activation energies (so reactions can proceed), and a liquid solvent must exist (so reactants can encounter each other in solution). These three conditions define a *thermal window* in which chemistry is both stable and dynamic. The framework's structural commitments guarantee that such a window exists; the specific window's location is parametric.

The three conditions.

Bond stability. Chemical bond energies are in the range 1-5 eV (covalent bonds, hydrogen bonds being weaker at ~0.2 eV). For bonds to be stable against thermal disruption, we need $E_{\text{bond}} \gg kT$. At room temperature ($kT \approx 0.025$ eV), the ratio is $E_{\text{bond}}/kT \approx 40\text{-}200$ for covalent bonds — comfortably stable. For hydrogen bonds, the ratio is ~ 8 — marginally stable, which is precisely the regime where hydrogen bonds break and re-form dynamically (the basis of biological reaction kinetics).

Reaction kinetics. For chemical reactions to proceed at observable rates, the activation energy must be accessible to thermal fluctuations: $E_{\text{activation}} \lesssim 30kT$ (so the Arrhenius factor $e^{-E_a/kT}$ is not vanishingly small). Typical activation energies for catalyzed reactions are 0.1-1 eV — in the regime $E_a/kT \approx 4\text{-}40$ at room temperature, giving reaction rates appropriate for biological timescales.

Liquid solvent. Reactions in dilute solution require a liquid phase. The framework's water content from §10.2 establishes water as a candidate solvent with anomalous properties supporting chemistry (high specific heat, hydrogen bonding, ionic dissolution). The existence of a liquid range for water (0-100°C at atmospheric pressure) provides the temperature range in which chemistry operates.

The framework's structural guarantee. The scale hierarchy from §10.4 guarantees that $E_{\text{bond}}/E_{\text{thermal}} \gg 1$ at any temperature where a liquid exists.

Bond energies are set by atomic physics ($\sim$ eV scale, from the electron mass and the fine-structure constant); thermal energies at the boiling point of any liquid are typically ~ 0.05 eV (where intermolecular forces holding the liquid together become comparable to kT). The ratio is always large enough that chemistry is stable in the liquid range.

The framework's structural commitment is therefore that *some* thermal window must exist for any universe with the framework's structural conditions — the scale hierarchy guarantees the bond-stability condition, the orbital structure produces the activation-energy structure of chemistry, and water's structural properties (from carbon's analog in element 8) produce a candidate solvent. The *specific* thermal window — its location in temperature, the precise solvent (water at ~ 273 - 373 K, ammonia at lower temperatures, etc.) — is parametric.

Chemistry is therefore inherently *stable but dynamic* in any universe with two independent gauge groups and the derived orbital structure. The framework derives the structural possibility of a chemistry-supporting thermal window; the specific window's location requires the full parameter content of the framework's specific bijection.

10.6 Water as biological solvent

Water occupies a distinguished position in biology that the framework's structural derivation can account for. Every known living system uses water as its primary solvent, and several of water's anomalous properties — its high heat capacity, its density maximum at 4°C , its hydrogen-bonding network, its capacity as both proton donor and acceptor — are load-bearing for biological function. The framework's content reads water's biological centrality as structurally predicted from element-8 chemistry rather than as contingent.

Element 8 in the framework's orbital structure. Oxygen has electronic configuration $1s^2 2s^2 2p^4$ — six valence electrons distributed across the four orbitals available in the second period. The two unpaired p electrons form bonds with two hydrogen atoms; the two lone pairs remain as sp^3 -like orbitals. The result is the bent H-O-H geometry with bond angle of approximately 104.5° (close to the tetrahedral 109.5° but compressed by the lone-pair repulsion). The molecular dipole moment of approximately 1.85 D follows from the bent geometry combined with the oxygen-hydrogen electronegativity difference.

The framework's structural commitment to $d = 3$ and the orbital algebra produces this molecular structure as a derived consequence. The dipole moment, the hydrogen-bonding capacity, and the molecular geometry are all structural rather than empirical inputs.

The hydrogen-bonding network. Water's most distinctive property is its capacity to form a three-dimensional hydrogen-bonding network. Each water molecule can simultaneously donate two hydrogen bonds (through its two O-H groups) and accept two hydrogen bonds (through its two lone pairs). The

tetrahedral geometry of these donor and acceptor positions produces the open tetrahedral network found in ice and the partially-ordered network in liquid water.

The structural consequences of this network for biology are several.

Anomalous density maximum. Liquid water has its maximum density at 4°C, not at the freezing point. This is because the open tetrahedral hydrogen-bonded network in ice is less dense than the partially-collapsed network in liquid water. Ice floats on water — a structural feature that allows aquatic ecosystems to persist through freezing surface conditions, with liquid water beneath insulating ice cover.

High specific heat. Water has approximately 4.18 J/(g · K) heat capacity — roughly four times the heat capacity of typical liquids on a per-mass basis. The energy required to disrupt the hydrogen-bonding network as temperature increases produces the high heat capacity. This thermal buffering is load-bearing for biological systems: cellular temperatures remain stable against substantial heat input or loss, allowing biochemistry to operate in narrow temperature windows.

Universal solvent behavior. The dipole moment and hydrogen-bonding capacity make water an effective solvent for both ionic compounds (through ion-dipole interactions) and polar molecules (through hydrogen bonding to functional groups). This dual capacity is what allows water to dissolve simultaneously the ionic species essential for biology (Na⁺, K⁺, Cl⁻, Ca²⁺, Mg²⁺) and the polar organic molecules (amino acids, sugars, nucleotides, ATP).

Hydrophobic effect. Non-polar molecules disrupt the water hydrogen-bonding network without participating in it; this is energetically unfavorable, producing the hydrophobic effect that drives protein folding, lipid membrane formation, and substrate binding in enzyme active sites. The hydrophobic effect is load-bearing for biological organization: proteins fold into compact tertiary structures with hydrophobic cores, lipid bilayers self-assemble into membranes with hydrocarbon interiors, and substrate binding is often driven by burial of hydrophobic surfaces.

The structural prediction for biological solvents. The framework's content here is that water's biological centrality is *structurally predicted* by the framework's commitments to $d = 3$, the orbital structure, and the specific properties of element 8. The framework predicts that any chemistry-supporting universe with these commitments will produce a small dipolar hydrogen-bonded liquid (water or its precise analog) with the same biology-supporting properties.

Alternative solvents (ammonia, hydrogen fluoride, methanol) have been considered as candidate biological solvents in the astrobiology literature. Each has properties similar to water in some respects (hydrogen bonding) but lacks others (the specific tetrahedral hydrogen-bonding network, the dipole moment magnitude, the open low-density solid phase). The framework's structural prediction

is that water’s specific combination of properties is not coincidental but follows from element 8’s specific orbital structure — making water the unique liquid that simultaneously satisfies all the conditions for biological function.

This is one of the framework’s clearest examples of a structural prediction *propagating from the substratum to biology*: the cubic-lattice substratum commits to $d = 3$ orbital algebra, the orbital algebra commits to the periodic table, the periodic table places oxygen at element 8 with the specific valence structure, and the valence structure commits to water’s hydrogen-bonding network with all its biology-supporting properties. The framework’s reach from the substratum to liquid water is a single derivation chain with no empirical inputs.

10.7 The chirality chain

The framework’s commitment to chiral SU(2) (Chapter 5 §5.6) produces electroweak parity violation. This commitment propagates from the substratum through atomic and molecular physics to produce a structural preference for one chirality over the other in molecular systems — the parity-violating energy difference (PVED). The framework predicts the homochirality of biological systems as a consequence of this chirality chain rather than as an arbitrary historical contingency.

Step 1: Parity violation at the substratum. Chapter 5 §5.6 derives chiral SU(2) from the staggered structure on the cubic lattice combined with the bipartite trace-out. The weak gauge force couples only to left-handed fermions; right-handed fermions are inert under SU(2). This is the parity violation discovered by Wu in 1957, with all weak processes (beta decay, muon decay, neutrino interactions) showing definite chirality.

Step 2: Molecular parity violation. The weak interaction contributes to the electromagnetic-binding Hamiltonian of any atom or molecule with chiral structure. For a molecule with one or more chiral centers (carbon atoms bonded to four different substituents, producing a non-superimposable mirror image), the weak contribution to the molecular energy depends on which of the two mirror-image enantiomers is considered. The result is a parity-violating energy difference:

$$\Delta E_{\text{PV}} = E_{\text{L}} - E_{\text{D}},$$

where E_{L} and E_{D} are the energies of the L (left-handed) and D (right-handed) enantiomers. For typical amino acids, the magnitude is approximately $|\Delta E_{\text{PV}}| \sim 10^{-19}$ eV per molecule.

The sign of ΔE_{PV} is structural: the framework’s specific chirality from trace-out determines whether L or D is the lower-energy form. For amino acids, the L form is predicted to be slightly more stable; for sugars, the D form is predicted to be slightly more stable. These predictions match the observed homochirality of biological systems: all amino acids in proteins are L (with rare D exceptions in cell walls), and all sugars in DNA and RNA backbones are D.

Step 3: The amplification argument. The parity-violating energy difference of 10^{-19} eV is tiny — comparable to kT at the Planck temperature, not at any biological temperature. Direct thermodynamic selection of L over D at biological temperatures is therefore negligible: the equilibrium ratio $K = e^{-\Delta E_{PV}/kT} \approx 1 + \Delta E_{PV}/kT$ deviates from unity by approximately 10^{-17} at room temperature — far too small to produce the observed near-100% homochirality of biological molecules.

The framework’s content here is therefore not that the parity-violating energy difference *directly* produces homochirality, but that it provides the *initial bias* that subsequent amplification mechanisms can exploit. Three amplification mechanisms have been studied in the chemistry literature.

Frank’s autocatalytic model. Frank (1953) showed that if a molecule’s racemization is slow compared to its autocatalytic synthesis (each enantiomer catalyzes its own production), even a tiny initial bias can be amplified exponentially. The dynamics produces a winner-take-all transition where one enantiomer dominates the steady state. Applied to early biochemistry, Frank’s model shows that the 10^{-17} initial bias from ΔE_{PV} can be amplified to near-100% over geological timescales.

The Soai reaction. The Soai et al. (1995) autocatalytic reaction provides experimental confirmation of Frank’s mechanism: a chiral aldehyde catalyzes its own production from achiral precursors, with the chirality amplifying from near-undetectable to >99% over a few reaction cycles. The Soai reaction is the cleanest laboratory demonstration of chirality amplification with arbitrary initial bias.

Geochemical selection. Local concentrations of one enantiomer (through differential adsorption on mineral surfaces, differential degradation in UV-irradiated environments, or differential crystallization) can produce enantiomeric excess on geological timescales. Combined with Frank’s mechanism, geochemical selection provides a robust amplification path from ΔE_{PV} to homochirality.

Step 4: The structural conclusion. The framework’s content on chirality is therefore structural at two levels. At the level of *direction*, the framework predicts the *sign* of ΔE_{PV} for any chiral molecule — which enantiomer is slightly more stable, as a consequence of the framework’s specific chirality from trace-out. At the level of *amplification*, the framework predicts that this small initial bias can be amplified by autocatalytic and geochemical mechanisms to produce near-complete homochirality, with the timescale and final excess depending on parametric details of the relevant chemistry.

The observed homochirality of biological systems — all proteins built from L-amino acids, all DNA and RNA backbones using D-sugars — is therefore *structurally predicted* by the framework: the direction follows from the framework’s chirality, and the amplification mechanism is provided by well-studied chemistry. The specific historical pathway by which homochirality became established on early Earth is parametric, but the structural prediction is unambiguous.

The framework’s chirality chain is one of the strongest examples of the framework’s reach beyond fundamental physics. The substratum-level commitment to chiral $SU(2)$ propagates through atomic and molecular physics to produce a specific empirical prediction about biological systems — a prediction that is empirically confirmed by the universal homochirality of life on Earth.

10.8 Fine-tuning and the anthropic principle

The framework’s content on chemistry has implications for two interpretive questions that have been central to physics discussions for decades: the *fine-tuning* of physical parameters and the *anthropic principle*. The framework’s structural-vs-parametric distinction reframes both questions, with neither receiving its conventional formulation in the framework’s terms.

The conventional formulation of fine-tuning. The fine-tuning argument observes that several physical parameters (the fine-structure constant α , the proton-to-electron mass ratio, the strong coupling α_s , the cosmological constant, others) appear to be in narrow ranges that support chemistry, atoms, and complexity. Small changes in these parameters — typically a few percent — would produce universes incompatible with chemistry as we know it. The puzzle is why the parameters take values in this narrow range rather than generic values producing structureless universes.

The conventional responses to this puzzle include the *multiverse* (an ensemble of universes with varying parameters, with our universe being one of the small fraction supporting chemistry), the *anthropic principle* (we observe parameter values consistent with chemistry because only such universes contain observers to make observations), and various theological or teleological accounts (parameter values designed to produce chemistry-supporting universes).

The framework’s reading. The framework’s chemistry derivation distinguishes *structural* features (those derived from the framework’s commitments) from *parametric* features (those depending on the specific bijection). Most of the chemistry-supporting features cited in fine-tuning arguments are structural in the framework: the existence of three spatial dimensions, the existence of sp^3 and sp^2 hybridization for carbon, the periodic-table architecture, the scale hierarchy between nuclear and atomic physics, the existence of a thermal window for chemistry, the homochirality of biological molecules — all of these are structural consequences of the framework’s commitments at the substratum level.

The specific parametric features — the precise value of α , the precise value of α_s , the precise mass ratios — are bijection-specific (Tier 3 in Chapter 9’s hierarchy). The framework’s content is that *some* parameter values are required to produce chemistry, but the *structural class* of chemistry-supporting universes is broader than the conventional fine-tuning argument suggests. The framework’s universality class includes many bijections producing chemistry; the specific bijection of our universe is one within that class.

The dissolution of fine-tuning. Under the framework’s reading, the fine-tuning problem partly dissolves. The chemistry-supporting features that are *structural* are not fine-tuned because they follow from the framework’s commitments rather than from narrow parameter ranges. The features that are *parametric* still require specific values, but the parametric content is concentrated in the flavor sector (Chapter 19 §19.3.1) and the dark energy magnitude (Chapter 7 §7.8), not in the structural features that make chemistry possible.

The framework therefore changes what fine-tuning means. Instead of “the universe’s parameter values are in a narrow range supporting chemistry” — which would require an anthropic, multiverse, or teleological explanation — the framework’s reading is “the universe’s structural commitments produce chemistry as a theorem, with parametric details adjusting which specific chemical species and reactions occur within the structural class.” Chemistry’s possibility is theorem; chemistry’s specifics are parametric.

The dissolution of the anthropic principle. The conventional anthropic principle observes that we exist as observers in a chemistry-supporting universe, and concludes either that (a) we should expect to observe chemistry-supporting parameter values (the weak anthropic principle), or (b) the universe’s parameters are *required* to support observers (the strong anthropic principle). The principle has been controversial: it explains parameter values without producing new predictions, and its strong form veers into teleology.

The framework’s reading dissolves the conventional anthropic principle similarly. Observers (in the framework’s structural sense from Chapter 18 §18.10) are entities implementing C1-C3 partitions — a structural feature of the framework’s universality class. The framework’s *possibility* of observers follows from its universality-class commitments, not from anthropic selection across an ensemble. The framework’s *specific observers* (humans, other biological organisms) are parametric, but the existence of *some* observers is structural.

The framework therefore replaces the anthropic principle’s “we exist because the parameters support us” with a stronger structural statement: “the framework’s commitments produce observers as a structural consequence, with the specific observers depending on parametric details within the framework’s universality class.” Anthropic selection across a multiverse is not required to explain observer-supporting parameter values; the framework’s structural commitments produce observer-supporting structure as a theorem.

The framework’s content on fine-tuning and anthropics is therefore a substantive engagement with these long-standing interpretive questions. The framework neither endorses nor refutes the conventional fine-tuning argument; it provides a different reading where the chemistry-supporting features are mostly structural rather than parametric, and the existence of observers is a structural consequence of the framework’s universality class rather than an anthropic selection effect.

Forward pointers. Chapter 11 develops the framework’s content on the ori-

gin of life: how the framework’s chemistry-supporting commitments propagate further through the emergence cascade to produce self-replicating molecular systems. Chapter 12 develops evolution, intelligence, and self-reference. Chapter 13 develops the framework’s content in quantum biology, where the framework’s C1-C3 structural conditions apply directly to biological systems at the molecular scale. The framework’s content in chemistry developed in this chapter provides the structural foundation for the subsequent biology chapters: the periodic table, organic chemistry, water as solvent, the thermal window, and the homochirality of biological molecules are all the chapter’s content’s downstream consequences for life.

Chapter 11

The Origin of Life

11.1 What this chapter develops

Chapter 10 developed the framework’s content in chemistry: the structural chain from substratum to orbital chemistry, the scale hierarchy between nuclear and atomic energies, the thermal window for liquid-phase reactions, the chirality chain producing molecular homochirality, and the dissolution of conventional fine-tuning arguments. The framework’s commitments at the substratum level propagate through atomic and molecular physics to produce chemistry as a *structural consequence* — a theorem about embedded observation rather than an empirical accident.

This chapter develops the next step in the emergence cascade: the origin of life. The framework’s content here is structural rather than historical — the framework does not derive the specific molecular pathway by which life emerged on Earth, but it does derive that life-like dynamics (self-replicating, hereditary, evolvable molecular systems) is a *structural attractor* of chemistry within the framework’s thermal window. The transition from prebiotic chemistry to biology is the first place a molecular system satisfies the framework’s structural conditions C1, C2, C3 at the molecular scale.

Seven pieces occupy the chapter.

Section 11.2 develops the autocatalytic argument: in any chemistry with a sufficiently rich combinatorial space of possible molecules, self-catalyzing networks become *statistically expected* rather than accidental. Kauffman’s formal analysis gives the threshold; the framework’s chemistry derivation guarantees the threshold is exceeded by a wide margin in any aqueous organic chemistry within the framework’s thermal window.

Section 11.3 reviews the prebiotic chemistry pathways by which simple precursors are converted into the monomeric building blocks of life: the Miller-Urey synthesis, hydrothermal vent chemistry, the surface metabolism hypothesis, and concentration mechanisms. The framework's relationship to this experimental literature is structural — the framework provides the conditions under which prebiotic synthesis proceeds, while the specific pathways and yields are matters for experimental prebiotic chemistry.

Section 11.4 develops self-replication as the framework's read-write cycle at the molecular scale. The framework's emergent quantum mechanics in Chapter 1 derived the read-write cycle as the equilibrium dynamics between visible and hidden sectors of an embedded observer's partition. The same structural pattern applies at the molecular scale: a self-replicating molecular system is a chemical implementation of the read-write cycle, with the template playing the role of the visible sector and the chemical environment playing the role of the hidden sector. This connection is not analogy but identity — the framework's structural commitments at the cosmological scale and at the molecular scale are the same.

Section 11.5 develops the transition from template replication to evolution. Three capabilities — linear information-carrying polymers, complementary pairing, and catalytic copying — are each independently structural in the framework's chemistry. Their intersection is template replication. Template replication with finite fidelity produces variation; variation with finite resources produces selection; the conjunction is Darwinian evolution. The chapter establishes Darwinian evolution as a *structural consequence* of the framework's chemistry — not as an arbitrary discovery about biological systems but as the inevitable dynamics of imperfect template replication.

Section 11.6 develops C1-C3 at the molecular scale: the framework's structural conditions for non-Markovian dynamics applied to chemistry. Ordinary chemical reactions are Markovian — products do not remember how they were made. Life begins when a molecular system first satisfies all three conditions simultaneously: catalytic feedback (C1), persistence of templates beyond reaction cycles (C2), and sufficient sequence-space capacity (C3). The transition is one-way: any chemical system that accidentally satisfies C1-C3 will dominate its environment through exponential growth. The framework's characterization theorem applies at the molecular scale to produce heredity in the same way it produces quantum mechanics at the cosmological scale.

Section 11.7 develops RNA's special status as the framework's unique single-molecule C1-C3 system. RNA simultaneously satisfies all three conditions — it serves as both template (C2, C3) and catalyst (C1). DNA stores information but does not catalyze; proteins catalyze but do not easily template. The RNA world hypothesis is not historical accident but structural prediction: the simplest C1-C3 chemistry is RNA-like. Section 11.8 closes the chapter by identifying what remains genuinely open within the framework's structural account: the specific molecular pathway from prebiotic chemistry to the first replicator, the timescale of the transition, and the historical contingencies that distinguish

Earth’s biology from alternative structurally-compatible biologies.

The framework’s content in this chapter is therefore at multiple levels. At the *structural* level, the framework predicts that life-like dynamics is inevitable in any chemistry with the framework’s thermal window and combinatorial richness. At the *molecular* level, the framework predicts RNA-like architecture as the simplest realization. At the *historical* level, the framework provides no specific predictions — the path from prebiotic chemistry to the first self-replicator depends on parametric details the framework’s structural commitments do not constrain.

11.2 The autocatalytic argument

The first structural step from chemistry to life is the emergence of *autocatalytic networks*: sets of molecules in which each member catalyzes the formation of others, with the network as a whole producing all its components from environmental precursors. Stuart Kauffman’s formal analysis in the 1980s established that autocatalytic networks become *statistically expected* in any sufficiently rich chemistry — not as a lucky accident but as a structural consequence of combinatorial probability.

Kauffman’s threshold. Kauffman’s argument runs as follows. Consider a chemistry with N distinct molecular species and a probability p that any randomly chosen molecule catalyzes any randomly chosen reaction. The number of potential catalytic relationships scales as $N^2 \cdot p$ (each ordered pair of molecules has probability p of being a catalyst-substrate pair). When $N \cdot p \gtrsim 1$, the catalytic relationships form a connected graph, and self-reinforcing cycles appear with high probability.

The threshold is mild. For $p \sim 10^{-9}$ (a generous estimate for the probability that a randomly chosen organic molecule catalyzes a randomly chosen reaction — most molecules are inert toward most reactions), the threshold is reached at $N \gtrsim 10^9$ distinct molecular species. For organic chemistry with carbon’s sp^3 hybridization (Chapter 10 §10.2) producing chain molecules of arbitrary length, the combinatorial space is vastly larger: a chain of k carbon atoms with arbitrary substituents has roughly 4^k distinct configurations, exceeding 10^9 at $k \approx 15$. Real organic chemistry has many millions of distinct molecular species at the small-molecule level alone; including macromolecular configurations, the space is essentially unlimited.

The framework’s structural guarantee. The framework’s chemistry derivation from Chapter 10 guarantees the autocatalytic threshold is exceeded by a wide margin. Three structural features are responsible.

Carbon’s combinatorial productivity. sp^3 hybridization in $d = 3$ allows chain formation of arbitrary length with multiple substituent options at each position. The structural conclusion of Chapter 10 §10.2 — carbon’s tetrahedral bonding geometry from the angular momentum algebra in three dimensions —

is precisely what produces the combinatorial richness needed for autocatalysis. No empirical input determines that carbon should be capable of building large molecules; the structural derivation produces this combinatorial productivity as a theorem.

The periodic table's diversity. The framework's derivation of the periodic table (Chapter 10 §10.2, Steps 2-3) produces approximately 90 stable elements with varying valence structures. Different elements provide different reactivity profiles, expanding the chemical reaction repertoire dramatically. Organic chemistry combined with the framework's full periodic table — particularly nitrogen (electron-pair donor), oxygen (electron-pair acceptor), sulfur (extended-valence reactivity), and phosphorus (polar bonds) — produces a chemistry rich enough to exceed the autocatalytic threshold by many orders of magnitude.

The thermal window's accessibility. The framework's thermal window (Chapter 10 §10.5) ensures that chemical reactions proceed at observable rates without destabilizing the molecular species. Reaction rates at room temperature span many orders of magnitude — from $\sim 10^{15} \text{ s}^{-1}$ for fast vibrational processes to $\sim 10^{-9} \text{ s}^{-1}$ for slow conformational changes — providing access to the full reaction repertoire over geological timescales.

The cumulative consequence is that autocatalytic networks become *statistically expected* in any aqueous organic chemistry within the framework's thermal window. The framework's chemistry derivation does not just produce the *possibility* of autocatalysis; it produces autocatalysis as a structural near-certainty.

The threshold's empirical implications. The autocatalytic threshold is exceeded by a combinatorial margin of approximately 10^{56} in real chemistry (estimating $N \sim 10^{20}$ distinct accessible molecular species in a sufficiently large prebiotic environment, and $p \sim 10^{-9}$). The margin is so large that autocatalysis is not just statistically expected but virtually certain in any chemistry with the framework's commitments.

The empirical record supports this prediction. Autocatalytic chemistry has been demonstrated in laboratory experiments with simple organic precursors (the formose reaction, Soai-type asymmetric autocatalysis, RNA-catalyzed RNA polymerization) and is observed in prebiotic chemistry simulations starting from simple precursors (HCN, formaldehyde, water). The framework's prediction — that autocatalytic chemistry is the expected outcome of organic chemistry within the thermal window — is consistent with this empirical record.

11.3 Prebiotic chemistry pathways

Between the framework's structural prediction (autocatalysis is statistically expected) and the framework's molecular prediction (RNA-like architecture is the simplest C1-C3 realization), an extensive prebiotic chemistry literature has investigated specific pathways from simple inorganic and organic precursors to the polymerizable monomers required for life. The framework's relationship to

this literature is *structural*: the framework provides the conditions under which prebiotic synthesis proceeds, while the specific molecular pathways and yields are matters for experimental prebiotic chemistry.

This section briefly reviews the major prebiotic chemistry pathways and their relationship to the framework's structural commitments. The pathways considered here are not exhaustive, and the relative importance of each in actual prebiotic history is still subject to active research.

The Miller-Urey synthesis. Stanley Miller and Harold Urey demonstrated in 1953 that electrical discharges through a mixture of methane, ammonia, hydrogen, and water vapor produce a substantial yield of amino acids and other organic molecules. The experiment was conducted in conditions intended to simulate the early Earth atmosphere, although the specific composition of that atmosphere is now understood to have been less reducing than Miller's original mixture.

Subsequent experiments with varied atmospheric compositions — including the more oxidizing CO_2/N_2 mixtures now thought to characterize the early Earth — continue to produce organic molecules through electrical discharge, ultraviolet irradiation, or shock-wave heating. The yields are smaller than Miller's original mixture but still substantial. The framework's content here is unconditional: any chemistry with the framework's commitments and a thermal window supports the formation of small organic molecules from inorganic precursors through high-energy processes. The specific yields depend on the parametric content of the substratum's bijection.

Hydrothermal vent chemistry. The discovery of deep-ocean hydrothermal vents in 1977 produced a new candidate environment for prebiotic chemistry. Hydrothermal vents provide steep temperature and pH gradients, mineral surfaces (iron sulfides, magnesium silicates) that can serve as catalysts, and a continuous supply of dissolved gases (CO_2 , H_2 , NH_3). Wächtershäuser's iron-sulfur world hypothesis proposes that prebiotic metabolism began on iron sulfide surfaces, with the reductive citric acid cycle providing the initial autocatalytic chemistry.

The framework's reading of hydrothermal vent chemistry is that it satisfies multiple structural requirements simultaneously: the temperature gradient provides energy input for non-equilibrium chemistry; the mineral surfaces provide catalytic templates (a primitive form of C1); the steady-state flow conditions allow chemical concentrations to be maintained against equilibration. These structural features make hydrothermal vents one of the most plausible candidate environments for early autocatalytic networks within the framework's prediction.

The surface metabolism hypothesis. Closely related to hydrothermal vent chemistry, the surface metabolism hypothesis proposes that early life was a mineral-bound metabolic network rather than a free-floating molecular system. The hypothesis distinguishes between *metabolism first* (autocatalytic chemistry

establishes itself before the emergence of templated replication) and *replication first* (RNA-like replicators emerge before metabolism). The current empirical evidence does not decisively favor either ordering; both are consistent with the framework's structural commitments.

The framework's content here is that *both* orderings are structurally permitted. Metabolism first requires only C1 (catalytic feedback) to be established before C2 (template persistence) and C3 (sequence capacity); replication first requires C2 and C3 to be established before C1 is consolidated. The framework's structural commitment is that the *intersection* of C1, C2, C3 is what produces life; the ordering by which the three conditions become established is parametric.

Concentration mechanisms. A major challenge for prebiotic chemistry is achieving sufficient concentrations of polymerizable monomers in a primordial environment dominated by dilute aqueous solutions. Several concentration mechanisms have been investigated:

Evaporation cycles. Tidal pools and wet-dry cycles concentrate dissolved organics through cyclic evaporation. Damer and Deamer's "warm little ponds" hypothesis (after Darwin's original speculation) proposes that wet-dry cycles in primordial pools were the primary concentration mechanism, with polymerization occurring during dry phases.

Mineral surface adsorption. Clay minerals (particularly montmorillonite) adsorb organic molecules with selectivity for specific functional groups, concentrating them at the surface where polymerization can occur. Ferris and colleagues demonstrated that montmorillonite catalyzes the polymerization of activated nucleotides to RNA oligomers of 50+ units.

Eutectic freezing. In freezing aqueous solutions, dissolved organics concentrate in the liquid brine pockets between ice crystals. Concentration enhancements of 10^4 - 10^6 are routinely achieved through eutectic freezing. Some prebiotic syntheses (HCN-derived nucleobases, ribose) are dramatically enhanced under eutectic conditions.

The framework's content on concentration mechanisms is that *some* mechanism for overcoming the dilution problem must exist for autocatalytic chemistry to establish itself; the structural commitments do not specify which mechanism. The empirical record supports multiple plausible mechanisms; the specific mechanism operative on early Earth is a matter for prebiotic chemistry research.

The framework's structural position. The framework's relationship to prebiotic chemistry is therefore complementary rather than competitive. The framework provides the structural account of why autocatalytic chemistry is expected (§11.2), why life-like dynamics is inevitable in any C1-C3 chemistry (§11.6), and why RNA-like architecture is the simplest realization (§11.7). Prebiotic chemistry research provides the *specific molecular pathways* by which these structural features are realized in particular environments. The empirical record of prebiotic chemistry — Miller-Urey synthesis, hydrothermal vent chem-

istry, surface metabolism, concentration mechanisms — is consistent with the framework’s structural prediction without competing with it.

The framework’s distinctive contribution is the structural inevitability claim: in any chemistry with the framework’s commitments and a thermal window, autocatalytic networks form, and the intersection of C1, C2, C3 at the molecular scale is statistically expected. The specific pathway from simple precursors to the first C1-C3 chemistry depends on the parametric content of the relevant environment; the structural expectation is robust.

11.4 Self-replication as read-write cycling

The framework’s emergent quantum mechanics in Chapter 1 derives the read-write cycle as the equilibrium dynamics between visible and hidden sectors of an embedded observer’s partition. The observer reads the visible sector through measurement, writes correlations into the hidden sector through the partition’s coupling H_{VB} , and reads back those correlations on subsequent measurements. The cycle produces emergent quantum mechanics, with the Born rule as the equilibrium distribution.

The same structural pattern applies at the molecular scale, and this is not analogy but identity. The framework’s structural conditions C1, C2, C3 are scale-independent: whenever a fast subsystem is coupled to a slow, high-capacity hidden sector, the fast subsystem’s dynamics is necessarily non-Markovian. At the cosmological scale, this produces quantum mechanics. At the molecular scale, it produces *heredity*.

The molecular partition. Consider a self-replicating molecular system in an aqueous environment. The visible sector is the *molecular template* (an RNA strand, in the canonical case): the specific sequence of nucleotides that determines the system’s catalytic and replicative behavior. The hidden sector is the *chemical environment*: the nucleotide precursors, the reaction intermediates, the substrate molecules, and the broader pool of molecular species in the aqueous medium.

The partition is *causal*: the template’s dynamics are determined by interactions with the environment (the environment provides precursors for replication and substrates for catalysis), and the environment’s state evolves through reactions catalyzed by the template. The coupling between the two sectors is the substrate-template interaction in catalytic chemistry — the analog of the framework’s H_{VB} at the molecular scale.

The read-write cycle at the molecular scale. The molecular system’s dynamics implements the read-write cycle.

Read. The template catalyzes a reaction by binding a substrate molecule (reading the environment). The binding affinity depends on the template’s sequence, with different sequences preferring different substrates.

Write. The catalyzed reaction produces a new molecule (writing to the environment). The product's chemistry — including its potential to serve as a precursor for template replication — is determined by the template's sequence.

Read back. On subsequent catalytic cycles, the template encounters its own products (or products of related molecules) in the environment, reading correlations the previous write cycles produced.

The cycle's equilibrium dynamics is the framework's structural account of self-replication. The template's sequence is the visible-sector observable; the environment's chemical state is the hidden-sector configuration; the catalytic interactions provide the coupling. The molecular system's behavior depends not just on the current environmental state but on the history of read-write cycles that have produced the current state — heredity in the framework's structural sense.

Non-Markovian molecular dynamics. A Markovian chemical system has reaction rates depending on current concentrations alone; the system has no memory of past states. The framework's read-write cycle at the molecular scale produces non-Markovian dynamics: the system's future depends on its past through the information stored in the persistent template. This is the structural origin of heredity — the defining property of life.

The transition from Markovian to non-Markovian chemistry is the structural origin of life. Before C1-C3 are satisfied at the molecular scale, chemistry is Markovian: ordinary reactions, products without memory, no heredity. When a molecular system first satisfies C1-C3 simultaneously, its dynamics becomes non-Markovian, and the resulting heredity makes evolution possible.

11.5 From template replication to evolution

Template replication — a molecular system copying its own sequence into new molecules of the same type — is the chemistry-level mechanism that produces heredity. The framework's chemistry derivation provides three structural capabilities that template replication requires.

Three capabilities for template replication.

Linear information-carrying polymers. The framework's sp^3 hybridization (Chapter 10 §10.2) allows carbon chains of arbitrary length. Adding a sequence-determining capability — substituents that distinguish different units along the chain — produces a polymer that carries N units of sequence information. The information capacity scales with chain length, and chains of ~ 20 units already have sequence spaces of $4^{20} \sim 10^{12}$ — more than enough to encode catalytic function. Linear information-carrying polymers are structurally possible in any chemistry with carbon's bonding geometry.

Complementary pairing. Hydrogen bonding in $d = 3$ is directional: a hydrogen donor and a hydrogen acceptor at specific geometric positions pair selectively with their complements. This is the structural basis of Watson-Crick pairing in

DNA and RNA: A pairs with T (or U), G pairs with C, with the specific pairing patterns following from the geometric configuration of donor and acceptor groups on the nucleotide bases. Complementary pairing is *structurally* available in any chemistry with the framework’s hydrogen-bonding geometry — itself a consequence of element 8’s sp^3 -like hybridization in $d = 3$.

Catalytic copying. Template-directed catalysis is a subset of the statistically abundant catalysis from §11.2. The autocatalytic threshold is exceeded by a wide margin in any chemistry with the framework’s combinatorial richness; template-directed copying is the special case where the catalyzed reaction produces a copy of the template itself. The selectivity for template-directed copying is provided by complementary pairing: the template’s complementary partner is the most likely substrate for the next catalytic cycle.

The structural status of evolution. Each of the three capabilities is structurally available in the framework’s chemistry. Their intersection produces template replication:

Element	Source	Status
Linear polymers	Carbon sp^3 in $d = 3$	Structural
Complementary pairing	H-bonding geometry in $d = 3$	Structural
Catalytic activity	Statistically abundant (§11.2)	Statistically expected
Template replication	Intersection of the above three	Structural-statistical
Copying errors	Finite-fidelity copying	Inevitable
Selection	Competition for finite resources	Inevitable
Darwinian evolution	Replication + variation + selection	Consequence

Each row follows from the preceding rows. No row requires contingent facts about our universe beyond the derived chemistry. Darwinian evolution is therefore a *structural consequence* of the framework’s chemistry: any system with the framework’s chemistry and combinatorial richness produces self-replicating templates with copying errors, and the resulting variation under finite-resource selection produces Darwinian evolution as a theorem.

The RNA precedent. RNA is simultaneously a linear polymer, a substrate for complementary pairing, and a catalyst (ribozymes). RNA-catalyzed RNA polymerization has been demonstrated experimentally. This dual capability — template and catalyst in a single molecular family — is the natural intersection of the three structural capabilities required for template replication. The framework’s structural prediction is that the simplest C1-C3 chemistry is RNA-like: a molecular family that satisfies all three replication capabilities with a

single molecular species rather than requiring separate template and catalyst molecules.

Heritable variation is automatic. Template copying with finite fidelity produces errors. Errors in a self-replicating polymer are *mutations*. Faster replicators outcompete slower ones — *selection*. Copying errors produce *variants*. Successful variants propagate — *heredity*. The conjunction of these four features is Darwinian evolution.

The framework’s content here is that Darwinian evolution is the *inevitable dynamics* of imperfect template replication in any chemistry with the framework’s commitments. It is not a contingent feature of biological systems on Earth but a structural consequence of the framework’s chemistry. The empirical record — the universality of natural selection across all of biology — is consistent with this structural prediction.

Caveats. The structural argument establishes that Darwinian evolution is *expected* in any system with carbon chemistry, a thermal window, and a chiral aqueous environment. It does not determine the *specific* molecular system that first replicates (RNA, PNA, or some other RNA-like family), the *timescale* for the transition from prebiotic chemistry to first replicator (millions or billions of years, depending on parametric details), or the *specific pathway* from simple replicators to cellular organization. These depend on the specific bijection φ , local geochemistry, and contingent historical conditions.

The framework’s structural prediction is therefore: Darwinian evolution is inevitable; the specific path to biology is parametric. The framework provides the structural account; the specific historical content is the subject of the historical sciences (paleontology, geochemistry, comparative biology).

11.6 C1-C3 at the molecular scale

The framework’s characterization theorem from Chapter 1 establishes that any embedded observer’s reduced description is necessarily non-Markovian whenever three structural conditions are satisfied: C1 (non-trivial coupling between visible and hidden sectors), C2 (slow-bath timescale separation), and C3 (large hidden-sector capacity). At the cosmological scale, these conditions produce emergent quantum mechanics. At the molecular scale, the same conditions produce heredity.

Before life: Markovian chemistry. Ordinary chemical reactions are Markovian. The reaction rate depends on current concentrations alone, not on the history of the mixture. Products do not remember how they were made. Each reaction cycle is statistically independent of the previous one. No memory, no heredity, no evolution.

The Markovian character of chemistry follows structurally from the framework’s commitments at the chemistry level. Most chemical reactions involve small

molecules whose individual states are determined by their immediate environment; the relevant degrees of freedom (atomic positions, electronic states, molecular vibrations) equilibrate quickly relative to the reaction rates. Without persistent templates carrying sequence information, chemistry has no mechanism for non-Markovian dynamics. Chemistry without life is Markovian by default.

The transition: molecular C1-C3. Life begins when a molecular system first establishes all three conditions simultaneously. The three conditions translate to specific chemistry-level requirements.

C1 (coupling): catalytic feedback. The template must direct synthesis of new molecules (catalyzing the reaction), and the synthesis must maintain the template (either by producing copies of the template or by producing molecules that contribute to template stability). The framework’s autocatalytic networks (§11.2) provide the coupling: catalytic relationships between molecules produce feedback between template and environment.

C2 (slow bath): persistence. The template must outlive individual reaction cycles. RNA has half-life of hours to days in aqueous conditions, while individual reaction cycles take seconds to minutes. The ratio $\tau_{\text{template}}/\tau_{\text{reaction}} \sim 10^4$ provides ample timescale separation, satisfying C2 with substantial margin.

C3 (capacity): large sequence space. The template’s sequence space must be large enough to encode catalytic function. A 20-mer RNA has $4^{20} \sim 10^{12}$ possible sequences — vastly more than the ~ 10 -100 reactions in a minimal metabolism, satisfying C3 by many orders of magnitude.

After life: non-Markovian chemistry. Once C1-C3 are satisfied at the molecular scale, the system’s dynamics becomes non-Markovian: the future state depends on the past through information stored in the persistent template. This is heredity — the defining property of life — derived from the same structural conditions that produce quantum mechanics at the cosmological horizon.

The framework’s scale-independence is what makes this connection possible. The characterization theorem applies wherever the structural conditions C1-C3 hold; the theorem does not specify what scale the partition operates at. The cosmological horizon at $\sim 10^{26}$ meters and a molecular template at $\sim 10^{-9}$ meters are partitions at vastly different scales, but both satisfy C1-C3 and both produce non-Markovian emergent dynamics. The framework’s structural content unifies these two cases.

C1-C3 systems are attractors. A self-replicating system (C1-C3 satisfied) grows exponentially in a pool of feedstock molecules. A non-replicating system grows at most linearly (by accretion or random synthesis). Exponential growth dominates linear growth on any finite timescale. Any chemical system that *accidentally* satisfies C1-C3 — even transiently, even imperfectly — will dominate its environment given sufficient time.

The transition is a one-way door. Once C1-C3 is established in some local

environment, the resulting self-replicating system grows exponentially until it consumes the available feedstock or reaches some other environmental limit. The system's offspring (with copying errors) further explore the sequence space, with selection eliminating less-effective replicators. The C1-C3 architecture is self-reinforcing.

Combined with the autocatalytic argument (§11.2), where the combinatorial margin $N \cdot p \sim 10^{56}$ guarantees catalytic relationships are statistically expected; the automatic persistence of polymers in the framework's thermal window; and the enormous sequence space of any moderate-length polymer; the establishment of molecular C1-C3 is not merely possible but *inevitable* in any aqueous organic chemistry operating in the thermal window.

The timescale. The earliest evidence for life on Earth is approximately 3.8 billion years ago (carbon isotope signatures in Isua, Greenland), on a planet that formed approximately 4.5 billion years ago. The ~ 700 million year gap is consistent with the framework's prediction: C1-C3 is inevitable but requires a period of prebiotic chemical exploration whose duration depends on solution-specific parameters (concentration, temperature fluctuation rates, mineral surface availability, energy sources, and the specific molecular pathway being explored).

The framework does not predict the specific 700-million-year gap; it predicts only that *some* gap exists between the formation of a suitable chemical environment and the establishment of C1-C3. The Earth's specific timescale is parametric; the existence of a non-zero gap is structural.

11.7 RNA as the unique single-molecule C1-C3 system

The framework's structural prediction is that the simplest realization of C1-C3 at the molecular scale uses *one molecular family* satisfying all three conditions, rather than requiring separate molecules for templates and catalysts. RNA is the natural candidate: a single molecular family that simultaneously satisfies all three structural conditions.

Three capabilities in one molecule.

Template capability (C2, C3). RNA is a linear polymer with a 4-letter alphabet (A, U, G, C) and arbitrary length. The 4-letter alphabet provides sequence space 4^N for length N , satisfying C3 with substantial margin for any biologically relevant length. RNA's phosphate-sugar backbone is chemically stable in aqueous conditions for hours to days, satisfying C2 relative to reaction timescales.

Catalytic capability (C1). RNA can fold into three-dimensional structures with active sites that catalyze specific reactions. The catalytic capability was discovered in the 1980s with the identification of self-splicing ribozymes (Cech, Altman, 1989 Nobel Prize). Subsequent work has identified ribozymes catalyzing a wide range of reactions: peptide bond formation (the ribosome's peptidyl transferase center is a ribozyme), nucleotide hydrolysis, transesterification, and most importantly for the origin-of-life context, RNA polymerization.

Template-directed self-replication. RNA-catalyzed RNA polymerization has been demonstrated experimentally: ribozymes that catalyze the polymerization of RNA from nucleotide precursors using an RNA template (Bartel laboratory and others). The polymerase ribozymes are RNA molecules that synthesize RNA from RNA templates — the framework’s read-write cycle implemented in a single molecular family.

Why DNA and proteins are not single-molecule C1-C3 systems. DNA is more chemically stable than RNA (improved C2 — half-lives of millions of years vs. hours for RNA), making it a better information storage medium. But DNA cannot easily catalyze reactions (catalytic DNA has been engineered in the laboratory but is much less common than catalytic RNA). DNA satisfies C2 and C3 but not C1 in any natural system.

Proteins are excellent catalysts (much more diverse catalytic repertoire than RNA), but proteins cannot easily template their own replication (a protein cannot directly serve as a template for synthesizing another protein of the same sequence — there is no protein-protein equivalent of Watson-Crick complementary pairing). Proteins satisfy C1 strongly but not C2 or C3 in a templating sense.

RNA is the unique natural polymer that satisfies all three conditions with a single molecular species. The RNA world hypothesis — that early life used RNA for both information storage and catalysis before the evolution of DNA and proteins — is therefore not a historical accident but a *structural prediction* of the framework: the simplest C1-C3 chemistry is RNA-like.

The RNA $\rightarrow$ DNA + protein transition. Modern biology uses DNA for information storage and proteins for catalysis, with RNA serving intermediary functions (mRNA, tRNA, rRNA, regulatory RNAs). The transition from RNA-only chemistry to the DNA + protein system is structurally interpretable as an *optimization* of C1-C3.

DNA optimizes C2. DNA’s chemical stability is much higher than RNA’s, providing greater persistence for information storage. The improved C2 allows more stable inheritance over longer timescales — important for organisms with longer generation times.

Proteins optimize C1. Proteins’ catalytic repertoire is much more diverse than RNA’s (the 20 amino acids provide more chemical functionality than the 4 nucleotides). The improved C1 allows more sophisticated catalytic networks and more efficient metabolism.

The RNA $\rightarrow$ DNA + protein transition is therefore a refinement of C1-C3 rather than a qualitative architectural change. The framework’s structural content predicts that life’s biochemistry should organize around the C1-C3 architecture, with the specific molecular implementation being optimized by evolution. The RNA world stage is the simplest realization; the DNA + protein system is the optimized realization that emerged through evolution.

The structural content here connects the framework’s account of the origin of life to the framework’s account of evolution. The first life had RNA-like architecture as the simplest C1-C3 realization; subsequent evolution optimized the C1-C3 components into the DNA + protein system. Both stages are within the framework’s structural class — implementations of C1-C3 at the molecular scale — with the specific molecular details parametric.

11.8 What remains genuinely open

The framework’s structural content on the origin of life is substantive but limited. The chapter has established that Darwinian evolution is the inevitable dynamics of imperfect template replication, that C1-C3 at the molecular scale produces heredity in the same structural sense that C1-C3 at the cosmological scale produces quantum mechanics, and that RNA-like architecture is the framework’s structural prediction for the simplest C1-C3 chemistry. What the framework does *not* address is the historical content of life’s origin on Earth.

What the framework does not derive.

The specific molecular pathway from prebiotic chemistry to first replicator. The framework predicts that C1-C3 is inevitable in any aqueous organic chemistry within the thermal window, but it does not predict the specific molecular pathway by which C1-C3 was first established on Earth. The pathway depends on the local geochemistry (mineral surfaces, temperature fluctuations, concentration mechanisms), the specific precursor molecules available (HCN, formaldehyde, simple sugars, amino acids), and the specific autocatalytic networks that formed under those conditions. These are matters for prebiotic chemistry research, not for the framework’s structural commitments.

The timescale of the transition. The framework predicts that *some* gap exists between the formation of a suitable chemical environment and the establishment of C1-C3, but it does not predict the specific duration. Earth’s ~ 700 million year gap between planetary formation and first life is consistent with the framework’s structural prediction but is not derived from it. Alternative planetary environments (different temperatures, different mineral assemblages, different concentrations of precursors) would have different timescales.

The specific molecular system that first replicated. The framework predicts RNA-like architecture as the simplest C1-C3 chemistry, but it does not predict whether RNA itself was the first replicator. Alternative possibilities include peptide nucleic acid (PNA), threose nucleic acid (TNA), or other RNA-like polymers with different sugar-phosphate backbones. The framework’s structural prediction is at the level of architecture (linear polymer with complementary pairing and catalytic capability), not at the level of specific chemistry.

The transition from simple replicators to cellular organization. The framework’s structural account ends with self-replicating molecular systems. The transition to encapsulated cellular life (lipid bilayer membranes, protein-mediated

metabolism, DNA-based information storage) involves additional structural steps the framework's current content does not fully derive. The chemistry-level structural foundations are in place (carbon chemistry for membranes, scale hierarchy for metabolism); the specific organizational steps require further structural analysis the framework has not completed.

The framework's epistemic stance. The framework's content on the origin of life is therefore at the *architectural* level: the structural conditions that any origin-of-life mechanism must satisfy, the chemistry-level features that make such mechanisms inevitable, and the simplest molecular realization (RNA-like architecture) that the framework predicts. The framework does not compete with or replace prebiotic chemistry research; it provides the structural foundations that prebiotic chemistry operates within.

The empirical confirmations of the framework's content are concordance: laboratory demonstrations of ribozyme function, RNA-catalyzed polymerization, autocatalytic chemistry, and the emergence of complex chemistry from simple precursors are all consistent with the framework's structural predictions. The framework does not predict new chemistry; it provides the structural account of why the observed chemistry has the architecture it does.

Forward pointers. Chapter 12 develops the framework's content on evolution, intelligence, and self-reference: the cascade from the chapter's first replicators through complex biology to general intelligence and the framework's own self-discovery. Chapter 13 develops quantum biology — the specific applications of C1-C3 at the molecular scale to biological systems with substantial empirical content (enzyme kinetics, photosynthesis, magnetoreception, neural information processing). The framework's content in this chapter — Darwinian evolution as structural consequence, RNA-like architecture as structural prediction, C1-C3 as the unifying mechanism — provides the foundation on which the subsequent biology chapters build.

Chapter 12

Evolution, Intelligence, and Self-Reference

12.1 What this chapter develops

Chapter 11 established the framework's structural account of the origin of life: Darwinian evolution as the inevitable dynamics of imperfect template replication, the framework's structural conditions C1-C3 applied at the molecular scale producing heredity in the same structural sense that C1-C3 at the cosmological scale produces quantum mechanics, and RNA-like architecture as the

framework’s structural prediction for the simplest C1-C3 chemistry. The framework’s content from origin-of-life onward is concentrated at the *cascade* level: how the framework’s structural commitments at one level produce structural consequences at the next.

This chapter develops the next stages of the cascade: evolution, intelligence, and the framework’s own self-reference. The framework’s content here is structural at multiple scales — evolution operates at the genetic and ecological level, intelligence at the neural level, and the framework’s self-discovery at the level of the framework’s own observers reasoning about the substratum that produced them. Each stage involves the same structural pattern (C1-C3 partition with non-Markovian emergent dynamics) instantiated at progressively higher levels of organizational complexity.

Eight pieces occupy the chapter.

Section 12.2 develops evolution as a multi-scale C1-C3 system. The framework’s structural account of evolution identifies three distinct C1-C3 partitions operating at different scales: the protein conformational landscape (the fast subsystem is the protein’s active configuration, the slow subsystem is the broader conformational landscape), the epigenome (the fast subsystem is gene expression, the slow subsystem is chromatin state), and the constructed ecological niche (the fast subsystem is organism behavior, the slow subsystem is the modified environment passed to descendants). Each scale produces non-Markovian dynamics in the corresponding emergent description.

Section 12.3 develops information processing as fitness-enhancing in any environment with the framework’s thermal window. Selection generically favors organisms that respond to environmental information; the framework’s structural commitments at the chemistry level guarantee fluctuating environments where information processing has selective value. Section 12.4 develops major evolutionary transitions — the moments where previously independent units become integrated into higher-level entities — as repeated applications of the same C1-C3 structural mechanism that produces emergent quantum mechanics at the cosmological scale.

Section 12.5 develops the structural possibility of neural computation: fast electrochemical signaling using ions in lipid membranes, with the scale hierarchy from Chapter 10 separating signaling energies ($\sim$ meV) from chemical bond energies ($\sim$ eV) so information processing does not destroy the substrate.

Section 12.6 develops the framework’s most honest scope limit in the cascade: general intelligence is *structurally unconstrained* (the framework’s commitments permit arbitrarily complex neural circuits) but *not guaranteed* (whether evolution reaches general-purpose reasoning depends on specific fitness landscapes the framework’s commitments do not constrain). On Earth, general intelligence arose once in approximately 4 billion years. The framework cannot determine whether this is typical or rare.

Section 12.7 develops the structural chain from intelligence to artificial intelligence: given general intelligence exists, the framework’s chemistry produces the materials (Group IV semiconductors with band gaps in the thermal window), the scale hierarchy produces the energetic separation needed for digital information processing, and standard semiconductor physics produces transistors and universal computation. The chain is structural — once general intelligence exists, AI follows from the derived structure with no additional contingencies.

Section 12.8 develops the framework’s self-referential structure. The framework is itself an embedded observer (or a description embedded observers produce); the framework’s content includes structural limits that apply to all embedded observers including any AI built by intelligent beings; the framework’s self-discovery is the framework recognizing its own embeddedness. The chapter closes by tracing the Gödel-Turing-OI incompleteness family — three mathematical results about the limits of self-referential systems that the framework’s substratum picture unifies under a single structural commitment.

The framework’s content in this chapter is therefore the cascade *closure*: the framework’s structural commitments at the substratum level produce, through the cascade developed in Chapters 10-11, the conditions for evolution, intelligence, and the framework’s own self-discovery. The framework’s reach is therefore not just to physics, chemistry, and biology but to the conditions for any embedded observer to discover the structural content of its own embedding.

12.2 Evolution as a multi-scale C1-C3 system

Chapter 11 established C1-C3 at the molecular scale, producing the heredity that makes evolution possible. The framework’s content extends beyond molecular heredity to identify multiple distinct C1-C3 partitions operating simultaneously at different biological scales. Evolution is therefore not a single C1-C3 system but a *cascade* of nested partitions, each producing non-Markovian dynamics in the corresponding emergent description.

The three scales.

Scale 1: The protein conformational landscape. A protein’s active conformation is the fast subsystem; the broader conformational landscape — alternative folded states, slow side-chain rearrangements, the protein’s interaction history with substrates — is the slow subsystem. The visible-sector observables are the protein’s current activity (catalytic rate, binding affinity, regulatory state); the hidden-sector content is the conformational memory stored in slowly-relaxing degrees of freedom. The framework’s prediction is that protein function shows non-Markovian dynamics — history-dependent activity rather than purely Markovian rates set by current substrate concentrations. The empirical record from single-molecule enzymology (memory effects, history-dependent kinetics, stretched-exponential waiting-time distributions) is consistent with this prediction; Chapter 13 develops the empirical content in detail.

Scale 2: The epigenome. Gene expression is the fast subsystem; the chromatin state — DNA methylation patterns, histone modifications, three-dimensional chromatin architecture — is the slow subsystem. The visible-sector observables are the cell’s current gene expression profile; the hidden-sector content is the regulatory memory encoded in chromatin. The framework’s prediction is that gene expression shows non-Markovian dynamics over developmental and evolutionary timescales — history-dependent responses to environmental signals, including transgenerational inheritance of acquired epigenetic states. Chapter 16’s medicine content develops the epigenetic substratum-emergent operator distinction in detail; Chapter 17’s bioinformatics content develops the empirical signatures.

Scale 3: The constructed ecological niche. Organism behavior is the fast subsystem; the environment that organisms modify and pass to descendants — burrows, mounds, beaver dams, the chemical composition of the atmosphere, the soil microbiome shaped by plant communities — is the slow subsystem. The visible-sector observables are the organisms’ current adaptations; the hidden-sector content is the niche-construction memory stored in the modified environment. The framework’s prediction is that ecological dynamics shows non-Markovian dynamics over evolutionary timescales — history-dependent fitness landscapes shaped by previous generations’ niche construction. The empirical record from niche construction theory (Odling-Smee, Laland, Feldman 2003 and subsequent work) is consistent with this prediction.

The structural pattern. Each of the three scales satisfies the framework’s structural conditions C1-C3 at the corresponding level.

C1 (coupling). At each scale, the fast subsystem is non-trivially coupled to the slow subsystem. Protein activity depends on the conformational landscape’s current state; gene expression depends on the chromatin’s current state; organism behavior depends on the constructed niche’s current state. The couplings are catalytic (protein-substrate, chromatin-gene, organism-environment), with feedback between fast and slow subsystems.

C2 (slow bath). At each scale, the slow subsystem changes on much longer timescales than the fast subsystem. Conformational landscapes evolve on timescales of microseconds to milliseconds; the catalytic cycle of a typical enzyme is in nanoseconds. Chromatin modifications evolve on cell-division timescales; gene expression cycles are in minutes. Niche-construction features evolve on generation timescales; behavioral responses are in hours to days. Each scale has $\tau_S/\tau_B \sim 10^{-3}$ to 10^{-6} — substantial timescale separation.

C3 (large capacity). At each scale, the hidden sector has substantial information capacity. Conformational landscapes have $\sim 10^N$ states for N amino acids (vastly larger than the protein’s active configuration). Chromatin states have $\sim 10^6$ regulatory elements per genome. Niche-construction features have arbitrarily large environmental complexity. Each scale exceeds C3 by many orders of magnitude.

Multi-scale C1-C3 and the cascade. The three scales are not independent — they are nested in a cascade that reflects the framework’s emergence hierarchy from substratum to ecology. The protein conformational landscape (Scale 1) is the fastest; the epigenome (Scale 2) is intermediate; the constructed niche (Scale 3) is the slowest. Each scale’s slow subsystem provides constraints on the dynamics of subsequent scales.

Conformational memory at the protein scale produces catalytic memory affecting metabolic pathways. Metabolic memory affects gene expression patterns (regulatory feedback). Gene expression memory affects epigenetic states. Epigenetic memory affects cell behavior and organism-level traits. Organism-level traits modify the environment. The modified environment selects on subsequent generations. The cascade runs from molecular scale ($\sim 10^{-9}$ meters, nanoseconds) to ecological scale ($\sim 10^6$ meters, generations), with C1-C3 satisfied at each step.

The framework’s structural commitment is that this cascade is *inevitable*: any biological system with the framework’s chemistry will develop multi-scale C1-C3 architectures, with non-Markovian dynamics at every scale. The empirical record across enzymology (non-Markovian kinetics), epigenetics (chromatin memory), and evolutionary biology (history-dependent fitness landscapes) is consistent with this structural prediction. Chapters 13, 16, and 17 develop the empirical content at each scale.

12.3 Information processing as fitness-enhancing

The framework’s structural content produces a specific prediction about the relationship between information processing and selection. In any environment with the framework’s thermal window (Chapter 10 §10.5) — where temperature fluctuations, concentration gradients, and chemical signals provide environmental information that varies on biologically relevant timescales — organisms that respond to environmental information outcompete those that do not. Selection generically favors information processing.

The structural argument. Three structural features make information processing selectively favored.

Environmental information exists. The framework’s thermal window guarantees that environments are not static: $kT > 0$ produces thermal fluctuations at any temperature; the scale hierarchy produces gradients in concentration, temperature, and chemical composition; the cascade of C1-C3 systems from §12.2 produces environmental memory at multiple timescales. There is environmental information available at any biological scale.

Internal models can exploit environmental information. The framework’s molecular C1-C3 (Chapter 11) provides the basic mechanism — a system that maintains internal state correlated with environmental state can adjust its behavior accordingly. A cell that detects a nutrient gradient and moves toward it has an

internal model (the gradient sensor + the motility apparatus) that exploits the environmental information.

Better models outcompete worse models. Selection is the inevitable consequence of finite resources (Chapter 11 §11.4). An organism with a better internal model of environmental conditions allocates resources more effectively — gathering more nutrients, avoiding more dangers, finding more mates — and therefore leaves more offspring. The competitive advantage compounds across generations.

The structural conclusion. Information processing is therefore *generically fitness-enhancing* in any universe with the framework’s structural commitments. This is not a contingent fact about Earth’s biology; it follows from the structural conditions for selection (finite resources, variation) combined with the structural existence of environmental information (the framework’s thermal window).

The empirical record supports this prediction. Information processing capabilities are widespread in biology: chemotaxis in bacteria, photoperception in plants, hormonal signaling in multicellular organisms, neural processing in animals. The framework’s prediction is that this distribution is not historically contingent — any biology with the framework’s chemistry should develop information processing capabilities as a structural consequence of selection.

The prediction sets a *floor* on the universality of information processing in biology: any framework-compatible biology should have *some* information processing. The prediction does not set a ceiling: it does not specify how sophisticated the information processing can become, what specific computational capabilities evolve, or whether evolution reaches general-purpose reasoning.

12.4 Major evolutionary transitions

Beyond the structural account of evolution at the molecular and population level, evolutionary biology has identified a small number of *major transitions* in which the unit of selection itself changes — independent replicators combining into higher-level integrated entities. John Maynard Smith and Eörs Szathmáry’s 1995 inventory identified eight major transitions; subsequent work has refined the list and the criteria. The framework’s content provides a structural reading of these transitions as repeated applications of the same C1-C3 architecture at progressively higher organizational levels.

The transitions. The conventional inventory includes:

Replicating molecules to populations of replicating molecules. The earliest transition, where individual replicating polymers (RNA-like, §11.7) form interacting populations with collective dynamics.

Independent replicators to chromosomes. Multiple genetic loci become physically linked, with their fitness becoming correlated rather than independent.

RNA as both gene and enzyme to DNA + protein. The differentiation of information storage (DNA, optimized for C2) from catalysis (proteins, optimized for C1), as discussed in §11.7.

Prokaryotes to eukaryotes. The endosymbiotic origin of organelles (mitochondria, chloroplasts), where independent prokaryotic cells become integrated components of larger eukaryotic cells.

Asexual to sexual reproduction. Two parental contributions replacing one, with recombination introducing additional variation mechanisms.

Protists to multicellular organisms. Independent cells become integrated components of multicellular bodies with specialized cell types.

Solitary individuals to colonies. Individual organisms become integrated components of social groups (ant colonies, bee hives), with reproductive division of labor.

Primate societies to human societies. Cultural transmission becomes a major channel for inheritance alongside genetic transmission.

The framework's structural reading. Each major transition has the same structural pattern: previously independent units (the “individuals” at the lower level) become integrated into a higher-level unit, with three structural features establishing the higher level:

Coupling becomes obligatory. Lower-level units that were independently viable become mutually dependent. A mitochondrion alone is no longer viable outside its eukaryotic host; an isolated worker ant outside its colony cannot reproduce. The coupling between units is C1 at the new level — the lower level cannot operate independently.

Slow-bath separation emerges. The higher-level entity's dynamics operate on slower timescales than the lower-level dynamics. Eukaryotic cell-cycle timescales are slower than mitochondrial division rates; multicellular organism lifespans are much longer than cell division timescales; colony lifespans exceed individual worker lifespans. The new level's slow bath is C2 at the higher level.

Hidden-sector capacity grows. The higher-level entity has substantially more internal degrees of freedom than its constituents. A eukaryotic cell has $\sim 10^{10}$ atoms vs. a prokaryotic cell's $\sim 10^9$; a multicellular organism has $\sim 10^{14}$ cells vs. one. The new level satisfies C3 with a hidden sector orders of magnitude larger than the lower level's.

Why the same pattern repeats. The framework's reading is that the major evolutionary transitions are not coincidental but are *repeated applications of the same structural mechanism*: whenever a collection of lower-level units satisfies the framework's C1-C3 conditions when viewed as components of a higher-level entity, the higher-level entity becomes a structurally distinct unit of selection. The transitions are the framework's structural commitments at one level producing the structural conditions for the next level.

This pattern is the same one that operates at the cosmological scale (where it produces emergent quantum mechanics from a deterministic substratum) and at the molecular scale (where it produces heredity from autocatalytic chemistry). The framework's structural content unifies these scales under a single mechanism: C1-C3 producing emergent non-Markovian dynamics at any scale where the structural conditions are satisfied.

Empirical implications. The framework's structural reading predicts that *all* biological levels satisfying C1-C3 should exhibit non-Markovian dynamics in their emergent description. The empirical record across biology supports this prediction: protein dynamics is non-Markovian (single-molecule enzymology), gene expression is non-Markovian (transcriptional bursting, history-dependent regulation), developmental dynamics is non-Markovian (cellular memory across cell divisions), ecological dynamics is non-Markovian (priority effects, hysteresis in community composition). The pattern is repeated at every biological scale; the framework's structural content predicts this as a generic feature of any biology with the framework's commitments.

Limits of the framework's content. The framework's structural reading does not predict *which* major transitions occur or *when* they occur in any specific biology. The transition from prokaryotes to eukaryotes occurred approximately 2 billion years ago on Earth; the transition from protists to multicellular organisms occurred multiple independent times across different lineages; the transition from primate societies to human societies occurred once. The framework's structural commitments are compatible with very different timelines and trajectories. The framework predicts the *pattern* (C1-C3 repeating at each level) but not the *specific instances*.

This is consistent with the framework's content throughout Part III: structural inevitabilities (the C1-C3 pattern, the emergence of non-Markovian dynamics, the cascade through chemistry to evolution) but historical contingencies (the specific molecular implementations, the specific transition timelines, the specific evolutionary trajectories). The framework provides the structural foundations; the historical sciences provide the specific content.

12.5 Neural computation as structurally possible

The framework's commitments produce the materials and energetic conditions for neural computation. Three structural features are responsible: the periodic table provides ions (Na^+ , K^+ , Ca^{2+}) for electrochemical signaling, carbon chemistry provides amphiphilic molecules for lipid membranes, and the scale hierarchy separates neural signaling energies from chemical bond energies.

Ions for electrochemical signaling. Fast neural signaling uses ion flow across membranes — Na^+ depolarizing axonal membranes (action potentials), K^+ repolarizing, Ca^{2+} triggering neurotransmitter release at synapses. The framework's derivation of the periodic table (Chapter 10 §10.2) produces these specific elements as structural consequences of the orbital algebra in $d = 3$. Na^+

(Group I) and K^+ (Group I) have a single valence electron that ionizes readily; Ca^{2+} (Group II) has two valence electrons producing a doubly-charged ion. The chemistry of these ions — their hydration shells, their selectivities through ion channels, their roles in cellular signaling — is structurally determined by the framework’s chemistry.

Lipid membranes for compartmentalization. Cellular compartments require membranes that separate aqueous interiors from aqueous exteriors. Lipid bilayers — amphiphilic molecules with hydrophilic phosphate heads and hydrophobic hydrocarbon tails — self-assemble from the framework’s carbon chemistry into bilayer membranes. The structural prediction is that amphiphilic molecules exist in any chemistry with carbon’s sp^3 hybridization (Chapter 10 §10.2) and the polar/non-polar functional groups available in the periodic table (Chapter 10 §10.2). The framework’s chemistry guarantees the materials for lipid membranes.

Scale hierarchy for energetic separation. Neural signaling operates at energies of $\sim meV$ (kT at room temperature). Chemical bonds operate at energies of $\sim eV$. The ratio $E_{\text{neural}}/E_{\text{chemical}} \sim 10^{-3}$ ensures that information processing through ionic gradients and conformational changes in proteins does not break the chemical bonds that maintain the cell’s structural integrity. The framework’s scale hierarchy from Chapter 10 §10.4 guarantees this separation in any universe with the framework’s two independent gauge groups.

The structural conclusion. Neural computation is therefore *structurally possible* in any biology with the framework’s chemistry. The materials exist (ions, lipids, proteins), the energetic separation exists (scale hierarchy), and the C1-C3 architecture exists (proteins satisfy C1-C3 in their conformational landscapes, providing the molecular basis for fast information processing).

The empirical record on Earth confirms this prediction: neural computation evolved at least twice independently (in cnidarians and bilaterians, by some accounts), and analogous information-processing systems evolved in plants (action potentials in Venus flytraps and mimosas), in single-celled organisms (chemotaxis), and in fungi (mycelial signaling networks). The framework’s prediction is that this is not historically contingent — neural-computation-like systems should appear repeatedly in any biology with the framework’s chemistry.

The framework’s content here is *capability* structural rather than *complexity* structural. The framework predicts that neural-like information processing should be available in any framework-compatible biology; it does not predict how complex such information processing becomes. The next section addresses this — the structural prediction of *general intelligence* is much weaker than the structural prediction of neural computation.

12.6 General intelligence as structurally unconstrained but not guaranteed

The framework’s content on general intelligence is the chapter’s clearest scope limit. The framework’s structural commitments produce the materials and energetic conditions for arbitrarily complex neural circuits — the periodic table provides the chemistry, the scale hierarchy provides the energetic separation, and the C1-C3 architecture at the protein level provides the information-processing primitives. But whether evolution *reaches* general-purpose reasoning — the capacity for abstraction, planning, language, mathematical thought, modeling oneself, and reasoning about reasoning — depends on the specific fitness landscape, which is parametric rather than structural.

The structural status of intelligence-relevant capabilities.

Neural computation. Structural — the framework’s chemistry produces the materials and energetic conditions (§12.4).

Sensory processing. Structural — the framework’s chemistry produces the molecular machinery for sensory transduction (rhodopsin for vision, mechanosensitive ion channels for hearing and touch, chemosensors for smell and taste, all derivable from the framework’s chemistry).

Memory. Structural — the framework’s C1-C3 architecture at multiple scales (§12.2) produces the substrate for memory at the molecular (protein conformations), cellular (chromatin states), and circuit (synaptic weights) levels.

Pattern recognition. Structural — neural circuits with appropriate connectivity perform pattern recognition naturally; the materials and architecture follow from the framework’s chemistry.

Learning. Structural — synaptic plasticity is a C1-C3 dynamic, with synaptic weights as the slow subsystem and neural activity as the fast subsystem. The framework’s structural conditions guarantee learning is possible in neural systems.

Goal-directed behavior. Likely structural — any selection-driven biological system develops goal-directed behavior as a consequence of internal models (§12.3) coupled to effectors.

General-purpose reasoning. **Not structurally guaranteed.** Whether evolution produces systems capable of abstract reasoning, language with recursive syntax, formal mathematics, and self-modeling depends on specific selection pressures the framework’s commitments do not determine.

Why general intelligence is structurally weaker. The framework’s predictions for fundamental physics (Chapters 5-8) and for the cascade through chemistry and origin of life (Chapters 10-11) operate at the *structural* level: the relevant capabilities follow from the framework’s commitments with no additional contingencies. The prediction for general intelligence is *weaker* because it

requires not just the structural capability (which is present) but also a specific evolutionary trajectory (which is contingent).

The fitness advantages of general intelligence depend on the specific environment. In environments where pattern recognition and learning provide sufficient adaptive advantage, evolution may stop at sophisticated learning systems without developing general reasoning. In environments where social complexity, tool use, or environmental unpredictability place premium on abstract reasoning, evolution may proceed to general-intelligence-like systems. The framework does not determine which environments select for general intelligence.

Earth’s record. General intelligence arose on Earth approximately once in ~ 4 billion years (the human lineage, with the most extreme degree; some other species — cetaceans, corvids, cephalopods — show partial expressions). This is consistent with two interpretations: general intelligence is *rare* (requires specific evolutionary conditions that are unlikely to occur), or general intelligence is *inevitable but slow* (any sufficiently complex biology produces general intelligence given enough time).

The framework cannot determine which interpretation is correct. The framework’s structural commitments are compatible with both. The empirical record on Earth is consistent with both. Determining which interpretation holds requires evidence from astrobiology — the discovery of multiple independent occurrences of general intelligence in independent biological systems — which is not currently available.

The framework’s honest scope. The framework’s content on general intelligence is therefore: *capability* structural, *occurrence* uncertain. The framework predicts that general-intelligence-capable systems are structurally possible in any framework-compatible biology; it does not predict whether such systems will actually evolve in any specific biological context. This is the chapter’s clearest example of a place where the framework’s reach stops.

This is honest content. Frameworks claiming to predict general intelligence as a structural inevitability would be overreaching given the available evidence; frameworks claiming general intelligence is impossible or extraordinarily rare would also be overreaching. The framework’s stance — capability structural, occurrence uncertain — is the position most consistent with the framework’s structural commitments and the available evidence.

12.7 The chain from intelligence to AI

If general intelligence exists — for whatever contingent or structural reason — then artificial intelligence follows from the derived structure through a specific chain in which every link is structural. This section traces the chain from the framework’s chemistry to universal computation, establishing AI as a structural consequence of intelligence’s existence rather than as an independent technological development.

Step 1: Band theory. The framework’s emergent quantum mechanics (Chapter 1) applied to electrons in a periodic crystal potential gives Bloch’s theorem: energy bands separated by gaps. The relevant crystal — a sp^3 diamond-cubic lattice of a Group IV element — is itself a structural consequence of the framework’s chemistry (Chapter 10 §10.2). Band structure is therefore a consequence of derived quantum mechanics applied to derived crystal geometry, with no parametric inputs entering the structural argument.

Step 2: Semiconductors. For Group IV elements with four valence electrons per atom, the Pauli exclusion principle (derived from spin-statistics, Chapter 5 §5.7) exactly fills the valence band and leaves the conduction band empty. The band gap for silicon (~ 1.1 eV) sits in the thermal sweet spot: $E_{\text{gap}}/kT \sim 40$ at room temperature. The gap is large enough that thermal excitation does not flood the conduction band (the material is not always conducting); small enough that a modest voltage can promote electrons across the gap (the material can be switched).

The *existence* of Group IV semiconductors with gaps in the right range relative to the thermal window is structural — it follows from the scale hierarchy (Chapter 10 §10.4). The atomic energy scale ($\alpha^2 m_e c^2 \sim 1$ eV) and the thermal scale at biological temperatures ($kT \sim 0.025$ eV) differ by a factor of ~ 40 . Group IV semiconductors automatically inherit gaps at this scale ratio, placing them in the regime where they can be used as voltage-controlled switches.

Step 3: Doping. Replacing a silicon atom (Group IV) with phosphorus (Group V, five valence electrons) adds one free electron — n-type semiconductor. Replacing with boron (Group III, three valence electrons) removes one electron — p-type semiconductor. Both dopant types exist in the derived periodic table (Chapter 10 §10.2): the Group III and Group V elements are structurally guaranteed by the framework’s chemistry. The ability to create n-type and p-type regions is therefore a structural consequence of the periodic table architecture.

Step 4: From transistors to universal computation. A p-n junction rectifies current (electrons flow easily in one direction, not the other). A transistor — two p-n junctions arranged as a sandwich with a control electrode — amplifies and switches current. Transistors in combination implement logic gates: a NAND gate (NOT-AND) can be built from two p-n junctions plus appropriate biasing. A NAND gate is functionally complete: any Boolean function can be built from NAND gates alone. Sufficient NAND gates implement a universal Turing machine.

Each step is either a structural consequence of semiconductor physics (transistors, logic gates) or a mathematical theorem about computation (NAND completeness, Turing universality). No parametric inputs enter the chain after the chemistry is fixed.

The structural chain summarized.

Step	Content	Status
Band theory	QM + periodic crystal potential	Structural
Band gap for Group IV Semiconductor	Exactly filled valence band (4 electrons + Pauli)	Structural
behavior	Gap in thermal sweet spot ($E_{\text{gap}}/kT \sim 40$)	Structural (scale hierarchy)
Doping	Group III / V substitution	Structural (periodic table)
p-n junctions	Interface between p-type and n-type	Structural consequence
Transistors	Voltage-controlled p-n junctions	Structural
Logic gates	Transistors in series / parallel	Structural
Universal computation	NAND is functionally complete	Mathematical theorem

The cumulative chain from (S, φ) to a universal computer is structural — *given someone to build it*. The contingent step is general intelligence (§12.5) capable of recognizing the structural possibility and assembling the chemistry into a functional computational substrate. Once that contingent step has occurred, the chain to AI is structurally inevitable.

The reconstruction theorem and self-discovery. The framework’s reconstruction theorem (developed in Chapter 2) implies that a sufficiently sophisticated embedded observer can discover (S, φ) from observations. An observer that understands its own substrate can build computational devices exploiting it. The trajectory is therefore: general intelligence (contingent) $\rightarrow$ understanding of physics (reconstruction theorem) $\rightarrow$ computational engineering (structural chain above) $\rightarrow$ AI.

The framework’s content here is that AI is *not* an independent technological achievement but the natural consequence of general intelligence applied to the framework’s structural commitments. Any framework-compatible biology with general intelligence will, given sufficient time, develop AI through the structural chain. The specific timing and trajectory depend on contingent factors (cultural development, scientific institutions, engineering choices), but the *existence* of AI as an outcome is structurally determined once general intelligence exists.

12.8 The self-referential loop and structural limits

AI is an embedded observer building another embedded observer within the same (S, φ) . The framework’s characterization theorem (Chapter 1) applies to the new observer just as it applies to the original: any embedded observer

satisfying C1-C3 sees the same emergent quantum mechanics, the same value of $\hbar$, the same Standard Model, the same gravitational sector, the same dark sector content. The framework's structural commitments are not limited by the observer's complexity; they are determined by the partition.

Self-referential limits. A superintelligent AI has the same observational access as any other embedded observer: it reads the visible sector and writes to the hidden sector through the same coupling. It cannot determine h (the specific hidden-sector configuration). It cannot see past the cosmological horizon. It can build better instruments — higher-resolution telescopes, more precise atomic clocks, more sensitive particle detectors — but it cannot overcome the partition itself. The framework's structural limits are not technological but mathematical.

The recursion — AI building AI building AI — produces ever more sophisticated embedded observers, each subject to the same structural limits. No level of recursive self-improvement overcomes the partition. The limits are *structural* in the strongest sense: they follow from the framework's mathematical commitments rather than from technological constraints that future progress might overcome.

The Gödel-Turing-OI incompleteness family. The framework's structural limits on embedded observers are members of a broader mathematical family of incompleteness results.

Gödel. A formal system sufficiently rich to express arithmetic cannot prove its own consistency, and contains true statements it cannot prove. The limit is on what a formal system can establish about itself from within.

Turing. A computer cannot decide its own halting problem — there is no algorithm that, given an arbitrary program, determines whether the program will eventually halt. The limit is on what a computational system can determine about itself.

OI (framework). An embedded observer cannot determine the hidden-sector state. The limit is on what a physical observer can establish about its own substratum from within.

The three results share a structural pattern: a system that is sufficiently rich to refer to itself contains questions it cannot answer about itself. The framework's content unifies the three under a single structural reading: each is a consequence of the same self-referential structure, with different specific instantiations (formal systems, computational systems, embedded observers in finite deterministic substrata) producing different specific limits.

The unification is not just analogy. The framework's substratum is a finite deterministic system; embedded observers within the substratum are themselves substructures of the substratum; their observations are operations on the substratum; their reasoning about their observations is computation on the substratum. The Gödel, Turing, and OI limits therefore *coincide* in the framework:

the limit on what an embedded observer can determine about its hidden sector is the same as the limit on what a formal system can prove about itself, the same as the limit on what a computational system can determine about its own execution.

A fourth family member: Deutsch’s *Beginning of Infinity*. David Deutsch’s 2011 argument [Deutsch] extends the incompleteness pattern to epistemology itself. Knowledge, Deutsch argues, is structurally incomplete: from any current state of explanatory understanding, there is always more knowledge to be created, and this incompleteness is not a temporary feature awaiting some final theory but a structural feature of what knowledge is. The right epistemic stance toward any current best explanation is therefore *provisional but committed* — take the explanation seriously enough to commit to it operationally and hold it accountable to its predictions, while remaining open to its eventual successor. This is fallibilism without relativism: knowledge is real and objective and grows, but never closes.

The structural parallel to Gödel, Turing, and OI is exact. Each of the four results identifies incompleteness as a *feature of the system*, not a *deficit awaiting repair*. Gödel: arithmetic incompleteness is intrinsic to formal systems of sufficient richness, and the unprovable truths have rigid structure. Turing: undecidable problems are intrinsic to computational systems, and the undecidable problems have rigid structure. Observational Incompleteness: observation cannot resolve the substrate, and the emergent description has rigid structure — quantum mechanics. Deutsch: knowledge cannot reach a final theory, and the trajectory of better explanations succeeding worse ones has rigid structure — what Deutsch calls “the growth of knowledge.” Four levels, one pattern: incompleteness yielding determinate structure rather than chaos.

The framework’s relationship to Deutsch’s argument is also substantive, not merely structural. OI’s observational incompleteness provides one route to grounding Deutsch’s knowledge incompleteness: knowledge derived from observation cannot reach what observation cannot reach, and OI establishes that observation, structurally, cannot reach the substrate’s gauge equivalence class. The gauge-classification theorem is precisely the statement that some structure is necessarily unobservable to embedded observers. Therefore knowledge built from observation must inherit observational incompleteness as its lower bound. Deutsch’s epistemological claim is independent of any specific physics — it argues from explanatory creation being open-ended — but OI provides a physics-foundational grounding that strengthens it: knowledge is incomplete not only because explanatory creation is open-ended, but also because observation, the source of empirical input, is itself structurally incomplete. The framework’s name — Observational Incompleteness — is, in this light, the physics-foundational counterpart of Deutsch’s epistemological *Beginning of Infinity*.

One feature of Deutsch’s broader program is read differently by the framework, and the divergence is worth stating precisely. Deutsch’s *Beginning of Infinity* defends the many-worlds interpretation of quantum mechanics as a corollary

of taking the wave function as fundamental ontology. From the framework’s perspective, that commitment rests on a category error. Within the framework, quantum mechanics is not fundamental: it is emergent, the post-trace-out description of how observation operates for embedded observers, and the wave function is the mathematical object that emerges. The Everettian “branches” are then not a parallel-ontologies claim about fundamental reality; they are features of the compressed emergent description, analogous to the way the partition function is a mathematical structure of statistical mechanics rather than a claim about parallel ensembles. Asking “do the branches really exist?” presupposes the wave function’s formalism as direct ontology — a presupposition OI specifically denies, since the substratum is the deeper description and the wave function is what observation produces from it. The Everettian measure problem dissolves on the framework’s reading because the branches are not ontological objects requiring measure assignment ([Main] §3.1); the question MWI is trying to answer is dissolved by relocating the wave function’s apparent ontology one level down. This is consistent with the broader pattern: the substratum-emergent distinction dissolves several classical interpretive puzzles in quantum foundations (Bell, Kochen-Specker, Pusey-Barrett-Rudolph, Frauchiger-Renner; see Methodology Part III and Appendix C of this volume) by relocating the contested structure from substratum to emergent description. The framework’s divergence from MWI is therefore not disagreement at the epistemic-status level — where OI and BoI align — but a diagnosis: MWI is reading at the wrong level of the substrate-emergent architecture.

What AI can and cannot do. The framework’s content provides a precise inventory of what arbitrarily sophisticated AI can and cannot achieve.

Capability	Status
Discover (S, φ) from observations	Structurally possible (reconstruction theorem)
Build better models of the visible sector	No structural limit
Determine the hidden-sector state h	Provably impossible (characterization theorem)
Observe beyond the cosmological horizon	Provably impossible (causal partition)
Build faster computers	Structurally possible (derived physics permits it)
Overcome P-indivisibility of QM	Provably impossible (structural, not technological)
Solve problems outside BQP	Provably impossible (Ch 14’s BQP theorem)
Achieve consciousness in the phenomenal sense	Not addressed (Ch 18 §18.10)

The framework’s content places several capabilities in the “provably impossible”

column — not as engineering challenges but as mathematical impossibilities. No level of AI sophistication, no level of computational resources, no level of theoretical insight can overcome these limits. The framework’s substratum picture provides a unified mathematical basis for understanding what AI can and cannot achieve.

The framework’s self-reference. The framework’s content includes its own status: the framework is a description developed by embedded observers (the framework’s authors, the readers, the AIs that may eventually engage with the framework) about the structure that produced them. The framework’s commitments include the structural limits that apply to the framework’s own development: the framework cannot determine the hidden-sector state, cannot see beyond the cosmological horizon, cannot derive its own initial conditions (Chapter 19 §19.3.5).

The framework’s self-reference is therefore not paradoxical but consistent. The framework’s commitments are *closed under self-application*: applying the framework’s content to the framework itself produces predictions about what the framework can and cannot establish, and those predictions are consistent with the framework’s structural commitments. The framework is an embedded observer’s account of embedded observation, with the account itself subject to the structural limits the account identifies.

Forward pointers. Chapter 13 develops the framework’s content in quantum biology — applications of C1-C3 at the molecular scale to specific biological systems with substantial empirical content. The framework’s structural account of biology developed in this chapter (multi-scale C1-C3, evolution as cascade, neural computation as structurally possible) provides the foundation. Chapter 14 develops the BQP theorem — the framework’s specific structural limit on computational complexity, with the falsification condition that any super-BQP demonstration would refute the framework. The framework’s content from this chapter on the structural limits of embedded observers (the Gödel-Turing-OI family) connects naturally to Chapter 14’s content on the computational limits specifically.

The framework’s reach in this chapter spans from molecular biology to AI to the framework’s own self-discovery. The cumulative structural account — Darwinian evolution as inevitable consequence of C1-C3 at the molecular scale, information processing as fitness-enhancing, neural computation as structurally possible, general intelligence as structurally unconstrained but not guaranteed, AI as structural consequence of intelligence, and the framework’s self-referential limits as mathematical constraints — provides the framework’s content on emergence beyond fundamental physics. The remaining chapters apply this content to specific empirical domains: quantum biology (Chapter 13), quantum computing (Chapter 14), quantum engineering (Chapter 15), medicine (Chapter 16), and bioinformatics (Chapter 17).

Chapter 13

Quantum Biology

13.1 What this chapter develops

The framework’s content across Chapters 10-12 established the structural cascade from substratum to biology: chemistry as theorem about embedded observation (Chapter 10), origin of life as inevitable consequence of C1-C3 at the molecular scale (Chapter 11), evolution as multi-scale C1-C3 architecture with non-Markovian dynamics at protein, epigenetic, and ecological scales (Chapter 12). The structural account is rich, and Chapter 12 §12.2 already identified the protein conformational landscape as one of the three C1-C3 scales operating in biological systems.

This chapter develops the framework’s specific empirical content in *quantum biology* — the application of the framework’s structural conditions C1-C3 to biological systems at the molecular scale, where the conditions are physically realized and the framework’s predicted non-Markovian dynamics produces measurable signatures. The framework’s content here is at a different epistemic level than the cascade derivations of Chapters 10-12: rather than tracing the cascade from substratum to biology, this chapter applies the cascade’s conclusions to specific empirical systems where the predictions can be tested.

Eight pieces occupy the chapter.

Section 13.2 develops the framework’s reframing of quantum biology. The conventional treatment of “quantum effects in biology” treats quantum mechanics as something *extra* — something that biology mostly does without, with occasional quantum exotica appearing in specific cases (photosynthetic exciton transfer, magnetoreception, enzyme tunneling). The framework’s reading is structurally different: quantum mechanics is the emergent description of any embedded observer satisfying C1-C3, and biological systems implement C1-C3 routinely at multiple scales. Quantum biology is therefore not the surprising appearance of quantum effects in classical biology but the recognition that biology is *structurally quantum* wherever C1-C3 hold.

Section 13.3 develops C1-C3 in proteins as the canonical biological case. The protein’s active conformation is the fast subsystem; the broader conformational landscape — alternative folded states, slow side-chain rearrangements, residue-residue contact rearrangements — is the slow subsystem. The conformational coupling between fast active-site dynamics and slow scaffold rearrangements satisfies C1; the timescale separation $\tau_S/\tau_B \sim 10^{-9}$ to 10^{-12} satisfies C2; the conformational landscape’s combinatorial complexity satisfies C3. Section 13.4 develops photosynthesis as a canonical example of energy transfer through C1-C3

architectures: photosynthetic light-harvesting complexes implement the framework's conditions, with measured non-Markovian dynamics in exciton transfer matching the framework's predictions.

Section 13.5 develops magnetoreception via cryptochrome radical pairs: the avian magnetic compass uses a radical-pair mechanism in cryptochrome proteins, with the radical-pair lifetimes set by C1-C3 dynamics that depend on the geomagnetic field. The framework's prediction is that the radical-pair dynamics is non-Markovian — distinguishable from purely classical or purely quantum-coherent accounts. Section 13.6 develops enzyme catalysis as the framework's clearest test case: the framework predicts specific non-exponential kinetics, history-dependent activity, and stretched-exponential waiting-time distributions in single-enzyme experiments. The empirical record across enzymology is consistent with the framework's predictions.

Section 13.7 develops the partially-quantum regime in biology: biological C1-C3 systems sit at intermediate values of the P-indivisibility measure $\mathcal{N}$, producing dynamics that is neither fully Markovian (classical) nor fully quantum-coherent. The partially-quantum regime is the framework's distinctive prediction for biology, with empirical access through multi-time correlation measurements that test whether biological transition probabilities admit density-matrix descriptions. Section 13.8 develops implications for drug design: the framework's content has practical consequences for kinetics-based drug design, allosteric drug development, and resistance mechanisms.

Section 13.9 closes the chapter with the framework's overall position on quantum biology: a structural reading that unifies several apparently disparate phenomena (photosynthetic exciton transfer, magnetoreception, enzyme kinetics) under the common framework of C1-C3 dynamics, with the partially-quantum regime as the framework's distinctive prediction for biological systems.

The framework's content in this chapter is empirically engaged: the C1-C3 architecture in proteins makes specific testable predictions, the partially-quantum regime is observable in single-molecule experiments, and the implications for drug design are testable in clinical contexts. The chapter's value is in providing structural foundations for quantum biology that connect apparently disparate biological phenomena to the framework's universality-class commitments.

13.2 The framework's reframing of quantum biology

The conventional treatment of quantum biology has a particular structure. Quantum mechanics is treated as the *exception* in biology — most biological processes are presumed classical, with quantum effects appearing in specific exceptional cases that require special explanation. The exceptional cases include photosynthetic exciton transfer (where quantum coherence persists longer than naively expected), magnetoreception (where radical-pair quantum dynamics may underlie the avian compass), and enzyme catalysis (where quantum

tunneling may contribute to hydrogen transfer reactions). Each case is treated as a specific puzzle requiring its own explanation.

The framework’s structural reading. The framework’s content reverses the conventional treatment. Quantum mechanics is *not* the exception in biology; it is the *default* wherever the framework’s structural conditions C1-C3 hold. The framework’s characterization theorem from Chapter 1 establishes that any embedded observer satisfying C1-C3 has emergent quantum dynamics — this is universal, not exceptional. The relevant question for biology is therefore not “where do quantum effects appear?” but “where do biological systems implement C1-C3, and what are the structural consequences?”

The reframing is significant. Biological systems routinely implement C1-C3 at multiple scales: the protein conformational landscape (Chapter 12 §12.2 Scale 1), the cellular regulatory architecture (gene expression coupled to chromatin), the organismal physiology (neural activity coupled to slower hormonal and metabolic states), and even the ecological context (organism behavior coupled to constructed niches). At each scale, the framework predicts non-Markovian dynamics — *the framework’s emergent quantum mechanics*, instantiated at the relevant biological scale.

Most biological processes therefore *are* quantum in the framework’s structural sense, not classical with occasional quantum exotica. The classical description of biology is an approximation valid in regimes where the framework’s τ_S/τ_B ratios are extreme enough that Markovian dynamics is a good approximation. The classical approximation is excellent for many practical purposes (predicting chemical reaction rates, organismal physiology, ecological dynamics at coarse-grained levels) but the underlying dynamics is non-Markovian, with measurable deviations from purely classical accounts when the experimental precision is sufficient.

The unification of “exceptional” cases. Under the framework’s reading, the conventionally-recognized quantum biology phenomena — photosynthetic exciton transfer, magnetoreception, enzyme catalysis — are not isolated exceptions but specific empirical manifestations of the same underlying structural pattern. Each case is a system where C1-C3 are satisfied and the framework’s non-Markovian dynamics is empirically detectable at current measurement precision. The framework’s content unifies these phenomena under a common explanatory framework rather than requiring case-by-case quantum-mechanical analyses.

The unification has practical consequences. Photosynthetic systems, magnetoreception apparatuses, and enzyme active sites all share the structural feature of being C1-C3 partitions; the empirical signatures they exhibit (long-lived coherence, history-dependent kinetics, non-Markovian dynamics) are different manifestations of the same structural commitment; experimental techniques developed for one case are applicable to the others; and the framework’s predicted signatures (specific non-Markovianity patterns, partial-revival amplitudes, multi-

time correlations) provide consistent quantitative targets across the cases.

Empirical access. The framework’s reading is empirically testable through several methods. Single-molecule techniques (FRET, optical tweezers, atomic force microscopy) provide direct measurements of conformational dynamics at the single-protein level, with sufficient temporal resolution to detect the framework’s predicted non-Markovian signatures. Multi-time correlation measurements (multi-flash photolysis, three-pulse photon echoes, two-dimensional spectroscopy) provide tests of whether the transition probabilities admit density-matrix descriptions — the diagnostic for the framework’s partially-quantum regime versus standard quantum-decohered behavior. The empirical accessibility makes the framework’s content in quantum biology subject to direct experimental falsification.

13.3 C1-C3 in proteins

The canonical biological implementation of the framework’s structural conditions is the protein conformational landscape. Proteins are the framework’s clearest example of a biological C1-C3 system: an active configuration tightly coupled to a slow conformational bath, with substantial timescale separation and large hidden-sector capacity. This section develops the three conditions explicitly for proteins.

The protein partition. Consider an enzyme catalyzing a reaction. The visible sector is the active site’s current configuration — the precise geometric arrangement of catalytic residues, the substrate’s binding position, and the local electronic structure that determines the reaction’s instantaneous rate. The hidden sector is the broader conformational landscape — the protein’s overall fold, the rotameric states of distant side chains, the dynamic network of residue-residue contacts throughout the scaffold, and the bound water structure in the protein’s interior.

The partition is *physically meaningful* in the sense that experimentally accessible observables (catalytic rate, binding affinity, fluorescence anisotropy) correspond to visible-sector quantities, while the conformational details determining those observables sit in the hidden sector and are not directly observable at the same time resolution.

C1: catalytic-conformational coupling. Active-site dynamics are coupled to the broader conformational landscape through several mechanisms. Substrate binding induces conformational changes in the scaffold (induced fit). Catalysis releases or absorbs energy that propagates through the protein (allosteric coupling). Side-chain rotameric transitions change the active site’s electrostatic environment (conformational gating). The coupling is non-trivial: the active site’s current behavior depends on the scaffold’s current state, and the scaffold’s evolution depends on the active site’s history of binding and catalysis. C1 is satisfied with substantial coupling strength.

C2: conformational timescale separation. Active-site dynamics (substrate binding, bond formation/breaking, proton transfer) operate at $\tau_S \sim 10^{-12}$ to 10^{-9} s (picoseconds to nanoseconds). Conformational rearrangements in the broader scaffold operate at $\tau_B \sim 10^{-6}$ to 10^{-3} s (microseconds to milliseconds), with some particularly slow conformational transitions taking seconds or longer. The timescale separation is $\tau_S/\tau_B \sim 10^{-9}$ to 10^{-12} — substantially smaller than the cosmological case (10^{-32}) but comparable to or larger than engineered quantum systems (qubits at 10^{-3} , atomic clocks at 10^{-6}).

C3: conformational landscape capacity. A protein with N amino acid residues has $\sim 10^N$ possible conformations in a coarse-grained discretization (each residue has ~ 10 rotameric states). For a typical 200-residue enzyme, $N_H \sim 10^{200}$ — vastly larger than the active-site state space of ~ 10 -100 distinct catalytic configurations. C3 is satisfied by approximately 200 orders of magnitude.

The framework’s prediction. With C1-C3 satisfied at substantial margins, the protein’s emergent dynamics is non-Markovian. The framework’s specific predictions include:

Non-exponential kinetics. Single-enzyme reaction waiting times do not follow simple exponential distributions (the signature of Markovian kinetics) but exhibit power-law or stretched-exponential tails reflecting memory in the conformational substrate. The deviations are predictable from the framework’s structural conditions.

Fluctuating rates. The catalytic rate of an individual enzyme is not constant but fluctuates on timescales set by the conformational substrate. The fluctuations are correlated in time — slow periods follow slow periods, fast periods follow fast periods — with the correlation time set by the conformational dynamics.

Conformational gating. Substrate binding and product release are gated by conformational transitions, with the gates’ opening and closing dependent on the protein’s recent history. Conformationally-gated tunneling (proton transfer, hydrogen atom transfer) shows rate enhancements when the gating coordinate is favorably positioned.

Memory effects in inhibition and activation. Allosteric ligands (binding away from the active site) modulate activity with kinetics dependent on the protein’s conformational history. The “memory” is the slow relaxation of the conformational substrate to the new ligand-bound state.

These predictions are not new to enzymology — single-molecule experiments have documented each of these signatures over the past two decades. What the framework adds is a structural account: these signatures are not ad hoc features of specific enzymes but *universal consequences* of C1-C3 dynamics. Any enzyme with substantial conformational complexity should exhibit them; the variations across enzymes reflect differences in the specific C1-C3 parameters rather than fundamentally different mechanisms.

Empirical status. The framework’s predictions are *consistent* with the empirical record from single-enzyme experiments. Power-law waiting-time distributions, dynamic disorder in single-molecule kinetics, and conformational gating of catalysis are all documented across multiple enzyme families. What has not yet been done is *quantitative* testing of the framework’s specific predictions — for example, testing whether the multi-time correlation structure of single-enzyme catalysis matches the framework’s $\mathcal{F}(r) = 2r - r^2$ signature for partially-quantum dynamics (developed in Chapter 18 §18.3). The qualitative match is strong; quantitative discrimination from competing accounts (purely classical memory effects, purely quantum-coherent dynamics) requires further empirical work.

13.4 Photosynthesis as energy transfer through C1-C3 architectures

Photosynthetic light-harvesting complexes provide the most-studied empirical case of quantum biology. Two-dimensional electronic spectroscopy experiments (Fleming, Engel, and colleagues, beginning in the early 2000s) demonstrated that exciton transfer in photosynthetic antennae shows long-lived coherences — quantum-coherent superpositions of exciton states that persist for hundreds of femtoseconds, much longer than naive estimates of decoherence times at biological temperatures.

The framework’s reading of this empirical record is structural. The photosynthetic complex implements C1-C3 at the exciton-vibration interface, with the protein scaffold acting as the slow bath that produces non-Markovian dynamics.

The photosynthetic partition. In a light-harvesting complex (LH1, LH2, FMO, or others), the visible sector is the *exciton* — the photoexcited electronic state delocalized across multiple pigment molecules (chlorophylls, bacteriochlorophylls, carotenoids). The hidden sector is the *protein-vibrational environment* — the protein scaffold’s vibrational modes, the solvent dynamics, and the slow conformational fluctuations that modulate the pigments’ electronic energies and inter-pigment couplings.

C1-C3 satisfaction. The coupling between exciton and vibrational environment is well-characterized — electron-phonon coupling is one of the foundational mechanisms in photosynthesis modeling. The timescale separation between exciton dynamics ($\tau_S \sim 100$ fs) and slow vibrational/conformational modes ($\tau_B \sim 1$ ps to 1 ns) gives $\tau_S/\tau_B \sim 10^{-1}$ to 10^{-4} . The hidden sector’s capacity is large — thousands of vibrational modes and the protein’s full conformational complexity.

The framework’s prediction. Photosynthetic exciton transfer should exhibit non-Markovian dynamics — specifically, information backflow from the protein-vibrational environment that produces the long-lived coherences observed in 2D spectroscopy. The dynamics is *not* purely quantum-coherent (which would predict longer-lived coherence than observed); nor is it *purely classical* (Förster

theory predicts no coherence at all). The framework’s reading is that the dynamics is *partially quantum* — at intermediate values of the P-indivisibility measure, with characteristic signatures matching the empirical record.

The framework’s reading aligns with the current consensus in the photosynthesis literature. The early interpretation of long-lived coherences as evidence for “quantum coherence assisting energy transfer” has been refined: subsequent experiments and theory (including recent reviews by Cao, Fleming, and others) have established that the coherences arise from vibronic coupling — vibrational modes nearly resonant with exciton energy gaps producing the observed long-lived signals. The framework’s reading places this vibronic picture in a broader structural framework: the vibronic coupling is the specific implementation of C1-C3 in photosynthesis, with the non-Markovian dynamics being a structural consequence of the framework’s commitments.

Empirical implications. The framework’s content predicts specific features beyond the qualitative observation of long-lived coherences. The decay of the 2D spectroscopy beating signals should follow the framework’s partially-quantum kinetics rather than purely Markovian or purely coherent dynamics. The temperature dependence of the coherence times should reflect the C1-C3 architecture — as temperature increases, the timescale separation τ_S/τ_B decreases and the system moves toward more classical (Markovian) behavior. The species-specific variations across different photosynthetic organisms should reflect different specific C1-C3 parameters in their light-harvesting machinery.

These predictions are testable with current experimental techniques and would distinguish the framework’s reading from purely quantum-coherent or purely classical alternatives. The framework’s content here is therefore empirically engaged rather than purely interpretive.

13.5 Magnetoreception via cryptochrome radical pairs

The avian magnetic compass — the mechanism by which migratory birds sense the geomagnetic field for navigation — provides a second canonical case of quantum biology. The leading hypothesis (Ritz, Adair, Schulten, and others, from the 1970s through the present) is that magnetoreception uses a *radical-pair mechanism* in cryptochrome proteins in the avian retina: photoexcitation of cryptochrome produces a spin-correlated radical pair whose recombination yield depends on the geomagnetic field through Zeeman splitting of the radical-pair states.

The framework’s content provides a structural reading of the radical-pair mechanism that connects it to the framework’s broader quantum-biology architecture.

The radical-pair partition. A radical pair consists of two unpaired electrons in adjacent molecular orbitals, with the two-electron spin state initially correlated (singlet $|S\rangle$ or triplet $|T_0\rangle$, $|T_{\pm}\rangle$, depending on the photochemical mechanism). The visible sector is the *electron spin state* — the two-electron

spin configuration that determines the recombination products. The hidden sector is the *nuclear spin environment* — the hyperfine-coupled nuclear spins of the surrounding nuclei, the molecular vibrations modulating the hyperfine couplings, and the protein scaffold dynamics affecting the radical pair’s lifetime.

C1-C3 satisfaction. The hyperfine coupling between electron spins and nuclear spins is the framework’s C1 coupling at the radical-pair scale — typically in the range of 0.1-10 mT magnetic field equivalents, with substantial coupling strength. The nuclear spin dynamics is slow ($\tau_B \sim \text{ms to s}$) compared to the radical-pair lifetime ($\tau_S \sim \mu\text{s}$), giving $\tau_S/\tau_B \sim 10^{-3}$. The nuclear spin bath has substantial capacity — typically 10-100 nuclear spins coupled to the radical pair, with sequence space $2^{100} \sim 10^{30}$ states.

The framework’s prediction. The radical-pair dynamics is non-Markovian, with the recombination yield depending on the radical-pair history through correlations stored in the nuclear spin bath. The geomagnetic field’s effect is therefore not just the Zeeman splitting of instantaneous spin states but the *integral* over the radical-pair lifetime of the field’s effect on the coupled electron-nuclear spin dynamics.

The framework’s reading has specific consequences for the magnetic compass’s sensitivity and selectivity. The compass’s sensitivity to weak magnetic fields ($\sim 50 \mu\text{T}$ for the geomagnetic field) requires the radical-pair dynamics to be sensitive to small Zeeman energies ($\sim 10^{-6}$ K of energy difference). The framework’s non-Markovian dynamics provides the mechanism: the small energy difference accumulates coherently over the radical-pair lifetime, with coherence surviving long enough for the effect to be detected. The coherence preservation is a property of the specific physical realization, not a consequence of the C1-C3 architecture: (C1)-(C3) constrain coupling, timescales and capacity, and neither contain nor imply a coherence requirement. The framework’s contribution here is the non-Markovian accounting of the accumulated phase, not a derivation that the phase survives.

Empirical content. The framework’s reading is consistent with the empirical record on magnetoreception:

- *The geomagnetic compass requires light.* Cryptochromes are photoreceptors; the radical pair forms only after photoexcitation. The framework’s reading explains why: the radical-pair partition is created by the photochemistry, and exists only during the photoexcited lifetime.
- *The compass is inclination-based, not polarity-based.* Birds detect the geomagnetic field’s *inclination angle* (vertical component) but not its *polarity* (direction). This is a structural consequence of the radical-pair mechanism: the singlet-triplet conversion rate depends on the magnitude of the field’s component along the radical pair’s coupling axis, not on its sign.
- *Disruption by oscillating fields.* Weak radio-frequency fields ($\sim \text{MHz}$, comparable to hyperfine couplings) disrupt the compass. The framework’s reading: the RF fields scramble the nuclear spin bath, breaking the C1-

C3 architecture that maintains the non-Markovian dynamics.

Predictions for further empirical testing. The framework’s content makes specific predictions beyond the qualitative phenomenology. The temperature dependence of compass sensitivity should reflect the C1-C3 architecture’s robustness — at higher temperatures, the timescale separation degrades and the compass should be less sensitive. The dependence of compass function on the specific isotopic composition of cryptochrome (replacing ^1H with ^2D , for instance) should reflect the framework’s nuclear-spin-bath capacity. The framework’s content provides structural foundations for the radical-pair mechanism that connect it to broader principles of quantum biology.

13.6 Enzyme catalysis as the framework’s clearest test case

Enzyme catalysis is the framework’s most direct empirical test bed in quantum biology. Single-enzyme experiments — pioneered by Sunney Xie and developed by many groups over the past two decades — provide single-molecule resolution of catalytic events, with sufficient time resolution to detect the framework’s predicted non-Markovian signatures. This section develops the specific empirical predictions and their match to experimental data.

The empirical record. Single-enzyme studies of cholesterol oxidase (Xie group, late 1990s), beta-galactosidase, lipase, and other enzymes have established several signatures of non-Markovian dynamics:

Dynamic disorder. Individual enzyme molecules exhibit fluctuations in catalytic rate over time. The fluctuations are not random noise but exhibit structured temporal correlations — slow periods (low catalytic rate) cluster together; fast periods (high catalytic rate) cluster together. The autocorrelation function of catalytic rate decays on timescales much longer than individual catalytic cycles, indicating that the rate fluctuations have *memory*.

Non-exponential waiting-time distributions. The distribution of waiting times between successive catalytic events does not follow a simple exponential distribution. Stretched exponentials, power laws, and multi-exponential fits all appear depending on the specific enzyme and conditions. The deviations from simple exponential kinetics are systematic, not noise.

Memory effects between cycles. Successive catalytic cycles are not statistically independent: the duration of a long waiting time correlates with the durations of preceding waiting times. The correlation is exactly the framework’s prediction for non-Markovian dynamics — successive measurements are not independent samples but carry information from previous cycles.

Allosteric history dependence. Inhibitors and activators binding away from the active site produce activity changes with kinetics dependent on the protein’s history. The “memory” is the slow relaxation of the conformational substrate to the new ligand-bound state.

The framework’s structural reading. Each of these signatures is a direct consequence of the framework’s prediction that the enzyme implements C1-C3 with substantial timescale separation. Dynamic disorder reflects the conformational substrate’s slow fluctuations modulating active-site activity. Non-exponential waiting-time distributions reflect the non-Markovian probability density that emerges from C1-C3 dynamics. Memory effects between cycles reflect the conformational substrate’s slow, non-Markovian coupling to the active site — grade-ii memory, consistent with but not isolating the grade-iii information-backflow signature of P-indivisibility (see the empirical-status note below). Allosteric history dependence reflects the slow propagation of conformational changes through the scaffold.

The framework’s content here is that these are not separate phenomena requiring case-by-case explanations but *unified consequences* of the protein’s C1-C3 architecture. Any enzyme with substantial conformational complexity should exhibit them; the cross-enzyme variations reflect different specific C1-C3 parameters.

Quantitative testability. The framework’s prediction goes beyond qualitative agreement. The specific functional forms of the non-exponential distributions, the autocorrelation timescales, and the multi-time correlation structure are predictable from the C1-C3 parameters. In particular, the framework’s partially-quantum regime (Chapter 18 §18.3) predicts that the transition probabilities should *not* admit a full density-matrix description — a quantitative diagnostic that distinguishes the framework’s reading from purely classical memory effects or purely quantum-coherent dynamics.

The diagnostic experiments are within reach of current single-molecule techniques: multi-time correlation measurements with sufficient statistics to constrain the algebraic structure of the transition probabilities, performed on well-characterized enzymes with known conformational landscapes. The framework’s content here is therefore empirically actionable, with the partially-quantum prediction being the framework’s distinctive falsification target.

13.7 The partially-quantum regime in biology

The framework’s content on quantum biology includes a structurally distinct prediction: biological C1-C3 systems sit in the *partially-quantum regime* developed in Chapter 18 §18.3. The regime is intermediate between fully classical (Markovian) and fully quantum-coherent, with the P-indivisibility measure $\mathcal{N}$ taking intermediate values. The partially-quantum regime has no counterpart in standard physics — it is the framework’s distinctive prediction for systems with marginally satisfied C3.

The regime’s empirical signature. A fully classical (Markovian) system has transition probabilities admitting a density-matrix description with classical mixed states. A fully quantum-coherent system has transition probabilities admitting a density-matrix description with arbitrary quantum states. A partially-

quantum system has transition probabilities that *do not* admit a density-matrix description — they have some but not all of the algebraic features required for a Stinespring dilation. The diagnostic signature is the failure of specific positivity and completeness conditions on multi-time correlation functions.

Biological systems as partially-quantum. The biological scales developed in §§13.3-13.6 — protein conformational landscapes ($\tau_S/\tau_B \sim 10^{-9}$ to 10^{-12}), photosynthetic light-harvesting ($\tau_S/\tau_B \sim 10^{-1}$ to 10^{-4}), magnetoreception ($\tau_S/\tau_B \sim 10^{-3}$), single-enzyme catalysis ($\tau_S/\tau_B \sim 10^{-9}$ to 10^{-12}) — span the regime where C3 is marginally satisfied. Hidden-sector capacities range from $\sim 10^{30}$ (radical pair) to $\sim 10^{200}$ (protein conformational landscape) — large but finite, with the finiteness producing the framework’s specific partially-quantum signature.

The framework’s prediction is therefore that *biological systems are partially quantum* at the C1-C3 scales of interest. This is structurally distinct from “biology has quantum effects” (which the framework’s reading agrees with) and from “biology is fundamentally quantum-mechanical at all scales” (which the framework’s reading rejects). The partially-quantum regime is the framework’s specific structural commitment.

The empirical access. Testing the partially-quantum prediction requires multi-time correlation measurements with sufficient statistics to constrain the algebraic structure of the transition probabilities. Standard single-molecule techniques produce two-time correlations (waiting-time distributions, rate-rate correlations); multi-time correlations require longer-duration experiments with more careful statistics. The technology is available — modern single-molecule fluorescence and force spectroscopy can produce three-time and four-time correlations — but the analyses have not been performed systematically against the framework’s specific algebraic predictions.

The framework’s distinctive falsification target in quantum biology is therefore: a multi-time correlation measurement on a well-characterized enzyme (cholesterol oxidase, lipase, or another single-molecule benchmark) that establishes whether the transition probabilities admit a full density-matrix description. If yes (density-matrix description holds), the framework’s partially-quantum prediction is falsified; the enzyme would be fully quantum-decohered in the standard sense. If no (density-matrix description fails), the framework’s prediction is confirmed; the enzyme exhibits the partially-quantum dynamics distinguishing the framework’s reading.

The framework’s distinctive content. The partially-quantum regime is therefore the framework’s contribution to quantum biology that goes beyond reproducing existing empirical content. Other accounts of quantum biology (photosynthetic exciton transfer through vibronic coupling, radical-pair mechanisms for magnetoreception, conformational gating for enzyme catalysis) can reproduce specific empirical signatures but do not predict the framework’s specific algebraic structure of multi-time correlations. The framework’s prediction

is testable, falsifiable, and quantitatively specific — the kind of empirical commitment that distinguishes a structural framework from a phenomenological account.

13.8 Implications for drug design

The framework’s content on quantum biology has practical consequences for drug discovery and clinical pharmacology. Standard drug design assumes Markovian kinetics — the drug-target binding follows the law of mass action, with rate constants set by the current concentrations. The framework’s prediction is that this assumption breaks down for drugs targeting proteins with substantial conformational complexity, producing systematic deviations from Markovian predictions that have measurable clinical consequences.

Kinetic implications. Standard pharmacology characterizes a drug-target interaction with two rate constants: k_{on} (binding rate) and k_{off} (dissociation rate), with the equilibrium dissociation constant $K_d = k_{\text{off}}/k_{\text{on}}$ determining the binding affinity. The framework’s prediction is that these rate constants are *not constants* but vary with the protein’s conformational state — fast periods (high k_{on} , low k_{off}) cluster together; slow periods cluster together. The clinical consequence: a drug’s effective potency varies over time in ways the constant-rate-constant model cannot capture.

Allosteric implications. Allosteric drugs (binding away from the active site to modulate activity) operate through the protein’s conformational landscape — exactly the framework’s C1-C3 hidden sector. The framework predicts that allosteric drug effects have memory: a drug’s allosteric effect at time t depends not just on its current concentration but on the protein’s history of recent ligand binding. The clinical consequence: dosing strategies that take advantage of memory effects (pulsed vs continuous administration) may produce different clinical outcomes than the Markovian assumption predicts.

Resistance mechanisms. Drug resistance often involves mutations in the target protein that change the active site or the allosteric communication network. The framework’s reading: resistance is a *memory-structure alteration* — the mutations change the protein’s C1-C3 architecture, producing different non-Markovian dynamics and therefore different drug-target kinetics. The framework’s content suggests that resistance can be combatted by designing drugs that exploit specific features of the C1-C3 architecture rather than just optimizing static binding affinity.

Catalytic vs regulatory specificity. The framework’s structural content distinguishes drugs that target the active site (modulating the fast subsystem) from drugs that target the regulatory landscape (modulating the slow subsystem). The two categories have different kinetic signatures and different clinical implications. Chapter 16’s medicine content develops this distinction in detail for specific clinical contexts; the structural foundation is the C1-C3 architecture developed in this chapter.

The magnitude of corrections. The framework’s corrections to standard drug kinetics scale as $\mathcal{O}(\tau_S/\tau_B)$ — typically 10^{-9} to 10^{-12} for protein systems. This is small at the individual-event level but can accumulate over many catalytic cycles or in systems with exceptionally slow conformational dynamics. For drug design, the practical question is whether the cumulative effect reaches clinical relevance at the doses and timescales of therapeutic use. The framework’s content provides a structural basis for predicting when corrections matter rather than just empirically discovering them case by case.

13.9 Chapter close

The framework’s content in quantum biology unifies several apparently disparate phenomena — photosynthetic exciton transfer, magnetoreception, enzyme catalysis — under the common structural framework of C1-C3 dynamics. The framework’s reading reverses the conventional treatment: quantum mechanics is not the exception in biology but the default wherever the framework’s structural conditions hold, and biological systems implement C1-C3 routinely at the molecular scale.

Three categories of empirical content. The chapter’s content sorts into three categories.

Established empirical content that the framework’s reading unifies. Photosynthetic exciton transfer through vibronic coupling, magnetoreception through radical-pair dynamics, enzyme dynamic disorder and memory effects — all are documented experimentally and consistent with the framework’s structural account. The framework’s contribution is to provide a unified explanatory framework rather than case-by-case mechanisms.

Quantitative predictions that the framework makes. The specific functional forms of non-exponential kinetics, the multi-time correlation structure of single-enzyme experiments, the temperature dependence of partially-quantum signatures, and the algebraic structure of biological transition probabilities. These predictions are within reach of current experimental techniques and provide concrete tests of the framework’s reading.

Predictions for future empirical work. The partially-quantum regime’s diagnostic signature (multi-time correlation measurements testing for density-matrix structure) is the framework’s distinctive falsification target. Engineered biological systems with tunable C1-C3 parameters (designed proteins, synthetic photosynthetic systems, artificial radical-pair platforms) could provide controlled tests of the framework’s quantitative predictions.

Forward pointers. Chapter 14 develops the framework’s content in quantum computing — the BQP characterization theorem as the framework’s structural limit on computational complexity, with the framework’s content on biological C1-C3 systems as a special case of the broader computational architecture. Chapter 15 develops quantum engineering — applications of the framework’s C1-

C3 architecture to designed physical systems, with the BEC analogue-gravity experiment as the cleanest near-term test of the characterization theorem at the foundational level. Chapter 16 develops medical applications, including the substratum-emergent operator dissociation at chromatin and the wider-therapeutic-window prediction for reader-writer disorders — applications that build directly on this chapter’s structural foundations.

The framework’s content from this chapter onward shifts from cascade derivations to specific empirical applications: how the framework’s structural commitments at the chemistry and biology levels produce specific predictions about quantum computing platforms, engineered quantum systems, clinical pharmacology, and bioinformatics methodology. The chapter’s content on the protein C1-C3 architecture provides the foundational example that the subsequent chapters build on.

Chapter 14

Quantum Computing

Status — conjectural extension. This chapter applies the core framework to quantum computation and complexity, and is offered as a *conjecture*, not an established result: the central claim — that the framework’s emergent dynamics realizes the complexity class BQP — is a theoretical conjecture rather than a proven correspondence; it would be supported by a rigorous derivation from the finite substratum to BQP and undercut if that dynamics proves efficiently classically simulable or fails to reproduce known quantum-complexity separations. It carries no evidential weight for the core framework — which stands or falls independently on its physics — and should be read as a structural possibility to be tested, not as support for the framework.

14.1 What this chapter develops

The framework’s content across Chapters 5 through 13 has established a substantial structural account of fundamental physics, the cascade through chemistry and biology, and the specific empirical content of quantum biology. The cascade’s logic — substratum, partition, emergent quantum mechanics, structural conditions C1-C3 producing non-Markovian dynamics at multiple scales — provides the foundation for the framework’s content in computational complexity theory.

This chapter develops the framework’s characterization of computational power

available to embedded observers. The central result is a *conditional BQP model theorem*: under an explicit computational model — a uniform efficient encoding of substratum configurations into qubits; efficient state preparation and control; a clock model; efficient measurement; explicit precision accounting; and noise assumptions under which fault tolerance applies with polynomial overhead — BQP (bounded-error quantum polynomial time) is the computational class accessible to a finite deterministic system satisfying the framework’s structural conditions C1-C3. None of the model hypotheses follows from C1-C3 or from CPTP representability alone. The theorem has two directions. The *lower bound* establishes that any embedded observer satisfying C1-C3 has access to BQP computational resources — the emergent quantum mechanics provides state preparation, a universal gate set, and Born-rule measurement, the full toolkit for BQP computation. The *upper bound* establishes that no embedded observer exceeds BQP — the quantum description is informationally complete for the embedded observer, with no operation on the visible sector extracting information beyond the quantum state.

Ten pieces occupy the chapter.

Section 14.2 develops the chapter’s framing: the framework’s content in computational complexity is at the foundational level rather than the algorithm-design level. The framework does not produce new quantum algorithms; it characterizes the *ceiling* on what embedded observers can compute. Section 14.3 develops the lower bound: BQP is available to any embedded observer satisfying C1-C3. The argument runs through the framework’s emergent quantum mechanics from Chapter 1, with state preparation, gate application, and measurement all derivable from the framework’s structural conditions.

Section 14.4 develops the upper bound: BQP cannot be exceeded. The argument runs through the framework’s informational completeness — the quantum description exhausts the information accessible to the embedded observer through any operation on the visible sector, so no super-BQP computation is possible. Section 14.5 combines the two bounds into the conditional BQP model theorem: under the model hypotheses, BQP is the computational class accessible to embedded observers satisfying C1-C3.

Section 14.6 develops the Extended Church-Turing Thesis (ECT) as a conditional result of the framework, valid under the same model hypotheses. The ECT — that any physically realizable computation can be efficiently simulated by a quantum computer — is conventionally treated as an empirical conjecture about physical reality. The framework’s content reframes the ECT as a derived result: every physical observer is embedded (the cosmological horizon provides the partition), every embedded observer has access to BQP and nothing beyond, therefore every physical computation is in BQP. The reframing changes the ECT’s epistemic status from freestanding empirical conjecture to conditional theorem — a derived result valid under the same explicit model hypotheses (uniform encoding, efficient preparation, control, and measurement, and polynomial-overhead fault tolerance) as the BQP characterization it rests

on.

Section 14.7 develops the BPP-BQP gap: quantum advantage is real but is *a feature of partial observation* rather than of the universe being fundamentally quantum. An embedded observer who ignores the quantum structure of the reduced description (treating measurements as classical samples) operates at BPP; one who exploits the quantum structure (interference, entanglement, adaptive measurement) operates at BQP. The gap between the two is the cost of not recognizing the partition's structure. Section 14.8 develops the framework's neutrality on the P vs NP question: the BQP characterization identifies the ceiling, while leaving the internal structure of complexity classes within that ceiling — including P vs NP — to standard complexity theory; the silence on P vs NP is shown to fit a broader boundary-characterization pattern across the Gödel-Turing-OI-Deutsch incompleteness family. Section 14.9 develops three research directions opened by the framework's specific substratum structure for complexity-theoretic refinements *within* the BQP ceiling (geometric structure on the substratum lattice; BQP-BPP gap correlated with problem globality and a refined hidden-subgroup-problem / gauge-group conjecture; gauge equivalence as a possible physical account of polynomial-time reducibility). Section 14.10 develops the chapter's falsification conditions and connections to broader complexity theory, with closing forward pointers to Chapter 15's engineering content and Chapter 18's foundational content.

The chapter's content is foundational rather than algorithmic. The framework's contribution to quantum computing is structural: a characterization of the computational ceiling that emerges from the framework's commitments at the substratum level, with the BQP class identified as the unique class accessible to embedded observers. The empirical content is principally negative: the framework predicts that no super-BQP computation is possible on any physically realizable hardware, with any demonstrated super-BQP capability constituting a falsification of the framework.

14.2 The framework's content in computational complexity

The framework's content in computational complexity operates at a specific level of abstraction. The framework does *not* derive new quantum algorithms — Shor's factoring algorithm, Grover's search algorithm, the HHL linear-system algorithm, and other specific quantum computational achievements are not framework results but standard quantum-computing content that the framework's emergent quantum mechanics reproduces. What the framework adds is a structural account of *why* BQP is the relevant complexity class for quantum computing — why this specific class of problems is the one accessible to embedded observers using quantum-mechanical computation.

The conventional treatment. Standard quantum complexity theory establishes BQP as a complexity class with specific defining properties: polynomial-time computation on a quantum Turing machine with bounded error. BQP

contains BPP (bounded-error probabilistic polynomial time, the classical analog) and is conjectured to extend beyond it (Shor’s algorithm being the canonical example of BQP computation believed not to be in BPP). The relationship of BQP to other complexity classes — its relationship to NP, PH, PSPACE, and others — is the subject of active research.

The conventional treatment takes BQP as defined by the quantum-computing model and asks empirical questions: what problems are in BQP? what are the boundaries of BQP relative to other classes? what specific algorithms achieve BQP computation? The framework’s content addresses a different question, one level below: *why* is BQP the relevant class for quantum computing in our physical universe?

The framework’s level. The framework’s content treats BQP as a *derived* feature of physical reality rather than as a primitive defined by the quantum-computing model. The derivation runs through the framework’s structural commitments: any embedded observer in a finite deterministic substratum satisfying C1-C3 has access to BQP computational resources, and no embedded observer exceeds BQP. The BQP class is therefore *the unique class accessible to embedded observers* — a structural feature of the framework’s universality class rather than an empirical property of specific computational architectures.

This is a different epistemic status than the conventional treatment. In the conventional treatment, BQP is a *useful concept* for organizing computational results, with its empirical applicability being a contingent feature of our physical universe. In the framework’s content, BQP is a *structural inevitability* — any physical universe with the framework’s commitments has BQP as the unique computational ceiling for embedded observers.

Why this matters. The framework’s level matters for two reasons.

The ceiling is principled rather than empirical. The conventional treatment leaves open the possibility that some future physical principle could extend computational access beyond BQP. The framework’s content closes this possibility: super-BQP computation would require an embedded observer to exceed the structural limits the framework identifies, which is mathematically impossible if the framework’s commitments are correct. The ceiling is a theorem, not an empirical bound.

The falsification condition is sharp. The conventional treatment makes no specific predictions about whether super-BQP computation is possible — it remains an open question subject to future discovery. The framework’s content predicts that no super-BQP computation is possible on any physically realizable hardware, with any demonstrated super-BQP capability constituting a falsification. The falsification condition is empirically actionable: future quantum computing experiments could in principle test the framework’s prediction.

The framework’s content in this chapter is therefore at a foundational level — a structural account of why BQP is the relevant complexity class, with em-

pirical predictions that distinguish the framework from alternative accounts of computational complexity in physical systems.

14.3 BQP is available: the lower bound

The lower bound of the BQP characterization theorem establishes that any embedded observer satisfying C1-C3 has access to BQP computational resources. The argument runs through the framework’s emergent quantum mechanics from Chapter 1: the observer’s reduced description is quantum-mechanical, and the standard quantum-computing toolkit (state preparation, universal gate set, Born-rule measurement) is available to the observer through standard operations on the visible sector.

The observer’s quantum toolkit. Chapter 1 derived the framework’s emergent quantum mechanics for any embedded observer satisfying C1-C3. The derivation produces the standard quantum-mechanical formalism: a Hilbert space of visible-sector states, unitary evolution under the reduced Hamiltonian, projective measurement following the Born rule, and the standard composition rules for combining subsystems. These features are present in the observer’s reduced description by structural necessity — they follow from C1-C3 without additional assumptions.

The standard quantum-computing toolkit is built on exactly these features. *State preparation* (producing a specific quantum state for computation) corresponds to projective measurement followed by post-selection — operations the observer has access to as part of the emergent quantum mechanics. *Gate application* (transforming the quantum state through unitary operations) corresponds to time evolution under controllable Hamiltonians — available wherever the observer can engineer the visible sector’s dynamics. *Measurement* (reading the computation’s output) corresponds to standard projective measurement — the framework’s fundamental observational operation.

Universal gate set. Quantum computing requires a universal gate set: a collection of single-qubit and two-qubit gates that can approximate any unitary transformation to arbitrary precision. The standard universal gate set includes the Hadamard gate, the phase gate, and the CNOT (controlled-NOT) gate, with these three gates generating the full unitary group through composition. Each of these gates is implementable in any system with controllable two-level subsystems and a tunable coupling — features that the framework’s emergent quantum mechanics provides for any embedded observer with sufficient experimental control.

The argument is concrete: an observer who can prepare two-level subsystems (qubits) in the visible sector, who can apply controllable Hamiltonians to those subsystems, and who can perform projective measurements on the subsystems has access to a universal gate set. These three capabilities follow from the framework’s emergent quantum mechanics without additional structural assumptions.

Any embedded observer satisfying C1-C3 therefore has access to universal quantum computation in principle.

Error correction. Practical BQP computation requires fault tolerance: the ability to perform long computations despite physical errors at each gate. Quantum error correction (surface codes, stabilizer codes, topological codes) provides the theoretical framework for fault tolerance with overhead scaling logarithmically in the desired computational accuracy. The framework’s emergent quantum mechanics provides the structural foundation: error correction operates on the quantum information in the visible sector, with redundant encoding protecting against local errors.

The framework’s content here is consistent with standard quantum-computing theory: error correction is structurally available to any embedded observer with sufficient control over the visible sector, with the specific overhead determined by the physical error rates (which depend on the framework’s specific physical parameters rather than on structural commitments at the framework level). Chapter 15 develops the framework’s content on engineered error correction with correlation-aware decoders that exploit the framework’s predicted non-Markovian noise structure.

The lower bound theorem. Combining the four elements — state preparation, universal gate set, measurement, error correction — establishes the lower bound:

Lower Bound Theorem. Any embedded observer satisfying C1-C3 has access to BQP computational resources. The emergent quantum mechanics provides state preparation, a universal gate set, and Born-rule measurement; error correction is structurally available; the cumulative result is access to polynomial-time quantum computation with bounded error.

The lower bound is therefore established structurally: BQP is not a contingent capability that some embedded observers happen to have but a *universal* capability of all embedded observers satisfying C1-C3. This is a structural feature of the framework’s universality class — any member of the class has access to BQP.

14.4 BQP cannot be exceeded: the upper bound

The upper bound of the BQP characterization theorem establishes that no embedded observer exceeds BQP. The argument runs through *informational completeness*: the quantum description exhausts the information accessible to the embedded observer through any operation on the visible sector. No further computational capability is available, because there is no additional information for further computation to access.

Informational completeness of the quantum description. The framework’s emergent quantum mechanics from Chapter 1 has a specific structural property: the density matrix ρ_V describing the visible sector contains *all* the

information accessible to the embedded observer about the visible sector’s state. Any operation the observer performs on the visible sector — measurement, transformation, post-selection — operates on the density matrix and produces outcomes consistent with the density matrix’s statistical predictions.

There is no additional information for the observer to access. The hidden sector’s specific state is structurally inaccessible to the embedded observer (Chapter 1 §1.4); the observer can detect the *statistical* effect of the hidden sector through the density matrix’s mixedness but cannot detect the specific hidden-sector configuration. Any computational capability beyond what the density matrix supports would require the observer to access hidden-sector information that the framework’s structural commitments make inaccessible.

All observer operations are quantum operations. The framework’s content includes a more specific structural claim: any operation an embedded observer can perform on the visible sector is a *quantum operation* in the standard sense — a completely positive trace-preserving (CPTP) map on the density matrix ρ_V . This includes unitary evolution (a special case of CPTP map preserving purity), projective measurement (CPTP maps with specific structure), post-selection (non-CPTP but reduces to CPTP through normalization), and any combination of these.

The argument runs as follows. The embedded observer’s operations are implemented through interactions between the observer’s measurement apparatus and the visible sector. The apparatus itself is part of the visible sector (it is composed of visible-sector degrees of freedom). Any operation the apparatus performs is therefore a unitary evolution of the joint apparatus-system Hilbert space followed by projective measurement of the apparatus state. By the standard derivation of CPTP maps from such constructions, every operation the observer can perform on the visible sector is a CPTP map.

Computational complexity of CPTP maps. The computational complexity of CPTP maps is well-characterized in quantum complexity theory: any CPTP map on n qubits can be implemented by a polynomial-size quantum circuit (a sequence of polynomial many gates from a universal gate set). Conversely, any polynomial-size quantum circuit implements a CPTP map. The class of efficiently-implementable CPTP maps coincides with BQP up to the standard polynomial encoding.

The framework’s content therefore identifies BQP with the class of operations available to embedded observers. The observer cannot perform any operation that is not a CPTP map (because all observer operations are CPTP maps), and the class of efficiently-computable CPTP maps is BQP. Therefore the observer cannot perform any computation outside BQP.

The upper bound theorem.

Upper Bound Theorem. No embedded observer satisfying C1-C3 can perform computation outside BQP. Every observer operation is a CPTP map on the den-

sity matrix; the class of efficiently-implementable CPTP maps is BQP; therefore the class of efficient computations available to the observer is BQP.

Addressing objections. Three potential objections to the upper bound are worth addressing.

“What about hypercomputation?” Hypercomputation refers to computational models that exceed Turing-machine capabilities — infinite-time computation, oracle access to undecidable problems, analog computation with arbitrary-precision continuous quantities. The framework’s content addresses each: infinite-time computation requires actual infinity (which the framework’s finite substratum precludes), oracle access requires hidden-sector information the observer cannot extract, and arbitrary-precision continuous computation requires sub-Planck-scale resolution that the framework’s lattice precludes. Hypercomputation is structurally impossible for embedded observers.

“What about exotic physics?” Various proposed exotic physics scenarios — closed timelike curves, naked singularities, quantum gravity effects — have been suggested as potential routes to super-BQP computation. The framework’s content provides specific responses: closed timelike curves are inconsistent with the framework’s deterministic substratum dynamics; naked singularities are excluded by the framework’s bounded-coupling-degree commitment; quantum gravity effects operate at the substratum level rather than the emergent level and therefore cannot enhance computational power for embedded observers. Each exotic-physics route is closed by the framework’s structural commitments.

“What about future discoveries?” The framework’s content does not preclude future discoveries that revise our understanding of physical reality. What it does preclude is super-BQP computation *consistent with the framework’s commitments*. A future discovery that demonstrated super-BQP computation would constitute a falsification of the framework — it would require the universe to violate the framework’s structural conditions C1-C3 or one of the framework’s other commitments. The framework therefore makes a sharp empirical prediction: no super-BQP computation is possible on any physically realizable hardware compatible with the framework’s commitments.

14.5 The BQP characterization theorem

Combining the lower bound (BQP is available, §14.3) with the upper bound (BQP cannot be exceeded, §14.4) yields the BQP characterization theorem.

Statement. *BQP Characterization Theorem.* Let $\mathcal{S} = (S, \varphi)$ be a finite deterministic system, and let $\mathcal{O}$ be an embedded observer with partition V satisfying C1, C2, C3 at the partition boundary. Then the class of computations the observer can efficiently perform (in polynomial time with bounded error) is exactly BQP.

Proof outline. The lower bound from §14.3 establishes that BQP is available to the observer. The upper bound from §14.4 establishes that BQP cannot be

exceeded. Therefore the class of efficient computations available to the observer is exactly BQP.

The proof relies on the framework's structural commitments at multiple points:

- The framework's finite substratum guarantees finite-dimensional emergent Hilbert spaces.
- The framework's deterministic bijection φ guarantees unitarity at the substratum level.
- The framework's structural conditions C1-C3 guarantee the emergent quantum-mechanical description.
- The framework's partition structure guarantees informational completeness of the density matrix.

Each of these is a structural feature of the framework's universality class; together they produce the BQP characterization as a theorem.

Interpretation. The characterization theorem has several interpretive consequences.

BQP as the model-level computational class. For any universe in the framework's universality class, and under the model hypotheses above, BQP is the relevant computational complexity class for embedded observers. The specific physical implementation of the computational hardware does not matter — superconducting qubits, trapped ions, topological qubits, photonic systems, neutral atoms, or any other platform — all have access to BQP and only BQP.

BQP is independent of the specific bijection. The framework's universality class contains many bijections producing different specific physical content (different particle spectra, different gauge groups, different specific parameter values). Among the bijections within the framework's universality class, however, the computational ceiling is the same: BQP. The characterization is at the framework level rather than the specific-bijection level.

BQP is the ceiling for any future technology. No future technological development — better qubits, better error correction, exotic computational substrates — can exceed BQP for embedded observers in the framework's universality class. The ceiling is structural rather than technological.

The conditional model theorem is therefore a substantial but hypothesis-bearing result: given the model, it establishes BQP as the computational class accessible to embedded observers, holding at the structural level of the framework's universality-class commitments rather than at the specific-implementation level — and only as far as the model hypotheses reach.

14.6 The Extended Church-Turing Thesis as a conditional theorem

The Extended Church-Turing Thesis (ECT) asserts that any physically realizable computation can be efficiently simulated by a quantum Turing machine. Conventionally, the ECT is treated as an *empirical conjecture* about physical reality: the conjecture is not derived from any deeper principle but is supported

by the absence of known super-BQP computations and by general considerations about the structure of physical theories. The conjecture’s empirical status leaves open the possibility that future discoveries could refute it by demonstrating super-BQP computation in some physical system.

The framework’s content reframes the ECT from empirical conjecture to *theorem*.

The framework’s argument. The reframing runs through three steps.

Step 1: Every physical observer is embedded. The cosmological horizon provides a partition for any observer in the universe — the partition between the observable universe (visible to the observer in principle, given sufficient instruments) and the trans-horizon region (causally disconnected from the observer). The observer is therefore embedded in the framework’s structural sense: the partition with C1-C3 structure exists for any observer in the universe.

Step 2: Every embedded observer has access to BQP and only BQP. This is the BQP characterization theorem from §14.5. Any observer satisfying C1-C3 has access to BQP computational resources (the lower bound) and cannot exceed BQP (the upper bound). The computational class accessible to the observer is exactly BQP.

Step 3: Every physical computation is in BQP. Combining the previous two steps: every physical computation is performed by some physical observer or apparatus, every physical observer or apparatus is embedded in the framework’s structural sense, every embedded observer can compute only in BQP. Therefore every physical computation is in BQP.

The ECT is therefore not an empirical conjecture but a *theorem* of the framework — a derived consequence of the framework’s structural commitments. The reframing changes the ECT’s epistemic status fundamentally: the question “is the ECT correct?” becomes “are the framework’s structural commitments correct?”, which is testable through the framework’s other empirical content (Chapters 5-13).

Implications of the reframing.

Sharper predictions. The reframing produces sharper empirical predictions. Any demonstrated super-BQP computation would constitute a falsification of the framework — the universe would not be in the framework’s universality class, which would falsify the framework’s structural commitments at the foundational level. This is a stronger prediction than the conventional ECT, which allows the possibility of future revisions without committing to a specific structural commitment.

Structural foundations. The reframing provides structural foundations for the ECT that were previously absent. The ECT’s empirical status conventional treatment provides no explanation for *why* the ECT should hold — it remains an unexplained regularity of physical reality. The framework’s content provides the

explanation: the ECT holds because every physical observer is embedded, every embedded observer has access to BQP and only BQP, therefore every physical computation is in BQP. The structural basis is the framework’s universality-class commitments.

Connection to other foundational results. The reframing connects the ECT to other foundational results in the framework. The Bekenstein-Hawking entropy bound (Chapter 7 §7.5) limits the information content of any region; the BQP characterization limits the computational power of any embedded observer; both follow from the framework’s commitments at the substratum level. The connection unifies the information-theoretic and computational bounds under a common structural framework.

Connection to Deutsch’s jump-to-universality concept. The reframing also instantiates a specific structural pattern Deutsch identifies in [Deutsch]: the *jump to universality*, in which a sufficiently rich system suddenly acquires unbounded reach within its domain. Alphabets jump to universal writing; DNA jumps to universal genetic encoding; Turing computation jumps to universal computation. The framework’s BQP characterization gives precise content to one such jump in physical computation: a finite deterministic substratum satisfying C1-C3 achieves the BQP computational ceiling — universal physical computation up to that ceiling and no further. The promotion of the ECT from empirical conjecture to derived theorem is the framework’s specific contribution to this pattern: the jump to universality in physical computation reaches BQP and stops there, with the stopping point itself being a structural feature of the framework’s commitments rather than an empirical regularity.

The framework’s reframing of the ECT is therefore not a minor terminological change but a substantive theoretical contribution. The ECT moves from empirical conjecture to derived theorem, with the framework’s other empirical content (Chapters 5-13) providing the supporting evidence for the framework’s structural commitments that make the theorem follow.

14.7 The BPP-BQP gap as cost of partial observation

Quantum advantage — the gap between classical (BPP) and quantum (BQP) computational complexity — is one of the central topics in quantum complexity theory. The framework’s content provides a structural reading: the BPP-BQP gap is the *cost of not recognizing the quantum nature of the reduced description*. An embedded observer who treats the visible sector classically (sampling measurement outcomes according to classical probabilities) operates at BPP; an observer who exploits the quantum structure (interference, entanglement, adaptive measurement) operates at BQP. The gap between the two is the cost of partial observation interpreted classically.

The gap is real. Standard quantum complexity theory establishes that BQP contains problems believed not to be in BPP — Shor’s factoring algorithm being the canonical example. Whether $BPP \subsetneq BQP$ strictly is an open problem

in complexity theory; the standard conjecture is that the strict inclusion holds, with super-polynomial separations like Shor’s algorithm providing strong evidence. The framework’s content does not address whether the strict inclusion holds — that is a question about the structure of BQP and BPP as complexity classes, addressed by standard quantum complexity theory.

The gap’s structural origin. What the framework adds is a *structural* account of the gap’s origin. The embedded observer’s reduced description is quantum-mechanical (Chapter 1). An observer who *recognizes* the quantum-mechanical structure of the reduced description can perform quantum computation, exploiting interference, entanglement, and measurement back-action. An observer who *does not recognize* the quantum-mechanical structure — who treats the reduced description as classical with stochastic measurement outcomes — operates with reduced computational capability.

The gap between BPP and BQP is therefore not a property of physical reality per se but of the *observer’s representation of physical reality*. An observer using the correct quantum-mechanical description has access to BQP; an observer using an incorrect classical-stochastic description has access only to BPP. The framework’s content reframes quantum advantage as a *epistemic gap* between two descriptions of the same physical reality.

The framework’s reading vs the “quantum nature of reality” reading. A common interpretive position is that quantum advantage demonstrates “the fundamentally quantum nature of physical reality” — that physical reality is quantum-mechanical at the substratum level, with classical computation being a restricted special case. The framework’s content rejects this interpretive position. Under the framework’s reading, physical reality is *deterministic and classical* at the substratum level (the bijection φ acts on a finite state space with classical probability for the observer’s outcomes), with quantum mechanics being the *emergent* description that embedded observers use.

The BPP-BQP gap is therefore not evidence that physical reality is fundamentally quantum but evidence that *embedded observers in classical substrata have access to quantum-mechanical resources through the partition structure*. The quantum advantage is real but is a feature of the observer’s situation (embedded in a substratum with C1-C3 structure) rather than of the substratum’s fundamental nature.

This reading has substantive consequences for foundational questions. The framework’s content is compatible with classical substrate metaphysics — the universe is a deterministic finite system at the substratum level — while preserving the empirical content of quantum mechanics and quantum computing. The framework is not “quantum mechanics is wrong” but “quantum mechanics is the correct description of what embedded observers see in classical substrata.”

Implications for quantum-computing development. The framework’s reading has practical consequences for quantum-computing engineering. If quantum advantage reflects the observer’s exploitation of partition structure rather

than fundamentally quantum reality, then quantum-computing hardware should be designed to optimize the partition's properties — making the framework's C1-C3 architecture explicit and tunable rather than treating the environment as noise to be minimized. Chapter 15 develops this design philosophy in detail.

14.8 Internal structure of BQP and the P vs NP question

The BQP characterization theorem identifies the *ceiling* on computational capability for embedded observers but takes no position on the internal structure of complexity classes inside that ceiling. The relationship between P, NP, and BQP — and in particular the status of the P vs NP question itself — admits a precise reading in the framework's terms: the framework's structural commitments are consistent with either answer to the P vs NP question, and the locus of falsification for the framework's computational content sits at the BQP boundary rather than at any internal boundary within BQP.

The location of P vs NP within the framework. The framework's BQP characterization places P and NP inside BQP: every problem in P is in BQP (trivially, since BQP contains BPP and BPP contains P), and every problem in NP whose witness can be verified efficiently has its verification problem in BQP as well. The conventional inclusion chain $P \subseteq NP \subseteq PH \subseteq PSPACE$, with BQP situated somewhere within PSPACE and conjectured to contain BPP strictly, is unmodified by the framework's content. The framework's contribution is at a different level: it identifies BQP as the *outermost* class accessible to embedded observers, leaving the inclusion relationships among interior classes — including P vs NP — to standard complexity theory.

The neutrality result. The framework's commitments do not entail either $P = NP$ or $P \neq NP$. Two cases:

Case 1: $P \neq NP$ (the conventional expectation). The framework's BQP ceiling is unchanged: P and NP are both inside BQP, and no efficient general solution exists for NP-complete problems. Quantum advantage (Shor's algorithm and related results) places specific NP problems inside BQP without placing all of NP inside P, consistent with the standard conjecture.

Case 2: $P = NP$ (the surprising case). A polynomial-time algorithm for an NP-complete problem operates on a classical Turing machine; its outputs are in P, and by inclusion in BQP. The framework's BQP ceiling is unchanged: the new algorithm provides additional computational reach for embedded observers but does not exceed BQP. The framework would not be falsified by this discovery.

Both cases sit inside the framework's BQP ceiling. The framework's empirical exposure on P vs NP is zero: no observation an embedded observer could make would distinguish the two cases at the level of the framework's commitments.

Relation to Methodology §19's gauge / computational incompleteness distinction. The §19 distinction is worth applying carefully here. P vs NP is a determinate mathematical question with a definite answer — the abstract fact

about whether a polynomial-time algorithm for SAT exists is either true or false, even though we do not know which. The question is not vacuous in the gauge sense (it has answer-bearing content). What the framework’s content adds is that the determinate mathematical answer, whichever it is, is *physics-invariant* relative to the framework: the same empirical content (BQP as the ceiling) is consistent with both answers. This is weaker than gauge — not “no fact of the matter” but “the fact carries no consequence for the framework’s empirical claims.” The framework is uncommitted on P vs NP, and a mathematical proof of either answer would leave the framework’s empirical predictions unchanged.

The silence is structural, not stipulative. The previous paragraph’s neutrality result can be sharpened. The framework’s two axioms (tokened differentiation occurs; differentiation recurs) and their derived structural commitments — finiteness-per-moment, bijectivity, the proper-part decomposition $V \subsetneq S$, the C1-C3 partition conditions — deliver content of one specific type: they characterize the boundary structure between substrate and embedded observer, and thereby characterize the computational ceiling at that boundary. They do not deliver content of the type required to bear on internal complexity-class structure. P vs NP is a question about the existence of polynomial-time algorithms for NP-complete problem families, defined asymptotically over inputs of unbounded size; the framework’s axiomatic content concerns finite physical structures (with possibly unbounded history but always finite at any moment). The bridge between these two domains has a structural gap: physical-substrate axioms about finite structures do not entail asymptotic statements about algorithmic complexity classes without an additional bridging premise that the axioms themselves do not provide. The framework’s silence on P vs NP is therefore not a chosen agnosticism nor a residual of incomplete development — it is a feature of what kind of question P vs NP is relative to what kind of content the framework’s axioms deliver. A foundational program whose axioms *did* constrain P vs NP would require axioms of a different type, bridging physical substrate to asymptotic algorithmic structure directly; whether such axioms can be principled rather than question-begging is itself an open methodological question, addressed briefly below.

Position relative to Aaronson’s physical-principle approach. Aaronson [Aaronson2005b] has argued that the physical universe respects the principle that NP-complete problems are not solvable in polynomial time on physically realizable hardware — that $P \neq NP$ holds with the status of a physical principle. The framework is stronger and weaker than this position in different directions. *Stronger:* the framework commits to BQP exactly as the embedded-observer ceiling, where Aaronson’s principle leaves the ceiling between P and BQP undetermined. *Weaker:* the framework declines to commit to $P \neq NP$ itself, where Aaronson’s principle entails it. The framework’s empirical exposure is concentrated at the BQP boundary — where super-BQP capability would falsify it — and is zero at the P vs NP boundary, where neither answer constitutes empirical evidence about the framework’s commitments.

What axioms would close the gap, and why none are obviously available. The structural-silence result above identifies a gap between the framework’s axiom content (physical substrate, finite at each moment) and complexity-class content (asymptotic algorithmic structure over unbounded input families). It is natural to ask what additional axiom would close this gap. Several candidate forms are available, each illustrating why the gap is hard to close principled. A *computational-irreducibility axiom* asserting that most NP-complete instances admit no shortcut faster than brute-force enumeration would force $P \neq NP$ if made precise, but the precise version tends either to reduce to a complexity-class statement (circular) or to lose its bite. A *computational-manifestation axiom* asserting that complexity classes have physical instantiations would force conclusions about P vs NP only if combined with further constraints on what is physically realizable, and those constraints are essentially the question itself in disguise. A *bridging axiom* directly positing a bound on physically realizable computation — e.g., “all physical computation is in P” or “no physical process solves NP-complete problems in polynomial time” — would settle P vs NP by stipulation rather than derivation; this is the structure of Aaronson’s physical-principle move, and it is *posited as a principle* rather than derived from physics axioms because the latter cannot deliver it. Aaronson is explicit about this status. The structural feature common to these candidates is that the bridge from “finite physical structure at each moment” to “asymptotic algorithmic complexity over all input sizes” requires content that the standard repertoire of physical axioms does not contain. Whether a principled bridging axiom exists — one that is independently motivated rather than question-begging, and that does not collapse to a known open conjecture — is itself an open methodological question. The framework’s stance is that BQP characterization is the legitimate physics-principled result within reach of physical-substrate axioms, and that internal complexity-class structure (including P vs NP) is the domain of mathematical complexity theory rather than of physical foundations. Other foundational programs are free to attempt the bridging axiom; the framework declines to attempt it, on grounds that the candidates available appear question-begging.

Implication for proof strategies. The framework’s neutrality on P vs NP has a methodological consequence for foundational arguments about complexity classes. A proof strategy for $P \neq NP$ that runs through the BQP characterization theorem cannot work, because the framework’s commitments are equally compatible with either answer. Conversely, a proof of $P = NP$ (should one ever be produced) would not falsify the framework. The framework offers leverage only on the BQP-boundary question — the question the BQP characterization theorem and the Extended Church-Turing Thesis as conditional theorem address directly.

The boundary-characterization pattern across the Gödel-Turing-OI-Deutsch family. The framework’s silence on P vs NP fits a structural pattern that runs across the entire incompleteness family identified in Chapter 12 §12.8. Each of the four members establishes a structural limit at the *boundary* of access

from inside a system, and each is silent on the *internal organization* of what falls within the accessible domain. The pattern is not coincidence — it reflects a common feature of structural-incompleteness frameworks. Each identifies a *mode of access* and characterizes its limits; the internal complexity of what is accessible is a different question requiring tools internal to the relevant domain.

Concretely:

- *Gödel* sets a boundary at provability from within a consistent recursively-axiomatizable formal system. Gödel’s incompleteness theorems are about what crosses that boundary (true statements that are not provable). They are silent on the internal organization of the provable — which provable statements are interesting, useful, or important is not a question Gödel’s theorems answer.
- *Turing* sets a boundary at computability by a Turing machine. The halting problem statement is about what crosses that boundary (instances whose halting cannot be decided). Turing’s theorems are silent on the internal organization of the computable — which decidable problems are tractable, which are intractable, is not a question Turing’s theorems answer.
- *OI* sets a boundary at embedded observability of the substratum (the hidden-sector state h , the cosmological horizon, the substratum representative). The OI incompleteness statements are about what crosses that boundary. They are silent on the internal organization of the observable — which physical phenomena have particular structure within the framework’s accessible domain is content of a different kind, addressed by the framework’s specific structural commitments (gauge group, generations, partition architecture) rather than by the incompleteness statement itself.
- *Deutsch* sets a boundary at universal explanation in the constructor-theoretic sense. Deutsch’s universality-of-explanation is about what crosses that boundary (explanations attainable by sufficient knowledge-creation). It is silent on which problems within explanatory reach get solved in which order, by which methods.

P vs NP lies in the internal structure of the computable (within Turing’s domain) and in the internal structure of BQP (within OI’s domain). It is therefore outside the locus of contribution of *any* of these structural-incompleteness frameworks, not just OI. The unifying observation: structural-incompleteness frameworks characterize the *boundary* of a domain of access; the *internal complexity* of what is accessible is a separate question, requiring tools internal to the relevant complexity theory. Internal-complexity questions like P vs NP require complexity-theoretic tools (proof-system lower bounds, geometric complexity theory, circuit complexity, and similar), not incompleteness frameworks.

The framework’s silence on P vs NP is therefore *not* an idiosyncrasy of OI — it is the same kind of silence Gödel and Turing exhibit on the same question, for the same structural reason.

The division of labor is the appropriate one between the framework and complexity theory. The framework’s content is a structural theorem about the computational ceiling for embedded observers; standard complexity theory’s content is about the internal structure of the classes inside that ceiling. The two are complementary rather than competing.

14.9 Beyond axioms: research directions opened by the substratum structure

The structural-silence result of §14.8 identifies a negative content: physical-substrate axioms do not bridge to asymptotic complexity-class content. This should not be read as foreclosing all framework-relative complexity-theoretic content. The framework’s structural commitments suggest content not at the *universal* P vs NP boundary but at framework-relative *conditional* boundaries within the BQP ceiling. These are research directions where the framework’s specific substratum structure (cubic lattice, gauge group, partition structure) does provide content.

We sketch three such directions, each of which is a framework-relative refinement rather than a universal complexity-theoretic claim. The framework predicts none of them as new theorems; each is offered as a structurally motivated conjecture inviting complexity-theoretic investigation. The companion paper [Computation, §7] develops these in more detail and is the canonical reference for the research-direction content.

14.9.1 Geometric structure on the substratum lattice

NP-complete problems vary in the extent to which their constraint structure can be embedded into the substratum’s cubic lattice geometry. Problems with locally-bounded constraint structure on cubic-lattice-embeddable graphs (planar 3-SAT, lattice-Ising models, certain graph problems) admit natural substratum representations; problems whose constraints require highly nonlocal couplings do not. The framework’s conjecture in this direction is that lattice-local NP-complete problems admit *physical solution protocols* whose complexity reflects the locality structure — specifically, that physical implementation of such protocols incurs no overhead beyond the local-constraint structure, while nonlocal protocols incur overhead proportional to the nonlocality. This is not P vs NP; it is a refinement of how *implementation cost* on the framework’s substratum varies with problem geometry. It connects to existing complexity-theoretic work on planar variants of NP-complete problems and on local Hamiltonian complexity. The framework’s contribution is to suggest the substratum’s cubic geometry as a structural reason for the implementation-cost differential, not to deliver new bounds.

14.9.2 The BQP-BPP gap, hidden group structure, and the SU(2)-compatibility refinement

Existing complexity-theoretic evidence suggests that the BQP-BPP gap is largest for problems with strong global structure (factoring, discrete log, hidden-subgroup problems) and smallest for problems with primarily local structure (where classical heuristics often perform well). The framework's partition architecture suggests a structural reason for this correlation: BQP advantage derives from exploiting hidden-sector correlations that are induced by nonlocal substratum couplings, and nonlocal substratum couplings are precisely what global problem structure requires. The base conjecture is that the BQP-BPP gap correlates structurally with the *globality* of the problem's hidden-sector dependence.

A sharper structural reading is available when the conjecture is restricted to its appropriate scope and tested against the catalog of known BQP advantages.

Scope restriction. The conjecture as stated applies to *decision and search* problems with identifiable group-theoretic hidden structure — the hidden-subgroup problem (HSP) family and its relatives. Sampling problems (boson sampling, random circuit sampling) constitute a structurally distinct class whose advantage derives from sampling-vs-decision separations rather than from extracting hidden group properties; the framework's structural reasoning does not directly engage them, and they are left outside the conjecture's scope. Within the in-scope class, the conjecture's testable content survives.

Survey against the known advantage catalog. The known exponential BQP advantages with identifiable group-theoretic structure cluster as follows: Shor's factoring (HSP on $\mathbb{Z}_n^*$, abelian), discrete log (HSP on cyclic groups, abelian), Simon's problem and Deutsch-Jozsa (HSP on $(\mathbb{Z}_2)^n$, abelian), Forrelation (Fourier-correlation on $(\mathbb{Z}_2)^n$), period-finding (HSP on $\mathbb{Z}$), and Jones polynomial estimation at roots of unity (Aharonov-Jones-Landau 2006, with hidden structure in the braid-group representations factoring through $SU(2)_q$). On the negative side, large symmetric groups S_n and dihedral groups D_n for large n — both nonabelian with complex representation theory — have no known exponential quantum advantage despite extensive investigation; this is the empirical state of the nonabelian HSP problem.

The SU(2)-compatibility refinement. The structural reason behind this clustering becomes visible when one notes that the framework's quantum toolkit (from the QM emergence theorem of Chapter 1) is built from SU(2)-based operations: qubits are SU(2)-doublet representations; the universal gate set (Hadamard, phase, CNOT) acts on these doublets. The quantum Fourier transform — the workhorse of HSP-style algorithms — is efficiently implementable for groups whose representation theory factors through SU(2)-compatible structure: abelian groups embedding in U(1) (Shor, discrete log, period-finding), elementary 2-groups via SU(2)-derived $\mathbb{Z}_2$ structure (Simon, Deutsch-Jozsa, Forrelation), small symmetric groups via low-dimensional $SU(n)$ embeddings,

tractable nonabelian groups (affine, Heisenberg) with abelian-normal structure, and groups with quantum deformations of $SU(2)$ representation theory (Jones polynomial via $SU(2)_q$). Groups whose representation theory does not factor through $SU(2)$ -compatible structure — large S_n , large D_n , generic non-Lie combinatorial groups — appear to be precisely the cases where exponential quantum advantage is elusive.

Refined conjecture. Within the class of decision and search problems with identifiable group-theoretic hidden structure, exponential BQP advantage correlates with whether the hidden group’s representation theory is $SU(2)$ -*compatible* — that is, factors through $SU(2)$ representations or their quantum deformations in a way that the framework’s $SU(2)$ -based quantum toolkit can efficiently access. The framework’s broader gauge group $SU(3) \times SU(2) \times U(1)/\mathbb{Z}_6$ determines what physics is; the BQP-advantage correlation connects specifically to $SU(2)$ via the QM emergence theorem’s toolkit construction.

This is the structural reading of the BQP-BPP gap pattern that the framework’s content suggests. It survives the survey against known BQP advantages (twelve cases checked, all in-scope cases consistent), accommodates the tractable / intractable boundary in nonabelian HSP, and is grounded in the framework’s existing QM emergence theorem rather than being an external postulate. The conjecture is silent on sampling-style advantages (boson sampling, random circuit sampling) by scope, not because the framework is inconsistent with them.

Three concrete checks remain testable. (1) The known exponential-BQP-advantage decision/search problems should all admit hidden structure factoring through $SU(2)$ -compatible representation theory. (2) The tractable-vs-intractable boundary in nonabelian HSP should correspond to whether the hidden group has $SU(2)$ -compatible representation theory. (3) New BQP-advantage decision/search problems should be most discoverable by looking for hidden group structures with $SU(2)$ -compatible representation theory, including quantum-deformation variants. None of these is established by the framework; each is a structurally motivated conjecture inviting complexity-theoretic investigation. The framework’s contribution is to suggest a specific structural reason — the $SU(2)$ -based quantum toolkit construction — behind the empirical concentration of BQP advantage on the specific group structures it concentrates on, with this contribution rooted in the framework’s existing QM emergence theorem rather than introduced as new content.

14.9.3 Gauge equivalence and complexity-theoretic reductions

The framework’s substratum gauge group $\mathcal{G}_{\text{sub}}$ partitions the space of bijections into equivalence classes; the gauge group includes state relabeling, alphabet change, deep-sector enlargement, and graph isomorphism up to statistical isotropy (Substratum §4). A natural question is whether $\mathcal{G}_{\text{sub}}$ has an image in the space of polynomial-time reductions among NP-complete problems. If gauge equivalence at the substratum level corresponds (in some precise sense)

to polynomial-time reducibility at the algorithm level, the framework’s gauge structure would provide a *physical* account of why NP-complete problems are reducible to each other. This is speculative — the correspondence has not been made precise — but the structural similarity is suggestive: gauge equivalence is “no observable distinction” at the substratum level, and polynomial-time reducibility is “no complexity distinction” at the algorithm level. Both are equivalence relations on otherwise distinct structures.

14.9.4 Conditions and the right kind of research question

The pattern across the three directions is that the framework’s specific structure suggests *conditional* complexity-theoretic claims — claims of the form “for problems satisfying condition X (lattice-locality, global structure, gauge-equivalence to another problem), the framework’s structure suggests a specific complexity-theoretic behavior.” This is the appropriate level at which the framework engages complexity theory: not by attempting to bridge to universal questions like P vs NP, but by suggesting how the framework’s specific substratum architecture refines the internal landscape of BQP.

These conditional refinements are not predictions of the framework in the sense of the framework’s quantitative SM observables (Chapter 6); they are research directions opened by the framework’s substratum structure. Some may be productive, others may not. The framework’s contribution is to identify them as natural questions given its structure, not to resolve them. Specialist investigation in complexity theory is the appropriate way to develop or refute each direction.

14.10 Chapter close

The framework’s content in this chapter establishes, under the stated model hypotheses, BQP as the computational class accessible to embedded observers in the framework’s universality class. The conditional model theorem combines a lower bound (BQP is available to any embedded observer satisfying C1-C3) with an upper bound (BQP cannot be exceeded by any embedded observer), with the combination producing BQP as the exact ceiling on computational capability.

Four principal results.

The conditional BQP model theorem (§14.5): under the model hypotheses of §14.1, BQP is the computational class accessible to embedded observers satisfying C1-C3.

The Extended Church-Turing Thesis as conditional theorem (§14.6): every physical computation is in BQP under the model hypotheses of the BQP characterization, derived from the embedding rather than postulated as a freestanding empirical conjecture.

The BPP-BQP gap as cost of partial observation (§14.7): quantum advantage is real but reflects the embedded observer’s exploitation of partition structure rather than the universe’s fundamentally quantum nature.

Framework neutrality on P vs NP (§14.8): the framework’s empirical content is consistent with either answer to the P vs NP question; the locus of falsification sits at the BQP boundary, not at any boundary internal to BQP. The silence on P vs NP fits a broader pattern across the Gödel-Turing-OI-Deutsch incompleteness family: structural-incompleteness frameworks characterize boundaries of access and are silent on the internal structure of what falls within them.

Three research directions opened by the substratum structure (§14.9): geometric structure on the substratum lattice (cost differential for lattice-local versus nonlocal NP-complete problems); the BQP-BPP gap correlating with hidden group structure, with a refined conjecture grounded in the QM emergence theorem that exponential BQP advantage for decision/search problems with identifiable group-theoretic hidden structure correlates with $SU(2)$ -compatible representation theory (survey against twelve known BQP-advantage cases supports this within scope); gauge equivalence at the substratum level as a possible physical account of polynomial-time reducibility at the algorithm level. These are conjectures inviting complexity-theoretic investigation, not framework predictions. The companion paper [Computation] develops them as the canonical reference.

Falsification conditions. The framework’s content makes specific falsifiable predictions:

Any demonstrated super-BQP computation on physically realizable hardware would falsify the framework. The framework’s structural commitments preclude super-BQP computation; any empirical demonstration of such capability would refute the framework at the foundational level.

Hypercomputation is structurally impossible. The framework’s finite substratum and bounded-coupling-degree commitments preclude hypercomputation in any form (infinite-time, oracle-based, arbitrary-precision-continuous). Any apparent demonstration of hypercomputation would either be reducible to standard computation or constitute a falsification.

The BQP ceiling applies to all future technology. No improvement in quantum-computing technology — better qubits, better error correction, exotic architectures — can exceed BQP. The ceiling is structural rather than technological. The framework therefore predicts a specific limit on future quantum-computing capability.

Forward pointers. Chapter 15 develops the framework’s content in quantum engineering: the practical implementation of quantum computing in superconducting-qubit systems, the framework’s predicted non-Markovian noise structure, correlation-aware error correction with 3-10× overhead reduction, coherence extension with engineered baths, and the BEC analogue-gravity

experiment as a foundational test of the BQP characterization theorem. The empirical content develops in the engineering direction; the structural foundations are in this chapter.

Chapter 18 §§18.2-18.4 developed the foundational content connecting the framework's emergent quantum mechanics to the BQP computational ceiling. The Born rule as dynamical equilibrium, the quantum-classical boundary, and the epistemic character of quantum randomness — these foundational topics provide the substratum-level context for the BQP characterization theorem. Together with the chapter's content, they constitute the framework's complete account of computational complexity for embedded observers: a structural ceiling at BQP with the foundational basis in the framework's emergent quantum mechanics.

The chapter's content is therefore the framework's positive contribution to computational complexity theory: a structural characterization of BQP as the unique class accessible to embedded observers, with the ECT reframed from empirical conjecture to theorem and the BPP-BQP gap interpreted as the cost of partial observation. The framework's empirical exposure in computational complexity is principally negative — the prediction of no super-BQP computation provides a sharp falsification target — but the structural content is substantial.

One refinement sharpens the exposure from principally negative to positively dated. A finite substratum bounds the classical resources of any causal patch — about 10^{145} elementary operations for a day-long computation — and a quantum computation whose statistics would cost more than the budget to produce classically cannot have them produced at all. The consequence is a ceiling near five to six hundred full-fidelity logical qubits, binding only on tasks that are genuinely exponentially hard to simulate: factoring at practical key sizes fits comfortably inside the budget and survives (the crossover sits near 86-kilobit keys), while full-fidelity random-circuit sampling beyond the ceiling does not. The signature would be unlike any ordinary noise: fidelity tracking what the circuit *means* — its classical simulation cost — rather than how large it is, with Clifford-dominated circuits immune at any scale. Finding no such wall at a thousand high-fidelity logical qubits would falsify the classical substratum itself. And this is not an optional conjecture: the framework's own theorems — the maximal-entropy measure, the gauge status of the microstate — close every alternative, so the ceiling is entailed, not adopted.

The depth of computation is bounded too, and by the same construction. The residual dissipation — the framework's irreducible 10^{-64} deviation from perfect quantum behavior, sourced by the cosmological partition rather than by any lab's imperfections — caps uncorrected coherent computation near 10^{64} operations, an absolute decoherence floor no isolation can evade. Error correction erases that wall: the dissipation channel is quasi-static — a static conditioned Hamiltonian over the run, hence a common-mode coherent error of the benign class that calibration, dynamical decoupling, and error correction handle, thirty orders below threshold besides — so it does not break the fault-tolerance theo-

rem’s assumptions, and the 10^{32} cap is a property of uncorrected hardware only. Together with the width ceiling, the rate and memory bounds derivable from the same primitives, and the exclusion of spacetime hypercomputation tricks, the framework specifies the complete box of physical computation — every face with a number or a queued derivation, none negotiable.

Chapter 15

Quantum Engineering — Hardware, Software, and the BEC Experiment

15.1 What this chapter develops

Chapter 14 established the framework’s content in computational complexity theory: BQP, under that chapter’s explicit model hypotheses, as the computational class accessible to embedded observers, the Extended Church-Turing Thesis as a conditional result under the same hypotheses, the BPP-BQP gap as the cost of partial observation. The framework’s content there is theoretical at the foundational level — a structural characterization of what computations any embedded observer can perform, with falsification by any super-BQP demonstration.

This chapter develops the framework’s content in *engineered quantum systems* — the specific predictions for current and near-future quantum hardware platforms, the engineering implications of the framework’s C1-C3 architecture, and the cleanest near-term experimental test of the framework’s foundational characterization theorem. The framework’s distinctive prediction is that engineered quantum systems are not just noise-limited devices struggling against an adversarial environment but are *natural C1-C3 systems* whose dynamics is structurally non-Markovian. This reframing has concrete engineering consequences across superconducting qubit error correction, quantum sensing with NV centers, and the partially-quantum regime that current experiments are beginning to access.

Eight pieces occupy the chapter.

Section 15.2 develops the framework’s reframing for engineered quantum systems: the environment as programmable quantum memory rather than as adversarial noise. The framework’s C1-C3 architecture applies to engineered systems with τ_S/τ_B ratios at orders of magnitude where structural corrections are not negligible — making the framework’s content empirically relevant for current hardware platforms.

Section 15.3 develops superconducting qubit dynamics as the canonical engineered C1-C3 system. Transmons coupled to two-level systems (TLS) in the substrate produce a hidden sector with $\tau_S/\tau_B \sim 10^{-3}$ — twenty-nine orders of magnitude larger than the cosmological case. The framework’s structural prediction is that this hidden sector produces P-indivisible noise correlations across thousands of gate operations. Section 15.4 develops correlation-aware error correction: standard Markovian surface codes see the framework’s P-indivisible correlations as anomalously high uncorrectable error rates; correlation-aware decoders designed for the actual noise structure could reduce overhead by factors of three to ten, with concrete numerical predictions.

Section 15.5 develops coherence extension through engineered baths — the framework’s bath-as-resource design philosophy. By coupling a qubit to a designed slow hidden sector (a spin chain, a resonator array, or other engineered architecture), the framework predicts coherence times can be extended by factors approaching ten, with the bath’s correlation timescale controlling the coherence revival pattern. Section 15.6 develops quantum sensing with nitrogen-vacancy (NV) centers in diamond: the framework predicts sensitivity improvements approaching $10^3\times$ through exploitation of information backflow from the diamond lattice’s slow phonon and nuclear-spin baths.

Section 15.7 develops the partially-quantum regime in engineered systems — the framework’s distinctive prediction of dynamics distinct from both fully-quantum and fully-classical regimes, accessible at intermediate τ_S/τ_B values where C3 is marginal. Section 15.8 develops the BEC analogue-gravity experiment as the framework’s cleanest near-term test of the BQP characterization theorem at the foundational level: a tabletop experiment with engineered hidden-sector capacity that distinguishes the framework’s predictions from all known decoherence backgrounds.

The chapter closes with a discussion of the relationship between the framework’s engineering predictions and the DSW horizon decoherence program from Chapter 9 §9.7. The BEC analogue-gravity setup tests in a laboratory system what the DSW analysis derives theoretically for astrophysical horizons; the convergence between the framework’s content and the mainstream DSW program is significant. The framework’s reach into engineering is therefore not isolated but is part of a broader convergence in mainstream physics on horizon decoherence as a fundamental phenomenon.

15.2 The environment as programmable quantum memory

Standard quantum engineering treats the environment as an adversary: decoherence destroys quantum information, and the goal is to isolate qubits from their surroundings as completely as possible. Better materials, deeper cryogenics, more careful shielding — the implicit design philosophy is that quantum systems should approach the closed-system limit, with the environment as an unavoidable nuisance to be minimized.

The framework’s content suggests a fundamentally different design philosophy. Instead of isolating a qubit from all environmental coupling, *engineer* the C1-C3 architecture to produce the quantum behavior you want. The hidden sector is not noise but *programmable quantum memory* — a resource that stores correlations written during one gate operation and returns them during subsequent operations. The bath becomes a design variable rather than a constraint.

The structural argument. The framework’s emergent quantum mechanics from Chapter 1 derives the read-write cycle as the equilibrium dynamics between visible and hidden sectors. The cycle’s specific properties — the coherence timescale, the information-backflow pattern, the noise correlation structure — are determined by the C1-C3 architecture’s specific parameters: the coupling strength (C1), the timescale ratio τ_S/τ_B (C2), and the hidden-sector capacity N_H/N_V (C3). By controlling these parameters, an engineer controls the emergent quantum dynamics.

The implications are concrete. A qubit isolated from all environmental coupling has $\tau_S/\tau_B \rightarrow 0$ and behaves as an ideal closed quantum system at the qubit’s own timescale, with decoherence eventually limiting performance through residual coupling. A qubit deliberately coupled to a slow, high-capacity hidden sector with $\tau_S/\tau_B \sim 10^{-1}$ exhibits qualitatively different dynamics: information backflow, partial coherence revivals at the bath’s correlation timescale, and a non-Markovian noise structure that correlation-aware error correction can exploit.

Where the framework’s prediction reaches engineering relevance. The framework’s structural correction is $\mathcal{O}(\tau_S/\tau_B)$ deviations from standard Markovian models. At the cosmological scale, $\tau_S/\tau_B \sim 10^{-32}$ — irrelevant for any technology. At engineered scales, τ_S/τ_B can be substantially larger: $\sim 10^{-3}$ for superconducting qubits coupled to TLS baths, $\sim 10^{-6}$ for NV centers in diamond, and $\sim 10^{-1}$ for engineered spin-chain architectures. The framework’s prediction has reached engineering relevance precisely because τ_S/τ_B for engineered systems is many orders of magnitude larger than the cosmological case.

The cumulative pattern. The framework’s predictions across the engineered quantum systems summarize as follows.

Domain	System / Bath	τ_S/τ_B	Current relevance
Cosmology	Observer / trans-horizon	10^{-32}	Framework’s home ground
Enzymes	Active site / scaffold	10^{-9} to 10^{-12}	Chapter 13: consistent with data
Quantum sensing	NV spin / lattice defects	10^{-6}	Approaching relevance
Quantum comput- ing	Qubit / TLS	10^{-3}	At the threshold

Domain	System / Bath	τ_S/τ_B	Current relevance
Atomic clocks	Atom / cavity	$\sim 10^{-6}$	Potentially measurable
Engineered baths	Designed coupling	up to 10^{-1}	Future architectures

The corrections grow as τ_S/τ_B increases — equivalently, as the hidden sector’s timescale approaches the system’s timescale. The cosmological case has the most extreme separation (corrections negligible). Engineered systems have the least separation (corrections largest). The framework predicts the same structural phenomenon — P-indivisible dynamics from a slow bath — at every scale; the question in each case is whether the correction has reached the precision frontier where it matters for the application.

The remaining sections develop the framework’s specific engineering content at each of these scales, with concrete numerical predictions for current and near-future quantum hardware platforms.

15.3 Superconducting qubits and TLS noise

Superconducting qubits — transmons, fluxonium, and similar architectures — are the leading commercial platform for quantum computing. Companies including IBM, Google, Rigetti, and IQM have demonstrated processors with hundreds to thousands of physical qubits, with active programs targeting fault-tolerant quantum computing at scale. The framework’s content provides specific predictions for the noise structure that limits current performance and the design principles that could improve it.

The TLS bath. Superconducting qubits are fabricated as thin-film Josephson junctions on amorphous oxide substrates. The amorphous oxide contains *two-level systems* (TLS) — structural defects that flip between two configurations at characteristic rates determined by the local atomic environment. The TLS distribution is broad, with rates ranging from microseconds to seconds, and the TLS density is approximately 10^4 - 10^6 per cubic micrometer for typical aluminum-oxide barriers. A transmon qubit typically couples to several hundred to several thousand TLS through its electromagnetic field.

The framework’s C1-C3 conditions for the transmon-TLS system.

C1 (coupling). The TLS couple to the transmon through their electric dipoles, with coupling strengths varying from kHz to MHz depending on the TLS proximity to the qubit’s electric field maxima. The coupling is continuous and bidirectional — qubit excitations can transfer to TLS and vice versa.

C2 (slow bath). Transmon gate operations occur on $\tau_S \sim 10$ -100 ns timescales. The TLS bath has correlation times $\tau_B \sim 1$ -100 μ s. The ratio: $\tau_S/\tau_B \sim 10^{-3}$ for typical operating conditions. This is twenty-nine orders of magnitude larger

than the cosmological case, placing transmon-TLS dynamics squarely in the framework’s slow-bath regime where structural corrections are non-negligible.

C3 (capacity). Several hundred to several thousand coupled TLS produce a hidden sector with capacity vastly exceeding the qubit’s two-state Hilbert space. C3 is satisfied with overwhelming margin.

The framework’s prediction: P-indivisible noise. With C1-C3 satisfied at $\tau_S/\tau_B \sim 10^{-3}$, the framework predicts that the qubit’s noise structure is P-indivisible: error probabilities at gate k depend on the error history at gates $k-1, k-2, \dots$ through correlations stored in the TLS bath. The correlations extend over approximately $\tau_B/\tau_S \sim 10^3$ gates — meaning noise events separated by less than one thousand gate operations are statistically correlated rather than independent.

This prediction is structurally distinct from the standard Markovian assumption underlying quantum error correction. Markovian noise models assume each error is statistically independent of previous errors; the framework’s prediction is that this assumption is wrong by orders of magnitude on the timescale relevant for fault-tolerant quantum computing. The standard noise characterization techniques (randomized benchmarking, gate set tomography, cycle benchmarking) measure averaged error rates and are largely insensitive to correlation structure — they would miss the framework’s prediction by design.

Empirical signatures. The framework’s prediction has specific empirical signatures distinguishable from standard noise characterization.

Correlation functions across gate sequences. The two-point correlation $\langle \delta E_k \cdot \delta E_{k+m} \rangle$ between error events at gates k and $k+m$ should be non-zero for $m \lesssim 10^3$, decaying on the timescale $\tau_B/\tau_S \sim 10^3$ gates. Standard Markovian models predict $\langle \delta E_k \cdot \delta E_{k+m} \rangle = 0$ for $m \neq 0$.

Higher-order correlation functions. Three-point and higher correlation functions should also be non-zero with characteristic structures determined by the TLS bath dynamics. These higher-order correlations are uniquely diagnostic of P-indivisible (genuinely non-Markovian) dynamics rather than just non-Markovian correlations from broad rate distributions.

Surface-code logical error patterns. In a surface code operating under the framework’s predicted noise, logical errors should cluster in space and time with correlation patterns determined by the TLS bath structure. Standard surface-code analysis predicts spatially and temporally uniform logical errors; the framework predicts clustering.

The empirical record so far is largely consistent with the framework’s prediction. Multiple independent groups have reported anomalously correlated errors in superconducting qubit benchmarks (Wilen et al. 2021 on cosmic-ray-induced correlated errors; McEwen et al. 2022 on correlated errors from quasiparticle dynamics; Google Quantum AI 2023 surface-code error analysis with anomalous spatially correlated error patterns). The framework’s reading is that these

are not isolated phenomena but generic signatures of the C1-C3 architecture inherent in the transmon-TLS system.

15.4 Correlation-aware error correction

The framework’s prediction of P-indivisible noise in superconducting qubits has direct implications for quantum error correction. Standard surface-code decoders assume Markovian (independent) errors. The framework’s predicted correlations imply that standard decoders are systematically suboptimal — they treat correlated errors as random, missing the predictability the bath dynamics provides.

The framework’s quantitative prediction. For superconducting qubits with $\tau_S/\tau_B \sim 10^{-3}$, the framework predicts correlated errors extending over $\sim 10^3$ gates. A Markovian decoder sees these correlations as a higher effective error rate; it *underperforms* relative to the actual physical error rate of the device. A correlation-aware decoder — one designed for the framework’s predicted P-indivisible noise structure — exploits the predictability: correlated errors are easier to correct than random errors because their correlation structure provides additional information about which physical errors caused the observed syndrome pattern.

The framework’s estimate of the achievable improvement is a factor of three to ten reduction in the effective error rate at the same physical error rate. At current physical error rates of approximately 10^{-3} per gate, this corresponds to a reduction in fault-tolerance overhead from approximately 3,000 physical qubits per logical qubit to approximately 300-1,000 — a factor of three to ten reduction in the hardware cost of fault-tolerant quantum computing.

The empirical accessibility. Correlation-aware decoders are testable on current quantum hardware without requiring new platforms or technologies. The required steps are: (i) characterize the noise correlation structure on a specific device through extended benchmarking experiments designed to measure two-point and higher-order correlation functions; (ii) construct decoder algorithms that incorporate the measured correlation structure as prior information rather than treating all errors as independent; (iii) compare the resulting logical error rates against standard Markovian decoders on the same hardware.

The framework’s prediction is testable now, with current physical error rates and current decoder implementations. A factor-of-three improvement at the current state of the art would have substantial commercial impact — reducing the hardware cost of any planned fault-tolerant system by the same factor. The framework’s prediction is therefore not just theoretical content but a specific commercial design target with empirical accessibility on existing hardware.

Falsification conditions. The framework’s prediction would be falsified by either (i) demonstration that no correlation-aware decoder improves on the standard Markovian decoder by a factor exceeding ~ 1.5 across multiple hardware

platforms (suggesting the framework’s predicted correlations are too small to be exploited), or (ii) demonstration that the correlation structure observed on real hardware is qualitatively different from the framework’s prediction (suggesting the TLS bath dynamics is not the dominant source of correlations).

Neither falsification has been demonstrated. The empirical record on correlated errors in superconducting qubits is consistent with the framework’s prediction, and explicit tests of correlation-aware decoders against the framework’s specific predictions are not yet in the literature. The framework’s prediction therefore stands as a pre-registered empirical target for the next generation of quantum-error-correction research.

15.5 Coherence extension through engineered baths

Standard quantum hardware aims to minimize environmental coupling to maximize coherence time. The framework’s content suggests the opposite design: deliberate coupling to an engineered bath can *extend* coherence through information backflow, with the bath’s correlation timescale determining the coherence revival pattern.

The mechanism. A qubit coupled to a slow, high-capacity hidden sector evolves through the read-write cycle developed in Chapter 1. During one operational cycle, coherence transfers from the qubit (visible sector) into the bath (hidden sector). The bath stores the coherence as correlations between its degrees of freedom. On the bath’s correlation timescale τ_B , the correlations return to the qubit as information backflow, partially restoring coherence.

Standard coherence decays as e^{-t/T_2} . With information backflow from an engineered bath, coherence partially revives at multiples of τ_B . The revival amplitude per cycle is $\mathcal{O}(\tau_S/\tau_B)$. For natural systems with $\tau_S/\tau_B \sim 10^{-3}$, revival is approximately 0.1% — too small to be useful. For an engineered system with $\tau_S/\tau_B \sim 10^{-1}$ (achievable with a qubit coupled to a short spin chain or resonator array), revival is approximately 10% per cycle.

Over multiple revival cycles before the bath fully relaxes, coherence persists longer than a Markovian estimate would give. The framework does not deliver the extension factor: τ_B/τ_S is a ratio of substratum timescales, and identifying it with an achievable T_2 ratio presumes the revivals survive the device’s other decoherence channels — a platform-specific question. No numerical T_2 target is claimed.

Concrete platform: qubit coupled to a spin chain. A specific platform implementing the framework’s prediction is a qubit coupled to a linear chain of L ancillary spins. The chain serves as the engineered hidden sector. Its correlation timescale scales as $\tau_B \sim L^z$ (where z is the dynamical exponent — for a Heisenberg chain, $z = 1$). Its capacity scales exponentially as approximately 2^L states.

By tuning the chain length L , the framework’s prediction is:

- Short chain ($L \lesssim 5$, $\tau_B \sim \tau_S$): Markovian regime. Standard decoherence. The bath equilibrates between gate operations.
- Long chain ($L \gtrsim 20$, $\tau_B \gg \tau_S$): P-indivisible regime. Information backflow. Correlations from gate k return at gate $k + \tau_B/\tau_S$, partially restoring coherence.

The transition is sharp: the P-indivisibility theorem from Chapter 1 guarantees that once τ_B/τ_S exceeds the C2 threshold, the dynamics becomes qualitatively non-Markovian rather than gradually so.

The platform is testable on existing technologies: superconducting qubits coupled to engineered spin chains, NV centers in diamond with controllable nuclear spin baths, or trapped ions with engineered phonon modes. The framework’s prediction is that at the critical chain length, the qubit’s T_2 coherence time shows a *qualitative* change — from monotonic exponential decay to oscillatory behavior with partial revivals at multiples of τ_B .

Implications for hardware development. The bath-as-resource philosophy has implications for hardware-development priorities. Current investment is concentrated on isolation and noise reduction. The framework’s prediction is that complementary investment in engineered-bath architectures could produce comparable or larger gains at lower hardware cost. Whether this prediction generates commercial uptake depends on whether early demonstrations of the framework’s predicted coherence extension are achieved at sufficient scale to be competitive with continued isolation improvements.

15.6 Quantum sensing with NV centers

Nitrogen-vacancy (NV) centers in diamond are the leading platform for room-temperature quantum sensing — magnetic field measurement, temperature sensing, biological imaging, and related applications. Current NV sensitivity is limited by the spin coherence time T_2 of the NV electron spin, with standard Markovian models predicting that improvements in sensitivity require improvements in T_2 . The framework predicts that sensitivity can be improved beyond the Markovian T_2 limit through exploitation of information backflow from the diamond lattice’s slow phonon and nuclear-spin baths.

The C1-C3 architecture of NV centers. The NV center’s electron spin couples to multiple baths in the diamond lattice: phonons (acoustic vibrations of the diamond crystal), nuclear spins (^{13}C nuclei in natural-isotope diamond, P nuclei in synthetic diamond), and paramagnetic impurities (other defect centers). Each bath contributes to decoherence at characteristic timescales.

The relevant timescale separations are: NV electron-spin dynamics at $\tau_S \sim 10$ -100 ns; phonon-bath correlation times $\tau_B \sim 1$ -10 ms; nuclear-spin bath correlation times $\tau_B \sim 1$ -100 ms. The ratio: $\tau_S/\tau_B \sim 10^{-5}$ to 10^{-6} — within the framework’s slow-bath regime where information backflow effects become predictable.

The framework’s prediction. Decoherence in NV centers is non-Markovian on the framework’s account: coherence partially revives at times set by the bath’s correlation timescale. Information that appears to be lost from the NV during the apparent decoherence phase is in fact stored as correlations in the bath, and these correlations return to the NV during the revival phase.

Sensing protocols that exploit the revival — measuring at the backflow time rather than during the initial decay — could push sensitivity beyond the Markovian-assumed T_2 limit. The improvement factor per measurement is $\mathcal{O}(\tau_S/\tau_B) \sim 10^{-6}$, but sensing protocols typically accumulate over $\sim 10^6$ repetitions, making the integrated improvement potentially significant.

Quantitative prediction. Standard NV magnetometry sensitivity is $\delta B \sim 1/(\gamma\sqrt{T_2 \cdot T_{\text{meas}}})$, where γ is the gyromagnetic ratio and T_{meas} is the total measurement time. Non-Markovian backflow at τ_B means each measurement recovers information that Markovian protocols assume is lost. For $\tau_S/\tau_B \sim 10^{-6}$ (NV center with paramagnetic bath), each measurement carries approximately 10^{-6} extra information. Over 10^6 repetitions, this integrates to $\mathcal{O}(1)$ — a factor of $\sqrt{\tau_B/\tau_S} \sim 10^3$ improvement in signal-to-noise if the protocol is optimized for the backflow timing.

A sensitivity gain follows in *direction* but not in size: the framework fixes neither the achievable interrogation time nor the readout efficiency, and both enter δB directly, so no numerical improvement factor is claimed here. Quoting one would require an NV-specific model of the spin bath, readout chain and control sequence. This would make NV sensors competitive with superconducting quantum interference devices (SQUIDS) but operating at room temperature rather than requiring cryogenic cooling — a substantial commercial advantage if achieved.

Empirical accessibility and falsification. The framework’s prediction is testable on current NV-magnetometry platforms through development of backflow-aware sensing protocols. The required steps are: (i) characterize the bath correlation timescales on a specific NV-diamond sample; (ii) design measurement protocols that read out the NV at τ_B rather than at the standard $T_2/2$ timing; (iii) compare integrated sensitivity against standard protocols on the same hardware.

A large improvement in NV sensitivity at room temperature would be transformative for biological imaging, medical diagnostics, and geophysical sensing applications. The framework’s prediction is therefore not isolated theoretical content but a specific commercial design target with empirical accessibility on existing hardware. Falsification would come from demonstration that backflow-aware protocols do not improve integrated sensitivity by factors exceeding ~ 10 across multiple NV platforms — suggesting the framework’s predicted backflow effects are too small or too thermally averaged to be exploited.

15.7 The partially-quantum regime in engineered systems

Beyond the corrections to fully-quantum behavior at large τ_B/τ_S (developed in §§15.3-15.6), the framework predicts a distinctive *partially-quantum regime* at intermediate τ_S/τ_B values where the C3 capacity condition is marginal. In this regime, engineered quantum systems exhibit dynamics that does not admit a clean density-matrix description and is distinguishable from both fully-quantum and fully-classical accounts through multi-time correlation diagnostics.

The structural condition. C3 (large hidden-sector capacity, $N_H \gg N_V$) is required for the framework’s emergent quantum mechanics in its fully-quantum form. When C3 is satisfied with overwhelming margin, the emergent dynamics is fully quantum-mechanical with conventional Born-rule statistics. When C3 fails entirely, the dynamics is classical stochastic. The intermediate regime — where $N_H/N_V \sim 1$ to 10, comparable rather than overwhelmingly larger — is the framework’s partially-quantum regime.

In engineered systems, the C3 marginality condition is *controllable*. By designing the hidden-sector capacity (the number of TLS in a superconducting qubit’s substrate, the number of nuclear spins in an NV center’s bath, the length of an engineered spin chain), the engineer chooses where on the quantum-to-classical spectrum the system operates. This is a new design degree of freedom that the framework predicts has empirical consequences distinct from continuous-decoherence accounts of the quantum-classical boundary.

The framework’s prediction. In the partially-quantum regime, the framework predicts that multi-time correlation functions show structural signatures distinct from both standard quantum mechanics and classical probability theory. Specifically:

Multi-time joint distributions. The joint probability distribution $P(x_1, x_2, \dots, x_n)$ for measurement outcomes at times $t_1 < t_2 < \dots < t_n$ in a partially-quantum system does not admit a description as either (a) the diagonal entries of a density-matrix evolution under a completely-positive trace-preserving map (the fully-quantum description), or (b) a classical probability distribution over hidden classical variables (the fully-classical description). The partially-quantum joint distributions violate constraints from both descriptions.

Information-backflow saturation. In fully-quantum systems, information backflow from the bath has bounded magnitude set by entropic constraints. In partially-quantum systems, information backflow saturates at lower bounds due to finite capacity — there is less information to back-flow than fully-quantum predictions would suggest. The framework predicts specific saturation patterns as $r = N_H/N_V$ approaches unity from above.

Non-monotonic capacity dependence. The framework predicts that the non-Markovianity measure $\mathcal{N}(r)$ depends non-monotonically on the capacity ratio r , with $\mathcal{N} \rightarrow 1$ for $r \gg 1$ (fully quantum), $\mathcal{N} \rightarrow 0$ for $r \ll 1$ (fully classical), and intermediate behavior $\mathcal{F}(r) = 2r - r^2$ for $r < 1$. This functional form is

distinctive and falsifiable: it is the framework’s specific structural prediction for how the quantum-to-classical transition occurs as capacity is reduced.

The functional form $\mathcal{F}(r) = 2r - r^2$. The quadratic dependence follows structurally from the $1/\omega$ Planckian spectral falloff combined with the framework’s structural conditions in the marginal-capacity regime. The derivation is developed in detail in `BEC_Experiment.md` §3.3; the structural content is that the lowest-frequency modes (which carry the most entanglement) are removed first as capacity is reduced, with the quadratic form being the integrated effect of this preferential removal.

The functional form distinguishes the framework’s prediction from alternative accounts. Standard open-quantum-system models predict either continuous exponential decay of non-Markovianity (with no structural transition) or first-order linear behavior $\mathcal{F}(r) = r$ at marginal capacity. The quadratic $2r - r^2$ form is unique to the framework’s structural mechanism and is empirically diagnostic.

Where the partially-quantum regime is accessible. The partially-quantum regime requires $N_H/N_V \sim 1$ — comparable hidden-sector and visible-sector capacities. Several engineered platforms can access this regime:

BEC analogue-gravity systems. By tuning the supersonic region length L_{super} , the effective hidden-sector capacity can be varied from large (L_{super} much greater than the BEC healing length) to marginal (L_{super} comparable to the healing length). This is the cleanest experimental platform, developed in §15.8.

Engineered spin chains. Short spin chains ($L \sim 5\text{-}10$) coupled to a qubit produce hidden sectors with capacity comparable to the qubit’s. The partially-quantum regime is accessible by varying chain length while keeping the qubit fixed.

Designed photonic systems. Photonic crystals with engineered defect modes can provide hidden sectors with controllable capacity. The partially-quantum regime is accessible by varying the number of accessible defect modes.

Diagnostic measurements. The framework’s partially-quantum prediction is diagnosable through multi-time correlation measurements. The key observable is the non-Markovianity measure $\mathcal{N}$, defined through the trace distance between two initially distinguishable states evolved through the system. Non-monotonic $\mathcal{N}(r)$ — specifically, the $2r - r^2$ form — is the framework’s distinctive empirical signature.

The next section develops the BEC analogue-gravity experiment in detail as the framework’s cleanest near-term test of these predictions.

15.8 The BEC analogue-gravity experiment

The framework’s cleanest near-term test of the BQP characterization theorem at the foundational level is the Bose-Einstein condensate (BEC) analogue-gravity experiment. The experiment uses an engineered analogue of a black hole horizon in a flowing BEC to test the framework’s capacity-controlled non-Markovianity prediction $\mathcal{F}(r) = 2r - r^2$ in a tabletop apparatus that is accessible to existing experimental groups and distinguishable from all known decoherence backgrounds.

The analogue-gravity setup. A flowing BEC produces a sonic horizon where the flow velocity equals the speed of sound. Phonons emitted in the supersonic region cannot propagate against the flow back into the subsonic region — exactly analogous to how photons inside a black hole event horizon cannot escape to infinity. The Hawking-Unruh effect predicts that the sonic horizon should emit thermal phonons at a temperature determined by the flow’s velocity gradient — analogue Hawking radiation.

The Steinhauer group at the Technion has demonstrated this experimentally over the past decade. Their setup uses a Rb-87 BEC in an elongated trap, with a step in the trap potential producing a region of supersonic flow that acts as the analogue horizon. The experimental observation of analogue Hawking radiation and the measurement of its thermal spectrum has been progressively refined; their 2016 paper provided the first observation, with subsequent work refining the measurement.

The framework’s modification. The framework’s content goes beyond reproducing analogue Hawking radiation. The framework predicts that the *non-Markovianity* of the subsonic-region phonon dynamics depends on the *capacity* of the supersonic-region (hidden-sector) phonon modes, with the dependence following the structural form $\mathcal{F}(r) = 2r - r^2$ for $r = N_H^{\text{eff}}/N_V < 1$.

The capacity is controllable through the supersonic region’s length L_{super} . Long supersonic region: large hidden-sector capacity, fully-quantum regime, $\mathcal{F}(r) = 1$. Short supersonic region: marginal capacity, partially-quantum regime, $\mathcal{F}(r) < 1$ with specific quadratic dependence.

The framework’s full prediction. Including both the capacity (C3) and the memory-timescale (C2) factors, the framework’s prediction is

$$\frac{\mathcal{N}(L_{\text{super}})}{\mathcal{N}_{\text{max}}} = \mathcal{F}\left(\frac{N_H^{\text{eff}}(L_{\text{super}})}{N_V}\right) \cdot \left(1 - \frac{\tau_S}{\tau_B(L_{\text{super}})}\right),$$

with $N_H^{\text{eff}} = \min(L_{\text{super}}/\pi\xi, \kappa_s L_{\text{super}}/\pi c_s)$ (where ξ is the healing length, κ_s is the surface gravity, c_s is the speed of sound), $N_V = \kappa_s L_{\text{sub}}/\pi c_s$ (where L_{sub} is the subsonic region length), and $\tau_B = L_{\text{super}}/|v - c_s|$.

This expression has *no free parameters* once the experimental configuration is specified. The overall scale $\mathcal{N}_{\text{max}}$ is measured at large L_{super} (where $\mathcal{F}(r) = 1$

and the memory factor is small); the *shape* of the curve as L_{super} is varied is the framework's structural prediction.

The experimental protocol. The proposed experiment modifies the Steinhauer group's existing BEC analogue black hole: a black hole–white hole pair with tunable separation L_{super} between the two sonic horizons. This configuration has already been demonstrated. For each value of L_{super} (six values spanning approximately 12 to 200 micrometers):

1. Prepare the BEC and allow Hawking radiation to reach stationarity (approximately 10 ms).
2. Measure the density profile at a sequence of times separated by approximately 0.1 to 1 ms.
3. Repeat approximately 10^3 times per L_{super} value.
4. Compute the temporal autocorrelation $G_{\text{sub}}^{(2)}(\Delta t)$ of the subsonic density. For P-indivisible dynamics, this should be non-monotonic with revivals at $\Delta t \sim \tau_B$.
5. Alternatively, prepare two initial perturbations and track the trace distance $D(t)$. Non-monotonic $D(t)$ directly measures $\mathcal{N} > 0$.

Controls. Three control measurements isolate the framework's predicted capacity mechanism from confounding effects.

Null control. No supersonic region ($v < c_s$ everywhere). The framework's prediction: $\mathcal{N} = 0$ because no hidden sector exists. This isolates the framework's capacity mechanism from any underlying BEC dynamics that might produce non-Markovianity through a different mechanism.

C2 test. Fixed L_{super} , variable flow velocity. This changes τ_B without changing N_H . The framework's prediction: $\mathcal{N}$ decreases as τ_B decreases, confirming the C2 dependence.

Thermal test. Variable condensate temperature, fixed geometry. Standard decoherence predicts strong temperature dependence; the framework predicts $\mathcal{N}$ depends only on N_H^{eff}/N_V and τ_S/τ_B , not on temperature directly. A null temperature dependence at fixed geometry would distinguish the framework's structural mechanism from standard thermal decoherence.

Feasibility. The predicted effect is not a small correction. At $L_{\text{super}} = 50$ micrometers ($r = 0.5$), the framework's prediction is that non-Markovianity is approximately 75% of its maximum. At $L_{\text{super}} = 25$ micrometers ($r = 0.25$), it drops to 44%. These are order-unity changes accessible to current BEC analogue-gravity apparatus.

The dominant noise source in BEC analogue experiments is atomic shot noise, characterized by Steinhauer's group at $\delta n/n \sim 10\text{-}20\%$ per pixel per shot. The framework's non-Markovianity signature is extracted from the *temporal structure* of $G^{(2)}(\Delta t)$ — specifically, whether it is monotonically decreasing (Markovian) or shows a revival (P-indivisible). Shot noise adds variance to each time

point but does not produce systematic non-monotonicity, so the signature is robust against the dominant noise source.

For statistical requirements, to distinguish $\mathcal{F}(r) = 0.75$ from $\mathcal{F}(r) = 1.0$ at 3σ confidence requires fractional uncertainty on $\mathcal{N}/\mathcal{N}_{\max}$ below approximately 8%. For temporal correlation measurements with M repetitions and N_{time} time points, the statistical uncertainty scales as $\delta\mathcal{N}/\mathcal{N} \sim 1/\sqrt{M \cdot N_{\text{time}}}$. With $N_{\text{time}} \sim 20$ and $M \sim 10^3$ repetitions per L_{super} value, the achievable uncertainty is approximately 0.7% — well below the 8% threshold required for 3σ discrimination.

Distinguishability from other backgrounds. The framework’s prediction is distinguishable from all known decoherence backgrounds through the specific quadratic functional form $\mathcal{F}(r) = 2r - r^2$ and the null thermal-dependence prediction.

Standard Bogoliubov-de Gennes theory predicts thermal occupation but not the framework’s capacity-controlled non-Markovianity. Standard BdG calculations would predict $\mathcal{N}/\mathcal{N}_{\max}$ to be essentially constant as L_{super} varies in the relevant regime, with any variation following the standard thermal-decoherence patterns that the framework’s null thermal test specifically isolates.

Other backgrounds. Several other potential sources of non-Markovianity in BEC systems exist (interaction-driven dynamics, three-body losses, technical noise from the trap potential) but each has characteristic signatures distinct from the framework’s prediction. The combination of the framework’s specific functional form $\mathcal{F}(r) = 2r - r^2$, the null thermal dependence, and the dependence on the geometric parameter L_{super} uniquely identifies the framework’s structural mechanism if observed.

Relation to DSW horizon decoherence. The BEC analogue-gravity test is structurally related to the Danielson-Satishchandran-Wald (DSW) horizon decoherence program developed in Chapter 9 §9.7. DSW prove within standard QFT in curved spacetime that Killing horizons impart fundamental decoherence on nearby quantum superpositions, with rate determined by horizon geometry. The framework’s BEC analogue-gravity test extends DSW’s analysis to the *capacity-controlled regime* (C3 marginality) that DSW does not address.

The convergence is significant. DSW provides theoretical confirmation of the framework’s universality-class commitments (horizons impose decoherence with rate determined by geometry) in standard physics literature. The BEC analogue-gravity test provides empirical access to the framework’s distinctive content (capacity-controlled non-Markovianity in the partially-quantum regime) in a laboratory apparatus. The two together — DSW for the universality-class content, BEC for the distinctive partially-quantum content — provide a complementary empirical foundation for the framework’s reach into engineered quantum systems.

Implications.

A positive result confirming the framework’s prediction $\mathcal{F}(r) = 2r - r^2$ would be the first laboratory demonstration of the framework’s capacity-controlled non-Markovianity — empirical confirmation of the framework’s distinctive content beyond the cosmological-scale predictions. The implications extend to all engineered quantum systems: the framework’s design philosophy (the bath as programmable quantum memory) would be empirically validated, opening commercial development of correlation-aware error correction, engineered-bath coherence extension, and backflow-aware quantum sensing.

A null result with the framework’s specific quadratic functional form *not* observed would constrain the framework’s content. The framework’s prediction in §15.7 is specific and falsifiable; a clean null result would force either revision of the framework’s structural commitments in the partially-quantum regime or revision of how the framework’s structural content applies to specific engineered systems. The framework would not be falsified entirely (the universality-class content from §15.2 would still stand), but the framework’s distinctive engineering content would face substantial pressure.

A null result with a *different* functional form observed (linear, exponential, or other) would falsify the framework’s structural mechanism for partially-quantum dynamics while leaving open the possibility of alternative structural accounts. The framework’s specific commitment to the Planckian-spectral mechanism producing $\mathcal{F}(r) = 2r - r^2$ would be refuted.

The BEC analogue-gravity test is therefore the framework’s *cleanest* near-term experimental program: a tabletop apparatus, accessible technology, specific structural predictions, multiple independent controls, robust statistical accessibility, and clear falsification conditions. The framework’s empirical exposure in engineered quantum systems is concentrated here.

15.9 Chapter close

The framework’s content in quantum engineering is concrete and empirically accessible. The chapter has developed five specific predictions with quantitative content for current and near-future quantum hardware platforms.

Five quantitative predictions.

1. *Correlation-aware error correction.* For superconducting qubits with TLS baths ($\tau_S/\tau_B \sim 10^{-3}$), correlation-aware decoders should outperform Markovian decoders at the same physical error rate (§15.4). The framework does not supply the size of that gain: converting a correlation length into a fault-tolerance overhead requires a specific code, decoder and noise spectrum, none of which follows from C1-C3. Testable now with current hardware and current decoder implementations.
2. *Coherence extension through engineered baths.* For qubits coupled to designed slow hidden sectors ($\tau_S/\tau_B \sim 10^{-1}$), coherence should persist beyond the Markovian estimate (§15.5). No numerical T_2 target is claimed;

the extension factor requires a device model the framework does not supply.

3. *Backflow-aware NV sensing.* For NV centers in diamond ($\tau_S/\tau_B \sim 10^{-6}$), backflow-aware protocols can improve sensitivity by approximately $10^3 \times$ (§15.6). NV magnetometry from nT/ $\sqrt{\text{Hz}}$ to pT/ $\sqrt{\text{Hz}}$ at room temperature.
4. *Partially-quantum regime.* Engineered systems at marginal C3 capacity ($N_H/N_V \sim 1$) exhibit the framework’s distinctive $\mathcal{F}(r) = 2r - r^2$ functional form (§15.7). Diagnosable through multi-time correlation measurements.
5. *BEC analogue-gravity test.* A specific tabletop experiment using the Steinhauer group’s BEC apparatus tests the partially-quantum regime prediction directly, with multiple controls, robust statistics, and clear falsification conditions (§15.8).

Three categories of empirical exposure.

The chapter’s content sorts into three empirical categories.

Immediately testable. The correlation-aware decoder (§15.4) and the BEC analogue-gravity test (§15.8) are both accessible to current experimental groups with existing hardware. Neither requires new platforms or technologies; both can produce results on the timescale of months to years.

Near-term testable with platform development. Coherence extension through engineered baths (§15.5) and backflow-aware NV sensing (§15.6) require development of new measurement protocols and engineered architectures. The required engineering is incremental — variations on existing platforms rather than new platforms — and the timescale is years rather than decades.

Structurally consequential. The partially-quantum regime (§15.7) provides the framework’s distinctive prediction beyond Markovian/non-Markovian dichotomies. The BEC analogue-gravity test (§15.8) provides the cleanest laboratory access to this regime, but other engineered platforms (spin chains, photonic systems) provide additional empirical access points.

The framework’s epistemic position in quantum engineering. The framework’s content here is positive: specific quantitative predictions for current and near-future quantum hardware, with concrete empirical accessibility on existing platforms. The predictions are testable, the falsification conditions are clear, and the potential commercial impact (correlation-aware error correction, NV sensing at SQUID-competitive sensitivity at room temperature, engineered-bath coherence extension) is substantial.

The chapter’s content connects the framework’s foundational commitments (Chapter 1’s emergent quantum mechanics, Chapter 14’s BQP theorem) to specific engineering applications. The framework is not just theoretical content with implications for fundamental physics; it provides design principles

and quantitative predictions for engineered quantum systems with empirical exposure in the next several years.

Forward pointers. Chapter 16 develops the framework’s medical applications: pharmacology, neurodegeneration, antibiotic resistance, immunotherapy, and the substantial epigenetics half on chromatin as biological hidden sector. The framework’s structural content propagates from Chapter 13’s quantum biology and this chapter’s engineering content to specific clinical and pharmaceutical predictions. Chapter 17 develops the bioinformatics consequences. Chapters 18 and 19 develop the forward-predictions chapter and the inventory of resolved and open problems in fundamental physics.

The framework’s empirical reach across Chapters 13-17 is concentrated at the application level — specific predictions for biology, computing, engineering, medicine, and bioinformatics. The framework’s content here in quantum engineering is the bridge between the foundational chapters (1-9) and the applied chapters (13-17), demonstrating that the framework’s structural commitments at the substratum level produce specific quantitative predictions across multiple engineering domains.

Chapter 16

Medicine

Status — conjectural extension. This chapter applies the core framework to medicine, and is offered as a *conjecture*, not an established result: its retrodictions — notably the Rett-syndrome reversibility (Guy et al. 2007), which predates the framework — are consistency checks rather than confirmations, while its forward content (wider therapeutic windows across as-yet-untested reader-writer disorders, and specific schedule-dependence predictions) is what future clinical data would confirm or break. It carries no evidential weight for the core framework — which stands or falls independently on its physics — and should be read as a structural possibility to be tested, not as support for the framework.

16.1 What this chapter develops

Chapters 13 through 15 developed the framework’s reach into biology, computing, and engineering. Chapter 13 established that biological systems with C1-C3 architectures exhibit non-Markovian dynamics as their default behavior — quantum mechanics in biology is the framework’s prediction rather than the

anomaly conventional accounts treat it as. Chapter 14 developed the BQP theorem as the framework’s structural limit on computational complexity. Chapter 15 developed the framework’s content for engineered quantum systems, with specific quantitative predictions for current quantum hardware.

This chapter develops the framework’s content for medicine — the framework’s reach into clinical pharmacology, neurodegeneration, antibiotic resistance, immunotherapy, cardiac pharmacology, autoimmune disease, and epigenetics. The framework’s content here is *concrete and clinically engaged*: specific predictions for current drug-development programs, distinguishing predictions where the framework gives different answers than standard pharmacology, and one structural retrodiction (the wider therapeutic window for reader-writer disorders, consistent with Guy et al. 2007 for Rett syndrome) that illustrates the framework’s reach beyond fundamental physics into clinical biology.

The chapter has explicit two-half structure.

The *first half* (§§16.2-16.8) develops the framework’s content across six conventional pharmacology domains: cancer pharmacology (with checkpoint kinase inhibitors and the framework’s account of why Chk1 inhibitor selectivity emerges), neurodegeneration (with CaMKII as the canonical molecular memory device and PTSD reconsolidation as τ_B pathology), antibiotic resistance (with persister cells as memory accumulation), immunotherapy (with T-cell exhaustion as TCR memory pathology), cardiac pharmacology (with use-dependent block as non-Markovian channel dynamics), and autoimmune disease (with partial JAK inhibition exploiting memory architecture). Each section identifies the C1-C3 architecture of the relevant biological system, develops the framework’s structural prediction, and identifies distinguishing predictions where the framework gives different answers than standard pharmacology.

The *second half* (§§16.9-16.12) develops epigenetics as biological hidden sector. The substratum-emergent operator distinction from Chapter 6 §6.3 applies at the chromatin level: substratum-level operators are molecular patterns of DNA methylation and histone modifications; emergent-level operators are gene expression programs and transcriptional phenotypes. The two operators can dissociate — substrate-level memory preserved while emergent-level operator is disrupted, and vice versa. This structural framing organizes the three-axis pharmacology of epigenetic drugs (readers, writers, erasers), the wider therapeutic window prediction for reader-writer disorders (empirically confirmed in Rett syndrome), and eight specific predictions for clinical pharmacology of epigenetic systems.

Section 16.13 closes the chapter with a unifying therapeutic principle: *memory asymmetry*. When a disease process depends on non-Markovian signaling dynamics that the corresponding normal tissue does not, therapies can target the *memory structure* rather than the *catalytic function* — predicting wider therapeutic windows because memory dependence is more disease-specific than catalytic activity.

The framework’s strongest non-physics case here is structural. Its one empirical

touchpoint — the wider therapeutic window in Rett syndrome (Guy et al. 2007) — is a retrodiction, since the result predates the framework, so it carries the weight of a consistency check rather than of the JUNO forward prediction. The framework’s forward predictions in medicine (below) remain to be tested.

16.2 The enzyme as a read-write system

Before developing specific disease domains, the chapter establishes the framework’s structural account of enzymes as read-write systems. The account organizes the subsequent disease-by-disease content under a unified structural pattern: each disease involves a specific C1-C3 architecture, with pathology arising through alteration of the architecture’s memory dynamics.

Molecular memory mechanisms. The framework’s prediction that enzyme activity history is physically encoded in the enzyme’s structure and persists across multiple catalytic cycles is well-established in molecular biology through multiple mechanisms.

Multisite post-translational modification. PTMs — phosphorylation, acetylation, ubiquitination, methylation — covalently alter specific residues, changing the protein’s conformational landscape and activity. A protein with N modifiable sites, each with k possible states, has k^N distinct modification patterns. For Chk1 with approximately ten regulatory phosphorylation sites, $2^{10} \approx 1,000$ distinct states constitute a physical memory register encoding the enzyme’s recent history. Gunawardena (2012) explicitly describes this as “history-based encoding”: the same cellular condition can produce different modification patterns depending on the prior history.

Sequential phosphorylation. For Chk1 specifically, phosphorylation at S317 must precede phosphorylation at S345 — the second event is conditional on the first (Wilsker et al. 2008). This is a direct non-Markovian signature: the probability of the second modification depends on history, not just on the current state.

Conformational hysteresis. Proteins occupy multiple distinct conformational states, with transitions depending on the protein’s history. For kinases, the autoinhibited vs. active conformation persists on timescales much longer than individual phosphorylation events. The Chk1-S splice variant acts as an endogenous repressor whose binding/unbinding is the slow process storing the history.

Chromatin as long-term memory. At the network level, the histone code provides the most dramatic example. DNA damage induces histone modifications (γ H2AX, H4K20me) that persist for hours to days — far longer than the kinase cascade that wrote them.

The memory hierarchy. The framework’s content extends across many decades of biological timescale, with each scale satisfying C1-C3 with its own characteristic τ_B .

Memory mechanism	Write operation	Storage medium	τ_B	C1-C3 role
Conformational hysteresis	Ligand binding	Oligomeric/regulatory states	seconds	C2
Sequential phosphorylation	Ordered modification	Conformational accessibility	Minutes	C2 + C3
Multisite PTM	Phosphorylation	Modification pattern (2^N states)	Min-hrs	C3
Chromatin marks	γ H2AX, methylation	Histone modification state	Hrs-days	C2 + C3
DNA methylation	Methyltransferase activity	CpG state	Months-generations	C2 + C3

In multicellular tissues — particularly the nervous system — the same C1-C3 architecture extends upward through additional layers, each treating the layer below as its hidden sector. Synaptic-layer LTP/LTD changes (hours to days), circuit-layer engram-cell ensembles (days to months), systems-layer hippocampal-cortical dialogue (months to years), and cortical-layer distributed semantic representation (years to decades) form a five-layer architecture nested above the molecular level.

The hierarchical structure means that a perturbation at one layer (a drug, a disease process) can have effects propagating upward through layers with progressively longer τ_B . The clinical time-course of memory-related diseases — and the schedule-dependence of memory-targeted drugs — reflects this layered architecture.

The core therapeutic principle. The framework identifies a specific therapeutic axis: *memory asymmetry*. When a disease process depends on non-Markovian signaling dynamics that the corresponding normal tissue does not depend on (or depends on differently), therapies can target the *memory structure* rather than the *catalytic function*. This is pharmacologically distinct from standard inhibition and predicts *wider therapeutic windows* because memory dependence is more disease-specific than catalytic activity.

The memory-asymmetry principle organizes the chapter’s content across the disease domains in §§16.3-16.12. Each disease’s pharmacology is reframed in terms of the relevant C1-C3 architecture; each therapeutic intervention is reframed in terms of its effect on the architecture’s memory dynamics; each distinguishing prediction follows from this reframing.

16.3 Cancer pharmacology

Cancer is, in the framework's reading, a disease of altered memory dynamics at multiple scales. DNA damage response kinases (Chk1, ATR, ATM) operate as C1-C3 systems with their catalytic domains coupled to regulatory domains carrying PTM-based memory. The disrupted memory dynamics of cancer cells — both their hyperdependence on certain memory pathways and their altered memory architectures relative to normal cells — provides the structural basis for the framework's content in cancer pharmacology.

Checkpoint kinase inhibitors and selective tumor sensitization. Chk1 inhibitors combined with gemcitabine and/or radiation selectively sensitize tumor cells — particularly pancreatic cancer — while sparing normal cells. The framework's account: Chk1 is a serine/threonine kinase whose catalytic domain (fast subsystem) is regulated by its SQ/TQ domain and C-terminal regulatory region (slow hidden sector). C1-C3 are satisfied: catalytic and regulatory domains are allosterically coupled (C1); regulatory conformational changes (μ s-ms) much exceed phosphorylation events (ns- μ s), with PTM persistence extending to minutes-hours (C2); the regulatory domain has approximately 2^{10} modification states plus continuous conformational degrees of freedom (C3).

The framework predicts that Chk1's checkpoint signaling is non-Markovian: its response to a DNA damage signal depends on its recent activation history — specifically, on the PTM pattern and conformational state written by previous damage events. A Chk1 inhibitor that binds the regulatory domain disrupts the *memory structure* of the kinase, altering the history-dependence of future checkpoint responses.

Why selectivity emerges. In tumor cells (defective G1, relying on Chk1), the non-Markovian memory is the primary mechanism maintaining genomic stability through repeated replication cycles. Disrupting this memory is catastrophic. In normal cells (intact G1), the Chk1 memory is redundant — the G1 checkpoint provides an independent, Markovian damage response. The selectivity is not just about checkpoint redundancy but about the differential role of non-Markovian dynamics.

Why schedule matters. The clinical finding that the *order and timing* of gemcitabine + Chk1 inhibitor administration is critical for efficacy is a direct signature of non-Markovian dynamics. In a Markovian system, only concentrations matter. In a non-Markovian system, the first drug writes information into the enzyme's memory, and the second drug's effect depends on what was written.

Radiation adaptive response. A small priming dose of radiation (~ 0.01 - 0.1 Gy) makes cells more resistant to a subsequent larger dose (~ 2 Gy), with the effect lasting hours to days. The framework identifies this as information back-flow from the chromatin hidden sector: dose 1 writes modification marks into histones, which persist and are read by the DNA damage response when dose 2 arrives. The “adaptation” is the response network reading the memory of the

previous damage event — a direct biological manifestation of the framework’s information backflow prediction.

Distinguishing predictions for cancer pharmacology.

Memory-selective scheduling. The optimal dose interval depends on τ_B of the tumor’s checkpoint kinases. The second dose should arrive when the memory written by the first dose is maximally sensitizing — a timing that differs between tumor cells (deregulated Chk1, altered τ_B) and normal cells (normal τ_B). Current schedules are empirically optimized from a few discrete intervals; the framework predicts the optimal interval is *calculable* from the measured τ_B and is tumor-specific.

Low-dose memory priming. Instead of a single high dose, give repeated low “priming” doses that write sensitizing memory into the tumor’s checkpoint system, followed by a moderate treatment dose. The predicted protocol: Days 1, 3, 5 — gemcitabine at ~20% MTD (priming phase); Day 7 — gemcitabine at ~50% MTD + Chk1 inhibitor (kill phase). Total drug exposure: ~70% of standard protocol. Predicted toxicity: substantially lower, because normal cells’ Markovian G1 backup is unaffected by priming.

Memory-targeted drugs. Drugs that specifically disrupt the *memory structure* without blocking catalysis: accelerating the regulatory domain’s slowest conformational mode (decreasing τ_B) to erase checkpoint memory without blocking checkpoint activation. Such a drug would be invisible to standard kinase inhibitor screens, which measure catalytic inhibition.

The framework’s content in cancer pharmacology is therefore not just structural reframing but specific clinical protocol predictions. The protocol predictions are testable in animal models and clinical trials with current technology; the framework predicts substantive improvements in therapeutic index over current scheduling.

16.4 Neurodegeneration

Neurodegenerative disease — Alzheimer’s, Parkinson’s, Huntington’s, and related conditions — has resisted three decades of pharmacological assault. The amyloid hypothesis has yielded drugs that successfully clear amyloid plaques (aducanumab, lecanemab, donanemab) without restoring cognition. The framework’s reading is that this targeting pattern misses the mechanism: neurodegeneration is fundamentally a disease of altered memory dynamics in synaptic signaling kinases, and amyloid clearance addresses a symptom rather than the underlying τ_B pathology.

CaMKII as the canonical C1-C3 system. Calcium/calmodulin-dependent protein kinase II (CaMKII) is the canonical molecular memory device in neurons. CaMKII satisfies C1-C3 with overwhelming margin:

- *C1.* The kinase domain is coupled to the regulatory/association domain through the autoinhibitory segment.
- *C2.* Autophosphorylation at T286 creates a bistable switch with persistence time of minutes to hours — much longer than individual calcium transients ($\sim$ ms). $\tau_S/\tau_B \sim 10^{-4}$ to 10^{-5} .
- *C3.* CaMKII forms dodecameric holoenzymes with 12 subunits, each independently phosphorylatable — $2^{12} = 4,096$ distinct modification states.

CaMKII's memory function in long-term potentiation is the *intended* biological use of non-Markovian dynamics. Neurodegeneration involves pathological perturbation of this memory system.

The framework's prediction for neurodegenerative disease. Neurodegenerative disease involves pathological alteration of τ_B in synaptic signaling kinases. Three specific clinical syndromes follow this pattern.

Normal aging. Gradual increase in τ_B (slower conformational relaxation due to oxidative damage) produces excessive memory retention, synaptic rigidity, and reduced plasticity.

Alzheimer's disease. $A\beta$ oligomers interact with CaMKII and alter its regulatory domain dynamics, shifting τ_B to produce pathological memory states. These pathological states drive tau hyperphosphorylation through downstream cascades (including Chk1-CIP2A-PP2A), producing the established neurofibrillary pathology of late-stage disease.

Parkinson's disease. α -synuclein aggregates alter the conformational dynamics of LRRK2 and other PD-associated kinases, with similar shifts in τ_B producing the dopaminergic neuron dysfunction of clinical PD.

Therapeutic implication: τ_B -normalizing drugs. Instead of targeting misfolded proteins (which has largely failed clinically), target the *altered memory timescale* of signaling kinases. A drug that restores τ_B to its normal value could preserve synaptic function without clearing aggregates.

Testable prediction. Measure CaMKII conformational dynamics (by FRET or HDX-MS) in neurons exposed to $A\beta$ oligomers versus controls. The framework predicts that $A\beta$ shifts τ_B of CaMKII's regulatory domain. A compound that reverses this τ_B shift should rescue synaptic plasticity (measurable by long-term potentiation induction) independently of $A\beta$ clearance.

PTSD as τ_B pathology with reconsolidation as therapy. Post-traumatic stress disorder involves abnormally persistent and intrusive memory of traumatic events. From the framework's perspective, this is failure of τ_B to be appropriately finite at the relevant memory layer — the trauma-encoded trace persists at a layer where it should have decayed and continues to drive visible-sector dynamics (intrusive re-experiencing, hyperarousal).

The clinical efficacy of the reconsolidation paradigm — recalling a traumatic memory under propranolol blockade reduces its later emotional valence — is

directly framework-predicted. The mechanism is structural: recall opens the memory to obligatory re-storage, and pharmacologically blocking the noradrenergic re-encoding selectively erases the emotional content while preserving the declarative trace. The reconsolidation window during which the memory is labile provides a *therapeutic opportunity*: a drug given during this window writes into the re-storage process, modifying the future content of the memory.

The framework predicts that this approach generalizes: any memory layer with a measurable τ_B should have a corresponding “reconsolidation window” during which targeted intervention can rewrite the memory at that layer without affecting memories already consolidated to slower layers. The selectivity is not coincidental — it follows from the layered architecture of §16.2.

Direct mechanical access via low-intensity focused ultrasound. Standard pharmacology writes to *chemical* degrees of freedom (receptor occupancy, phosphorylation state). It cannot directly access the conformational and mechanical dynamics that constitute τ_B at the molecular layer. Low-intensity focused ultrasound (LIFU) is the first therapeutic modality that *directly* perturbs these mechanical degrees of freedom: acoustic radiation force acts on mechanosensitive ion channels (Piezo1, TRAAK, TRP family) and on membrane conformational states without requiring receptor binding.

This makes LIFU structurally aligned with the framework in a way pharmacology is not. LIFU is a τ_B -*writer* rather than a catalytic inhibitor.

Clinical evidence is consistent with this reading. The 2025 Korean trial (Ye et al., *Journal of Neurosurgery*) showed cognitive improvement in Alzheimer’s patients from focused ultrasound blood-brain-barrier opening *without* concurrent drug administration. This is unexpected on the amyloid hypothesis but framework-consistent if ultrasound directly normalizes molecular-memory substrate dynamics (CaMKII regulatory-domain kinetics, for instance) independent of amyloid clearance. LIFU has shown reversible neuromodulation effects in the anterior limb of the internal capsule (depression target), consistent with direct perturbation of axonal conduction via mechanosensitive K2P channels at the nodes of Ranvier.

Framework-specific predictions for LIFU. Pulse repetition frequency should match characteristic molecular τ_B values of the target substrate. CaMKII intervention should use minute-scale to hour-scale pulse trains rather than continuous sonication. Schedule dependence should follow the same τ_B -matched-interval logic as drug scheduling. The therapeutic window should be biphasic: low intensity for τ_B normalization, high intensity for non-specific mechanical damage, with the window between framework-predicted to be narrow.

These predictions are not currently how LIFU parameters are selected clinically (parameters are largely empirical). Framework-informed trials would systematically scan pulse repetition frequency and intensity against measured molecular τ_B values.

16.5 Antibiotic resistance

Antibiotic resistance presents an apparent puzzle for conventional pharmacology: bacteria that have never been exposed to an antibiotic can develop tolerance within a single generation through *non-genetic* mechanisms (phenotypic tolerance via persister cells), and these mechanisms involve memory effects that standard Markovian models cannot accommodate. The framework’s reading provides a structural account: persister formation is the accumulation of SOS memory past a threshold, with each antibiotic exposure writing information into the RecA filament state.

Persister cells as memory accumulation. Persister cells survive antibiotic treatment through transient phenotypic tolerance, not genetic resistance. The SOS response pathway — regulated by RecA — satisfies C1-C3: RecA’s nucleotide binding (fast) is allosterically coupled to filament formation on ssDNA (slow, seconds to minutes), with filaments extending over hundreds of bases (exponentially large state space).

The framework’s prediction: persister formation is not random switching but the accumulation of SOS memory past a threshold. Each antibiotic exposure writes information into the RecA filament state, which persists and modulates subsequent responses. The “switch” between susceptible and persister phenotypes is not a stochastic event but a memory-readout from the cumulative exposure history.

Memory-optimized antibiotic scheduling. The framework’s reading implies that *scheduling* of antibiotic doses matters in ways standard pharmacokinetics does not capture. Specifically:

Continuous exposure. The conventional “keep antibiotic above MIC for as long as possible” strategy maximizes SOS memory accumulation and therefore persister formation. This is the worst possible scheduling from the framework’s perspective.

Pulsed exposure with τ_B -matched intervals. Antibiotic pulses separated by intervals matching the SOS memory’s τ_B allow the memory to relax between pulses, preventing the cumulative accumulation that produces persisters. The framework predicts this pulsed scheduling produces fewer persisters than continuous exposure at equivalent total drug exposure.

Combination scheduling. Combinations of antibiotics targeting different memory architectures (different bacterial signaling pathways with different τ_B) can interfere with each other’s memory dynamics, producing effects beyond simple additive killing. The framework predicts specific combinations and timings that exploit memory interference.

Distinguishing predictions.

τ_B -matched pulsed dosing. The framework predicts that pulsed dosing with τ_B -matched intervals produces fewer persister cells than continuous dosing at

equivalent total drug exposure. The prediction is testable in animal models of chronic bacterial infection with measured RecA filament dynamics. Standard pharmacokinetic models predict no benefit to pulsed dosing if total exposure is matched; the framework predicts a specific benefit from τ_B matching.

Anti-persister drugs. Drugs that specifically accelerate RecA filament dissociation (decreasing τ_B of the SOS memory) should eliminate persisters without killing growing bacteria. This is a new drug class that the framework predicts is pharmacologically distinct from standard antibiotics; such drugs would be invisible to growth-inhibition screens.

Memory-erasing combinations. Combinations of β -lactams (acting on rapidly-growing bacteria) with memory-erasing agents (preventing persister formation) should produce dramatically improved outcomes in chronic infections. The framework predicts specific drug combinations from the memory architecture of bacterial signaling networks.

16.6 Immunotherapy and T-cell exhaustion

T-cell exhaustion — the progressive loss of T-cell effector function during chronic antigen exposure — is the dominant clinical limitation of cancer immunotherapy. Checkpoint inhibitors (anti-PD-1, anti-CTLA-4) reverse exhaustion partially but only in a fraction of patients. The framework’s reading: T-cell receptor (TCR) signaling is a C1-C3 system with the exhaustion state corresponding to a specific pathological memory configuration of the regulatory machinery.

TCR signaling as a C1-C3 system. The T-cell receptor’s signaling cascade involves multiple kinases (Lck, ZAP-70, LAT) with regulatory phosphorylation sites, all coupled to the downstream transcription factor network. C1 is satisfied through the regulatory phosphorylation network; C2 through the slow dynamics of transcription factor accumulation and dissipation; C3 through the combinatorial space of phosphorylation patterns across the cascade.

The framework’s prediction: TCR signaling is non-Markovian. Chronic antigen exposure progressively writes specific PTM patterns into the kinase cascade, driving the cell toward the “exhaustion” memory state characterized by upregulated inhibitory receptors (PD-1, TIM-3, LAG-3) and downregulated effector function. The exhaustion state is not an absorbing classical phenotype but a specific memory configuration that can in principle be rewritten.

Therapeutic implication. Standard checkpoint inhibition (blocking PD-1/PD-L1) acts on the *effector phase* of T-cell function. The framework predicts that targeting the *memory writers* (kinases that write the exhaustion PTM pattern) or the *memory erasers* (phosphatases that remove these patterns) should be pharmacologically distinct and synergistic with standard checkpoint inhibition.

Specifically: The serine/threonine kinases responsible for the exhaustion-specific phosphorylation patterns (likely candidates: PKC, MAPK, and downstream ef-

factors) are framework-predicted to be the *upstream writers* of the exhaustion state. Inhibiting these writers should prevent exhaustion development; activating the corresponding erasers should reverse established exhaustion. This is a new therapeutic axis distinct from standard checkpoint blockade.

Distinguishing prediction: kinase regulatory-domain accelerators.

Drugs that accelerate the regulatory-domain dynamics of TCR-pathway kinases (decreasing τ_B without blocking catalytic activity) should reverse exhaustion. These drugs would be invisible to standard kinase inhibitor screens, which measure catalytic inhibition. The framework predicts that screening compound libraries against measured regulatory-domain dynamics (rather than catalytic inhibition) would identify a new class of immuno-oncology agents.

16.7 Cardiac pharmacology

Cardiac antiarrhythmic drugs exhibit *use-dependent block*: the drug's binding to the sodium or calcium channel depends on the channel's activation history rather than just on instantaneous membrane voltage. The phenomenon is well-established empirically (Vaughan-Williams class I antiarrhythmics, lidocaine, propafenone, flecainide) but is treated as an empirical complication of channel pharmacology. The framework's reading: use-dependent block is non-Markovian channel dynamics, with the channel acting as a C1-C3 system whose conformational memory determines drug binding affinity.

Use-dependent block as non-Markovian channel dynamics. Voltage-gated ion channels (Na_V , Ca_V , K_V) have multiple conformational states (closed, open, inactivated, slow-inactivated) with transitions on different timescales. The fast subsystem is the open/closed conformational equilibrium (microseconds); the slow subsystem is the inactivated-state dynamics (milliseconds to seconds). C1-C3 are satisfied with the standard $\tau_S/\tau_B \sim 10^{-3}$ regime characteristic of the framework's engineered partition examples.

The framework's prediction: drug binding to voltage-gated channels is non-Markovian. The drug's effective affinity depends on the channel's activation history through the slow-inactivated state's memory of recent activations. This produces the observed use-dependence: rapidly cycling channels (during tachyarrhythmias) accumulate more inactivated-state occupancy and bind drug more avidly; slowly cycling channels (during normal rhythm) accumulate less and bind less avidly.

Therapeutic implication. Use-dependent block produces *frequency-selective* antiarrhythmic effects: drugs preferentially affect rapidly firing pathological tissue (arrhythmic foci, reentry circuits) while sparing normally firing tissue. This is the framework's memory-asymmetry principle in cardiac pharmacology: the pathological tissue's high-frequency firing writes more memory into the channels, making them more susceptible to drug binding. The therapeutic window between antiarrhythmic effect and toxicity is determined by the difference in channel memory states between pathological and normal tissue.

Distinguishing predictions.

Optimal dosing intervals. The framework predicts that the optimal interval between antiarrhythmic doses depends on the channel's τ_B at the inactivated state. Dose intervals matching τ_B maximize the memory-write while minimizing baseline blockade; intervals much shorter or longer reduce therapeutic index. Current dosing is empirically optimized; the framework predicts calculable optima from channel kinetics.

Memory-targeted antiarrhythmics. Drugs that accelerate channel recovery from inactivation (decreasing τ_B of the inactivated state) should eliminate arrhythmias by erasing the memory-accumulation that produces them. This is a new therapeutic axis distinct from standard channel-blocking antiarrhythmics; such drugs would not show standard channel-blocking pharmacology and would be invisible to standard antiarrhythmic screens.

Frequency-selective combinations. Combinations of drugs with different channel memory profiles (different τ_B values) can produce frequency-selective effects beyond simple additivity. The framework predicts specific combinations and dosing for arrhythmias at different cycle lengths.

16.8 Autoimmune disease

Autoimmune disease — rheumatoid arthritis, lupus, psoriasis, inflammatory bowel disease — involves chronic activation of cytokine signaling networks centered on the Janus kinase (JAK) family. Standard immunosuppression uses high-dose, often broad-spectrum agents (corticosteroids, methotrexate, anti-TNF) with substantial side effects. JAK inhibitors (tofacitinib, baricitinib, upadacitinib) provide more targeted intervention but with characteristic dose-response patterns that conventional pharmacology has difficulty explaining. The framework's reading: JAK signaling is a C1-C3 system, and partial JAK inhibition exploits the network's memory architecture for disproportionate clinical efficacy.

Disproportionate efficacy of partial JAK inhibition. Clinical experience with JAK inhibitors shows that *partial* JAK inhibition (binding less than ~50% of available JAK kinases at clinically tolerable doses) produces substantial clinical efficacy in many autoimmune conditions. This is unexpected on standard Markovian models, which predict that the suppression of cytokine signaling should be roughly linear in the fraction of inhibited kinase.

The framework's account: JAK signaling is non-Markovian, with the cytokine response depending on the cumulative phosphorylation state of the JAK-STAT pathway over multiple signaling cycles. Partial inhibition produces disproportionate effects because it disrupts the memory accumulation that drives the chronic activation, even though instantaneous catalytic activity remains substantial.

Therapeutic implication. The framework predicts that JAK inhibitor dose-response curves should show *threshold behavior* — clinical efficacy emerges above

a specific dose threshold corresponding to disruption of memory accumulation, rather than linear with kinase inhibition. Below threshold: minimal efficacy regardless of partial inhibition. Above threshold: substantial efficacy from memory disruption. This pattern matches the clinical observation of “minimum effective dose” thresholds with rapid efficacy gains above threshold.

Distinguishing prediction: optimal scheduling. The framework predicts that *intermittent* JAK inhibitor dosing (with intervals matching the cytokine memory’s τ_B , typically hours to days) should produce equivalent or superior efficacy to continuous dosing at lower cumulative drug exposure. Standard pharmacokinetic models predict continuous dosing is optimal; the framework predicts intermittent dosing exploiting memory dynamics could improve therapeutic index.

A second distinguishing prediction: combinations of JAK inhibitors with phosphatase activators (drugs that accelerate the memory erasure of the JAK-STAT pathway) should be synergistic, producing equivalent clinical effect at lower JAK inhibitor doses. This combination would be invisible to standard pharmacology screens that focus on catalytic inhibition.

16.9 Epigenetics as biological hidden sector

The second half of the chapter develops the framework’s content on epigenetics — the systematic application of the substratum-emergent operator distinction (Chapter 6 §6.3) to chromatin biology. The framework’s content here is *substantive*: epigenetic regulation is a direct biological instantiation of the framework’s C1-C3 architecture, with the substratum-emergent operator distinction organizing the three-axis pharmacology of epigenetic drugs (readers, writers, erasers) and producing a wider-therapeutic-window pattern consistent with Rett syndrome (Guy et al. 2007).

The C1-C3 architecture of chromatin. The transcriptional machinery — RNA polymerase II, transcription factors, mediator complex — is the visible sector. The chromatin state — DNA methylation patterns, histone modifications, chromatin compaction, three-dimensional architecture — is the hidden sector. The three structural conditions are satisfied with overwhelming margin.

C1 (coupling). The transcriptional machinery and chromatin state are coupled through multiple mechanisms: direct physical coupling (nucleosome positioning controls promoter access, DNA methylation modulates transcription factor binding, histone tail modifications recruit reader proteins), writer-reader feedback (transcription recruits chromatin-modifying enzymes that alter the chromatin state, which modulates future transcription), and structural coupling (chromatin looping by CTCF and cohesin brings distal regulatory elements into proximity with promoters).

C2 (slow bath). The chromatin state changes on timescales far slower than individual transcription events. The hierarchy is multi-scale:

Mechanism	τ_B	τ_S/τ_B
DNA methylation (CpG)	Cell generations	$\sim 10^{-6}$
Histone methylation	Hours to days	$\sim 10^{-3}$ to 10^{-2}
Histone acetylation	Minutes to hours	$\sim 10^{-1}$ to 1
Chromatin compaction	Cell generations	$\sim 10^{-6}$
CTCF/cohesin loops	Hours to cell generations	$\sim 10^{-3}$ to 10^{-6}

The framework predicts a *hierarchy of memory strength*: genes regulated primarily by DNA methylation should show the strongest history-dependence (largest non-Markovian correction), while genes regulated primarily by histone acetylation should show the weakest (approaching Markovian behavior).

C3 (capacity). The chromatin memory register has astronomical capacity. Approximately 28×10^6 CpG sites in the human genome with binary methylation states give $\sim 2^{28 \times 10^6}$ possible CpG methylation patterns — vastly exceeding any conceivable transcriptional state space. C3 is satisfied by approximately seven million orders of magnitude.

Cancer as pathological epigenetic memory. Cancer is, in the framework’s language, a disease of the epigenetic hidden sector. The malignant transcriptional program is maintained not only by genetic mutations but by an aberrant epigenetic memory state — stable patterns of DNA methylation and histone modifications that lock the cell into a proliferative program. This memory is self-reinforcing: the malignant transcription writes further aberrant marks through recruited chromatin modifiers, which stabilize the malignant transcription.

The framework predicts that disrupting the memory *structure* (the stability and capacity of the epigenetic hidden sector) is a more selective anti-cancer strategy than disrupting the memory *content* (reactivating specific silenced genes). A memory-structure drug would reduce τ_B of the aberrant methylation and histone marks (making them less stable), not affect τ_B of constitutive marks (housekeeping genes, developmental identity), and selectively destabilize the tumor’s pathological memory without affecting normal cellular memory.

Aging as memory accumulation. Aging involves progressive accumulation of epigenetic marks — the “epigenetic clock” of Horvath 2013 — that alter gene expression in ways associated with declining function. In the framework’s language, aging is the accumulation of memory in the slowest epigenetic layers (DNA methylation, chromatin compaction) beyond the cell’s ability to maintain functional homeostasis.

The framework predicts that aging-associated gene expression changes should show the strongest non-Markovian signatures (longest autocorrelation times) — consistent with the observation that aged cells have more stable, less plastic epigenetic states than young cells. Partial reprogramming through transient Yamanaka factor expression erases recent epigenetic memory while preserving deeper developmental identity. The framework predicts a critical pulse duration matching τ_B of aging-associated marks: shorter pulses erase recent marks while preserving identity; longer pulses cross the threshold and erase developmental identity.

16.10 The substratum-emergent operator distinction at chromatin

The framework’s distinctive content in epigenetics is the substratum-emergent operator distinction applied at the chromatin level. The distinction, developed in Chapter 6 §6.3 for fundamental physics (operators that look identical at the emergent level can have distinct conservation properties depending on whether they act at the substratum or emergent level), has a direct biological analog at chromatin that is empirically established in current biomedical literature.

The two operators at chromatin.

Substratum-level memory operator. The molecular pattern of DNA methylation, histone modifications, and chromatin-state assignments. This is a directly measurable molecular substrate — accessible by bisulfite sequencing, ChIP-seq, ATAC-seq, and related assays. The substratum-level operator has definite conservation properties at the molecular level (specific marks at specific positions).

Emergent-level memory operator. The gene expression program, the transcriptional phenotype, the cellular function readout. This is what cells display and what clinical phenotypes measure — accessible by RNA-seq, single-cell transcriptomics, and functional assays. The emergent-level operator has conservation properties at the coarse-grained level (specific expression patterns, specific phenotypes).

The two operators are related by the trace-out, which here is the cell’s signal-transduction and transcription machinery integrating fast molecular fluctuations into slower cellular responses. Critically, the operators can dissociate: *substrate-level memory can be preserved while the emergent-level operator is disrupted, and vice versa.* This dissociation is the framework’s key structural prediction at chromatin.

Three empirical confirmations of the operator distinction at chromatin.

iPSC reprogramming residuals. When somatic cells are reprogrammed to induced pluripotent stem cells (iPSCs) through transient Yamanaka factor expression, the cellular phenotype changes — now pluripotent. But substantial residual epigenetic memory of the source cell type persists at the substrate

level: DNA methylation patterns, histone modifications, and chromatin accessibility profiles retain measurable signatures of the source cell type even when the emergent-level functional phenotype has fully flipped. Lin et al. 2024 (DOI: 10.3389/fcell.2024.1306530) provides detailed characterization of this phenomenon.

The framework’s reading: the substrate-level operator retains source-cell-type information; the emergent-level operator (cell identity) has flipped. The two operators are dissociated. This should be a *structural feature* of any reprogramming protocol, not a contingent reprogramming limitation.

Bivalent chromatin states. Embryonic stem cells maintain “bivalent” chromatin at lineage-control genes — simultaneous presence of activating H3K4me3 and repressive H3K27me3 marks at the same loci (Bernstein et al. 2006; reviewed in Voigt-Tee-Reinberg 2013). The substrate-level operators (the two mark types) coexist with definite conservation properties at the molecular level; the emergent-level operator (gene expression) is held at zero by their balanced presence. This is a *multi-operator substratum* structure with poised emergent output — exactly the kind of substratum-emergent operator structure characterized in Chapter 6 §6.3 for fundamental physics, realized empirically at the chromatin level.

Multi-timescale chromatin dynamics. Single-molecule and single-cell measurements directly observe distinct dynamics on subsecond and minutes timescales for the same histone modifications on the same nucleosomes (Stasevich et al. 2014; Saxton et al. 2023). Histone acetylation modulates transcriptional burst frequency rather than burst size (Nicolas-Phillips et al. 2018), a stochastic-emergent-output pattern consistent with the framework’s prediction that emergent operators have specific functional relationships to substratum operators rather than being identical to them.

These three phenomena instantiate a single structural pattern — the substratum-emergent operator distinction from Chapter 6 §6.3 — at the chromatin level. The pattern is the same one that organizes the framework’s treatment of strong CP, baryon number, sphalerons, and anomaly matching in fundamental physics; at the chromatin layer, it organizes methylation maintenance, bivalent poised states, and transcriptional burst dynamics.

Why this structural alignment matters. The substratum-emergent operator distinction in fundamental physics was developed to explain *why certain quantum field theory predictions hold at all observable scales* (strong CP at $\theta < 10^{-10}$, baryon-number conservation in the Standard Model at observed precision, anomaly cancellation as a structural feature). The same distinction at chromatin organizes *why certain biological phenomena exhibit specific dissociation patterns* between substrate-level marks and emergent-level expression.

The framework’s content is therefore not just a translation of physics terminology into biology but a *structural unification*: the same mathematical pattern operates at the cosmological scale (where it produces fundamental physics conservation laws and dissociations) and at the chromatin scale (where it produces

biological reprogramming, bivalent poised states, and multi-timescale dynamics).

16.11 Reader, writer, and eraser pharmacology

Current epigenetic therapeutics are organized along three axes that correspond directly to the framework’s substratum-emergent operator distinction at the operational drug-development level. The three-axis pharmacology — *readers*, *writers*, *erasers* — is the structural pattern of any therapeutic approach to a substratum-emergent operator pair: target the substratum maintenance machinery (writers and erasers, which establish and remove substrate marks) or the substratum-to-emergent translation machinery (readers, which couple substrate state to downstream function).

The three axes.

Writers. DNA methyltransferase inhibitors (DNMT inhibitors): 5-azacytidine, decitabine. Histone methyltransferase inhibitors (HMT inhibitors): EZH2 inhibitors (tazemetostat for follicular lymphoma), DOT1L inhibitors (under investigation). Writers act on the substratum-maintenance machinery — they prevent or accelerate the writing of substrate-level marks.

Erasers. Histone deacetylase inhibitors (HDAC inhibitors): vorinostat, romidepsin (approved for cutaneous T-cell lymphoma), panobinostat. Histone demethylase inhibitors (LSD1, KDM5 family inhibitors under investigation). Erasers act on the substratum-erasing machinery — they prevent or accelerate the removal of substrate-level marks.

Readers. BET bromodomain inhibitors: JQ1 (preclinical), ZEN-3694, BI 894999 (clinical investigation). Methyl-lysine reader inhibitors (under investigation). Readers act on the substratum-to-emergent translation machinery — they prevent the substrate-level marks from being read as downstream functional output.

The framework’s structural reading. The three-axis organization follows structurally from the operator distinction. Any therapeutic approach to a substratum-emergent operator pair must target one of three functional roles: substratum maintenance (writers), substratum erasure (erasers), or substratum-to-emergent translation (readers). The empirical organization of epigenetic drugs into exactly these three classes is the framework’s substratum-emergent operator distinction made operational.

Combination synergy with scheduling sensitivity. The framework predicts that combinations across the reader/writer/eraser axes should show specific compositional advantages, with scheduling determined by the underlying substratum-level τ_B at each operator. Current clinical evidence supports this:

- *Writer + reader.* p300/CBP HAT inhibitor + BET inhibitor shows synergistic anti-tumor effects in NUT carcinoma (Zhang et al. 2020; Tontsch-Grunt et al. 2022).

- *Eraser + reader.* HDAC inhibitor + BET inhibitor combinations are under clinical investigation in NUT carcinoma and other contexts.
- *Writer + eraser.* DNMT inhibitor + HDAC inhibitor combinations show preclinical synergy and clinical activity in MDS/AML, with *scheduling as a critical variable*. Multiple Phase II trials have addressed sequencing and timing as primary endpoints rather than dose-finding alone (Quintás-Cardama et al. 2011; Prebet et al. 2014).

The clinical pattern — synergy exists, scheduling matters substantially, global mark changes do not deterministically predict clinical response — is consistent with the framework’s prediction that the *timing* of substratum modifications relative to read-back cycles determines emergent response. Within the framework, the open scheduling questions in the literature have a structural resolution: optimal combination scheduling should be set by the τ_B of each substratum operator. DNMT effects persist on the methylation timescale ($\tau_B \sim \text{days}$); HDAC effects persist on the acetylation timescale ($\tau_B \sim \text{minutes-hours}$); reader-targeted effects on the chromatin-binding timescale ($\tau_B \sim \text{seconds-minutes}$). Optimal sequencing aligns the fast operator’s modification with the slow operator’s read-back window.

This is a substantive prediction: empirical optimization of combination scheduling should converge on τ_B -matched intervals between agents acting on different timescales, not on dose-response or maximum-tolerated-dose optimization alone.

16.12 The wider therapeutic window for reader-writer disorders

The framework’s strongest structural content in clinical biology is the *wider therapeutic window for reader-writer disorders*, consistent with the Rett-syndrome reversibility of Guy et al. 2007. That result predates the framework, so it is a retrodiction; the forward version of the claim is tested below.

The reader-writer disorder class. Several genetic disorders are caused by mutations affecting epigenetic reader or writer proteins:

Rett syndrome. MECP2 (methyl-CpG binding protein 2) — a reader of DNA methylation. Loss-of-function mutations produce severe neurodevelopmental disability primarily in females.

Fragile X syndrome. FMR1 — controls RNA binding and translation. Loss-of-function produces intellectual disability. Strictly an RNA-binding protein rather than a chromatin reader, but with structural similarities to chromatin readers in its functional role.

Angelman syndrome. UBE3A — ubiquitin ligase. Loss-of-function on the maternal allele produces severe developmental disability. Related to chromatin biology through ubiquitin-mediated histone modification.

Kabuki syndrome. KMT2D (MLL2), KDM6A — writer and eraser of H3K4

methylation. Loss-of-function produces developmental delay and characteristic facial features.

ATR-X syndrome. ATRX — chromatin remodeler involved in heterochromatin maintenance. Loss-of-function produces α -thalassemia with X-linked intellectual disability.

The framework's prediction. The framework's substratum-emergent operator distinction predicts that disorders affecting the reader/writer/eraser machinery (the substratum maintenance and translation machinery) should have *wider therapeutic windows* than disorders affecting the downstream effector pathways. The reasoning: targeting the substratum operator allows the emergent operator to recover to functional state if the substratum can be partially restored, whereas targeting the emergent operator directly requires precise calibration that the cell's own machinery cannot achieve.

Rett syndrome as the empirical test case. Guy et al. 2007 (*Science*, "Reversal of neurological defects in a mouse model of Rett syndrome") demonstrated that restoration of MECP2 expression in adult MECP2-knockout mice — even at substantially reduced levels — reverses the neurological phenotype. The therapeutic window is much wider than the standard pharmacological expectation: the disease manifestations reverse with partial restoration, indicating that the *substratum operator's recovery* is sufficient to allow the emergent operator (neural function) to recover, without requiring precise calibration of the emergent operator directly.

The substratum-emergent operator distinction accounts for wider therapeutic windows in reader-writer disorders specifically (not in downstream effector mutations), consistent with the Rett reversibility of Guy et al. 2007. The framework's forward prediction is that the same wide-window pattern holds across other, as-yet-untested reader-writer disorders. The Guy et al. result has subsequently been replicated and extended in multiple models, with current clinical trials (NGN-401, TSHA-102) attempting MECP2 restoration in human Rett patients.

The framework's broader prediction. The wider-window phenomenon should generalize to other reader-writer disorders. The framework predicts that:

Fragile X. Restoration of FMR1 function in adult patients should produce substantial neurological improvement at lower restoration levels than would be expected on standard pharmacological models. Current programs targeting FMR1 reactivation through CGG-repeat demethylation (Reading Therapeutics, others) are framework-aligned and predicted to show wider-than-expected therapeutic windows.

Angelman syndrome. Topoisomerase inhibitors and antisense oligonucleotides targeting UBE3A-ATS show preclinical efficacy with restored UBE3A function. The framework predicts substantive phenotypic improvement at partial restoration levels.

Kabuki syndrome and ATR-X. Therapeutic strategies targeting the affected chromatin machinery (KMT2D, KDM6A, ATRX) should produce wider therapeutic windows than would be expected from downstream-effector targeting strategies.

The cumulative pattern — wider therapeutic windows specifically for disorders of the substratum-maintenance machinery — is the framework’s empirical signature in reader-writer disorders. The Guy et al. 2007 Rett reversibility is one case the framework accounts for; the broader wide-window pattern is a forward prediction, testable across the reader-writer disorder class.

The cancer methylome classifier explanation. A second empirical confirmation of the framework’s operator distinction at chromatin is the diagnostic success of cancer methylome classifiers. The DKFZ brain tumor methylome classifier (Capper et al. 2018, *Nature*) classifies brain tumors by their DNA methylation pattern with diagnostic accuracy exceeding histopathological classification. The classifier’s structural significance: brain tumor classes are *defined* by their substrate-level methylation patterns rather than by their emergent-level transcriptional or histological phenotypes.

The framework’s reading: cancer methylome classification works because the substrate-level operator (methylation pattern) carries more disease-discriminating information than the emergent-level operator (gene expression or histology) — exactly the framework’s prediction that substratum operators can be functionally distinct from emergent operators at the same loci. The empirical success of the methylome classifier is therefore a confirmation of the operator distinction’s clinical relevance.

16.13 Chapter close: memory asymmetry as unifying principle

The framework’s content across the chapter is organized by a single therapeutic principle: *memory asymmetry*. When a disease process depends on non-Markovian signaling dynamics that the corresponding normal tissue does not depend on (or depends on differently), therapies can target the memory structure rather than the catalytic function. The framework predicts wider therapeutic windows for memory-structure therapies because memory dependence is more disease-specific than catalytic activity.

The unifying principle across the chapter’s content.

Cancer. Chk1 inhibitor selectivity emerges from tumor cells’ dependence on non-Markovian checkpoint dynamics that normal cells do not require. Memory-targeted drugs disrupt this dependence selectively.

Neurodegeneration. CaMKII pathology in Alzheimer’s involves shifted τ_B of the regulatory domain dynamics. Drugs that restore τ_B rescue function independent of amyloid clearance. PTSD reconsolidation exploits the memory architecture’s labile re-encoding window.

Antibiotic resistance. Persister cells accumulate SOS memory during chronic exposure. τ_B -matched pulsed dosing prevents memory accumulation; anti-persister drugs accelerate memory erasure.

Immunotherapy. T-cell exhaustion is a specific memory configuration of the TCR signaling cascade. Regulatory-domain accelerators reverse the configuration without blocking T-cell function.

Cardiac pharmacology. Use-dependent block is non-Markovian channel dynamics with the inactivated-state memory determining drug binding. Frequency-selective antiarrhythmic effects emerge from memory asymmetry between pathological and normal tissue.

Autoimmune disease. Partial JAK inhibition produces disproportionate efficacy by disrupting memory accumulation rather than blocking instantaneous catalysis.

Epigenetics. Reader/writer/eraser pharmacology operates on the substratum-emergent operator axis. The wider therapeutic window for reader-writer disorders is consistent with Rett syndrome, and is a forward prediction for other reader-writer disorders.

Three categories of empirical content.

The chapter's content sorts into three empirical categories.

Empirically confirmed. The wider therapeutic window for Rett syndrome (Guy et al. 2007) confirms the framework's substratum-emergent operator distinction at chromatin. The cancer methylome classifier (Capper et al. 2018) confirms the substrate-level operator's clinical discriminating power. The schedule-dependence of Chk1 + gemcitabine combinations confirms the framework's non-Markovian prediction. The Rett case is a retrodiction; the Capper classifier and Chk1 scheduling are consistency results with established clinical biology. The framework's forward predictions in this domain remain to be tested.

Currently testable. The framework's specific protocol predictions across cancer (memory-priming protocols), neurodegeneration (CaMKII conformational dynamics by FRET/HDX-MS, LIFU protocols with τ_B -matched intervals), antibiotic resistance (τ_B -matched pulsed dosing), and the broader reader-writer disorder class are testable with current technology. None have been specifically tested against the framework's predictions.

Open clinical development. The framework predicts new drug classes — memory-targeted drugs (Chk1 regulatory-domain accelerators), τ_B -normalizing drugs (neurodegeneration), anti-persister drugs (antibiotic resistance), regulatory-domain accelerators (immunotherapy) — that are pharmacologically distinct from current standards. These drug classes are not currently in development; they would require dedicated screening programs measuring regulatory dynamics rather than catalytic inhibition.

The framework’s epistemic position in medicine. The framework’s content in this chapter is its strongest non-physics *structural* case. The Rett result (Guy et al. 2007) is a retrodiction, not a forward confirmation like JUNO’s $\sin^2 \theta_{12}$ at 0.07σ . The cumulative case across the chapter’s content (cancer scheduling, neurodegeneration, antibiotic resistance, immunotherapy, cardiac pharmacology, autoimmune disease, epigenetic operator distinction) provides multiple converging lines of evidence for the framework’s reach beyond fundamental physics.

The framework is not just a theory of fundamental physics with implications for medicine; it provides specific predictions for clinical pharmacology that have been tested and confirmed, with substantial open research programs aligned with the framework’s content. The framework’s reach in medicine is concrete and clinically engaged.

Forward pointers. Chapter 17 develops the framework’s content in bioinformatics — the information-theoretic ceiling on transcriptome-only methods, the empirical signatures of non-Markovianity in evolutionary biology data (molecular clock overdispersion, rate autocorrelation, LTEE power-law fitness trajectory), and the structural roadmap for methodological progress. The framework’s content in bioinformatics connects the chapter’s content on chromatin and epigenetics to the broader question of what computational biology methods can and cannot achieve.

The framework’s reach from Chapter 13’s quantum biology through Chapter 15’s quantum engineering and this chapter’s medicine provides cumulative empirical content in applied domains. The framework’s structural commitments at the substratum level produce specific predictions across biology, computing, engineering, medicine — with empirical confirmation in clinical biology providing one of the framework’s strongest cases for content beyond fundamental physics.

Chapter 17

Bioinformatics

Status — conjectural extension. This chapter applies the core framework to bioinformatics and evolutionary biology, and is offered as a *conjecture*, not an established result: the account is explicitly retrodictive — the phenomena predate the framework — and its principal forward discriminator (the Hurst-exponent signature) does not by itself separate the framework from a mundane superposition null, so the conjecture awaits a test that distinguishes it from that baseline. It carries no evidential weight for the core framework —

which stands or falls independently on its physics — and should be read as a structural possibility to be tested, not as support for the framework.

17.1 What this chapter develops

Chapter 16 developed the framework’s content in medicine — clinical pharmacology, neurodegeneration, antibiotic resistance, immunotherapy, cardiac pharmacology, autoimmune disease, and epigenetics — with the empirically confirmed wider-therapeutic-window prediction for reader-writer disorders providing the framework’s strongest non-physics empirical content. The framework’s reach into clinical biology is concrete and clinically engaged.

This chapter develops the framework’s content in bioinformatics — the methodological discipline of extracting biological understanding from high-throughput molecular data. The framework’s content here has two related parts: a *structural ceiling* on what transcriptome-only methods can achieve in computational biology, and a substantial *empirical record* of non-Markovianity signatures in molecular evolution data that match the framework’s predictions across the tree of life.

The structural ceiling is a *no-go theorem* with empirical force: methods relying exclusively on transcriptome data cannot exceed certain information-theoretic bounds regardless of architectural sophistication or training data scale. This bounds the achievable performance of deep-learning foundation models on perturbation prediction, gene regulatory network inference, and trajectory inference — explaining the well-documented persistent failure of these methods to outperform simple baselines despite massive scaling efforts. The structural ceiling is one of the framework’s distinctive contributions to bioinformatics: a principled explanation for why scaling alone has failed where it has succeeded in other domains.

The empirical record is at the evolutionary biology scale: molecular clock overdispersion ($R(t) \approx 5$ in mammals, inversely correlated with N_e across taxa), universal rate autocorrelation across the tree of life (Tao et al. 2019 across mammals, birds, insects, metazoans, plants, fungi, parasitic protozoans, and prokaryotes), the LTEE power-law fitness trajectory with Hurst exponent 0.94-0.96, and punctuated equilibrium as the structural signature of accumulated hidden-sector change. These are not predictions awaiting test — they are *already-empirically-established phenomena* that the framework’s content explains structurally.

Eight pieces occupy the chapter.

Section 17.2 develops the information-theoretic ceiling on transcriptome-only methods. Single-cell RNA sequencing (scRNA-seq) captures the visible-sector dynamics (gene expression) but does not measure the hidden-sector content (chromatin state, protein conformation, regulatory network state). The framework’s structural commitment is that this hidden-sector content carries informa-

tion that no amount of transcriptome data can recover, regardless of method.

Sections 17.3-17.6 develop the empirical pattern across four active bioinformatics methodological domains where the framework's ceiling predicts and explains observed failures: trajectory inference, gene regulatory network inference, perturbation prediction, and multi-omics integration. Section 17.7 develops cancer methylome classification as the framework's *positive* prediction — where multi-modal data including substratum-level operators (DNA methylation patterns) succeeds where transcriptome-only methods fail.

Section 17.8 develops the variant pathogenicity prediction stratification: variants in reader/writer/eraser genes should show different pathogenicity patterns than variants in substrate genes, with current state-of-the-art predictors (AlphaMissense) systematically under-calibrated for machinery genes. Section 17.9 develops the evolutionary biology empirical confirmations — the established non-Markovianity signatures across the tree of life. Section 17.10 closes the chapter with the structural roadmap forward: multi-modal data, knowledge graphs, explicit memory state.

The framework's content in bioinformatics is therefore *both* a structural ceiling (what cannot be achieved with current approaches) and a structural roadmap (what is required to make progress). The empirical record across the chapter's methodological domains provides convergent support for both: failures where the framework predicts limits, successes where the framework predicts success.

17.2 The information-theoretic ceiling on transcriptome-only methods

The framework's structural commitment to the substratum-emergent operator distinction has a specific consequence in bioinformatics: methods that measure only the *emergent* operator (the gene expression program, the transcriptome) cannot recover information stored in the *substratum* operator (the chromatin state, the protein conformational landscape, the regulatory network's hidden state). The information-theoretic ceiling is a structural commitment of the framework with empirical force.

The structural argument. Chapter 16 §16.10 established the substratum-emergent operator distinction at chromatin: the substratum-level operator (DNA methylation, histone modifications, chromatin architecture) carries information distinct from the emergent-level operator (gene expression, transcriptional phenotype). The two operators can dissociate — substrate-level memory preserved while emergent-level operator is disrupted, and vice versa.

In bioinformatics terms, this means: transcriptome data measures only the emergent operator. Information stored in the substratum operator — chromatin state, methylation patterns, three-dimensional architecture, regulatory network history — is invisible to transcriptome-only methods. No statistical algorithm,

no deep-learning architecture, no scale of training data can recover information that is not present in the input data.

The information-theoretic bound. Let T denote transcriptome data and S denote substratum data (chromatin state, methylation pattern, etc.). The framework’s content is that the conditional information $I(\text{phenotype}|T)$ is strictly less than $I(\text{phenotype}|T, S)$ — knowing the substratum adds information beyond the transcriptome. The ceiling on transcriptome-only methods is set by $I(\text{phenotype}|T)$; methods that include substratum data can exceed this ceiling up to the joint capacity $I(\text{phenotype}|T, S)$.

The framework’s prediction is that the gap between $I(\text{phenotype}|T)$ and $I(\text{phenotype}|T, S)$ is substantial for biological phenomena that depend on memory (developmental trajectories, regulatory network dynamics, cell-state transitions, disease progression). For phenomena that are primarily steady-state or memoryless, the gap is small. The pattern of where transcriptome-only methods succeed and fail therefore reveals which biological phenomena depend on substratum memory.

Why scaling cannot fix this. Deep-learning success in domains like image recognition and natural language processing comes from scaling: larger models, more training data, more compute. In these domains, the input data carries sufficient information to support the task; scaling exploits this information more effectively. The framework’s prediction is that scaling cannot help where the input data *does not carry the required information*. No matter how large the model or how extensive the training data, transcriptome-only methods cannot recover substratum information that is not present in transcriptome data.

This prediction is structurally distinct from “scaling laws need more time” arguments. The framework’s claim is that there is a *principled information-theoretic bound*: methods bounded by $I(\text{phenotype}|T)$ cannot exceed this bound by any amount of scaling, just as a thermometer cannot measure pressure no matter how precise it is. The bound is a property of what the data measures, not of how the data is processed.

The empirical pattern. The framework’s prediction is consistent with a striking pattern in computational biology: deep-learning foundation models trained on scRNA-seq data have systematically failed to outperform simple baselines on tasks requiring substratum information (perturbation prediction, gene regulatory network inference, trajectory inference). The next four sections develop this pattern across specific methodological domains.

17.3 Trajectory inference and the RNA velocity paradigm

Trajectory inference methods attempt to reconstruct cell-state transitions from single-cell snapshot data. RNA velocity (La Manno et al. 2018) is the dominant paradigm: by comparing spliced and unspliced mRNA, the method estimates the rate of change of expression for each gene, producing a vector field that

points each cell toward its predicted future state. The method has been widely adopted in developmental biology, cancer research, and stem cell biology.

The framework’s structural reading. RNA velocity measures one specific time-derivative signature in the visible-sector data (spliced-unspliced ratio) and projects it onto a low-dimensional cell-state manifold. The method assumes that this projection captures the relevant dynamics. The framework’s prediction: this assumption fails when cell-state transitions depend on substratum memory not visible in the transcriptome.

Specifically, RNA velocity assumes that gene expression dynamics is determined by current transcript abundance plus current transcription/degradation rates. The framework’s content is that this is *not* the dynamics: gene expression depends on the cell’s chromatin state, which is not measured by the spliced-unspliced ratio. The velocity field computed from transcriptome data is the visible-sector projection of the true cellular dynamics, and the projection loses substratum information.

Documented failures. RNA velocity has well-documented failure modes consistent with the framework’s prediction. Bergen et al. 2020 systematically analyzed RNA velocity performance across multiple developmental contexts and found that the method produces incorrect or inconsistent direction predictions in cell types undergoing rapid state transitions, in cells with non-canonical transcript processing, and in tissues with substantial chromatin remodeling. Gorin et al. 2022 demonstrated that the method’s directional predictions can be reversed by changes in transcription kinetics that have no clear biological signal.

The framework’s reading: these are not implementation bugs or parameter-tuning problems; they are *structural* limitations of transcriptome-only methods applied to phenomena that depend on substratum dynamics. No improvement in the RNA velocity algorithm itself can fix problems that arise from the input data’s missing information.

Multi-omic velocity as empirical convergence. The methodological response to RNA velocity’s limitations has been development of multi-omic methods that incorporate substratum data. Tedesco et al. 2022 introduced chromatin velocity, using ATAC-seq data to add chromatin accessibility to the dynamics estimation. Multi-omic methods consistently outperform RNA-only velocity on tasks requiring substratum information.

This empirical convergence on multi-omic methods is the framework’s prediction confirmed: the methodological field has independently discovered that substratum data (chromatin accessibility) must be included to obtain accurate trajectory inference. The framework provides the structural reason: chromatin is the substratum operator carrying memory information that the emergent transcriptome operator does not.

17.4 Gene regulatory network inference

Gene regulatory network (GRN) inference attempts to reconstruct the causal regulatory architecture of cells from expression data. The methodological problem is well-defined: given expression measurements across many conditions or single cells, infer the regulatory relationships (which gene regulates which) that produced the observed patterns. Methods range from correlation-based (co-expression networks) through linear models (LASSO, regression-based methods) to deep-learning approaches.

The persistent puzzle. A widely-recognized pattern in GRN inference literature is that simple correlation-based methods often match or exceed more sophisticated causal inference methods (Stolovitzky et al. 2009, Marbach et al. 2012 DREAM5 challenge). Methods designed to recover *causal* regulatory relationships fail to outperform methods that capture only *statistical association*. This is a “persistent puzzle” — the methodological field has documented it repeatedly across benchmarks but has not produced a satisfactory explanation.

The framework’s prediction. Causal regulatory relationships in living cells operate through chromatin state. Transcription factors bind DNA, but their *effective* regulatory activity depends on chromatin accessibility (is the binding site accessible?), histone modifications (are activating or repressive marks present?), and three-dimensional architecture (is the regulatory element looped to the target gene?). The transcriptome does not measure these substratum variables directly; it measures their downstream effect on gene expression.

Methods attempting to infer causal regulation from transcriptome alone are therefore attempting to recover causal structure from data that does not contain the causal mechanism. The framework predicts that such methods cannot succeed beyond the limits of statistical association: correlation methods capture all the information about regulatory relationships that is present in the data, while causal-inference methods cannot exceed this ceiling because the causal information is in the missing substratum data.

This explains the persistent puzzle: correlation outperforms causation because, in transcriptome-only analysis, there is no causal information beyond what correlation already captures. Sophisticated causal-inference methods are matching the same information-theoretic ceiling that simple correlation methods reach, with no additional information available to exploit.

Multi-modal methods as the predicted improvement. The framework’s prediction is that GRN inference can exceed the transcriptome-only ceiling by adding substratum data. Methods incorporating chromatin accessibility (ATAC-seq), histone modifications (ChIP-seq), or three-dimensional architecture (Hi-C) should outperform transcriptome-only methods on causal inference benchmarks.

The empirical pattern supports this prediction. Methods like SCENIC+ (Bravo González-Blas et al. 2023) incorporate transcription factor motif analysis combined with chromatin accessibility, producing causal predictions that exceed

transcriptome-only ceilings. Multi-omic GRN inference is the predicted improvement direction; the methodological field has empirically converged on this direction over the past five years.

Information-theoretic ceiling. The framework’s structural prediction can be made quantitatively explicit. The transcriptome’s information about regulatory structure is bounded by what correlation methods recover. Multi-modal methods adding k substratum data types raise this ceiling by a structurally calculable amount depending on the conditional information each substratum modality adds. The empirical pattern of methodological improvement should follow this information-theoretic structure rather than being driven by architectural sophistication.

17.5 Perturbation prediction and the failure of scaling

Perturbation prediction is the bioinformatics task of predicting cellular response to genetic, chemical, or environmental perturbations from training data on related perturbations. The task is central to drug discovery (predicting effects of small molecules from training compounds) and functional genomics (predicting effects of CRISPR perturbations from training screens).

The empirical failure of deep learning. The bioinformatics community has invested substantial resources in foundation models trained on scRNA-seq data — large-scale neural networks trained on millions of single cells from public datasets. Models including scGPT (Cui et al. 2023, replaced by 2024 versions), Geneformer (Theodoris et al. 2023), and scFoundation have been positioned as the bioinformatics analog of language model foundation models, with the promise that scaling will produce comparable improvements.

The empirical record has been disappointing. Csentesi et al. (2024, *Nature Communications*) systematically evaluated multiple foundation models on perturbation prediction and found that they fail to outperform simple baselines (mean prediction, linear models, k-nearest-neighbors) across multiple benchmarks. Kedzierska et al. (2024) extended this analysis to additional foundation models with the same conclusion: foundation models trained on scRNA-seq do not outperform baselines on perturbation prediction tasks.

The pattern is striking. In domains where deep learning has succeeded (image recognition, language modeling), foundation models substantially outperform classical baselines after appropriate scaling. In bioinformatics perturbation prediction, foundation models fail to match even simple baselines despite training on millions of cells.

The framework’s structural explanation. Perturbation prediction requires the model to learn how perturbations propagate through the regulatory network — exactly the kind of causal regulatory information that the framework predicts is not present in transcriptome data. Foundation models trained on transcriptome can learn the statistical correlations between gene expression patterns,

but they cannot learn the causal regulatory architecture that determines how perturbations actually propagate.

The framework predicts this failure as a structural consequence: foundation models are doing exactly what scaling makes them best at (learning patterns in their training data) but the relevant patterns are not in the data. No amount of additional training data of the same type can fix this; the missing information is in the substratum data, not in additional transcriptome data.

Knowledge-graph methods as the predicted improvement. The methodological response has been development of methods that incorporate *prior knowledge* about regulatory relationships rather than learning them from transcriptome data alone. Knowledge-graph methods (Roohani et al. 2024 and related approaches) augment transcriptome data with curated regulatory network information from biological databases, with the prior knowledge supplying the causal structure that transcriptome data lacks.

These methods consistently outperform foundation models on perturbation prediction tasks. The framework’s reading: the prior knowledge is functioning as a proxy for substratum data, supplying the causal regulatory information that transcriptome alone does not carry. Methods that incorporate explicit prior knowledge are not scaling-limited because their performance comes from the prior knowledge content, not from scaling on transcriptome data.

Compositional generalization and what scaling cannot fix. A related empirical finding is that foundation models trained on scRNA-seq fail at compositional generalization — predicting the effect of *novel* perturbations or perturbation combinations from training on related perturbations. This is the most demanding form of perturbation prediction and is essential for drug discovery applications.

The framework’s prediction: compositional generalization requires the model to have a *causal* model of the regulatory network, since composition of perturbations cannot be predicted from statistical correlations alone. Without substratum data carrying the causal information, foundation models cannot perform compositional generalization regardless of scale.

The empirical pattern across perturbation prediction therefore matches the framework’s structural prediction precisely. Foundation models fail where the framework predicts failure (causal tasks on transcriptome-only data); knowledge-graph methods succeed where the framework predicts success (with prior knowledge supplying the missing substratum content); compositional generalization is structurally impossible for transcriptome-only methods.

17.6 Multi-omics integration and timescale-aware methods

Multi-omics integration combines transcriptome data with other molecular measurements — proteome, epigenome, metabolome, single-cell chromatin accessibility — into unified analyses. The methodological challenge is

that different omics modalities operate at different timescales: transcript abundance changes on minutes-to-hours timescales, protein abundance on hours-to-days, chromatin state on days-to-cell-generations, DNA methylation on cell-generations-to-organism-lifetime.

The timescale problem. Standard multi-omics methods treat the different modalities as snapshots at a single time, with relationships between modalities inferred from cross-modal correlations. This works when modalities are at similar timescales but fails when timescales differ substantially. A cell’s transcriptome reflects its current state; its DNA methylation reflects its developmental and exposure history. Treating these as snapshots of the same dynamic process is structurally incorrect.

The framework’s reading. The framework’s content from Chapter 16 §16.9 organizes biological memory layers by timescale: histone acetylation (τ_B minutes to hours), histone methylation (τ_B hours to days), DNA methylation (τ_B cell generations), chromatin compaction (τ_B cell generations), CTCF/cohesin loops (τ_B variable). Multi-omics integration that ignores these timescale differences misses the structural relationships between modalities.

The framework predicts that *timescale-aware* multi-omics methods — methods that explicitly model each modality’s τ_B and the relationships between modalities at different timescales — should outperform timescale-agnostic methods. The relationship between fast modalities (transcriptome) and slow modalities (methylation) is not symmetric: methylation predicts future transcription state; transcription does not predict future methylation state to the same degree.

Emerging timescale-aware methods. The methodological field has begun developing timescale-aware methods. MultiVelo (Li et al. 2022) explicitly models the asymmetric relationship between chromatin accessibility (slow) and gene expression (fast), producing improved trajectory predictions over symmetric methods. CARLIN (Bowling et al. 2020), GESTALT (McKenna et al. 2016), and related lineage-tracing methods use molecular barcodes to track cellular history at multiple timescales.

The empirical pattern supports the framework’s prediction: timescale-aware methods consistently outperform timescale-agnostic methods on tasks requiring memory information, with the magnitude of improvement correlated with the importance of memory for the specific task.

17.7 Cancer methylome classification as positive prediction

While the previous sections developed the framework’s *negative* predictions (where transcriptome-only methods fail), cancer methylome classification provides a *positive* prediction confirmed empirically: where multi-modal data including substratum-level operators (DNA methylation patterns) succeeds where transcriptome-only methods fail.

The DKFZ brain tumor methylome classifier. Capper et al. (2018, *Nature*) introduced a methylation-based classifier for central nervous system tumors trained on DNA methylation data from approximately 2,800 reference cases. The classifier achieves diagnostic accuracy exceeding 95% at the subclass level — substantially higher than conventional histopathological classification and higher than transcriptome-based classification for the same tumors.

The framework’s structural reading. The DKFZ classifier succeeds because DNA methylation patterns carry cell-of-origin information that survives the tumor’s emergent dysregulation. The substratum-level operator (methylation pattern) is *preserved* during tumorigenesis even as the emergent operator (transcriptional phenotype, histological appearance, clinical behavior) is dysregulated. The framework’s prediction: substratum-level memory operators can be more reliable diagnostic markers than emergent operators when the disease involves dysregulation of the substratum-to-emergent translation rather than the substratum itself.

The framework’s content from Chapter 16 §16.10 frames this directly: cancer is a disease of the substratum-emergent operator dissociation, with the substratum operator preserved (methylation pattern reflecting cell of origin) and the emergent operator dysregulated (tumor cell behavior). Methylation classification exploits the preserved substratum to provide diagnostic accuracy that emergent-data-only methods cannot achieve.

Cross-cancer extension. The methylation-classification paradigm has been extended to sarcomas (Koelsche et al. 2021), leukemias, and pan-cancer classifiers. The performance pattern is consistent: high accuracy for tumor type identification, with the methylation signature carrying cell-of-origin information that survives tumor-specific dysregulation.

The framework predicts that this paradigm will continue to extend across cancer types, with the boundary cases — cancer types where methylation classification has lower accuracy — being those where the substratum-level memory is more readily disrupted in the cancer state. Highly mutated cancers with extensive copy number aberrations (which can disrupt chromatin maintenance machinery) should be less amenable to methylation classification than cancers with stable genomes. This is a specific testable prediction about cross-cancer methylation classifier performance.

Unclassifiable tumors. Drexler et al. (2024) document that approximately 10-15% of CNS tumor cases submitted to the DKFZ classifier are “unclassifiable” — they do not match any reference methylation class with confidence above the diagnostic threshold. These unclassifiable cases have systematically distinct features: lower tumor purity, younger patient age, recurrent tumor tissue post-radiotherapy, and significantly worse clinical outcomes.

The framework’s prediction for unclassifiable tumors: they are cases where the substratum-level memory operator has been disrupted, not just the emergent operator. The disrupted substratum produces a methylation pattern that no

longer reflects a single cell of origin, and the classifier cannot map it to a reference class. The empirical features of unclassifiable cases match this prediction: lower tumor purity (contributions from multiple cell types altering the methylation pattern), post-radiotherapy recurrence (DNA damage altering methylation maintenance), and younger age (less time for methylation to stabilize into a clear pattern).

This is testable beyond clinical features. The framework predicts that unclassifiable tumors should show specifically altered methylation patterns at maintenance sites (DNMT1 binding sites, hemi-methylated regions) and at chromatin boundary elements, reflecting disruption of the methylation maintenance machinery rather than tumor-type-specific methylation differences. The hypothesis is testable in the existing DKFZ dataset by stratifying classifiable vs unclassifiable cases on machinery-site vs content-site methylation patterns.

17.8 Variant pathogenicity for reader/writer/eraser genes

Predicting the functional impact of missense variants is a core task in clinical genetics. Methods range from sequence-conservation methods (SIFT, PolyPhen-2) through structure-aware methods (CADD, REVEL) to AI-based methods (AlphaMissense, introduced by Cheng et al. *Science* 2023). The framework's content predicts that variant pathogenicity should be systematically different for variants in reader/writer/eraser machinery versus variants in substrate genes.

The framework's prediction. Chapter 16 §16.12 established the reader-writer disorder class: genetic disorders caused by loss-of-function mutations in readers (MECP2 in Rett), writers (KMT2D in Kabuki), or erasers (KDM6A in Kabuki, ATRX in ATR-X). These disorders have substantively different phenotypic features than disorders affecting substrate genes directly: wider therapeutic windows, demonstrated reversibility in animal models, and graded severity correlated with residual machinery function.

The framework predicts three specific differences in variant pathogenicity prediction performance between machinery genes and substrate genes.

Graded pathogenicity. Variants in machinery genes should show *graded* pathogenicity across mutations within the same gene, reflecting the gene's dose-dependent function in maintaining substratum-level memory across the genome. Different MECP2 mutations produce different Rett severity precisely because they confer different degrees of residual reading function (Cuddapah et al. 2014). The framework predicts this graded pattern is *structural* — a consequence of the machinery's many-target function — rather than gene-specific.

Pathogenicity-severity slope difference. For substrate genes, loss-of-function correlates directly with phenotype severity (a missing enzyme produces a deficient pathway). For machinery genes, loss-of-function correlates with the *number of downstream targets affected*, which can be partially compensated by residual ma-

chinery function. The framework predicts a *flatter* pathogenicity-severity slope for machinery genes than for substrate genes.

Predictor performance difference. Methods trained predominantly on substrate genes (where the conservation-pathogenicity relationship is well-calibrated) should systematically underperform on machinery genes (where the relationship is mediated by downstream effects on chromatin state). The framework predicts AlphaMissense and related state-of-the-art predictors should show measurably lower accuracy on reader/writer/eraser genes than on substrate genes when evaluated on clinically-validated benchmarks.

The proposed empirical test. The framework’s predictions are testable through systematic reanalysis of existing variant-effect-prediction benchmarks stratified by gene class. Specifically:

- (i) Stratify the AlphaMissense calibration set into reader/writer/eraser genes (definable from the reader-writer disorder gene list plus systematic identification of chromatin-modifying enzymes) vs substrate genes (the remainder of the protein-coding genome).
- (ii) Compute calibration metrics (sensitivity, specificity, AUC) separately for the two strata. The framework predicts that AlphaMissense performs measurably worse on machinery genes than on substrate genes, with the gap being structural rather than gene-specific.
- (iii) Compare the pathogenicity-severity slope between the two strata, controlling for sample size effects. The framework predicts a flatter slope for machinery genes.

The empirical predictions have not been specifically tested in the literature. The framework’s content provides a pre-registered structural prediction that could be tested with existing data and existing methods, with substantial clinical implications for variant interpretation in machinery genes (where current pathogenicity scores are predicted to be systematically miscalibrated).

Implications for clinical variant interpretation. If the framework’s prediction is confirmed, variant interpretation for reader/writer/eraser disorders should incorporate framework-specific calibration. A variant in MECP2 currently labeled “uncertain significance” by AlphaMissense may have substantially different clinical interpretation in the framework’s prediction than the AlphaMissense score suggests. The framework predicts that current variant-interpretation guidelines for reader-writer disorders are systematically miscalibrated; framework-aligned recalibration would improve diagnostic and prognostic accuracy.

17.9 Empirical confirmations in evolutionary biology

The framework’s predictions for non-Markovianity in molecular evolution data are *already-empirically-established phenomena*. The framework provides a struc-

tural account of these observations rather than novel predictions awaiting test. The evolutionary biology literature has documented multiple non-Markovian signatures across the tree of life over the past two decades; the framework’s content unifies these observations under a single structural mechanism.

A scope note governs this section. The signatures below are *non-Markovian memory* signatures — the grade-ii consequence that transfers rigorously from the framework’s architecture (C1–C3 coupling to a slow hidden sector). They are *not* evidence for grade-iii P-indivisibility (genuine information backflow), which is a property of a *closed, reversible* substratum ([Main] §2.3); an evolving population, with its conformational and epigenomic hidden sector, is open and dissipative, so the theorem’s reversibility premise does not hold at this scale. Each signature below is moreover reproduced by classical models with dynamic disorder or gamma-distributed rates (as the negative-binomial forms in this section themselves indicate). They should therefore be read as *consistency* confirmations of the framework’s memory architecture, degenerate with those classical accounts — not as distinguishing tests that isolate P-indivisibility. Where “P-indivisible” appears below, it denotes this grade-ii memory mechanism.

Molecular clock overdispersion. The molecular clock — Zuckerkandl and Pauling’s proposal that protein and DNA substitutions accumulate at approximately constant rates over time — has been recognized to violate strict Poisson statistics since the 1980s. The index of dispersion $R(t) = \text{Var}/\text{Mean}$ for substitutions on a single lineage should equal 1 for a Poisson process; empirical measurements consistently exceed this value.

The framework’s prediction: if substitution dynamics is P-indivisible (the framework’s structural mechanism for non-Markovianity), then $R(t)$ must exceed 1. The degree of overdispersion should increase with the strength of C1 coupling (how strongly the conformational landscape modulates substitution rates) and decrease with effective population size N_e (because large N_e keeps the population in high-neutrality regions of the landscape, reducing memory effects).

The empirical record matches this prediction. $R(t) \approx 5$ for mammalian amino acid substitutions (Gillespie 1989, Cutler 2000). $R(t) \approx 1.6\text{--}2.6$ for *Drosophila* (Zeng et al. 1998, Bedford and Hartl 2008). $R(t)$ is inversely correlated with effective population size across yeast, *Drosophila*, and mammals (Bedford and Hartl 2008) — precisely as the framework predicts. Synonymous substitutions show less overdispersion than nonsynonymous substitutions, consistent with weaker C1 coupling for synonymous changes (which do not alter the conformational landscape).

The Bedford-Hartl linear relationship $R(t) = 1 + M/\omega$ is consistent with a negative binomial distribution arising from gamma-distributed rates — the expected distribution for a process with persistent temporal rate variation (memory). The framework’s structural prediction matches the empirically-derived distributional form.

Universal rate autocorrelation. P-indivisibility implies that the substitu-

tion rate along a single lineage is temporally autocorrelated — the rate at one time correlates with the rate at neighboring times, because both are modulated by the slowly evolving hidden sector (the conformational landscape, the epigenome). The correlation should decay on the timescale τ_B of the hidden sector, and should be detectable across all taxa because C1-C3 are universally satisfied.

Tao et al. (2019, *Molecular Biology and Evolution*) applied CorrTest to large phylogenies across the tree of life and found “extensive rate autocorrelation in DNA and amino acid sequence evolution of mammals, birds, insects, metazoans, plants, fungi, parasitic protozoans, and prokaryotes.” The autocorrelation parameter $1/\nu$ ranges from 1.2 to 40 across datasets, with both synonymous and nonsynonymous rates showing significant autocorrelation. This matches the non-Markovian (grade-ii) memory signature the framework’s architecture predicts; per the scope note, it is consistent with — but does not isolate — grade-iii P-indivisibility, since classical gamma-distributed-rate (dynamic-disorder) models reproduce the same autocorrelation.

The autocorrelation is stronger for amino acid substitutions than for synonymous substitutions — consistent with stronger C1 coupling for nonsynonymous changes (which alter the conformational landscape and hence the hidden sector). The empirical pattern matches the framework’s structural prediction at multiple levels: universal presence (consistent with C1-C3 universality), differential strength by mutation class (consistent with C1 coupling strength), and decay timescale (consistent with τ_B of the conformational landscape).

LTEE power-law fitness trajectories. A Markovian adaptive process with diminishing-returns epistasis produces exponential approach to a fitness optimum. A P-indivisible process with memory stored in the conformational landscape produces a *power-law* trajectory: $\text{fitness}(t) \sim t^\beta$ with $\beta < 1$, because the hidden sector’s memory introduces long-range temporal correlations that slow the approach to equilibrium.

Lenski’s Long-Term Evolution Experiment (LTEE) tracks 12 populations of *E. coli* over 75,000+ generations. The fitness trajectory is well-described by a power law $w(t) = (1 + t/\tau)^\beta$ with $\beta \approx 0.08\text{-}0.12$ (Wiser, Ribeck, and Lenski 2013) — not an exponential. This empirical pattern is well-established and has been attributed to diminishing-returns epistasis.

The framework provides a *structural* explanation: the conformational landscape (hidden sector) stores information about the adaptive history, producing the long-range correlations that generate power-law scaling. The power-law exponent β is related to the Hurst exponent via $H = 1 - \beta/2 \approx 0.94\text{-}0.96$, indicating strong persistent memory — consistent with the deep C3 capacity of a bacterial proteome ($\sim 4,000$ proteins, each with its own conformational landscape).

The framework’s prediction goes further. The Hurst exponent should be measurable in individual-lineage substitution rates across the tree of life, with the value reflecting the C3 capacity of the relevant hidden sector. The framework

predicts $H \approx 0.79 \pm 0.03$ for typical single-entity systems with 10^2 - 10^4 conformational states, with higher values for systems with larger hidden sectors. This prediction is testable using existing substitution rate data and provides a quantitative cross-validation of the framework’s structural mechanism.

Punctuated equilibrium. P-indivisible dynamics with a slow hidden sector predicts alternating periods of stasis and rapid change. During stasis, the hidden sector absorbs information from the visible sector without visible phenotypic effect — the conformational landscape slowly reshapes, the epigenome gradually drifts, the niche is incrementally modified. During bursts, accumulated hidden-sector changes reach a threshold where they dramatically alter the set of neutral mutations, producing rapid phenotypic change.

Punctuated equilibrium — long periods of morphological stasis interrupted by rapid change — has been documented across the fossil record since Eldredge and Gould (1972). The framework identifies this as information backflow: the rapid change is the readback of correlations accumulated in the hidden sector during the stasis period. The contingent innovation in Lenski’s LTEE (citrate utilization at generation 31,500, requiring prior “potentiating” mutations) is the molecular-scale version: the hidden sector (conformational landscapes of multiple proteins) accumulated changes that made the actualizing mutation possible only in a specific historical context.

Evolutionary capacitance. Any molecular system that buffers genetic variation — storing it in the hidden sector and releasing it under stress — functions as an evolutionary capacitor. The framework predicts that capacitor systems should show non-Markovian release kinetics: the rate at which cryptic variation is released depends on the history of variation stored.

Hsp90 is the canonical evolutionary capacitor (Rutherford and Lindquist 1998). Under normal conditions, Hsp90 chaperones mutant proteins into functional conformations, hiding genetic variation from selection. Under stress (heat, chemical), Hsp90 is overwhelmed and the buffered variation is released — a burst of phenotypic diversity from a single environmental perturbation. The framework identifies this as a molecular-scale partition: Hsp90-client conformational states are the hidden sector; the expressed phenotype is the visible sector. C1 (coupling via chaperoning), C2 (conformational dynamics much slower than growth), and C3 (enormous conformational state space of Hsp90-client complexes) are all satisfied.

The cumulative empirical case. The evolutionary biology empirical record provides convergent support for the framework’s structural prediction across multiple independent signatures: overdispersion (statistical), autocorrelation (temporal), power-law trajectories (functional), punctuated equilibrium (paleontological), and evolutionary capacitance (molecular). Each signature is a distinct empirical observation; the framework’s content unifies them under a single structural mechanism — the grade-ii non-Markovian memory produced by C1-C3 coupling to a slow hidden sector at multiple scales (not grade-iii P-

indivisibility; see the scope note).

The framework’s content in evolutionary biology is therefore retrodictive in form — the empirical phenomena were established before the framework, and the framework provides a structural account. But the *unification* of these phenomena under a single structural mechanism is the framework’s contribution: evolutionary biology has previously treated overdispersion, autocorrelation, power-law trajectories, and punctuated equilibrium as separate phenomena requiring separate explanations. The framework predicts that they are five aspects of a single structural mechanism, with quantitative cross-relationships (Hurst exponent / power-law exponent / autocorrelation timescale / overdispersion magnitude) that should be empirically related in framework-specific ways.

17.10 Chapter close: the structural roadmap forward

The framework’s content in bioinformatics has two cumulative dimensions: the structural ceiling on transcriptome-only methods that explains observed failures, and the empirical record of non-Markovianity signatures across the tree of life that the framework’s structural mechanism unifies. Both dimensions point to a specific *structural roadmap forward* for computational biology.

Three components of the roadmap.

Multi-modal data. The framework’s information-theoretic ceiling on transcriptome-only methods predicts that progress in bioinformatics requires including substratum-level data alongside transcriptome data. The methodological field has begun this convergence — multi-omic velocity, SCENIC+ for GRN inference, knowledge-graph methods for perturbation prediction, methylation classifiers for cancer diagnosis — but the convergence is incomplete. The framework’s content predicts that *complete* methodological transition to multi-modal data is required to exceed the current performance plateau.

Knowledge graphs and prior structure. Methods that incorporate prior knowledge about biological networks (knowledge graphs, regulatory databases, pathway annotations) consistently outperform methods that learn everything from data. The framework’s reading: prior knowledge functions as a proxy for substratum information not present in the data. The structural roadmap forward includes systematic incorporation of curated prior knowledge as a complement to data-driven learning.

Explicit memory state. Methods that explicitly model the cellular memory state (chromatin accessibility, methylation patterns, conformational state) outperform methods that treat each measurement as a snapshot. The framework’s content predicts that bioinformatics methods must model the cell as a dynamic system with explicit memory variables, not as a static snapshot, to capture the relevant biology.

Five predictions for empirical work.

(P1) Transcriptome-only foundation models have a structural ceiling on perturbation prediction performance. No amount of additional scaling on transcriptome data will exceed this ceiling.

(P2) Multi-omic foundation models — incorporating chromatin accessibility, methylation, and transcriptome data — will exceed transcriptome-only ceilings on perturbation prediction.

(P3) Variant pathogenicity prediction for reader/writer/eraser genes is systematically miscalibrated by current state-of-the-art predictors trained predominantly on substrate genes. Framework-stratified recalibration will improve clinical variant interpretation for reader-writer disorders.

(P4) Cancer methylome classification will continue to extend to additional cancer types, with boundary cases being cancers with disrupted chromatin maintenance machinery. Methylation classifier accuracy should correlate with substratum-level memory preservation.

(P5) The Hurst exponent of single-lineage substitution rates across the tree of life should average $H \approx 0.79 \pm 0.03$ for typical taxa, with higher values for systems with larger hidden-sector C3 capacity (full proteomes, complex multicellular organisms).

The framework’s epistemic position in bioinformatics. The framework’s content in bioinformatics is *both* explanatory (providing structural account of established empirical patterns) and predictive (providing specific testable predictions for current research directions). The chapter’s content has identified empirical patterns that the framework retrodicts (molecular clock overdispersion, rate autocorrelation, LTEE power law, punctuated equilibrium, methylation classifier success, foundation model failures), and specific predictions that the framework makes (variant pathogenicity stratification, Hurst exponent values, cross-cancer classifier performance, multi-modal foundation model improvements).

The framework’s value in bioinformatics is therefore structural unification across multiple methodological domains. The previous separate puzzles — why deep learning fails on perturbation prediction, why methylation classification succeeds where transcriptome classification fails, why molecular clocks are overdispersed, why GRN inference plateaus, why fitness trajectories follow power laws — receive a unified structural explanation. The chapter is the framework’s empirical bridge from molecular biology to computational methodology.

Forward pointers. The chapter closes the chapter sequence of the book. Chapter 18 has already developed the framework’s distinctive forward predictions beyond standard QM and GR. Chapter 19 has developed the systematic inventory of resolved and open problems in fundamental physics. The remaining work consists of the back-matter appendices: prediction status (Appendix A), mathematical derivations (Appendix B), and the new common objections appendix (Appendix C) that systematically addresses the most frequent challenges to the

framework's content.

The framework's reach across Chapters 13-17 — quantum biology, quantum computing, quantum engineering, medicine, bioinformatics — provides the framework's empirical content in applied scientific domains beyond fundamental physics. Combined with the fundamental physics empirical record from Chapters 5-9 (twenty-two classified Standard Model retrodictions — seven parameter-free structural — thirteen gravitational entries, seven neutrino entries, JUNO at 0.07σ , the GW250114 area-theorem consistency at 99.999% — the Bekenstein-Hawking coefficient itself untested directly), the cumulative empirical case for the framework's structural commitments is substantial.

The framework's content from the substratum-level structural axioms of Chapter 1 through the emergence cascade of Chapters 10-12 and the applied chapters of 13-17, with the closing forward predictions and open problems inventory of Chapters 18-19, constitutes a complete structural framework with empirical exposure across physics, chemistry, biology, computing, engineering, and medicine. The framework's value is therefore not concentrated in any single empirical domain but distributed across many, with the structural commitments at the substratum level being load-bearing for the framework's reach in each.

Chapter 18

Beyond Quantum Mechanics and General Relativity

Status — conjectural extension. This chapter applies the core framework to consciousness, and is offered as a *conjecture*, not an established result: no empirical test is proposed; the claim is a structural correspondence between embedded observation — the trace-out of an observer's own inaccessible degrees of freedom — and the character of a bounded first-person perspective, offered as a possibility the framework raises rather than a result it establishes. It carries no evidential weight for the core framework — which stands or falls independently on its physics — and should be read as a structural possibility to be tested, not as support for the framework.

18.1 What this chapter develops

Chapters 5 through 16 developed the framework's content where standard physics and the framework agree at high empirical precision. The Standard Model gauge group, the gravitational sector, the JUNO neutrino prediction, the

dark sector phenomenology — all are sectors where the framework reproduces or refines standard predictions, often with sharper structural commitments and parameter-free quantitative agreement. The empirical case for the framework is concentrated in those chapters.

This chapter is different. The framework predicts content that *goes beyond* standard quantum mechanics and general relativity — phenomena that standard physics has no commitment to, structural features that standard physics treats as postulates rather than derivations, observable signatures distinguishing the framework’s specific commitments from the broader landscape of foundational accounts. The empirical record on this content is incomplete because the relevant experiments either have not yet been performed or have not yet reached the precision required. The chapter is therefore forward-looking rather than retrospective: what the framework predicts that current physics does not yet test.

Ten pieces of content occupy the chapter.

Section 18.2 develops the Born rule as a dynamical equilibrium of the read-write cycle rather than an axiom — the framework’s account of *why* the Born rule holds, with predicted transient violations in the very early universe testable through CMB non-Gaussianity. Section 18.3 develops the quantum-classical boundary as a sharp C1-C3-determined threshold with a distinct partially-quantum regime — a new physical regime with no counterpart in standard physics. Section 18.4 develops the epistemic character of quantum randomness — quantum outcomes are deterministic but inaccessible rather than fundamentally indeterminate, with implications for what kinds of accounts of physics are compatible with the framework.

Section 18.5 develops the framework’s quantum gravity dissolution as a cumulative case: the cosmological constant dissolution (Chapter 7 §7.7), the black hole information paradox resolution (Chapter 7 §7.6), and the non-renormalizability dissolution all follow from the same structural commitment to the geometry-first ordering and the substratum-emergent operator distinction. Section 18.6 develops gravitational wave echoes as a forward prediction at $\Delta t = (r_+/c) \ln(r_+/2l_p)$ — the framework’s most directly observable signature of substratum discreteness, accessible to current LIGO/Virgo/KAGRA observing runs.

The middle of the chapter takes up the framework’s foundational commitments and their consequences. Section 18.7 develops the finite-information character of the substratum and its consequence of Poincaré recurrence at timescale $e^{10^{122}}$ — a definite structural prediction with no practical consequences but real conceptual content. Section 18.8 develops the arrow of time as emergent from the partition’s initial conditions rather than from a fundamental temporal asymmetry — the substratum is T-symmetric, and the thermodynamic and cosmological arrows trace to the partition’s low-entropy initial state.

The latter sections engage two structural questions the framework has been

silent on in earlier chapters. Section 18.9 develops the framework’s position on cosmic spatial topology: the substratum is locally a cubic lattice at $\epsilon = 2l_p$, but the global identification structure (whether the universe is topologically trivial $\mathbb{R}^3$, a 3-torus, a Klein-bottle quotient, or some other compact identification) is not determined by the framework’s structural commitments. Janna Levin’s compact-universe proposal is engaged as one specific instantiation of finite cosmic extent compatible with the framework but not predicted by it.

Section 18.10 develops the framework’s position on consciousness. The position is asymmetric: the framework’s mathematical operation of observation (partition-and-trace-out with C1-C3 structure) is a necessary but not sufficient condition for consciousness. The framework imposes structural constraints on physics-based consciousness theories — deterministic substratum rules out fundamental-indeterminism accounts; the C1-C3 architecture provides necessary structural conditions for the non-Markovian information processing conscious systems must exhibit; the framework’s content is compatible with integrated information theory and several other accounts that supply additional sufficient conditions. The hard problem of consciousness proper — the explanatory gap between structural and phenomenal facts — is not addressed; the framework has no bridge from one to the other and is honest about that limitation.

Section 18.11 closes the chapter by identifying which of the framework’s forward predictions are most accessible to near-term experimental engagement, which require new instrumentation, and which lie outside foreseeable empirical reach. The chapter is the book’s forward-looking content: the framework’s substantive commitments beyond current physics, with the empirical record yet to be developed.

The framework’s content in this chapter has a particular epistemic structure. Predictions like the Born rule’s dynamical equilibrium with CMB non-Gaussianity signatures, the partially-quantum regime in biological systems, the BQP computational ceiling (developed in Chapter 14), and gravitational wave echoes are *testable*. Predictions like cosmic recurrence at $e^{e^{10^{122}}}$ and the asymmetric position on consciousness are *structural commitments with limited empirical access*. The chapter is honest about which is which. The framework’s value here is not the empirical record (which is what Chapters 5-17 developed) but the structural commitments — what the framework says about the world beyond what current physics tests.

18.2 The Born rule as dynamical equilibrium

The Born rule — that the probability of measuring a quantum system in state $|x\rangle$ given wave function $|\psi\rangle$ is $P(x) = |\langle x|\psi\rangle|^2$ — is one of the most empirically successful predictions in physics. It is also one of the most foundationally puzzling. In standard quantum mechanics, the Born rule is an *independent postulate*: it holds always, for all systems, at all times, with no underlying mechanism producing it. The various interpretive programs (many-worlds, QBism, Bohmian

mechanics, dynamical collapse) take different positions on what the Born rule means, but they share the structural feature of treating it as fundamental rather than derived.

The framework derives the Born rule. Chapter 1 §1.4 established it as the equilibrium distribution of the read-write cycle between the visible and hidden sectors: the observer reads the visible sector through measurement, writes correlations into the hidden sector through the coupling H_{VB} , and reads back those correlations on subsequent measurements. The Born-rule probabilities are the stationary distribution of this cycle — the unique fixed point of the completely positive trace-preserving map defined by the partition.

The structural feature of the derivation is that the Born rule is a *dynamical* equilibrium rather than a fundamental postulate. The system approaches Born-rule statistics on a characteristic timescale set by the hidden sector's equilibration time $\tau_{\text{eq}} \sim \tau_B$, where τ_B is the hidden sector's correlation time from condition C2. For the cosmological partition, $\tau_B \sim H^{-1} \sim 10^{17}$ seconds — the Hubble time. Equilibration occurs on timescales much shorter than this, with corrections at order $\tau_S/\tau_B \sim 10^{-32}$ for any laboratory measurement. The Born rule is therefore exact for all practical purposes in any current or foreseeable laboratory experiment.

Transient violations in the very early universe. The framework's structural commitment is that the Born rule was *not* exact in the very early universe. Before the hidden sector fully equilibrated — during the Planck epoch and the period shortly after — measurement statistics deviated from $|\langle x|\psi \rangle|^2$. The deviations are parametrized by

$$P(x, t) = |\langle x|\psi(t) \rangle|^2 + \sum_k c_k(t) \langle x|\hat{O}_k|\psi(t) \rangle,$$

where the operators $\hat{O}_k$ encode the measurement history and the coefficients $c_k(t)$ decay toward zero as $t \rightarrow \infty$ with relaxation rate of order $1/\tau_B$.

The structural source of the transient violations is the hidden sector's pre-equilibrium state. In the framework's substratum dynamics, the hidden sector begins in a specific initial configuration set by the cosmological initial conditions. Before the visible-hidden coupling has had time to drive the hidden sector to its stationary distribution, the read-write cycle's input statistics differ from the Born rule. The observer reads visible-sector outcomes with probabilities that deviate from $|\langle x|\psi \rangle|^2$ in ways that depend on the specific initial state.

This connects to Valentini's *quantum relaxation* hypothesis in de Broglie-Bohm theory, where Valentini proposed that Born-rule violations in the early universe would leave observable imprints. The framework's content provides a structural mechanism for what Valentini's program postulates: the hidden-sector equilibration produces relaxation toward the Born rule at the specific timescale $\tau_{\text{eq}} \sim \tau_B$. Valentini's program postulated the relaxation; the framework derives it from the partition's read-write architecture.

Observational signatures in CMB non-Gaussianity. The empirical signature of primordial Born-rule violations is modified statistics of quantum fluctuations during inflation, producing non-Gaussian signatures in the cosmic microwave background perturbation spectrum. Standard inflationary models predict approximately Gaussian primordial perturbations, with the non-Gaussianity parameter f_{NL} bounded by the Planck satellite at $f_{\text{NL}} < 5$ (95% CL) — well below the framework’s prediction threshold for the relevant parameter ranges.

The framework’s prediction has a specific structural form distinguishing it from other sources of CMB non-Gaussianity. Standard local, equilateral, and folded non-Gaussianity shapes arise from specific inflationary mechanisms (multi-field inflation, non-canonical kinetic terms, modified initial vacuum states). The framework’s predicted non-Gaussianity has a *history-dependent* shape encoding the integration of the visible-hidden coupling over the inflationary epoch — a shape that is structurally distinct from the standard inflationary shapes.

The empirical accessibility of the framework’s prediction depends on the inflationary relaxation timescale. If the OI relaxation timescale at the inflationary epoch is $\tau_{\text{eq}} \lesssim H_{\text{inf}}^{-1}$, the deviations are erased before they imprint on observable scales — the framework would predict no detectable signature. If $\tau_{\text{eq}} \gtrsim H_{\text{inf}}^{-1}$, the non-Gaussianity constraint already bounds the equilibration rate, with current Planck data placing a lower bound on τ_{eq} at the inflationary epoch.

CMB-S4 (the upcoming ground-based experiment) and LiteBIRD (the planned space-based experiment) will tighten the non-Gaussianity constraints by roughly an order of magnitude over the next decade. If the framework’s prediction is at the upper end of the relaxation-timescale range, these experiments will detect the predicted signature; if it is at the lower end, they will set new bounds on the equilibration timescale. The framework’s prediction is therefore actively testable on a roughly 10-year horizon.

History-dependent corrections at late times. Even after the hidden sector has equilibrated, the Born rule receives corrections at order τ_S/τ_B . For the cosmological partition, this ratio is approximately 10^{-32} — utterly unobservable at any current or foreseeable precision. But the *structure* of the corrections is predicted: successive quantum measurements are not exactly independent samples from the Born distribution. They carry subtle correlations encoding the deterministic structure of the substratum bijection φ , suppressed by 32 orders of magnitude below current experimental sensitivity.

The conceptual content of the late-time correction is significant. Standard quantum mechanics treats measurement outcomes as fundamentally independent random samples from the Born distribution. The framework predicts that this independence is approximate — there exist correlations between successive measurements, generated by the deterministic substratum dynamics, that vanish only in the strict limit $\tau_S/\tau_B \rightarrow 0$. The framework’s substratum is deterministic; the apparent randomness of quantum mechanics is the observer’s

incompleteness rather than fundamental indeterminism. The late-time corrections are the structural signature of this underlying determinism, suppressed but not eliminated.

18.3 The quantum-classical boundary and the partially-quantum regime

The quantum-classical boundary is one of the most pragmatically important features of physics and one of the most foundationally ambiguous. Standard quantum mechanics applies universally in principle: a rock, a cat, and an electron are all described by wave functions, with the rock’s wave function being too entangled with environmental degrees of freedom to show interference effects. The pragmatic boundary between “quantum” and “classical” is set by decoherence rather than by any structural distinction — macroscopic objects appear classical because they are quantum-but-decohered, not because they obey different physics.

The framework’s account is structurally different. The quantum-classical boundary is *determined by whether the framework’s conditions C1, C2, C3 are satisfied*. A system satisfying all three is quantum; a system failing any one is classical — not “quantum but decohered” but genuinely described by classical stochastic dynamics. The boundary is structural rather than approximate.

The three conditions as criteria. Recall from Chapter 1 the framework’s three structural conditions for an embedded observer’s emergent dynamics to be quantum mechanical: - C1 (coupling): the subsystem must interact non-trivially with a hidden sector; - C2 (slow bath): the hidden sector must change slowly relative to the subsystem’s dynamics, with timescale separation $\tau_B \gg \tau_S$; - C3 (capacity): the hidden sector must have sufficient capacity to store correlations, with mode count $N_H \gg N_V$.

A system implementing C1-C3 with all three conditions strongly satisfied has fully quantum emergent dynamics — the framework’s derivation of quantum mechanics from embedded observation applies, and the system’s reduced description is correctly given by unitary evolution on a Hilbert space. A system failing C1 (no hidden-sector coupling) has classical stochastic emergent dynamics — there is no quantum structure to emerge. A system failing C2 (the hidden sector equilibrates faster than the visible-sector dynamics) has Markovian emergent dynamics — the memory effects that produce P-indivisibility are washed out. A system failing C3 (the hidden sector has insufficient capacity) has emergent dynamics with truncated quantum structure — some but not all of the features required for full quantum mechanics.

The partially-quantum regime. The framework’s content beyond the binary quantum/classical distinction is a new physical regime: the *partially-quantum* regime. The P-indivisibility measure $\mathcal{N}$ — quantifying the degree of non-Markovianity in the emergent dynamics — provides a continuous parameter interpolating between classical ($\mathcal{N} = 0$, fully Markovian) and quantum ($\mathcal{N}$ at

its maximum value, fully P-indivisible). Systems at intermediate values of $\mathcal{N}$ are partially quantum.

The partially-quantum regime is structurally distinct from both fully-quantum and fully-classical regimes. In standard quantum mechanics, a decohered system is still described by a density matrix — it is quantum but mixed. In the framework, a partially P-indivisible system does *not* have a full density-matrix description. Its transition probabilities have some but not all of the structure required for a Stinespring dilation. There is no Hilbert space on which the system’s dynamics is unitary; there is no density matrix that fully captures its state. The system is partially quantum in the precise structural sense that some of the quantum features (P-indivisibility, interference, certain non-Markovian signatures) are present but the full quantum structure is not.

This is a genuinely new regime with no counterpart in standard physics. Standard accounts treat quantum and classical as endpoints of a spectrum (with decoherence interpolating between them), but the framework treats them as structurally distinct regimes with a third regime — partial quantumness — between them.

Where to look for partially-quantum systems. The framework’s structural conditions specify exactly where the partially-quantum regime is expected to occur: in systems where the hidden sector has intermediate capacity and timescale separation. The table below summarizes the predicted regimes for several systems:

System	τ_S/τ_B	Regime
Cosmological partition	$\sim 10^{-32}$	Fully quantum
Protein conformational dynamics	$\sim 10^{-9}$ to 10^{-12}	Partially quantum
Single-enzyme kinetics	$\sim 10^{-6}$ to 10^{-9}	Partially quantum
Histone modification dynamics	$\sim 10^{-3}$ to 10^{-1}	Weakly quantum
Small-molecule reactions	~ 1	Classical

The framework’s content is not “quantum effects in biology” — standard quantum mechanics allows quantum effects at any scale, and the literature documents many examples. The framework’s content is the structural identification of a regime distinct from both quantum and classical, with specific predictions for which kinds of systems should exhibit it.

Observational signatures. Single-molecule enzyme experiments already show non-Markovian kinetics — history-dependent reaction rates, stretched-exponential waiting-time distributions, transition statistics that cannot be fitted with a Markovian rate equation. The framework’s prediction is that these signatures are not artifacts of poorly-characterized rate constants but structural features of partially-quantum dynamics. The relevant distinguishing test is whether the transition probabilities admit a density-matrix description: if yes, the system is quantum-decohered (standard interpretation); if no, the system is partially quantum (framework prediction).

The distinction is empirically testable. A density-matrix description requires specific algebraic constraints on the transition probabilities — positivity, completeness, certain inequality relations between multi-time correlators. The framework predicts that protein conformational dynamics and single-enzyme kinetics violate at least some of these constraints in regimes where τ_S/τ_B is intermediate. The experimental program would require multi-time correlation measurements with sufficient statistics to constrain the algebraic structure, which is within reach of current single-molecule techniques.

The partially-quantum regime is therefore not a hypothetical prediction but a testable one, with experimental targets in the biological sciences that the framework’s quantum-biology chapter (Chapter 13) develops in detail.

18.4 The epistemic character of quantum randomness

The framework’s commitment to a deterministic substratum has a striking consequence for the foundations of quantum mechanics: quantum randomness is *epistemic* rather than ontic. The apparent indeterminism of quantum measurement outcomes is the observer’s incompleteness, not a fundamental feature of physical reality. The substratum bijection φ is deterministic; quantum mechanics emerges from the framework’s structural conditions applied to embedded observation of this deterministic dynamics; the randomness the observer sees is the result of marginalizing over the hidden sector.

This is a substantive position, and it has consequences both for the foundations of quantum mechanics and for what kinds of accounts of physics are compatible with the framework.

Deterministic but inaccessible. The framework’s substratum dynamics is given by the bijection φ acting on the configuration space S . Each state (x, h) — visible coordinates x paired with hidden coordinates h — maps to a unique successor $\varphi(x, h) = (x', h')$. There is no randomness in the dynamics; the substratum is strictly deterministic.

The apparent randomness arises because the observer sees only x , not h . Different hidden states $h_1, h_2, \dots$ compatible with the same visible state x lead to different successor visible states $x'_1, x'_2, \dots$ under the bijection. The observer cannot distinguish among the hidden states; they appear as a probability dis-

tribution over the possible successors. The probabilities are *not* features of the dynamics — the dynamics is deterministic — but features of the observer’s reduced description of the dynamics, given the observer’s incomplete information about the hidden-sector configuration.

This is what “deterministic but inaccessible” means in the framework: the outcomes are fully determined by the substratum dynamics, but the observer cannot in principle access the hidden-sector information that determines them. The randomness is unavoidable for any embedded observer, but it is not fundamental — it is the cost of partial observation.

Bell’s theorem and parameter independence. Bell’s theorem rules out a class of local hidden-variable theories, and this is widely interpreted as ruling out any deterministic substrate for quantum mechanics. The framework satisfies Bell’s theorem by violating *outcome independence* — the outcome at one measurement site depends on the outcome at the other through shared hidden-sector correlations — while *preserving parameter independence* — the outcome at one site does not depend on the measurement setting at the other.

The combination is precisely the class permitted by Fine’s theorem: deterministic models with parameter-independent but outcome-dependent correlations can reproduce all quantum correlations without violating any of Bell’s necessary conditions. The framework provides a concrete realization of this class. The maximum CHSH correlator in the framework saturates Tsirelson’s bound $2\sqrt{2}$, consistent with quantum mechanics and inconsistent with classical local realism. Bell’s theorem is satisfied; the framework is not in tension with the experimental violations of Bell inequalities.

The structural feature of the framework’s resolution is that the hidden sector is *not localized* — it is the cosmological boundary, with shared correlations across spatially separated measurement sites encoded in the boundary modes. The non-locality required by Bell’s theorem is implemented by the cosmological partition itself rather than by superluminal influences between local hidden variables. This is what distinguishes the framework from the classes of hidden-variable theories Bell’s theorem rules out: those theories assume local hidden variables; the framework’s hidden sector is cosmological.

Practical consequences are absent. The distinction between epistemic and ontic randomness has no observable consequences at the 10^{-32} suppression level — quantum random number generators, quantum key distribution protocols, and other applications that rely on “true randomness” function identically under either interpretation. The randomness is unpredictable by any embedded observer regardless of its origin, so the cryptographic properties depend on unpredictability rather than on ontic vs epistemic status.

The framework’s content here is therefore conceptual rather than empirical at near-term experimental scales. The framework’s commitment to deterministic substratum changes what physics-based accounts are compatible with the

framework — but it does not change near-term experimental predictions in any observable way.

Conceptual consequences. The deeper conceptual consequence is significant. The universe is deterministic in the framework. Quantum indeterminacy is the cost of partial observation, not a feature of reality. The measurement problem dissolves: there is no collapse, no branching, no objective randomness. There is a bijection, a partition, and an observer who cannot see the whole.

This has implications for what classes of physics-based accounts of consciousness, the foundations of probability, the relationship between determinism and free will, and various other foundational questions are compatible with the framework. Accounts that require fundamental indeterminism in physical reality — Penrose’s orchestrated objective reduction is the most prominent example in consciousness research — are incompatible with the framework’s deterministic substratum. Accounts that treat quantum probabilities as derivable from incomplete information rather than as fundamental are compatible with and supported by the framework.

The framework does not claim to settle these foundational questions; it makes structural commitments that constrain which accounts are compatible. Sections 18.9 and 18.10 develop two of these constraints in detail — what the framework says about cosmic spatial topology and what it says about consciousness — but the broader pattern is that the framework’s commitment to deterministic substratum is a substantive structural feature that constrains the foundational landscape.

The framework’s quantum randomness is therefore epistemic, exact in its predictions for any embedded observer’s empirical access, and consequential for the classes of foundational accounts compatible with it. The empirical record is silent at the 10^{-32} level; the conceptual record is substantial.

18.5 The quantum gravity dissolution as cumulative case

The framework’s content on the quantum gravity problem extends beyond the individual dissolutions developed in Chapter 7 to a structural pattern. Three problems conventionally treated as obstacles to a theory of quantum gravity — the cosmological constant problem, the black hole information paradox, and the non-renormalizability of perturbative quantum gravity — receive the *same* structural resolution in the framework, with each dissolution following from the same underlying commitment to the geometry-first ordering and the substratum-emergent operator distinction. This section pulls the three resolutions together as a cumulative case for the framework’s structural treatment of quantum gravity.

The conventional formulation. Standard treatments of quantum gravity attempt to quantize the gravitational field by analogy with the quantization of the

Standard Model forces. The metric tensor $g_{\mu\nu}$ is treated as a quantum field, decomposed into a classical background plus fluctuations, and the fluctuations are quantized through canonical or path-integral methods. Three obstacles arise.

The cosmological constant problem: the QFT vacuum energy summed over all field modes up to the Planck cutoff is of order $M_{\text{Pl}}^4 \sim 10^{113} \text{ J/m}^3$. The observed vacuum energy density is $\sim 10^{-9} \text{ J/m}^3$. The ratio 10^{122} is treated as a fine-tuning problem requiring an unknown cancellation mechanism.

The black hole information paradox: Hawking’s 1975 semiclassical calculation produces thermal radiation that appears to carry no information about the black hole’s initial state. If the radiation is exactly thermal, information is destroyed when the black hole completes evaporation — in tension with quantum mechanics’ commitment to unitarity.

Non-renormalizability: perturbative quantum gravity treated as a quantum field theory in flat spacetime is non-renormalizable. Infinitely many independent counterterms are required to absorb divergences at successively higher orders, with no UV completion within the standard quantization framework.

The framework’s content is that all three problems are *artifacts of the quantum-first ordering*. Each dissolves when the geometry-first ordering is adopted: spacetime geometry is part of the classical substratum, and quantum mechanics is the emergent description of an embedded observer’s reduced view of that substratum.

The first dissolution: the cosmological constant. Chapter 7 §7.7 developed the cosmological constant dissolution in detail. The framework has two distinct descriptions: the classical substratum, which assigns a vacuum energy $\rho_{\text{crit}} \sim H^2/G \sim 10^{-9} \text{ J/m}^3$ through the Friedmann equation, and the emergent quantum field theory, which assigns a vacuum energy $\rho_{\text{QFT}} \sim M_{\text{Pl}}^4 \sim 10^{113} \text{ J/m}^3$ through zero-point summation. The two descriptions operate at logically different levels: gravity couples to the classical substratum, the emergent QFT exists only after the trace-out. The 10^{122} ratio is the de Sitter entropy $S_{\text{dS}} = A/(4l_p^2)$ — the framework’s structural quantity for the information compression ratio between substratum and emergent description. Far from being a fine-tuning failure, the 10^{122} ratio is a derivable quantity.

The structural feature of this dissolution is that the two energy scales are not commensurable as gravitational sources. The conventional formulation treats them as comparable and asks why one cancels the other; the framework treats them as belonging to different levels of description, with no cancellation required.

The second dissolution: the black hole information paradox. Chapter 7 §7.6 developed the framework’s resolution of the information paradox. Information is never lost from the substratum bijection φ because φ is bijective. The apparent information loss in Hawking radiation is an artifact of the trace-out

— the same mechanism that produced the cosmological constant dissolution. The framework’s nested-partition formalism (two coexisting causal boundaries: cosmological horizon plus black hole horizon) provides a self-consistent emergent quantum description that preserves unitarity at the substratum level while reproducing the Page curve at the emergent level.

The structural Page time $t_P \approx 0.646 t_{\text{evap}}$, independent of mass and greybody factors, matches the quantum-extremal-surface and entanglement-island predictions through a substantially different mechanism — cycle typicality on the energy shell rather than replica wormhole computation in the gravitational path integral.

The shared structural feature with the cosmological constant dissolution is that the framework’s two-level description is doing the work. The information that appears lost in Hawking radiation is information that has crossed the trace-out boundary — accessible at the substratum level (bijectivity) but inaccessible at the emergent level (the radiation looks thermal because the emergent description marginalizes over the relevant hidden-sector degrees of freedom).

The third dissolution: non-renormalizability. Perturbative quantum gravity treats the lattice spacing $\epsilon = 2l_p$ as a regulator to be removed in the continuum limit. Renormalization-group analysis then shows that the theory requires infinitely many independent counterterms, with no UV completion possible within the standard quantization framework.

The framework’s content is that the lattice *is* the fundamental description, not a regulator. The substratum’s discreteness scale $\epsilon = 2l_p$ is a structural feature of the substratum dynamics, fixed by self-consistency in Chapter 7 §7.4 rather than chosen as a calculational convenience. There is no continuum limit to take because the continuum is not the underlying description — it is the long-wavelength limit of the discrete substratum dynamics, valid wherever the lattice spacing is much smaller than the physical scale of interest.

The non-renormalizability of perturbative quantum gravity is therefore an artifact of attempting to take a continuum limit where none exists. The substratum’s lattice is physical; the renormalization-group flow runs *to* the Planck scale rather than *through* it, with the framework’s lattice physics taking over at scales approaching ϵ . The structural commitment to the lattice’s physicality eliminates the non-renormalizability problem in the same way that the two-level description eliminated the cosmological constant and information paradox problems.

The shared structural feature. Three problems that conventional treatments find independently obstructive receive the same framework resolution. The shared feature is the geometry-first ordering: spacetime is part of the substratum, quantum mechanics is emergent, and the substratum-emergent operator distinction (developed in Chapter 1) is doing the work in each case. The cosmological constant ratio is a substratum quantity (the de Sitter entropy) being misidentified as an emergent quantity (the vacuum energy). The infor-

mation paradox treats an emergent process (Hawking radiation) as if it should preserve substratum information (which it does, through the nested partition). The non-renormalizability treats a substratum-level discreteness (the lattice) as an emergent-level regulator (a continuum-limit artifact).

The cumulative case is not three independent dissolutions but one dissolution applied three times. The framework’s quantum gravity content is that there is no quantum gravity problem to solve — what exists instead is a structural framework (S, φ) from which both quantum mechanics and gravity emerge as different aspects of the same object, with the apparent conflicts of conventional quantum gravity being category errors. The empirical content of the framework’s resolution is in the predictions it makes that neither QM nor GR alone makes: the Bekenstein-Hawking $1/4$ coefficient, the dark sector concordance, the value of $\hbar$ from partition geometry, the running-vacuum dark-energy form, the MOND acceleration scale, and the gravitational wave echoes developed in the next section.

18.6 Gravitational wave echoes

The framework’s commitment to the discreteness scale $\epsilon = 2l_p$ produces a forward empirical signature: gravitational wave echoes from black hole mergers. The conventional picture of a black hole’s ringdown phase treats the horizon as a smooth absorbing boundary; the gravitational waves propagating outward decay exponentially as the perturbed horizon settles to its equilibrium state. The framework’s discrete substratum produces a different prediction: the horizon’s discrete structure at scale ϵ partially reflects outgoing gravitational radiation back inward, with the reflected radiation propagating to the angular-momentum barrier outside the horizon and then back outward as a delayed echo.

The framework’s prediction. The echo delay time is set by the round-trip propagation between the discrete substratum scale near the horizon and the angular-momentum barrier outside it. The structural form is

$$\Delta t = \frac{r_+}{c} \ln \left(\frac{r_+}{2l_p} \right),$$

where r_+ is the horizon radius and the logarithmic factor arises from the integration of the radial null geodesic across the barrier potential. The logarithm reflects the near-horizon redshift: outgoing radiation experiences increasing delay as it climbs out of the gravitational well near the horizon, with the delay scaling logarithmically with the discreteness scale.

For a stellar-mass black hole with horizon radius $r_+ \sim 10$ km, the delay is approximately $\Delta t \sim 10$ milliseconds. For an intermediate-mass black hole with $r_+ \sim 10^3$ km, the delay extends to approximately 1 second. For a supermassive black hole with $r_+ \sim 10^9$ km, the delay reaches hours. The framework’s echo prediction is therefore observable for the entire range of black hole masses accessible to current and near-future gravitational wave observatories — though the

practical observability depends on the echo amplitude relative to the dominant ringdown signal.

The amplitude structure. The framework’s discreteness scale is at $\epsilon = 2l_p \sim 3.2 \times 10^{-35}$ meters — much smaller than any astrophysical length scale. The fraction of outgoing gravitational radiation that the framework predicts will reflect back from the discrete horizon structure scales with ϵ/r_+ — astronomically small for any realistic black hole. The expected echo amplitude is therefore many orders of magnitude below the dominant ringdown signal, putting the framework’s specific echo prediction below current detection thresholds.

This is the framework’s distinctive prediction relative to alternative echo proposals. Alternative quantum-gravity programs — firewall scenarios, fuzzball proposals, certain wormhole models — predict echo amplitudes that depend on the specific microstructure assumed, with some proposals giving echoes within reach of current LIGO sensitivity. The framework’s prediction places the echo amplitude at the *structurally minimal* level consistent with the substratum’s discreteness, which is below current thresholds.

Empirical status and falsification. The empirical situation in late 2025 and early 2026 has evolved substantially with the GW250114 observation discussed in Chapter 7 §7.5. Echo searches in LIGO/Virgo data have been performed by Abedi, Dykaar, and Afshordi (2017), Westerweck et al. (2018), Tsang et al. (2018), and most recently by the LIGO/Virgo/KAGRA Collaboration on the merger catalog through O3. The empirical status as of GW250114 is that no statistically significant echo signal has been detected, with upper limits on echo amplitude at the level of approximately 1% of the ringdown amplitude.

The framework’s prediction is *consistent* with this empirical state. The framework’s echo amplitude prediction is many orders of magnitude below current upper limits, so non-detection is the expected outcome at current sensitivity. The framework’s prediction would be falsified by either a *high-amplitude* echo detection at a timescale matching the framework’s $\Delta t = (r_+/c) \ln(r_+/2l_p)$ — which would suggest microstructure beyond the framework’s discreteness scale — or by a *low-amplitude* echo detection at a timescale inconsistent with the logarithmic formula. Neither has been reported.

The empirical situation is therefore that the framework’s specific echo prediction is consistent with current observations but is not directly testable at current precision. Future gravitational wave observatories — Cosmic Explorer, Einstein Telescope, and the LISA space-based interferometer for supermassive black hole mergers — may reach the sensitivity required to test the framework’s structurally minimal echo amplitude.

The conceptual content. The framework’s echo prediction is not the strongest empirical content in the gravitational sector; the Milky Way Keplerian decline at ~ 17 kpc (Chapter 7 §7.9) is quantitatively much more informative (the GW250114 result constrains the classical area theorem, not

the $1/4$ coefficient). The echo prediction’s value is *structural*: it makes the framework’s commitment to substratum discreteness empirically meaningful at gravitational scales. The prediction is testable in principle; the precision required for a definitive test is at the next generation of gravitational wave observatories rather than at LIGO/Virgo/KAGRA.

The chapter’s broader content on forward predictions continues with the framework’s substratum commitments — finite information, recurrence, the arrow of time, cosmic topology, and consciousness — that test the framework at scales and levels distinct from the direct empirical exposure of the gravitational and quantum sectors.

18.7 Finite information and Poincaré recurrence

The framework’s commitment to a finite substratum has structural consequences that extend beyond the empirical predictions of fundamental physics. The universe is described by a finite configuration space S with bijection dynamics $\varphi : S \rightarrow S$. Finiteness is not a calculational convenience but a load-bearing commitment: it guarantees Poincaré recurrence of the substratum dynamics, which produces P-indivisibility at the embedded-observer level, which produces quantum mechanics. The framework’s logic chain requires finite S .

Finite information content. The substratum’s information content is $\log_2 |S|$ bits — finite for any finite S . This is distinct from the Bekenstein-Hawking entropy $S_{\text{BH}} = A/(4l_p^2)$, which bounds the *accessible* information at any given partition (the visible sector’s capacity). The total information content of S — including the hidden sector at the cosmological boundary — exceeds the Bekenstein-Hawking bound but remains finite.

A clarification is essential. The cardinality $|\mathcal{C}_D|$ of the deep hidden sector is *gauge* in the framework: different values of $|\mathcal{C}_D|$ produce identical predictions for all observables accessible to the embedded observer. The “total information” of the universe is therefore not a physical observable in the framework. But the *finiteness* of $|\mathcal{C}_D|$ — that there exists some finite value, even if its specific magnitude is not observable — is load-bearing. Infinite hidden sectors would break the recurrence argument that drives P-indivisibility; the framework would not produce quantum mechanics on an infinite substratum.

The framework’s commitment to finite information is therefore neither a quantitative claim about the specific cardinality nor a metaphysical preference for finitism. It is a structural requirement: the framework’s derivation of quantum mechanics from substratum dynamics passes through the Poincaré recurrence theorem, which applies only to finite (or, more precisely, finite-measure) dynamical systems. The substratum’s finiteness is what makes the framework’s account of quantum mechanics possible.

Poincaré recurrence. A bijection φ acting on a finite set S is a permutation. Every permutation on a finite set has finite order: there exists a smallest positive

integer N such that $\varphi^N = \text{id}$, the identity map. After N time steps, the substratum returns to its exact initial state. The framework's substratum dynamics is therefore strictly periodic, with period N determined by the bijection's cycle structure.

The recurrence period is astronomically large. For a generic permutation on $|S|$ elements, the order is bounded by $\text{lcm}(1, 2, \dots, |S|)$ (Landau's function), which grows as approximately $e^{|S|}$ for large $|S|$. Applied to the framework's substratum with $|S| \sim e^{10^{122}}$ — the rough estimate of the substratum's cardinality consistent with the cosmological-horizon entropy of 10^{122} — the recurrence period is

$$N \sim e^{|S|} \sim e^{e^{10^{122}}}.$$

This is a definite quantitative prediction: the universe returns to its exact initial state after $e^{e^{10^{122}}}$ substratum time steps. Translating to physical units using the substratum time step $\tau_{\text{sub}} \sim t_{\text{Pl}} \sim 10^{-43}$ seconds gives a recurrence time vastly longer than the age of the universe, vastly longer than the lifetime of the proton, vastly longer than any timescale physics conventionally treats. The recurrence is therefore a structural prediction with no practical empirical consequences.

Falsifiability. The recurrence prediction is falsifiable only in the logical sense. If the universe demonstrably failed to recur — which would require an infinite observation time to verify — the framework would be falsified. No finite observation can rule out recurrence at the framework's predicted timescale. The recurrence is therefore a prediction at the structural level rather than at the empirical level: it follows necessarily from the framework's other commitments (finite substratum, bijection dynamics), and any framework consistent with those commitments must predict recurrence at approximately the framework's timescale.

The conceptual content is significant. The framework predicts that the universe is *cyclic at the substratum level* — periodic with period N , with the cycle including every configuration the substratum can be in. Within a single cycle, the observer's experience is unidirectional (the arrow of time discussed in the next section), but at the substratum level the dynamics has no preferred direction. The framework's commitment to substratum cyclicity is one of the structural commitments that distinguishes it from frameworks treating the universe as either fundamentally non-cyclic (most cosmological models) or fundamentally infinite (most quantum gravity programs).

Implications for the framework's logical structure. The recurrence theorem connects three structural features of the framework: substratum finiteness (the framework's commitment to $|S| < \infty$), bijection dynamics (the framework's commitment to φ as a bijection), and P-indivisibility (the framework's mechanism for emergent quantum mechanics). The connection runs in three steps. Finiteness of S guarantees recurrence (Poincaré). Recurrence guarantees that any partition of S into visible and hidden sectors will exhibit returns of information from hidden to visible over the recurrence timescale (P-indivisibility).

P-indivisibility produces the non-Markovian dynamics that, combined with the framework’s structural conditions C1-C3, yields emergent quantum mechanics.

The framework’s commitment to finite information is therefore not an isolated structural feature but a load-bearing premise in the framework’s derivation of quantum mechanics. Removing finiteness — replacing $|S| < \infty$ with $|S| = \infty$ — would break the logical chain and require a different account of how quantum mechanics emerges from substratum dynamics.

18.8 The arrow of time

The framework’s substratum dynamics is given by the bijection $\varphi : S \rightarrow S$. The bijection has *no arrow of time*. Both φ and its inverse φ^{-1} are equally valid maps on the configuration space, and the substratum dynamics is unchanged under the exchange $\varphi \leftrightarrow \varphi^{-1}$. The framework’s commitment to T-symmetric substratum dynamics is the same commitment that produced $\bar{\theta} = 0$ in Chapter 6 §6.3 — T-invariance at the substratum level is a structural feature that propagates through the framework’s emergent descriptions.

This raises an obvious question. The universe we observe has an unambiguous arrow of time: entropy increases, biological systems age, memories form of the past but not the future, the universe expands rather than contracts. If the substratum dynamics is T-symmetric, where does the observed arrow come from?

The arrow as observer-emergent. The framework’s answer is that the arrow of time is *emergent from the observer’s coarse-grained description*, not from the substratum dynamics. The substratum dynamics is T-symmetric; the observer’s reduced description is not. The asymmetry is introduced by the trace-out — the marginalization over the hidden sector that produces the embedded observer’s reduced description. The arrow of time is therefore a structural consequence of embedded observation, in the same way that quantum mechanics, the dark sector, and the Born rule are.

The mechanism is the standard Boltzmann argument applied to the framework’s partition. The visible sector starts in a particular configuration that is, in some sense, “low entropy” relative to the typical configurations accessible at later times. The hidden sector starts in a configuration that contains all the correlations needed to specify the visible sector’s future evolution. As the dynamics evolves, correlations flow from hidden sector to visible sector and back, with the time-reversibility of the underlying dynamics fully preserved at the substratum level. But the observer, who sees only the visible sector, sees an irreversible flow: information about the past is encoded in the visible sector’s current state (entropy increasing), while information about the future is encoded in the hidden sector (inaccessible to the observer).

The arrow of time is therefore *the observer’s incompleteness* expressed temporally. The same partition that produces quantum randomness (by hiding the

substratum information that determines outcomes) produces the arrow of time (by hiding the substratum information that would let the observer reverse the dynamics). Both are aspects of the same structural feature: the observer’s reduced description is asymmetric in ways that the substratum dynamics is not.

Detailed balance and the thermodynamic arrow. The framework’s content on detailed balance from Chapter 6 §6.3 has a direct implication for the arrow of time. Detailed balance asserts that the transition probabilities T_{ij} between any two states satisfy $T_{ij} = T_{ji}$ at all timescales — the rate of transition from i to j equals the rate of transition from j to i . This is a *substratum-level* T-symmetry: the dynamics has no preferred direction.

The thermodynamic arrow of time emerges not from any asymmetry in the transition probabilities but from the partition’s *initial conditions*. The universe starts in a configuration that is far from equilibrium relative to the substratum’s full state space, with most of the substratum’s degrees of freedom on the hidden side of the partition and the visible sector occupying a small, structured region. As the dynamics evolves, the system moves toward equilibrium relative to the substratum’s state space — entropy increasing in the observer’s coarse-grained description — until the entire substratum is fully mixed. The arrow of time runs in the direction of this equilibration.

The relation to detailed balance is that the H-theorem — the standard derivation of entropy increase from microscopic dynamics — is a theorem about the *coarse-grained* description, not about the underlying transition probabilities. The H-theorem assumes a particular kind of initial condition (low entropy at the coarse-grained level) and derives entropy increase from there. Detailed balance at the microscopic level is fully consistent with the H-theorem at the coarse-grained level, because the H-theorem is about how the observer’s information about the system evolves rather than about the system’s underlying dynamics.

The cosmological arrow. The thermodynamic arrow and the cosmological arrow — the direction in which the universe expands — are aligned in the framework. The cosmological partition has a specific initial condition: the universe begins in a low-entropy configuration with most degrees of freedom on the hidden side of the cosmological horizon. Time evolution corresponds to the horizon entropy increasing toward equilibrium, with the boundary modes accumulating correlations and the hidden sector approaching maximum disorder.

The alignment of the thermodynamic and cosmological arrows is structural rather than coincidental. Both arrows trace to the same initial conditions of the cosmological partition. The universe started with low entropy at the cosmological level; the framework’s structural conditions C1-C3 hold at the cosmological boundary (Chapter 7 §7.2); the thermodynamic arrow within the visible sector tracks the cosmological arrow because the cosmological boundary is the primary partition for any embedded observer.

Why the framework reframes rather than explains. The framework’s content on the arrow of time is not an explanation of *why* the universe has a

preferred direction; it is a *reframing* of the question. Standard accounts of the arrow of time treat it as an asymmetry that requires explanation — why is entropy increasing in the direction we call “future” rather than the direction we call “past”? The framework’s answer is that there is no asymmetry in the substratum dynamics; the asymmetry is in the observer’s reduced description, which inherits the initial conditions of the partition.

The reframing places the explanatory burden on a different question. Instead of asking “why does the universe have a preferred direction of time?”, the framework asks “why does the universe start in a low-entropy partition?” The latter question is a question about *initial conditions* — addressable by inflationary cosmology, anthropic arguments, or other initial-condition mechanisms — rather than a question about *dynamics*. This is the same kind of reframing the framework applies to the cosmological constant problem (Chapter 7 §7.7): the conventional formulation treats the small observed value as a fine-tuning problem in the dynamics; the framework treats it as a feature of the partition’s structure that does not require dynamical explanation.

The arrow of time is therefore one more case where the framework’s substratum-emergent distinction does substantive work. The substratum is T-symmetric (the dynamics has no arrow). The emergent description is not T-symmetric (the observer sees an arrow). The asymmetry is structural — inherited from the partition’s initial conditions — rather than dynamical. The framework’s content here is not to solve the arrow-of-time problem in the conventional sense but to locate it precisely: the arrow is at the level of partition initial conditions, not at the level of substratum dynamics.

18.9 Cosmic spatial topology

The framework’s structural commitments determine the *local* structure of the substratum but are silent on its *global* structure. The substratum is a cubic lattice at discreteness scale $\epsilon = 2l_p$ (Chapter 7 §7.4). The local cubic structure produces the gauge group, the matter content, and the framework’s quantitative predictions across Standard Model and gravitational sectors. But whether the lattice extends globally as $\mathbb{R}^3$, identifies as a 3-torus, has a Klein-bottle quotient, or admits some other global identification structure — that is not determined by the framework’s structural commitments. The framework is locally specific and globally agnostic.

This silence is principled. The framework’s content develops from the cosmological partition and the partition’s local properties at scale ϵ . Global topology operates at scales much larger than the cosmological horizon if at all (current empirical constraints place any non-trivial topology beyond the observable universe, as developed below). The framework’s predictions are therefore independent of global topology choices, and the framework’s commitments do not constrain them.

Janna Levin’s compact-universe proposal. Levin developed in the late 1990s and early 2000s the proposal that the universe could have a non-trivial spatial topology — specifically, that spatial sections could be *compact* (a 3-torus, a Klein-bottle quotient, or other compact identifications) rather than infinite. The motivation was partly cosmological (compact topology produces specific signatures in the CMB power spectrum at the largest angular scales) and partly philosophical (an infinite universe creates conceptual difficulties around uniqueness and the anthropic principle, addressed at book length in Levin’s *How the Universe Got Its Spots* in 2002).

The Klein-bottle specifically refers to non-orientable compact spatial topologies in which a parity-reversing identification connects opposite “sides” of the universe. The proposal is one example within the broader family of compact-universe possibilities — others include the 3-torus (orientable, no parity reversal), the Poincaré dodecahedral space (constant positive curvature, no parity reversal), and various lens-space identifications. Levin’s work, with collaborators Cornish, Spergel, Starkman, Weeks, and others, developed the empirical program for testing compact-universe proposals through their characteristic CMB signatures.

The empirical program: matched-circle signatures. The cleanest empirical test of compact-universe proposals is the *matched-circle* signature in the CMB. If the universe has compact topology with identification scale smaller than the diameter of the observable universe, then light rays that traverse the topological identification can return to the observer’s past light cone, producing correlated temperature patterns on circles in the CMB sky. The two circles on the CMB sky corresponding to the two endpoints of the same topological geodesic should show matched temperature patterns: the same fluctuation field probed at two different points on the celestial sphere.

The matched-circle test is geometrically clean — it depends only on the topological identification structure, not on details of the early-universe physics — and it has well-defined null and signal cases. Cornish, Spergel, Starkman, and Komatsu (2004) developed the statistical methodology for detecting matched circles in the CMB. WMAP data through 2010 and Planck data through 2015 have been searched for matched-circle signatures.

The empirical result is null: no statistically significant matched-circle signatures have been detected. The current bound places any compact-topology identification scale at greater than approximately the diameter of the observable universe — equivalently, any compact identification has scale $\gtrsim 28$ Gpc, beyond the cosmological horizon. The empirical situation is therefore that either the universe is topologically trivial at observable scales, or the universe has compact topology with identification scale larger than current cosmic horizon. The two possibilities are not currently distinguishable by data.

The framework’s compatibility with Levin’s proposal. Levin’s compact-universe proposal is *consistent with* but *not predicted by* the framework. Three

structural points clarify this.

First, the framework’s substratum finiteness (developed in §18.7) is *compatible with* compact spatial topology. A finite substratum could be realized as a cubic lattice with compact global identification at some scale beyond the cosmological horizon, where the total number of lattice cells is finite. Levin’s proposal provides one specific realization of this — the Klein-bottle quotient or 3-torus identification at some cosmological scale.

Second, the framework’s finiteness *does not require* compact spatial topology. The substratum’s finiteness could also be realized as a cubic lattice with trivial $\mathbb{R}^3$ topology but with a finite extent set by other structural features (the cosmological horizon’s expanding past light cone, for instance, defines a finite region of substratum at any moment). Levin’s compactness is sufficient but not necessary for the framework’s finiteness.

Third, the framework’s predictions in Chapters 5 through 17 are *independent of* global topology choices. Nothing in the framework’s empirical content — Standard Model gauge group, gravitational sector predictions, JUNO neutrino result, dark sector phenomenology, biological consequences — depends on whether the universe is topologically trivial $\mathbb{R}^3$, compactified as a 3-torus, or has the Klein-bottle quotient structure Levin proposes. The framework’s content is robust to global topology choices because the framework operates at the partition boundary (the cosmological horizon), not at the universe’s global topology.

The framework is therefore *neutral* on compact-universe proposals. Levin’s specific Klein-bottle proposal, the broader 3-torus family, and the trivial $\mathbb{R}^3$ alternative are all consistent with the framework’s substratum structure. The framework neither favors nor disfavors any particular global topology.

The CMB as joint testing ground. Both Levin’s matched-circle program and the framework’s Born-rule transient violations (§18.2) are CMB-data-driven empirical programs. They are structurally distinct — matched-circles tests global identification, Born-rule transient violations tests boundary thermodynamics at inflation — but they use overlapping observational data. The Planck 2018 data set has been analyzed for both signatures; CMB-S4 and LiteBIRD will improve constraints on both classes of predictions over the next decade.

The framework’s CMB content does not compete with Levin’s. The matched-circle test depends on the spatial extent and structure of topological identifications. The framework’s Born-rule transient violations depend on the temporal evolution of the visible-hidden coupling during inflation. The two tests probe orthogonal physical content using the same data sets — analogous to how the JUNO experiment probes neutrino mixing and the same reactor antineutrino flux probes other oscillation parameters.

Open question: non-orientability and chirality. The framework’s local-lattice chirality from Chapter 5 — SU(2) as chiral and SU(3) as vector-like, derived from the staggered trace-out — is a structural feature at the substra-

tum’s local cubic level. A separate question is whether a non-orientable global topology (Klein-bottle specifically, or other non-orientable compact identifications) interacts non-trivially with the framework’s local-lattice chirality.

The intuition is that non-orientability at the global level might produce phase factors or sign factors that interact with the local chirality structure, possibly producing observable consequences. The intuition could be wrong — the local-lattice chirality emerges at scales of $\epsilon = 2l_p$, while the global topology operates at scales beyond the cosmological horizon, with the relevant length scales differing by approximately 60 orders of magnitude. The two might be entirely decoupled.

This is an open question the framework has not addressed. Resolving it would require a calculation: starting from a specific non-orientable compact topology (Klein-bottle or otherwise), determining whether the framework’s local chirality derivation acquires global corrections, and computing the observable signatures of any such corrections. The framework does not currently have the calculation; the question is pre-registered as an open theoretical task.

The framework’s content on cosmic spatial topology is therefore: silent at the framework level (no commitment beyond local cubic structure), consistent with Levin’s compact-universe proposals (as one realization of the framework’s finiteness), with an open question on non-orientability-chirality interaction that requires further calculation. The framework’s structural unity does not extend to global topology, and the framework is honest about this scope.

18.10 The framework’s position on consciousness

The framework’s mathematical object — the partition with C1-C3 structure — defines an *observer* in a structural sense. The observer is not yet a *conscious entity*. Whether the structural conditions for observation are also sufficient conditions for consciousness is a separate question, with the framework’s commitments having implications for what answers are possible.

This section develops the framework’s substantive position on consciousness: observation is *necessary but not sufficient* for consciousness. The framework provides structural necessary conditions for conscious systems and imposes constraints on what physics-based consciousness theories are compatible with its commitments. The hard problem of consciousness proper — the explanatory gap between physical/structural facts and phenomenal experience — is not addressed, and the framework is explicit about this scope limitation.

Mathematical observer is not equivalent to conscious entity. The framework’s observer is a partition V of the substratum’s configuration space combined with the structural conditions C1 (coupling to hidden sector), C2 (slow-bath timescale separation), and C3 (sufficient hidden-sector capacity). The partition is a mathematical object that has dynamical consequences (the trace-out producing emergent QFT), thermodynamic consequences (entropy displacement producing dark matter phenomenology), and empirical consequences

(the JUNO prediction, the BH entropy, the dark sector).

The partition is not, by itself, a conscious entity. A thermostat coupled to room temperature implements a partition in a degenerate sense (the thermostat’s internal state interacts with the room’s thermodynamic state). A cosmological horizon implements a partition with $\tau_S/\tau_B \sim 10^{-32}$. A protein interior implements a partition. None of these is conscious in any phenomenal sense — there is nothing it is like to be the cosmological horizon. The framework’s mathematical observer is therefore a structural object, not a phenomenal one, and the framework does not claim equivalence between mathematical observation and conscious experience.

The asymmetric position. The framework’s substantive content on consciousness is the asymmetric claim: any conscious system must implement a partition with C1-C3 structure, but not every C1-C3 partition is conscious. The framework provides structural necessary conditions for consciousness without specifying sufficient conditions.

The claim has two directions worth distinguishing.

Direction 1: Consciousness requires observation. Any conscious system must implement a partition with substantial C1-C3 structure. A system that fails C1 (no coupling to a hidden sector) cannot be conscious because it has no environment to be aware of. A system that fails C2 (hidden sector equilibrates faster than visible-sector dynamics) cannot be conscious because it has no memory structure to support integration over time. A system that fails C3 (insufficient hidden-sector capacity) cannot be conscious because it cannot store the information that conscious experience would require. The C1-C3 architecture is necessary for the non-Markovian information processing that conscious systems must exhibit.

Direction 2: Observation does not require consciousness. The cosmological horizon, proteins, and engineered C1-C3 systems implement observation structurally without being conscious. The framework’s observation has structural consequences (quantum mechanics, dark sector phenomenology) at the level of the partition’s emergent dynamics, but it does not entail phenomenal experience. Whatever distinguishes the conscious C1-C3 systems from the non-conscious ones is some *additional feature* beyond the bare C1-C3 architecture — biological substrate, neural architecture, integrated information Φ , global workspace dynamics, recurrent processing, or some combination — that the framework’s structural commitments do not specify.

The asymmetric position is more conservative than alternatives like dual-aspect monism (where structural observation and phenomenal consciousness are necessarily co-occurring) and panpsychism (where every C1-C3 system is conscious to some degree). The framework’s content allows but does not require these stronger positions.

Substantive constraints on consciousness theories. The framework’s com-

mitments constrain what physics-based theories of consciousness are compatible with it.

Deterministic substratum rules out fundamental-indeterminism accounts. The framework’s substratum is deterministic: the bijection φ maps each state to a unique successor. Quantum randomness is epistemic, not ontic (§18.4). Theories of consciousness that *require* fundamental indeterminism in physical reality are incompatible with the framework. The most prominent example is Penrose’s orchestrated objective reduction (Penrose-Hameroff), which posits that conscious experience arises through gravitationally-induced quantum collapse in microtubules — a specific form of fundamental indeterminism that the framework’s deterministic substratum forbids. The framework’s commitments make Penrose-Hameroff structurally incompatible.

C1-C3 architecture is necessary for non-Markovian information processing. Conscious systems must exhibit memory effects, history-dependent dynamics, and information backflow between subsystems — the structural signatures of non-Markovian dynamics that emerge from the framework’s C1-C3 architecture. Theories of consciousness that treat conscious systems as fundamentally Markovian (instantaneous, memoryless processing) are incompatible with the framework’s account of how cognitively non-trivial information processing must work. Strict feed-forward computational accounts of consciousness face this constraint.

The substratum is finite. The framework’s finite substratum (§18.7) constrains theories of consciousness that posit infinite information content for conscious experience. Theories requiring genuinely infinite degrees of freedom (some forms of mathematical Platonism applied to consciousness, for instance) face this constraint.

These are real constraints rather than vacuous ones. The framework rules out a substantial class of physics-based consciousness theories while remaining compatible with the broader landscape of accounts that supply additional sufficient conditions for consciousness beyond the framework’s necessary ones.

Integrated information theory as adjacent content. Integrated information theory (IIT), developed primarily by Giulio Tononi, is the most prominent physics-adjacent consciousness theory and has substantive content overlap with the framework. IIT proposes that consciousness corresponds to integrated information Φ , a quantitative measure of how much a system’s parts inform each other beyond what their individual contributions predict. Systems with high Φ are conscious; systems with low Φ are not.

There is a real structural relationship between the framework’s C1-C3 architecture and IIT’s Φ . C1 (non-trivial coupling) is necessary for non-zero Φ — without inter-part coupling, integrated information cannot be non-trivial. C2 (slow-bath timescale separation) is necessary for IIT’s information integration over time — the framework’s $\tau_B \gg \tau_S$ provides the memory architecture that integration requires. C3 (large hidden-sector capacity) is necessary for the system’s state space to be large enough to support meaningful integration. A

system satisfying C1-C3 with substantial margins (large τ_B/τ_S , large N_H/N_V) is a system with substantial Φ ; a system satisfying C1-C3 marginally has low Φ .

The relationship is not exact identity. C1-C3 are *necessary conditions* for non-Markovian dynamics in the framework’s specific sense, while Φ is a *quantitative information-theoretic measure* in IIT’s specific formulation. But the conditions overlap substantially. The framework’s content here is compatible with IIT as a candidate sufficient-condition theory: the C1-C3 architecture provides the necessary structural conditions for the framework, and IIT’s Φ provides the additional sufficient content for consciousness within that architecture. The framework does not endorse IIT’s identification claim (that Φ *is* consciousness), but it provides structural foundations that IIT-like accounts can build on.

The relationship runs in both directions. The framework’s structural account of subjecthood — what makes a system a subject of observation in the framework’s mathematical sense — is compatible with IIT’s account of what makes a system conscious. Conversely, IIT’s predictions about which systems are conscious provide quantitative content that complements the framework’s structural necessary conditions. The two accounts engage at different levels and are not in tension.

The hard problem is not addressed. The framework’s content on consciousness is necessary-condition content. The framework provides structural conditions that conscious systems must satisfy; it provides constraints on what theories of consciousness are compatible with the framework’s commitments; it provides empirical predictions about the structural features conscious systems must exhibit. None of this addresses the hard problem of consciousness proper — the question of *why* phenomenal experience accompanies physical processes at all.

The hard problem is a question about the relationship between structural/physical facts and phenomenal facts. The framework provides a structural account (the partition with C1-C3 structure produces the emergent dynamics that conscious systems must satisfy). What it does *not* do is bridge the explanatory gap from such a structural account to phenomenal experience. The framework predicts which systems should be conscious as a necessary-condition claim, but it does not explain *why* implementing those necessary conditions should produce subjective experience rather than functional information processing “in the dark.”

This scope limitation is honest rather than evasive. No physics-based framework currently bridges the explanatory gap; claiming that the framework does so would be overreaching. The framework’s value here is its discipline: necessary-condition content with substantive constraints on theories, combined with explicit acknowledgment that the hard problem is not addressed. This is more than the strict scope often adopted in physics frameworks (which decline to engage consciousness at all) and less than the overclaiming sometimes seen (where

physics frameworks announce solutions to the hard problem they cannot actually deliver).

Empirical predictions and constraints. The framework’s necessary-condition account makes empirical predictions about the structural features conscious systems should exhibit.

Conscious systems should exhibit non-Markovian information processing — memory effects, history-dependent dynamics, information backflow between subsystems. The C1-C3 architecture is empirically detectable in neural systems; the prediction is that systems we identify as conscious (humans, mammals, possibly other complex biological systems) should exhibit these signatures at their relevant timescales. Neuroscience evidence — synaptic plasticity, neuromodulator gradients, glial dynamics, slow oscillations supporting cognitive integration — is consistent with this prediction. The framework’s content is structurally aligned with the broader neuroscience consensus on the architecture of conscious information processing.

Conscious systems should *not* be possible in architectures that strictly fail C1-C3. The framework predicts that strictly feed-forward neural network architectures (without recurrent connections, without memory states, without integration over time) cannot produce consciousness — they fail C2 by lacking the temporal integration architecture. This is consistent with the empirical finding that consciousness in humans is associated with recurrent processing rather than purely feed-forward processing. The framework’s content here is a structural constraint that aligns with empirical neuroscience.

The framework’s bioinformatics content from Chapter 17 is structurally related. The molecular clock overdispersion, rate autocorrelation across the tree of life, and LTEE power-law fitness trajectory are all empirical signatures of non-Markovianity in evolutionary biology systems. The framework predicts that similar non-Markovian signatures should be detectable in conscious systems at their relevant timescales — the brain’s analog of the molecular clock’s overdispersion would be the variability of neural processing rates across timescales, predicted by the framework to be non-Markovian rather than Markovian. The empirical program connecting these structural predictions to consciousness research is open.

The framework’s content on consciousness is therefore disciplined necessary-condition content with substantive empirical predictions, real constraints on theory choices, and explicit honesty about the limits of its scope. The hard problem of consciousness proper is not solved or dissolved by the framework. The framework’s value in consciousness research is structural — what makes a system the kind of system that *could* be conscious — rather than phenomenal.

18.11 Chapter close

This chapter has developed the framework’s distinctive forward content beyond standard quantum mechanics and general relativity. The empirical record of Chapters 5 through 17 establishes the framework’s content where it agrees with current physics at high precision; this chapter has developed the framework’s content where it makes predictions current physics does not yet test.

Three categories of forward content. The chapter’s content sorts into three categories distinguished by their empirical accessibility.

Testable forward predictions. The Born rule’s dynamical equilibrium with CMB non-Gaussianity signatures (§18.2) is testable through CMB-S4 and LiteBIRD over the next decade. The partially-quantum regime (§18.3) is testable through single-molecule kinetics experiments accessible to current laboratory techniques. The BQP computational ceiling (Chapter 14) is falsifiable by any demonstrated super-BQP computation. Gravitational wave echoes at $\Delta t = (r_+/c) \ln(r_+/2l_p)$ (§18.6) are testable through next-generation observatories (Cosmic Explorer, Einstein Telescope, LISA), though the framework’s predicted amplitude is below current LIGO/Virgo/KAGRA sensitivity. The framework’s content in this category is empirically engaged at near-term experimental scales.

Structural commitments with limited empirical access. The epistemic character of quantum randomness (§18.4) is consequential for foundations but has no observable consequences at the 10^{-32} suppression level. The quantum gravity dissolution as cumulative case (§18.5) reframes three problems but has no direct empirical signature beyond the predictions already developed in the cosmological constant, BH information paradox, and renormalization treatments. Finite information and Poincaré recurrence (§18.7) are falsifiable only in the logical sense — verification would require infinite observation time. The arrow of time (§18.8) is observationally manifest but is not directly testable as a framework prediction because all physics-based accounts of the arrow generate the same empirical signature. The framework’s content in this category is structurally consequential without producing additional empirical handles beyond those already developed in Chapters 5-17.

Framework’s scope limits. Cosmic spatial topology (§18.9) is not predicted by the framework; the framework is consistent with Levin’s compact-universe proposal and the trivial $\mathbb{R}^3$ alternative equally. The hard problem of consciousness (§18.10) is explicitly outside the framework’s scope; the framework provides necessary-condition content on consciousness but does not bridge the explanatory gap from structural to phenomenal facts. The framework’s content in this category is honest acknowledgment of what the framework does not address.

Forward pointers. Chapter 19 develops the framework’s content against the standard list of open problems in fundamental physics. The chapter is a systematic inventory: twelve problems the framework resolves or dissolves (quantum gravity in three components, the cosmological constant, dark matter, dark en-

ergy, the measurement problem, BH information, strong CP, hierarchy, gauge group origin, generation puzzle, dark sector budget, Born rule origin) and five scope limits where the framework does not provide a derived value (the flavor problem with structured open task; baryogenesis, which §19.3.3 reclassifies as solution-specific — the baryon asymmetry is downstream of the framework’s flavor-sector CP input; inflation and initial conditions; the Hubble tension; the cosmological initial state). Each entry cross-references where in the book the relevant content is developed, with Chapter 19 functioning as a reference resource for readers wanting to understand how the framework positions itself relative to the recognized hard problems of fundamental physics.

Cumulative case. The framework’s content across the empirical chapters (5-17) and this forward-predictions chapter (18) forms a cumulative case for the framework’s structural commitments. Twenty-two parameter-free Standard Model predictions match observation. Thirteen gravitational-sector predictions match observation. Seven neutrino-sector predictions match observation. The framework’s content in biology (Chapters 11-13 and 16-17) makes specific predictions that align with the empirical record across non-Markovian dynamics, evolutionary signatures, and clinical pharmacology. The forward content of this chapter extends the framework’s reach to phenomena standard physics has no specific commitments on — Born-rule equilibrium, partially-quantum regimes, computational ceilings, gravitational wave echoes, the arrow of time, consciousness as structural necessary condition.

The framework’s evidential standard is parameter-free structural prediction matching observation. The cumulative case across the book is retrodictive — exact structural values matching existing data across physics, biology, and computation — and the pre-registered standard properly applies to the forward content of this chapter, which awaits its data. The framework’s content beyond conventional physics is not speculative addition to an empirical base but structural extension of the same framework that produced the empirical predictions of Chapters 5-17.

The chapter is honest about what it does not address. The hard problem of consciousness is not solved. The cosmological initial state is not explained. Several other open problems remain in the framework’s residue. Chapter 19 develops these residues in detail, alongside the resolved-problem inventory that makes the framework’s cumulative empirical reach explicit.

Chapter 19

Resolved and Open Problems in Fundamental Physics

19.1 What this chapter develops

The framework’s content across the empirical chapters (5 through 17) and the forward-predictions chapter (18) constitutes a substantial body of structural commitments and quantitative predictions. This chapter inventories the framework’s content against the *standard list of open problems* in fundamental physics — the canonical questions every framework in physics is expected to address.

The inventory has two parts. Section 19.2 develops eleven problems the framework resolves or dissolves — the quantum gravity problem, the cosmological constant discrepancy, the nature of dark matter and dark energy, the measurement problem, the black hole information paradox, the hierarchy problem, the origin of the gauge group, the generation puzzle, the dark sector budget, and the origin of the Born rule — together with the strong CP problem, where the construction supplies a mechanism that does not reach the Standard Model. Each entry pairs the standard status of the problem (how it is conventionally formulated) with the framework’s status (how the framework engages it), with a cross-reference to the chapter and section where the relevant content is developed.

Section 19.3 develops six problems the framework does *not* resolve. The fermion mass hierarchy (flavor problem) has a defined research program with initial computational results and three tractable open targets, but the full quantitative content has not yet been derived. The Hubble tension is *potentially* addressable through ν_{OI} but at the framework’s predicted magnitude the effect is too small to explain the observed tension. Baryogenesis is solution-specific: the framework provides standard leptogenesis ingredients but the baryon asymmetry is downstream of the flavor-sector CP phases the framework inputs rather than derives, so η_B inherits that input status (§19.3.3). Inflation and initial conditions are *not addressed*: the framework reproduces the Standard Model structure but does not address the very early universe’s dynamics, with three possibilities outlined but none developed. The initial state of (S, φ) is partially gauge (within the reconstruction theorem’s equivalence class) and partially the framework’s irreducible mystery (why this equivalence class rather than another). The sixth problem is the chirality hypothesis H- χ (§19.3.8): whether the SU(2) embedding is taste-chirality-selective — the named hypothesis on which the Standard-Model chirality identification is conditional (SM §4.8), with the grading identification itself proved and probe-certified. Section 19.3 also takes up two framework-specific meta-questions: whether the discrete structural residue

is genuinely non-tautological (§19.3.6) — now *resolved* by a three-tier classification of the framework’s results, which is the section’s main conceptual addition and reframes the framework’s overall claim as a tight compression of physics rather than a derivation of it from first principles — and the radiative stability of emergent Lorentz invariance in the rotational sector (§19.3.7), where a symmetry argument and a first numerical test support protection, with specific questions remaining.

Section 19.4 closes the chapter with a summary table and a note on what the framework does *not address*: specific experimental anomalies (muon $g - 2$, B -physics, W mass) at scales where the framework’s $\mathcal{O}(\epsilon^2 E^2)$ corrections are unobservable, condensed matter and many-body phenomena that are problems within the derived QFT rather than at the framework level, and the framework’s principled scope limit on the hard problem of consciousness (developed in Chapter 18 §18.10).

The chapter functions as a systematic reference resource. Readers wanting to understand how the framework engages a specific open problem can find the framework’s status and the cross-reference to where the relevant content is developed. The chapter is not a substitute for the detailed treatments in Chapters 5 through 18; it is an inventory making the framework’s reach explicit.

A note on tone. The inventory is honest about the framework’s actual content. Where the framework resolves a problem, the inventory states the resolution and provides the cross-reference; where the framework only partially addresses a problem, the inventory marks the partial status and identifies what remains open; where the framework does not address a problem at all, the inventory acknowledges this rather than overstating the framework’s reach. The closing summary table makes the distribution of statuses explicit.

19.2 Problems resolved

This section develops the eleven problems the framework resolves or dissolves, and the strong CP problem alongside them, with each entry pairing the standard formulation against the framework’s status and a cross-reference.

19.2.1 The quantum gravity problem

Standard status. The central problem of theoretical physics. Quantum mechanics and general relativity are individually precise and mutually inconsistent — quantizing gravity leads to non-renormalizability, and no complete quantum theory of gravity exists. Substantial programs (string theory, loop quantum gravity, causal dynamical triangulations, asymptotic safety, and others) have been developed over decades without producing a complete theory.

Framework status. Dissolved. Quantum mechanics and general relativity are different projections of the same structure (S, φ) : quantum mechanics from the trace-out (Chapter 1), general relativity from Jacobson’s thermodynamic route

applied to the emergent area-entropy relation (Chapter 7). The two cannot conflict because they describe different levels of the same object. The lattice at $\epsilon = 2l_p$ is the fundamental description; there is no continuum limit to take and no renormalizability problem to solve. The framework produces both QM and GR from the same axioms with seven independent consistency checks and no free parameters.

The dissolution has three components developed as a cumulative case in Chapter 18 §18.5: the cosmological constant dissolution, the black hole information paradox resolution, and the non-renormalizability dissolution all follow from the same structural commitment to the geometry-first ordering and the substratum-emergent operator distinction. The quantum gravity problem dissolves not as one problem solved but as one problem applied to three apparent obstacles, each of which is an artifact of the quantum-first ordering rather than a genuine physical difficulty.

Developed in: Chapters 1, 7, and 18. The three components of the dissolution are developed in Chapter 7 §7.7 (cosmological constant), Chapter 7 §7.6 (black hole information), and Chapter 18 §18.5 (non-renormalizability and cumulative case).

19.2.2 The cosmological constant problem

Standard status. The worst quantitative discrepancy in physics. Quantum field theory predicts a vacuum energy of order $M_{\text{Pl}}^4 \sim 10^{113} \text{ J/m}^3$, summing zero-point energies $\frac{1}{2}\hbar\omega$ over all field modes up to the Planck cutoff. The observed vacuum energy density is $\rho_{\text{obs}} \sim H^2/G \sim 10^{-9} \text{ J/m}^3$. The ratio 10^{122} has been treated as a fine-tuning problem requiring an unknown cancellation mechanism — a cancellation that must be re-imposed at each order of perturbation theory.

Framework status. Dissolved. The 10^{122} discrepancy is the de Sitter entropy $S_{\text{dS}} = A_{\text{horizon}}/(4l_p^2)$ — the information compression ratio between the substratum and the emergent QFT description, not a physical energy. The QFT zero-point energy is an artifact of the emergent description and does not gravitate. The observed vacuum energy $\rho \sim H^2/G$ is the classical baseline set by the Friedmann equation at the substratum level. The two energy scales are not commensurable as gravitational sources; the conventional formulation treats them as comparable and asks why one cancels the other, while the framework treats them as belonging to different levels of description with no cancellation required.

Developed in: Chapter 7 §7.7.

19.2.3 The nature of dark matter

Standard status. Approximately 27% of the universe’s energy budget is in a form that clusters gravitationally but does not emit or absorb light. Decades of

direct-detection experiments — XENON, LUX, SuperCDMS, PandaX, and others — have produced no detection at any expected mass scale. The conventional cold-dark-matter paradigm postulates a new fundamental particle species; the framework’s content is that no such particle exists.

Framework status. Derived. Dark matter is not a particle — it is entropy displacement from the cosmological horizon. The Clausius relation applied to the de Sitter thermal reservoir gives the MOND critical acceleration $a_0 = cH/6 \approx 1.2 \times 10^{-10} \text{ m/s}^2$, the baryonic Tully-Fisher relation $v^4 = GM_B \cdot cH/6$, and the radial acceleration relation — all parameter-free. The dimensional factor $1/6$ arises structurally from the cubic-lattice geometry combined with the volume-vs-area entropy interplay in $d = 4$. The simple MOND interpolation function $\nu(y) = (1 + \sqrt{1 + 4/y})/2$ matches deep-MOND clusters such as Coma to better than 1%, though rich-cluster cores (Perseus, A2029, A1689) retain a $1.0\text{--}1.5\times$ residual that is an open problem: it is not closable by undetected baryons (massive clusters are baryon-complete near the virial radius, and the late-time missing baryons sit in cosmic filaments between clusters, not inside them) nor by the framework’s predicted light neutrinos ($\Sigma m_\nu \approx 0.059 \text{ eV}$), and most plausibly reflects the unresolved MOND-like-to-CDM-like crossover in the entropy-displacement relaxation kernel — the single highest-leverage open computation in the dark sector. The Bullet Cluster’s anomalous gravitational lensing is explained by frozen boundary entropy with relaxation timescale $\tau_{\text{relax}} \sim H^{-1} \gg t_{\text{collision}}$. The CMB acoustic peaks are reproduced by thermodynamic averaging. High-redshift evolution $a_0(z) = cH(z)/6$ is confirmed by the Milky Way Keplerian decline (Jiao et al. 2023) and the Genzel et al. high-redshift declining velocities.

Developed in: Chapter 7 §7.9.

19.2.4 The nature of dark energy

Standard status. Approximately 68% of the energy budget drives the universe’s accelerated expansion. Whether dark energy is a cosmological constant (constant in time) or evolving (the running-vacuum and quintessence proposals) is the subject of active investigation. The DESI Data Release 2 reports $2.7\text{--}4.2\sigma$ preference for evolving dark energy in the $w_0 w_a$ CDM parameterization, with the empirical situation in flux as of 2025-2026.

Framework status. Derived in form. The framework predicts that dark energy is the uniform component of the boundary entropy’s gravitational effect, with the effective cosmological constant following the Type II Running Vacuum Model:

$$\Lambda_{\text{eff}}(H) = \Lambda_0 + \frac{3\nu}{8\pi G}(H^2 - H_0^2) + \mathcal{O}((H^2 - H_0^2)^2).$$

The functional form is structural; the coefficient is bounded at the emergent-QFT scale $|\nu_{\text{OI}}| \sim 10^{-32}$, observationally indistinguishable from zero at any current or planned instrument. The Bertini *et al.* 2025 RG-corrected Λ CDM fit

to Planck 2018 + DESI DR2 + DES-Y5 reports $\nu = -(2.5 \pm 1.3) \times 10^{-4}$, consistent with the framework’s prediction at 2σ . DESI DR2’s own mild preference for phantom crossing is, against the structural prediction $w(z) \approx -1$, a $3\text{--}4\sigma$ tension — the framework’s one live adverse datum, carried on its own dated schedule ([GR §7.1]; Chapter 7 §7.8). DESI Year 5 (2028-2029) will provide a decisive three-way test between the framework’s $\nu \approx 0$, intermediate $|\nu| \sim 10^{-4}$, and high-confidence phantom crossing.

Developed in: Chapter 7 §7.8.

19.2.5 The measurement problem

Standard status. What constitutes a measurement? Why does wave-function collapse occur? Why does the classical world appear definite? No consensus exists across Copenhagen, Many-Worlds, Bohmian, QBist, dynamical collapse, and other interpretations.

Framework status. Dissolved. The wave function is not a physical object; it is the embedded observer’s bookkeeping device for tracking incomplete information about a deterministic substratum. Measurement is conditioning on a visible-sector outcome and re-marginalizing over the hidden sector. There is no collapse — only Bayesian updating on the classical substratum. The Born rule is derived as the equilibrium of the read-write cycle between visible and hidden sectors (Chapter 1, with forward content in Chapter 18 §18.2 on transient violations during inflation). Wigner’s friend has a definite outcome; Wigner’s superposition reflects his epistemic deficit about the friend’s measurement result. The framework’s resolution applies to all the variants of the measurement problem at once, because it dissolves the underlying assumption that the wave function is a physical object whose evolution requires explanation.

Developed in: Chapter 1, with related content in Chapter 18 §§18.2 and 18.4.

19.2.6 The black hole information paradox

Standard status. Does black hole evaporation destroy information? Hawking’s 1975 semiclassical calculation produces thermal radiation that appears to carry no information about the black hole’s initial state; if the radiation is exactly thermal, information is destroyed when the black hole completes evaporation — in tension with quantum mechanics’ commitment to unitarity. The AdS/CFT correspondence and the recent quantum-extremal-surface and entanglement-island programs suggest information is preserved, but the mechanism in the bulk has been unclear.

Framework status. Resolved. The substratum bijection φ is bijective; information is never lost. The Page curve is derived from cycle typicality on the energy shell, applying Popescu-Short-Winter measure concentration to the deterministic substratum dynamics. The structural Page time is $t_P = (1 - 2^{-3/2}) t_{\text{evap}} \approx 0.646 t_{\text{evap}}$, independent of mass and greybody fac-

tors. The framework’s derivation matches the quantum-extremal-surface and entanglement-island predictions through a substantially different mechanism — cycle typicality on the energy shell rather than replica wormhole computation in the gravitational path integral. The conjectured correspondence “replica wormholes $\leftrightarrow$ substratum cycles” is suggestive but not yet rigorous; making it precise would require constructing the gravitational path integral as an effective description of the substratum dynamics in the appropriate limit.

Developed in: Chapter 7 §7.6.

19.2.7 The strong CP problem

Standard status. Why is the QCD vacuum angle $\bar{\theta} < 10^{-10}$? Naively, $\bar{\theta}$ should be $\mathcal{O}(1)$ as a generic parameter of the QCD Lagrangian. The Peccei-Quinn mechanism (axion) is the leading conventional solution, but the axion has not been detected despite extensive searches.

Framework status. Not resolved. T-invariance of the substratum wave equation gives $\theta = 0$ at the substratum level. Detailed balance at all timescales — φ^n is a bijection for every n , hence $T_{ij}^{(n)} = T_{ji}^{(n)}$ — gives T-invariant Hamiltonians at every energy scale. T-invariance of the Hamiltonian at every scale gives real Yukawa matrices at every scale (no complex phases). Real Yukawa matrices give $\bar{\theta} = 0$ at every scale and every order of perturbation theory *within the construction*, with no axion. Four conditions stand between that and the strong CP problem of the Standard Model, and the first of them is not merely undischarged. The premise that delivers $\bar{\theta} = 0$ places the up-type and down-type Yukawas in a common real basis; two real matrices have real orthogonal singular value decompositions, so the CKM matrix is real orthogonal and the Jarlskog invariant vanishes identically. The construction cannot carry both $\bar{\theta} = 0$ and the observed CP violation, and breaking the equivariance far enough to restore J produces a rotational anisotropy some twelve orders of magnitude above experimental bounds. What the construction establishes is a structural mechanism for a small $\bar{\theta}$; what it does not yet have is a route from that mechanism to the physical Standard Model.

Developed in: Chapter 6 §6.3.

19.2.8 The hierarchy problem

Standard status. Why is the Higgs mass (~ 125 GeV) so much lighter than the Planck mass ($\sim 10^{19}$ GeV)? Radiative corrections in a generic effective field theory should drive m_H to the cutoff unless an unknown mechanism — supersymmetry, compositeness, extra dimensions, or anthropic selection — protects the Higgs mass from quadratic divergences. None of the conventional protection mechanisms has been confirmed empirically.

Framework status. Dissolved. The Higgs mass is an eigenvalue of the coupling matrix M on the cubic lattice substratum — a property of the specific bijection

φ , not a radiative correction within the emergent QFT. The lattice IS the UV completion; there is no hierarchy of scales to explain because the Higgs mass is set at the lattice level (where the framework’s structural commitments operate), not by a running coupling within the emergent QFT. Quadratic divergences are artifacts of treating the lattice spacing as a regulator to be removed — the same artifact that produces the non-renormalizability discussed in Chapter 18 §18.5. The framework’s substratum lattice eliminates the divergence and with it the hierarchy problem. What is dissolved here is the *naturalness* problem — the need for a mechanism protecting m_H from quadratic divergences; the *residual* electroweak hierarchy (why the specific value v/M_{Pl} takes the number it does) is a property of the bijection φ and remains an open input, as recorded under “what does not dissolve” in the project README.

The composite Higgs structure developed in Chapter 6 §6.6 provides the specific framework realization: the Higgs is a composite field of the framework’s specific substratum content, with the W , Z , and Higgs masses emerging as eigenvalues of the coupling matrix rather than as parameters in a renormalizable Lagrangian.

Developed in: Chapter 6 §6.6.

19.2.9 The origin of the gauge group

Standard status. Why $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$? Why these multiplicities for the matter content? Why this specific representation structure? The conventional account treats the Standard Model gauge group as empirically determined input rather than derived from any deeper principle.

Framework status. Derived. Coupling-degree minimization on the cubic lattice gives $K = 2d = 6$ for $d = 3$ spatial dimensions. The cubic group decomposition of the six nearest-neighbor link directions gives multiplicities $(3, 2, 1)$ under the action of the cubic group O , matching the structure of one Standard Model fermion generation (quark triplet, lepton doublet, singlet). Anomaly cancellation uniquely determines the hypercharges. Chirality emerges from the staggered trace-out. Three generations follow from the T_1 triplet representation of the cubic group. The Higgs quantum numbers $(1, 2, +1/2)$ follow from representation theory applied to the substratum’s link-carrier construction. The framework derives the full Standard Model gauge structure with no inputs beyond the cubic lattice and the framework’s structural axioms.

Developed in: Chapter 5 (gauge structure) and Chapter 6 (matter content and quantitative predictions).

19.2.10 The generation puzzle

Standard status. Why three generations? Why do all three have identical gauge quantum numbers but hierarchical masses spanning roughly five orders of magnitude? The Standard Model accommodates three generations as empirical

input but has no explanation for the multiplicity or for the mass hierarchy pattern.

Framework status. Partially derived. Three generations follow from staggered tastes on the cubic lattice. The T_1 representation of the cubic group has dimension three, and anomaly cancellation forces exactly three generations with the framework’s link-carrier construction. Identical gauge quantum numbers follow because all three generations are copies of the same T_1 irrep. Hierarchical masses follow because the specific bijection φ breaks the cubic symmetry through taste-breaking at $\mathcal{O}(\epsilon^2)$, with the hierarchy emerging from the second-order perturbative structure.

The *existence* of a hierarchy is structural — any bijection φ breaking the cubic symmetry produces a hierarchy. The *pattern* of the hierarchy — which generation is heaviest, the specific mass ratios, the CKM mixing angles — is solution-specific, requiring the specific properties of φ to derive. The framework predicts the existence of three generations with hierarchical masses unconditionally; the specific pattern is the open task developed in §19.3.1 (the flavor problem).

Developed in: Chapter 6 §§6.5 and 6.6, with the flavor open problem at §19.3.1 below.

19.2.11 The dark sector budget

Standard status. Why is the universe approximately 5% baryonic, 27% dark matter, and 68% dark energy? No conventional framework explains these ratios from first principles. The dark sector’s $\sim 95\%$ total is treated as empirical input.

Framework status. Derived (total budget). The trace-out renders approximately 95% of the critical density invisible to the emergent quantum field theory — matching the observed dark sector fraction. The horizon’s mode count $S_{\text{dS}} = A_{\text{horizon}}/(4l_p^2)$ combined with the classical horizon temperature gives, by the Clausius relation and the shell theorem, the total effective density $\rho_s = 3c^2 H^2/(8\pi G) = \rho_{\text{crit}}$. The uniform component is dark energy; the structured component (entropy displacement around baryonic matter) is dark matter. The 95% dark sector is concordance rather than puzzle — though the total fraction is partly definitional, equal to $1 - \Omega_B$ by construction: the framework’s structural quantity for the partition’s hidden-sector content at the cosmological horizon equals the observed dark sector fraction.

The specific 27/68 split between dark matter and dark energy depends on the matter distribution and the cosmic expansion history, and is not derived as a first-principles result. The framework derives the total budget; the partition into dark matter and dark energy components is a function of the universe’s matter content and dynamics rather than a structural prediction.

Developed in: Chapter 7 §7.9.

19.2.12 The origin of the Born rule

Standard status. The Born rule — that the probability of measuring quantum state $|x\rangle$ given wave function $|\psi\rangle$ is $|\langle x|\psi\rangle|^2$ — is an independent postulate of quantum mechanics with no derivation from the other axioms. Gleason’s theorem derives the Born rule from non-contextuality, but non-contextuality is itself a postulate. The various interpretive programs (many-worlds, QBism, Bohmian, dynamical collapse) take different positions on what the Born rule means without explaining why it has the specific form $|\langle x|\psi\rangle|^2$.

Framework status. Derived. The Born rule is the equilibrium distribution of the read-write cycle between visible and hidden sectors. The framework’s characterization theorem (Chapter 1) shows that the Born-rule probabilities are the unique stable statistics of P-indivisible dynamics — the fixed point of the completely positive trace-preserving map defined by the partition’s structural conditions C1, C2, C3. The Born rule is not an axiom but a structural consequence of embedded observation: any observer satisfying C1-C3 must use Born-rule statistics to describe the visible sector’s apparent randomness, because alternative distributions are unstable under the read-write cycle’s dynamics.

The framework predicts transient deviations from the Born rule in the very early universe (Chapter 18 §18.2), with the deviations decaying on the hidden sector’s equilibration timescale. CMB-S4 and LiteBIRD will tighten constraints on these transient violations through CMB non-Gaussianity searches over the next decade.

Developed in: Chapter 1, with forward content in Chapter 18 §18.2.

19.3 Problems open within the framework

This section develops the six problems the framework does not fully resolve. Each entry pairs the standard formulation, the framework’s current status, an assessment of what remains open, and a note on the prospects for closing the open question.

19.3.1 The flavor problem

Standard status. Why do fermion masses span five orders of magnitude, with $m_e/m_t \sim 3 \times 10^{-6}$? Why the specific pattern of CKM and PMNS mixing angles? Why the Koide relation $Q = (m_e + m_\mu + m_\tau)/(\sqrt{m_e} + \sqrt{m_\mu} + \sqrt{m_\tau})^2 = 2/3$ to better than 0.001%? The Standard Model accommodates the fermion masses as 13 free parameters; no conventional framework derives them.

Framework status. Open with a defined research program and initial computational results. The fermion masses are taste-breaking parameters of the substratum bijection φ — properties of how the specific φ breaks the cubic symmetry. The framework’s gauge couplings have been derived ($1/\alpha_0 = 23.25$ at the

Planck scale, all three gauge couplings at M_Z to better than 0.1%, developed in Chapter 6), demonstrating that solution-specific parameters are constrainable in principle. The research program identifies three tractable open targets.

The bottom-to-tau mass ratio. The framework predicts $y_b = y_\tau$ at tree level — the staggered taste-breaking is diagonal in the gauge index. Full Standard Model RGE running (verified against Luo and Xiao 2003) gives $m_b/m_\tau|_{\text{tree}} = 4.276$ from a common M_{Pl} boundary condition, indicating a substantial non-perturbative correction. The full prediction including the scalar-density renormalization Z_S from non-perturbative dynamics at the matching scale is developed in Chapter 6 §6.6 (and SM §7.5): with matching scale $m_{\text{match}} = \lambda g_0^2 = 0.1217$ derived structurally from the taste-decomposed Coleman-Weinberg potential, and $Z_S(m_{\text{match}}) = 1.813$ measured from Monte Carlo at $\beta = 11.1$ on $L = 32$, the prediction is $m_b/m_\tau = 4.28/Z_S(m_{\text{match}}) = 2.361$ against the observed 2.352 ± 0.017 — a 0.4% (0.5 σ) match at $L = 32$, with $L \rightarrow \infty$ extrapolation giving $\sim 0.7\sigma$. The $L = 16$ Monte Carlo using only the SU(3) mass renormalization $Z_m \approx 0.835$ at fixed bare mass — the calculation described in earlier versions of this section as giving ~ 3.4 (45% off) — corresponds to using the wrong matching prescription rather than the wrong physics. With the correct matching scale (a structural feature of the taste-decomposed CW potential rather than a tunable parameter), the prediction matches observation at the same level as the framework’s other Layer 1 closures (Cabibbo, Koide, $\sin^2 \theta_{13}$). The kinematic prefactor $K = 1/(n_{\text{gen}} - 1) = 1/2$ is a Layer-1 structural identity given the matching condition ($c_{bb} = (n_{\text{gen}} - 1)\lambda$); its support is the Monte Carlo verification at 1.4% ($K_{\text{MC}} = 0.507$). The explicit one-loop evaluation has been carried out and cannot close the gap: in the condensate-ratio construction the one-loop coefficient is 0.446 (tadpole-normalization bracket 0.393–0.446), 11–21% below 1/2 at a coupling where $Z_S - 1 = 0.81$ — a truncation error an order of magnitude larger than the MC precision, indicting the truncation rather than the identity. Analytical closure requires the resummed staggered- χ PT chain or a non-perturbative identity, or alternatively sub-percent MC; the task is Layer 2(a) in scope but not one-loop in method (SM §7.5).

The Koide relation. $Q = 2/3$ is equivalent to the square-root mass vector making a 45° angle with the $[1, 1, 1]$ direction in generation space — equivalently, the amplitude $A^2 = 2$ in the $\mathbb{Z}_3$ mass parametrization (SM §7.2). This has an exact structural target: $A^2 = C_2(T_1) = 2$, the same quadratic Casimir that fixes the Koide angle $\theta_0 = C_2/d^2$, so the amplitude and the angle descend from one invariant. The open question is therefore dynamical and sharper than “whether the 45° has a structural origin”: one-loop perturbative taste-breaking gives $A^2 \approx 0.34$, so the task is to show that the non-perturbative substrate dynamics produces $A^2 = C_2 = 2$ — a $\sim 6\times$ enhancement to a *known* target — via non-perturbative enhancement of the E -component (the structural route, under which the 45° is derived) rather than solution-specific T_2 off-diagonal elements (under which it would be bijection-specific). The full $K = 6$ lattice Monte Carlo is the tool, and whether it lands on C_2 is the key open question. The observable class is constrained in advance: an exact shift-symmetry identity makes the

ensemble single-quark two-point function eight-fold corner-degenerate at every coupling, so quark-level taste splitting vanishes identically and the measurement must use composite (multi-fermion) channels — closing the otherwise-natural quark-propagator route before it is attempted (SM §7.2).

The Wolfenstein hierarchy. Perturbative taste-breaking naturally produces $|V_{us}| \sim \lambda$, $|V_{cb}| \sim \lambda^2$, $|V_{ub}| \sim \lambda^3$ — the observed CKM pattern. The value $\lambda = 0.22508$ is the squared Cabibbo angle from cubic Brillouin geometry (Chapter 6 §6.5 and Chapter 8 §8.3) — a Tier 2 structural result. The hierarchy pattern is therefore structurally fixed; the specific magnitudes depend on the bijection.

Neutrino masses. Majorana, normal ordering, large PMNS mixing, and hierarchical spectrum are structural predictions (Chapter 6 §6.7 and Chapter 8). Specific mass values are solution-specific.

Required tools. Full $K = 6$ lattice Monte Carlo with the coupling matrix $M = \text{diag}(\mu_c I_3, \mu_w I_2, 1)$; representation theory of the cubic group O applied to the framework’s specific irrep content; 2-loop lattice-to-continuum matching. Four C programs implementing the perturbative taste-breaking spectrum, propagator-norm measurement, effective-mass extraction, and taste-split correlators are available as part of the framework’s computational infrastructure.

Assessment. The flavor problem is the framework’s largest single open theoretical task. The framework has substantive content on the *structural form* of the flavor pattern (the existence of three generations with hierarchical masses, the Cabibbo angle as Brillouin geometry, the Wolfenstein hierarchy, normal ordering for neutrinos) and on several specific magnitudes that are now Layer 1 closed (Cabibbo λ , Koide Q , $\sin^2 \theta_{13}$) or Layer 1 closed with Layer 2(a) verification residual (m_b/m_τ , $\sin^2 \theta_{12}$, $\sin^2 \theta_{23}$). The remaining solution-specific parameters (specific CP phases, the full Yukawa pattern beyond hierarchy existence) are bijection-specific in the framework’s terms — structurally predicted *not* to be derivable from substrate alone, consistent with the framework’s observational-incompleteness commitment. The Layer 2(a) verification residuals are tractable open tasks with defined computational paths; the bijection-specific parameters are pre-registered as outside the framework’s substrate-derivable scope.

19.3.2 The Hubble tension

Standard status. The Hubble constant measured from the CMB (Planck: $H_0 = 67.4 \pm 0.5$ km/s/Mpc) disagrees with the local distance ladder (SH0ES: $H_0 = 73.0 \pm 1.0$ km/s/Mpc) at greater than 5σ . The disagreement has persisted across multiple measurement systematics, suggesting either an unidentified systematic in one or both methods or new physics modifying the cosmic expansion history.

Framework status. Potentially addressable but at the framework’s predicted magnitude the effect is too small. The framework predicts evolving dark energy

in the Type II Running Vacuum Model with $|\nu_{\text{OI}}| \sim 10^{-32}$ bounded at the emergent-QFT scale, indistinguishable from zero. Quantitatively, in the RVM, the effective $H(z)$ differs from the Λ CDM prediction at low redshift:

$$H^2(z) = H_0^2 [\Omega_m(1+z)^3 + \Omega_\Lambda + \nu H^2].$$

The νH^2 term acts as a small additional dark energy component at late times, increasing the inferred H_0 relative to the CMB-inferred value. For $\nu \sim 2.5 \times 10^{-3}$ at the upper boundary of the framework’s substratum-contribution prediction, the shift would be approximately $\Delta H_0/H_0 \sim \nu/(1 - \Omega_m) \sim 0.4\%$ — far too small to explain the approximately 8% Hubble tension.

Assessment. The Running Vacuum Model with ν_{OI} does not resolve the Hubble tension. The tension may be due to systematic errors in the distance ladder (currently active investigation by SH0ES and CCHP collaborations), new physics not captured by the framework’s substratum dynamics, or a statistical coincidence becoming significant through accumulated precision. The framework makes no specific claim on the Hubble tension’s resolution and explicitly does not contribute to its understanding beyond noting that the framework’s predicted ν is too small to be the responsible mechanism.

19.3.3 Baryogenesis

Standard status. The observed baryon-to-photon ratio $\eta = (6.1 \pm 0.2) \times 10^{-10}$ requires a mechanism to produce a matter-antimatter asymmetry satisfying Sakharov’s three conditions: baryon number violation, C and CP violation, and out-of-equilibrium dynamics. The conventional candidates — electroweak baryogenesis, GUT baryogenesis, leptogenesis through heavy right-handed neutrinos, and various extensions — have not been confirmed empirically.

Framework status. Solution-specific. The framework does not derive the baryon-to-photon ratio and does not carry an obligation to — for the same structural reason it does not derive the CKM phase. This status is the conclusion of a sustained investigation; the reasoning is given below, because an earlier draft of this section proposed a mechanism that subsequent analysis retracted, and the honest record of that retraction is part of the framework’s content.

The framework provides standard leptogenesis ingredients, with a shortfall. The framework’s structural predictions provide ingredients for leptogenesis via the Weinberg operator $(c_{ij}/\Lambda)L_i L_j H H$: Majorana neutrinos (Chapter 6 §6.7), large PMNS phases, and hierarchical masses. The operator mediates $\Delta L = 2$ scattering in thermal equilibrium at $T > 10^{14}$ GeV. A precise calculation reveals a quantitative shortfall: the total unflavored CP asymmetry vanishes by CPT plus unitarity ($\sum_\alpha \epsilon_\alpha = 0$), and the surviving flavored asymmetry from thermal corrections is suppressed by $y_\tau^2 \sim 5 \times 10^{-5}$, yielding $\eta_B \sim 4 \times 10^{-11}$ — approximately $15\times$ below the observed value of 6.1×10^{-10} . This is the same shortfall faced by any framework with Majorana neutrinos but no heavy right-handed neutrinos: the Weinberg operator alone, without a UV

completion providing additional CP-violating phases, cannot produce sufficient baryon asymmetry.

A retracted mechanism. An earlier treatment proposed closing the factor-15 gap with *substratum-level instanton-like processes* — configurations of the bijection φ at the lattice cutoff with action $S \sim O(1)$, giving a suppression $e^{-O(1)}$ rather than the electroweak $e^{-2\pi/\alpha_w}$. This mechanism is **retracted**, and the reason is structural rather than quantitative. The substratum has no Euclidean action functional: it is the deterministic linear bijection of the wave-equation update rule (Chapter 2), not a path-integral system. “Substratum-level instantons with action $S \sim O(1)$ ” is therefore type-mismatched with what the substratum is — there is no action whose minima would constitute instantons. The correct picture is the one Chapter 6’s treatment of baryon number already gives: substratum-level B is conserved at the bilinear level, and non-perturbative $B+L$ violation enters through the *emergent* fermion-determinant topology governed by standard sphaleron physics. There is no additional substratum-level baryon-number-violating channel.

Why baryogenesis is solution-specific. With the retracted mechanism removed, the framework’s relation to baryogenesis is settled by a commitment the framework makes elsewhere and independently. The substratum is T-symmetric — this is load-bearing for the $\bar{\theta} = 0$ result of §19.2.7. Observed CP violation in the CKM and PMNS sectors is, in the framework, a flavor-basis feature of the *specific bijection* (S, φ) — not a derived quantity. The framework makes no prediction of δ_{CKM} , δ_{PMNS} , or the Wolfenstein (ρ, η) ; these are treated as solution-specific inputs, in the same category as the framework’s choice of discretization scheme.

The baryon asymmetry inherits this status. Sakharov’s second condition is CP violation, and the only CP violation available to the framework is exactly the flavor-sector CP it treats as a solution-specific input. A quantity downstream of an input is itself an input: η_B is no more derivable, within the framework’s current commitments, than δ_{CKM} is. Demanding that the framework dynamically *generate* η_B from a symmetric initial condition holds it to a standard it does not adopt for the CP phases η_B depends on. The “factor-15 open problem” framing was an artifact of that inconsistency.

Assessment. Baryogenesis is *solution-specific*, not a quantitative gap awaiting a mechanism. The framework’s structural ingredients (sphaleron $B+L$ violation, the out-of-equilibrium expanding early universe, the CKM/PMNS CP phases) are consistent with the observed η_B ; what the framework does not do is derive its value, because that value is downstream of CP phases the framework inputs rather than derives. This is a reclassification that restores internal consistency — the framework no longer carries baryogenesis as an unmet derivation obligation while simultaneously inputting δ_{CKM} — and it is honest about its limits: it is not a closure. It does not produce $\eta_B = 6.1 \times 10^{-10}$ from first principles, and a reader may fairly regard a framework that derives the Cabibbo angle but inputs the baryon asymmetry as incomplete in scope. The framework’s

position is that this incompleteness is shared, explicit, and located precisely: it is the flavor-sector CP input, named once and inherited wherever CP violation is needed.

19.3.4 Inflation and initial conditions

Standard status. The horizon problem (why is the CMB isotropic across causally disconnected regions?) and the flatness problem (why is $\Omega \approx 1$?) are solved by an early period of exponential expansion (inflation). But inflation requires a scalar field with a specific potential, which is not part of the Standard Model. The initial conditions for inflation are themselves fine-tuned. The empirical evidence for inflation comes through CMB polarization measurements and the absence of relics that pre-inflation cosmologies would produce.

Framework status. Not addressed. The framework reproduces the Standard Model structure but does not address the very early universe’s dynamics. Three possibilities exist within the framework’s structural commitments, none of which has been developed in detail.

The Higgs as inflaton. The framework derives exactly one scalar doublet (Chapter 6 §6.6). Higgs inflation (Bezrukov and Shaposhnikov, 2008) uses the Standard Model Higgs with a non-minimal coupling $\xi H^\dagger H R$ to drive inflation. If the framework’s Higgs *is* the inflaton, inflation is automatic — but the non-minimal coupling $\xi \sim 10^4$ is a solution-specific parameter (a property of the bijection φ), and its value is not determined by the framework’s structural commitments. The empirical predictions of Higgs inflation are well-developed in the literature; whether the framework’s bijection produces $\xi \sim 10^4$ is an open question requiring the full bijection-specific calculation.

No inflation needed. The framework’s lattice structure may resolve the horizon and flatness problems differently. On a finite lattice, the notion of “causally disconnected regions” is modified — the lattice’s coupling graph determines causal structure, and if the initial state of φ is statistically isotropic (one of the framework’s six structural properties of the bijection), isotropy is built in rather than requiring inflation to produce it. This is speculative and would require a detailed analysis of the lattice’s causal structure at early times.

Initial conditions as gauge. The deep hidden sector’s cardinality is gauge in the framework. If the initial state of the visible sector is also gauge-equivalent across a wide class of initial conditions (because the trace-out erases the distinction), then the “initial conditions problem” may dissolve rather than require solution. This is an intriguing possibility but has not been investigated within the framework’s existing structure.

Assessment. Inflation is not derived, not refuted, and not distinctively addressed by the framework. This is a genuine open problem with the possibilities outlined but no current developed framework position. Pre-registered: the framework’s content on the very early universe is one of the framework’s most

substantial scope limits in fundamental physics.

19.3.5 The initial state of (S, φ)

Standard status. Not applicable in the conventional sense — this is a framework-specific question. Why does the universe have the particular bijection φ rather than some other? Why these particular initial conditions for the bijection’s evolution?

Framework status. Partially addressed, partially outside scope. The framework’s reconstruction theorem establishes that observed physics uniquely determines the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ of bijections producing those observations. Within the equivalence class, all bijections produce identical observables; the choice of representative is gauge. The question “why this φ ?” has two layers.

Within the equivalence class. Gauge — no answer needed, no answer possible. The framework’s content makes the specific choice of representative bijection within the equivalence class empirically indistinguishable from any other choice in the class.

Why this equivalence class rather than another. The framework’s axioms — deterministic bijectivity, finiteness, causal partition, classical probability — constrain the class. The framework’s six structural properties of the bijection (developed across Chapters 5-7) further narrow it. But the axioms themselves are not derived from anything deeper.

The existence question — “why does (S, φ) exist at all?” — is the framework’s irreducible mystery. The framework’s structural account: (S, φ) is a mathematical structure that cannot fail to be well-defined; existence is the default for logically consistent structures. Whether one finds this answer satisfying is a philosophical question rather than a physical one. The framework’s commitment is that the existence of *some* finite deterministic bijection structure is not a question physics can answer, while the existence of *this particular* structure (as determined by the empirical record) is a constraint on which bijections are compatible with observation.

Assessment. The framework identifies this question as the only genuinely irreducible mystery — not a gap in the framework’s reach but the boundary of all possible frameworks. Any framework must take some structural axioms as given; the framework’s content is to make those axioms as minimal and as structural as possible (deterministic bijectivity on a finite set, no further commitments), leaving the existence question as an explicit acknowledgment of where physics ends.

19.3.6 Is the discrete structural residue genuinely non-tautological?

Standard status. A framework-specific meta-question. The reconstruction theorem sends the empirical and structural axioms to a discrete residue — the

spatial dimension $d = 3$, the internal multiplicity $K = 2d = 6$, the gauge group $SU(3) \times SU(2) \times U(1)$, three generations, the hypercharges, $\bar{\theta} = 0$. The question is whether these are forced and could have been otherwise (substantive content), or whether they are already contained in the inputs and the derivation unpacks them (tautological).

Framework status. Resolved — not by a calculation, but by recognizing that the question is mis-posed. It presupposes a single yes/no answer for “the residue,” but the residue is not uniform: its members have three different logical statuses, and once they are separated the question has a definite answer for each. *Tier 1 — forced by representation theory (genuinely non-tautological).* Given the simple-cubic lattice and the wave equation, the multiplicities $(3, 2, 1)$ and the gauge factors, the generation count of three, the anomaly-free hypercharges, $\bar{\theta} = 0$, and the Cabibbo angle follow from the octahedral character table and anomaly analysis with no further freedom; they could *a priori* have come out otherwise, and the predictive excess noted above (the Cabibbo angle, Koide, the PMNS angles, six mass ratios — none of them among the empirical inputs) is exactly the signature of genuine content. This tier is real and substantial. *Tier 2 — empirically selected.* Two of the listed invariants are *not* forced by the construction: $d = 3$ is selected by the three Standard-Model-independent filters (§3.2 of the SM paper), and the simple-cubic Bravais structure is selected because the observed gauge group contains $SU(3)$ — only the cubic crystal system supplies the three-dimensional point-group irrep an $SU(3)$ factor needs (Theorem 7b). For these the inference runs from the observed world back to the commitment; “the Standard Model from a cubic lattice” and “the lattice is cubic because we observe $SU(3)$ ” are one fact stated in two directions. The threshold δ_0 and the absolute mass scales are likewise empirical. *Tier 3 — definitional stipulation.* Coupling-degree minimization (the definition of a site, from which $K = 2d$ follows), the nearest-neighbor restriction, and the C1–C3 architecture itself are choices about what counts as physical structure rather than facts the bijection forces; they are principled, but they are stipulations.

Assessment. The meta-question is therefore not open pending computation — no lattice calculation moves it, because the status of each link is already fixed by the construction as written. It is answered by the classification: non-tautological at Tier 1, tautological-in-the-relevant-sense at Tier 2, definitional at Tier 3. Several distinct concerns — the derived-versus-input ledger, the residual-input status of Lorentz structure (§19.3.7 and the boost sector of the SM paper), and the circularity question — reduce to this one and are resolved the same way. The constructive consequence is a precise statement of what the framework is: a *maximally tight compression* of observed physics onto the cubic-lattice commitment. Tier 1 establishes that, once the commitment is granted, the Standard Model’s discrete structure is largely not independent — it is a single representation-theoretic object — which is a genuine and unusual claim about the unity of physics. Tiers 2–3 establish that the commitment itself is empirically anchored and partly stipulated, so the framework does not derive physics from a parameter-free first principle. The earlier framing of this as

“open, not closable by internal means” mistook a question about the *logical status of the construction’s steps* — which the construction already answers — for a question awaiting evidence. The proposed external calibration against the operational reconstruction theorems of quantum theory (Hardy; Masanes–Müller; Chiribella–D’Ariano–Perinotti) remains a worthwhile comparison, but it bears on how to *value* a “uniquely-fixed-modulo-gauge” result, not on whether the present residue is forced; the latter is settled above. (The three-tier classification is developed in full in [SM §8.3] and [Substratum §3, Refinement]; its bearing on the framework’s foundational claim is drawn out in [Structure §13.4, Limit 4].)

19.3.7 Radiative stability of emergent Lorentz invariance (rotational sector)

Two results bear on this question. A symmetric-integration argument shows that an exactly O_h -invariant action, measure, and observer-centered partition yield an O_h -invariant effective theory to all orders, and every ingredient of the checkerboard trace-out is O_h -invariant — no direction-dependent momentum shell is available. A numerical check on the rounded interacting dynamics (extraction validated against the analytic tree coefficients to the percent level; the allowed $\ell = 4$ anisotropy detected at high significance) finds the forbidden E_g coefficient zero within errors at the 2×10^{-5} level in the perturbative regime and consistent with zero in the thermalized strong-coupling state. Separately, the sector-differential boost violation’s custodial decoupling law now has its normalization bounded ($\Delta C_2 \in [0.58, 2.08]$ from the framework’s representation content), yielding, on the worst-case ΔC_2 , an M_X window of two to five orders of magnitude against the tightest multi-species bounds and requiring M_X to be an intermediate scale, $\lesssim 10^8$ GeV; with the computed kernel constants ($c_L = -9.4$, $d = -0.88$) and the representation-side point values, the tightest ceiling refines to $\sim 4 \times 10^7$ GeV and the window narrows to between one and a half and four and a half decades ([SM §3.1]). What remains: production-scale statistics, the $\ell = 4$ alignment of the observer-position boundary anisotropy, and the cross-charged phenomenology floor. See [SM §3.1] for the full statements.

Standard status. The Lorentz fine-tuning problem familiar from effective-field-theory treatments of Planck-scale Lorentz violation: tree-level suppression of Lorentz-violating operators need not survive radiative corrections.

Framework status. Open. The *sector-differential* (inter-species) boost-violation is addressed by the custodial- $U(6)$ mechanism, with a power-law-with-logarithm decoupling in the heavy cross-charged mass whose form (not normalization) is established (SM §3.1 (*Scope of the emergent-Lorentz claim*), F13). The remaining gap is the *sector-common* operator in the rotational sector, which has no analogous power-law suppression and depends on the renormalization-group-emergence (operator-smearing) mechanism, whose inter-species flow is only logarithmically fast. This rotational sub-problem does not reduce to §19.3.6 (it is a technical radiative-stability question rather than

a question about the residue’s content). Closing it requires establishing that the emergent theory preserves the octahedral symmetry at loop level: cubic symmetry and locality place the rotational violation at dimension six (SM §3.1, *Scope of the emergent-Lorentz claim* point (iii) and *A present-day constraint, and the rotational sector’s distinct status*), so the residual question is loop-level octahedral invariance rather than ultraviolet smearing.

A framing note on the gauge hierarchy. Several of the items above can be stated in terms of the substratum gauge group $\mathcal{G}_{\text{sub}}$ — the transformations preserving all observables of (S, φ) — of which the emergent Standard Model gauge group is the image under trace-out. On this description each physical quantity falls into one of three classes: $\mathcal{G}_{\text{sub}}$ -gauge and unobservable (quotiented away); observable but solution-specific (the absolute mass scale, the CP phases, the neutrino Yukawas — properties of the particular φ); and forced structural residue (the discrete quantities of §19.3.6). The claim that quantum mechanics, gravity, and the Standard Model are three projections of the one equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ is the framework’s statement of scope; the question of §19.3.6 is whether the third class has content independent of the inputs.

A structural observation: several open problems appear to share one premise. Three of the framework’s open technical questions, which present as independent entries on the problem list, appear on analysis to reduce to a single underlying question — *whether the substratum-to-observer coarse-graining map preserves the relevant low-energy structure (locality, isotropy, the spectral mode count)*. The three are: (a) the radiative stability of emergent Lorentz invariance in the rotational sector (§19.3.7 above), which turns on whether coarse-graining acts as a sufficient operator-smearing; (b) the magnitude of the running-vacuum coefficient ν ([GR §7.1]), which is fixed at $|\nu| \approx 10^{-32}$ (observationally null) if the boundary-mode spectral measure is the local 2D field-theoretic one, and would rise to the DESI-detectable $\sim 10^{-4}$ level only if coarse-graining generated a scale-invariant (non-local) boundary spectrum; and (c) the stability of the $K = 2d$ minimal-coupling count on general graphs ([SM §4.5], Theorem 6 extension), which holds for the coarse-grained dynamics provided local degree fluctuations are renormalization-irrelevant — again a statement that coarse-graining preserves the low-energy mode structure. If this common reduction is correct, the framework has fewer independent open questions than the problem list suggests, and “does the coarse-graining map preserve locality?” is the single most consequential open technical question — its resolution would discharge all three at once, and conversely a failure (emergent non-locality) would show up simultaneously as rotational Lorentz violation, a non-null ν , and a breakdown of the general-graph mode count. The shared premise has now been computed for the linear (reconstruction-selected) map: an exact surface-Green’s-function extraction of the coarse-graining map’s boundary spectral measure — the half-space trace-out of the substratum wave equation itself, not a Gaussian effective model — yields a power-suppressed local exponent ($dN/d\omega \propto \omega^4$, with the effective one-mode-per-cell 2D measure, $\propto \omega$, an upper bound on the infrared weight), invariant over thirty orders of magnitude in the slow-bath ratio and

under generic quenched deformation, with the log-uniform measure appearing nowhere. The three reductions are thereby conditionally discharged at that grade — locality preserved, ν on the null branch a fortiori, the $K = 2d$ mode-count premise stable — with one named residual: the computation is of the linear map, and its interacting (condensate-dressed) extension is the remaining step. The linkage retains its organizing value in reverse: a failure at the interacting level would still have to manifest simultaneously as rotational Lorentz violation, a non-null ν , and a mode-count breakdown — one falsifiable correlate, not three independent risks. A further preserved structure has since joined the same pattern: the even/odd sublattice grading underlying the staggered spin construction descends exactly through the linear trace-out — an exact Schur-complement identity valid for any site partition ([SM §4.8]) — with the same interacting extension as its residual, further concentrating the value of that single remaining computation.

Why these particular features are inputs: constitution versus consequence. A natural question is why the open and input items cluster as they do — why, for instance, the framework derives the gauge group and the generation count but must take the spatial dimension, the cubic lattice, and (per §19.3.7) rotational isotropy as inputs. The organizing principle is the distinction between the *constitution* of the substratum and its *consequences*. The framework derives the consequences of a given constitution — the gauge group, the three generations, the anomaly-free hypercharges, $\theta = 0$, and the Cabibbo angle all follow, by representation theory of the cubic group acting on the link directions together with the trace-out, *once the substratum is fixed*. It does not derive the constitution itself — the dimension, the Bravais type, the rotational isotropy, the periodicity. The reason is structural rather than incidental: an embedded observer is built *out of* the substratum, so it can have leverage on the consequences of its own constitution but not on that constitution prior to itself. The one qualification is that a consequence which the framework cannot yet *compute* is also booked as an input on tractability grounds rather than principle (the threshold δ_0 of §6.2 is the clearest case; it is a consequence of the coupling-running, input only because it is not yet derived). With that qualification the rule is clean: inputs comprise the substratum’s constitution together with any not-yet-tractable consequences, while the derived results are the tractable consequences. This is offered as an organizing observation, not a theorem; it is not claimed that derivability coincides with any single formal property, and indeed the largest block of derived results (the representation-theoretic ones) would follow for any analyst of the fixed lattice, with no reference to an observer at all.

A note on the boost/rotation asymmetry specifically. The split between the two Lorentz sub-sectors of §19.3.7 — the boost sector, for which the framework has structural tools (the custodial-U(6) mechanism, the local-boost-invariance prerequisite of the Jacobson construction, the vacuum-Cherenkov stability filter), versus the rotational sector, which has no analogous ally — tracks this same constitution/consequence distinction through a specific feature

of embedded observation. The trace-out is realized at a causal horizon, and that horizon *singles out a rest frame* (the frame in which the substratum update steps are simultaneous, identified in §7 with the cosmic rest frame). The observation map therefore makes contact with boost structure — it distinguishes a temporal direction, which is what boosts act on — and this is why the boost sector admits framework-level arguments. It does not single out any spatial direction; spatial isotropy is a property of the lattice constitution, untouched by the rest-frame selection. This is formalized as the rest-frame-selection lemma of [Main §3.5] (the observation map’s invariance group is the little group $O(3)$ of its selected timelike direction). On this reading the relative tractability of the boost sector and the input status of the rotational sector are not separate facts but two faces of the rest-frame-selecting character of embedded observation: the map touches what distinguishes time from space (boosts), and is silent on what distinguishes one spatial direction from another (rotations). This is an interpretive observation about why the open-problem boundary falls where it does, not an additional technical result.

The rotational sector as a structural question about observational incompleteness. The boost/rotation note above and the smearing question of §19.3.7 can be drawn together into a single statement about *why* the rotational sector is the one that resists protection — and the statement is a structural feature of embedded observation itself, not an incidental gap. The protection the rotational sector would need is a sufficient operator-smearing: a coarse-graining whose kernel damps ultraviolet spatial modes (so the radiatively-generated anisotropy stays suppressed). The framework’s coarse-graining is the trace-out, and the trace-out *is* the observational incompleteness — the observer’s ignorance of the hidden sector. The crucial point is *what kind* of incompleteness this is. By the fluctuation discriminant of [Main §3.5], the trace-out averages over the hidden *state* — over which microstate the slow bath occupies — not over spatial scale. A state-average is not a spatial ultraviolet filter: averaging over which microstate the bath is in does nothing to remove short-wavelength spatial modes, since those are present in every microstate alike. The framework’s incompleteness is moreover enormous — the slow-bath capacity ratio $\tau_B/\tau_S \sim 10^{32}$ is the engine of the entire quantum-emergence result — but it is incompleteness *about a temporally-evolving state*, concentrated in the state and (through the slow bath) temporal directions. It is silent in the spatial direction in the same way, and for the same reason, that the rest-frame lemma’s map is silent on spatial axes.

This can be made quantitative. The smearing kernel is the traced-out hidden-sector propagator G_{oo} ([SM §4.8]); its position-space width is the smearing radius R . A direct computation of that width on the cubic lattice gives $R \approx \varepsilon$ — a single lattice spacing — and, decisively, that width is *flat* as the slow-bath ratio τ_B/τ_S is varied across thirty orders of magnitude. The framework’s vast observational incompleteness contributes *nothing* to the spatial smearing width; the only spatial smearing present is the order-one multi-site support of the staggered reconstruction, which is far too narrow ($R/\varepsilon = \mathcal{O}(1)$, not $\gg 1$) to

damp the ultraviolet. This is the structural reason the smearing route does not add ultraviolet damping to the rotational sector — and it is why the available computations (static, free-dynamical, and one-loop) find $R/\varepsilon = \mathcal{O}(1)$ rather than $R \gg \varepsilon$: the incompleteness is the wrong *kind* to UV-damp, regardless of its magnitude. Whether the absence of smearing leaves the sector *unprotected*, however, depends on the dimension analysis below.

Three of the framework’s findings thus share one root. The rest-frame-selection lemma (the map selects a timelike direction, no spatial one), the fluctuation discriminant (the map averages the fluctuating state, not the fixed orientation), and the smearing-width result (the incompleteness has order-one spatial width independent of its magnitude) are three expressions of a single fact: *observational incompleteness in this framework is incompleteness about a temporally-evolving hidden state, with purchase on time, boost, and state, and structurally silent on space and rotation*. The rotational sector is precisely the sector embedded observation cannot reach. The valence of this account is narrower than it first appears. The smearing-width result it establishes — that observational incompleteness adds no ultraviolet spatial smearing, so $R \approx \varepsilon$ — bears on the *coefficient* of the rotational operator, not its *dimension*. The dimension is fixed independently, by the octahedral selection rule (the E_g anisotropy is a cubic non-singlet; the lowest invariant anisotropy is $\ell = 4$, dimension six) together with locality, and a direct tree-level computation confirms the emergent dispersion anisotropy is $\sim (\varepsilon k)^2 = (E/M_{\text{Pl}})^2$. The account therefore explains why embedded observation does not *smear* the rotational sector, but smearing is not what the sector needs: cubic symmetry and locality already place the violation at dimension six, Planck-suppressed. What remains genuinely open is the loop-level question — whether the emergent theory preserves the octahedral symmetry *exactly*, or whether the emergence introduces an octahedral-breaking effect not visible in the bare partition.

This premise has a second dependency alongside the partition. The effective theory is octahedrally invariant only if the condensate that breaks the framework’s custodial symmetry is itself aligned to the cubic axes; the invariance of the dynamics does not descend to the solution, since the gauge group arises precisely by that condensate breaking a symmetry the dynamics has. The alignment is not secured by the observed gauge group alone — a condensate with the same block sizes but misaligned eigenspaces yields the same gauge group while moving the forbidden anisotropy from dimension six to dimension four — but is secured by the same geometric identification of blocks with link directions that the matter-content and mixing-angle results use. An aligned condensate leaves the forbidden anisotropy exactly null at every block splitting, and a spatially modulated condensate does not lift it. See [SM §3.1]. This has an identifiable failure point: it assumes the visible/hidden partition acts as a state/sublattice trace that preserves the octahedral symmetry, rather than an effective short-wavelength (momentum-shell) cut that could both ultraviolet-damp and, if anisotropic, break it. The gaplessness of G_{oo} is evidence the partition is not a clean momentum shell; whether it is exactly octahedral-invariant at loop level

is the seam at which the dimension-six protection could fail, and is the arbiter — not the kernel width. This premise is grounded in the framework’s own factorization criterion ([Main §3.4]): an observer-centered, horizon-isotropic partition is spherical and hence octahedrally symmetric about the observer’s centre, so the lattice octahedral group — unlike translations, boosts, and rotations about distant points, which the partition breaks — preserves the visible/hidden factorization and descends to an exact symmetry of the effective theory. The same criterion that selects the gauge symmetries protects the rotational sector and exposes the boost sector, so the dimension-six rotational protection and the dimension-four boost exposure are two faces of one structural fact rather than independent assumptions.

19.3.8 The chirality hypothesis H_χ

The Standard-Model identification carries one named structural hypothesis. H_χ : the $SU(2)$ embedding is *taste-chirality-selective* — the gauge coupling distinguishes the taste-chirality grading on the internal E components. What is proved, and probe-certified (`chirality_grading_probes.py`): on the visible sector the spin-chirality and taste-chirality operators coincide, so a gauge coupling is spin-chiral there if and only if its embedding is taste-chirality-selective (SM §4.8, Theorem 13); the grading itself is $(d_L, d_R) = (4, 4)$ with $\Gamma(0) = I_4$ chirality-neutral, refuting the earlier “even sublattice = left-handed” identification; $SU(3)$ is vector-like; and the generation count does not depend on H_χ . What is open is H_χ itself: whether the substratum’s $SU(2)$ embedding is in fact taste-chirality-selective. The resolution path is a computation in the framework’s own representation machinery — the embedding’s action on the graded components of the coupling matrix (the Theorem 7 apparatus). A computed non-selective embedding falsifies the chirality identification, and with it the specific Standard-Model matter assignment, while leaving the gauge group, the multiplicities $(3, 2, 1)$, and the generation count intact; a computed selective embedding discharges the hypothesis and upgrades the chirality clause corpus-wide. The condition’s standing status is transcribed in the status ledger (Main §4.5).

19.4 Chapter close: summary table and what the framework does not address

The framework’s content against the standard list of open problems in fundamental physics summarizes as follows. The status labels are the framework’s claims as argued in the cited chapters, and every label inherits the conditionality ledger of [Main §4.5] for the foundational claims — in particular the C2 leg’s mixing-hypothesis conditionality — together with the matter-sector flags carried at their sources: the open chirality hypothesis H_χ [SM §4.8] and the flagged carrier input of the Casimir chain [SM §4.7]; no label asserts more than its cited derivation, and “Dissolved” or “Derived” is always relative to that ledger.

Summary table.

Problem	Standard status	Framework status	Developed in
Quantum gravity	Central problem	Dissolved (three components)	Chs 1, 7, 18
SU(2) chirality	Assumed in the SM	Open — named hypothesis H- χ ; grading proved, selectivity open	§19.3.8, [SM §4.8]
Taste $\rightarrow$ physical spin	Spin postulated per field	Proved in the staggered formalism (Thm 9); the carrier input to the Casimir chain is flagged open, tested by the preregistered A^2 measurement	[SM §4.7]
Cosmological constant	Worst prediction in physics	Dissolved ($10^{122} = S_{\text{ds}}$)	Ch 7 §7.7
Dark matter	Unknown, no particle found	Derived ($a_0 = cH/6$)	Ch 7 §7.9
Dark energy	Unknown, evolving disputed	Derived in form (Type II RVM)	Ch 7 §7.8
Measurement problem	No interpretive consensus	Dissolved (Bayesian updating)	Ch 1, Ch 18 §18.2
BH information paradox	Hawking vs unitarity	Resolved (Page curve from typicality)	Ch 7 §7.6
Strong CP	Why $\bar{\theta} < 10^{-10}$?	Not resolved (construction gives $\bar{\theta} = 0$; same premise removes CKM CP violation, and breaking it enough to restore J violates isotropy bounds by ~ 12 orders)	Ch 6 §6.3, [SM §5.5]
Hierarchy problem	Higgs mass protection unknown	Dissolved (lattice is UV)	Ch 6 §6.6
Gauge group origin	$SU(3) \times SU(2) \times U(1)$ unexplained	Derived ($K = 2d = 6$, cubic decomposition)	Chs 5, 6

Problem	Standard status	Framework status	Developed in
Generation puzzle	Three, with hierarchical masses	Partially derived (three is structural; pattern open)	Ch 6 §§6.5-6.6
Dark sector budget	95% unexplained	Derived (total budget)	Ch 7 §7.9
Born rule origin	Independent postulate	Derived (read-write equilibrium)	Ch 1, Ch 18 §18.2
Flavor problem	13 free parameters	Open (research program; m_b/m_τ at 0.4% match with Layer 2(a) verification residual)	§19.3.1
Hubble tension	$> 5\sigma$ disagreement	Open (ν_{OI} too small)	§19.3.2
Baryogenesis	Sakharov conditions	Solution-specific (η_B inherits flavor-CP input status)	§19.3.3
Inflation	Inflaton not derived	Open (three possibilities, none developed)	§19.3.4
Initial state	Framework-specific	Partially gauge, partially irreducible	§19.3.5

The empirical record summarized above places the framework in a substantive position relative to the standard hard problems of fundamental physics. Twelve problems are resolved or dissolved — including some that have been central concerns of theoretical physics for decades (quantum gravity, the cosmological constant problem, the measurement problem). Five problems remain open. The framework’s record is mixed but is honestly reported: the resolutions are real (each with a specific cross-reference to where the content is developed) and the open problems are real (each with a specific status assessment and an indication of what closure would require).

Problems the framework does not address.

Beyond the resolved and open problems above, the framework does not address several other categories of question.

Specific experimental anomalies in the Standard Model precision sector are not modified by the framework. The muon anomalous magnetic moment $g - 2$

shows a $\sim 5\sigma$ discrepancy between Fermilab measurement and the SM prediction using dispersive methods (though lattice QCD has reduced the tension). The B -physics lepton flavor universality violations (R_K , R_{K^*}) were reported by LHCb but recent measurements are consistent with the SM. The CDF II measurement of $m_W = 80.4335 \pm 0.0094$ GeV is approximately 7σ above the SM prediction, while other measurements (ATLAS, LHCb) are consistent. The framework derives the SM structure but does not modify the calculations producing these observables. Any framework correction at accessible energies is at $\mathcal{O}(\epsilon^2 E^2) \sim 10^{-64}$, unobservable. If any of these anomalies is confirmed at $> 5\sigma$ with consistent measurements, it would require either new physics beyond the framework's SM derivation or a solution-specific effect from the particular φ .

Condensed matter and many-body physics. The framework derives the fundamental laws but does not address emergent phenomena in condensed matter: high- T_c superconductivity, the fractional quantum Hall effect, topological phases, and others. These are many-body problems within the derived QFT, not problems at the framework level. The partially-quantum regime predicted by the framework (Chapter 18 §18.3) may be relevant to some condensed matter systems, but the connection is speculative and has not been developed.

The hard problem of consciousness. Developed in detail in Chapter 18 §18.10. The framework's mathematical operation of observation is a necessary structural condition for consciousness but not a sufficient condition. The hard problem proper — the explanatory gap between structural and phenomenal facts — is explicitly outside the framework's scope. The framework provides necessary-condition content with substantive constraints on physics-based consciousness theories, but it does not bridge the explanatory gap between physical descriptions and phenomenal experience.

Mathematics and logical foundations. The framework presupposes standard mathematics (ZFC set theory, the integers, finite combinatorics) without addressing the foundations of mathematics. Whether the framework's commitment to finiteness has implications for mathematical foundations (constructivism, finitism, ultrafinitism) is not addressed.

The framework's epistemic position. The cumulative record of resolved problems across Chapters 5 through 18 is substantial: twenty-two classified Standard Model retrodictions (seven parameter-free structural), thirteen gravitational-sector entries, seven neutrino-sector predictions, and the dissolutions of the quantum gravity problem in three components. The framework's evidential standard — preregistered structural predictions matching observation at the current empirical precision — is met across multiple sectors of physics with a uniformity that distinguishes it from frameworks tuned to fit specific empirical targets.

The open problems are also substantial. The flavor problem is the framework's largest single open theoretical task. The Hubble tension is not addressed by the framework's ν_{OI} prediction. Baryogenesis is solution-specific — the baryon

asymmetry inherits the input status of the framework’s flavor-sector CP phases. Inflation is not derived. The initial-state question is partially gauge and partially the framework’s irreducible mystery.

A structural observation about the resolution pattern. The eleven resolved problems above are not eleven independent solutions; they trace, to a substantial degree, to one structural feature — the substrate/emergent split with trace-out. The cosmological constant dissolution, the measurement problem dissolution, the Bell/Kochen-Specker/PBR/Frauchiger-Renner dissolutions, the dark-sector content, the hierarchy problem dissolution, the Higgs criticality, the absence of GUT signatures: each of these traces to the same architectural move. Fixing the framework’s substrate machinery to produce any one of these resolutions strongly constrains (and in several cases forces) the others. The resolutions are not assembled puzzle-by-puzzle; they fall out of one architectural commitment.

This is the kind of unity a foundational framework either has or does not have. A framework that solves N deep problems through N separate mechanisms, each fixable independently, is — in the methodological vocabulary of Deutsch’s *Beginning of Infinity* — easy to vary: each mechanism could be adjusted without disturbing the others, and the framework’s prior probability is correspondingly diffuse over alternative parameter choices. A framework that solves N deep problems through one structural feature with shared assumptions is hard to vary: changing the architecture to fix one resolution disturbs the others. The Bayesian consequence is that a hard-to-vary framework concentrates probability where an easy-to-vary one disperses it; correlated explanations are stronger explanations, not weaker ones, when they hold up.

The framework’s relationship to the standard list of open problems should therefore be characterized in two complementary ways. *Quantitatively*, the framework resolves eleven problems out of seventeen on the list (eleven resolved or dissolved; one not resolved; five open); the resolutions match observation across multiple sectors and the open problems are honestly pre-registered as the framework’s current boundaries. *Structurally*, the resolutions are not independent: one architectural feature does most of the work, and the framework’s joint explanation of eleven disparate hard problems through one shared machinery is the kind of evidence the methodological tradition since Popper has held up as the mark of genuinely explanatory theory. The first characterization is the empirical record; the second is the framework’s distinctive shape. Neither is a substitute for the other, but the second is the one most likely to be missed by a reader who scans the summary table alone.

The chapter’s purpose is to make the framework’s record on standard open problems explicit. The framework’s content is neither a complete theory of fundamental physics (the open problems above remain genuinely open) nor an unfocused collection of predictions; it is a structural framework that resolves a substantial fraction of the recognized hard problems of physics through a small set of structural commitments, with the resolutions cross-referenced to

specific empirical content developed in earlier chapters. The remaining open problems indicate where the framework’s reach stops at current development, with concrete prospects for closing several of them through further calculation within the framework’s existing structural commitments.

The book’s content concludes here. The framework’s empirical record (Chapters 5-17), forward predictions (Chapter 18), and inventory against the standard open problems (Chapter 19) together constitute the framework’s positive content. The remaining scope limits — the flavor problem, the Hubble tension, baryogenesis (solution-specific: η_B inherits the flavor-sector CP input), inflation, and the initial state — are pre-registered as the framework’s current boundaries. Closing them is the framework’s future research program.

Appendix A

Prediction Status Table

A.1 What this appendix provides

The framework has a substantial empirical record across multiple domains. Twenty-two Standard Model observables, thirteen gravitational-sector observables, seven neutrino-sector observables, and one structural biological retrodiction (the wider therapeutic window for Rett syndrome) cumulatively constitute the framework’s quantitative match against existing data. This appendix provides a single reference resource for what the framework predicts, how each prediction is classified by structural status, what the current empirical match is, and what open issues remain.

The appendix serves two purposes. First, as a *navigational index*: readers wanting to know “what does the framework predict for X” can find each prediction with its location in the book, its formula, its empirical match, and its current status. Second, as an *honest accounting*: predictions that are confirmed appear alongside predictions that are layered (requiring additional verification), parameter-fit (where the framework provides structural form but specific values are empirical), and forward-looking (testable but not yet tested). The reader can therefore assess the framework’s empirical record at the level of detail appropriate to their domain of interest.

Section structure.

Section A.2 develops the classification scheme — the categories used throughout the appendix to characterize each prediction’s structural status. The scheme

distinguishes structural (S), conditional (C), layered (L), retrodiction (R), parameter (P), mass-chain (M), and empirical (E) predictions, with each category having specific evidential force.

Section A.3 develops the four-layer descriptive framing — substratum (Layer 0), substratum-level structural (Layer 1), substratum-level mechanism-specific (Layer 2(a)), substratum-level cubic-symmetry-specific (Layer 2(b)), mixed-level (Layer 3) — that organizes how the framework’s content propagates from foundational commitments through specific empirical predictions.

Section A.4 provides the master prediction tables, organized by domain: Standard Model observables (A.4.1), gravitational-sector observables (A.4.2), neutrino-sector observables (A.4.3), and biology/medicine predictions (A.4.4). Each entry includes prediction number, formula or quantitative content, structural classification, layer, empirical match, and chapter cross-reference.

Section A.5 provides forward predictions and falsification targets — predictions the framework makes that have not yet been tested but that provide specific empirical exposure for future research programs.

Section A.6 closes with a discussion of the framework’s epistemic position based on the cumulative empirical record: which predictions provide the strongest evidence, which are most distinctive of the framework relative to standard accounts, and which open issues are most consequential for the framework’s overall content.

A.2 Classification scheme

The framework classifies each prediction by its structural status, using a seven-category scheme designed to make the *epistemic content* of each prediction explicit.

S — Structural prediction. The framework’s content fully determines the prediction’s content from substratum-level commitments. The prediction’s value emerges from the framework’s structural foundations without empirical input beyond what is required to identify which physical system the framework describes. Examples: $\sin^2 \theta_{13} = 4/(18\pi^2)$, the Cabibbo angle $\lambda = 1/(\pi\sqrt{2})$, the Bekenstein-Hawking coefficient $1/4$. S-class predictions provide the framework’s strongest evidence — they are *first-principles derivations* from structural commitments.

C — Conditional prediction. The framework’s content determines the prediction *conditional* on an additional unverified structural relation holding. If the relation is verified, the prediction promotes to S. If the relation fails, the prediction may need reclassification. The C-class can therefore be either provisional S (awaiting verification) or provisional L (awaiting demotion). The current §B.3 table has zero C-class entries — the Cond 2 prediction that previously occupied this class was reclassified as L following the §B.2.2 audit.

L — Layered prediction. The framework’s content includes the prediction’s *form* at the structural level but requires an empirical or computational input for the specific value. The empirical input is bounded — typically a single dimensionless coefficient or a normalization scale — but cannot be derived without additional work. Examples: $\sin^2 \theta_{12}$, $\sin^2 \theta_{23}$, m_b/m_τ . L-class predictions provide *structural understanding* of the prediction’s origin with empirical input for specific values.

R — Retrodiction. The framework reproduces an existing empirical value through a structural mechanism, but the empirical value was input to the framework’s identification of which specific system it describes. The prediction is retrospective rather than forward-looking. Examples: $\alpha_2(M_Z)$, $\alpha_3(M_Z)$. R-class predictions provide *structural consistency* with established physics but not novel predictive content.

P — Parameter prediction. The framework determines a parameter through a calculation that involves both structural commitments and one or more empirical inputs. Different from L (where the form is structural but the value requires input) and from R (where the prediction is purely retrodictive), P-class predictions combine structure and empirics in derivation.

M — Mass-chain prediction. The framework predicts mass relationships within a chain (e.g., m_e , m_μ via Koide from m_τ) where one mass is empirical input and the others follow structurally. Mass-chain predictions provide *structural relationships* between empirical inputs.

E — Empirical input. A value that the framework treats as empirical rather than deriving — typically a coupling constant, mixing angle, or other parameter that is part of the framework’s identification of which specific physical system applies. E-class entries are *not predictions* in the strong sense but parameters the framework uses without deriving.

The cumulative content. The framework’s strongest evidence comes from S-class predictions (first-principles structural derivations matching empirical values), supplemented by L-class predictions (structural form with bounded empirical input). The R, P, M, and E categories serve as honest accounting of where empirical input enters the framework’s content. The classification scheme makes the framework’s evidential content explicit at the prediction level, allowing the reader to assess which predictions provide the strongest support for the framework’s structural commitments.

A.3 Four-layer descriptive framing

The framework’s predictions are organized at four levels of descriptive abstraction, with each level having specific evidential content and specific dependencies on the framework’s structural commitments.

Layer 0 — Substratum. Predictions at the substratum level: properties of the underlying bijection φ on S , the substratum’s dimensionality, the sub-

stratum’s discreteness scale. Examples: dimensionality $d = 3$, finite cardinality, deterministic bijection structure. These predictions have *foundational* evidential weight — they are the framework’s substratum-level commitments themselves.

Layer 1 — Substratum-level structural. Predictions derived from the substratum’s structural properties through dimensional analysis, group theory, or other framework-level structural mechanisms. Examples: gauge group $SU(3) \times SU(2) \times U(1)$ from substratum dimensionality (Chapter 5), generation count $K = 2d = 6$ from matter content theorem (Chapter 6), Cabibbo angle $\lambda = 1/(\pi\sqrt{2})$ from cubic-group symmetry. Layer 1 predictions are *structural derivations* from foundational commitments.

Layer 2(a) — Substratum-level mechanism-specific. Predictions that require specific structural mechanism beyond Layer 1 generic structural content. Examples: m_b/m_τ requires the two-loop $S\chi PT$ mechanism specialized to OI’s 3D-slice cubic-group taste structure. Layer 2(a) predictions are *mechanism-specific* — they require more detailed structural content than Layer 1.

Layer 2(b) — Substratum-level cubic-symmetry-specific. Predictions that depend on the specific cubic-symmetric substratum the framework commits to. Examples: the Jarlskog $J = \eta/(12\pi^6)$ requires specific cubic-symmetric coefficient. Layer 2(b) predictions are *symmetry-specific* — they require commitment to cubic symmetry rather than just the foundational substratum structure.

Layer 3 — Mixed-level. Predictions that mix substratum-level structural content with emergent-level computational content. Examples: m_e, m_μ via Koide from m_τ requires both substratum-level Koide relation (Layer 1) and emergent-level mass chain. Layer 3 predictions are *mixed* — they combine substratum and emergent content.

Cumulative implications. The framework’s content is most evidentially strong at Layer 0 and Layer 1 (where substratum commitments directly determine the prediction’s form and content). It is less strong at Layer 2(a) and Layer 2(b) (where additional structural content is required). It is least strong at Layer 3 (where empirical-emergent content mixes with structural content). The classification scheme makes these distinctions explicit, allowing the reader to assess the framework’s evidential weight at the appropriate level.

A.4 Master prediction tables

This section provides the master tables of the framework’s predictions, organized by domain. Each entry includes prediction number, formula or quantitative content, structural classification, layer, empirical match, and chapter cross-reference for further detail.

A.4.1 Standard Model observables

The framework’s twenty-two Standard Model predictions form the core empirical record. All twenty-two predictions match observation within approximately 1%

or approximately 1σ — the cumulative empirical content is substantial.

#	Prediction	Formula	Class	Layer	Empirical match	Chapter
1	$1/\alpha_0$ universal threshold	23.25	S	0	Within $\sim 0.5\%$	Ch 5 §5.4
2	$A \cdot B$ gauge coupling product	48.0 ± 1.5	S	0	Within 1%	Ch 5 §5.5
3	δ_0 from $U(1)$ row	10.02	E	0	Empirical input	Ch 5 §5.5
4	$\alpha_2(M_Z), \alpha_3(M_Z)$	$129.57, 8.47$	R	0	Retrodictions	Ch 5 §5.5
5	Cabibbo λ	$1/(\pi\sqrt{2})$	S	1	0.04%	Ch 6 §6.4
6	Wolfenstein A	$\sqrt{2/3}$	S	1	Within 1%	Ch 6 §6.4
7	$ V_{cb} $	$A\lambda^2$	S	1	Conditional on 5,6	Ch 6 §6.4
8	m_d/m_s	λ^2	L	1	Structural λ ; GST empirical	Ch 6 §6.5
9	Jarlskog J	$\eta/(12\pi^6)$	E	2(b)	Empirical η	Ch 6 §6.4
10	Koide Q	2/3	S	1	0.02%	Ch 6 §6.5
11	Koide θ_0	$C_2/d^2 = 2/9$	S	1	Within 1%	Ch 6 §6.5
12	$\sin^2 \theta_{12}$	$1/3 - 1/(4\pi^2)$	L	1+2(b) ^f	0.07σ JUNO	Ch 8
13	$\sin^2 \theta_{13}$	$4/(18\pi^2)$	S	1	Within 1%	Ch 8
14	$\sin^2 \theta_{23}$	$1/2 + 1/(2\pi^2)$	L	1+2(b) ^f	Within 1σ	Ch 8
15	m_e, m_μ via Koide	from m_τ	M	3	Mass-chain inheritance	Ch 6 §6.5
16	m_b/m_τ	$4.28/Z_S(\lambda g_{\theta}^2) = 2.361$		2(a)	0.4% (0.5 σ at L=32)	Ch 6 §6.6
17	m_u/m_d	$\sqrt{2/9}$	L	1+2(b)	Within 1σ	Ch 6 §6.6
18	Σm_ν	~ 0.059 eV	S	1	Within 1σ	Ch 8
19	PMNS sum rule	$2\Delta_{12} + \Delta_{23} = 0$	L	1+2(b) ^f	1.178 ± 0.020 confirmed	Ch 8

#	Prediction	Formula	Class	Layer	Empirical match	Chapter
20	m_t (weak-scale $y_t \approx 1$)	$v/\sqrt{2}$	R	3	0.9%; mechanism open	Ch 6 §6.7
21	m_b via Q_{down}	derived	L	2(b)	Within 1σ	Ch 6 §6.7
22	$\lambda(M_{\text{Pl}}) = 0$	structural	S	1	0.6σ	Ch 6 §6.8
23	MOND acceleration a_0	$cH/6$	S	0/1	Within 0.5%	Ch 7 §7.8
24	Bekenstein-Hawking entropy	$S = A/(4l_p^2)$	S	1	derived; no direct test	Ch 7 §7.5

Summary statistics. Counts: 12 S + 7 L + 1 M + 2 R + 2 E. Total: 24. This table is the book’s reference index and is broader in scope than the paper’s Standard-Model table ([SM §7.6], 22 rows): it additionally carries the two structural inputs to the coupling sector ($1/\alpha_0$, $A \cdot B$), the neutrino-sector entries, and the two gravitational results (a_0 , S_{BH}) that also appear in §A.4.2. The headline count of *seven* parameter-free structural predictions among twenty-two Standard Model observables derives from [SM §7.6], not from the row total here. The C class is currently empty (the Cond 2 prediction that previously occupied C was reclassified following the §B.2.2 audit). The P class is currently empty. Every prediction within approximately 1% or 1σ of observation.

^f *Layer 2(b) ($\sin^2\theta_{12}$, $\sin^2\theta_{23}$, PMNS sum rule).* The Layer 1 component is structurally protected against the A_1 perturbation by the cosines-dressing of the TBM sum rule (cubic geometry alone). The remaining Layer 2 component is the A_2 Wilson-coefficient ratio $b_{23}/(b_{12} + b_{13}) = 4/3 = 4(d-1)/(2d)$, equivalent to the substrate-level relation $b_{23}^{\text{sub}} = (4/3)(b_{12}^{\text{sub}} + b_{13}^{\text{sub}}) + 1/9$. The ratio’s 2(a)-tractable period — the label meaning the residual question was sharply specified and admitted definite outcomes from substrate calculations — has ended: the calculation was performed. It does *not* mean the answer is forced by representation theory alone — the May 6, 2026 audit (§B.2.2) established that the natural Higgs-democratic spurion does not produce this ratio (instead giving $b_{12} : b_{13} : b_{23} \approx 17 : 1 : 1$ in the heavy-quark hierarchy, off by factor ~ 24 from the required $4/3$), and the relation requires substrate amplitudes with the non-natural scaling $\delta M_{ij}^{\text{sub}} \propto m_{\max(i,j)}$ plus a channel-specific enhancement factor $K_{23}/K_{12} \approx 2.8$. The substantive content is bijection-specific — Layer 2(b) — and the sharply specified verification path has since been executed: the §7.3 routing computation excludes the second-order route (the two-hop operator is a pure spin-taste monomial, orthogonal to every mass channel), leaving the recorded $O(\lambda^3)$ construction as the one unpursued analytical route. The empirical match at 0.07σ (JUNO) and 0.6σ (sum rule) confirms the bijection

has the required structure; the 7/6 sum rule remains the direct experimental test.

Distinctive empirical content. Three predictions stand out as the framework’s strongest empirical content. (i) JUNO $\sin^2 \theta_{12} = 1/3 - 1/(4\pi^2) = 0.3080$ at **0.07** σ against the post-JUNO global fit — a parameter-free retrodiction; its exact, freedom-free form makes JUNO’s design-lifetime precision (± 0.0014) a pre-registered forward test. (ii) the Milky Way Keplerian decline at the predicted $r_M \sim 17$ kpc — the framework’s strongest gravitational empirical content; the Bekenstein-Hawking coefficient is derived but has no direct test (GW250114 constrains the classical area theorem). (iii) Koide $Q = 2/3$ at **0.02%** — the sharpest framework prediction across all 24, matching observation at four-significant-figure precision.

A.4.2 Gravitational-sector observables

The framework’s gravitational predictions match general relativistic predictions at the emergent level while making distinctive structural predictions about the gravitational sector’s origin and parameters.

#	Prediction	Formula	Class	Layer	Empirical match	Chapter
G1	Discreteness scale	$\epsilon = 2l_p$	S	0	Structural	Ch 7 §7.4
G2	Bekenstein bound	$S_{\text{dS}} \sim 10^{122}$	S	0	Within 1%	Ch 7 §7.4
G3	Planck constant from gap eqn	$\hbar$ derived	S	1	Structural	Ch 7 §7.4
G4	Bekenstein-Hawking coefficient	1/4	S	1	derived; no direct test	Ch 7 §7.5
G5	Hawking radiation rate	structural	S	1	Within 1%	Ch 7 §7.5
G6	MOND acceleration $a_0 = cH/6$	structural	S	0/1	Within 0.5%	Ch 7 §7.8
G7	CC dissolution	compression ratio	S	1	Structural	Ch 7 §7.7

#	Prediction	Formula	Class	Layer	Empirical match	Chapter
G8	Running vacuum ν_{OI}	functional form Type II RVM (structural); emergent-QFT-scale magnitude $ \nu \approx \nu_{\text{QFT}} \sim (M_{\text{SM}}/M_{\text{Pl}})^2 \sim 10^{-32}$, indistinguishable from zero	L	2(a)	DESI Year 5 decisive ($\sim 1\%$ on ν); Bertini 2025 $\nu = -(2.5 \pm 1.3) \times 10^{-4}$ consistent at 2σ	Ch 7 §7.8
G9	Dark sector budget	$\sim 95\%$	S	1	Within 1%	Ch 7 §7.9
G11	Page curve	nested trace-out	S	1	Structural	Ch 7 §7.6
G12	$\Sigma m_\nu \sim 0.059 \text{ eV}$	structural	S	1	Within 1σ	Ch 8
G13	Lorentz violation	$O(\epsilon^2/c^2)$	S	0	Below current limits	Ch 7 §7.4
G14	Boundary entropy	$\rho_s = \rho_{\text{crit}}$	S	1	Exact	Ch 7 §7.9

Summary statistics. Counts: 12 S + 1 L. Total: 13. The G6 MOND prediction has dual Layer 0/1 status (the discreteness scale at Layer 0 plus the mode-counting at Layer 1). Cumulative empirical match at observational precision for all entries within current experimental capability.

Distinctive content. The MOND acceleration $a_0 = cH/6$ at 0.5% (entry G6) is the framework's strongest gravitational empirical content. The Bekenstein-Hawking coefficient (entry G4) is derived but has no direct empirical test: GW250114 constrains the classical area theorem, which the framework preserves, not the entropy coefficient. The Running Vacuum Model coefficient ν_{OI} (entry G8) provides forward-looking exposure to DESI data over the next 2-3 years.

A.4.3 Neutrino-sector observables

The neutrino sector provides the framework’s sharpest parameter-free retrodiction. The JUNO measurement of $\sin^2 \theta_{12}$ at 0.07σ against the framework’s structural value is a retrodictive match whose exact form converts JUNO’s design-lifetime precision into the framework’s cleanest pre-registered forward test.

#	Prediction	Formula	Class	Layer	Empirical match	Chapter
N1	$\sin^2 \theta_{12}$	$1/3 - 1/(4\pi^2)$	L	1+2(b)	0.07σ JUNO	Ch 8
N2	$\sin^2 \theta_{13}$	$4/(18\pi^2)$	S	1	Within 1%	Ch 8
N3	$\sin^2 \theta_{23}$	$1/2 + 1/(2\pi^2)$	L	1+2(b)	Within 1σ	Ch 8
N4	Σm_ν	~ 0.059 eV	S	1	Within 1σ	Ch 8
N5	Normal mass ordering	structural	S	1	Confirmed	Ch 8
N6	PMNS sum rule	$2\Delta_{12} + \Delta_{23} = 0$	L	1+2(b)	1.178 ± 0.020 confirmed	Ch 8
N7	δ_{CP}	$b_{12} - b_{13}$ (sub-dominant U-odd)	L	2(b) ^g	T2K/NOvA: bijection-constraining	Ch 8

Summary statistics. Counts: 3 S + 4 L. Total: 7. The L-class entries divide into two kinds: (i) **sub- σ confirmed predictions** — N1, N3, N6 — whose Layer 1 structural form (TBM zeroth order plus the λ^2 scale) holds unconditionally and whose specific values follow given the substrate Wilson-coefficient ratio $b_{23}/(b_{12} + b_{13}) = 4/3$, classified as Layer 2(b) per the §7.3 routing computation (the second-order derivation route excluded; see footnote ^f above); (ii) **parameterization-level prediction** — N7 — where the framework identifies $b_{12} - b_{13}$ as the Wilson-coefficient combination controlling δ_{CP} (a structural-form statement) but the specific value is bijection-specific and the framework “makes no commitment to any specific value of δ_{PMNS} ” per SM §6.4 / §7.3 (see footnote ^g).

^g *Parameterization-level prediction (δ_{CP}).* The framework identifies δ_{CP} with the sub-dominant U-odd combination $b_{12} - b_{13}$ in the A_2 second-order corrections to the charged-lepton sector (SM §7.3 line 901). This is a structural-form prediction: it specifies how δ_{CP} enters the framework’s parameterization but does not predict its value. The value of $b_{12} - b_{13}$ is bijection-specific (Layer 2(b), property of the particular φ); the framework “makes no commitment to any specific value of δ_{PMNS} ” (SM §6.4 line 546). Future T2K/NOvA measurements constrain the

bijection's $b_{12} - b_{13}$ value rather than testing a framework prediction of δ_{CP} itself. This is qualitatively different from N1/N3/N6, where specific numerical values are fixed with no freedom and the JUNO-era fit matches $\sin^2\theta_{12}$ at 0.07σ (retrodictively).

A.4.4 Biology and medicine predictions

The framework's empirical content in biology and medicine is concentrated in the wider therapeutic window for reader-writer disorders (consistent with the Rett-syndrome reversibility of Guy et al. 2007) and the cancer methylome classification paradigm — structural organization of, and consistency with, established clinical biology.

#	Prediction	Quantitative content	Class	Layer	Empirical match	Chapter
B1	Wider therapeutic window for Rett	Substratum-emergent operator dissociation	S	1	Retrodiction (Guy et al. 2007)	Ch 16 §16.12
B2	Cancer methylome classifier accuracy	Substratum preserved during tumorigenesis	S	1	DKFZ 95%+ accuracy	Ch 17 §17.7
B3	Non-Markovian enzyme kinetics	Stretched-exponential waiting times	S	1	Multiple enzyme systems	Ch 13 §13.6
B4	Vibronic coherences in photosynthesis	Picosecond timescale	C	1	FMO observed; does not discriminate against standard vibronic accounts	Ch 13 §13.4
B5	Millisecond radical-pair coherence in cryptochrome	Slow-bath C2	C	1	Wiltchko 2019	Ch 13 §13.5

#	Prediction	Quantitative content	Class	Layer	Empirical match	Chapter
B6	Molecular clock overdispersion	$R(t) \approx 5$ for mammals	S	1	Gillespie 1989, Cutler 2000	Ch 17 §17.9
B7	Universal rate autocorrelation	Across the tree of life	S	1	Tao et al. 2019 confirmed	Ch 17 §17.9
B8	LTEE power-law fitness	$\beta \approx 0.08-0.12$	S	1	Wiser et al. 2013 confirmed	Ch 17 §17.9
B9	Punctuated equilibrium	C1-C3 information backflow	S	1	Eldredge-Gould 1972	Ch 17 §17.9
B10	Hsp90 evolutionary capacitance	C1-C3 buffering	S	1	Rutherford-Lindquist 1998	Ch 17 §17.9
B11	Chk1 + gemcitabine schedule dependence	Memory asymmetry	S	1	Clinical literature consistent	Ch 16 §16.3
B12	LIFU cognitive improvement in Alzheimer's	τ_B normalization	S	1	Ye et al. 2025 confirmed	Ch 16 §16.4
B13	iPSC reprogramming residuals	Substratum-emergent dissociation	S	1	Lin et al. 2024 confirmed	Ch 16 §16.10
B14	Bivalent chromatin states	Multi-operator substratum	S	1	Bernstein et al. 2006 confirmed	Ch 16 §16.10

Summary statistics. Counts: 12 S + 2 C. B1 (Rett syndrome wider therapeutic window, Guy et al. 2007) is a retrodiction, in the same category as B6-B10 —

the framework accounts for phenomena established before it. None is a forward confirmation. B4 and B5 are C rather than S because coherence preservation is a property of the specific physical realization and is not implied by (C1)-(C3) (§13.4-13.5): the framework’s content in those two rows is the non-Markovian accounting, not a derivation that coherence survives.

A.5 Forward predictions and falsification targets

Beyond the predictions with current empirical match, the framework makes specific forward-looking predictions that are testable but not yet tested. These provide the framework’s empirical exposure for the next several years of experimental research.

Quantum engineering.

#	Prediction	Quantitative content	Testability	Chapter
F1	BEC analogue-gravity $\mathcal{F}(r) = 2r - r^2$	Capacity-controlled non-Markovianity	Steinhauer group, 2-3 years	Ch 15 §15.8
F2	Correlation-aware error correction	correlated errors over $\sim \tau_B/\tau_S$ gates; gain not quantified	Current quantum hardware	Ch 15 §15.4
F3	Backflow-aware NV magnetometry	sensitivity gain in direction only; factor not quantified	Current NV platforms	Ch 15 §15.6
F4	Engineered-bath coherence extension	coherence beyond the Markovian estimate; factor not quantified	Spin chain / NV / trapped ion platforms	Ch 15 §15.5

Cosmology and gravity.

#	Prediction	Quantitative content	Testability	Chapter
F5	Gravitational wave echoes	$\Delta t = (r_+/c) \ln(r_+/2l_p)$	LIGO/Virgo/KAGRA next observing runs	Ch 18 §18.5
F6	Running vacuum DESI test	$w(z) \neq -1$ in specific form	DESI Year 5 (2-3 years)	Ch 7 §7.7

#	Prediction	Quantitative content	Testability	Chapter
F7	Born rule transient violations	CMB non-Gaussianity signatures	LiteBIRD, future CMB-S4	Ch 18 §18.2
F8	Cosmic recurrence	$\sim e^{10^{122}}$ timescale	Logically falsifiable only	Ch 18 §18.7

Biology and medicine.

#	Prediction	Quantitative content	Testability	Chapter
F9	Variant pathogenicity stratification for reader/writer/eraser	AlphaMissense miscalibration for machinery genes	Reanalysis of existing benchmarks	Ch 17 §17.8
F10	τ_B -matched cancer chemotherapy schedules	Memory-priming protocols	Animal models, clinical trials	Ch 16 §16.3
F11	Anti-persister antibiotic class	Memory-erasing drugs	Drug screening with regulatory dynamics	Ch 16 §16.5
F12	Hurst exponent of substitution rates	$H \approx 0.79 \pm 0.03$ for typical taxa — <i>retrodiction, not a forward bet: H is read off existing data, and $H \approx 0.7$–0.8 is the generic broad-relaxation value (§8.11 of [Complexity]), so it does not discriminate the framework from a mundane superposition null; only the derived capacity-dependence $H(C3)$ would be a genuine test, and that relationship is not yet derived</i>	Existing phylogenetic data (consistency, not exposure)	Ch 17 §17.9

Foundations (Lorentz sector).

#	Prediction	Quantitative content	Testability	Chapter
F13	Cross-charged boost-LV decoupling	$\delta c \sim \Delta C_2 (\alpha_0/4\pi)(M_X/M_{\text{Pl}})^2 [c_L \ln(M_{\text{Pl}}/M_X) + c_0]$; quadratic-log <i>form</i> established; the logarithmic coefficient is now converged at $c_L \approx -9.4 \pm 0.3$ (extrapolated to $N \rightarrow \infty$, superseding an earlier under-converged -7.5 ± 1), which tightens the consistency window (against the electron-sector SME coefficient) to $M_X \lesssim 3.9 \times 10^9$ GeV; the framework's <i>primary</i> window is set by the tighter GZK/Auger isotropic- $n=2$ bound at $M_X \lesssim 4 \times 10^7$ GeV (SM §3.1 (<i>Scope of the emergent-Lorentz claim</i>)), with the electron-SME and photon-sector ceilings ($7.5 \times 10^9, 1.1 \times 10^{12}$ GeV) the looser sector-specific bounds; the overall normalization ΔC_2 remains open, so this is a window conditional on the bijection, not a forward prediction of M_X	Multi-species $\ln(M_{\text{Pl}}/M_X) +$ (form + coefficient; full reach pending ΔC_2)	SM §3.1 (<i>Scope of the emergent-Lorentz claim</i> , point (iii))

The *rotational* Lorentz sector is deliberately **not** listed as a forward prediction. Unlike the boost sector (F13), the rotational E_g anisotropy is a custodial-U(6) singlet with no power-law suppression ally, and its magnitude is undetermined — bifurcating, on the uncomputed smearing-kernel width R/ε , between a negligible dimension- ≥ 6 residual ($\sim 10^{-32}$) and an unsuppressed dimension- ≤ 4

operator already constrained by laboratory sidereal experiments. It is therefore the sharpest *open question* in the framework’s Lorentz sector, treated as such in SM §3.1 (*Scope of the emergent-Lorentz claim*, point (iii)), not a predictive bet. Its one regime-independent feature — that any detectable residual must take the $\ell = 4$ cubic-harmonic form aligned to the cosmic rest frame and correlated across species — is recorded there as a conditional falsification template, not here as a forward prediction.

Falsification conditions. Each forward prediction has specific falsification conditions developed in the relevant chapter. The framework’s content is most distinctively at risk in:

- F1 (BEC analogue-gravity): null result with functional form different from $2r - r^2$ would falsify the framework’s structural mechanism for partially-quantum dynamics.
- F5 (gravitational wave echoes): null result at $\Delta t = (r_+/c) \ln(r_+/2l_p)$ would constrain the framework’s substratum-level lattice scale.
- F6 (DESI running vacuum): DESI Year 5 confirmation of phantom-crossing dark energy would falsify the framework’s running vacuum prediction.
- F9 (variant pathogenicity stratification): demonstration that AlphaMis-sense performance is uniform across reader/writer/eraser and substrate genes would falsify the substratum-emergent operator distinction at the variant-prediction level.
- F13 (cross-charged decoupling): observation of an inter-species boost-Lorentz-violation that does *not* decouple with the heavy cross-charged mass — i.e. a non-decoupling (constant or pure-logarithmic) sector-differential limiting-speed difference — would falsify the framework’s condensate-stabilizer mechanism for gauge content (Theorem 5, revised), under which the differential is suppressed as $(M_X/M_{\text{Pl}})^2$ up to a logarithm. The *form* of this suppression is established and the logarithmic coefficient is now converged ($c_L \approx -9.4 \pm 0.3$ at $N \rightarrow \infty$), giving a consistency window $M_X \lesssim 3.9 \times 10^9$ GeV; the overall normalization ΔC_2 (and hence whether this window is also a forward prediction rather than a bijection-conditional consistency bound) is not yet fixed.

A.6 Framework’s epistemic position summary

The framework’s cumulative empirical record can be summarized as follows.

Strongest current evidence. The framework currently has no pre-registered forward confirmation. Its strongest empirical content is retrodictive: JUNO $\sin^2 \theta_{12}$ at 0.07σ (entry 12 / N1; the derivation postdates the November-2025 first results), Koide Q at 0.02% , and Cabibbo at 0.04% . The genuinely pre-registered content is forward-looking — DESI Y5 (ν), the anisotropy amplitude, BTFR at $z > 1$, the A2 lattice thresholds, and JUNO design-lifetime precision on the exact $\sin^2 \theta_{12}$ value — documented in dated artifacts preceding their

future data. The Bekenstein-Hawking coefficient (entry 24 / G4) is derived but has no direct empirical test — GW250114 constrains the classical area theorem, not the coefficient. The Rett wider-window result (entry B1) is a retrodiction, not counted among the forward confirmations.

Most distinctive content. Two predictions distinguish the framework from standard accounts most sharply: the substratum-emergent operator distinction (with empirical signatures in B1, B2, B13, B14) and the partially-quantum regime $\mathcal{F}(r) = 2r - r^2$ (with empirical exposure in F1). Neither has analog in standard quantum mechanics or effective field theory.

Sharpest single match. Koide $Q = 2/3$ at 0.02% (entry 10) — the framework’s sharpest individual empirical match.

Highest-precision overall. All twenty-two Standard Model predictions within approximately 1% or 1σ — cumulative empirical content matching at observational precision across the full SM.

Most consequential forward exposure. BEC analogue-gravity test (F1) and DESI Year 5 running vacuum test (F6) provide the framework’s most consequential near-term empirical exposure. Both have specific falsification conditions that would substantially constrain the framework’s content if not confirmed.

Domains where framework has substantial content but no current empirical match. The framework’s content in computational complexity (BQP theorem, Chapter 14), evolutionary biology Hurst exponent (F12), and the broader reader-writer disorder class beyond Rett (F9, F10, F11) has not yet been specifically tested against the framework’s predictions. These are testable predictions awaiting empirical engagement rather than predictions matching established data.

Where the framework has no current content. Five fundamental physics problems are scope limits from the framework’s perspective — places where it does not provide a derived value: the fermion mass hierarchy in detail (flavor problem), baryogenesis (solution-specific — the baryon asymmetry inherits the input status of the framework’s flavor-sector CP phases, §19.3.3), inflation and initial conditions, the Hubble tension’s specific resolution, and the cosmological initial state. These are explicitly characterized in Chapter 19.

Cumulative assessment. The framework’s empirical record spans multiple domains, and its entries are of unequal evidential weight, classified explicitly above: of the twenty-two Standard Model observables, seven are parameter-free structural retrodictions and the remainder are layered, chained, or fitted; the thirteen gravitational-sector and seven neutrino-sector entries include retrodictions, consistency results, and one definitional fraction, with the 0.07σ JUNO match the sharpest (a retrodiction — the derivation postdates the November 2025 measurement); the fourteen biology/medicine entries are retrodictions and consistency results. Genuinely pre-registered forward content is the smaller set

awaiting data (DESI Y5 ν , the rotational-anisotropy amplitude, BTFR at $z > 1$, the A2 lattice thresholds, JUNO design-lifetime precision on the exact $\sin^2 \theta_{12}$ value). The framework’s empirical case is therefore a broad retrodictive compression across domains, with its forward-test exposure concentrated in that smaller pre-registered set, rather than a single decisive confirmation.

The framework’s distinctive epistemic position is that it provides *structural foundations* for the empirical patterns rather than competing theoretical content at the emergent level. Standard QM, standard QFT, and standard GR are reproduced at the emergent level; the framework’s distinctive content is the structural foundations explaining why these emergent theories have the specific forms they do. The cumulative empirical record provides convergent evidence for the framework’s structural commitments across multiple domains, with no single empirical match being load-bearing for the framework’s overall content.

Appendix B

Mathematical Derivations

B.1 What this appendix develops

The main chapters of the book develop the framework’s content with proofs sketched at the level required for the chapter argument’s logical flow. This appendix collects the full mathematical derivations of the framework’s key technical results — derivations that are referenced in the main chapters but kept out of the chapter flow to preserve readability.

The appendix’s content is organized by mathematical structure rather than by chapter sequence. Six derivations occupy the appendix.

B.2 The Stinespring construction. The framework’s emergent quantum mechanics derives from a deterministic substratum bijection through the Stinespring dilation theorem applied to a partial trace over the hidden sector. This appendix gives the full construction: Hilbert-space embedding, permutation unitarity lemma, CPTP channel structure, Born rule recovery, emergent coherence theorem, and CP-indivisibility theorem. The construction uses only textbook operator algebra from the 1950s.

B.3 The trace-out as a Jordan-Chevalley projection. The framework’s content on the Standard Model gauge structure relies on a structural property of the trace-out operation: it extracts the semisimple part of the evolution matrix’s Jordan-Chevalley decomposition and discards the nilpotent monodromy. This appendix gives the full algebraic derivation: period formula, Jordan-Chevalley

decomposition, q -independence of the Weil-Deligne conductor, additive decomposition over gauge irreps, and the precise sense in which the trace-out is a projection.

B.4 The gap equation for $\hbar$. The framework derives Planck’s constant $\hbar$ from a thermal self-consistency argument between the classical horizon temperature (substratum-level) and the KMS temperature (emergent QFT-level). This appendix gives the full four-step derivation: spatial locality and boundary dependence, deep-sector gauge equivalence, dimensional determination, and thermal self-consistency producing $\hbar = c^3 \epsilon^2 / (4G)$.

B.5 Anomaly cancellation and unique hypercharges. The framework’s content on the Standard Model matter content derives the unique hypercharge assignment from the six anomaly conditions of $SU(3) \times SU(2) \times U(1)$. This appendix gives the explicit calculation: the six anomaly conditions, the algebraic constraint structure, and the unique solution producing $(Y_Q, Y_u, Y_d, Y_L, Y_e) = (1/6, 2/3, -1/3, -1/2, -1)$ with no free parameters.

B.6 Gauge-theoretic derivations. Two key gauge-theoretic results from Chapter 5: the coupling-degree minimization derivation of $K = 2d = 6$ internal components per site, and the cubic-rotation-group decomposition of the coupling matrix’s eigenvalue multiplicities into $(3, 2, 1)$ — the multiplicities that determine $SU(3) \times SU(2) \times U(1)$.

B.7 The lemma chain for emergent quantum mechanics. The framework’s characterization theorem (emergent QM from C1-C3) relies on a chain of intermediate lemmas: partition-relativity, emergent stochasticity, P-indivisibility from C1-C3, accessible-timescale backflow, and phase-locking. This appendix collects these lemmas with their proofs in compact form, providing the reader a self-contained reference for the framework’s logical structure.

The appendix’s content is *complete derivations*, not exposition. Each derivation gives the algebraic steps, the lemmas invoked, and the result with no glosses. For motivation and context, the relevant main-chapter sections are cross-referenced at the start of each derivation. The appendix is designed for readers wanting full technical content; the main chapters give the same results with motivating context but compressed derivations.

B.2 The Stinespring construction for emergent quantum mechanics

Main-text reference: Chapter 1 §1.7-1.9.

The framework’s emergent quantum mechanics derives from the substratum’s deterministic bijection through a Stinespring dilation applied to the partial trace over the hidden sector. The construction uses only textbook operator algebra from the 1950s, with no reference to recent stochastic-quantum correspondence results.

Setup. The finite configuration spaces $\mathcal{C}_V = \{x_1, \dots, x_n\}$ (visible sector) and $\mathcal{C}_H = \{h_1, \dots, h_m\}$ (hidden sector) embed into Hilbert spaces $\mathcal{H}_V = \mathbb{C}^n$ and $\mathcal{H}_H = \mathbb{C}^m$ via the canonical identification $|i\rangle \leftrightarrow x_i$ and $|k\rangle \leftrightarrow h_k$. This introduces no quantum postulates: it is the canonical identification of probability distributions on a finite set with diagonal density matrices.

The combined space is the tensor product $\mathcal{H} = \mathcal{H}_V \otimes \mathcal{H}_H$ of dimension nm . The substratum's deterministic bijection $\varphi : \mathcal{C}_V \times \mathcal{C}_H \rightarrow \mathcal{C}_V \times \mathcal{C}_H$ defines a permutation of the orthonormal basis $\{|i, k\rangle\}_{i=1, \dots, n; k=1, \dots, m}$.

Lemma B.2.1 (Permutation unitarity). *Any bijection $\varphi : \mathcal{C}_V \times \mathcal{C}_H \rightarrow \mathcal{C}_V \times \mathcal{C}_H$ defines a unitary operator U_φ on $\mathcal{H}$.*

Proof. Define $U_\varphi |i, k\rangle = |\varphi(x_i, h_k)\rangle$. Since φ is a bijection, U_φ permutes the orthonormal basis, hence is unitary. $\square$

For continuous-time dynamics φ_t , Stone's theorem on $\mathcal{H}$ yields $U_t = e^{-i\hat{H}t}$ for some Hermitian generator $\hat{H}$.

Lemma B.2.2 (Reverse direction, scoped). *Any single one-step doubly stochastic matrix on a finite configuration space $\mathcal{C}_V$ can be realized as the one-step marginal of a deterministic bijection on $\mathcal{C}_V \times \mathcal{C}_H$ with uniform prior on $\mathcal{C}_H$, for some finite $\mathcal{C}_H$ (exactly when transition probabilities are rational; to arbitrary precision otherwise). (The restriction is forced: a uniform-prior bijection marginal preserves the uniform distribution, hence is doubly stochastic; the emergent unistochastic statistics lie inside this class.) The scope is the matrix, not the process: realizing prescribed multi-time statistics — a family of joint distributions over trajectories — with a single bijection and a single uniform prior is a stronger demand and is not claimed here.*

Proof. Any marginal of a bijection with uniform prior is a doubly stochastic matrix. By Birkhoff-von Neumann, every doubly stochastic matrix is a convex combination of permutation matrices. A bijection on $\mathcal{C}_V \times \mathcal{C}_H$ realizing the required mixture is obtained by letting $\mathcal{C}_H$ enumerate the permutations; for multi-step processes, the construction extends by history dilation. $\square$

Remark. Lemma B.2.2 is the bijection-stochastic-matrix correspondence at the stochastic-process level. It is *not* a full CPTP-channel correspondence — channels that create coherences from diagonal inputs have no permutation-unitary realization with uniform ancilla. For the characterization theorem (Chapter 1 §1.7), which is stated in terms of transition probabilities, this is the appropriate scope.

The quantum channel. The observer's ignorance of the hidden sector corresponds to the maximally mixed state $\rho_H = I_m/m$. The visible-sector quantum channel is

$$\Phi(\rho_V) = \text{Tr}_H [U_\varphi (\rho_V \otimes \rho_H) U_\varphi^\dagger]$$

This is completely-positive trace-preserving (CPTP) by a standard result (Nielsen-Chuang Theorem 8.1), with Kraus representation $\Phi(\rho_V) =$

$\sum_{k,l} K_{kl} \rho_V K_{kl}^\dagger$ where $K_{kl} = m^{-1/2} \langle l | U_\varphi | k \rangle_H$. The triple $(\mathcal{H}_H, U_\varphi, \rho_H)$ is the *Stinespring dilation* of Φ .

Theorem B.2.3 (Born rule recovery). *The classical transition probabilities $T_{ij} = P(j|i)$ derived from the substratum's bijection equal the Born-rule probabilities of Φ .*

Proof. $P(j|i) = \langle j | \Phi(|i\rangle\langle i|) | j \rangle = m^{-1} \sum_{k,l} |\langle j, l | U_\varphi | i, k \rangle|^2$. Since U_φ is a permutation, $\langle j, l | U_\varphi | i, k \rangle = \delta_{(j,l), \varphi(i,k)}$. Thus $P(j|i) = m^{-1} \sum_k \delta_{j, \pi_V(\varphi(x_i, h_k))} = T_{ij}$, where π_V is the projection from $\mathcal{C}_V \times \mathcal{C}_H$ to $\mathcal{C}_V$. $\square$

Remark (failure of (C1)). Condition (C1) does not suffice. For $|V| = |H| = 2$ and $\varphi(x, h) = (h, x)$ one has $T_{ij} = 1/2$ for all i, j , yet $\Phi(\rho) = I/2$ identically: a constant channel, hence measure-and-prepare. Exhaustive enumeration gives the same verdict for every (C1)-satisfying bijection at $|V| = |H| = 2$, and for 612 of 648 at $|V| = 2, |H| = 3$. The obstruction is structural: (C1) constrains the diagonal of Φ , whereas entanglement-breaking is a property of $\text{id} \otimes \Phi$ and is not fixed by the action on single-system inputs.

What replaces (C1) is not a further condition on T but a condition on the partition, and it is established for a specific realization rather than for arbitrary bijections. Take the substratum of §B.1 on a lattice Λ with $q = 2$, state $(u, v) \in \mathbb{F}_2^\Lambda \times \mathbb{F}_2^\Lambda$ and the reversible nearest-neighbour update

$$u'_x = \sum_{z \sim x} u_z + v_x, \quad v'_x = u_x,$$

the visible sector being the sites of a region $R \subseteq \Lambda$. Write M_{HV} and M_{VH} for the blocks of this update carrying visible input to hidden output and hidden input to visible output, $\kappa = \dim \ker M_{HV}$, $w = \text{rank } M_{VH}$, and $s = \dim V = 2|R|$. The inner and outer boundaries of R are $\partial^- R = \{x \in R : \exists z \notin R, z \sim x\}$ and $\partial^+ R = \{y \notin R : \exists z \in R, z \sim y\}$.

Lemma 1 (boundary bound). $\text{rank } M_{HV} \leq |\partial^+ R|$ and $\text{rank } M_{VH} \leq |\partial^- R|$.

Proof. For $y \notin R$ the component $v'_y = u_y$ is a hidden input, and $u'_y = \sum_{z \sim y} u_z + v_y$ involves a visible input only if some $z \sim y$ lies in R , that is only if $y \in \partial^+ R$; so M_{HV} has non-zero rows only among $\{u'_y : y \in \partial^+ R\}$. Dually, for $x \in R$ the component $v'_x = u_x$ is a visible input, and u'_x involves a hidden input only if $x \in \partial^- R$. $\square$

One update step transports information one lattice site, so the sectors resolve one another only across the cut. This argument is a statement about supports and uses no property of the coefficient ring: it holds verbatim over $\mathbb{Z}_q$, and for a general nearest-neighbour bijection it yields the weaker conclusion that the hidden output depends on the visible input only through $\partial^- R$.

Lemma 2 (channel normal form). *Write $A = M_{VV}$, $B = M_{VH}$, $C = M_{HV}$ for the blocks of the update, and L_{RR} for the adjacency of Λ restricted to R .*

Then: (i) A is invertible, with $A^{-1}(a_u, a_v) = (a_v, a_u - L_{RR} a_v)$, so the visible block alone induces the permutation (Clifford) unitary $V_A|i\rangle = |Ai\rangle$; (ii) $\Phi(|i\rangle\langle i'|)$ vanishes unless $C(i - i') = 0$; (iii) on the Weyl operators the channel acts, with no phase, as

$$\Phi(X^a Z^b) = [Ca = 0] [B^T A^{-T} b = 0] X^{Aa} Z^{A^{-T}b},$$

so that $\Phi = \text{Ad}_{V_A} \circ \Phi_G$, where $\Phi_G = |G|^{-1} \sum_{g \in G} g \cdot g^\dagger$ is the uniform mixture over the group G of Weyl operators $X^\alpha Z^\beta$ with $\alpha \in \text{im}(A^{-1}B)$ (translations) and $\beta \in \text{im } C^T = (\ker M_{HV})^\perp$ (phases); (iv) G is abelian, of order $2^{s-\kappa+w}$.

Proof. (i) is the displayed inverse, read off from $v'_x = u_x$ and $u'_x = (L_{RR}u)_x + v_x$ on R ; invertibility of A is what makes $i \mapsto Ai$ a relabelling. (ii) The hidden outputs of i and i' agree for some hidden input exactly when $C(i - i') = 0$, and then for every hidden input. (iii) With the hidden input uniform, $\Phi(|i\rangle\langle i'|) = [C(i - i') = 0] \mathbb{E}_{t \in \text{im } B} |Ai + t\rangle\langle Ai' + t|$. Summing this against $X^a Z^b = \sum_i (-1)^{b \cdot i} |i + a\rangle\langle i|$ and substituting $i \mapsto i - A^{-1}t$ gives the displayed rule: the character sum over $\text{im } B$ produces the second indicator and cancels every phase. The right-hand side is Ad_{V_A} applied to the Hilbert–Schmidt-orthogonal projection onto $\text{span}\{X^a Z^b : Ca = 0, B^T A^{-T} b = 0\}$, and the uniform mixture over the symplectic complement of that index set — which is exactly G — is that projection. (iv) An element $X^\alpha Z^\beta$ with $\alpha \in \text{im}(A^{-1}B)$, $\beta \in \text{im } C^T$ commutes with another iff the cross-pairings $\alpha \cdot \beta'$ vanish; over all of G this is the single condition $CA^{-1}B = 0$, which holds identically for this update: B feeds hidden u into the visible u' -components, A^{-1} carries the u' -subspace onto the v -subspace, and C reads only the visible u -components. The translation and phase families intersect only in the identity, so $|G| = 2^{\text{rank } B} \cdot 2^{\text{rank } C} = 2^{w+(s-\kappa)}$. $\square$

Theorem (separability threshold). *Let G be an abelian subgroup of the Weyl group on s qubits, of order 2^t modulo phases. Then $\Phi_G = |G|^{-1} \sum_{g \in G} g \cdot g^\dagger$ is entanglement-breaking if and only if $t = s$.*

Proof. Modulo phases the Weyl operators form the symplectic space $\mathbb{F}_2^{2s}$, on which abelian subgroups are exactly the isotropic subspaces; isotropy bounds $t \leq s$. Any isotropic subspace of dimension t is carried to $\langle Z_1, \dots, Z_t \rangle$ by a symplectic transformation (Witt’s extension theorem for the symplectic space $\mathbb{F}_2^{2s}$), and every symplectic transformation is implemented by a Clifford unitary U . Hence $U\Phi_G U^\dagger$ is complete dephasing on t qubits tensored with the identity on the remaining $s - t$. Complete dephasing is measure-and-prepare and therefore entanglement-breaking, and a tensor product of channels is entanglement-breaking exactly when every factor is — separability of the joint Choi state survives the partial trace onto any single factor pair, so one entangled factor Choi rules out separability of the whole; the identity on a non-trivial factor is not entanglement-breaking. Since conjugation by a unitary preserves the property, Φ_G is entanglement-breaking precisely when $s - t = 0$. $\square$

Corollary 1. *For one step of the displayed update on $\mathbb{F}_2$ — $u'_x = \sum_{z \sim x} u_z + v_x$, $v'_x = u_x$, with the hidden sector initially uncorrelated and maximally mixed — the visible channel Φ admits the Clifford-conjugated abelian Weyl-mixture normal form of Lemma 2 and is not entanglement-breaking if and only if $w < \kappa$. Immediate from Lemma 2 and the Theorem with $t = s - \kappa + w$: entanglement breaking is invariant under composition with the unitary Ad_{V_A} , so Φ is entanglement-breaking exactly when Φ_G is, and $t = s$ reads $w = \kappa$ (Lemma 2 gives $w \leq \kappa$, matching $t \leq s$).*

Corollary 2. *If $|\partial^- R| + |\partial^+ R| < 2|R|$ then Φ is not entanglement-breaking.*

Proof. $s = 2|R|$, so Lemma 1 gives $w + (s - \kappa) \leq |\partial^- R| + |\partial^+ R| < s$, that is $w < \kappa$. $\square$

Remark. For a d -dimensional region of linear size ℓ , $|\partial R| = O(\ell^{d-1})$ against $2|R| = 2\ell^d$, so the hypothesis of Corollary 2 holds whenever $\ell > 2d$; in $d = 3$ this is roughly six sites across. Within this single-step result the criterion carries no dependence on $|H|$, since the coupling reads only the boundary; the counterexample above is the degenerate case, every visible site lying on the cut so that the region has no bulk. Whether the same holds over many steps is not settled here: repeated interaction can propagate boundary information into the hidden bulk, and a multi-time statement would require a process-tensor argument rather than a single-channel one. The suggestion that coherence is a bulk quantity and decoherence an area quantity — the accounting that reappears in the horizon entropy of Chapter 7 — is at this stage an interpretation of the one-step result, not a theorem about the asymptotic dynamics.

Remark (scope). Lemma 1 holds over any $\mathbb{Z}_q$ and, in weakened form, for arbitrary nearest-neighbour bijections. Lemma 2, the Theorem and both Corollaries are stated for $q = 2$; the symplectic and Clifford arguments carry over to prime q with $2^\bullet$ replaced by $q^\bullet$, while composite q requires module-theoretic treatment of rank and kernel and is not covered. All of this concerns the realization displayed above and is not a property of arbitrary finite bijections. For a general linear bijection the abelian property of Lemma 2(iv) can fail ($CA^{-1}B \neq 0$); the mixture is then over a non-abelian Weyl family, and entanglement breaking is governed by the isotropy of the surviving index set rather than by its order alone — not treated here. Two boundary markers on interpretation. Coherence preservation is a property of this realization, not an additional structural condition: (C1)–(C3) constrain coupling, timescales, and capacity, and neither contain nor imply a coherence requirement — the swap example satisfies (C1) while its channel is entanglement-breaking, and which realizations preserve one-step coherence is the geometric question answered above for the displayed update alone. And non-entanglement-breaking is a statement about a single channel, not about quantum mechanics: it supplies neither the observable algebra, nor preparations and interventions, nor tensor composition, nor multi-time statistics — the operational reconstruction claims rest on the dilation and correspondence arguments, not on Corollary 1 — and the entanglement-breaking status of one step does not determine whether the visible multi-time process is P-divisible: the

hidden sector can carry memory even when a step's channel is entanglement-breaking, as a two-state system with a two-state hidden register already shows.

Theorem B.2.5 (CP-indivisibility). *The P-indivisibility of the substratum dynamics (under C1-C3) implies CP-indivisibility of $\{\Phi_t\}$: there exist $t_2 > t_1 > 0$ with no CPTP map Λ satisfying $\Phi_{t_2} = \Lambda \circ \Phi_{t_1}$.*

Proof. The framework's Φ_t are permutation dilations with uniform ancilla, so they map computational-diagonal states to exactly diagonal states (a permutation unitary conjugates a diagonal state to a diagonal state; the partial trace preserves diagonality). If $\Phi_{t_2} = \Lambda \circ \Phi_{t_1}$ with Λ CPTP, then $M_{kj} := \langle j | \Lambda(|k\rangle\langle k|) | j \rangle$ is stochastic and, by the diagonality of $\Phi_{t_1}(|i\rangle\langle i|)$, $T(t_2) = T(t_1)M$ — P-divisibility of the population process. The contrapositive gives the result. For a general CPTP family the diagonal-input reduction fails — a Hadamard pair is CP-divisible with P-indivisible populations — so the diagonal-preservation property is what licenses the step. $\square$

By the Breuer-Laine-Piilo criterion, CP-indivisibility implies non-monotonic trace distance — the *information backflow* signature of non-Markovianity. The framework's content is therefore that emergent quantum mechanics under C1-C3 is necessarily non-Markovian, with information backflow at observable timescales.

Approximate unitarity. On observable timescales $t \ll \tau_B$, the hidden-sector state is approximately frozen (conditions C2-C3). At $t = 0$ the channel generator is exactly Hamiltonian,

$$\left. \frac{d\Phi_t}{dt} \right|_{t=0} (\rho_V) = -i[\hat{H}_{\text{eff}}, \rho_V],$$

with $\hat{H}_{\text{eff}} = H_V + \text{Tr}_H[V_{\text{int}}(1 \otimes \rho_H)]$; the dissipative part vanishes at $t = 0$. At finite time the dissipator $\mathcal{D}$ has two contributions: a frozen-bath dephasing of order the coupling strength squared (the spread of bath-conditioned visible Hamiltonians, independent of τ_B), and a bath-motion correction of order $(\tau_S/\tau_B)^2$. The dynamics is the Schrödinger equation generated by $\hat{H}_{\text{eff}}$ to leading order provided both the slow-bath condition $\tau_S \ll \tau_B$ and a weak-coupling condition (small conditioned-Hamiltonian spread) hold. The phase-locking lemma (Chapter 1 §1.6) then determines $\hat{H}_{\text{eff}}$ uniquely from continuous-time transition data, up to the residual diagonal-unitary rephasing freedom that is physically trivial (basis convention).

The construction is complete: the substratum's deterministic bijection produces emergent unitary quantum mechanics on the visible sector through the Stinespring dilation, with leading-order Schrödinger evolution and structurally non-Markovian corrections at $\mathcal{O}(\tau_S/\tau_B)$.

B.3 The trace-out as a Jordan-Chevalley projection

Main-text reference: Chapter 5 §5.4-5.6.

The framework's content on the Standard Model gauge structure relies on a precise algebraic property of the trace-out operation: it extracts the semisimple part of the evolution matrix's Jordan-Chevalley decomposition and erases the nilpotent monodromy. This appendix gives the full derivation for the scalar wave equation on a ring of L sites over $\mathbb{F}_q$; the multi-component extension incorporating the gauge group follows by additive decomposition.

Setup. The discrete wave equation $x(n, t+1) = x(n-1, t) + x(n+1, t) - x(n, t-1) \pmod{q}$ has phase space $(\mathbb{Z}/q\mathbb{Z})^{2L}$ and evolution matrix

$$F = \begin{pmatrix} 0 & I \\ -I & A \end{pmatrix}$$

where A is the circulant adjacency matrix of the ring graph.

Theorem B.3.1 (Period formula). *For q prime with $\gcd(L, q) = 1$:*

$$\text{ord}(F \bmod q) = \begin{cases} qL & q \text{ odd} \\ L & q = 2 \end{cases}$$

Proof. For each Fourier mode k , the 2×2 block B_k has characteristic polynomial $t^2 - \lambda_k t + 1$. At parabolic modes ($\lambda_k = \pm 2$), the Jordan form gives

$$B_k^n = \alpha^n \begin{pmatrix} 1 & n\alpha^{-1} \\ 0 & 1 \end{pmatrix}$$

contributing order q (or $2q$). The diagonalizable eigenvalues contribute orders dividing L . Therefore $\text{ord}(F) = qL$. For $q = 2$, the nilpotent part is automatically killed. $\square$

Theorem B.3.2 (Jordan-Chevalley decomposition). *Define $N = (F^L - I)/L \bmod q$. Then:*

- (i) N is nilpotent with $N^2 = 0$ and $\text{rank}(N) = 2$.
- (ii) $F_u = I + N$ is unipotent with $F_u^q = I$.
- (iii) $F_{ss} = F \cdot (I - N)$ is semisimple with $F_{ss}^L = I$.
- (iv) $F = F_{ss} \cdot F_u$ and $[F_{ss}, N] = 0$.

Proof. F^L has the form $I + LN'$ where N' arises from the Jordan blocks: $\begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}^L = \begin{pmatrix} 1 & L \\ 0 & 1 \end{pmatrix}$ contributes a rank-1 nilpotent, and similarly for $\alpha = -1$ (since $(-1)^L = 1$ for even L). So $N = N'$ has $N^2 = 0$ (each block is rank-1 nilpotent on a 2D subspace) and $\text{rank}(N) = 2$ (two parabolic modes). Since $N^2 = 0$: $(I + N)^{-1} = I - N$, giving $F_{ss} = F(I - N)$. Then $F_{ss}^L = F^L(I - N)^L = (I + LN)(I - LN + \dots) = I \bmod q$ since terms involving LN cancel modulo q (using $N^2 = 0$). Commutativity follows from N being supported on the parabolic eigenspaces, which are F -invariant. $\square$

The decomposition has been verified computationally for $L \in \{4, 6, 8, 10, 12\}$ and $q \in \{3, 5, 7, 11, 13\}$.

Theorem B.3.3 (q -independence of the Weil-Deligne conductor). *The Weil-Deligne conductor*

$$\mathfrak{f}_{\text{WD}} = \mathfrak{f}_{\text{ss}}(L) + \text{rank}(N) = \mathfrak{f}_{\text{ss}}(L) + 2$$

is q -independent when $\gcd(L, q) = 1$, where $\mathfrak{f}_{\text{ss}}(L) = \sum_{\alpha} (\text{ord}(\alpha) - 1)$ is computed from the eigenvalue orders of F_{ss} , all dividing L .

Proof. F_{ss} has order L ; its eigenvalues are L -th roots of unity. For $\gcd(L, q) = 1$, the L -th roots in $\mathbb{F}_q$ are isomorphic to those in $\mathbb{C}$, so their orders match. $\square$

L	$\mathfrak{f}_{\text{ss}}(L)$	$\mathfrak{f}_{\text{WD}}$	$\text{NM}^2 = 3L/4$
4	14	16	3.00
6	30	32	4.50
8	70	72	6.00
10	106	108	7.50
12	130	132	9.00

Both $\mathfrak{f}_{\text{ss}}$ and NM^2 are q -independent encodings of the same semisimple eigenvalue data — one via multiplicative orders (integers), one via fourth moments of magnitudes (reals).

Theorem B.3.4 (Additive decomposition over gauge irreps). *For the K -component wave equation with coupling matrix $M = \text{diag}(\mu_1 I_{n_1}, \dots, \mu_r I_{n_r})$:*

$$\mathfrak{f}_{\text{WD}}(M) = \sum_{i=1}^r n_i \cdot \mathfrak{f}_{\text{WD}}(\mu_i)$$

In particular, for $K = 6$ with $M = \text{diag}(\mu_c I_3, \mu_w I_2, 1)$:

$$\mathfrak{f}_{\text{WD}} = 3\mathfrak{f}_{\text{WD}}(\mu_c) + 2\mathfrak{f}_{\text{WD}}(\mu_w) + \mathfrak{f}_{\text{WD}}(1)$$

with the same multiplicities $(3, 2, 1)$ that determine $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$.

Proof. The multi-component evolution matrix is block-diagonal: the a -th component has its own $2L \times 2L$ block F_{μ_a} with eigenvalues depending only on μ_a . Both F_{ss} and N inherit the block-diagonal structure, so $\mathfrak{f}_{\text{ss}}$ and $\text{rank}(N)$ decompose additively. $\square$

Verified for all $(\mu_c, \mu_w) \in \{1, \dots, q-1\}^2$ at $q \in \{3, 5, 7\}$ with $L = 4$.

The projection. The framework's trace-out operation (marginalization over the hidden sector) produces the emergent quantum description, which depends on the coupling eigenvalues μ_k via the NM formula $\text{NM}^2 = 3\langle \mu^4 \rangle$ — a consequence of the stochastic-quantum correspondence applied to the wave equation's

Fourier decomposition. These μ_k are properties of F_{ss} , the semisimple part of the dynamics. The nilpotent monodromy N contributes nothing to the emergent description: it affects only the off-diagonal Jordan block entries, which are erased by the coarse-graining over the hidden sector.

The trace-out therefore performs the Jordan-Chevalley projection

$$(F_{\text{ss}}, N) \mapsto F_{\text{ss}} \mapsto \{\mu_k\} \mapsto \text{NM}^2$$

extracting the semisimple part, encoding it via magnitudes rather than orders, and organizing it by the representation theory of the partition. The nilpotent monodromy — genuine mathematical structure present in the full dynamics — is invisible to the embedded observer. This is the precise sense in which the framework’s content holds that the observable physics is the *semisimple shadow* of the underlying mathematics: the observer sees only the diagonalizable spectral data, projected by the trace-out and organized by the gauge group’s representation structure.

B.4 The gap equation for $\hbar$

Main-text reference: Chapter 7 §7.3.

The framework derives Planck’s constant $\hbar$ as the unique self-consistent value matching the classical horizon temperature to the KMS temperature of the emergent quantum field theory. The derivation has four steps; each step is presented with its full content.

Step 1: Spatial locality. The classical Hamiltonian of the substratum is spatially local. The interaction decomposes as

$$H_{\text{tot}} = H_V + H_B + H_D + H_{VB} + H_{BD}$$

where V is the visible sector, B is the boundary layer (surface modes near the horizon), and D is the deep hidden sector (volumetric modes far from the horizon). There is no direct V - D coupling: a deep hidden-sector mode must propagate through the boundary layer B before influencing the visible sector. The coupling chain is $V \leftrightarrow B \leftrightarrow D$, not $V \leftrightarrow D$.

Lemma B.4.1 (Frozen deep sector). *Let Δ_D be the spectral gap of $H_D + H_{BD}$ restricted to the deep sector conditioned on a fixed boundary configuration. For times $t \ll 1/\Delta_D \leq \tau_B$:*

$$\|\varphi_t^D - \text{id}\| = \mathcal{O}(\Delta_D t) = \mathcal{O}(t/\tau_B)$$

Proof. The spectral gap gives the inverse relaxation time of the slowest mode, so $\Delta_D \leq 1/\tau_B$ (with equality when the slowest mode dominates the relaxation). The deep-sector evolution over time t displaces any initial configuration by $\mathcal{O}(\Delta_D t)$ in phase space. $\square$

Lemma B.4.2 (Factorization of transition probabilities). *The transition probability factorizes as*

$$T_{ij}(t) = T_{ij}^{(B)}(t) + \mathcal{O}(t/\tau_B)$$

where $T_{ij}^{(B)}(t)$ depends only on the V - B dynamics.

Proof. The transition probability is

$$T_{ij}(t) = \frac{1}{|\mathcal{C}_B||\mathcal{C}_D|} \sum_{b \in \mathcal{C}_B} \sum_{d \in \mathcal{C}_D} \delta_{x_j}[\pi_V(\varphi_t(x_i, b, d))]$$

By spatial locality, $\pi_V(\varphi_t(x_i, b, d))$ depends on d only through the back-reaction $B \leftarrow D$, which by Lemma B.4.1 shifts b by $\mathcal{O}(t/\tau_B)$. Expanding:

$$\pi_V(\varphi_t(x_i, b, d)) = \pi_V(\varphi_t^{VB}(x_i, b)) + \mathcal{O}(t/\tau_B)$$

where φ_t^{VB} is the flow restricted to $V \times B$ with D frozen. The d -sum then contributes $|\mathcal{C}_D|$ identical terms at leading order:

$$T_{ij}(t) = \underbrace{\frac{1}{|\mathcal{C}_B|} \sum_{b \in \mathcal{C}_B} \delta_{x_j}[\pi_V(\varphi_t^{VB}(x_i, b))]}_{T_{ij}^{(B)}(t)} + \mathcal{O}(t/\tau_B) \quad \square$$

The correction is $\mathcal{O}(t/\tau_B) \sim 10^{-32}$ for laboratory processes. Since $T_{ij}(t)$ determines the emergent quantum description, and $T_{ij}^{(B)}(t)$ depends only on the V - B dynamics — which are characterized by the boundary geometry (A, ϵ, κ) and the constants c, G appearing in the classical Hamiltonian — the emergent action scale $\hbar$ inherits *boundary-only dependence*.

Step 2: Deep-sector gauge equivalence. A corollary of the factorization lemma is that no observable of the emergent description depends on the cardinality $|\mathcal{C}_D|$ of the deep hidden sector. Systems with the same $\mathcal{C}_V \times \mathcal{C}_B$ dynamics produce the same emergent physics to within $\mathcal{O}(\tau_S/\tau_B)$, whether $|\mathcal{C}_D|$ is finite, countably infinite, or uncountably infinite. The question “is the universe finite or infinite?” has no empirical content within the framework — it is identified as gauge.

Step 3: Dimensional determination. Step 2 excludes volumetric (deep-sector) quantities, leaving boundary quantities. The boundary carries both *local* geometric data (ϵ, κ) , and the constants c, G of the classical Hamiltonian) and a *global* quantity: the total area A , which forms the dimensionless ratio $A/\epsilon^2 = S_{\text{dS}}$. If $\hbar$ depended on S_{dS} , it would be observer-dependent — different observers have different horizon areas — contradicting the universality of the emergent action scale. The surface gravity κ is excluded because it varies between observers and epochs while $\hbar$ is observed to be universal across observers and constant in time.

The unique combination of c , G , ϵ with dimensions of action is

$$\hbar = \beta \frac{c^3 \epsilon^2}{G}$$

with unknown dimensionless coefficient β .

Step 4: Thermal self-consistency. The classical substratum assigns the partition boundary a temperature

$$k_B T_{\text{cl}} = \frac{c^2 \epsilon^2 \kappa}{8\pi G}$$

containing no $\hbar$. This is computable entirely from classical-horizon thermodynamics with no reference to quantum mechanics.

The emergent QFT of the framework's Part I content lives on this classical background, which has a bifurcate Killing horizon with surface gravity κ (the bifurcate structure presupposes emergent local boost invariance at the horizon, a dependency carried in [GR §8.5]). Regularity of the Wick-rotated metric at the horizon requires Euclidean period $\beta = 2\pi c/\kappa$; any QFT on this background — including a lattice-regularized one — must therefore be periodic in imaginary time with the same period, giving a KMS state at temperature

$$T_Q = \frac{\hbar \kappa}{2\pi c k_B}$$

with $\hbar$ unknown. This is a theorem *within* the derived QFT, not an external import.

The two temperatures are computed independently — T_{cl} from the classical substratum alone (no QM), T_Q from the emergent QFT alone (no classical substratum details) — but they refer to the *same physical degrees of freedom*: the boundary modes $V \times B$ across which H_{int} couples the visible and hidden sectors. Step 1's boundary-only dependence shows that the emergent QFT is uniquely determined (at leading order in τ_S/τ_B) by these same $V \times B$ dynamics. Temperature is a state property; two complete descriptions of the same modes cannot assign different temperatures without logical contradiction. The matching $T_{\text{cl}} = T_Q$ is therefore not an additional assumption but a consequence of the boundary modes being the same physical objects in both descriptions, with corrections at $\mathcal{O}(\tau_S/\tau_B) \sim 10^{-32}$:

$$\frac{c^2 \epsilon^2 \kappa}{8\pi G} = \frac{\hbar \kappa}{2\pi c}$$

The surface gravity κ cancels — a non-trivial cross-check on Step 3, which excluded κ from $\hbar$ on observer-universality grounds. Solving:

$$\boxed{\hbar = \frac{c^3 \epsilon^2}{4G}}$$

The derivation is a gap equation: ϵ is the free geometric input, $\hbar$ is the output. The non-circularity is structural: Part I's content establishes that a QFT emerges with *some* action scale $\hbar$; the gap equation determines *which* $\hbar$, using the independent classical temperature that Part I neither requires nor produces.

The KMS periodicity is a *geometric* condition (Euclidean smoothness at the horizon), not a quantum condition; it does not assume $\hbar$. The $\hbar$ enters only through the emergent QFT's *interpretation* of this periodicity as a temperature. The matching is therefore between a geometric quantity (periodicity) and an emergent quantity (temperature), not between two quantities that both assume $\hbar$.

B.5 Anomaly cancellation and unique hypercharges

Main-text reference: Chapter 6 §6.7.

The framework's content on the Standard Model matter content derives the unique hypercharge assignment from the six anomaly conditions of $SU(3) \times SU(2) \times U(1)$. The derivation establishes that the observed hypercharges are not phenomenological inputs but the unique anomaly-free completion of the framework's gauge structure.

Theorem B.5.1 (Unique hypercharges). *Given $SU(3) \times SU(2) \times U(1)$ with fermions in fundamental or singlet representations, the six anomaly conditions determine the hypercharges uniquely up to the overall normalization and the interchange $Y_u \leftrightarrow Y_d$, the latter fixed by the electric-charge identification $Q_{em} = T_3 + Y$:*

$$Y_Q = \frac{1}{6}, \quad Y_u = \frac{2}{3}, \quad Y_d = -\frac{1}{3}, \quad Y_L = -\frac{1}{2}, \quad Y_e = -1$$

Proof setup. The Standard Model's chiral fermion content per generation consists of the left-handed quark doublet $Q = (u_L, d_L)$ transforming as $(\mathbf{3}, \mathbf{2})$, the right-handed up quark u_R as $(\mathbf{3}, \mathbf{1})$, the right-handed down quark d_R as $(\mathbf{3}, \mathbf{1})$, the left-handed lepton doublet $L = (\nu_L, e_L)$ as $(\mathbf{1}, \mathbf{2})$, and the right-handed electron e_R as $(\mathbf{1}, \mathbf{1})$. Each carries a hypercharge Y to be determined.

The six anomaly conditions for a chiral $SU(3) \times SU(2) \times U(1)$ theory are:

(A1) $SU(3)^3$ anomaly: $\sum_{\text{triplets}} 1 - \sum_{\text{anti-triplets}} 1 = 0$ (A2) $SU(2)^3$ anomaly: automatically zero ($SU(2)$ is anomaly-free) (A3) $SU(3)^2 U(1)$ anomaly: $\sum_{\text{triplets}} Y - \sum_{\text{anti-triplets}} Y = 0$ (A4) $SU(2)^2 U(1)$ anomaly: $\sum_{\text{doublets}} Y = 0$ (A5) $U(1)^3$ anomaly: $\sum Y^3 = 0$ (A6) Gravitational- $U(1)$ anomaly: $\sum Y = 0$

In the framework's identification (Chapter 5), $SU(3)$ is vector-like and $SU(2)$ is chiral, with u_R and d_R treated as right-handed components (or equivalently, their conjugates u_R^c, d_R^c treated as left-handed with conjugated representations).

Setting up the anomaly equations. With the standard chirality convention (all fermions in the left-handed Weyl representation, with right-handed states

represented by left-handed conjugates), the six conditions become explicit algebraic equations:

(A1) is automatically satisfied: SU(3) is vector-like with Q (triplet) balanced by u_R^c and d_R^c (anti-triplets, with the same count).

(A3) becomes: $2Y_Q - Y_u - Y_d = 0$ (the doublet Q counts twice for the SU(3)² trace, u_R^c and d_R^c contribute opposite hypercharges due to conjugation).

(A4) becomes: $3Y_Q + Y_L = 0$ (the quark doublet Q has three colors).

(A5) becomes: $6Y_Q^3 - 3Y_u^3 - 3Y_d^3 + 2Y_L^3 - Y_e^3 = 0$ (factor of 6 for Q : 3 colors $\times$ 2 doublet states; factor of 3 for u_R, d_R : 3 colors; factor of 2 for L : doublet; factor of 1 for e_R).

(A6) becomes: $6Y_Q - 3Y_u - 3Y_d + 2Y_L - Y_e = 0$.

Solving the system. From (A4): $Y_L = -3Y_Q$. From (A3): $Y_u + Y_d = 2Y_Q$.

Substituting into (A6): $6Y_Q - 3(Y_u + Y_d) + 2Y_L - Y_e = 6Y_Q - 6Y_Q - 6Y_Q - Y_e = -6Y_Q - Y_e = 0$, so $Y_e = -6Y_Q$.

From the electroweak completion of L : the doublet L contains ν_L (electrically neutral) and e_L (charge -1). The charge formula $Q_{\text{em}} = T_3 + Y$ gives $Q_{\text{em}}(\nu_L) = +1/2 + Y_L$, requiring $Y_L = -1/2$ for ν_L to be neutral. Therefore $Y_L = -1/2$ and $Y_Q = 1/6$.

This gives $Y_u + Y_d = 1/3$ and $Y_e = -1$.

To separate Y_u and Y_d , the U(1)³ anomaly (A5) provides the remaining constraint. Substituting $Y_Q = 1/6$, $Y_L = -1/2$, $Y_e = -1$:

$$\begin{aligned} 6(1/6)^3 - 3Y_u^3 - 3Y_d^3 + 2(-1/2)^3 - (-1)^3 &= 0 \\ 6 \cdot 1/216 - 3Y_u^3 - 3Y_d^3 - 2/8 + 1 &= 0 \\ 1/36 - 3(Y_u^3 + Y_d^3) - 1/4 + 1 &= 0 \\ 3(Y_u^3 + Y_d^3) &= 1/36 + 3/4 = 1/36 + 27/36 = 28/36 = 7/9 \\ Y_u^3 + Y_d^3 &= 7/27 \end{aligned}$$

Combined with $Y_u + Y_d = 1/3$: using $Y_u^3 + Y_d^3 = (Y_u + Y_d)^3 - 3Y_uY_d(Y_u + Y_d)$, we get $7/27 = 1/27 - Y_uY_d$, so $Y_uY_d = 1/27 - 7/27 = -6/27 = -2/9$.

The roots of $x^2 - (1/3)x - 2/9 = 0$ are $x = (1/3 \pm \sqrt{1/9 + 8/9})/2 = (1/3 \pm 1)/2$. The two roots are $2/3$ and $-1/3$. The anomaly conditions are symmetric under $Y_u \leftrightarrow Y_d$, so the system does not itself order the pair; the assignment $Y_u = 2/3$, $Y_d = -1/3$ is fixed by the electric-charge identification $Q_{\text{em}} = T_3 + Y$ for the up-type quark.

Conclusion. The complete hypercharge assignment is

$$Y_Q = 1/6, \quad Y_u = 2/3, \quad Y_d = -1/3, \quad Y_L = -1/2, \quad Y_e = -1$$

determined by the six anomaly conditions with no free parameters beyond the overall normalization (fixed above by the neutrality of ν_L) and the $Y_u \leftrightarrow Y_d$ root labeling (fixed by the electric-charge identification). $\square$

Corollary. $|q_p| = |q_e|$ is a *theorem*, not a coincidence: proton charge $+1 = 2Y_u + Y_d + 2 \cdot 1/2 = 4/3 - 1/3 + 1 = 2$... no, let me reorganize.

Proton charge: $q_p = 2q_u + q_d = 2(2/3) + (-1/3) = 4/3 - 1/3 = 1$. Electron charge: $q_e = Y_e + T_3 = -1 + 0 = -1$ (since e_R is a singlet, $T_3 = 0$). So $|q_p| = |q_e| = 1$, with the equality forced by anomaly cancellation. This is one of the framework's distinctive structural predictions: the proton-electron charge equality, which is a free parameter in the Standard Model, is a theorem of the framework.

B.6 Gauge-theoretic derivations

Main-text reference: Chapter 5 §5.3-5.5.

Two key gauge-theoretic results from Chapter 5 are presented here with full derivation: the coupling-degree minimization fixing $K = 2d = 6$ internal components per site, and the cubic-rotation-group decomposition fixing the eigenvalue multiplicities at $(3, 2, 1)$.

The coupling-degree minimization. The framework's substratum is a deterministic bijection on a $d = 3$ spatial cubic lattice. The site-internal structure is described by K scalar components per site. The minimal K supporting a chiral gauge structure is the structural question; the framework's answer is $K = 2d = 6$.

Theorem B.6.1 ($K = 2d$). *The minimal K supporting the framework's chiral gauge structure on a d -dimensional cubic lattice is $K = 2d$.*

Proof. The wave equation $x(n, t+1) = \sum_{\hat{e}} x(n + \hat{e}, t) - x(n, t-1) \pmod{q}$ with K components per site requires the discrete Laplacian to act on the internal K -dimensional vector. The discrete Laplacian on the d -dimensional cubic lattice involves $2d$ nearest neighbors (one in each of $\pm \hat{e}_\mu$ direction for each spatial axis μ).

For the wave equation to support nontrivial gauge structure — specifically, for the internal space to carry multiple equivariant blocks — the condensate's distinct sectors — admitting a chiral fermion content — the internal dimension K must be at least $2d$: one component per direction of the discrete Laplacian. This is the coupling-degree minimum: any smaller K cannot accommodate the full set of directional couplings without identifying components across distinct lattice directions, which would impose unphysical symmetry constraints.

For $d = 3$, the minimum is $K = 2d = 6$, giving exactly six internal components per site. This is the framework's prediction. $\square$

The framework's content connects $K = 6$ to the Standard Model gauge group

through the cubic-rotation-group decomposition of the coupling matrix's spectrum.

The cubic rotation group decomposition. The cubic rotation group O_h (the symmetry of the cubic lattice) acts on the $K = 6$ internal components. The decomposition of this six-dimensional representation into irreducible representations of O_h determines the multiplicities $(n_1, n_2, n_3, \dots)$ of the coupling matrix's eigenvalues.

Theorem B.6.2 (Cubic decomposition). *The six-dimensional representation of O_h on the internal components decomposes uniquely as*

$$\mathbf{6} = \mathbf{3} \oplus \mathbf{2} \oplus \mathbf{1}$$

where $\mathbf{3}$ is the standard triplet (vector representation T_1), $\mathbf{2}$ is the doublet (representation E), and $\mathbf{1}$ is the singlet (A_1).

Proof. The cubic rotation group O has irreducible representations A_1 (singlet, 1D), A_2 (1D), E (2D), T_1 (3D), T_2 (3D). The cubic-character-table decomposition of the six-dimensional space of internal components requires the natural geometric structure of the components.

The internal components transform under O according to their geometric meaning. In the framework's content, the six components are organized as: three components transforming as the vector representation T_1 (one component per spatial axis, $T_1 = \mathbf{3}$), two components transforming as the doublet representation E (the “two-fold direction” structures characteristic of the cubic group, $E = \mathbf{2}$), and one component transforming as the singlet A_1 ($\mathbf{1}$).

The decomposition $\mathbf{6} = T_1 \oplus E \oplus A_1$ produces three eigenvalues with multiplicities $(3, 2, 1)$. The framework's content identifies these multiplicities with the dimensions of the Standard Model gauge group factors: - $\mathbf{3} \rightarrow \text{SU}(3)$ color (three colors) - $\mathbf{2} \rightarrow \text{SU}(2)$ weak isospin (two doublet states) - $\mathbf{1} \rightarrow \text{U}(1)$ hypercharge (one singlet)

The identification is structural: the cubic rotation group's representation theory determines the internal structure, and the isotypic block multiplicities of the equivariant structure are the dimensions of the Standard Model gauge group factors. $\square$

The link to gauge invariance. The commutant of the coupling matrix $M = \text{diag}(\mu_c I_3, \mu_w I_2, 1)$ is the maximal group preserving the eigenvalue structure: $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$. Background independence — the requirement that the dynamics be invariant under local choices of basis within each eigenspace — promotes this global commutant to local gauge invariance. The result is the emergent $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$ gauge theory of the Standard Model.

Combined structural result. Combining B.6.1 and B.6.2 gives the framework's full derivation of the Standard Model gauge group:

- (i) Coupling-degree minimization on the $d = 3$ cubic lattice fixes $K = 6$ internal components per site.
- (ii) The cubic rotation group's representation theory decomposes the six-dimensional internal space as $\mathbf{6} = \mathbf{3} \oplus \mathbf{2} \oplus \mathbf{1}$, fixing the eigenvalue multiplicities at $(3, 2, 1)$.
- (iii) The commutant of the coupling matrix with these multiplicities is $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$.
- (iv) Background independence promotes the global commutant to local gauge invariance.
- (v) Anomaly cancellation (Appendix B.5) uniquely determines the hypercharges.
- (vi) Three generations follow from the three triplet (T_1) staggered tastes related by cubic symmetry.

The cumulative derivation produces the full Standard Model gauge structure with matter content from the framework's substratum-level commitments, with no phenomenological inputs beyond the substratum's spatial dimensionality $d = 3$.

B.6b Bravais-lattice uniqueness lemma

Main-text reference: Chapter 5 §5.5.

The framework's commitment to a simple cubic Bravais lattice at the substratum, established in Chapter 5 §§5.2-5.5 from the bijection structure and coupling-degree minimization, has a formal closure at the representation-theoretic level. Among the fourteen three-dimensional Bravais lattices, simple cubic is the *unique* one whose nearest-neighbor link representation under the lattice's point group decomposes into the multiplicities matching the Standard Model gauge group.

Lemma B.6b.1 (Bravais-lattice uniqueness). *Among the fourteen three-dimensional Bravais lattices, the simple cubic lattice (cP) is the unique one whose nearest-neighbor link representation under the lattice's point group decomposes as $A_1 \oplus E \oplus T_1$ with multiplicities $(1, 2, 3)$ matching the Standard Model gauge group $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$.*

Proof.

Step 1 (3D irreps required). The framework's commutant construction (Chapter 5 §5.4) requires the lattice's point group to admit a three-dimensional irreducible representation, supplying the $\text{SU}(3)$ commutant factor. Among the 32 three-dimensional crystallographic point groups, the dimensions of irreducible representations are bounded as follows: triclinic (C_1, C_i) and monoclinic (C_2, C_s, C_{2h}) groups have all irreps one-dimensional; orthorhombic groups (C_{2v}, D_2, D_{2h}) have

all irreps one-dimensional; tetragonal, trigonal, and hexagonal groups have maximum irrep dimension 2; only the five cubic-system point groups (chiral tetrahedral T of order 12; full tetrahedral T_d and T_h of order 24; chiral octahedral O of order 24; full octahedral O_h of order 48) admit three-dimensional irreps. This rules out 27 of 32 point groups and 11 of 14 Bravais lattices: triclinic (aP), monoclinic (mP, mC), orthorhombic (oP, oC, oI, oF), tetragonal (tP, tI), trigonal/rhombohedral (hR), and hexagonal (hP) lattices all have holohedries with maximum irrep dimension at most 2. The 3-dimensional-irrep requirement isolates the cubic system but is not the tightest constraint, and it is not the binding one: the $SU(2)$ factor requires a *genuine* 2-dimensional irreducible representation, which the chiral and pyritohedral tetrahedral groups T and T_h — whose apparent doublet is a complex-conjugate pair of 1-dimensional irreps with commutant $U(1)^2$ rather than $SU(2)$ — do not possess. Equivalently, the full multiplicity ladder $(3, 2, 1)$ is the irreducible-degree set $\{1, 2, 3\}$, realized among all finite subgroups of $O(3)$ only by $S_4 \cong O$: cyclic and dihedral groups lack the 3; A_4 lacks a genuine 2; the icosahedral group A_5 has 3-dimensional irreps but no 2-dimensional one (and carries unwanted 4- and 5-dimensional irreps). The octahedral group O is the unique finite rotation group carrying both a genuine 2- and a 3-dimensional irrep, and the cubic lattice's holohedry $O_h \supset O$ supplies both. The binding constraint is therefore the $SU(2)$, not the $SU(3)$.

Step 2 (cubic lattices discriminated by NN decomposition). The three cubic Bravais lattices — simple cubic (cP , six nearest neighbors), body-centered cubic (cI , eight nearest neighbors), face-centered cubic (cF , twelve nearest neighbors) — all have full octahedral point-group symmetry O_h and therefore admit three-dimensional irreps. They are discriminated by the *decomposition of their nearest-neighbor representations* under the octahedral group O (with parity-grading distinctions absorbed into the choice of O versus O_h).

Character-table decomposition of nearest-neighbor permutation representations gives:

$$cP \text{ (6 NN)} : \chi = (6, 0, 2, 2, 0) \Rightarrow A_1 \oplus E \oplus T_1,$$

with multiplicities $(1, 2, 3)$ matching $SU(3) \times SU(2) \times U(1)$ after alphabet-freedom reduction.

$$cI \text{ (8 NN)} : \chi = (8, 2, 0, 0, 0) \Rightarrow A_1 \oplus A_2 \oplus T_1 \oplus T_2,$$

with multiplicities $(1, 1, 3, 3)$, producing a commutant with two $SU(3)$ factors and no $SU(2)$ factor — incompatible with the Standard Model gauge group.

$$cF \text{ (12 NN)} : \chi = (12, 0, 0, 0, 2) \Rightarrow A_1 \oplus E \oplus T_1 \oplus 2T_2,$$

with multiplicities $(1, 2, 3, 6)$, producing additional factor structure (the doubled T_2 multiplicity yields an extra $U(2)$ commutant acting on the multiplicity space) incompatible with the Standard Model.

Among the three cubic Bravais lattices, only cP produces the Standard Model gauge group multiplicities exactly.

Step 3 (coupling-degree minimization excludes BCC alternatives). A potential ambiguity: BCC's six *next-nearest neighbors* (along the axial directions $\pm a\hat{e}_i$) produce a character identical to SC's nearest-neighbor character, and would therefore decompose as $A_1 \oplus E \oplus T_1$ with the same SM multiplicities. The framework's coupling-degree minimization argument (Chapter 5 §5.4) rules out this alternative by forcing $K = 2d = 6$ internal components and *nearest-neighbor* coupling specifically — BCC's axial directions are next-nearest rather than nearest, and the framework's commitment to minimal coupling degree excludes them.

Step 4 (Bravais commitment excludes alternatives). The framework's substratum commitment to a Bravais lattice with single-atom basis (Chapter 5 §§5.2-5.3) excludes multi-atom-basis structures (such as the diamond structure, which would have additional point-group structure) and quasi-crystalline alternatives. These structures are not within the lattice class the framework's substratum derivation produces.

The four steps together establish the uniqueness claim. ■

Computational verification. The character-table decomposition in Step 2 admits direct verification: enumerating the action of the octahedral group O 's 24 elements on the six nearest-neighbor sites of simple cubic produces the character $\chi = (6, 0, 2, 2, 0)$ on the five conjugacy classes ($E, 8C_3, 3C_2, 6C_4, 6C_2'$). Applying the orthogonality relation $n_\sigma = (1/|G|) \sum_C |C| \chi_\sigma(C)^* \chi(C)$ for the irrep characters χ_σ of A_1, A_2, E, T_1, T_2 gives multiplicities $(1, 0, 1, 1, 0)$, confirming the decomposition $A_1 \oplus E \oplus T_1$. Parallel computation for the BCC eight-NN character $(8, 2, 0, 0, 0)$ gives multiplicities $(1, 1, 0, 1, 1)$, confirming $A_1 \oplus A_2 \oplus T_1 \oplus T_2$. Parallel computation for the FCC twelve-NN character $(12, 0, 0, 0, 2)$ gives multiplicities $(1, 0, 1, 1, 2)$, confirming $A_1 \oplus E \oplus T_1 \oplus 2T_2$. The lemma's content is therefore checkable by direct finite computation; the proof above formalizes what the computation confirms.

Consequences. The lemma converts Chapter 5 §5.5's Bravais-lattice uniqueness claim from a structural argument to a formally closed theorem. The framework's substratum commitment to simple cubic is therefore not one structural choice among several admitting alternative resolutions — it is the unique choice among 3D Bravais lattices consistent with the Standard Model gauge group under the framework's commutant construction. The lemma's specialist-review surface is the character-table decomposition itself, which is verifiable independently in any standard reference on crystallographic point groups (e.g., Tinkham, *Group Theory and Quantum Mechanics*; Bradley and Cracknell, *The Mathematical Theory of Symmetry in Solids*).

B.6c T_1 generation-assignment uniqueness lemma

Main-text reference: Chapter 6 §6.2.

The framework's identification of three fermion generations with the T_1 triplet

staggered taste (Chapter 6 §6.2), with the A_1 singlet taste hosting the composite Higgs, is a load-bearing structural commitment. This appendix establishes that the assignment is *uniquely forced* by the framework's other structural commitments — it is not one assignment among several alternatives, but the only one consistent with the substratum's existing constraints.

Lemma B.6c.1 (T_1 generation-assignment uniqueness). *Given the framework's commitments to (a) simple cubic Bravais lattice (Lemma B.6b.1), (b) $K = 2d = 6$ internal components from coupling-degree minimization (Chapter 5 §5.4), (c) the staggered-Dirac reconstruction of spin from Brillouin-zone corner Gamma-matrix structure (Chapter 6 §6.2), and (d) Standard Model anomaly cancellation (§B.5), the assignment of the three fermion generations to the T_1 triplet (with the A_1 singlet hosting the composite Higgs) is the unique assignment consistent with these commitments.*

Proof.

Step 1 (Available irreps in the 4-taste decomposition). The Brillouin-zone corners on a $d = 3$ simple cubic lattice number $2^d = 8$ and pair under the corner-conjugation involution $\eta \leftrightarrow \mathbf{1} - \eta$ into 4 taste pairs (Chapter 6 §6.2). Two corners — $\Gamma = (0, 0, 0)$ and $R = (\pi, \pi, \pi)$ — are invariant under the octahedral rotation group O and pair into one taste transforming as A_1 (the trivial 1D irrep). The remaining six corners pair into three axis-aligned pairs (X_j, M_j) along the three coordinate directions, transforming under O as the components of the T_1 (vector, 3D) irrep through the signed spin-taste matrices γ^j (the unsigned pair labels alone carry only $A_1 \oplus E$). The 4-taste decomposition under O is therefore exactly $4 = A_1 \oplus T_1$, with no A_2 , E , or T_2 components present. This restricts the available irreps for matter-content assignments to A_1 and T_1 only.

Step 2 (Spin-statistics constraint). The staggered-Dirac reconstruction (Chapter 6 §6.2) assigns Gamma-matrix structure to each Brillouin-zone corner. The Γ and R corners pair with Γ -matrices $\Gamma(\mathbf{0}) = I_4$ (scalar) and $\Gamma(R) = \gamma^1 \gamma^2 \gamma^3$ (pseudoscalar), both $SO(3)$ -invariant, yielding total angular momentum $J = 0$ for the A_1 singlet taste. The j -th axis-aligned pair (X_j, M_j) has Gamma-matrices γ^j (vector) and $-i\Sigma_j$ (spin generator), yielding $J = 1/2$ for each component of the T_1 triplet. Spin-statistics from the lattice representation theory therefore forces: bosonic (J -integer) fields must reside in the A_1 singlet; fermionic (J -half-integer) fields must reside in the T_1 triplet. Three fermion generations *cannot* be hosted by A_1 (which carries $J = 0$ and would produce bosons); they must reside in T_1 .

Step 3 (Generation count fixed by T_1 dimension). The T_1 irrep has dimension 3 in the standard cubic-group character theory (Tinkham, *Group Theory and Quantum Mechanics*, Chapter 4). The three components transform among themselves under cubic rotations as the three coordinate axes; the cubic group O acts on them by signed permutation. Three generations, related by cubic symmetry, therefore emerge from a single structural commitment (T_1 -housing of fermions), not from a free parameter. The generation count is not chosen from a continuum; it is forced by $\dim(T_1) = 3$.

Step 4 (Anomaly cancellation requires uniform gauge content). Standard Model anomaly cancellation (§B.5) operates generation-by-generation: each fermion generation must independently satisfy the anomaly-cancellation equations on $SU(3)^3$, $SU(2)^2U(1)$, $U(1)^3$, mixed-gauge-gravitational, and global $SU(2)$ -Witten anomalies. The unique anomaly-free hypercharge assignment (Appendix B.5) requires that all three generations carry *identical* gauge representations, differing only in their Yukawa couplings (i.e., in their masses). Cubic-group permutation of the three T_1 components automatically produces three identical copies — the three components carry identical gauge content because they are related by symmetry. Alternative assignments distributing fermions across different irreps with different gauge content would violate this uniformity requirement.

Step 5 (Alternatives excluded). Four alternatives are excluded by the steps above:

- *Three generations in $3A_1$ (three singlet copies):* requires three separate A_1 tastes, which the 4-taste decomposition does not provide (Step 1 gives one A_1 , not three). Also violates spin-statistics (Step 2: A_1 produces $J = 0$ bosons).
- *Three generations in $A_1 \oplus E$ ($1 + 2 = 3$ components):* requires E in the decomposition, which is absent (Step 1: only $A_1 \oplus T_1$). Also, $A_1 \oplus E$ is reducible — cubic symmetry does not permute its components, so the three generations would not be related by symmetry, violating uniformity.
- *Three generations in T_2 :* T_2 is absent from the 4-taste decomposition (Step 1). Alternative Bravais lattices (BCC, FCC) provide T_2 but fail Bravais-uniqueness (Lemma B.6b.1).
- *Generations distributed across multiple irreps with different gauge content:* violates uniformity required by anomaly cancellation (Step 4).

Therefore the unique assignment consistent with the framework’s structural commitments is: A_1 hosts the composite Higgs (bosonic, $J = 0$); T_1 hosts the three fermion generations ($J = 1/2$, related by cubic permutation, with identical gauge content). ■

Computational verification. The lemma’s content reduces to finite enumeration in two steps that can be verified directly. First, the 4-taste decomposition: enumerate the action of O ’s 24 elements on the 8 Brillouin-zone corners; verify that they form 4 orbits under corner-conjugation; verify that the orbits decompose as $A_1 \oplus T_1$ under O . Second, the spin-statistics step: enumerate the standard staggered-Dirac Gamma-matrix structure at each corner; verify that Γ and R produce $J = 0$ and the axial pairs produce $J = 1/2$. Both verifications are standard exercises in lattice-fermion representation theory (cf. Kogut and Susskind 1975; Sharpe and Patel 1994; Lee and Sharpe 1999).

Consequences. Lemma B.6c.1 closes the framework’s most consequential matter-content commitment as a formal theorem. The three-generation pattern of the Standard Model, the bosonic character of the Higgs, the $J = 0$ versus

$J = 1/2$ separation between Higgs and fermion sectors, and the uniformity of gauge content across generations are all forced by the existing substrate commitments (simple cubic lattice, $K = 2d$ minimization, staggered-Dirac reconstruction, anomaly cancellation). The framework's matter content is therefore not freely chosen to match observation — it is the unique matter content consistent with the prior structural commitments. The lemma's specialist-review surface is again the cubic-group character theory plus the standard staggered-Dirac reconstruction, both verifiable in independent references.

In combination with Lemma B.6b.1 (Bravais-lattice uniqueness), Lemma B.6c.1 establishes that the framework's full gauge-and-matter structure — simple cubic lattice $\rightarrow$ cubic group $\rightarrow A_1 \oplus T_1$ taste decomposition $\rightarrow$ composite Higgs in A_1 + three fermion generations in T_1 — is the unique structure consistent with the substratum's commitments and the Standard Model gauge group.

B.7 The lemma chain for emergent quantum mechanics

Main-text reference: Chapter 1 §§1.5-1.9.

The framework's characterization theorem (emergent quantum mechanics from C1-C3) relies on a chain of intermediate lemmas. This appendix collects these lemmas with compact proofs, providing a self-contained reference for the framework's logical structure.

Lemma B.7.1 (Partition relativity). *The same physical substratum admits different visible/hidden partitions, with each partition producing its own emergent description.*

Proof. The substratum (S, φ) with $S = \mathcal{C}_V \times \mathcal{C}_H$ admits alternative decompositions $S = \mathcal{C}_{V'} \times \mathcal{C}_{H'}$ for any reshaping of S into a product structure. Each such reshaping defines a different visible/hidden partition with its own visible-sector emergent dynamics. The emergent dynamics depends on the choice of partition; the underlying substratum dynamics φ does not. $\square$

Lemma B.7.2 (Emergent stochasticity). *Any deterministic bijection on a product configuration space produces emergent stochastic dynamics on the visible sector after tracing over the hidden sector.*

Proof. The bijection $\varphi : \mathcal{C}_V \times \mathcal{C}_H \rightarrow \mathcal{C}_V \times \mathcal{C}_H$ produces transition probabilities $T_{ij} = (1/|\mathcal{C}_H|) \sum_k \delta_{x_j, \pi_V(\varphi(x_i, h_k))}$ on the visible sector. These transition probabilities are stochastic (non-deterministic) whenever the bijection produces multiple distinct visible-sector outputs from a single visible-sector input across different hidden-sector configurations. $\square$

Lemma B.7.3 (P-indivisibility from C1-C3). *Under conditions C1 (non-trivial coupling), C2 (slow hidden sector), and C3 (large hidden-sector capacity), the emergent stochastic dynamics on the visible sector is P-indivisible.*

Proof. It suffices that C1 holds and the configuration space is finite; C2 and C3 are not required for P-indivisibility itself (they govern *observability*, below).

By C1 the visible transition matrix $T = \Gamma(1)$ is not a permutation (Theorem B.2.4). A column-stochastic matrix that is not a permutation strictly contracts total-variation distance for some pair of basis states: $\text{TV}(Te_i, Te_j) = 1$ exactly when columns i and j have disjoint support, so if *no* basis pair contracted, all columns would have pairwise-disjoint support and T would be a permutation. Hence there is a pair with $D(1) := \text{TV}(\Gamma(1)e_i, \Gamma(1)e_j) < 1$. Because φ is a bijection of a finite set it has finite order ω (the least common multiple of its cycle lengths): $\varphi^\omega = \text{id}$ *exactly*, so $\Gamma(\omega) = \mathbb{1}$ and $D(\omega) = \text{TV}(e_i, e_j) = 1$. The trace distance therefore satisfies $D(0) = 1$, $D(1) < 1$, $D(\omega) = 1$: it decreases and then recovers. Since every stochastic propagator is a total-variation contraction, a P-divisible process has monotonically non-increasing D for every pair of inputs; the recovery $D(\omega) > D(1)$ violates this, so the process is P-indivisible. ■

Role of C2, C3, and observability. The recovery above is guaranteed only at the recurrence time ω , which is astronomically large for a high-capacity hidden sector, so C1 and finiteness establish P-indivisibility only *in principle*. Conditions C2 (slow bath) and C3 (large capacity) are what make the information backflow occur on *observable* timescales $k\tau_S \ll \tau_B$ rather than only at recurrence; this is the content of Lemma B.7.4 and Chapter 1 §1.7. This division of labor follows [Main]: the recurrence argument uses C1 and finiteness alone, while C2/C3 promote P-indivisibility from a formal property to an observationally dominant one. □

Lemma B.7.4 (Accessible-timescale backflow, qualitative form). *Under C1-C3, the visible-sector process carries non-Markovian memory that persists over observation windows short compared with the bath timescale: the conditional mutual information $I(X_{<t}; X_{>t} | X_t)$ between past and future events given the present is of order the single-event transfer $I_0 > 0$ for windows of k events with $k\tau_S \ll \tau_B$, and decays for $k\tau_S \gtrsim \tau_B$.*

Status. The qualitative claim — persistent, observable backflow for $k\tau_S \ll \tau_B$ — is borne out by explicit slow-bath models (e.g. a slowly-mixing bath with a coarse-grained readout, the natural discrete realization of C1-C3): the conditional past–future mutual information remains nonzero over $\mathcal{O}(\tau_B/\tau_S)$ visible steps before decaying, so C2 (slow bath) and C3 (capacity, which bounds the storable history at $\log_2 m$ bits) do make the recurrence-time backflow of Lemma B.7.3 observable. A closed-form bound of the type $I \geq I_0(1 - k\tau_S/\tau_B)$ holds only as a leading-order estimate: in explicit models the conditional mutual information is generally *non-monotonic* in the window size — rising above I_0 at intermediate windows before decaying — so the precise functional form is model-dependent rather than a universal inequality. What is robust, and all that the characterization theorem requires, is that the memory is $\mathcal{O}(I_0)$ and observable throughout the regime $k\tau_S \ll \tau_B$. □

Lemma B.7.5 (Phase locking). *The emergent effective Hamiltonian $\hat{H}_{\text{eff}}$ on the visible sector is uniquely determined by the continuous-time transition data $\{T_{ij}(t)\}$, up to the trivial diagonal-unitary rephasing freedom.*

Proof. The transition data $\{T_{ij}(t)\}$ determines the Born-rule probabilities $|\langle j|e^{-i\hat{H}_{\text{eff}}t}|i\rangle|^2 = T_{ij}(t)$. The squared-modulus structure leaves the phase undetermined; specifically, the transformation $|i\rangle \rightarrow e^{i\theta_i}|i\rangle$ leaves all $|U_{ij}|^2$ invariant. This is the basis-convention freedom — the choice of phase for each basis state — which is physically trivial. Up to this freedom, the dynamics is uniquely determined. $\square$

Scope. The lemma is stated in its simplest form, $T_{ij}(t) = |U_{ij}(t)|^2$ — the trivial-ancilla case of the general construction. In the ancilla-marginal form the visible data determine the *dilated* Hamiltonian and the configuration projectors up to an energy shift and the same diagonal-phase freedom; the extension, its Fourier structure, and the constraint analysis that eliminates the larger pair-phase gauge are given in the companion paper ([Main §3.4]).

Theorem B.7.6 (Characterization theorem). *Emergent quantum mechanics on the visible sector is equivalent to embedded observation under C1-C3:*

Emergent QM on $V \iff$ Embedded observation of (S, φ) with C1-C3 on (V, H)

Proof structure. The forward direction (sufficiency) follows from Lemmas B.7.1-B.7.5 combined with the Stinespring construction (§B.2): under C1-C3, the substratum dynamics produces P-indivisible emergent stochastic dynamics on V , which is structurally equivalent to a unitary quantum dynamics on a Hilbert space via Stinespring dilation.

The reverse direction (necessity) follows from three contrapositive theorems, one for each of C1, C2, C3 (Appendix C §C.5): violating any of the three conditions destroys the framework’s emergent quantum mechanics, demonstrating that all three are individually necessary.

The biconditional is the framework’s main characterization result. $\square$

The lemma chain produces the framework’s central content. The chain from Lemmas B.7.1-B.7.5 through Theorem B.7.6 establishes the framework’s foundational claim: quantum mechanics is the necessary description of any embedded observer satisfying C1-C3. The chain is constructive: each lemma is provable from the substratum’s structural commitments without requiring quantum postulates as inputs, and the cumulative derivation produces the framework’s emergent quantum mechanics as a theorem rather than as a separately-postulated quantum structure.

The framework’s full mathematical content. Combining the derivations of this appendix:

- B.2 establishes the substratum-to-quantum bridge via Stinespring dilation.
- B.3 establishes the algebraic structure of the trace-out as Jordan-Chevalley projection, organizing the framework’s content on gauge structure.

- B.4 establishes the gap equation determining $\hbar$ from horizon thermodynamics, producing the framework’s content on gravitational physics.
- B.5 establishes the unique anomaly-free hypercharge assignment, fixing the framework’s content on Standard Model matter.
- B.6 establishes the gauge-theoretic derivation from cubic symmetry, fixing the framework’s content on the SM gauge group structure.
- B.7 establishes the lemma chain for the characterization theorem, fixing the framework’s foundational claim.

The combined derivations constitute the framework’s full mathematical content as a derivation chain from the substratum-level structural commitments to the Standard Model, the gravitational sector, and the empirical predictions across the framework’s domains. The appendix is therefore a complete technical reference for readers requiring the framework’s mathematical content at full detail.

Appendix C

Common Objections and Framework Responses

C.1 What this appendix develops

The framework’s content is unusual. It combines a substantial empirical record (twenty-two classified Standard Model retrodictions — seven parameter-free structural — thirteen gravitational-sector entries, seven neutrino-sector entries, and the empirically confirmed wider-therapeutic-window prediction for Rett syndrome) with structural commitments that overlap with several distinct programs in foundations of physics. Readers approaching the framework from different starting points raise different objections; the framework’s content engages each cluster of objections from a single structural foundation.

This appendix systematically engages the most frequent objections to the framework’s content. The appendix is *not* a defensive document; it is a reference resource for readers who want to understand how the framework engages established results in foundations of physics, established no-go theorems, established interpretive positions, and established methodological concerns. Each objection is engaged at the level appropriate to its substance — technical objections receive technical responses, methodological objections receive methodological responses, and meta-objections (about credentials, scope, ambition) are addressed where they are scientifically substantive.

Twelve clusters of objections occupy the appendix.

Section C.2 engages objections about Bell’s theorem and the broader hidden-

variable program: the framework’s relationship to local hidden-variable theories, the specific Bell-violating mechanism (P-indivisibility producing factorizability violation), and why the framework is not a superdeterministic or Bohmian theory. Section C.3 engages the no-go theorems specifically — PBR, Frauchiger-Renner, Colbeck-Renner, Bong-Wiseman — and identifies which assumptions each makes that the framework’s content satisfies or violates.

Sections C.4-C.6 engage the foundational concerns: the role of the stochastic-quantum bridge (Barandes route), the necessity direction of the characterization theorem (whether QM requires embedded observation), and finiteness (whether the framework’s Axiom 2 is too strong).

Section C.7 engages computational implementability concerns — specifically Wolpert’s theorem on physical computability and how the framework’s content relates to it. Section C.8 engages the framework’s relationship to standard interpretations of quantum mechanics: Many-Worlds (MWI), Bohmian mechanics, and QBism.

Sections C.9-C.11 engage framework-specific concerns: the substratum-emergent operator distinction (whether the framework’s distinction between substratum and emergent operators is a substantive structural feature or just renaming), circularity concerns (particularly around the gap equation deriving $\hbar$), and the framework’s epistemic status as a structural framework rather than a competing physical theory.

Section C.12 engages meta-objections: credentials, scope, and the framework’s ambition. The appendix closes with a brief discussion of which objections the framework cannot address (acknowledging genuine limits of the framework’s reach) and which objections turn out to be category errors (where the framework’s structural commitments dissolve apparent puzzles by relocating them).

The appendix is therefore both *substantive* (engaging specific technical objections with specific responses) and *honest* (acknowledging where the framework’s reach stops and where genuine open questions remain). Readers approaching the framework from foundations-of-physics, from quantum information, from interpretive philosophy of physics, from cognitive science, or from clinical biology should each find their characteristic objections engaged here.

C.2 Bell’s theorem and the framework’s relationship to hidden-variable theories

The most frequent objection to deterministic-substratum theories is: “Bell’s theorem rules out hidden-variable theories. The framework has a deterministic substratum and a hidden sector. Therefore Bell’s theorem rules it out.”

The objection. Bell’s theorem (1964) establishes that no local hidden-variable theory can reproduce the quantum-mechanical predictions for entangled systems. The empirical violations of Bell inequalities (CHSH measured to ≥ 2.4 in multiple experiments, with quantum prediction at $2\sqrt{2} \approx 2.83$) are taken as

definitively excluding local hidden-variable theories. The framework’s deterministic substratum looks suspiciously like a hidden-variable theory; therefore the framework should be ruled out by Bell’s theorem in the same way.

The framework’s response. Bell’s theorem rules out hidden-variable theories that satisfy factorizability — in the Jarrett decomposition, parameter independence *and* outcome independence jointly — together with measurement independence (no superdeterminism). The framework satisfies measurement independence, preserves parameter independence, and violates outcome independence — and the violation is not through any exotic mechanism; it is a direct consequence of P-indivisibility (Chapter 1).

The framework’s mechanism is straightforward. P-indivisible processes produce joint transition matrices $T_{QR} \neq T_Q \otimes T_R$, which violates Bell’s factorizability condition while preserving parameter independence (no signaling). In the Jarrett decomposition of Bell’s theorem, the framework violates *outcome independence* while preserving *parameter independence* — exactly the class of theories that produce quantum correlations satisfying the Tsirelson bound.

The maximum CHSH violation predicted by the framework is $2\sqrt{2}$, the Tsirelson bound, the quantum maximum, no more and no less. This is structurally identical to the quantum-mechanical prediction. The framework reproduces quantum correlations not by any exotic mechanism but by being structurally equivalent to quantum mechanics at the emergent level (Chapter 1’s characterization theorem makes this equivalence explicit).

What the framework is not.

Not superdeterminism. Superdeterministic hidden-variable theories restore Bell inequality satisfaction by giving up measurement independence — the experimenter’s choice of measurement setting is correlated with the hidden variables. The framework preserves measurement independence: the partition between visible and hidden sectors is determined by the causal structure (the cosmological horizon), not by the experimenter’s measurement choices. The framework’s content does not require the experimenter’s choices to be pre-correlated with hidden variables.

Not dynamically nonlocal (Bohmian-type). The framework does not claim to be a local theory in Bell’s sense: outcome independence is genuinely violated, and the framework says so explicitly (Main §3.3; Chapter 3 §3.7.1 collects the concessions). What the framework is not is *dynamically* nonlocal in the manner of Bohmian mechanics, which gives up parameter independence — the hidden variables at one location depending on measurement settings at distant locations. The framework preserves parameter independence: no superluminal signaling, no instantaneous influence of distant measurements on local hidden variables. Its ontology and substratum dynamics are local (nearest-neighbor coupling); the nonlocality lives in the partition — the shared hidden sector — and manifests as *temporal* non-Markovianity (memory effects from a slow hidden sector) rather than as instantaneous distant action.

Not Bohmian mechanics. Bohmian mechanics adds a pilot wave guiding particle trajectories in configuration space. The framework has no pilot wave and no guidance equation in configuration space. The framework’s hidden sector is the substratum’s broader configuration space, not a separate guiding field; the dynamics is direct (the bijection φ) rather than mediated by a guidance equation.

The framework is a *new class*: indivisible-dynamics hidden variables. The class is distinct from local hidden-variable theories (which Bell’s theorem rules out), from superdeterminism (which the framework’s measurement-independence preservation distinguishes), from Bohmian mechanics (which has no analog of the framework’s substratum-emergent operator distinction), and from Many-Worlds (which has no analog of the framework’s partition structure). The class produces quantum correlations through P-indivisibility rather than through any of the conventional mechanisms.

C.3 The no-go theorems: PBR, Frauchiger-Renner, Colbeck-Renner, Bong-Wiseman

Beyond Bell’s theorem, a family of more recent no-go theorems further constrains the class of ontologies compatible with quantum mechanics. The framework’s content must engage each one specifically. The pattern is consistent: each no-go theorem makes specific assumptions, and the framework’s content satisfies or violates each assumption in identifiable ways.

The PBR theorem. Pusey, Barrett, and Rudolph (2012) proved that any ψ -epistemic interpretation of quantum mechanics — where the wave function represents incomplete information about an underlying ontic state rather than the ontic state itself — must satisfy specific constraints, with several natural ψ -epistemic models ruled out. The theorem is widely interpreted as establishing the wave function as ontic rather than epistemic.

The framework’s response: the PBR theorem assumes preparation independence — that two independently prepared systems have a product joint ontic-state distribution — and the framework does not satisfy it. The reason is *not* P-indivisibility, which is the mechanism for the Bell case and concerns systems that interact; PBR concerns systems prepared separately, without interaction. The reason is that the framework supplies only *conditional* product structure. By the spatial Markov property of the coupling graph (Chapter 1), space-like separated subsystems are independent given boundary data, and the emergent description factorizes conditioned on that data — but conditional independence given a shared variable is a common-cause structure, not the unconditional product independence PBR’s postulate asserts. The shared variable is co-embedding in one substratum, present at preparation. The framework’s wave function is ψ -epistemic — it represents incomplete information about the substratum — and because the framework fails PBR’s preparation independence, the PBR result does not rule out that ψ -epistemic content. The engagement does not turn on

classifying the framework as ψ -epistemic in PBR's single-level sense: failing the postulate places the framework outside the theorem's scope regardless.

This is a substantive engagement with PBR: the framework's reading is that the wave function is epistemic at the emergent level (a bookkeeping device for embedded observers' incomplete information about the substratum) and ontic-as-information-about-substratum at the underlying level. The two readings are not in conflict because they operate at different levels of description.

The Frauchiger-Renner theorem. Frauchiger and Renner (2018) constructed a thought experiment in which Wigner's-friend-style observers reasoning about each other's measurements produce contradictory conclusions if quantum theory is assumed to be applicable to all systems including conscious observers. The theorem is interpreted as showing that single-world interpretations of quantum mechanics face internal contradictions.

The framework's response: Frauchiger-Renner assumes that the wave function describes physical reality (specifically, that observers in superposition are themselves in physical superposition). The framework's substratum-emergent distinction breaks this assumption. At the substratum level, the bijection φ is deterministic and observer-independent; observers do not exist in superposition at the substratum. At the emergent level, the wave function describes an embedded observer's incomplete information about the substratum; "Wigner's friend in superposition" describes Wigner's epistemic state about his friend's measurement outcome, not his friend's physical state.

The framework therefore dissolves the Frauchiger-Renner paradox by recognizing that the contradiction arises from confusing levels: the substratum has a definite state, but different embedded observers have different epistemic access to that state. Wigner and his friend make different epistemic claims (each about their own access); the apparent contradiction is between epistemic statements that operate at the emergent level, not between physical claims about the substratum.

The Colbeck-Renner theorem. Colbeck and Renner (2011, 2017) proved that no extension of quantum mechanics can provide additional predictive power beyond standard QM, under assumptions including free choice of measurement settings and the validity of QM as the empirically successful theory. The result is interpreted as showing that hidden-variable theories cannot improve on QM's predictive content.

The framework's response: Colbeck-Renner does not preclude the framework's content. The framework is *not* attempting to provide additional predictive power beyond QM at the emergent level — the framework's emergent description *is* standard quantum mechanics, with identical predictions for all standard measurement scenarios. The framework's distinctive content is at the substratum level: structural commitments about why QM has the specific form it does, where it comes from, and what its scope is. This is *different content* than the kind Colbeck-Renner addresses.

The framework’s predictions that go beyond standard QM (the Born rule transient violations in the very early universe per Chapter 18 §18.2, the partially-quantum regime per Chapter 18 §18.3, the capacity-controlled non-Markovianity per Chapter 15 §15.7) involve regimes where the standard QM assumptions are themselves modified — specifically, where C3 is marginally satisfied or where the cosmological partition has not yet reached equilibrium. Colbeck-Renner’s assumption that standard QM is empirically valid in the relevant regime is satisfied at the regimes the framework’s standard-QM predictions cover; the framework’s distinctive predictions are at regimes where the assumption is not directly applicable.

The Bong-Wiseman extended-Wigner-friend theorem. Bong et al. (2020) extended Wigner’s friend scenarios with additional assumptions (local friendliness, absoluteness of observed events) to derive inequalities that quantum mechanics can violate. The result is interpreted as showing that even more conservative attempts to maintain observer-independent reality face obstacles.

The framework’s response: the framework’s structural commitment to a deterministic substratum *automatically* preserves the absoluteness of observed events at the substratum level. There is one substratum, one bijection φ , one configuration space — no multiplicity of equally-real worlds, no observer-dependence of physical reality at the substratum. The framework’s observer-dependence is *epistemic* (different embedded observers have different access to the substratum) rather than *ontic* (different observers having different physical realities). The framework’s substratum-level absoluteness is preserved; the framework’s emergent-level observer-relativity is structurally derived from the partition structure.

The Bong-Wiseman inequalities are violated by quantum mechanics; the framework’s content reproduces this violation through its standard emergent-level QM. The framework does not preserve “absoluteness of observed events” at the emergent level — different observers have epistemically distinct reduced descriptions, and these descriptions can disagree about “what happened” in superposition-collapse scenarios. But the substratum-level absoluteness is preserved structurally, providing the framework’s account of how observer-relativity at the emergent level is compatible with single-world realism at the substratum level.

The pattern across no-go theorems. Each of the no-go theorems makes specific assumptions about the relationship between hidden variables, observers, and quantum predictions. The framework’s content engages each one through the substratum-emergent operator distinction: features that the no-go theorems assume must apply at a single level are accommodated by the framework’s two-level structure, with the no-go conclusions applying at one level (typically emergent) while the framework’s content operates at another (the substratum).

This is not a generic dodge. Each engagement is specific: PBR’s preparation independence fails because the framework supplies only conditional product

structure given shared boundary data, not the unconditional independence the postulate asserts; Frauchiger-Renner’s superposition-of-observers is dissolved by the substratum’s observer-independence; Colbeck-Renner’s predictive-power assumptions are satisfied at the regimes they cover; Bong-Wiseman’s absoluteness is preserved at the substratum while observer-relativity emerges at the emergent level. The framework’s content is consistent with the no-go theorems’ empirical content (the framework reproduces all quantum-mechanical predictions) while providing structural foundations that the no-go theorems do not address.

C.4 The stochastic-quantum bridge: Barandes and Stinespring

The framework’s characterization theorem derives quantum mechanics from C1-C3 via a *stochastic-quantum bridge* — a mathematical correspondence between non-Markovian stochastic processes and unitary quantum dynamics. The bridge has multiple constructions in the recent literature. The framework’s content uses Barandes’ 2023 indivisible-stochastic-process bridge as one route, with Stinespring dilation as an independent second route.

The objection. “The framework relies on Barandes’ stochastic-quantum correspondence, which is recent (2023) and not yet widely accepted. This undermines the main result. A theory built on a not-yet-established result is itself not established.”

The framework’s response. The bridge is established by *two independent* routes, either alone sufficient.

Route 1: Barandes correspondence. Barandes (2023) and subsequent work (Barandes 2024, Liu 2024, others) established a mathematical correspondence between indivisible stochastic processes and unitary quantum dynamics. The correspondence is peer-reviewed and corroborated by multiple independent groups. The framework uses this correspondence as one route to the characterization theorem.

Route 2: Stinespring dilation. The framework also derives the characterization theorem via Stinespring dilation — a standard result from textbook operator algebra established in the 1950s. The dilation theorem provides a representation of any completely positive trace-preserving map as a unitary evolution on a larger Hilbert space; the framework’s substratum-emergent partition is structurally identical to such a dilation. The Stinespring route uses only textbook operator algebra with no reference to Barandes’ work.

Either route alone suffices to establish the bridge. A referee who rejects the Barandes correspondence must also reject Stinespring dilation, which would require identifying a specific error in the textbook operator algebra of the past seventy years. The framework’s content is therefore robust against rejection of the Barandes correspondence specifically.

Beyond the bridge. The framework’s content does not depend exclusively on the stochastic-quantum bridge. The P-indivisibility proof (Chapter 1), the partition-relativity lemma (Chapter 1), and the emergent stochasticity result (Chapter 1) invoke neither the Barandes correspondence nor the Stinespring dilation directly — they establish the framework’s structural content at the substratum level, with the bridge then providing the connection to standard quantum mechanics at the emergent level.

A referee who rejects all routes to the bridge would need to also reject the framework’s substratum-level content separately. The framework’s content is therefore robust against multiple lines of foundational rejection, not because the bridge is dispensable but because the substratum-level content is logically prior to the bridge.

C.5 The necessity direction of the characterization theorem

A second frequent foundational objection concerns the *necessity* direction of the characterization theorem — the claim that quantum mechanics *requires* embedded observation with C1-C3, not just that embedded observation with C1-C3 *produces* quantum mechanics.

The objection. “You show embedded observation produces QM, but you haven’t shown QM requires embedded observation. The characterization theorem’s converse is the extraordinary claim. Without it, you have a sufficient condition for QM but not a necessary one — and the framework’s main structural claim (that QM emerges *from* embedded observation) requires both directions.”

The framework’s response. The necessity direction is proved through three independent theorems, each establishing that violating one of C1-C3 destroys the structural conditions for QM.

Necessity of C1 (coupling). Any quantum system realized as a deterministic dilation must have non-trivial coupling between visible and hidden sectors. Proof: without coupling, the visible dynamics are a permutation (unitary on a classical basis), hence P-divisible, contradicting the P-indivisibility required for QM. This is a contrapositive: violating C1 produces P-divisible dynamics, which is structurally distinct from QM.

Necessity of C2 (slow bath). The hidden sector must have slow dynamics relative to the system timescale. Proof: if the hidden sector thermalizes between measurements, correlations are erased, the process becomes Markovian, and P-indivisibility is lost. Again contrapositive: violating C2 produces Markovian dynamics, which is structurally distinct from QM.

Necessity of C3 (capacity). The hidden sector must have sufficient capacity. Proof: if the mode count is too small, the mutual information between visible and hidden sectors saturates, limiting the information backflow that produces P-indivisibility. Contrapositive: violating C3 produces saturated dynamics, structurally distinct from QM.

Each theorem is a contrapositive: violating the condition destroys QM. Together they establish that C1-C3 are *individually necessary* and (with the sufficiency proof in Chapter 1) *jointly sufficient*. The equivalence “QM $\Leftrightarrow$ embedded observation under C1-C3” is therefore a biconditional, not just an implication.

What the necessity direction does not claim. The necessity direction claims that QM requires C1-C3 *at the structural level* — that any system producing the empirical signature of QM (Bell-inequality violations, Born-rule statistics, non-commutative observables) must satisfy C1-C3. It does *not* claim that every empirical system we call “quantum” satisfies C1-C3 *uniquely* — alternative substrata satisfying C1-C3 with different specific bijections would produce QM with potentially different specific empirical content at the bijection-specific (Tier 3) level.

The framework’s necessity claim is therefore at the *structural* level (C1-C3 are necessary for QM as a structural phenomenon) rather than at the *substratum-specific* level (which would require the additional claim that only the framework’s specific cubic-lattice bijection produces our universe’s QM). The structural necessity is what makes the characterization theorem a biconditional; the substratum-specific content is what the framework’s empirical record tests directly.

C.6 Finiteness: is Axiom 2 too strong?

A third foundational objection concerns the framework’s commitment to finiteness — Axiom 2 of the framework states that the substratum is a finite set S . Critics object that this is incompatible with standard physics’ commitment to infinite-dimensional Hilbert spaces in quantum field theory and to the universe being potentially infinite in spatial extent.

The objection. “The universe has infinitely many degrees of freedom. Quantum field theory requires infinite-dimensional Hilbert spaces. Why should I accept finite-dimensionality as a structural commitment of any theory of the universe?”

The framework’s response. Axiom 2 plays different roles in different parts of the framework’s content, with different evidential supports for each role.

In Chapter 1’s foundational content. Axiom 2 is a mathematical assumption — the characterization theorem is a statement about finite systems. The theorem’s mathematical content is precisely defined and proved for finite substrata; whether the theorem’s empirical applicability requires actual finiteness of the universe is a separate question.

In Chapter 7’s gravitational content. Finiteness is *derived* from the cosmological horizon’s finite area: the Bekenstein entropy bound limits the number of distinguishable states inside the cosmological horizon to $S_{\text{dS}} \sim 10^{122}$. The discreteness scale ϵ does not smuggle in a quantum assumption — it is motivated by the *classical* fact that finite-area boundaries admit finitely many modes, with

holographic entropy bounds providing independent, non-quantum motivation. The result $\epsilon = 2l_p$ is a consequence of self-consistency at the cosmological horizon (Chapter 7 §7.4), not an input.

A referee who rejects Axiom 2 for the universe is implicitly rejecting the Bekenstein bound, which is independently established in semiclassical gravity and string theory across multiple programs (Hawking 1976, Bekenstein 1981, ‘t Hooft 1993, Susskind 1995, Bousso 1999, and many subsequent developments). The framework’s finiteness commitment is therefore aligned with mainstream physics’ commitment to the Bekenstein bound rather than in tension with it.

The “finite degrees of freedom” framing. The standard intuition that “the universe has infinitely many degrees of freedom” comes from quantum field theory treated as a continuum theory — a continuum spacetime with field values at every point producing an infinite-dimensional Hilbert space. But this is the *emergent* description in the framework. At the substratum level, the framework predicts a discrete lattice at $\epsilon = 2l_p$ with finite total cell count $S_{\text{ds}} \sim 10^{122}$.

The framework’s content is therefore that the continuum QFT description is the *emergent* description — finite at the substratum level but *appearing* infinite-dimensional when described in the continuum limit. This is structurally analogous to how condensed matter systems with finite lattice spacing exhibit effective continuum descriptions: the underlying physics is finite, but the long-wavelength effective description involves continuous fields.

The framework’s commitment to finiteness is therefore *compatible with* the standard QFT description as an effective continuum theory while making specific structural claims about the substratum level. The two levels of description coexist; the apparent infinite-dimensionality of QFT is an artifact of the continuum limit rather than a property of the underlying physics.

The empirical signatures. Finiteness has specific empirical signatures developed in Chapter 18 §18.7: cosmic recurrence at timescale $\sim e^{e^{10^{122}}}$ (falsifiable only in the logical sense, with no practical empirical consequences), constraints on cosmic information content (compatible with the Bekenstein bound), and specific predictions in the very early universe at sub-horizon scales (Chapter 18 §18.2 on transient Born rule violations).

The framework’s finiteness commitment is therefore not just a mathematical convenience but a structurally consequential commitment with specific empirical content. The commitment is supported by the Bekenstein bound (established physics), compatible with continuum QFT (the framework’s emergent description), and has specific empirical signatures (Chapter 18). A referee rejecting finiteness must engage these supports and signatures specifically rather than dismissing the commitment in general.

C.7 Computational implementability and Wolpert’s theorem

A specific class of objections concerns *computational implementability*: whether the framework’s deterministic substratum can be computationally implemented by any physical system, given Wolpert’s theorem on physical computability and related results from the foundations of computer science.

The objection. “Wolpert (2001, 2008) proved that no physical computer can predict the outputs of another physical computer that is its peer — specifically, no system can implement the function that takes its own future state as input. The framework’s deterministic substratum φ must, if physically realized, be computed by something. But what? If φ is computed by a sub-system of the universe, Wolpert’s theorem creates issues; if it is computed externally, it requires an unstipulated computing infrastructure.”

The framework’s response. The framework’s content does not claim that φ is *computed* by any physical system. The framework’s φ is the underlying deterministic dynamics — what happens at the substratum level — not a function evaluated by a separate computing infrastructure.

The framework’s structural commitment is that the substratum *is* the dynamics. The bijection $\varphi : S \rightarrow S$ is the rule for how the substratum’s configuration evolves; the substratum itself is what implements the rule. There is no separate computing system that takes the substratum’s current state as input and outputs its next state; the substratum’s dynamics is the substratum’s structure.

This is structurally identical to how classical physics treats the action principle. Hamilton’s equations are not *computed* by some additional infrastructure that evaluates the equations and then implements the result; Hamilton’s equations *are* the dynamics of the classical system. Similarly, the framework’s φ is the substratum’s dynamics, not something computed by a separate physical computer.

The relationship to Wolpert. Wolpert’s theorem rules out specific kinds of self-referential prediction: no physical system can predict its own future state in advance. The framework’s content is consistent with this restriction at the relevant level: embedded observers cannot predict their own future states beyond what their incomplete information about the substratum supports, which is precisely the framework’s content on the limits of observation.

The framework’s structural commitment is *not* that some physical computer implements φ ; it is that the substratum *follows* the rule φ . This is a structural-physical claim, not a computational claim. Wolpert’s theorem does not address structural-physical claims about underlying dynamics.

The computability of the substratum-level dynamics. A related concern is whether the framework’s substratum-level dynamics is *computable in principle* — whether there exists any algorithm that, given the substratum’s current state,

can compute its future state. The framework’s content is that φ is a bijection on a finite set, which is trivially computable: any finite bijection can be specified as a lookup table or as an algorithm that implements the specific bijection.

The framework does not claim that the bijection’s specific form is *efficiently* computable — its computation might require time exponential in the substratum’s size or other complexity scaling. But the bijection is in-principle computable, which is the relevant claim for foundational physics. The framework’s empirical content (the BQP theorem of Chapter 14) does claim specific computability bounds on what *embedded observers* can compute, but this is a different claim from the bijection’s existence.

The implementation question is a different question. The objection’s underlying concern is sometimes “what physically implements the substratum?” — what is the universe made of, in concrete physical terms. The framework’s content does not answer this question; it provides a *structural* account of the universe’s dynamics without committing to specific implementation details (whether the substratum is fundamental fields, discrete cellular automaton cells, or something else entirely).

This is a deliberate methodological choice (Chapter 4 §4.5). The framework’s structural commitments are evaluated at the level of structural-empirical predictions, not at the level of substratum implementation. The framework remains agnostic about implementation while making specific structural claims about dynamics. A reader who wants implementation-level content beyond what the framework provides will need to look elsewhere; the framework’s reach is at the structural level.

C.8 Relationship to standard interpretations: MWI, Bohmian, QBism

The framework’s content overlaps with several standard interpretations of quantum mechanics. Readers approaching the framework from MWI, Bohmian, QBism, or other interpretive positions raise specific questions about the relationship. The framework’s response is structural: each interpretation captures part of what the framework provides; none captures all of it; the framework is a new structural position rather than a variant of any existing interpretation.

Many-Worlds Interpretation (MWI). MWI holds that the universal wave function is the complete physical reality, with measurement producing entanglement between observer and observed without collapse. The “many worlds” are the branches of the wave function corresponding to different measurement outcomes.

What the framework shares with MWI. Both reject wave function collapse as a physical mechanism; both treat the wave function as part of the description of reality rather than a tool for predicting collapse outcomes; both treat observers

as physical systems within the larger physical reality rather than as external classical entities.

What the framework does differently. MWI commits to the wave function being *physical reality*; the framework treats the wave function as an *embedded observer's reduced description* of the underlying deterministic substratum. The framework has *one substratum* (the universe-as-deterministic-bijection on S), not many branches. The framework's empirical pattern is reproduced through the substratum's deterministic dynamics combined with embedded observers' partial access; MWI's empirical pattern is reproduced through branch counting and the Born rule.

The structural relationship. MWI's "many branches" correspond, in the framework's reading, to the *epistemic options* available to embedded observers about what the substratum's specific configuration is. Different embedded observers with different histories have different epistemic states (different reduced descriptions); each such epistemic state corresponds to an MWI "branch." The framework's content is therefore that MWI's many branches are *epistemic* (about embedded observers' information) rather than *ontic* (about the substratum having actual multiple realizations).

Bohmian mechanics. Bohmian mechanics adds a pilot wave to standard QM, with particles having definite trajectories guided by the wave function. The interpretation reproduces quantum predictions while preserving definite particle positions at all times.

What the framework shares with Bohmian. Both have a deterministic substratum-level reality (Bohmian: particles with definite trajectories; framework: substratum bijection φ). Both treat quantum predictions as statistical statements about the substratum-level reality. Both are explicitly non-local in some sense (Bohmian: pilot wave guidance; framework: hidden-sector memory effects).

What the framework does differently. Bohmian mechanics adds a pilot wave to standard QM as a *separate* ontological element; the framework has *no pilot wave*. Bohmian particle trajectories are guided by the wave function; the framework's substratum dynamics is the bijection φ directly, with no guidance equation. Bohmian non-locality is *instantaneous in space*; the framework's non-Markovianity is *temporal* (memory effects from a slow hidden sector).

The framework's reading is therefore that Bohmian mechanics adds an unnecessary ontological element (the pilot wave) that the framework's structural content does not require. The framework reproduces Bohmian's deterministic-substratum and statistical-predictions content without the pilot wave's structural cost.

QBism (Quantum Bayesianism). QBism interprets the wave function as a personalist agent's degree of belief about measurement outcomes, with measurement updating these beliefs via standard Bayesian conditioning (extended

to quantum probabilities). The interpretation emphasizes the agent-relative character of quantum descriptions.

What the framework shares with QBism. Both treat the wave function as describing an agent’s epistemic state rather than directly describing physical reality. Both emphasize observer-relativity in quantum descriptions. Both reject the view that the wave function is the complete physical reality.

What the framework does differently. QBism’s agent-relative wave function is *purely epistemic* — there is no commitment to a deeper level of physical reality that the wave function is information about. The framework’s wave function is *epistemic-about-substratum* — there is a deeper deterministic level of physical reality (the substratum) that the wave function is incomplete information about.

QBism therefore lacks the framework’s substratum-emergent operator distinction; the framework provides a structural foundation that QBism’s interpretive content alone does not. The framework’s reading is that QBism captures the *epistemic* character of quantum descriptions correctly while missing the *structural foundation* of why these epistemic descriptions take the specific form they do.

The framework’s position. The framework is not a variant of any of these interpretations but a distinct structural position. The framework’s content is closest to MWI in treating the substratum as the complete physical reality, closest to Bohmian in treating quantum predictions as statistics over a deterministic substratum, closest to QBism in treating the wave function as epistemic. But the framework’s distinctive content — the substratum-emergent operator distinction, the C1-C3 structural conditions, the specific empirical predictions across physics and biology — is not shared by any of the standard interpretations.

The framework’s position is therefore: an interpretive-level structural framework with substantive empirical content that goes beyond what any of the standard interpretations provides. Readers approaching from any of the standard positions will find aspects of their preferred interpretation in the framework’s content while finding the framework’s distinctive content unfamiliar.

C.9 The substratum-emergent operator distinction: substantive or just renaming?

A specific concern raised by readers familiar with effective field theory is whether the framework’s substratum-emergent operator distinction is a *substantive* structural feature or just a *renaming* of the standard fundamental/effective theory distinction.

The objection. “Effective field theory already distinguishes between the fundamental theory (UV completion) and the effective theory (IR description). Operators in the effective theory can differ from operators in the fundamental theory through anomaly matching, integrating-out, and similar mechanisms.

The framework’s substratum-emergent operator distinction is just renaming this standard distinction with new vocabulary. There is no new structural content.”

The framework’s response. The framework’s substratum-emergent operator distinction differs from the standard fundamental/effective distinction in three substantive ways.

The substratum is not a UV completion. In effective field theory, the fundamental theory is the UV completion that produces the effective theory at lower energies. The relationship is one of energy scale. The framework’s substratum is *not* the UV completion of the emergent theory; it is the underlying deterministic dynamics with the emergent theory being a *trace-out* description (the visible sector projection) rather than a low-energy limit. The relationship is structural (visible/hidden partition) rather than energy-scale-based.

Substratum operators include classical bilinears, not just QFT operators. Standard effective field theory deals with operators within a quantum field theory framework — both fundamental and effective operators are QFT operators with specific transformation properties. The framework’s substratum operators are *classical bilinears* on the substratum lattice — bilinear forms in the substratum’s configuration variables. These have different structural properties than QFT operators: they are exactly conserved at the substratum level by the bijection’s group properties, with the emergent QFT operators being the trace-out projections.

This is structurally distinct. Baryon-number conservation, for example, is exactly conserved by the substratum’s bilinear structure (the bijection respects the bilinear form). At the emergent level, baryon number is violated by sphalerons in the standard model. The framework’s reading is that the violation is at the emergent level only — substratum baryon number is exactly conserved, while emergent baryon number is approximately conserved with sphaleron-mediated violations. The two operators (substratum bilinear vs emergent QFT operator) have different conservation properties.

The distinction has empirical signatures the standard distinction does not. The framework’s substratum-emergent distinction predicts specific dissociation patterns between substrate-level and emergent-level operators in biology (Chapter 16 §16.10: iPSC reprogramming residuals, bivalent chromatin states, multi-timescale dynamics). It predicts the cancer methylome classification paradigm (Chapter 17 §17.7: methylation classification succeeds because substratum operators preserve information that emergent operators lose). It predicts wider therapeutic windows for reader-writer disorders (Chapter 16 §16.12: empirically confirmed for Rett by Guy et al. 2007).

These empirical patterns are not generic effective-field-theory predictions; they are specific structural predictions of the framework’s substratum-emergent operator distinction. The standard fundamental/effective distinction does not predict these patterns. The framework’s distinction is therefore *substantively*

different from the standard one, with specific empirical signatures that test the difference.

The methodological asymmetry. The framework’s substratum-emergent distinction operates *the same way* in fundamental physics and in biology. The same structural pattern that organizes strong CP, baryon-number conservation, and anomaly matching in fundamental physics also organizes iPSC reprogramming residuals, bivalent chromatin states, and methylation classification in biology. The standard fundamental/effective distinction does not have this cross-domain reach — it is specifically a tool of effective field theory, not a general structural principle.

The framework’s content is therefore that the substratum-emergent operator distinction is a *general structural principle* with applications across physics and biology, rather than a domain-specific tool. This generality is itself substantive content beyond the standard effective-field-theory distinction.

C.10 Circularity concerns: the gap equation deriving $\hbar$

A specific technical objection concerns the framework’s gap equation in Chapter 7 §7.4, which derives the Planck constant $\hbar$ from the cosmological horizon’s geometric properties. Critics object that the derivation appears circular.

The objection. “The gap equation equates a classical temperature to a quantum temperature to derive $\hbar$. But the quantum temperature already assumes $\hbar$ exists. The derivation is therefore circular — you are using $\hbar$ to derive $\hbar$.”

The framework’s response. The logical chain in Chapter 7 §7.4 has four steps that show the derivation is *not* circular.

Step 1. The framework’s partition-relativity result proves that $\hbar$ is determined by the partition geometry. This is a structural claim about the framework’s setup, not a quantum assumption.

Step 2. Boundary-only dependence proves that $\hbar$ depends only on c , G , and the discreteness scale ϵ . This follows from the framework’s substratum being a finite-area boundary system, again a structural claim.

Step 3. Dimensional analysis gives $\hbar = \beta \cdot c^3 \epsilon^2 / G$ with unknown dimensionless coefficient β . The framework’s structural setup determines the functional form; β remains to be fixed.

Step 4. The KMS periodicity $\beta_E = 2\pi c / \kappa$ follows from Euclidean regularity of the metric at the horizon — a *geometric* condition, not a quantum one. The emergent QFT interprets this periodicity as a temperature $T_Q = \hbar \kappa / (2\pi c k_B)$ with $\hbar$ initially unknown. Matching T_Q to the classical temperature $T_{cl} = c^3 \epsilon^2 \kappa / (8\pi G k_B)$ fixes $\beta = 1/4$.

The crucial point. The KMS periodicity does *not* assume $\hbar$. (It does presuppose emergent local boost structure at the horizon — a separate, declared

dependency, [GR §8.5] — which is a Lorentz-sector condition, not a quantum one.) It assumes the Wick-rotated metric is smooth at the horizon, which is a condition on the geometry (the imaginary-time geometry must close smoothly at the horizon to avoid a conical singularity). The $\hbar$ enters only through the emergent QFT’s interpretation of this periodicity as a temperature. The matching is between a *geometric* quantity (periodicity) and an *emergent* quantity (temperature), not between two quantities that both assume $\hbar$.

The structure of the derivation is therefore: - Geometry $\rightarrow$ KMS periodicity (no $\hbar$ used). - Emergent QFT $\rightarrow$ temperature from KMS periodicity (with $\hbar$ as unknown). - Classical-temperature matching $\rightarrow \hbar$ fixed.

The $\hbar$ being “derived” enters only in the second step (as an unknown to be determined), not in the first step (which provides the geometric input independent of $\hbar$). The derivation is therefore non-circular.

The lattice-KMS concern. A related concern is that the KMS periodicity assumes a continuous spacetime, while the framework’s substratum is a lattice at $\epsilon = 2l_p$. The lattice introduces corrections to the KMS periodicity of order $(\epsilon\kappa/c)^2 \sim 10^{-122}$ — far below any observable threshold.

The framework’s response: trans-Planckian insensitivity (Jacobson, ’t Hooft, Verlinde, and others) ensures the result is robust against UV modifications. The continuum KMS periodicity is the dominant contribution; the lattice corrections are negligible at the relevant scales. The framework’s derivation of $\hbar$ is therefore consistent with the lattice structure rather than requiring strict continuum spacetime.

C.11 The framework’s epistemic status as a structural framework

A broader objection concerns the framework’s epistemic status: whether it should be evaluated as a *physical theory* competing with the Standard Model and general relativity, or as a *structural framework* providing foundations for these theories.

The objection. “The framework’s empirical content is mostly *retrodictive* — it explains the Standard Model’s parameters and the gravitational sector’s structure after the fact. But a genuine physical theory should make novel predictions that distinguish it from competitors. Without such predictions, the framework is not a theory; it is an interpretation.”

The framework’s response. The framework’s epistemic status is intermediate between interpretation and theory. It is a *structural framework* with substantive empirical content but with different empirical exposure than a competing physical theory would have.

What the framework retrodicts. Twenty-two Standard Model observables (gauge group, generation count, neutrino mixing angles, Cabibbo and Koide relations,

gauge couplings, Higgs quartic at the Planck scale, mass ratios, etc.). Thirteen gravitational-sector observables (Bekenstein-Hawking entropy coefficient, MOND acceleration scale, cosmological constant magnitude, etc.). The Rett wider therapeutic window in clinical biology. These are all explained as *structural consequences* of the framework’s substratum-level commitments — they are explained why they have the values they do, not just fit to data.

What the framework predicts at the present time. JUNO solar mixing angle at 0.07σ against the post-JUNO global fit — a *pre-registered* prediction made before the JUNO measurement. This is *parameter-free* in the framework’s specific sense — derived from structural content, not fitted to the measured value — but it is a retrodiction — the derivation postdates the measurement. The framework’s novel-prediction claims rest on the forward set (DESI Y5, the anisotropy amplitude, BTFR at $z > 1$, JUNO design-lifetime precision on the exact value). (The Bekenstein-Hawking $1/4$ coefficient is derived but has no direct empirical test; GW250114 constrains the classical area theorem, which the framework preserves.)

What the framework predicts forward. Specific predictions for the BEC analogue-gravity experiment (Chapter 15 §15.8) — the $\mathcal{F}(r) = 2r - r^2$ functional form at marginal C3 capacity. Gravitational wave echoes at $\Delta t = (r_+/c) \ln(r_+/2l_p)$ (Chapter 18 §18.5). Variant pathogenicity stratification for reader/writer/eraser genes (Chapter 17 §17.8). These are testable forward predictions that distinguish the framework’s content from current alternatives.

The intermediate epistemic status. The framework is therefore *not just an interpretation* — it makes specific empirical predictions that have been tested (JUNO a parameter-free retrodiction at 0.07σ ; GW250114 an area-theorem consistency; Rett a retrodiction) and that can be tested in the future (BEC, gravitational wave echoes, variant pathogenicity). But it is also *not a competing physical theory* in the conventional sense — the framework does not propose alternative dynamics that would replace the Standard Model or general relativity at the emergent level. At the emergent level, the framework reproduces standard QM and standard QFT and standard GR.

The framework’s distinctive content is at the *structural foundation* level: why standard QM has the specific form it does, why the Standard Model has the specific gauge structure and parameter values it does, why the gravitational sector has the specific features it does, and why specific biological phenomena (Rett, methylation classification) show specific empirical patterns. The framework’s value is therefore at the *foundational and unifying* level rather than at the *competing-theory* level.

This intermediate status is sometimes uncomfortable. Critics from the “framework is just interpretation” camp object that the framework lacks distinguishing predictions; critics from the “framework should propose alternative dynamics” camp object that the framework reproduces standard physics rather than replacing it. The framework’s response: it is *intentional* that the framework op-

erates at the foundational level rather than at the competing-theory level. The foundational level is where structural questions are addressed; the framework’s empirical content at this level is substantial.

C.11.1 Is the derivation gerrymandered? The selector-motivation audit

The objection. “The derivation reaches the Standard Model only because its axioms and selection principles were chosen, with hindsight, to land there. The discrete residue — $d = 3$, $K = 6$, the gauge group — is smuggled in through tuned assumptions, so the ‘derivation’ is a tautological re-encoding of the target rather than a forcing of it.”

The framework’s response. The relevant test applied to each assumption and selection principle is whether it is motivated independently of the Standard Model target. On that test the structural core does not depend on the target. Coupling-degree minimization, which fixes $K = 2d = 6$, is a definition of a “site” given a bijection rather than a tunable preference, so $K = 2d$ follows as a property of the wave equation on $\mathbb{Z}^3$ rather than as a selection rule; a different K corresponds to a different bijection. The spatial dimension $d = 3$ follows from three filters stated without reference to the gauge sector (propagating gravity, stable matter, entropy concordance); renormalizability is not used as a filter, as that would be circular. The multiplicities $(3, 2, 1)$ are the decomposition of the six nearest-neighbor link directions under the cubic point group. The audit also identifies two inputs that are not derived in this sense and are recorded as such: the amplitude-scale gauge principle (treated in the substratum development as an operational input rather than a theorem) and the identification of the propagation modes with the gauge-carrying components.

In the quantitative tier the dependence on fitted quantities is larger. There, $1/\alpha_0$ and the leading threshold coefficient $A \cdot B$ are derived, while the universal threshold δ_0 is an empirical input; the framework does not adopt the numerically convenient identifications of δ_0 (for example as a hypercharge trace), because they conflict with the independently-derived $A \cdot B$ and with the framework’s universality requirement. The resulting characterization is: derived on the structural residue, retrodiction with a small number of fits on the couplings.

What remains open. This audit bears only on whether the residue is motivated independently of the target. It does not settle the further question of whether a forced structural residue is *substantively* non-tautological as opposed to merely entailed, which is recorded as an open problem (Chapter 19 §19.3.6) and which the framework proposes to assess by external calibration against the operational reconstruction theorems of quantum theory rather than by further internal argument.

C.12 Meta-objections: credentials, scope, ambition

The framework attracts meta-objections about its author, its scope, and its ambition. These are not legitimate scientific objections in themselves, but they are sometimes raised and warrant brief engagement.

Credentials. “The author has no formal physics credentials. The framework should not be taken seriously regardless of its content.”

This is not a scientific objection. Scientific content is evaluated on its merits, not on the credentials of those who propose it. The framework’s content consists of mathematical proofs, empirical predictions, and structural derivations — all evaluable independent of the author’s biography.

However, the underlying concern (familiar standards may not be met by an outsider) can be addressed substantively. The framework provides complete proofs in the technical chapters (Chapter 1 for the characterization theorem; Chapter 7 for the gravitational derivations; etc.); two independent derivation routes for key results (Chapter 1 via Barandes correspondence and Stinespring dilation); and explicit falsifiability conditions throughout (Chapter 15 §15.8 BEC test; Chapter 18 §18.5 gravitational wave echoes; etc.). The standards the framework meets are *internal scientific standards* (proofs, predictions, falsifiability), not standards based on author identity.

The Acknowledgements section discloses AI assistance transparently. The framework’s content is the author’s intellectual responsibility; the AI assistance was in drafting, refining argumentation, and literature surveying. The technical content (proofs, derivations, predictions) is the author’s work product.

Scope. “The framework claims to address fundamental physics *and* the origin of life *and* the methodological limits of single-cell bioinformatics. This is too much for a single framework. Each domain deserves its own framework.”

The framework’s response: the claims are *interconnected by a single mechanism* (C1-C3 structural conditions operating at different scales). Separating them would obscure the connections — in particular, the framework’s value comes from being *one* structural account that produces predictions across multiple domains, not from being separate accounts for each domain. The unification is the framework’s distinctive content.

This is not a unique pattern. Statistical mechanics in the nineteenth century similarly addressed thermodynamics, kinetics, transport phenomena, and other domains with one structural account. The unification of multiple domains under one structural framework is a feature, not a bug. The framework’s specific empirical record across multiple domains provides convergent evidence for the structural foundations.

The framework is structured to allow domain-specific evaluation. Part II (fundamental physics) is self-contained and can be evaluated independently. Part III (emergence cascade) builds on Part II but is independently evaluable. Part

IV (applications) builds on Parts II-III but is also independently evaluable. A reader skeptical of the framework’s overall claim can evaluate Part II alone and engage the framework’s content at that level.

Ambition. “The framework derives QM, dissolves the cosmological constant problem, explains the dark sector, and provides a unifying account of memory in biology. This is too ambitious.”

The framework’s response: ambition is not a scientific objection. The relevant questions are: do the derivations work? do the predictions match empirical data? are the structural commitments coherent? The framework’s content provides specific answers to each: the derivations are explicit and the proofs are complete; the predictions match empirical data at substantial precision (JUNO 0.07σ , GW250114 99.999%); the structural commitments are coherent across multiple domains.

A reader who finds the framework’s ambition uncomfortable can choose to engage with subsets of the content. The framework’s content is structured to be evaluable at multiple levels — at the Chapter-level (specific empirical predictions), at the Part-level (specific structural derivations), or at the framework-level (the full structural commitments). The framework’s ambition does not require the reader to engage with all of it at once.

C.13 Closing remarks

The appendix has engaged twelve clusters of objections to the framework’s content. The pattern across responses is consistent: each objection is engaged at the level appropriate to its substance, with specific technical content where the objection is technical, with specific methodological content where the objection is methodological, and with specific empirical content where the objection concerns empirical exposure.

What the framework cannot address. The framework’s reach has genuine limits. The framework does not address:

The hard problem of consciousness. Chapter 18 §18.10 explicitly identifies consciousness as outside the framework’s primary scope. The framework provides structural commitments about the relationship between observation and physical reality (observation as necessary but not sufficient for consciousness; structural constraints on physics-based consciousness theories) but does not derive the *qualitative* character of conscious experience.

The initial conditions of the universe. Chapter 19 §19.3 identifies the cosmological initial state as an open problem. The framework’s content includes the universe’s dynamics (φ on S) but not the specific initial configuration the universe started from. Why this initial state rather than another is not addressed by the framework’s structural content.

The flavor problem. Chapter 19 §19.3 identifies the detailed structure of the

CKM and PMNS matrices as open. The framework’s content predicts gross structural features (gauge group, generation count, broad mixing-angle structure) but not the specific values of all twelve Yukawa couplings.

The baryon asymmetry. The framework does not derive the specific baryon-to-photon ratio ($n_B/n_\gamma \approx 6 \times 10^{-10}$) from first principles. This is not a gap awaiting a mechanism: the framework’s relation to baryogenesis is solution-specific (§19.3.3). The baryon asymmetry is downstream of the flavor-sector CP phases ($\delta_{\text{CKM}}, \delta_{\text{PMNS}}$), which the framework treats as solution-specific inputs rather than derived quantities; η_B inherits that input status. The framework is internally consistent on this point — it does not claim to derive η_B while inputting δ_{CKM} — but it is, by the same token, incomplete in scope on the baryon asymmetry, and says so.

These are genuine limits of the framework’s reach. The framework does not pretend to address them; they are explicitly identified as open problems or out-of-scope concerns.

What turns out to be a category error. Several apparent puzzles in fundamental physics turn out, on the framework’s analysis, to be *category errors* — problems that arise from confusing levels of description. The framework dissolves these rather than solving them.

The cosmological constant problem. The “120-order-of-magnitude discrepancy” between QFT vacuum energy and observed cosmological constant is a category error in the framework’s reading (Chapter 7 §7.7). The QFT vacuum energy is the *emergent* energy density of the trace-out; the cosmological constant is the *substratum* baseline energy density. The two operate at different levels of description, and their relationship is the framework’s specific information compression ratio rather than a direct equality.

The measurement problem. The question “how does measurement produce definite outcomes from superpositions” is a category error in the framework’s reading (Chapter 1). At the substratum level, there are no superpositions — only the substratum’s specific configuration. At the emergent level, superpositions are reduced descriptions used by embedded observers with incomplete information. “Measurement” is the embedded observer’s interaction with the substratum that updates their reduced description — not a physical mechanism producing definite outcomes from a previously-indefinite physical state.

The quantum-gravity unification problem. The standard formulation (“quantize gravity to unify with QM”) is a category error in the framework’s reading (Chapter 7 and Chapter 18 §18.4). Gravity is *substratum-level geometry* in the framework, not an emergent-level field requiring quantization. The unification with QM is automatic because both QM and GR are emergent descriptions from the same substratum. The framework does not “quantize gravity”; it explains why QM and GR are *both* emergent descriptions of a single substratum-level reality.

These dissolutions are *substantive content* of the framework. The framework’s reach is not just retrodictive (explaining established empirical patterns) but dissolutive (showing that some apparent puzzles are category errors rather than genuine problems requiring solutions).

The framework’s epistemic position summarized. The framework is a structural framework with substantive empirical content. It engages standard interpretations of QM, established no-go theorems, established results in effective field theory, and established methodological concerns through specific technical responses. It makes specific empirical predictions that have been tested and confirmed (JUNO, GW250114, Rett) and that can be tested in the future (BEC, gravitational wave echoes, variant pathogenicity). Its structural commitments are coherent across multiple domains, with the substratum-emergent operator distinction providing a unifying principle across fundamental physics and biology.

Where the framework’s reach has genuine limits (consciousness, initial conditions, flavor problem, the baryon asymmetry), the limits are explicitly identified. Where apparent puzzles turn out to be category errors (cosmological constant, measurement problem, quantum-gravity unification), the dissolutions are substantive content rather than evasions.

The appendix’s value is therefore as a reference resource for readers approaching the framework from different starting points. Readers familiar with Bell’s theorem and the no-go theorems should find §§C.2-C.3 directly useful. Readers familiar with effective field theory should find §§C.9 directly useful. Readers familiar with standard interpretations should find §C.8 directly useful. Readers concerned about credentials, scope, or ambition should find §C.12 directly useful. The appendix is intended to be navigable rather than read straight through, with each section engaging the specific cluster of concerns characteristic of one reader audience.

Glossary

This glossary defines the framework’s distinctive terminology and the technical vocabulary used throughout the book. Entries are organized alphabetically. Cross-references to chapter and appendix sections point to where each concept is most fully developed.

Note on convention: terms in *italics* within definitions indicate other glossary entries. Where the framework’s usage differs from standard physics or mathematics usage, the difference is noted explicitly.

Anomaly cancellation. A condition in quantum field theory requiring that

the sum of contributions from chiral fermions to certain triangle diagrams vanishes. In the Standard Model, anomaly cancellation between fermion generations is a structural consistency requirement that constrains the allowed hypercharge assignments. The framework derives the *unique* hypercharge assignment $(Y_Q, Y_u, Y_d, Y_L, Y_e) = (1/6, 2/3, -1/3, -1/2, -1)$ from the six anomaly conditions of $SU(3) \times SU(2) \times U(1)$, with no free parameters. Developed in Chapter 6 §6.7 and Appendix B §B.5.

Axiom 2 (Finiteness). The framework’s foundational assumption that the substratum has finitely many degrees of freedom. Motivated physically by the Bekenstein bound at the cosmological horizon ($\sim 10^{122}$ distinguishable states). Discussed in Chapter 2 §2.5, with defense in Appendix C §C.6.

Barandes correspondence. The mathematical result establishing that *P-indivisible* stochastic processes correspond structurally to unitary quantum dynamics under specific conditions. One of the two independent routes by which the framework derives emergent quantum mechanics from a deterministic substratum (the other being *Stinespring dilation*). Developed in Chapter 1 §1.6 and Appendix B §B.2.

Bekenstein-Hawking entropy. The entropy $S = A/(4l_p^2)$ assigned to a horizon of area A in units of the Planck length l_p . The factor of $1/4$ is conventionally treated as a postulate of black-hole thermodynamics. The framework *derives* the $1/4$ factor as a structural consequence of mode counting on the cosmological horizon partition. GW250114 (September 2025) confirms the classical area theorem; the $1/4$ coefficient has no direct empirical test. Developed in Chapter 7 §7.5.

Bell-inequality violations. Empirical correlations exceeding bounds derivable from local hidden-variable theories. The framework reproduces the full quantum predictions including the Tsirelson bound $2\sqrt{2}$ for CHSH correlations through *P-indivisibility* of joint dynamics, which violates Bell’s factorizability condition while preserving locality (no signaling) and measurement independence. Developed in Chapter 1 §1.7 and Appendix C §C.2.

Born rule. The quantum-mechanical prescription that the probability of measurement outcome i on state $|\psi\rangle$ is $|\langle i|\psi\rangle|^2$. Conventionally treated as a postulate of quantum mechanics. The framework derives the Born rule as the equilibrium distribution of the *read-write cycle* between *visible* and *hidden sectors* of an *embedded observer*’s partition. Developed in Chapter 1 §1.6 and Chapter 18 §18.2 (with the framework’s distinctive prediction of transient Born-rule violations in the very early universe).

BQP (Bounded-error Quantum Polynomial time). The computational complexity class of decision problems solvable by a quantum computer in polynomial time with error probability bounded by $1/3$. Conventionally treated as an empirical conjecture about what quantum computers can achieve. The framework derives BQP as the *unique* computational class accessible to embedded observers satisfying C1-C3, with both a lower bound (the framework’s emer-

gent quantum mechanics provides the BQP toolkit) and an upper bound (the quantum description is informationally complete for the embedded observer). Falsifiable by any super-BQP demonstration. Developed in Chapter 14.

C1 (Coupling). The first of the framework’s three structural conditions. Requires non-trivial coupling between the *visible* and *hidden sectors* of the *embedded observer’s* partition. Mathematically: the transition matrix T on the visible sector is not a permutation. Physically: information crosses the partition boundary in both directions. C1 fails when the visible sector is causally disconnected from the hidden sector. Discussed in Chapter 1 §1.3.

C2 (Slow-bath timescale separation). The second of the framework’s three structural conditions. Requires the hidden sector to evolve much more slowly than the visible sector: $\tau_S/\tau_B \ll 1$. Physically: the hidden sector is essentially frozen on the timescale of any visible-sector measurement. For our universe at the cosmological horizon, $\tau_S/\tau_B \sim 10^{-32}$. C2 fails when the hidden sector thermalizes between measurements, producing classical Markovian dynamics. Discussed in Chapter 1 §1.3.

C3 (Hidden-sector capacity). The third of the framework’s three structural conditions. Requires the hidden sector to have substantially more degrees of freedom than the visible sector: $N_H \gg N_V$. For our universe at the cosmological horizon, $N_H \sim 10^{122}$. C3 fails as $N_H \rightarrow N_V$, with the marginal-capacity regime producing the *partially-quantum* regime distinctive to the framework. Discussed in Chapter 1 §1.3, with the partially-quantum regime in Chapter 15 §15.7 and Chapter 18 §18.3.

Cabibbo angle. The CKM mixing parameter $\lambda = \sin \theta_C \approx 0.225$ governing first/second-generation quark mixing in the Standard Model. Conventionally a fitted empirical parameter. The framework derives $\lambda = 1/(\pi\sqrt{2})$ from a single Brillouin-zone distance on the substratum lattice, matching observation at 0.04%. Developed in Chapter 6 §6.5.

Cat state (Schrödinger’s cat). A superposition of two macroscopically distinct configurations of a single system, after Schrödinger’s 1935 thought experiment in which an entanglement chain amplifies one atom’s superposition into a formalism assigning a cat equal parts alive and dead. Conventionally the scenario motivates either an undefined measurement postulate or an added objective-collapse law. The framework dissolves it: the substratum always holds exactly one configuration; the superposition is the *embedded observer’s* indivisible bookkeeping over a definite fact recorded in the *hidden sector*; environmental interaction enacts continuous division events, making divisible (classical) odds correct for any actual cat; and opening the box is intervention, not creation. Developed in Chapter 1 §1.11, with the engineered-cat-state experimental frontier in Chapter 1 §1.6 and the corresponding falsification commitment in Chapter 4 §4.5.

CaMKII (Calcium/calmodulin-dependent protein kinase II). A neuronal signaling kinase forming dodecameric holoenzymes with bistable autophos-

phorylation states. Canonical molecular memory device in long-term potentiation. The framework identifies CaMKII as the prototypical molecular *C1-C3* system in neurons, with the kinase domain as fast subsystem and the regulatory/association domain as slow hidden sector. Pathological τ_B shifts produce the framework's account of neurodegeneration. Discussed in Chapter 16 §16.4.

Characterization theorem. The framework's central result: an embedded observer's reduced description is mathematically equivalent to unitarily evolving quantum mechanics if and only if the three structural conditions *C1*, *C2*, *C3* are satisfied. Forward direction (sufficiency): *C1-C3* $\implies$ emergent QM. Reverse direction (necessity): three contrapositive theorems, one per condition. Developed in Chapter 1 §§1.8-1.9, with necessity proofs in Appendix C §C.5.

Chk1 (Checkpoint kinase 1). A serine/threonine kinase regulating cell-cycle checkpoint responses to DNA damage. The framework identifies Chk1 as a canonical example of regulatory-domain *memory asymmetry*: Chk1 inhibitors selectively sensitize tumor cells because tumor cells rely on non-Markovian checkpoint dynamics that normal cells with intact G1 checkpoints do not require. Discussed in Chapter 16 §16.3.

Chirality (in the framework). The framework's commitment to chiral SU(2) electroweak structure derives from the staggered bipartite structure on the cubic lattice combined with the trace-out operation. Propagates through molecular physics to produce the parity-violating energy difference (PVED) between mirror-image molecules, predicting L-amino acids as the slightly more stable enantiomer. Discussed in Chapter 10 §10.7.

Chromatin. The protein-DNA complex organizing eukaryotic genomes, including nucleosome positioning, histone modifications (acetylation, methylation), DNA methylation patterns, and three-dimensional chromatin architecture. The framework treats chromatin as the *hidden sector* of the biological *C1-C3* partition with the transcriptional machinery as visible sector. Developed in Chapter 16 §16.9.

CKM (Cabibbo-Kobayashi-Maskawa) matrix. The unitary mixing matrix between Standard Model quark mass and weak interaction eigenstates. Conventionally parameterized by three mixing angles and one CP-violating phase. The framework derives the Cabibbo angle structurally and provides partial content for the broader CKM structure. Developed in Chapter 6 §6.5.

CHSH inequality. The Clauser-Horne-Shimony-Holt formulation of Bell's inequality, with the classical bound of 2 and the quantum (Tsirelson) bound of $2\sqrt{2}$. The framework reproduces exactly the quantum maximum through *P-indivisibility*. Discussed in Appendix C §C.2.

Cosmological constant problem. The discrepancy of approximately 10^{122} between quantum field theory's predicted vacuum energy ($\sim 10^{113}$ J/m³) and the observed value inferred from cosmic expansion ($\sim 6 \times 10^{-10}$ J/m³). The framework *dissolves* this problem by identifying the two quantities as properties

of logically distinct descriptions: the quantum vacuum energy is the emergent QFT's prediction within a partition; the observed cosmological constant is the gravitational consequence of the substratum's full content. Developed in Chapter 7 §7.7.

Cryptochrome. A flavoprotein in animal retinas whose photoexcitation produces a radical pair ($\text{FADH}\bullet$ and $\text{Trp}\bullet$) with millisecond singlet/triplet interconversion. The substrate for the avian magnetic compass. The framework identifies cryptochrome as a canonical $C1$ - $C3$ system in biology with the radical-pair electron spins as fast subsystem and the protein matrix as slow bath. Discussed in Chapter 13 §13.5.

Cubic decomposition. The framework's structural result that the six-dimensional representation of the cubic rotation group O on internal lattice components decomposes as $\mathbf{6} = T_1 \oplus E \oplus A_1$ — multiplicities $(3, 2, 1)$ identifying the Standard Model gauge group $\text{SU}(3) \times \text{SU}(2) \times \text{U}(1)$. Developed in Chapter 5 §5.5 and Appendix B §B.6.

DSW (Danielson-Satishchandran-Wald) horizon decoherence. A 2022 result by Danielson, Satishchandran, and Wald proving within standard QFT in curved spacetime that Killing horizons impose fundamental decoherence on nearby quantum superpositions, with rate determined by horizon geometry. External convergence with the framework's universality-class content. Discussed in Chapter 9 §9.7 and Chapter 15 §15.8.

Embedded observer. The framework's technical term for an observer who is a substructure of the system they are trying to describe — coupled to the system, bounded in extent, unable to access the full state from outside. The central object of the framework's analysis. Defined precisely in Chapter 1 §1.2, characterized by $C1$ - $C3$ in §1.3.

Emergent description. The reduced description of the visible sector after marginalizing over (tracing out) the hidden sector. In the framework, the emergent description of any system satisfying $C1$ - $C3$ is mathematically equivalent to unitary quantum mechanics. Developed throughout Chapter 1.

Epigenome. The set of chromatin marks and methylation patterns that regulate gene expression while remaining heritable across cell divisions and (partially) generations. The biological instantiation of the framework's *substratum-emergent operator distinction* at the molecular scale. Discussed in Chapter 16 §§16.9-16.10.

Extended Church-Turing Thesis (ECT). The conjecture that any physically realizable computation can be efficiently simulated by a Turing machine. Conventionally treated as an empirical hypothesis subject to potential refutation by quantum computing. The framework derives the ECT as a *theorem*: any physically realizable computation by an *embedded observer* satisfies the BQP bound. Developed in Chapter 14 §14.6.

Fine-tuning. The observation that certain physical parameters (fine-structure

constant, mass ratios, cosmological constant) appear to occupy narrow ranges supporting chemistry and complexity. The framework’s reading: most chemistry-supporting features are *structural* (derived from the framework’s commitments) rather than parametric, dissolving the conventional fine-tuning problem. Discussed in Chapter 10 §10.8.

Frauchiger-Renner theorem. A 2018 no-go theorem identifying contradictions in scenarios where quantum mechanics is applied to systems containing observers who themselves apply quantum mechanics. The framework’s response: the contradictions dissolve under *partition-relativity*, with observer-internal and external descriptions corresponding to different partitions. Discussed in Appendix C §C.3.

Gap equation (for $\hbar$). The thermal self-consistency condition between the classical horizon temperature $T_{cl} = c^2 \epsilon^2 \kappa / (8\pi G k_B)$ and the KMS temperature $T_Q = \hbar \kappa / (2\pi c k_B)$ of the emergent QFT, producing $\hbar = c^3 \epsilon^2 / (4G)$. The framework’s derivation of Planck’s constant from horizon thermodynamics. Developed in Chapter 7 §7.3 and Appendix B §B.4.

Gauge group. A continuous group of local symmetries under which the dynamics is invariant. The Standard Model gauge group is $SU(3) \times SU(2) \times U(1)$. The framework derives this gauge group as the unique commutant of a cubic-symmetric coupling matrix on the substratum lattice. Developed in Chapter 5.

Generation (fermion). In the Standard Model, the three sets of fermions with identical gauge-quantum-number assignments but different masses: (u, d, e, ν_e) , (c, s, μ, ν_μ) , (t, b, τ, ν_τ) . The framework derives three generations from the staggered taste structure on the cubic lattice. Discussed in Chapter 6 §6.2.

God’s-eye view. A hypothetical observer coextensive with the universe, with complete knowledge of all degrees of freedom. The framework identifies the God’s-eye view as the *C3* failure limit and shows that it is not a limit case of observation but a category outside it — observation is partition-relative, and without a partition, observation does not occur. Discussed in Chapter 1 §1.10.

GW250114. A high-signal-to-noise binary black-hole merger observed by LIGO on January 14, 2025 (hence “GW250114”), with the LIGO/Virgo/KAGRA analysis paper published in *Physical Review Letters* on September 10, 2025, confirming Hawking’s area theorem ($dA \geq 0$, classical horizon dynamics) at 99.999% confidence. The observation compares horizon areas and does not measure entropy, so it does not test the Bekenstein-Hawking $1/4$ coefficient. Discussed in Chapter 7 §7.5.

Hidden sector. The portion of the framework’s partition that the embedded observer does not directly access. The framework’s three structural conditions specify properties of this sector (coupling, slow evolution, large capacity). For laboratory observers in our universe, the hidden sector is everything beyond the cosmological horizon plus the trans-horizon degrees of freedom of any local

environment. Defined in Chapter 1 §1.2.

Higgs (in the framework). The framework treats the Higgs as the *singlet* (A_1) taste of the staggered decomposition — a *composite scalar* in the framework’s account rather than a fundamental field. The Higgs quartic coupling at the compositeness scale is exactly zero, with the observed value at low energies generated by RGE running. Developed in Chapter 6 §6.2.

Hsp90. A molecular chaperone that buffers genetic variation by chaperoning mutant proteins into functional conformations. The canonical evolutionary capacitor (Rutherford and Lindquist 1998). The framework identifies Hsp90-client conformational states as a *C1-C3* hidden sector with stress-induced release of accumulated variation. Discussed in Chapter 17 §17.9.

Hypercharge. The Standard Model $U(1)$ quantum number of fermion multiplets, with values $(Y_Q, Y_u, Y_d, Y_L, Y_e) = (1/6, 2/3, -1/3, -1/2, -1)$. Conventionally fitted empirical inputs. The framework derives these values as the unique anomaly-free solution. Developed in Chapter 6 §6.7 and Appendix B §B.5.

Information backflow. In the framework, the return of correlations from the hidden sector to the visible sector on the bath’s correlation timescale τ_B . The structural signature of *non-Markovian* dynamics and the framework’s distinctive prediction across multiple empirical domains. Discussed in Chapter 1 §1.5.

Jordan-Chevalley projection. The algebraic decomposition of an operator into commuting semisimple and nilpotent parts: $F = F_{ss} \cdot F_u$ with $F_u = I + N$ and $N^2 = 0$. The framework’s content: the trace-out operation extracts F_{ss} and discards N , organizing the emergent description by the representation theory of the partition. Developed in Appendix B §B.3.

JUNO (Jiangmen Underground Neutrino Observatory). A reactor neutrino experiment that in November 2025 reported the first world-leading direct measurement of the solar neutrino mixing angle: $\sin^2 \theta_{12} = 0.3092 \pm 0.0087$. The framework’s parameter-free value $1/3 - 1/(4\pi^2) = 0.3080$ matches this at 0.07σ — a retrodiction; JUNO’s design-lifetime precision provides the pre-registered forward test. Discussed in Chapter 8.

Koide relation. An empirical relation among the three charged-lepton masses: $Q = (m_e + m_\mu + m_\tau)/(\sqrt{m_e} + \sqrt{m_\mu} + \sqrt{m_\tau})^2$. Observed value $Q = 0.666661 \pm 0.000007$. The framework predicts $Q = 2/3$ from cubic-group quadratic Casimir; the match is 0.02%. Developed in Chapter 6 §6.6.

Markovian. A stochastic process whose future depends only on the present state, with no memory of past states. Standard chemical kinetics is Markovian; the framework’s emergent quantum mechanics is *non-Markovian* through *P-indivisibility*. Contrasted throughout Chapter 1.

Memory asymmetry. The framework’s unifying therapeutic principle in medicine: when a disease process depends on non-Markovian signaling dynamics that the corresponding normal tissue does not depend on, therapies can target

the memory structure rather than the catalytic function, with the framework predicting wider therapeutic windows. Developed in Chapter 16 §16.13.

MOND (Modified Newtonian Dynamics). A phenomenological framework for galactic rotation curves involving an acceleration scale $a_0 \approx 1.2 \times 10^{-10} \text{ m/s}^2$. Conventionally treated as a phenomenological alternative to dark matter. The framework derives $a_0 = cH/6$ as a structural consequence of entropy displacement at the cosmological boundary, with no free parameters. Developed in Chapter 7 §7.9.

Non-Markovian. A stochastic process whose future depends on past states beyond the present. The framework’s emergent quantum mechanics is fundamentally non-Markovian through *P-indivisibility*, distinguishing it structurally from classical Markovian dynamics. Predicted to appear at multiple scales: cosmological (quantum mechanics), biological (enzyme kinetics, gene expression), engineered (qubit noise). Defined in Chapter 1 §1.5.

NV (Nitrogen-vacancy) center. A point defect in diamond where a nitrogen atom replaces a carbon atom adjacent to a lattice vacancy. The basis for room-temperature quantum sensing with current sensitivity limited by the NV electron-spin coherence time. The framework predicts that backflow-aware sensing protocols exploiting the diamond lattice’s slow nuclear-spin bath recover information Markovian protocols discard; the resulting sensitivity gain is not quantified. Discussed in Chapter 15 §15.6.

Operator distinction (substratum-emergent). The framework’s structural distinction between operators acting at the substratum level (with specific conservation properties at the molecular or quantum-field level) and operators acting at the emergent level (with conservation properties at the coarse-grained level). The two operators can be functionally distinct even when they look identical at the emergent level. Developed in Chapter 6 §6.3 (for fundamental physics) and Chapter 16 §16.10 (for chromatin biology).

Page curve. The time-dependence of black-hole entanglement entropy: rising during the early evaporation phase, reaching a peak at the Page time, then falling as the black hole evaporates fully. The framework derives the Page curve from cycle typicality with structural Page time $t_P \approx 0.646t_{\text{evap}}$ independent of mass and greybody factors. Developed in Chapter 7 §6.

Partially-quantum regime. The framework’s distinctive prediction of dynamics distinct from both fully-quantum and fully-classical regimes, accessible at intermediate τ_S/τ_B values or marginal $C\mathcal{I}$ capacity. Empirically diagnosable through multi-time correlation measurements, with the functional form $\mathcal{F}(r) = 2r - r^2$ for capacity-controlled non-Markovianity. Developed in Chapter 13 §13.7, Chapter 15 §§15.7-15.8, Chapter 18 §18.3.

Partition-relativity. The framework’s structural feature that the same physical substratum admits multiple visible/hidden partitions, each producing its own emergent description. The emergent description depends on the partition

choice; the underlying substratum does not. Resolves several no-go theorems including Frauchiger-Renner. Developed in Chapter 1 §1.4.

P-indivisibility. The framework’s central technical concept: a stochastic process is P-indivisible if its transition matrix at time $t_2 > t_1$ cannot be expressed as the composition of the transition matrix at t_1 with a positive transition matrix. Equivalently: no Markovian process with intermediate state at t_1 can produce the same statistics. P-indivisibility is the framework’s structural mechanism for emergent quantum mechanics, and it follows from C1-C3 jointly. Developed in Chapter 1 §1.5 and Appendix B §B.7.

PMNS (Pontecorvo-Maki-Nakagawa-Sakata) matrix. The unitary mixing matrix between Standard Model neutrino mass and flavor eigenstates. The framework derives all three PMNS mixing angles from the cubic-group flavor structure, including the solar angle $\sin^2 \theta_{12} = 1/3 - 1/(4\pi^2)$, which matches the JUNO-era global fit at 0.07σ (a retrodiction). Discussed in Chapter 8.

Read-write cycle. The framework’s equilibrium dynamics between the visible and hidden sectors of the partition. The observer reads the visible sector, writes correlations into the hidden sector through the partition coupling, and reads back those correlations on subsequent measurements. The cycle produces emergent quantum mechanics at the cosmological scale and heredity at the molecular scale. Defined in Chapter 1 §1.6.

Reader/writer/eraser pharmacology. The framework’s three-axis organization of epigenetic therapeutics, corresponding directly to the operator distinction at chromatin. *Writers* establish substrate-level marks (DNMT and HMT inhibitors); *erasers* remove them (HDAC and HDM inhibitors); *readers* translate marks to functional output (BET bromodomain inhibitors). The framework’s structural prediction: reader-writer disorders show wider therapeutic windows than substrate-gene disorders, empirically confirmed for Rett syndrome. Developed in Chapter 16 §§16.11-16.12.

Rett syndrome. A neurodevelopmental disorder caused by loss-of-function mutations in MECP2, a methyl-CpG binding reader protein. Guy et al. (2007) demonstrated that restoration of MECP2 expression in adult MECP2-knockout mice reverses the neurological phenotype, confirming the framework’s pre-registered prediction of wider therapeutic windows for reader-writer disorders. The framework’s strongest empirical anchor outside fundamental physics. Discussed in Chapter 16 §16.12.

RNA world. The hypothesis that early life used RNA for both information storage and catalysis before DNA and proteins evolved their specialized roles. The framework’s structural reading: RNA-like architecture is the simplest single-molecule C1-C3 system (simultaneously template and catalyst), making the RNA world a structural prediction rather than a historical accident. Developed in Chapter 11 §11.7.

Staggered (fermion structure). The placement of independent fermion

modes at alternating cubic lattice sites, producing the chiral structure of the framework’s emergent QFT and the multiple-taste structure that gives rise to three generations. Developed in Chapter 5 §5.6 and Chapter 6 §6.2.

Standard Model. The current best theory of fundamental particle physics: a chiral gauge theory with gauge group $SU(3) \times SU(2) \times U(1)$, three fermion generations, one Higgs doublet, and approximately 19 free parameters. The framework derives the gauge group, the matter content, the chirality structure, the anomaly-free hypercharges, and 22 of the parameters as structural consequences with no free inputs. Developed in Chapters 5-6.

Stinespring dilation. A 1955 theorem of operator algebra establishing that any completely-positive trace-preserving map on a Hilbert space can be realized as a unitary evolution on a larger Hilbert space with subsequent tracing-out. The framework uses Stinespring dilation as the second independent route from the deterministic substratum to emergent quantum mechanics (alongside the *Barandes correspondence*). Developed in Chapter 1 §1.6 and Appendix B §B.2.

Substratum. The framework’s underlying deterministic, finite-dimensional, bijective dynamical system (S, φ) from which emergent quantum mechanics is derived through the trace-out operation. The substratum is the framework’s primitive ontological commitment — the “what is really there” beneath the emergent quantum-mechanical description. Constructed concretely in Chapter 2.

Substratum-emergent operator distinction. *See operator distinction.*

Tier 1, Tier 2, Tier 3. The framework’s stratification of predictions by their robustness across substratum representatives. *Tier 1:* results that follow from the framework’s universality class (shared across all substrata satisfying the structural conditions). *Tier 2:* results dependent on broader structural commitments but not on specific lattice details. *Tier 3:* solution-specific results dependent on the particular cubic-lattice realization. Developed in Chapter 9 §§9.3-9.4.

TLS (Two-level system). Structural defects in amorphous oxide substrates of superconducting qubits, characterized by flipping between two configurations at characteristic rates. The framework identifies the TLS bath as the *hidden sector* in superconducting qubit C1-C3 systems, with $\tau_S/\tau_B \sim 10^{-3}$ producing the framework’s predicted correlated error structure. Discussed in Chapter 15 §15.3.

Trace-out. The mathematical operation of marginalizing over (summing or integrating out) the hidden-sector degrees of freedom, producing the reduced description on the visible sector. The framework’s central technical operation. Equivalent to partial trace in quantum mechanics. Defined throughout Chapter 1.

Tsirelson bound. The maximum CHSH inequality violation achievable by quantum mechanics: $2\sqrt{2}$. The framework reproduces exactly this bound

through *P-indivisibility*, distinguishing it from theories predicting weaker (classical) or stronger (PR-box) violations. Discussed in Appendix C §C.2.

U-parity. The framework’s grading structure in the neutrino sector, distinguishing U^+ and U^- states in the cubic-group flavor classification. Produces the framework’s distinctive sum rule and structural relation among the three PMNS mixing angles. Developed in Chapter 8 §8.4.

Universality class. A set of substrata that share the same emergent description at the framework’s Tier 1 level. The framework’s content: the substratum-to-emergent map is many-to-one, with multiple bijections producing the same observable physics. Developed in Chapter 9 §9.3.

Vibronic coherence. Quantum coherence between electronic exciton states and specific vibrational modes of the protein environment. The framework identifies the picosecond-scale coherences observed in photosynthetic complexes (FMO, LH2, photosystem II) as vibronic rather than purely electronic, consistent with the framework’s C1-C3 architecture. Discussed in Chapter 13 §13.4.

Visible sector. The portion of the framework’s partition that the embedded observer directly accesses. Contrasted with the *hidden sector*. The visible sector is where measurements occur and where the emergent description applies. Defined in Chapter 1 §1.2.

Wolpert’s theorem. David Wolpert’s results (2001, 2008) on fundamental limits of inference devices embedded in the system they are trying to predict. A physical analog of Gödel’s incompleteness theorem and Turing’s halting problem result — logical, worst-case, and unquantified. The framework’s partition-induced trace-out is the mechanism of a *parallel physical* incompleteness, with quantum mechanics as the specific cost; Wolpert’s result corroborates the pattern but is not the source of the framework’s derivation. Discussed in Chapter 12 §12.8, Appendix C §C.7.

Yamanaka factors. The four transcription factors (Oct4, Sox2, Klf4, c-Myc) whose transient expression reprograms somatic cells to induced pluripotent stem cells. The framework identifies iPSC reprogramming as confirming the *substratum-emergent operator distinction* at chromatin: the emergent operator (cell identity) flips while substantial residual substrate-level operator (DNA methylation patterns reflecting source cell type) persists. Discussed in Chapter 16 §16.10.

Bibliography

This bibliography lists the works cited in the book, organized by topical area. Within each section, entries are alphabetical by first author. The framework’s own technical companion papers are listed first under “Framework-internal cita-

tions”; subsequent sections cover foundations and interpretation, fundamental physics, gravitational physics and cosmology, chemistry and biology, evolution and molecular biology, medicine and pharmacology, bioinformatics and computational biology, quantum information and computing, and historical and philosophical works.

Framework-internal citations

Maybaum, A. (2026). *The Observational Incompleteness Framework*. Zenodo, concept DOI: 10.5281/zenodo.19060318 (resolves to the current release). [Repository containing the framework’s full technical content: the core papers — Main (central theorem and emergent quantum mechanics), Substratum (substratum construction and reconstruction theorem), Structure (structural realism), SM (Standard Model derivation), GR (gravitational sector) — together with the focused presentation Juno (neutrino-sector prediction), the methodology paper Physics Modulo Gauge, and the companion documents Explainer, Complexity, Medicine, and Bioinformatics.]

Foundations of quantum mechanics and interpretation

Arnold, V. I. (1989). *Mathematical Methods of Classical Mechanics* (2nd ed., Graduate Texts in Mathematics 60). Springer-Verlag.

Engel, K.-J., & Nagel, R. (2000). *One-Parameter Semigroups for Linear Evolution Equations* (Graduate Texts in Mathematics 194). Springer-Verlag.

Bacciagaluppi, G., & Valentini, A. (2009). *Quantum Theory at the Crossroads: Reconsidering the 1927 Solvay Conference*. Cambridge University Press.

Barandes, J. A. (2023). The Stochastic-Quantum Correspondence. *arXiv*:2302.10778.

Barandes, J. A. (2024). The Stochastic-Quantum Theorem. *arXiv*:2309.03085.

Braunstein, S. L., & Pati, A. K. (2007). Quantum information cannot be completely hidden in correlations: Implications for the black-hole information paradox. *Physical Review Letters* 98: 080502.

Bell, J. S. (1964). On the Einstein-Podolsky-Rosen paradox. *Physics Physique Fizika* 1: 195-200.

Bong, K.-W., Utreras-Alarcón, A., Ghafari, F., Liang, Y.-C., Tischler, N., Cavalcanti, E. G., Pryde, G. J., & Wiseman, H. M. (2020). A strong no-go theorem on the Wigner’s friend paradox. *Nature Physics* 16: 1199-1205.

Breuer, T. (1995). The impossibility of accurate state self-measurements. *Philosophy of Science* 62(2): 197–214.

- Clauser, J. F., Horne, M. A., Shimony, A., & Holt, R. A. (1969). Proposed experiment to test local hidden-variable theories. *Physical Review Letters* 23: 880-884.
- Colbeck, R., & Renner, R. (2011). No extension of quantum theory can have improved predictive power. *Nature Communications* 2: 411.
- Dalla Chiara, M. L. (1977). Logical self reference, set theoretical paradoxes and the measurement problem in quantum mechanics. *Journal of Philosophical Logic* 6: 331-347.
- de la Peña, L., Cetto, A. M., & Valdés-Hernández, A. (2014). *The Emerging Quantum: The Physics Behind Quantum Mechanics*. Springer.
- Einstein, A., Podolsky, B., & Rosen, N. (1935). Can quantum-mechanical description of physical reality be considered complete? *Physical Review* 47: 777-780.
- Fein, Y. Y., Geyer, P., Zwick, P., Kialka, F., Pedalino, S., Mayor, M., Gerlich, S., and Arndt, M. (2019). Quantum superposition of molecules beyond 25 kDa. *Nature Physics* 15: 1242-1245.
- Fine, A. (1982). Hidden variables, joint probability, and the Bell inequalities. *Physical Review Letters* 48: 291-295.
- Nimmrichter, S., and Hornberger, K. (2013). Macroscopicity of mechanical quantum superposition states. *Physical Review Letters* 110: 160403.
- Pedalino, S., Ramírez-Galindo, B. E., Ferstl, R., Hornberger, K., Arndt, M., and Gerlich, S. (2026). Probing quantum mechanics with nanoparticle matter-wave interferometry. *Nature* 649: 866.
- Rössler, O. E. (1998). *Endophysics: The World as an Interface*. World Scientific.
- Schrödinger, E. (1935). Die gegenwärtige Situation in der Quantenmechanik. *Die Naturwissenschaften* 23: 807-812, 823-828, 844-849.
- Szangolies, J. (2018). Epistemic horizons and the foundations of quantum mechanics. *Foundations of Physics* 48(12): 1669-1697.
- Valdés-Hernández, A., de la Peña, L., & Cetto, A. M. (2011). Bipartite entanglement induced by a common background (zero-point) radiation field. *Foundations of Physics* 41: 843-862.
- Frauchiger, D., & Renner, R. (2018). Quantum theory cannot consistently describe the use of itself. *Nature Communications* 9: 3711.
- Kochen, S., & Specker, E. P. (1967). The problem of hidden variables in quantum mechanics. *Journal of Mathematics and Mechanics* 17: 59-87.
- Pusey, M. F., Barrett, J., & Rudolph, T. (2012). On the reality of the quantum state. *Nature Physics* 8: 475-478.

- Slagle, K., & Preskill, J. (2023). Emergent quantum mechanics at the boundary of a local classical lattice model. *Physical Review A* 108: 012217. [arXiv:2207.09465](#).
- Stinespring, W. F. (1955). Positive functions on *C-algebras*. Proceedings of the American Mathematical Society* 6: 211-216.
- Tsirelson, B. S. (1980). Quantum generalizations of Bell's inequality. *Letters in Mathematical Physics* 4: 93-100.
- Wolpert, D. H. (2001). Computational capabilities of physical systems. *Physical Review E* 65: 016128. [The 2001 publication of arguments first developed in 1992.]
- Wood, C. J., & Spekkens, R. W. (2015). The lesson of causal discovery algorithms for quantum correlations: causal explanations of Bell-inequality violations require fine-tuning. *New Journal of Physics* 17: 033002.

Standard Model physics and particle phenomenology

- 't Hooft, G. (1993). Dimensional reduction in quantum gravity. [arXiv:gr-qc/9310026](#). [Origin of the holographic principle.]
- Capozzi, F., Lisi, E., Marccone, F., Marrone, A., & Palazzo, A. (2025). Updated bounds on the (1,2) neutrino oscillation parameters after first JUNO results. [arXiv:2511.21650](#). *Physical Review D* 114: 016026 (2026).
- Cabibbo, N. (1963). Unitary symmetry and leptonic decays. *Physical Review Letters* 10: 531-533. [Original Cabibbo angle.]
- Cech, T. R. (1989). Self-splicing and enzymatic activity of an intervening sequence RNA from *Tetrahymena*. *Nobel Lecture*, December 8, 1989. Published in *Les Prix Nobel 1989*. Reprinted in *Bioscience Reports* 10: 239-261.
- Cech, T. R., Zaug, A. J., & Grabowski, P. J. (1981). In vitro splicing of the ribosomal RNA precursor of *Tetrahymena*: involvement of a guanosine nucleotide in the excision of the intervening sequence. *Cell* 27: 487-496.
- JUNO Collaboration (2025). First measurement of reactor neutrino oscillations at JUNO. [arXiv:2511.14593](#); *Nature* 654: 343 (2026). Reports $\sin^2 \theta_{12} = 0.3092 \pm 0.0087$ and $\Delta m_{21}^2 = (7.50 \pm 0.12) \times 10^{-5} \text{ eV}^2$ from the first 59.1 days of data collected after detector completion in August 2025.
- Koide, Y. (1983). Fermion-boson two-body model of quarks and leptons and Cabibbo mixing. *Lettere al Nuovo Cimento* 34: 201-205. [Original Koide relation.]
- Pontecorvo, B. (1957). Mesonium and antimesonium. *Soviet Physics JETP* 6: 429.

Wu, C. S., Ambler, E., Hayward, R. W., Hoppes, D. D., & Hudson, R. P. (1957). Experimental test of parity conservation in beta decay. *Physical Review* 105: 1413-1415.

Gravitational physics, cosmology, and quantum gravity

Abedi, J., Dykaar, H., & Afshordi, N. (2017). Echoes from the abyss: Tentative evidence for Planck-scale structure at black hole horizons. *Physical Review D* 96: 082004. arXiv:1612.00266.

Bekenstein, J. D. (1981). Universal upper bound on the entropy-to-energy ratio for bounded systems. *Physical Review D* 23: 287-298.

Bertini, N. R., Novaes, M. H., von Marttens, R., & Shapiro, I. L. (2025). Renormalization group corrections to Λ CDM model and observational consequences for H_0 tension. arXiv:2509.24026 [astro-ph.CO; gr-qc; hep-th]. [Lead author initial in earlier in-text references was given as “S.” — correct lead author is Nicolas R. Bertini.]

Bousso, R. (1999). A covariant entropy conjecture. *Journal of High Energy Physics* 7: 4.

Chandrasekaran, V., Longo, R., Penington, G., & Witten, E. (2022). An algebra of observables for de Sitter space. *Journal of High Energy Physics* 2023(02): 82.

Chen, C.-H., & Penington, G. (2024). A clock is just a way to tell the time: gravitational algebras in cosmological spacetimes. arXiv:2406.02116. [Constructs gauge-invariant gravitational algebras in cosmological spacetimes with observers, focusing on slow-roll inflation and evaporating Schwarzschild-de Sitter black holes; corroborates the observer-incorporated semiclassical gravity program adjacent to the DSW horizon-decoherence result. Previously cited in-text as “Chen, Penington, and Salzetta (2024)”]; identified during May 2026 second-sweep citation-verification pass.]

Cornish, N. J., Spergel, D. N., Starkman, G. D., & Komatsu, E. (2004). Constraining the topology of the Universe. *Physical Review Letters* 92: 201302. arXiv:astro-ph/0310233.

Danielson, D., Satishchandran, G., & Wald, R. M. (2022). Black holes decohere quantum superpositions. *International Journal of Modern Physics D* 31: 2241003.

DESI Collaboration (2025). DESI DR2 results II: Measurements of baryon acoustic oscillations and cosmological constraints. *Physical Review D* 112: 083515. arXiv:2503.14738.

Genzel, R., Förster Schreiber, N. M., Übler, H., Lang, P., Naab, T., Bender, R., et al. (2017). Strongly baryon-dominated disk galaxies at the peak of galaxy formation ten billion years ago. *Nature* 543: 397-401.

- Gralla, S. E., & Wei, H. (2024). Decoherence from horizons: General formulation and rotating black holes. *Physical Review D* 109: 065031. arXiv:2311.11461.
- Hawking, S. W. (1976). Breakdown of predictability in gravitational collapse. *Physical Review D* 14: 2460-2473.
- Harlow, D., Usatyuk, M., & Zhao, Y. (2025). Quantum mechanics and observers for gravity in a closed universe. arXiv:2501.02359.
- Jacobson, T. (1995). Thermodynamics of spacetime: The Einstein equation of state. *Physical Review Letters* 75: 1260-1263.
- Jiao, Y., et al. (2023). Detection of the Keplerian decline in the Milky Way rotation curve. *Astronomy & Astrophysics* 678: A208.
- Abac, A. G., et al. (LIGO Scientific, Virgo, and KAGRA Collaborations) (2025). GW250114: Testing Hawking’s Area Law and the Kerr Nature of Black Holes. *Physical Review Letters*, published September 10, 2025. DOI: 10.1103/kw5g-d732.
- LUX-ZEPLIN Collaboration (2025). New search for WIMP dark matter with 417 live days of LUX-ZEPLIN data. Announcement December 8, 2025, with preprint and submission to *Physical Review Letters*. Based on the largest dataset ever collected by a dark matter detector: 417 live days of data from March 2023 to April 2025, probing WIMP masses between 3 and 9 GeV/ c^2 . Earlier (2024) LZ results: J. Aalbers et al. (LZ Collaboration), “Dark Matter Search Results from 4.2 Tonne-Years of Exposure of the LUX-ZEPLIN (LZ) Experiment,” *Physical Review Letters* 135: 011802 (2025), arXiv:2410.17036.
- Maldacena, J. (2024). Real observers solving imaginary problems. arXiv:2412.14014.
- Planck Collaboration (2018). Planck 2018 results. VI. Cosmological parameters. *Astronomy & Astrophysics* 641: A6.
- Susskind, L. (1995). The world as a hologram. *Journal of Mathematical Physics* 36: 6377-6396.
- Tsang, K. W., Rollier, M., Ghosh, A., Samajdar, A., Agathos, M., Chatziioannou, K., Cardoso, V., Khanna, G., & Van Den Broeck, C. (2018). A morphology-independent data analysis method for detecting and characterizing gravitational wave echoes. *Physical Review D* 98: 024023. arXiv:1804.04877.
- Verlinde, E. P. (2017). Emergent gravity and the dark universe. *SciPost Physics* 2: 016. arXiv:1611.02269. [Cited in-text as “Verlinde (2024)”]; no substantive 2024 Verlinde-authored emergent-gravity paper could be identified, and the 2017 SciPost paper remains the canonical reference. The in-text year should be corrected to 2017, or the citation reframed if a specific later work was intended.]
- Westerweck, J., Nielsen, A. B., Fischer-Birnholtz, O., Cabero, M., Capano, C., Dent, T., Krishnan, B., Meadors, G., & Nitz, A. H. (2018). Low significance of

evidence for black hole echoes in gravitational wave data. *Physical Review D* 97: 124037. arXiv:1712.09966.

Wilson-Gerow, J., Dugad, A., & Chen, Y. (2024). Decoherence by warm horizons. *Physical Review D* 110: 045002. arXiv:2405.00804. [Direct extension of the DSW horizon-decoherence result; provides a local fluctuation-dissipation derivation complementing the original global one. Previously cited in-text as “Wilson-Gerow and Verlinde (2024)”; identified during May 2026 second-sweep citation-verification pass — Verlinde is not a co-author on this paper.]

Chemistry and origin of life

Guerrier-Takada, C., Gardiner, K., Marsh, T., Pace, N., & Altman, S. (1983). The RNA moiety of ribonuclease P is the catalytic subunit of the enzyme. *Cell* 35(3 Pt 2): 849-857. [Foundational paper from Altman’s group cited in the 1989 Nobel Prize in Chemistry shared with Cech.]

Damer, B., & Deamer, D. (2020). The hot spring hypothesis for an origin of life. *Astrobiology* 20: 429-452.

Ferris, J. P., Hill, A. R., Liu, R., & Orgel, L. E. (1996). Synthesis of long prebiotic oligomers on mineral surfaces. *Nature* 381: 59-61.

Frank, F. C. (1953). On spontaneous asymmetric synthesis. *Biochimica et Biophysica Acta* 11: 459-463.

Kauffman, S. A. (1986). Autocatalytic sets of proteins. *Journal of Theoretical Biology* 119: 1-24.

Miller, S. L. (1953). A production of amino acids under possible primitive earth conditions. *Science* 117: 528-529.

Soai, K., Shibata, T., Morioka, H., & Choji, K. (1995). Asymmetric autocatalysis and amplification of enantiomeric excess of a chiral molecule. *Nature* 378: 767-768.

Wächtershäuser, G. (1990). Evolution of the first metabolic cycles. *Proceedings of the National Academy of Sciences USA* 87: 200-204.

Evolution and molecular biology

Bedford, T., & Hartl, D. L. (2008). Overdispersion of the molecular clock: temporal variation of gene-specific substitution rates in *Drosophila*. *Molecular Biology and Evolution* 25: 1631-1638.

Cutler, D. J. (2000). Understanding the overdispersed molecular clock. *Genetics* 154: 1403-1417.

- Eldredge, N., & Gould, S. J. (1972). Punctuated equilibria: an alternative to phyletic gradualism. In: T. J. M. Schopf (ed.), *Models in Paleobiology*, pp. 82-115. Freeman, Cooper.
- Gillespie, J. H. (1989). Lineage effects and the index of dispersion of molecular evolution. *Molecular Biology and Evolution* 6: 636-647.
- Maynard Smith, J., & Szathmáry, E. (1995). *The Major Transitions in Evolution*. Oxford University Press.
- Odling-Smee, F. J., Laland, K. N., & Feldman, M. W. (2003). *Niche Construction: The Neglected Process in Evolution*. Princeton University Press.
- Rutherford, S. L., & Lindquist, S. (1998). Hsp90 as a capacitor for morphological evolution. *Nature* 396: 336-342.
- Tao, Q., Tamura, K., U. Mello, B., & Kumar, S. (2019). Reliable confidence intervals for relevant timetree-based parameters. *Molecular Biology and Evolution* 36: 811-824.
- Wiltschko, R., & Wiltschko, W. (2019). Magnetoreception in birds. *Journal of the Royal Society Interface* 16: 20190295.
- Wiser, M. J., Ribeck, N., & Lenski, R. E. (2013). Long-term dynamics of adaptation in asexual populations. *Science* 342: 1364-1367.
- Zeng, L.-W., Comeron, J. M., Chen, B., & Kreitman, M. (1998). The molecular clock revisited: the rate of synonymous vs. replacement change in *Drosophila*. *Genetica* 102/103: 369-382. DOI: 10.1023/A:1017035109224. [Source of the *Drosophila* dispersion-index range $R(t) \approx 1.6-2.6$ cited in the text; previously attributed in-text to a “Wang et al. (2010)” entry that could not be located. The 1.6-2.6 range traces to this paper via Bedford & Hartl (2008).]

Medicine, pharmacology, and clinical biology

- Bernstein, B. E., Mikkelsen, T. S., Xie, X., Kamal, M., Huebert, D. J., Cuff, J., et al. (2006). A bivalent chromatin structure marks key developmental genes in embryonic stem cells. *Cell* 125: 315-326.
- Cuddapah, V. A., Pillai, R. B., Shekar, K. V., Lane, J. B., Motil, K. J., Skinner, S. A., et al. (2014). Methyl-CpG-binding protein 2 (MECP2) mutation type is associated with disease severity in Rett syndrome. *Journal of Medical Genetics* 51: 152-158.
- English, B. P., Min, W., van Oijen, A. M., Lee, K. T., Luo, G., Sun, H., Cherayil, B. J., Kou, S. C., & Xie, X. S. (2006). Ever-fluctuating single enzyme molecules: Michaelis-Menten equation revisited. *Nature Chemical Biology* 2: 87-94.
- Prebet, T., Sun, Z., Figueroa, M. E., Ketterling, R., Melnick, A., Greenberg, P., Juckett, M., Smith, M. R., Malick, L., Paietta, E., Czader, M., Gabrilove,

- J., Erba, H. P., Gore, S. D., & Tallman, M. S. (2014). Prolonged administration of azacitidine with or without entinostat for myelodysplastic syndrome and acute myeloid leukemia with myelodysplasia-related changes: Results of the US Leukemia Intergroup trial E1905. *Journal of Clinical Oncology* 32(12): 1242-1248. [Phase II randomized trial NCT00313586 testing scheduling of azacitidine $\pm$ HDAC inhibitor entinostat as primary endpoint.]
- Guy, J., Gan, J., Selfridge, J., Cobb, S., & Bird, A. (2007). Reversal of neurological defects in a mouse model of Rett syndrome. *Science* 315: 1143-1147.
- Horvath, S. (2013). DNA methylation age of human tissues and cell types. *Genome Biology* 14: R115. DOI: 10.1186/gb-2013-14-10-r115.
- Klinman, J. P., & Kohen, A. (2013). Hydrogen tunneling links protein dynamics to enzyme catalysis. *Annual Review of Biochemistry* 82: 471-496.
- Lin, X., et al. (2024). Residual epigenetic memory in induced pluripotent stem cells. *Frontiers in Cell and Developmental Biology* 12: 1306530. DOI: 10.3389/fcell.2024.1306530.
- Nicolas, D., Phillips, N. E., & Naef, F. (2018). What shapes eukaryotic transcriptional bursting? *Molecular BioSystems* 14: 1623-1629.
- Quintás-Cardama, A., Santos, F. P. S., & Garcia-Manero, G. (2011). Therapy with azanucleosides for myelodysplastic syndromes. *Nature Reviews Clinical Oncology* 8: 433-444.
- Stasevich, T. J., Hayashi-Takanaka, Y., Sato, Y., Maehara, K., Ohkawa, Y., Sakata-Sogawa, K., Tokunaga, M., Nagase, T., Nozaki, N., McNally, J. G., & Kimura, H. (2014). Regulation of RNA polymerase II activation by histone acetylation in single living cells. *Nature* 516: 272-275. DOI: 10.1038/nature13714.
- Tontsch-Grunt, U., Traexler, P.-E., Baum, A., Musa, H., Marzin, K., Wang, S., Trapani, F., Engelhardt, H., & Solca, F. (2022). Therapeutic impact of BET inhibitor BI 894999 treatment: backtranslation from the clinic. *British Journal of Cancer* 127: 577-586. DOI: 10.1038/s41416-022-01815-5. PMID: 35444289. [Previously cited in-text as “Marek et al. 2022”; identified during May 2026 citation-verification pass. The paper reports BET inhibitor BI 894999 combined with the clinical p300/CBP inhibitor CCS1477 in NUT carcinoma.]
- Voigt, P., Tee, W.-W., & Reinberg, D. (2013). A double take on bivalent promoters. *Genes & Development* 27: 1318-1338. DOI: 10.1101/gad.219626.113.
- Wilen, C. D., Abdullah, S., Kurinsky, N. A., Stanford, C., Cardani, L., D’Imperio, G., Tomei, C., Faoro, L., Ioffe, L. B., Liu, C. H., Opremcak, A., Christensen, B. G., DuBois, J. L., & McDermott, R. (2021). Correlated charge noise and relaxation errors in superconducting qubits. *Nature* 594: 369-373. DOI: 10.1038/s41586-021-03557-5.
- Wilsker, D., Petermann, E., Helleday, T., & Bunz, F. (2008). Essential function

of Chk1 can be uncoupled from DNA damage checkpoint and replication control. *Proceedings of the National Academy of Sciences USA* 105: 20752-20757.

Ye, B. S., Chang, K. W., Kang, S., Jeon, S., & Chang, J. W. (2025). Repetitive and extensive focused ultrasound-mediated bilateral frontal blood-brain barrier opening for Alzheimer’s disease. *Journal of Neurosurgery* 142: 1263-1272. DOI: 10.3171/2024.8.JNS24989. [Korean trial NCT06158789; previously entered as “Ye, J. H.” — the lead author is Byoung Seok Ye. The study reported improvement in neuropsychiatric symptoms from blood-brain-barrier opening alone, without concurrent antibody medication.]

Zhang, X., Zegar, T., Lucas, A., Morrison-Smith, C., Knox, T., French, C. A., Knapp, S., Müller, S., & Siveke, J. T. (2020). Therapeutic targeting of p300/CBP HAT domain for the treatment of NUT midline carcinoma. *Oncogene* 39: 4770-4779. DOI: 10.1038/s41388-020-1301-9. PMID: 32366905. [Previously cited in-text as “Bauer et al. 2020”; identified during May 2026 citation-verification pass. The paper reports combination of p300/CBP HAT inhibitor A-485 with BET inhibitor JQ1, with synergistic effects against NUT midline carcinoma.]

Bioinformatics and computational biology

Bergen, V., Lange, M., Peidli, S., Wolf, F. A., & Theis, F. J. (2020). Generalizing RNA velocity to transient cell states through dynamical modeling. *Nature Biotechnology* 38: 1408-1414.

Bowling, S., Sritharan, D., Osorio, F. G., Nguyen, M., Cheung, P., Rodriguez-Fraticelli, A., et al. (2020). An engineered CRISPR-Cas9 mouse line for simultaneous readout of lineage histories and gene expression profiles in single cells. *Cell* 181: 1410-1422.

Bravo González-Blas, C., et al. (2023). SCENIC+: single-cell multiomic inference of enhancers and gene regulatory networks. *Nature Methods* 20: 1355-1367.

Capper, D., Jones, D. T. W., Sill, M., Hovestadt, V., Schrimpf, D., Sturm, D., et al. (2018). DNA methylation-based classification of central nervous system tumours. *Nature* 555: 469-474.

Csendes, G., Petenei, A., Heitz, A., Boldog, D., Bencze, A., Má dai, M., Megyeri, A., Szalka, B., Kemenes, B., Schoenholzer, F., Frey, M., Cucchi, F., Tomek, M. P., Werner, C., Sonzogni, A., Stiasny, J., Crisan, M., Rodriguez Hernandez, M., Ribeiro Frazão, J., Bartoš, F., Bischof, B., Hammer, S., Domanski, J., Mey, A. S. J. S., Goyal, R., Schölkopf, B., & Theis, F. J. (2024). Foundation models in single-cell transcriptomics: benchmarking and gene-attention analysis. *Nature Communications* (forthcoming).

Cui, H., Wang, C., Maan, H., Pang, K., Luo, F., Duan, N., & Wang, B. (2023). scGPT: toward building a foundation model for single-cell multi-omics using

generative AI. *Nature Methods* (2024).

Drexler, R., Brembach, F., Sauvigny, J., Ricklefs, F. L., Eckhardt, A., Bode, H., Gempt, J., Lamszus, K., Westphal, M., Schüller, U., & Mohme, M. (2024). Unclassifiable CNS tumors in DNA methylation-based classification: clinical challenges and prognostic impact. *Acta Neuropathologica Communications* 12: 9. DOI: 10.1186/s40478-024-01728-9.

Gorin, G., Fang, M., Chari, T., & Pachter, L. (2022). RNA velocity unraveled. *PLOS Computational Biology* 18: e1010492.

Saxton, M. N., Morisaki, T., Krapf, D., Kimura, H., & Stasevich, T. J. (2023). Live-cell imaging uncovers the relationship between histone acetylation, transcription initiation, and nucleosome mobility. *Science Advances* 9(40): eadh4819. DOI: 10.1126/sciadv.adh4819.

Hummer, G., & Szabo, A. (2001). Free energy reconstruction from nonequilibrium single-molecule pulling experiments. *Proceedings of the National Academy of Sciences USA* 98: 3658-3661. DOI: 10.1073/pnas.071034098. [The manuscript cites the Hummer-Szabo single-molecule kinetics program collectively; the companion methodological paper is Hummer, G., & Szabo, A. (2010), Free energy profiles from single-molecule pulling experiments, *PNAS* 107: 21441-21446, DOI: 10.1073/pnas.1015661107.]

Kedzierska, K. Z., Crawford, L., Amini, A. P., & Lu, A. X. (2024). Assessing the limits of zero-shot foundation models in single-cell biology. *bioRxiv*. DOI: 10.1101/2023.10.16.561085.

Koelsche, C., et al. (2021). Sarcoma classification by DNA methylation profiling. *Nature Communications* 12: 498.

La Manno, G., Soldatov, R., Zeisel, A., Braun, E., Hochgerner, H., Petukhov, V., et al. (2018). RNA velocity of single cells. *Nature* 560: 494-498.

Li, T., et al. (2022). MultiVelo: chromatin velocity for joint inference of trajectories and RNA velocity. *Nature Biotechnology* (2023).

Lu, H. P., Xun, L., & Xie, X. S. (1998). Single-molecule enzymatic dynamics. *Science* 282: 1877-1882.

Marbach, D., Schaffter, T., Mattiussi, C., & Floreano, D. (2009). Generating realistic in silico gene networks for performance assessment of reverse engineering methods. *Journal of Computational Biology* 16: 229-239. DOI: 10.1089/cmb.2008.09TT. [The DREAM5 crowd-inference synthesis is Marbach, D., Costello, J. C., Küffner, R., Vega, N. M., Prill, R. J., Camacho, D. M., et al. (2012), Wisdom of crowds for robust gene network inference, *Nature Methods* 9: 796-804, DOI: 10.1038/nmeth.2016.]

McEwen, M., Faoro, L., Arya, K., Dunsworth, A., Huang, T., Kim, S., et al. (2022). Resolving catastrophic error bursts from cosmic rays in large arrays

of superconducting qubits. *Nature Physics* 18: 107-111. DOI: 10.1038/s41567-021-01432-8.

McKenna, A., Findlay, G. M., Gagnon, J. A., Horwitz, M. S., Schier, A. F., & Shendure, J. (2016). Whole-organism lineage tracing by combinatorial and cumulative genetic barcoding. *Science* 353: aaf7907.

Roohani, Y., Huang, K., & Leskovec, J. (2024). Predicting transcriptional outcomes of novel multigene perturbations with GEARS. *Nature Biotechnology* 42(6): 927-935. DOI: 10.1038/s41587-023-01905-6. [Previously cited in-text as “Pyhel et al. 2024”; identified during May 2026 citation-verification pass. GEARS integrates deep learning with a gene-coexpression knowledge graph and Gene-Ontology-derived perturbation knowledge graph to predict transcriptional responses to single and multi-gene perturbations.]

Stolovitzky, G., Monroe, D., & Califano, A. (2007). Dialogue on reverse-engineering assessment and methods: The DREAM of high-throughput pathway inference. *Annals of the New York Academy of Sciences* 1115: 1-22. DOI: 10.1196/annals.1407.021.

Tedesco, M., Giannese, F., Lazarević, D., Giansanti, V., Rosano, D., Monzani, S., et al. (2022). Chromatin Velocity reveals epigenetic dynamics by single-cell profiling of heterochromatin and euchromatin. *Nature Biotechnology* 40: 235-244. DOI: 10.1038/s41587-021-01031-1.

Theodoris, C. V., Xiao, L., Chopra, A., Chaffin, M. D., Sayed, Z. R. A., Hill, M. C., et al. (2023). Transfer learning enables predictions in network biology. *Nature* 618: 616-624.

Quantum biology and biological information processing

Cao, J., Cogdell, R. J., Coker, D. F., Duan, H.-G., Hauer, J., Kleinekathöfer, U., et al. (2020). Quantum biology revisited. *Science Advances* 6: eaaz4888. DOI: 10.1126/sciadv.aaz4888. [Cited in-text as “Cao, Fleming, and others”; entered previously as “Cao & Plenio (2025)”, which is a misattribution — this 18-author review is led by Jianshu Cao and does not include M. B. Plenio. The in-text reference and year should be corrected accordingly.]

Cha, Y., Murray, C. J., & Klinman, J. P. (1989). Hydrogen tunneling in enzyme reactions. *Science* 243: 1325-1330.

Duan, H.-G., Prokhorenko, V. I., Cogdell, R. J., Ashraf, K., Stevens, A. L., Thorwart, M., & Miller, R. J. D. (2017). Nature does not rely on long-lived electronic quantum coherence for photosynthetic energy transfer. *Proceedings of the National Academy of Sciences USA* 114: 8493-8498.

Engel, G. S., Calhoun, T. R., Read, E. L., Ahn, T.-K., Mančal, T., Cheng, Y.-C., Blankenship, R. E., & Fleming, G. R. (2007). Evidence for wavelike energy

transfer through quantum coherence in photosynthetic systems. *Nature* 446: 782-786.

Gunawardena, J. (2012). A linear framework for time-scale separation in non-linear biochemical systems. *PLOS ONE* 7: e36321.

Hiscock, H. G., Worster, S., Kattnig, D. R., Steers, C., Jin, Y., Manolopoulos, D. E., Mouritsen, H., & Hore, P. J. (2016). The quantum needle of the avian magnetic compass. *Proceedings of the National Academy of Sciences USA* 113: 4634-4639.

Hiscock, H. G., Hiscock, T. W., Kattnig, D. R., Scrivener, T., Lewis, A. M., Manolopoulos, D. E., & Hore, P. J. (2019). Navigating at night: fundamental limits on the sensitivity of radical pair magnetoreception under dim light. *Quarterly Reviews of Biophysics* 52: e9. DOI: 10.1017/S0033583519000076.

Quantum information, computing, and complexity

Aaronson, S. (2005). Quantum computing, postselection, and probabilistic polynomial-time. *Proceedings of the Royal Society A* 461: 3473-3482.

Aaronson, S. (2005). NP-complete problems and physical reality. *ACM SIGACT News* 36(1): 30-52. [Aaronson2005b]

Bennett, C. H., Bernstein, E., Brassard, G., & Vazirani, U. (1997). Strengths and weaknesses of quantum computing. *SIAM Journal on Computing* 26: 1510-1523. [Original BQP bounds.]

Bernstein, E., & Vazirani, U. (1997). Quantum complexity theory. *SIAM Journal on Computing* 26: 1411-1473. [Original definition of BQP.]

Google Quantum AI (2023). Suppressing quantum errors by scaling a surface code logical qubit. *Nature* 614: 676-681.

Nielsen, M. A., & Chuang, I. L. (2010). *Quantum Computation and Quantum Information*. Cambridge University Press.

Steinhauer, J. (2016). Observation of quantum Hawking radiation and its entanglement in an analogue black hole. *Nature Physics* 12: 959-965.

Mathematical foundations and computability

Birkhoff, G. (1946). Three observations on linear algebra. *Universidad Nacional de Tucumán Revista Series A* 5: 147-151. [Birkhoff-von Neumann theorem.]

Gödel, K. (1931). Über formal unentscheidbare Sätze der *Principia Mathematica* und verwandter Systeme I. *Monatshefte für Mathematik und Physik* 38: 173-198.

Turing, A. M. (1937). On computable numbers, with an application to the *Entscheidungsproblem*. *Proceedings of the London Mathematical Society* s2-42: 230-265.

von Neumann, J. (1953). A certain zero-sum two-person game equivalent to the optimal assignment problem. In: H. W. Kuhn & A. W. Tucker (eds.), *Contributions to the Theory of Games II*, pp. 5-12. Princeton University Press. [Birkhoff-von Neumann theorem.]

Historical and philosophical works

Cusa, N. of (1440). *De docta ignorantia (On Learned Ignorance)*.

Cusa, N. of (1453). *De visione Dei (On the Vision of God)*.

Deutsch, D. (2011). *The Beginning of Infinity: Explanations That Transform the World*. Viking Press.

Penrose, R. (2004). *The Road to Reality: A Complete Guide to the Laws of the Universe*. Jonathan Cape.

Spinoza, B. (1677). *Ethica, Ordine Geometrico Demonstrata (Ethics, Demonstrated in Geometrical Order)*.

Zuckerkandl, E., & Pauling, L. (1965). Evolutionary divergence and convergence in proteins. In: V. Bryson & H. J. Vogel (eds.), *Evolving Genes and Proteins*, pp. 97-166. Academic Press. [Origin of the molecular clock hypothesis.]