

The Incompleteness of Observation

Why the Universe's Biggest Contradiction Might Not Be a Mistake

Alex Maybaum — February 2026

The Problem

Physics has a contradiction it cannot resolve. Its two most successful theories — quantum mechanics and general relativity — flatly disagree about the most basic property of empty space: how much energy it contains.

Quantum mechanics says the vacuum is seething with energy. Add up the zero-point fluctuations of every quantum field and you get roughly 10^{113} joules per cubic meter — an unimaginably large number.

General relativity measures the vacuum's energy through its gravitational effect — the accelerating expansion of the universe — and gets about 6×10^{-10} joules per cubic meter. A very tiny number.

The ratio is roughly 10^{122} — conventionally rounded to 10^{120} in the physics literature, and by any accounting the largest disagreement between theory and observation in all of science. For context, the number of atoms in the observable universe is only about 10^{80} .

For decades, physicists have assumed something has gone badly wrong — that one or both calculations must contain an error, and that finding the mistake will lead to a "theory of everything" unifying quantum mechanics and gravity.

The opposite is true. **Neither calculation is wrong. They disagree because they're answering different questions about the same thing.** And they *have* to disagree, for a reason that has nothing to do with the specific physics involved.

The Two Calculations

Quantum mechanics calculates the vacuum's **fluctuation power spectrum**. It measures the absolute magnitude of zero-point activity for every quantum field. You add up all those active fluctuations — from the longest wavelengths to the shortest ones allowed by physics (the "Planck scale" cutoff) — and get roughly 10^{113} joules per cubic meter.

General relativity calculates the vacuum's **macroscopic expectation value**. It measures how the vacuum's energy actually pushes on spacetime (the stress-energy tensor). Work backwards from the observed cosmic acceleration and you get roughly 6×10^{-10} joules per cubic meter.

$$\frac{\rho_{\text{QM}}}{\rho_{\text{grav}}} \sim 10^{122}$$

They disagree because they are extracting *different statistical properties* from the same underlying reality.

The Core Idea: You Can't See Everything From Inside

Imagine you want to understand the energy of a calm glass of water. You have two ways to measure it:

A **thermometer** measures the total thermal fluctuation — every molecule's kinetic energy contributes positively, regardless of direction. The reading is enormous because nothing cancels out. This is a *variance-type* measurement.

A **suspended dust speck** (Brownian motion) reveals the net mechanical push the water exerts. At any given microsecond, millions of molecules strike from the left and millions from the right. Their impacts mostly cancel. The net push is just the tiny statistical residual left over. This is a *mean-type* measurement.

These aren't giving contradictory information. They're measuring fundamentally different statistical properties. The thermometer measures total unsigned activity (the variance). The dust speck measures the net, canceled-out push (the mean). For a system with trillions of molecules pushing in random directions, the total unsigned activity is naturally enormous compared to the tiny canceled-out residual.

Quantum mechanics and general relativity are just like the thermometer and the dust speck. Quantum mechanics measures the *fluctuation content* of the vacuum — the total, unsigned activity. General relativity measures the *net mechanical effect* — the aggregate, canceled-out push on spacetime. The 10^{122} ratio is the difference between an unsigned total and a canceled-out residual for a system with an astronomically large number of degrees of freedom.

Why This Isn't Just an Analogy

There's a deeper reason this works. In 2008, physicist David Wolpert proved a set of theorems showing that **any observer that is part of the system it's trying to measure faces irreducible limits on what it can know**. These limits don't depend on technology, intelligence, or computational power. They follow purely from the mathematical structure of being inside the thing you're measuring.

Wolpert's Framework

The observer has a mapping from the complete state of the universe to what they can access — a *projection*. The critical property is that this projection is **many-to-one**: many different complete universe states look the same through the observer's limited window.

Wolpert proved two key results:

(a) **The "Blind Spot" Theorem.** There is always at least one fact about the universe that the observer simply *cannot* determine — no matter how clever they are or how much computing power they have.

(b) **The "Mutual Inference" Impossibility.** If two methods use genuinely different projections to study the same target, they cannot fully reconstruct each other's conclusions.

These theorems aren't about technology limitations. They're about *logical structure* — in the same family as Gödel's incompleteness theorem and Turing's halting problem.

Crucially, in Wolpert's framework, an "inference device" doesn't have to be a conscious scientist or a supercomputer. **Spacetime itself is an inference device.** The local

gravitational field physically couples to the vacuum, and the expansion of the universe is the physical "record" of that coupling. Similarly, **localized matter** (like a hydrogen atom or a particle detector) is an inference device physically coupling to the vacuum's local fluctuations. Because both spacetime and localized matter are physical subsystems embedded *inside* the universe, they are bound by the exact same strict, mathematical limits of observation.

The Hidden Sector

What prevents the observer (or these physical inference devices) from accessing these hidden degrees of freedom? The answer is the **speed of light**. Nothing transmits information faster than light — a structural feature of spacetime, not a technological limitation. This creates a boundary around every observer, beyond which information cannot reach them.

The universe's degrees of freedom split into everything the observer *can* access and the **hidden sector** — everything they can't. The hidden sector isn't exotic. It consists of standard physics rendered inaccessible by spacetime's causal structure:

- **Beyond the cosmological horizon.** The universe is 13.8 billion years old and expanding. Light from beyond ~ 46 billion light-years has not had time to reach us. Beyond that boundary: standard matter governed by the same physics, but causally disconnected from us.
- **Inside black holes.** Every black hole's event horizon is a boundary beyond which the escape velocity exceeds the speed of light. Matter that crosses it exits our observable projection permanently.
- **Below the Planck scale.** Below about 10^{-35} meters, the energy required to probe would create a micro black hole, swallowing the information you're trying to extract.

Because our entire universe shares a common causal origin — the Big Bang and early cosmic inflation — the degrees of freedom that were pushed beyond these horizons are intimately correlated with the ones we can still see. They were once in direct causal contact, and exact conservation laws (energy, momentum, charge) permanently link them.

How much can an observer access? There is a fundamental result in physics (the Bekenstein-Hawking entropy) that sets the answer: the maximum information available to an observer is proportional to the *area* of their causal boundary. For our cosmological horizon, this works out to $\sim 10^{122}$ degrees of freedom in the visible sector. The hidden sector is astronomically larger, as we will calculate below.

Two Projections, Two Answers

The mathematical structure of quantum mechanics and gravity forces them to extract different statistics from this hidden sector.

Projection 1: Quantum Mechanics Measures Fluctuation Content

While the formal mathematics of QFT contains both positive and negative contributions, its operational reality — the phenomena that actually verify the vacuum's seething nature, like the Lamb shift or spontaneous emission — is driven by the vacuum's *fluctuation power spectrum*. This is an unsigned, variance-type measurement. Every active fluctuation contributes positively, no cancellation is possible, and the sum grows linearly with the number of modes.

Projection 2: Gravity Measures the Net Effect

Gravity is unique: it does not couple to the absolute fluctuation power. The Einstein field equations couple spacetime curvature to the macroscopic expectation value of the stress-energy tensor. Unlike the QFT fluctuation spectrum, this is a signed sum. Bosonic fields, fermionic fields, and vacuum condensates contribute with different signs.

This involves a substantive physical assumption. The hidden sector does not possess an unbroken symmetry (like perfect supersymmetry) that would align all contributions with a common sign. Under this assumption, the vast hidden sector acts like a high-dimensional, complex statistical system. A huge number of uncoordinated contributions with varying signs means the macroscopic residual (the mean) is naturally tiny compared to the total unsigned activity (the variance). Providing a first-principles argument for why hidden-sector contributions carry quasi-random signs remains an open problem.

The Observational Incompleteness Theorem

The Informal Version

The quantum-mechanical and gravitational descriptions of vacuum energy are structurally incompatible projections that cannot be unified into a single embedded description.

The two projections require **contradictory operations** on the hidden sector. The quantum projection works by *tracing out* the hidden sector — treating it as inaccessible. The gravitational projection works by *coupling to* it — feeling its mechanical presence. One hides it. The other feels it. No single description available to an embedded observer (or an embedded hardware device) can do both simultaneously.

The Formal Argument

Status note: The structural identification of QM and GR vacuum energies as complementary projections is well-defined, and the structural form of Wolpert's constraints is sufficient to establish the incompatibility. The specific quantitative bound remains an open problem, and the theorem's validity depends on the sign-structure assumption described above.

Step 1: Define the setup. Split the universe into visible and hidden sectors. Define two target functions: the total fluctuation content (a variance-type quantity) and the net mechanical effect (a mean-type quantity).

Step 2: Identify the observers as Wolpert inference devices. The **local gravitational field** is an embedded physical subsystem whose curvature dynamically infers the state of the vacuum's macroscopic mean. For quantum mechanics, the inference device isn't a human physicist — it is **localized matter** (like a hydrogen atom or a particle detector) whose physical state is perturbed by the absolute magnitude of local vacuum fluctuations. Both are physical "hardware" restricted by Wolpert's many-to-one setup functions.

Step 3: Check independent configurability. Wolpert's mutual inference impossibility requires that the two targets can be changed independently. The mean depends on the net sign balance, while the variance depends on absolute amplitudes. Unless the ultimate laws of physics rigidly and exactly lock the sign structure to the excitation spectrum across the entire universe, independent configurability holds.

Step 4: Apply Wolpert's structural constraint. Wolpert's framework establishes that two inference targets defined on the same system impose joint constraints on any embedded device — the device's own state must be consistent with both measurements simultaneously, creating a fundamental trade-off on joint accuracy. The specific quantitative product bound on the mean-squared errors of the variance-type and mean-type targets depends on the coupling structure and remains an open problem (identified as a priority in the formal paper). What the framework establishes is the structural claim: these are complementary projections of a shared sector, and their disagreement is a signature of the projection rather than an error in either theory.

The Ratio as Measurement

Because the 10^{122} discrepancy is a variance-to-mean ratio, we can work backwards to find the hidden sector's true size.

With N independent degrees of freedom, the quantum projection sums all contributions without regard to sign. If we look at the linear energy density scalings derived from these statistics, the variance-type scaling grows directly with the number of modes ($V \propto N$). The gravitational projection sums them with their signs. By the rules of statistical typicality for large, complex systems without unbroken symmetries, the residual macroscopic mean's linear scaling grows as the square root ($M \sim \sqrt{N}$). Their ratio is a function of \sqrt{N} alone:

$$\frac{V}{M} \sim \frac{N}{\sqrt{N}} = \sqrt{N}$$

Setting this equal to the observed value of cosmic acceleration:

$$\sqrt{N} \sim 10^{122} \implies N \sim 10^{244}$$

This is exactly S_{ds}^2 — the square of the Bekenstein-Hawking entropy of our cosmological horizon ($S_{\text{ds}} \sim 10^{122}$). Independently, physicist Rafael Sorkin derived $\Lambda \sim N^{-1/2}$ from causal set theory — the exact same functional form — before the 1998 discovery of cosmic acceleration.

The cosmological constant problem isn't a problem. It's the most precise measurement we have of the dimensionality of the parts of reality we cannot see.

Where Does Quantum Mechanics Come From?

The paper now makes a considerably stronger claim: quantum mechanics itself is a *consequence* of observational incompleteness. The two claims are logically independent. The Observational Incompleteness Theorem stands or falls on its own merits regardless of whether the derivation of quantum mechanics succeeds, and vice versa. If the first theorem is right but the second is wrong, the cosmological constant is still a measurement, not a mistake. If the second is right but the first is wrong, quantum mechanics still emerges from classical tracing-out, just without the specific cosmological constant interpretation.

When you "trace out" the hidden sector — mathematically discard the hidden degrees of freedom — the resulting description of what you *can* see has a very specific mathematical structure. It's not classical. It's not random noise. It's **quantum mechanics**.

The Barandes Stochastic-Quantum Correspondence

In 2023, Jacob Barandes proved that **any indivisible stochastic process is exactly equivalent to a quantum system**.

A stochastic process is **divisible** if you can break a long transition into independent shorter ones. An **indivisible** process cannot be decomposed this way — the system has *temporal memory* that can't be captured by the present state alone. Think of it like a conversation: if you walk in mid-sentence, you can't reconstruct what was said before from the current word alone — the history matters in a way that can't be faked.

Barandes proved that if a process is indivisible, it *automatically* reproduces interference, entanglement, the Born rule, and superposition. These emerge mathematically from indivisibility.

Independent Corroboration

This conclusion is reinforced by independent convergence: physicist Christof Wetterich has separately shown that when a subsystem possesses only "incomplete statistics" about a larger classical system, the necessary description of that subsystem automatically takes the form of the quantum formalism — the same conclusion reached here via a completely different mathematical route.

Why Tracing Out Produces Indivisibility

When you trace out part of a system, the remaining part's evolution acquires a **memory kernel** — encoding how the hidden sector's past states influence the visible sector's present. If the hidden sector has temporal correlations (perturbations persist rather than vanishing instantly), the memory kernel is nonzero and the visible sector's dynamics becomes non-Markovian — its future depends on its past, not just its present.

But there is a subtlety. Having memory is necessary but not sufficient. There is an intermediate regime — dynamics with memory that can still be decomposed into independent steps. Indivisibility is the stronger condition: the steps *themselves* fail, not just

the memorylessness. A generic mathematical trace-out doesn't guarantee you land in the indivisible regime rather than the intermediate one.

This is where the specific physics of our universe matters. The hidden sector isn't a generic mathematical abstraction — it has three concrete physical properties, each of which independently pushes the dynamics away from divisibility:

- **The boundary is maximally correlated.** Because the visible and hidden sectors share a causal origin — the Big Bang and early cosmic inflation — the degrees of freedom on either side of the horizon were once in direct causal contact. This shared history, combined with exact conservation laws (energy, momentum, charge), ensures the joint probability distribution over visible and hidden sectors is strongly non-factorizable. Tracing out a sector with such deep classical correlations produces extreme information backflow — the hidden sector's past rigidly constrains the visible sector's future.
- **The hidden sector is a fast scrambler.** Causal horizons scramble information across all their degrees of freedom exponentially fast. Any imprint the visible sector makes on the hidden sector is rapidly and chaotically distributed across $\sim 10^{244}$ degrees of freedom. A memory kernel generated by a fast-scrambling environment is densely correlated over time, resisting the clean decomposition into independent steps.
- **Conservation laws prevent reset.** Dynamics usually only becomes divisible when the environment can instantly "forget" each interaction. But the universe is a closed system with strict conservation laws. The hidden sector cannot reset to equilibrium after each interaction without violating those laws. Because it must remember where the energy went, the resulting dynamics cannot be memoryless.

Closing the Loop: The Mathematical Argument

The following argument is presented as the **Trace-Out Conjecture**, derived through a chain of established mathematical results with clearly stated assumptions.

First, by utilizing Random Matrix Theory, we can mathematically model the hidden sector's astronomically large 10^{244} degrees of freedom as a chaotic, high-dimensional system. Why random matrix theory? Because the hidden sector is a vast, complex system whose detailed structure is unknown to any embedded observer. In physics, when you have a large, complicated system whose specifics you can't know, random matrix theory gives the statistically universal answer — this was first discovered in nuclear physics (where it describes the energy levels of complex nuclei) and has since been proven to apply to broad classes of sufficiently complex systems. The specific choice of matrix ensemble (the "GOE") reflects the time-reversal symmetry of classical dynamics, but the key result doesn't depend on the details: any sufficiently large, generically coupled system gives the same conclusion. This rigorous statistical approach shows that the visible sector's transition probabilities are forced into a state of "full-rank covariance" — meaning that the dynamics are indivisible with probability 1. By Wigner-Dyson universality, this result extends to physically realistic Hamiltonians of sufficient complexity, though extending the rigorous proof to the full Liouvillian matrix structure (rather than just spectral statistics) remains an open mathematical problem.

Second, the argument explicitly bridges the gap to the continuous world of quantum mechanics. When we translate this history-dependent "memory kernel" into a continuous, present-moment equation, classical real numbers are mathematically insufficient to store the integrated history of the hidden sector. We are forced to introduce a complex phase (the "imaginary" part of a quantum wave function) to act as a mathematical storage buffer. The classical memory kernel seamlessly transforms into the complex phase gradient, naturally birthing the exact Schrödinger equation.

There is a known technical objection here. In 1994, physicist Timothy Wallstrom showed that the fluid-like equations underlying this derivation don't *quite* reproduce quantum mechanics on their own — they're missing a topological condition (single-valuedness of the wavefunction). In this framework, the objection is addressed by the fact that the discrete stochastic foundation (the indivisible stochastic matrices from the RMT argument) already carries the full quantum structure *before* the continuous fluid equations are ever invoked. The continuous Schrödinger equation is the continuum limit of dynamics whose quantum character was established at the discrete level first. The fluid representation is used only to *display* an already-quantum result, not to derive it from scratch. The Wallstrom objection is therefore substantially mitigated, but not fully eliminated for field theories — showing that the discrete-to-continuous passage preserves the required topological structure in infinite-dimensional field-theoretic settings remains an open mathematical problem.

The Complete Chain

1. You're inside the universe (embedded observer)
2. Some degrees of freedom are inaccessible (hidden sector)
3. Your description discards them (trace out)
4. The trace-out generates a history-dependent memory kernel
5. The extreme complexity of the hidden sector forces that memory into the indivisible regime with probability 1 (established via Random Matrix Theory and extended by Wigner-Dyson universality)
6. Indivisible stochastic processes *are* quantum systems (Barandes)
7. The classical memory kernel requires a complex phase representation, directly yielding the Schrödinger equation (with the Wallstrom gap closed for finite-dimensional systems and substantially mitigated for field theories)

Quantum mechanics isn't a fundamental law — it's what any embedded observer sees after tracing out a maximally correlated, fast-scrambling hidden sector constrained by conservation laws.

Explaining the Quantum World

The counterintuitive phenomena of quantum mechanics all have natural readings — not as irreducible mysteries, but as features of what an embedded observer sees.

Interference. The double-slit pattern arises because the particle’s journey involves hidden-sector degrees of freedom that retain correlations between passage and arrival. Adding a detector forces the hidden sector to relinquish that information, disrupting the correlations and destroying the pattern.

Superposition and measurement. Superposition is what incomplete information looks like when the incompleteness has the mathematical structure of indivisibility. The full state of the universe is perfectly definite — your projected view of it is irreducibly fuzzy.

Entanglement. When two entangled particles are measured, the correlation arises because the hidden sector mediates correlations between them — like two thermometers reading the same temperature because they’re immersed in the same water. No information travels between the particles.

Can We Test This?

This framework isn’t just a philosophical reinterpretation; it makes concrete physical predictions that can be falsified or confirmed by future astronomical observations. Because General Relativity is treated as an effective ”mean-field” theory, its smooth descriptions must eventually break down when pushed to their limits:

Gravitational Wave Echoes. When two black holes merge, the extreme disruption should interact with the granularity of the hidden sector. The hidden sector acts as a ”fast scrambler” — a system that redistributes information as quickly as physics allows. The timescale for this scrambling depends logarithmically on the sector’s entropy. Taking the hidden-sector entropy as $\sim 10^{122}$ (consistent with the holographic bound on our cosmological horizon), and a 30 solar-mass post-merger remnant, the echo delay works out to roughly 0.00008 seconds — placing it in the 10^{-5} to 10^{-4} second range that current LIGO/Virgo/KAGRA instruments can already probe in post-merger data. The echo should be faint (its amplitude is an open problem), but its timescale is a hard prediction.

A Universal Noise Floor. The irreducible fluctuations of the hidden sector must source a stochastic gravitational wave background. However, a naive dimensional estimate produces a problematic result — an energy density of order the critical density of the universe, which would violate existing observational constraints from LIGO/Virgo, pulsar timing arrays, and Big Bang nucleosynthesis. The naive estimate must therefore be an overestimate, likely failing due to spectral dilution across ~ 61 decades of frequency and quadrupole coupling suppression for sub-Planckian modes. Computing the precise spectral shape $\Omega_{\text{gw}}(f)$ across the MHz–GHz band (where astrophysical foregrounds are minimal) is the most pressing empirical test of the framework. If the spectral decomposition, once computed, cannot be reconciled with observational upper bounds for *any* physically reasonable hidden-sector mode structure, the framework is falsified.

What This Means

The century-long search for a unified theory has been asking the wrong question. It’s like asking for a single instrument that simultaneously measures both temperature and

pressure by being a thermometer and a barometer at the exact same time. The request is structurally impossible.

This also elegantly explains why the most famous mathematical unification of quantum mechanics and gravity — the AdS/CFT correspondence — works flawlessly. AdS/CFT is formulated by placing an observer on the *asymptotic boundary* of a hypothetical universe, looking inward. Because they are on the outside looking in, they are not an embedded observer, and Wolpert's limits don't apply. But our actual universe is expanding (de Sitter space); it has no outer boundary. We are permanently embedded inside it.

The universe is not broken. We are just observing it from within, which sets fundamental limits on our ability to unify certain projections of reality. The remaining open problems — particularly the quantitative Wolpert bound, the sign structure of the hidden sector, and the rigorous universality extension beyond the GOE ensemble — define the path toward complete closure.

In 1926, Einstein wrote to Max Born: "I, at any rate, am convinced that He does not throw dice." For a century, this has been read as Einstein being wrong. There is a more sympathetic reading. The full state of the universe, including its hidden sector, *is* definite. The dice are real, but they belong to the projection, not to reality itself. What Einstein called "the secret of the Old One" is not randomness. It is the structural fact that no observer inside the universe can see the whole game.

This is a simplified overview of "The Incompleteness of Observation: Why Quantum Mechanics and General Relativity Cannot Be Unified From Within" (Maybaum, February 2026), which presents the formal arguments with detailed derivations and experimental predictions.

Acknowledgment of AI-Assisted Technologies: The author acknowledges the use of **Claude Opus 4.6** and **Gemini 3.1 Pro** to assist in synthesizing technical concepts and refining clarity. The final text and all scientific claims were reviewed and verified by the author.