

Structural Consequences of Observational Incompleteness

From Chemistry to Evolution

Author: Alex Maybaum

Date: April 2026

Status: DRAFT PRE-PRINT

Classification: Theoretical Physics / Foundations / Astrobiology

Abstract

The Observational Incompleteness (OI) framework derives the Standard Model’s gauge structure from a single definition of embedded observation [1, 2]. We trace the consequences of this derived structure from atomic physics through biological evolution. The structural chain is entirely parameter-free: $d = 3$ determines the orbital algebra $SO(3)$, the periodic table, carbon’s tetrahedral bonding geometry, the nuclear-atomic scale hierarchy, water’s solvent properties, and the thermal window. Three generations guarantee CP violation; the Higgs guarantees a mass hierarchy; chiral $SU(2)$ produces a parity-violating energy difference whose sign is fixed by the partition. The viable fraction of the 18-dimensional parameter space is estimated at $\sim 16\%$.

The combinatorial diversity of carbon chemistry ($\sim 20^N$ peptides of length N) exceeds both Kauffman’s autocatalytic threshold and von Neumann’s self-reproducing automaton threshold by $\sim 10^{56}$. Template replication follows from three independently structural capabilities (linear polymers, complementary pairing, catalysis). The origin of life is identified as the first molecular system to satisfy C1–C3: RNA is the unique molecule that simultaneously serves as template (C2, C3) and catalyst (C1), making the RNA world a structural prediction. C1–C3 systems are exponential-growth attractors in chemical space; the transition from prebiotic chemistry to life is inevitable and irreversible. Darwinian evolution follows from imperfect template replication. The chain has a natural terminus at biological evolution; the continuation from evolution through general intelligence, AI, the BQP characterization theorem for embedded observers, and the framework’s position on the P vs NP question is the subject of the companion paper *Embedded Computation in the Incompleteness of Observation Framework* [38].

The same C1–C3 conditions that produce QM cosmologically are satisfied in proteins ($\tau_S/\tau_B \sim 10^{-9}$ to 10^{-12}), predicting non-Markovian enzyme kinetics consistent with single-molecule experiments. The non-Markovian corrections grow as τ_S/τ_B increases, approaching engineering relevance in quantum computing (10^{-3}) and quantum sensing (10^{-6}). The framework identifies a new design space — engineered C1–C3 architectures — in which the environment is

programmable quantum memory rather than noise.

Status — conjectural extension. This paper applies the core framework to quantum computation and complexity, and is offered as a *conjecture*, not an established result: the central claim — that the framework’s emergent dynamics realizes the complexity class BQP — is a theoretical conjecture rather than a proven correspondence; it would be supported by a rigorous derivation from the finite substratum to BQP and undercut if that dynamics proves efficiently classically simulable or fails to reproduce known quantum-complexity separations. It carries no evidential weight for the core framework — which stands or falls independently on its physics — and should be read as a structural possibility to be tested, not as support for the framework.

1. Introduction

The fine-tuning problem in physics rests on a counterfactual: if the laws of physics or the values of fundamental constants were different, the universe would not permit complex structures, and we would not exist to observe it. The anthropic principle elevates this observation into an explanatory principle — our existence selects the laws from an ensemble of possibilities.

The OI framework [1, 2] challenges the premise. The companion papers prove that an embedded observer in a deterministic system with conditions C1–C3 (coupling, slow bath, sufficient capacity) necessarily describes the visible sector using quantum mechanics [1], and that the dynamics uniquely selected by this requirement determines the Standard Model’s complete structure [2]. If the laws are derived rather than contingent, the largest source of anthropic variation — the structure of the laws themselves — is eliminated.

This paper traces the consequences. §2 establishes the structural chain from (S,) to the full preconditions for organic chemistry — orbital structure, the periodic table, carbon’s bonding geometry, the nuclear-atomic scale hierarchy, water’s solvent properties, and the thermal window for dynamic chemistry. §3 estimates the viable fraction of the parameter space using existing literature. §4 derives the chirality chain from the partition to biological homochirality. §5 extends the chain through autocatalytic networks, template replication, the origin of life (identified as the first molecular C1–C3 system), and Darwinian evolution — closing with a forward-pointer to the companion paper [38] that continues the chain through general intelligence, AI, and the BQP characterization for embedded observers. §6 addresses the existence question. §7 applies C1–C3 to biological systems, predicting non-Markovian enzyme kinetics with implications for drug design. §8 identifies where OI corrections reach engineering relevance — quantum computing, quantum sensing, and precision metrology — and proposes three engineering frontiers: engineered partitions (the bath as

programmable memory), non-Markovian quantum computation (P-indivisible gate sequences), and quantum materials with designed C1–C3 architectures.

2. The Structural Chain

2.1 From (S,) to orbitals

The framework derives $d = 3$ [2, §3.2] from three independent self-consistency filters (propagating gravity, stable matter, and the boundary entropy concordance), with renormalizability of the emergent Yang-Mills coupling as a downstream consequence rather than an independent filter. In three spatial dimensions, the rotation group is $SO(3)$, whose irreducible representations are labeled by angular momentum quantum number $l = 0, 1, 2, 3, \dots$ with $(2l + 1)$ states each. These are the atomic orbitals: s ($l = 0$, spherical, 1 orientation), p ($l = 1$, dumbbell, 3 orientations), d ($l = 2$, 5 orientations), f ($l = 3$, 7 orientations).

This is the *only* orbital algebra consistent with embedded observation. In $d = 2$, angular momentum has a single component — only circular harmonics exist, with no three-dimensional bonding geometry. In $d \geq 4$, the Coulomb potential $V(r) \sim r^{2-d}$ does not support stable bound states [3, 4]: classical orbits spiral into the nucleus, and the quantum Hamiltonian is unbounded below.

2.2 From orbitals to the periodic table

The framework derives spin-1/2 fermions from the staggered structure on the $d = 3$ lattice [2, §4.2]. The spin-statistics theorem — a consequence of the emergent QFT’s Lorentz invariance [2, §3.1] — gives the Pauli exclusion principle: each orbital holds at most 2 electrons (spin up and spin down). The shell capacities follow:

$$s : 2, \quad p : 6, \quad d : 10, \quad f : 14$$

The periodic table’s period lengths — 2, 8, 8, 18, 18, 32 — are determined by the filling order of these shells. This architecture is entirely structural: it depends on $d = 3$, spin-1/2, and the Pauli principle, all of which are derived.

2.3 Carbon’s bonding geometry

Element 6 (carbon) has electron configuration $1s^2 2s^2 2p^2$ — four valence electrons. In $d = 3$, four valence electrons admit three hybridization modes:

sp³ hybridization: Four equivalent orbitals pointing to the vertices of a regular tetrahedron (bond angle 109.5°). This is the geometry of methane (CH₄), diamond, and the backbone of all organic macromolecules — proteins, nucleic acids, lipids, carbohydrates.

sp² hybridization: Three planar orbitals at 120° plus one bond. This is the geometry of graphite, benzene, and aromatic chemistry.

The *possibility* of these hybridizations — and hence the possibility of three-dimensional macromolecular architecture — is a structural consequence of the angular momentum algebra in $d = 3$.

2.4 Beyond carbon: the extended structural chain

Four additional structural results extend the chain beyond orbital chemistry.

Matter-antimatter asymmetry. Three generations [2, §4.7] give the CKM matrix with a physical CP-violating phase — one of the three Sakharov conditions for baryogenesis [5]. The partition breaks P [2, §4.8, Theorem 13]. Combined with CP violation, baryogenesis is structurally possible. Without three generations, the CKM matrix has no CP-violating phase and the standard electroweak baryogenesis mechanism does not operate.

Mass hierarchy. The Higgs mechanism [2, §4.7] gives fermion masses through Yukawa couplings to a single doublet. The taste-breaking mechanism [2, §4.7] produces generation-dependent couplings generically, ensuring different generations have different masses. This guarantees the existence of light fermions (electron, u and d quarks) — not as a coincidence, but as a structural feature of the staggered lattice.

Nuclear binding. SU(3) color confinement [2, §4.6] produces bound hadrons. The existence of multi-nucleon bound states (nuclei) requires specific parameter values (§3), but the *mechanism* — color confinement binding quarks into hadrons, residual strong force binding hadrons into nuclei — is structural.

Chiral chemistry. Chiral SU(2) [2, §4.8] produces electroweak parity violation in all weak processes. This propagates to molecular physics as a parity-violating energy difference between mirror-image molecules (§4).

2.5 The scale hierarchy

The framework derives two independent gauge groups — SU(3) and U(1) — with independent coupling strengths corresponding to independent eigenvalues of the coupling matrix M [2, §4.6]. This guarantees two widely separated energy scales:

$$E_{\text{nuclear}} \sim \Lambda_{\text{QCD}} \sim \text{MeV}, \quad E_{\text{atomic}} \sim \alpha^2 m_e c^2 \sim \text{eV}$$

The ratio $E_{\text{nuclear}}/E_{\text{atomic}} \sim 10^6$ is partly parametric (it depends on the specific couplings), but the *existence* of two independent scales is structural — it follows from having two independent gauge groups. The universal induced coupling $1/\alpha_0 = 23.25$ at the Planck scale is itself structural — determined by $N_f = 6$ and the Dynkin index, not by the eigenvalues of M — and the framework reproduces

all three SM gauge couplings at M_Z through non-perturbative gauge self-energy corrections and RG running [2, §6]. This separation has a consequence: nuclei are stable against chemical reactions. Chemistry rearranges electrons (eV) without touching nuclei (MeV). Atoms are permanent building blocks that can be assembled, disassembled, and reassembled without being destroyed.

2.6 Water

Element 8 (oxygen): configuration $1s^2 2s^2 2p$ — six valence electrons, two unpaired. The sp^3 -like hybridization in $d = 3$, distorted by two lone pairs, gives a bond angle of $\sim 104.5^\circ$. This is the geometry of water (H O).

Water’s anomalous properties follow from this geometry. The bent structure creates a permanent dipole moment. The lone pairs enable hydrogen bonding — a secondary interaction (~ 0.2 eV) intermediate between covalent bonds (~ 3 eV) and thermal energy (~ 0.025 eV at 300K). Hydrogen bonding produces: high specific heat (thermal buffering for chemical reactions), density maximum near 4°C (ice floats, insulating liquid water below), and exceptional solvent capability (dissolving ionic and polar molecules for solution-phase chemistry).

The *possibility* of a molecule with these properties is structural: element 8’s electron configuration in $d = 3$ with the derived orbital algebra guarantees a bent dihydride with lone pairs capable of hydrogen bonding. The *existence* of a liquid phase at specific temperatures is parametric.

2.7 The thermal window

For chemistry to support complexity, a thermal window must exist where bond energies ($\sim 1\text{--}5$ eV) $\gg kT$ (bonds are stable), $kT > E_{\text{activation}}$ (reactions proceed), and a liquid solvent exists. The scale hierarchy (§2.5) guarantees that $E_{\text{bond}}/E_{\text{thermal}} \gg 1$ at any temperature where a liquid exists — chemistry is inherently *stable but dynamic* in any universe with two independent gauge groups and the derived orbital structure.

2.8 Summary of the structural chain

$(S, \varphi) \xrightarrow{\text{derived}} d = 3 \xrightarrow{\text{SO}(3)} \text{orbitals} \xrightarrow{\text{Pauli}} \text{periodic table} \xrightarrow{\text{element 6}} \text{tetrahedral carbon}$

$(S, \varphi) \xrightarrow{\text{derived}} 3 \text{ generations} \xrightarrow{\text{CKM}} \text{CP violation} \xrightarrow{\text{Sakharov}} \text{baryogenesis possible}$

$(S, \varphi) \xrightarrow{\text{derived}} \text{Higgs} \xrightarrow{\text{taste-breaking}} \text{mass hierarchy} \rightarrow \text{light fermions} \rightarrow \text{stable atoms}$

$$(S, \varphi) \xrightarrow{\text{derived}} \text{SU}(3) \xrightarrow{\text{confinement}} \text{hadrons} \xrightarrow{\text{residual force}} \text{nuclei possible}$$

$$(S, \varphi) \xrightarrow{\text{derived}} \text{SU}(3) + \text{U}(1) \text{ independent} \rightarrow \text{scale hierarchy} \rightarrow \text{atoms are permanent building blocks}$$

$$(S, \varphi) \xrightarrow{\text{derived}} \text{element 8 in } d = 3 \rightarrow \text{bent H}_2\text{O} \xrightarrow{\text{H-bonding}} \text{anomalous solvent}$$

$$(S, \varphi) \xrightarrow{\text{derived}} \text{scale hierarchy} + \text{orbital structure} \rightarrow \text{thermal window exists}$$

$$(S, \varphi) \xrightarrow{\text{derived}} \text{chiral SU}(2) \xrightarrow{\text{PVED}} \text{L-amino acid preference}$$

Every arrow is either a theorem from [1, 2] or a standard textbook result applied to the derived structure. No parameter values enter. The possibility of a universe with atoms, nuclei, a periodic table, organic chemistry, matter-antimatter asymmetry, and chirally selected molecules is a theorem about embedded observation.

3. The Viable Parameter Space

3.1 The question

The structural chain establishes that organic chemistry is *possible*. For it to be *actual*, the 18 free parameters of the SM must fall in a viable region. How large is this region?

Traditional fine-tuning studies [6, 7, 8] vary both the structure and the parameters. The OI framework fixes the structure, reducing the question to parameters alone. This sharpens the problem considerably: the gauge group, dimension, generation count, and discrete symmetries are no longer free variables.

3.2 The most constrained parameters

Not all 18 parameters are equally constrained. The literature identifies four as critical for complex chemistry:

The fine-structure constant . Governs atomic binding, molecular bond strengths, and the balance between electromagnetic and thermal energies. Stable atoms require $\alpha \lesssim 1$ (above this, relativistic effects destabilize inner shells).

Complex chemistry requires bond energies $\gg kT$ at habitable temperatures, giving $\alpha \gtrsim 10^{-3}$ [8]. The viable range spans roughly two orders of magnitude: $10^{-3} \lesssim \alpha \lesssim 10^{-1}$, with our universe at $\alpha \approx 1/137 \approx 7.3 \times 10^{-3}$ — within the viable range but not near its center.

The electron-to-proton mass ratio m_e/m_p . Determines the separation of nuclear and atomic scales. Stable atoms with a rich periodic table require $m_e/m_p \ll 1$ (so electrons orbit far from the nucleus). Complex chemistry breaks down for $m_e/m_p \gtrsim 10^{-1}$ (electron clouds too compact for directional bonding). The viable range is roughly $10^{-5} \lesssim m_e/m_p \lesssim 10^{-1}$, spanning four orders of magnitude. Our value: $m_e/m_p \approx 5.4 \times 10^{-4}$.

The quark mass ratio $m_d - m_u$. Controls the proton-neutron mass difference and hence nuclear stability. If $m_d - m_u$ changes sign (proton heavier than neutron), hydrogen is unstable — catastrophic for chemistry. If $m_d - m_u$ is too large, no bound nuclei beyond hydrogen exist. The viable range is approximately $1 \lesssim (m_d - m_u)/\text{MeV} \lesssim 5$ [9], about a factor of 5. Our value: $(m_d - m_u) \approx 2.5$ MeV.

The strong coupling α_s . At the QCD scale (~ 1 GeV), α_s determines nuclear binding. If α_s is $\sim 10\%$ smaller, the deuteron is unbound and big bang nucleosynthesis fails to produce helium — but stellar nucleosynthesis could still produce carbon [8]. If α_s is $\sim 30\%$ larger, di-protons are bound, stars burn through hydrogen catastrophically fast. The viable range is roughly $0.8\alpha_s^{\text{obs}} \lesssim \alpha_s \lesssim 1.3\alpha_s^{\text{obs}}$, about a factor of 1.6 [10].

3.3 The viable fraction

The four critical parameters span roughly:

Parameter	Viable range (orders of magnitude)	Log-fraction
	~ 2 orders within ~ 3 available	$\sim 60\%$
m_e/m_p	~ 4 orders within ~ 6 available	$\sim 65\%$
$m_d - m_u$	factor of ~ 5 within ~ 10 available	$\sim 50\%$
α_s	factor of ~ 1.6 within ~ 2 available	$\sim 80\%$

The joint viable fraction, treating the parameters as independent (a conservative assumption — correlations generally enlarge the viable region), is roughly:

$$f_{\text{viable}} \sim 0.6 \times 0.65 \times 0.5 \times 0.8 \approx 16\%$$

This is a rough estimate, but the order of magnitude is robust: the viable fraction is $\mathcal{O}(10\%)$, not $\mathcal{O}(10^{-50})$. The remaining 14 SM parameters (heavy quark masses, lepton masses, mixing angles) are less constrained — they affect the details of nuclear and atomic physics but not the existence of carbon chemistry.

3.4 Comparison to fine-tuning claims

Traditional fine-tuning arguments quote numbers like 10^{-120} (for the cosmological constant) or 10^{-50} (for the Higgs mass). These numbers arise from varying the *structure* — asking what happens if the gauge group is different, or if there are different numbers of generations, or if $d \neq 3$. The OI framework eliminates all of these variations. What remains is a 16% viable fraction in the critical parameter subspace — not fine-tuned by any reasonable standard.

4. The Chirality Chain

4.1 From the partition to parity violation

The OI partition breaks parity: the trace-out over the hidden sector treats left-handed and right-handed spinor components asymmetrically ($D_{LL} = 0$, $D_{RR} \neq 0$) [2, §4.8, Theorem 13]. This is structural — any observer in any (S,) sees a chiral weak force. The result: the W and Z bosons couple only to left-handed fermions, and parity is maximally violated in all weak processes.

4.2 Molecular parity violation

The weak neutral current (Z boson exchange between electrons and nuclei) produces a parity-violating potential in atoms and molecules. For chiral molecules — molecules that are not superimposable on their mirror image — this creates an energy difference between left-handed (L) and right-handed (D) enantiomers [11, 12]:

$$\Delta E_{\text{PV}} \sim G_F \alpha Z^5 m_e c^2 \left(\frac{a_0}{R} \right)^3$$

where G_F is the Fermi constant, Z the nuclear charge, and R the bond length. For amino acids: $\Delta E_{\text{PV}} \sim 10^{-14}$ eV $\sim 10^{-17} kT$ at 300 K [11]. The sign is universal — it favors L-amino acids — and is fixed by the partition: the same partition that makes SU(2) left-handed makes L-amino acids lower in energy.

4.3 The amplification argument

The energy difference $10^{-17} kT$ is far too small to directly select one handedness. The enantiomeric excess at thermal equilibrium would be $ee \sim 10^{-17}$. The question is whether amplification mechanisms can promote this bias to complete homochirality.

Three mechanisms are experimentally demonstrated. Autocatalytic amplification (the Soai reaction [13]) amplifies $ee \sim 10^{-5}$ to $> 99\%$. Crystallization (Viedma ripening [14]) drives racemic crystal mixtures to homochirality from

small initial biases. Chiral polymerization with cross-inhibition [15] amplifies small monomer ee to polymer homochirality.

The gap between 10^{-17} (PVED) and 10^{-5} (Soai threshold) is 12 orders of magnitude, and bridging it from PVED alone is non-trivial. In a small prebiotic pool of N chiral molecules, random statistical fluctuations produce $ee \sim 1/\sqrt{N}$. For $N \sim 10^6$ (a microdroplet), statistical $ee \sim 10^{-3}$ — within the Viedma threshold. The PVED’s structural role is to fix the *sign* of the energy difference between L and D enantiomers; how this sign propagates to biological homochirality through the amplification mechanisms cited above is an active research question in origins-of-life. With current PVED magnitude ($\sim 10^{-17}$ ee) versus statistical fluctuations in prebiotic pools ($\sim 10^{-3}$ ee for $N \sim 10^6$ monomers), a single-pool calculation gives an L-preference probability of only $\sim 10^{-14}$ above 50%. Delivering global biological homochirality from this requires either (i) cross-pool selection mechanisms that propagate small per-pool L-biases globally, or (ii) extended timescales of prebiotic chemistry with cumulative bias amplification, or (iii) other mechanisms not yet fully characterized in the literature [13, 14, 15]. The framework’s structural result is the *sign* of the PVED; the bridge to biological homochirality requires additional empirical mechanisms that remain partially open.

4.4 The structural conclusion

If the amplification chain is robust — and every step is either derived, known physics, or experimentally demonstrated — then the *sign* of the molecular parity-violating energy difference is fixed by the partition (Theorem 13 of [SM §4.8]); biological homochirality plausibly inherits this sign through amplification mechanisms whose details remain partially open in origins-of-life research (§4.3). The handedness of every amino acid in every protein in every organism is *consistent* with the structural sign result; whether it is *uniquely traceable* to (S,) depends on which amplification mechanism actually operated in Earth’s prebiotic chemistry.

Step	Content	Status
1	Partition breaks P	Theorem [2, §4.8]
2	SU(2) _L → electroweak PV	Structural consequence
3	Electroweak PV → PVED	Known physics [11, 12]
4	Statistical fluctuation + directional bias → homochirality	Demonstrated mechanisms [13, 14, 15]; sign from Step 1
5	Biological homochirality	Observed

5. Implications

5.1 The dissolution of fine-tuning

The fine-tuning problem assumes two freedoms: the laws could have been different, and the parameters could have been different. The OI framework removes the first — the structure is derived. §3 shows the second is not as constrained as traditionally claimed: the viable parameter fraction is $\mathcal{O}(10\%)$, not $\mathcal{O}(10^{-50})$. The apparent fine-tuning was dominated by structural variation that the framework proves doesn’t exist.

The substratum gauge group $\mathcal{G}_{\text{sub}}$ [33, §4] makes this precise. The structural chain of §2 — $d = 3$, orbitals, the periodic table, carbon’s tetrahedral bonding, water, the thermal window, chirality — is invariant under the generators of $\mathcal{G}_{\text{sub}}$: state relabeling, alphabet change, deep-sector enlargement, and graph isomorphism up to statistical isotropy (as specified in [Substratum §4, Theorem 24(iv)]: for the gauge sector, restricted to isomorphisms preserving the cubic octahedral structure). These preconditions for complexity are not contingent features that happen to hold in our universe. They are gauge-invariant properties of the equivalence class $[(S, \varphi)]/\mathcal{G}_{\text{sub}}$ — they hold for every (S, φ) in the class. The only quantities that could have been different are the solution-specific parameters (gauge couplings, fermion masses, mixing angles), and §3 shows these have a viable fraction of $\mathcal{O}(10\%)$.

5.2 The dissolution of the anthropic principle

The anthropic principle assumes there is an ensemble of possible universes from which observation selects. The framework removes the ensemble: the structure is the only structure consistent with embedded observation (characterization theorem [1, §3.4]; reconstruction theorem [33, §3]). The preconditions for organic chemistry — atoms, carbon bonding geometry, the scale hierarchy, water, the thermal window, chiral selection — are structural (§2). The parameters have a viable fraction of $\mathcal{O}(10\%)$ (§3). There is nothing for observation to select, because there is nothing that could have been otherwise.

A closely related concern — the Boltzmann-brain problem, where thermal fluctuations in a long-lived finite universe should overwhelmingly produce observers with random memories — is also dissolved at the framework level, by a theorem rather than by an anthropic filter. The structural observer-selection theorem [1, §4.6] shows that C1–C3-satisfying observers cannot exist in the equilibrium phase of (S, φ) : the hidden-sector correlation time τ_B is bounded by the local mixing time on equilibrium microstates, and C2 fails there. Boltzmann brains are by construction local fluctuations in a thermally equilibrated environment, so their τ_B is bounded by local thermal decorrelation and is at least six orders of magnitude too short for C2. They do not see emergent quantum mechanics; they are not observers in the framework’s sense. The “past hypothesis” of low initial entropy — usually imported as a separate axiom to explain why we observe a low-entropy universe — is replaced by the same theorem, since observers

necessarily exist in the non-equilibrium phase.

The completeness of $\mathcal{G}_{\text{sub}}$ [33, Theorem 24] closes this argument: the four generator families are proved to exhaust all observables-preserving transformations (via Stinespring uniqueness applied at the substratum level). No finer distinction between substrata is empirically accessible. The structural preconditions for complexity are therefore as immutable as the gauge group itself — they are not drawn from an ensemble but fixed by the equivalence class.

5.3 The autocatalytic argument

The structural chain derives the preconditions for life. This section argues that the transition from prebiotic chemistry to self-replication is not merely possible but statistically expected, given the derived structure.

Combinatorial diversity is structural. Carbon’s sp^3 hybridization in $d = 3$ (§2.3) produces chain-forming chemistry with four bonding directions. The number of distinct molecules constructible from C, H, O, N grows exponentially with chain length: $\sim 20^N$ for amino acid sequences, $\sim 4^N$ for nucleotide sequences. For chains of length 50: $\sim 10^{65}$ possible peptides, $\sim 10^{30}$ possible RNA sequences. This combinatorial explosion is structural — it follows from the orbital algebra in $d = 3$ with four valence electrons. No other element in the derived periodic table produces comparable diversity: silicon has similar valence but weaker bonds (~ 2.3 eV for Si-Si vs. ~ 3.6 eV for C-C), insufficient for stable long chains at thermal window temperatures.

The autocatalytic threshold. Kauffman’s theorem [16]: in a network of N molecule types where each molecule has probability p of catalyzing any given reaction, autocatalytic sets — closed networks that collectively catalyze their own production from simple feedstocks — emerge with probability approaching 1 when $N \times p$ exceeds a critical threshold of $\mathcal{O}(1)$. The framework provides $N \sim 10^{65}$ (structural). The question reduces to whether $p > 10^{-65}$ — whether catalysis is not astronomically rare among organic molecules.

Catalysis is not rare. Empirically, short random peptides show catalytic activity at rates suggesting $p \sim 10^{-9}$ to 10^{-7} per pair [17]. RNA fragments catalyze their own splicing and polymerization [18, 19]. Simple amino acids catalyze aldol reactions (the Hajos-Parrish reaction). Even individual metal ions in aqueous solution show catalytic activity for a range of organic reactions. With $p \sim 10^{-9}$ and $N \sim 10^{65}$: $N \times p \sim 10^{56}$, exceeding Kauffman’s threshold by 56 orders of magnitude.

The role of chirality. The PVED (§4) selects L-amino acids, halving the combinatorial space and — critically — eliminating cross-chiral parasitic reactions. In a racemic mixture, L-monomers can be incorporated into D-chains and vice versa, poisoning both. Homochirality, driven by the PVED’s directional bias (§4.3), ensures that the combinatorial exploration proceeds within a self-consistent chemical subspace. This is not a minor refinement — it may

be the difference between a productive autocatalytic network and a parasitized one [15].

What’s structural vs. parametric.

Element	Status
Combinatorial diversity ($N \sim 20^L$)	Structural (carbon in $d = 3$)
Energy-driven exploration	Structural (thermal window, §2.7)
Chirality eliminates cross-inhibition	Structural (PVED sign, §4)
Catalytic probability p	Parametric (depends on bond energies, activation barriers)
$p > 1/N$	Empirically satisfied by ~56 orders of magnitude

The framework derives N (structural), the thermal window (structural), and the chiral bias (structural). The catalytic probability p is parametric but empirically enormous relative to the threshold. The autocatalytic emergence of self-replication is therefore not a lucky accident but a statistical consequence of the combinatorial diversity that carbon chemistry in $d = 3$ necessarily produces.

5.4 Self-replication as read-write cycling

The structural parallel. The framework derives QM from read-write cycling: the observer reads the visible sector and writes correlations into the hidden sector through the coupling H_{VB} ; the hidden sector stores those writes and returns them on subsequent reads; the resulting statistics are quantum mechanics [33, §2]. Template replication is the same structural pattern at the molecular level: a polymer reads its own sequence (through complementary hydrogen bond pairing) and writes a copy (through catalytic polymerization). Self-replication is not a special biological capability — it is molecular read-write cycling, the chemical instantiation of the same information-theoretic pattern that produces quantum mechanics.

Von Neumann’s threshold. Von Neumann’s self-reproducing automaton theorem [20] establishes that in any computational system above a complexity threshold, self-reproducing configurations must exist. The threshold requires three capabilities: (a) a universal constructor — a subsystem that reads instructions and assembles a product; (b) a copying mechanism — that duplicates the instructions; (c) a control mechanism — that sequences construction and copying.

The derived chemistry provides all three:

(a) *Universal construction.* Any catalytic polymer that reads a template sequence and assembles a corresponding product from a monomer alphabet. The autocatalytic argument (§5.3) establishes that catalysis is statistically abundant

($N \times p \sim 10^{56}$). Template-directed assembly follows from complementary pairing (structural, §2.6) combined with catalytic activity.

(b) *Instruction copying.* Template replication — a polymer directing the synthesis of its complement through complementary pairing. Three independently structural capabilities whose intersection is template replication (§5.5, below).

(c) *Control.* The thermal window (§2.7) provides environmental cycling — temperature fluctuations, wet-dry cycles, concentration gradients — that drives alternation between construction and copying phases. This is structural: the thermal window guarantees that the environment fluctuates on timescales relevant to chemistry.

Von Neumann’s threshold is exceeded by the same combinatorial margin ($\sim 10^{56}$) that drives the autocatalytic argument. Self-reproducing molecular configurations are not a lucky accident in the derived chemistry — they are a mathematical consequence of the system’s computational richness exceeding von Neumann’s threshold.

5.5 From template replication to evolution

Template replication requires three capabilities, each independently structural: (i) *Linear information-carrying polymers* — carbon’s sp^3 hybridization allows chain formation carrying N units of sequence information (structural). (ii) *Complementary pairing* — hydrogen bonding in $d = 3$ is directional, enabling selective donor-acceptor pairing between complementary sequences — the structural basis of Watson-Crick pairing (structural). (iii) *Catalytic copying* — template-directed catalysis is a subset of the statistically abundant catalysis established in §5.3, with selectivity provided by complementary pairing.

The RNA precedent. RNA is simultaneously a linear polymer, a substrate for complementary pairing, and a catalyst (ribozymes [18]). RNA-catalyzed RNA polymerization has been demonstrated [19]. This dual capability is the natural intersection of the three structural capabilities in a single molecular family.

Heritable variation is automatic. Template copying with finite fidelity produces errors. Errors in a self-replicating polymer are mutations. Faster replicators outcompete slower ones (selection), copying errors produce variants (variation), and successful variants propagate (heredity). This is Darwinian evolution — the inevitable dynamics of imperfect template replication.

The structural status of evolution.

Element	Status
Linear polymers	Structural (carbon sp^3 in $d = 3$)
Complementary pairing	Structural (H-bonding geometry in $d = 3$)
Catalytic activity	Statistically expected (§5.3)

Element	Status
Template replication	Intersection of the above three
Copying errors	Inevitable (finite-fidelity copying)
Selection	Inevitable (competition for finite resources)
Darwinian evolution	Consequence of replication + variation + selection

Each row follows from the preceding rows. No row requires contingent facts about our universe beyond the derived structure and the ~16% viable parameter fraction.

Caveats. The argument establishes that Darwinian evolution is *structurally expected* in any system with carbon chemistry, a thermal window, and a chiral aqueous environment. It does not determine: the specific molecular system that first replicates (RNA, PNA, or something else), the timescale for the transition (millions or billions of years), or the specific pathway from simple replicators to cellular organization. These depend on the specific , local geochemistry, and contingent history.

5.6 The origin of life: C1–C3 at the molecular scale

The OI framework’s characterization theorem is scale-independent: whenever a fast subsystem is coupled to a slow, high-capacity hidden sector (C1–C3), the fast subsystem’s dynamics is necessarily non-Markovian — history-dependent. At the cosmological scale, this produces quantum mechanics. At the molecular scale, it produces *heredity*: a chemical system whose future depends on its past through information stored in a persistent template. The transition from prebiotic chemistry to life is the first time a molecular system satisfies all three conditions.

Before life: Markovian chemistry. Ordinary chemical reactions are Markovian — the reaction rate depends on current concentrations alone, not on the history of the mixture. Products do not remember how they were made. Each reaction cycle is statistically independent of the previous one. No memory, no heredity, no evolution.

The transition: molecular C1–C3. Life begins when a molecular system first establishes: (C1) catalytic feedback — the template directs synthesis and synthesis maintains the template; (C2) persistence — the template outlives individual reaction cycles (RNA half-life of hours–days vs. reaction times of seconds–minutes); (C3) capacity — the template’s sequence space is large enough to encode catalytic function ($4^{20} \sim 10^{12}$ for a 20-mer RNA, vastly exceeding the \$10–100 reactions in a minimal metabolism).

After life: non-Markovian chemistry. The system’s future depends on its history, stored in the template. This is heredity — the defining property of life

— derived from the same C1–C3 conditions that produce quantum mechanics at the cosmological horizon.

RNA as the unique single-molecule C1–C3 system. RNA is the only natural polymer that satisfies all three conditions with a single molecular species: it is both template (C2: persistent; C3: 4-letter alphabet with arbitrary length) and catalyst (C1: ribozymes catalyze RNA polymerization [18, 19]). DNA stores information but cannot catalyze; proteins catalyze but cannot easily template. The RNA world hypothesis is not a historical accident — it is the structural prediction that the simplest C1–C3 system is an RNA-like molecule. DNA and proteins are later optimizations: DNA improves C2 (greater chemical stability), proteins improve C1 (more diverse catalysis). The RNA $\rightarrow$ DNA+protein transition refines C1–C3 without changing the qualitative architecture.

C1–C3 systems are attractors. A self-replicating system (C1–C3 satisfied) grows exponentially in a pool of feedstock molecules. A non-replicating system grows at most linearly (by accretion or random synthesis). Exponential growth dominates linear growth on any finite timescale. Any chemical system that *accidentally* satisfies C1–C3 — even transiently, even imperfectly — will dominate its environment. The transition is a one-way door: once C1–C3 is established, it is self-reinforcing. Combined with the autocatalytic argument (§5.3: $N \times p \sim 10^{56}$, guaranteeing C1), the automatic persistence of polymers (C2), and the enormous sequence space (C3), the establishment of molecular C1–C3 is not merely possible but *inevitable* in any aqueous organic chemistry operating in the thermal window.

The timescale. The earliest evidence for life on Earth is \$ \$3.8 Gya (carbon isotope signatures in Isua, Greenland [32]), on a planet that formed \$ \$4.5 Gya. The \$ \$700 Myr gap is consistent with the framework’s prediction: C1–C3 is inevitable but requires a period of prebiotic chemical exploration whose duration depends on solution-specific parameters (concentration, temperature fluctuation rates, mineral surface availability).

5.7 From evolution onward: companion to *Computation.md*

The structural chain established through §5.6 reaches biological evolution as a natural terminus. The next link — from evolved biological observers to information-processing systems, neural computation, semiconductor-based universal computation, and the embedded-computation theory that follows — is the subject of the companion paper *Embedded Computation in the Incompleteness of Observation Framework* [38]. There, the chain is continued through general intelligence, AI as embedded-observer construction, the Gödel-Turing-OI incompleteness family applied recursively, the BQP characterization theorem for embedded observers, and the framework’s position on the P vs NP question. The dividing line between this paper and the companion reflects the qualitative shift from *material-structural* consequence-tracing (chemistry-biology) to *informational-structural* consequence-tracing (computation-complexity).

6. The Existence Question

6.1 The last question

The structural chain derives everything from (S, φ) — physics, chemistry, life, intelligence, AI. But it does not derive (S, φ) itself. The reconstruction theorem [33, §3] establishes:

$$\text{Observed physics} \longleftrightarrow [(S, \varphi)] / \sim$$

This is conditional: *if* observation exists, *then* (S, φ) exists. The question is whether the framework can say anything unconditional about *why* (S, φ) exists at all.

6.2 Mathematical existence

(S, φ) is a finite set and a bijection. This is a mathematical object — a structure in the space of logical possibilities. It requires no physical substrate, no creator, and no cause. It exists in the same sense that the integers exist or that the permutation group S_n exists: as a logically consistent structure that cannot fail to be well-defined.

If mathematical structures are taken as necessarily existent — the Platonic position — then (S, φ) *necessarily* exists. Not “happens to exist” or “was created” but cannot fail to exist, because there is no logical contradiction in its definition. On this reading, “why is there something rather than nothing?” dissolves: the question assumes non-existence is the default and existence requires explanation. For mathematical structures, existence is the default. Non-existence would require a logical inconsistency, and (S, φ) has none.


6.3 Observable mathematics

The framework makes this sharper than generic Platonism. Tegmark’s Mathematical Universe Hypothesis [7] asserts that all consistent mathematical structures are equally real. The OI framework says something more specific: among all mathematical structures, only those containing embedded observers are *observable from within*. The characterization theorem [1, §3.4] identifies which structures contain embedded observers: exactly those of the form (S, φ) with conditions C1–C3. The reconstruction theorem [33, §3] identifies what those observers see: exactly our physics.

There is no ensemble of equally real structures from which observation selects. There is one equivalence class — the one compatible with embedded observation — and it is the one we see. The framework does not need a multiverse, an ensemble, or a selection mechanism. It needs only the observation that (S, φ) is a well-defined mathematical structure and that its observable content is unique.

6.4 The self-referential closure

The complete chain:



One step is contingent (general intelligence). Everything else is structural, statistical, or inevitable. The chain is closed: (S, φ) is a mathematical structure that contains observers who discover (S, φ) .

6.5 What cannot be answered

One question is outside the framework’s scope — not contingently but necessarily:

Why is there mathematical structure at all? The framework takes mathematical existence as given. It does not explain why logical consistency is a property that structures can have, or why “exists” and “is logically consistent” are related. This is not a gap in the framework — it is the boundary of all possible frameworks. Any explanation of mathematical existence would itself be a mathematical structure, requiring explanation in turn. The regress has no bottom.

The framework’s contribution is not to answer this question but to make it the *only* question. Everything else — the laws, the constants, the dark sector, the periodic table, life, intelligence, the fine-tuning problem, the anthropic principle — is derived, dissolved, or identified as generic. The mystery of existence is real. But it is the *only* mystery. And the framework proves it is the *same* mystery for every possible observer in every possible structure.

7. Quantum Biology: C1–C3 Inside Living Systems

7.1 The reframing

Quantum biology — the study of quantum coherence in biological systems — is an active and contested field. The central question is usually framed as: “How do biological systems maintain quantum coherence at warm, wet, noisy conditions?” The framework reframes this. QM is not a fragile state that biology must protect from decoherence. QM is the *necessary description* of any subsystem coupled to a slow, high-capacity hidden sector (the characterization theorem [1, §3.4]). If biological subsystems satisfy C1–C3 internally, their dynamics is necessarily P-indivisible — not because evolution engineered quantum coherence, but because the coupling architecture mandates it.

7.2 C1–C3 in proteins

Consider an enzyme’s active site as the visible sector and the protein scaffold’s conformational degrees of freedom as the hidden sector.

C1 (coupling). The active site is covalently bonded to the scaffold — coupling is continuous and bidirectional. Electronic transitions at the active site induce conformational strain in the scaffold; conformational changes in the scaffold modulate the active site’s electronic structure. This is the standard picture of allostery.

C2 (slow bath). Electronic transitions at the active site occur on femtosecond timescales ($\tau_S \sim 10^{-15}$ s). Conformational changes of the protein scaffold occur on microsecond to millisecond timescales ($\tau_B \sim 10^{-6}$ to 10^{-3} s). The ratio: $\tau_S/\tau_B \sim 10^{-12}$ to 10^{-9} . This is the *inverse* of the Markovian regime — exactly the slow-bath condition. The scaffold retains correlations for $\sim 10^9$ electronic transitions before relaxing.

C3 (capacity). A typical enzyme has $\sim 10^3$ to 10^4 atoms, each with multiple conformational degrees of freedom. The conformational state space is exponentially large: $\sim q^{N_{\text{conf}}}$ with $N_{\text{conf}} \sim 10^3$. This exceeds the active site’s electronic state space (~ 10 to 100 relevant states) by tens of orders of magnitude.

All three conditions are satisfied. The characterization theorem predicts: the active site’s reduced dynamics is P-indivisible.

7.3 The prediction

The protein scaffold is a slow bath. Correlations written into the scaffold during one catalytic event persist through subsequent events. The active site’s transition probabilities are history-dependent. This produces:

Information backflow. The scaffold returns correlations to the active site on conformational timescales (s to ms), modulating catalytic rates in a way that depends on the enzyme’s recent history. Standard rate theory predicts exponential kinetics; P-indivisibility predicts non-exponential kinetics with a specific

signature: the rate “constant” oscillates or shows memory effects on the conformational timescale.

Non-Markovian signatures. The mutual information $I(X_{<t}; X_{>t} | X_t)$ between past and future catalytic events, conditioned on the present state, is $\mathcal{O}(1)$ — not exponentially suppressed. This is the accessible-timescale backflow lemma [1, §2.3] applied to the protein system.

7.4 Existing evidence

Single-molecule enzyme kinetics. Individual enzyme molecules show fluctuating catalytic rates with correlations persisting for milliseconds to seconds (English et al. 2006 [21]; Lu et al. 1998 [22]) — precisely the signature of a slow bath modulating active site dynamics. The multi-timescale structure has been characterized at single-molecule resolution for T4 Lysozyme [34], where three distinct conformational states with ns-to-ms exchange times have been resolved, each playing a different role in the catalytic cycle.

Dispersed kinetics. Many enzymes show stretched-exponential waiting-time distributions rather than the single-exponential predicted by Markovian theory — generic in P-indivisible systems. The trace-out structure has been formalized within the single-molecule biophysics literature [35]: low-dimensional projections of high-dimensional molecular dynamics generically exhibit memory effects and produce non-Markov processes, independent of the specific dynamical model.

Photosynthetic energy transfer. Long-lived coherences in the FMO complex were first reported by Engel et al. [23]. Subsequent work (Duan et al. 2017 [36]; Cao-Plenio 2025 [37]) has refined the interpretation: the long-lived coherences are vibronic — superpositions of electronic exciton states with the protein’s vibrational environment — rather than purely electronic. Purely electronic coherences dephase within ~ 240 fs; the picosecond-scale coherences observed in 2D spectra arise from coupling between the electronic states and the slow vibrational protein scaffold. This is the C1–C3 structure operating: the protein scaffold (slow bath, C2) mixes with the electronic transitions at FMO chromophores (fast, visible) to produce coherences that persist for $\sim \tau_B$ rather than τ_S . The picosecond coherence timescales correspond to vibrational mode periods of the protein scaffold, consistent with the framework’s prediction that protein-scale C1–C3 systems exhibit dynamics on the slow-bath timescale.

Cryptochrome radical-pair magnetoreception. Avian magnetic compass orientation operates through cryptochrome radical pairs in the retinal photoreceptors. Radiofrequency fields at the electron Larmor frequency disrupt magnetic orientation in migratory birds (Wiltschko-Wiltschko 2019 [39]), diagnostically confirming radical pair mechanism. The crucial radical pairs have millisecond lifetimes [40]. Millisecond coherence in a protein-embedded quantum system at warm physiological conditions sits at the upper end of the protein conformational τ_B range identified in §7.2. The avian compass is therefore an

operating biological quantum sensor with coherence times the framework predicts as structurally required by C1–C3 inside protein scaffolds.

Conformationally gated tunneling. Hydrogen tunneling in enzymes shows temperature-independent kinetic isotope effects [24] — inconsistent with transition state theory but consistent with quantum tunneling modulated by slow conformational dynamics (C2).

7.5 Implications for drug design

Allosteric drugs. The framework predicts that allosteric effects carry temporal correlations through the scaffold — the drug’s binding event writes conformational information that is returned to the active site over the conformational timescale. Optimal design should account for these temporal correlations, not just equilibrium binding affinities.

Drug resistance. Resistance mutations often occur far from the active site. The framework predicts they cluster in regions that maximally alter the scaffold’s slowest conformational modes — altering the *memory structure* of the hidden sector. Testable via normal mode analysis.

Enzyme engineering. The framework suggests optimizing the *coupling architecture* (C1–C3 between active site and scaffold), not just the active site’s electronic structure. Better-tuned τ_S/τ_B ratios should yield both higher catalytic rates and more robust performance.

7.6 Epistemic status

The application of C1–C3 to protein systems is a *structural prediction* of the framework — it follows from the same characterization theorem that produces QM at the cosmological scale. The specific signatures (non-exponential kinetics, fluctuating rates, conformational gating of tunneling) are consistent with existing experimental data but have not been quantitatively tested against the framework’s specific predictions ($I(X_{<t}; X_{>t}|X_t) = \mathcal{O}(1)$, P-indivisibility of the catalytic cycle statistics). The medical implications (allosteric temporal correlations, resistance as memory-structure alteration) are consequences of the prediction but have not been independently tested.

The summary: the framework predicts that enzyme kinetics is non-Markovian for the same structural reason that cosmological observation is quantum-mechanical. The prediction is consistent with existing evidence and makes specific testable claims. Whether the non-Markovian corrections are large enough to matter for practical drug design is an empirical question — the corrections are $\mathcal{O}(\tau_S/\tau_B) \sim 10^{-9}$ to 10^{-12} for individual catalytic events, but may accumulate over many cycles or in systems with exceptionally slow conformational dynamics.

8. The GPS Analogy: Where OI Corrections Reach Engineering Relevance

8.1 The precedent

Nobody building a satellite navigation system in the 1950s would have predicted they would need general relativity. GPS requires relativistic corrections because satellite clocks experience different gravitational potentials — 45 s/day from GR, 7 s/day from SR. Without corrections, position drifts ~ 10 km/day. The practical application emerged when engineering precision reached the scale where the correction mattered.

The framework’s specific correction is non-Markovian dynamics: $\mathcal{O}(\tau_S/\tau_B)$ deviations from standard Markovian models. At the cosmological scale, $\tau_S/\tau_B \sim 10^{-32}$ — irrelevant for any technology. But τ_S/τ_B is scale-dependent. The question: where does engineering precision reach the scale where OI corrections matter?

8.2 Quantum computing

The Markovian assumption in error correction. Standard quantum error correction (surface codes, stabilizer codes) assumes Markovian noise — each error is statistically independent of previous errors. The framework predicts this is wrong whenever the qubit’s environment satisfies C2 (slow bath).

The physical system. Superconducting qubits (transmons) are coupled to two-level systems (TLS) in the substrate — structural defects in the amorphous oxide layer. TLS dynamics is slow ($\tau_B \sim \mu\text{s}$) relative to gate operations ($\tau_S \sim \text{ns}$). The ratio: $\tau_S/\tau_B \sim 10^{-3}$. This is 29 orders of magnitude larger than the cosmological case.

The prediction. The noise is P-indivisible: error probabilities at gate k depend on the error history at gates $k-1, k-2, \dots$ through correlations stored in the TLS bath. Standard Markovian error correction sees these correlated errors as an anomalously high uncorrectable error rate. Non-Markovian error correction protocols — designed for the actual P-indivisible noise structure — could exploit the correlations rather than treating them as random, improving logical error rates.

The scale. A 0.1% correction per gate. Small — but in a quantum algorithm with 10^4 gates, the cumulative effect is $\sim 10\%$. Current quantum computers are at the threshold where this matters. As gate counts increase toward practical quantum advantage (10^6+ gates), the correction becomes dominant.

8.3 Quantum sensing

Nitrogen-vacancy (NV) centers. NV centers in diamond detect magnetic fields with sensitivity approaching the standard quantum limit, bounded by

the coherence time T_2 . Standard models assume Markovian decoherence — coherence decays exponentially and is gone.

The OI prediction. The diamond lattice’s phonon bath has slow modes (paramagnetic impurities, strain fields with relaxation times $\sim\text{ms-s}$). C2 is satisfied: τ_S (electronic spin dynamics, $\sim\text{ns}$) $\ll \tau_B$ (bath relaxation, $\sim\text{ms}$). The decoherence is non-Markovian — coherence partially revives at times set by the bath’s correlation time. This is information backflow: correlations written into the lattice during one measurement are returned during a later measurement.

The application. Sensing protocols that exploit the revival — measuring at the backflow time rather than during the initial decay — could push sensitivity beyond the Markovian-assumed T_2 limit. The improvement factor is $\mathcal{O}(\tau_S/\tau_B) \sim 10^{-6}$ per measurement, but sensing protocols accumulate over $\sim 10^6$ repetitions, making the integrated improvement potentially significant.

8.4 Precision metrology and atomic clocks

Current situation. Optical lattice clocks achieve fractional frequency uncertainty $\sim 10^{-18}$ — sensitive enough that GR corrections from a 1 cm height difference are measurable. At this precision, the atoms’ interaction with the lattice light field may satisfy C1–C3: the lattice photons provide coupling (C1), the optical cavity has slow modes (C2), and the photon field has high capacity (C3).

The prediction. If the atom-cavity system satisfies C2 with sufficient separation, the atomic transition dynamics is non-Markovian. Clock instability would show correlations at the cavity correlation time — not the white frequency noise assumed by standard Allan variance analysis. Whether the correction is measurable at 10^{-18} fractional uncertainty depends on the specific cavity parameters, but the framework predicts its existence as a structural consequence.

8.5 The pattern

Domain	System / Bath	τ_S/τ_B	Current relevance
Cosmology	Observer / trans-horizon	10^{-32}	Framework’s home ground
Enzymes	Active site / scaffold	10^{-9} to 10^{-12}	§7: consistent with data
Quantum sensing	NV spin / lattice defects	10^{-6}	Approaching relevance
Quantum computing	Qubit / TLS	10^{-3}	At the threshold
Atomic clocks	Atom / cavity	$\sim 10^{-6}$	Potentially measurable

The corrections grow as τ_S/τ_B increases — equivalently, as the hidden sector’s timescale approaches the system’s timescale. The cosmological case is the most extreme separation (corrections negligible). Engineered quantum devices have the least separation (corrections largest). The framework predicts the same structural phenomenon — P-indivisible dynamics from a slow bath — at every scale. The question in each case is whether the correction has reached the precision frontier.

The GPS analogy is exact: GR corrections were irrelevant until clocks got precise enough. OI corrections are irrelevant until quantum devices get precise enough. The framework predicts that as quantum technologies continue improving, non-Markovian corrections will transition from negligible to measurable to design-relevant — following the same trajectory that relativistic corrections followed from Newtonian mechanics to GPS.

8.6 Engineered partitions: the bath as a resource

Standard quantum engineering treats the environment as the enemy — decoherence destroys quantum information, and the goal is to isolate qubits from their surroundings. The framework suggests a fundamentally different design philosophy: *engineer* the C1–C3 architecture to produce the quantum behavior you want.

The design principle. Instead of isolating a qubit from all environmental coupling, deliberately couple it to a *slow, high-capacity hidden sector* with controllable properties. The hidden sector stores correlations written during one gate operation and returns them during a later operation. The bath is not noise — it is programmable quantum memory.

Concrete platform: qubit + spin chain. A qubit coupled to a linear chain of L ancillary spins. The chain is the engineered hidden sector. Its correlation time scales as $\tau_B \sim L^z$ (where z is the dynamical exponent — for a Heisenberg chain, $z = 1$). Its capacity scales exponentially: $\sim 2^L$ states. By tuning L :

- Short chain ($L \lesssim 5$, $\tau_B \sim \tau_S$): Markovian regime. Standard decoherence. The bath equilibrates between gate operations.
- Long chain ($L \gtrsim 20$, $\tau_B \gg \tau_S$): P-indivisible regime. Information backflow. Correlations from gate k return at gate $k + \tau_B/\tau_S$, partially restoring coherence.
- The transition is sharp: the P-indivisibility theorem [1, §2.3] guarantees that once τ_B/τ_S exceeds the C2 threshold, the dynamics becomes qualitatively non-Markovian.

This is testable in existing platforms: superconducting qubits coupled to engineered spin chains [25], NV centers in diamond with controllable nuclear spin baths, or trapped ions with engineered phonon modes. The prediction: at the critical chain length, the qubit’s T_2 coherence time shows a qualitative change — from monotonic exponential decay to oscillatory behavior with partial revivals

at multiples of τ_B .

8.7 Non-Markovian quantum computation

The standard assumption. Quantum circuits are sequences of unitary gates, each assumed to act independently of past gates. Error correction adds redundancy to protect against independent errors. The entire architecture assumes Markovian noise.

The OI alternative. If gates share a slow bath (a common hidden sector with τ_B longer than the circuit depth $\times \tau_S$), the gates carry temporal correlations. Gate k 's effect depends on what gates $k-1, k-2, \dots$ did, through correlations stored in the bath.

Built-in error correction. P-indivisible dynamics can *increase* the distinguishability of quantum states over time — information backflow reverses some decoherence without requiring additional overhead. In Markovian evolution, distinguishability only decreases (data processing inequality). In P-indivisible evolution, the bath returns information that was temporarily inaccessible. This is a form of dynamical error correction built into the physics rather than added as a computational layer.

The open question. Whether P-indivisible gate sequences can perform computations that no Markovian circuit of equal depth can is a well-posed mathematical question. The framework provides the tools — P-indivisible stochastic processes, the accessible-timescale backflow lemma [1, §2.3] — but the computational complexity implications are unexplored. If the answer is yes, it would constitute a new model of quantum computation beyond the standard circuit model, with the bath playing the role of a shared temporal resource analogous to entanglement as a spatial resource.

8.8 Quantum materials by design

Heavy-fermion materials as natural P-indivisible systems. Heavy-fermion compounds (CeCoIn_5 , YbRh_2Si_2 , UPt_3) contain localized f-electrons forming a slow bath ($\tau_B \sim \text{ps to ns}$) coupled to itinerant conduction electrons ($\tau_S \sim \text{fs}$). The ratio: $\tau_S/\tau_B \sim 10^{-3}$ to 10^{-6} . These materials exhibit anomalous transport — non-Fermi liquid resistivity ($\rho \sim T^\alpha$ with non-integer α), logarithmic specific heat, and memory effects in magnetoresistance. The standard explanation invokes Kondo physics and quantum criticality. The OI framework provides a structural reframing: the conduction electrons satisfy C1–C3 with the f-electron bath, so their dynamics is P-indivisible. The anomalous transport is a *consequence* of non-Markovian electronic dynamics, not a deviation from normal behavior requiring a special mechanism.

The design principle. In natural materials, τ_S/τ_B is fixed by the chemistry. In engineered metamaterials, it can be tuned. Candidate platforms: (i) *Molecular junctions* — single molecules bridging electrodes, with floppy side chains

as slow bath ($\tau_S/\tau_B \sim 10^{-9}$); prediction: history-dependent conductance. (ii) *Cold-atom lattices* — two species with independently tunable tunneling rates; prediction: non-thermal momentum distributions at large τ_B/τ_S . (iii) *Photonic crystals with quantum dots* — engineered bandgap suppresses fast decay while maintaining coupling; prediction: non-exponential fluorescence with partial revivals.

8.9 Quantitative performance estimates

The engineering gains from P-indivisible design are not incremental corrections. They scale with τ_B/τ_S — the ratio the framework identifies as the fundamental design parameter.

Quantum error correction: 3–10 $\times$ overhead reduction. Standard surface codes assume independent (Markovian) errors. Correlated errors from a slow bath extend over $\sim \tau_B/\tau_S$ gates. A Markovian decoder sees these correlations as a higher effective error rate — it *underperforms*. A correlation-aware decoder exploits the predictability: correlated errors are easier to correct than random errors. For superconducting qubits with TLS baths ($\tau_S/\tau_B \sim 10^{-3}$, correlations over $\sim 10^3$ gates), a correlation-aware decoder could reduce the effective error rate by a factor of 2–10 $\times$ at the same physical error rate. At current physical error rates ($\sim 10^{-3}$), this reduces the overhead from $\sim 3,000$ physical qubits per logical qubit to ~ 300 –1,000 — a factor of 3–10 $\times$ reduction in the hardware cost of fault-tolerant quantum computing.

Coherence extension: up to 10 $\times$ via engineered bath. Standard coherence decays as e^{-t/T_2} . With information backflow from an engineered bath, coherence partially revives at multiples of τ_B . The revival amplitude per cycle is $\mathcal{O}(\tau_S/\tau_B)$. For natural systems ($\tau_S/\tau_B \sim 10^{-3}$), revival is $\sim 0.1\%$ — negligible. For an engineered system ($\tau_S/\tau_B \sim 10^{-1}$, achievable with a qubit coupled to a short spin chain), revival is $\sim 10\%$ per cycle. Over multiple revival cycles before the bath relaxes, the effective coherence time extends by a factor of $\sim \tau_B/\tau_S$. Current superconducting qubit $T_2 \sim 100$ s could be pushed to ~ 1 ms without improving the qubit itself — the gain comes from the bath.

Quantum sensing: up to 10 $^3\times$ sensitivity improvement. Standard NV magnetometry sensitivity: $\delta B \sim 1/(\gamma\sqrt{T_2 \cdot T_{\text{meas}}})$. Non-Markovian backflow at τ_B means each measurement recovers information that Markovian protocols assume is lost. For $\tau_S/\tau_B \sim 10^{-6}$ (NV center with paramagnetic bath), each measurement carries $\sim 10^{-6}$ extra information. Over 10^6 repetitions, this integrates to $\mathcal{O}(1)$ — a factor of $\sqrt{\tau_B/\tau_S} \sim 10^3$ improvement in signal-to-noise if the protocol is optimized for the backflow timing. This would push NV magnetometry from nT/ $\sqrt{\text{Hz}}$ to pT/ $\sqrt{\text{Hz}}$ — competitive with SQUIDs but at room temperature.

Quantum materials: qualitative regime change. For engineered materials with $\tau_S/\tau_B \sim 10^{-1}$, non-Markovian corrections are not perturbative — they dominate. Standard band theory gives qualitatively wrong predictions. The

material’s electronic properties are determined by P-indivisible dynamics, not by Markovian approximations. This is not a percentage improvement — it is a new regime of matter. The heavy-fermion precedent (non-Fermi liquid transport, unconventional superconductivity, hidden order) suggests that qualitatively new phenomena emerge when the P-indivisible corrections become $\mathcal{O}(1)$.

Domain	τ_S/τ_B	Estimated gain	Status
Error correction	10^{-3}	3–10× overhead reduction	Testable now with correlation-aware decoders
Coherence extension	10^{-1} (engineered)	Up to $10\times T_2$	Requires engineered spin chain coupling
Quantum sensing	10^{-6}	Up to $10^3\times$ sensitivity	Requires backflow-optimized protocols
Quantum materials	10^{-1} (engineered)	Qualitatively new regime	Requires metamaterial fabrication

8.10 Existing experimental evidence

The framework’s predictions about non-Markovian effects are not speculative — they are corroborated by existing experimental data across multiple domains. The literature has been documenting these effects for over two decades without a unifying structural explanation. The OI framework provides one: wherever C1–C3 are satisfied, P-indivisibility is mandatory.

Superconducting qubits. Agarwal et al. [26] found that purely Markovian noise models cannot reproduce experimental data from driven superconducting qubits. The non-Markovian dynamics arises from two-level system (TLS) interactions in the substrate — precisely the slow bath (C2) the framework identifies. White et al. [27] performed the first full multi-time quantum process tomography on superconducting processors and found non-Markovian noise present in all cases measured, with a significant fraction originating from genuine quantum correlations across time. Notably, Burkard and collaborators [28] found that QAOA algorithm performance *improves* as the noise correlation time increases at fixed local error probability — direct evidence that correlated noise is a resource, not merely an obstacle, exactly as §8.7 predicts. More recent work has begun to operationalize this resource: Puviani et al. (Phys. Rev. Lett. 2025)

[41] used reinforcement learning to train recurrent neural networks providing quantum error correction schemes based on memory — responding to the full history of measurement outcomes rather than only the most recent one — and demonstrated that these non-Markovian feedback schemes significantly outperform standard Markovian strategies. Mandayam and collaborators [42] developed Petz-recovery-map QEC schemes adapted to non-Markovian noise, with the non-Markovian-adapted Petz map outperforming standard stabilizer codes. The transition from documenting non-Markovian noise as a problem (Agarwal, White) to engineering it as a resource (Burkard, Puviani, Mandayam) is the empirical realization of §8.7’s structural prediction.

Single-molecule enzymology. Edman and Rigler [29] directly measured “memory landscapes” of single horseradish peroxidase molecules, extracting non-Markovian behavior from the catalytic cycle. The enzyme’s activity fluctuates over timescales from milliseconds to seconds — the signature of a slow conformational bath (C2) modulating the active site’s electronic dynamics. Kou and Xie [30] showed that slow conformational interconversion produces memory effects in successive enzymatic turnover times: the waiting time for turnover $k + 1$ is correlated with the waiting time for turnover k , with correlation strength decaying on the conformational timescale. This is precisely the information backflow predicted by the accessible-timescale lemma [1, §2.3] applied to the protein system.

The unifying explanation. These experimental results — from quantum computing hardware and from single-enzyme biophysics — are conventionally treated as unrelated phenomena requiring separate theoretical frameworks. The OI framework unifies them: both are instances of P-indivisible dynamics arising from a slow, high-capacity hidden sector coupled to a fast observable subsystem. The TLS bath in a superconducting chip and the conformational bath in a protein scaffold play the same structural role — they satisfy C2 (slow) and C3 (high capacity), producing the same qualitative phenomenon (information backflow, memory effects, non-exponential dynamics) at different scales.

8.11 Cross-domain observation: universal memory strength in single-entity systems

The non-Markovian dynamics documented in §8.10 is conventionally quantified by domain-specific measures. These can be converted to the Hurst exponent H ($H = 0.5$: Markovian; $H > 0.5$: persistent memory) via $H \approx 1 - \beta/2$ for stretched exponential processes and $H = (1 + \alpha)/2$ for $1/f^\alpha$ noise.

A known objection to universality claims is that any ensemble of many independent exponentially relaxing modes with a broad rate distribution generically produces $H \approx 0.7\text{--}0.8$ [31]. This superposition argument applies to bulk/aggregate measurements (glasses, financial markets, river flows, EEG) where many components are averaged. It does *not* apply to single-molecule and single-system measurements, where there is no ensemble to superpose. Non-exponential ki-

netics in a single molecule must arise from the molecule’s internal dynamics — specifically, from coupling between the observed process and slow internal degrees of freedom (C1–C3).

Restricting to single-entity measurements where superposition is excluded by construction:

System	Measurement	Memory depth	H	Source
β -galactosidase	Single-molecule turnover	$\sim 10^{2.6}$	0.80	English et al. 2006 ($\beta = 0.4$)
Cholesterol oxidase	Single-molecule turnover	$\sim 10^{1.6}$	0.75	Lu et al. 1998 ($\beta \approx 0.5$)
Horseradish peroxidase	Single-molecule turnover	$\sim 10^3$	0.83	Edman/Rigler 2000 ($\beta \approx 0.35$)
Lipase B	Single-molecule turnover	$\sim 10^2$	0.80	Flomenbom et al. 2005 ($\beta \approx 0.4$)
Myoglobin CO rebinding	Single-molecule kinetics	$\sim 10^6$	0.95	Austin et al. 1975 ($\beta \approx 0.1$)
SC qubit (Lorentzian TLS)	Single-qubit trajectory	$\sim 10^2$	0.80	Noise spectroscopy ($\alpha \approx 0.6$)
Ion channel (membrane-coupled)	Single-channel recording	$\sim 10^3$	0.80	Dwell-time autocorrelation

For the four enzymes and the qubit — all with “typical” C3 capacity (conformational state spaces of $\sim 10^2$ – 10^4 states) — $H = 0.79 \pm 0.03$. The superposition objection is excluded: these are individual molecules or devices, not ensembles. The memory is intrinsic, arising from coupling to a slow internal hidden sector (protein scaffold conformational dynamics for enzymes, TLS bath for the qubit).

The myoglobin outlier ($H = 0.95$) has an anomalously deep conformational hierarchy — Frauenfelder’s “protein energy landscape” involves $\sim 10^6$ – 10^9 conformational substates organized in a fractal-like tier structure, far exceeding typical enzyme C3 capacity. This suggests H increases with C3: typical capacity gives $H \approx 0.8$; deep capacity gives $H \rightarrow 1$. The predicted relationship $H(C3)$ — derivable in principle from the characterization theorem’s dependence on hidden-sector dimension — is a quantitative test.

Acknowledgements

During the preparation of this work, the author used Claude Opus 4.6 (Anthropic) to assist in drafting, refining argumentation, and surveying the relevant literature in nuclear physics, prebiotic chemistry, origin-of-life research, quantum biology, and quantum computing. The author reviewed and edited all content and takes full responsibility for the publication.

References

- [1] A. Maybaum, “The Incompleteness of Observation,” (2026).
- [2] A. Maybaum, “The Standard Model from a Cubic Lattice,” (2026).
- [3] P. Ehrenfest, “In what way does it become manifest in the fundamental laws of physics that space has three dimensions?” *Proc. Amsterdam Acad.* **20**, 200 (1917).
- [4] F. R. Tangherlini, “Schwarzschild field in n dimensions and the dimensionality of space problem,” *Nuovo Cim.* **27**, 636 (1963).
- [5] A. D. Sakharov, “Violation of CP invariance, C asymmetry, and baryon asymmetry of the universe,” *JETP Lett.* **5**, 24 (1967).
- [6] C. J. Hogan, “Why the Universe is just so,” *Rev. Mod. Phys.* **72**, 1149 (2000).
- [7] M. Tegmark, “Is ‘the theory of everything’ merely the ultimate ensemble theory?” *Ann. Phys.* **270**, 1 (1998).
- [8] F. C. Adams, “The degree of fine-tuning in our universe — and others,” *Phys. Rep.* **807**, 1 (2019).
- [9] A. Jaffe, R. L. Jaffe, and F. Wilczek, “Quark masses: an environmental impact statement,” *Phys. Rev. D* **79**, 065014 (2009).
- [10] E. Epelbaum, H. Krebs, T. A. Lähde, D. Lee, and U.-G. Meißner, “Viability of carbon-based life as a function of the light quark mass,” *Phys. Rev. Lett.* **110**, 112502 (2013).
- [11] M. Quack, “How important is parity violation for molecular and biomolecular chirality?” *Angew. Chem. Int. Ed.* **41**, 4618 (2002).
- [12] P. Schwerdtfeger, “The search for parity violation in chiral molecules,” in *Relativistic Methods for Chemists* (Springer, 2010).
- [13] K. Soai, T. Shibata, H. Morioka, and K. Choji, “Asymmetric autocatalysis and amplification of enantiomeric excess of a chiral molecule,” *Nature* **378**, 767 (1995).

- [14] C. Viedma, “Chiral symmetry breaking during crystallization: complete chiral purity induced by nonlinear autocatalysis and recycling,” *Phys. Rev. Lett.* **94**, 065504 (2005).
- [15] G. F. Joyce, G. M. Visser, C. A. A. van Boeckel, J. H. van Boom, L. E. Orgel, and J. van Westrenen, “Chiral selection in poly(C)-directed synthesis of oligo(G),” *Nature* **310**, 602 (1984).
- [16] S. A. Kauffman, *The Origins of Order: Self-Organization and Selection in Evolution* (Oxford University Press, 1993).
- [17] M. R. Ghadiri, J. R. Granja, R. A. Milligan, D. E. McRee, and N. Khazanovich, “Self-assembling organic nanotubes based on a cyclic peptide architecture,” *Nature* **366**, 324 (1993).
- [18] T. R. Cech, “Self-splicing of group I introns,” *Annu. Rev. Biochem.* **59**, 543 (1990).
- [19] W. K. Johnston, P. J. Unrau, M. S. Lawrence, M. E. Glasner, and D. P. Bartel, “RNA-catalyzed RNA polymerization: accurate and general RNA-templated primer extension,” *Science* **292**, 1319 (2001).
- [20] J. von Neumann, *Theory of Self-Reproducing Automata*, ed. A. W. Burks (University of Illinois Press, 1966).
- [21] B. P. English, W. Min, A. M. van Oijen, K. T. Lee, G. Luo, H. Sun, B. J. Cherayil, S. C. Kou, and X. S. Xie, “Ever-fluctuating single enzyme molecules: Michaelis-Menten equation revisited,” *Nature Chemical Biology* **2**, 87 (2006).
- [22] H. P. Lu, L. Xun, and X. S. Xie, “Single-molecule enzymatic dynamics,” *Science* **282**, 1877 (1998).
- [23] G. S. Engel, T. R. Calhoun, E. L. Read, T.-K. Ahn, T. Mančal, Y.-C. Cheng, R. E. Blankenship, and G. R. Fleming, “Evidence for wavelike energy transfer through quantum coherence in photosynthetic systems,” *Nature* **446**, 782 (2007).
- [24] J. P. Klinman, “An integrated model for enzyme catalysis,” *FEBS Letters* **589**, 467 (2015).
- [25] P. Roushan et al., “Spectroscopic signatures of localization with interacting photons in superconducting qubits,” *Science* **358**, 1175 (2017).
- [26] A. Agarwal, L. P. Lindoy, D. Rocchetto, M. Sherbert, and I. Rungger, “Modelling non-Markovian noise in driven superconducting qubits,” arXiv:2306.13021 (2023).
- [27] G. A. L. White, C. D. Hill, F. A. Pollock, L. C. L. Hollenberg, and K. Modi, “Multi-time quantum process tomography of a superconducting qubit,” arXiv:2308.00750 (2023).

- [28] J. Beinert, F. Burkard, J. Olle, D. S. Wang, and F. K. Wilhelm, “Ability of error correlations to improve the performance of variational quantum algorithms,” *Phys. Rev. A* **107**, 042426 (2023).
- [29] L. Edman and R. Rigler, “Memory landscapes of single-enzyme molecules,” *Proc. Natl. Acad. Sci. USA* **97**, 8266 (2000).
- [30] S. C. Kou and X. S. Xie, “Generalized Langevin equation with fractional Gaussian noise: subdiffusion within a single protein molecule,” *Phys. Rev. Lett.* **93**, 180603 (2004).
- [31] D. C. Johnston, “Stretched exponential relaxation arising from a continuous sum of exponential decays,” *Phys. Rev. B* **74**, 184430 (2006).
- [32] M. Mojzsis, G. Arrhenius, K. D. McKeegan, T. M. Harrison, A. P. Nutman, and C. R. L. Friend, “Evidence for life on Earth before 3,800 million years ago,” *Nature* **384**, 55 (1996). [33] A. Maybaum, “The Substratum Construction: Reconstruction, the Substratum Gauge Group, and the Synthesis of Quantum Mechanics with General Relativity,” (2026).
- [34] H. Sanabria et al., “Resolving dynamics and function of transient states in single enzyme molecules,” *Nature Communications* **11**, 1231 (2020).
- [35] A. M. Berezhkovskii and D. E. Makarov, “Memory effects and static disorder reduce information in single-molecule signals,” *bioRxiv* 2022.01.13.476256 (2022).
- [36] H.-G. Duan et al., “Nature does not rely on long-lived electronic quantum coherence for photosynthetic energy transfer,” *Proc. Natl. Acad. Sci. USA* **114**, 8493 (2017).
- [37] E. Cao and M. B. Plenio, “Full microscopic simulations uncover persistent quantum effects in primary photosynthesis,” *Science Advances* (2025). doi:10.1126/sciadv.ady6751.
- [38] A. Maybaum, “Embedded Computation in the Incompleteness of Observation Framework: From Evolution to BQP and the P vs NP Question,” (2026).
- [39] R. Wiltschko and W. Wiltschko, “Magnetoreception in birds,” *Journal of the Royal Society Interface* **16**, 20190295 (2019).
- [40] M. Liedvogel et al., “Chemical magnetoreception: Bird cryptochrome 1a is excited by blue light and forms long-lived radical-pairs,” *PLoS ONE* **2**, e1106 (2007).
- [41] M. Puviani, S. Borah, R. Zen, J. Olle, and F. Marquardt, “Non-Markovian feedback for optimized quantum error correction,” *Phys. Rev. Lett.* (2025); arXiv:2312.07391.
- [42] D. Biswas, S. Utagi, and P. Mandayam, “Noise-adapted quantum error correction for non-Markovian noise,” arXiv:2411.09637 (2024).