

Non-Markovian Dynamics in Biology and Medicine

Molecular Memory as a Unifying Framework for Disease Mechanisms and Therapeutic Design

Author: Alex Maybaum

Date: April 2026

Status: DRAFT PRE-PRINT

Classification: Theoretical Biology / Pharmacology

Abstract

The Observational Incompleteness (OI) framework [1] proves that any fast subsystem coupled to a hidden sector realizing conditions C1–C4 — non-trivial coupling, record persistence (slowness one realization), sufficient capacity, and history readback (C1–C4) — exhibits, in the memory-bearing sector those conditions characterize, non-Markovian dynamics (P-indivisibility the stronger temporal-indivisibility diagnostic, established at the recurrence scale by C1 and finiteness alone) — history-dependent transition probabilities arising from information stored in the hidden sector and returned on subsequent interactions. Originally developed for fundamental physics, the theorem’s architectural conditions (C1–C4) are scale-independent; although biological subsystems are open and dissipative rather than the closed, reversible substratum the full theorem assumes, the model-independent consequence — that tracing out a persistent, high-capacity sector whose coupling reads back its own record yields history-dependent (non-Markovian) dynamics — applies to them directly. We identify biological systems — from single enzymes through signaling cascades to the epigenome — as natural instantiations of this architecture. The fast catalytic process (enzyme active site, ion channel gating, kinase phosphorylation) is coupled (C1) to slow conformational or post-translational modification dynamics (C2) with exponentially large state spaces (C3), producing history-dependent behavior that standard Markovian models cannot capture.

We develop this structural observation into a unified framework for seven medical domains: cancer pharmacology (checkpoint kinase memory and schedule-dependent sensitization), neurodegeneration (Alzheimer’s, Parkinson’s, and PTSD as disorders of molecular memory timescale, with reconsolidation paradigms and low-intensity focused ultrasound as the first directly τ_B -targeted therapeutic modality), antibiotic resistance (persister cells as SOS memory accumulation), immunotherapy (T cell exhaustion as accumulated TCR signaling memory), cardiac pharmacology (use-dependent ion channel block as gating memory), and autoimmune disease (disproportionate efficacy of partial JAK inhibition as memory disruption), and the treatment management

of genetic disorders (replacement therapy scheduling, inhibitor prevention, gene therapy durability). In each case, the framework identifies a specific therapeutic axis — **memory asymmetry** between disease and normal tissue — that is pharmacologically distinct from standard catalytic inhibition and predicts wider therapeutic windows. We extend the analysis to epigenetic regulation, identifying the chromatin state as the biological hidden sector with a hierarchical memory architecture spanning minutes (histone acetylation) to generations (DNA methylation), and further upward through synaptic, circuit, systems, and cortical layers in the nervous system. The framework’s central derivations — the necessarily reconstructive nature of memory access, the structural requirement of finite τ_B at every layer for functional dynamics, and the layer-specific reconsolidation window — provide unified accounts of phenomena previously treated as contingent features of biological memory. Twenty-nine main testable predictions (with three additional layer-specific PTSD/LIFU sub-predictions in neurodegeneration) are presented, each distinguishing the non-Markovian framework from standard Markovian pharmacology. Several predictions are already supported by existing data; the remainder are experimentally accessible with current techniques.

Status — conjectural extension. This paper applies the core framework to medicine, and is offered as a *conjecture*, not an established result: its retrodictions — notably the Rett-syndrome reversibility (Guy et al. 2007), which predates the framework — are consistency checks rather than confirmations, while its forward content (wider therapeutic windows across as-yet-untested reader-writer disorders, and specific schedule-dependence predictions) is what future clinical data would confirm or break. It carries no evidential weight for the core framework — which stands or falls independently on its physics — and should be read as a structural possibility to be tested, not as support for the framework.

1. Introduction

Standard pharmacological models treat enzyme catalysis, receptor signaling, and ion channel gating as Markovian processes — each event is independent of prior history, and the system’s future depends only on its current state. Michaelis-Menten kinetics, Hodgkin-Huxley channel models, and Hill-equation dose-response curves all embed this assumption. When history-dependent behavior is observed (use-dependent drug block, adaptive radiation responses, schedule-dependent chemotherapy efficacy), it is accommodated by adding internal states to the Markov model — often many states, without a unifying principle for when they are needed or how many to include.

The Observational Incompleteness (OI) framework [1] provides such a principle. Originally developed for fundamental physics — where it proves that the descrip-

tion of any embedded observer with partial access to a deterministic system is necessarily non-Markovian — quantum for processes in the correspondence’s class, under the translation hypothesis (T) — the framework’s core theorem is abstract and scale-independent. It identifies four conditions under which any fast subsystem necessarily exhibits non-Markovian dynamics:

C1 (Coupling). The fast subsystem (visible sector) is dynamically coupled to a slow subsystem (hidden sector) through bidirectional interactions.

C2 (Memory persistence). The written record persists to readback — realized here by a slow hidden sector, its relaxation timescale τ_B greatly exceeding the fast subsystem’s event timescale τ_S ($\tau_S/\tau_B \ll 1$), one sufficient route. The hidden sector retains correlations from past interactions across many fast events.

C3 (Capacity). The hidden sector has many more accessible states than the fast subsystem, providing sufficient room to store the full interaction history without saturation.

When C1–C4 are satisfied, the characterization theorem [1, §3.4] proves that the fast subsystem’s dynamics is memory-bearing — non-Markovian, with P-indivisibility the stronger diagnostic — so that transition probabilities at time t depend on the system’s history, stored in the hidden sector’s state and returned through the coupling. The strength of the history-dependence is controlled by τ_S/τ_B — the ratio of fast to slow timescales.

The biological relevance is immediate. Enzymes, kinases, ion channels, and receptors are composed of a fast catalytic domain coupled to slow regulatory domains, post-translational modification (PTM) sites, and conformational degrees of freedom. The catalytic cycle operates on nanosecond-to-microsecond timescales; the regulatory domain’s conformational changes persist for microseconds to milliseconds; PTM patterns persist for minutes to hours; chromatin modifications persist for days to generations. At every scale, C1–C4 are satisfied, and the theorem predicts non-Markovian dynamics.

This paper develops the medical implications of this observation across seven domains, identifies a unifying therapeutic principle (memory asymmetry), and presents twenty-nine testable predictions that distinguish the non-Markovian framework from standard Markovian pharmacology.

2. The Enzyme as a Read-Write System

2.1 Molecular memory mechanisms

The OI prediction requires that an enzyme’s activity history is physically encoded in its structure and persists across multiple catalytic cycles. This is well-established through multiple mechanisms.

Multisite post-translational modification. PTMs — phosphorylation,

acetylation, ubiquitination, methylation — covalently alter specific residues, changing the protein’s conformational landscape and activity. A protein with N modifiable sites, each with k possible states, has k^N distinct modification patterns. For Chk1 (with ~ 10 regulatory phosphorylation sites): $2^{10} \approx 1,000$ distinct states — a physical memory register encoding the enzyme’s recent history. Gunawardena (2012) explicitly describes this as “history-based encoding” — the same cellular condition can produce different modification patterns depending on the prior history.

Sequential phosphorylation. For Chk1 specifically, phosphorylation at S317 must precede phosphorylation at S345 — the second event is conditional on the first (Wilsker et al. 2008). This is a direct non-Markovian signature: the probability of the second modification depends on history, not just the current state.

Conformational hysteresis. Proteins occupy multiple distinct conformational states, with transitions depending on the protein’s history. For kinases, the autoinhibited vs. active conformation persists on timescales much longer than individual phosphorylation events. The Chk1-S splice variant acts as an endogenous repressor whose binding/unbinding is the slow process (C2) storing the history.

Chromatin as long-term memory. At the network level, the histone code provides the most dramatic example. DNA damage induces histone modifications (γ H2AX, H4K20me) that persist for hours to days — far longer than the kinase cascade that wrote them.

2.2 The key parameter

The strength of the non-Markovian correction is set by τ_S/τ_B . When this ratio is very small (e.g., 10^{-12} for enzyme active site electronics vs. scaffold motions), the correction per event is tiny but accumulates. When the ratio approaches 1 (e.g., for allosteric sites where regulatory dynamics is comparable to the catalytic cycle), non-Markovian effects dominate and Markovian models give qualitatively wrong predictions.

The memory hierarchy spans many decades of timescale. At the molecular level:

Memory mechanism	Write operation	Storage medium	τ_B	C1–C4 role
Conformational hysteresis	Ligand binding	Oligomeric/regulatory state	seconds	C2
Sequential phosphorylation	Ordered modification	Conformational accessibility	Minutes	C2 + C3

Memory mechanism	Write operation	Storage medium	τ_B	C1–C4 role
Multisite PTM	Phosphorylation	Modification pattern (2^N states)	Min–hrs	C3
Chromatin marks	γ H2AX, methylation	Histone modification state	Hrs–days	C2 + C3
DNA methylation	Methyltransferase activity	CpG state	Months–generations	C2 + C3

In multicellular tissues — particularly the nervous system — the same C1–C4 architecture extends upward through additional layers, each treating the layer below as its hidden sector:

Memory layer	Substrate	τ_B
Synaptic Circuit	LTP/LTD weight changes (AMPA trafficking)	hr–days
Systems	Engram-cell ensembles, recurrent dynamics	days–months
Cortical	Hippocampal–cortical consolidation dialogue	months–years
	Distributed semantic representation	years–decades

The hierarchical structure means that a perturbation at one layer (a drug, a disease process) can have effects propagating upward through layers with progressively longer τ_B . The clinical timecourse of memory-related diseases — and the schedule-dependence of memory-targeted drugs — reflects this layered architecture.

2.3 The core therapeutic principle

In disease contexts, the framework identifies a specific therapeutic axis: **memory asymmetry**. When a disease process depends on non-Markovian signaling dynamics that the corresponding normal tissue does not depend on (or depends on differently), therapies can target the *memory structure* rather than the *catalytic function*. This is pharmacologically distinct from standard inhibition and predicts wider therapeutic windows because the target (memory dependence) is more disease-specific than the target (catalytic activity).

2.4 Reconstructive consequences of memory access

A direct consequence of the framework’s core mechanism: *every act of accessing a memory necessarily alters the substrate that stores it*. Observation of a coupled hidden sector cannot be passive — the visible-sector readout requires interaction with the hidden sector, and this interaction back-acts on the hidden state. The

re-stored trace is therefore a function of both the original hidden state and the present visible-sector context at the moment of access.

This is the framework’s derivation of the *reconstructive nature of recall* — a phenomenon long established in cognitive science but typically treated as a contingent feature of biological memory. The framework makes it structural: any C1–C4 memory system must exhibit constructive re-storage on access. Pure address-store memory (where access leaves the substrate unchanged) is incompatible with C1–C4.

Clinical consequence: the reconsolidation paradigm. When a memory is recalled, the recalled trace becomes labile and must be re-stored. The window during which it is labile provides a therapeutic opportunity. A drug given during the recall window writes into the re-storage process, modifying the future content of the memory. Propranolol given during recall of a traumatic memory selectively blocks the noradrenergic component of re-encoding, reducing the emotional valence of subsequent retrievals — a result well-established in PTSD trials (Brunet et al., reviewed extensively).

The framework predicts that this approach generalizes: any memory layer with a measurable τ_B should have a corresponding “reconsolidation window” during which targeted intervention can rewrite the memory at that layer without affecting memories already consolidated to slower layers. The selectivity is not coincidental — it follows from the layered architecture of §2.2.

2.5 Forgetting as functional necessity

Standard accounts treat forgetting as a failure of memory. The framework reframes it: *finite τ_B at every layer is structurally required*. A system with infinite τ_B at every layer would be unable to discriminate present from past — it would have no temporal structure, no “now,” no capacity to register change. Such a system is not a memory device; it is a frozen state.

The brain operates in the parameter regime where each layer’s τ_B is matched to that layer’s functional role. Pattern separation requires fast forgetting at the molecular layer; generalization requires consolidation transfer from fast to slow layers; fear extinction and mood regulation require active erasure at intermediate layers.

Clinical consequence: memory disorders are τ_B disorders. From the framework’s perspective, the relevant clinical question for any memory-related condition is not “is the memory intact or damaged?” but “which layer’s τ_B has shifted, and in which direction?” This reframes:

- **Excessive retention** (PTSD, OCD intrusive thoughts, depressive rumination) $\rightarrow \tau_B$ too long at the affected layer
- **Failure of new encoding** (anterograde amnesia, attention deficits) $\rightarrow \tau_B$ too short, or write-rate too low

- **Selective loss of old content** (retrograde amnesia, semantic dementia) → loss of substrate at affected layer
- **Failure of cross-layer transfer** (Korsakoff's syndrome, sleep-deprivation-induced consolidation failure) → disruption of layer-coupling mechanisms
- **Mixed pattern** (normal aging, Alzheimer's) → τ_B shifts at multiple layers in opposite directions

The therapeutic axis follows directly: τ_B -normalizing interventions targeting the specific affected layer, rather than broad-spectrum receptor pharmacology.

3. Cancer Pharmacology

3.1 Checkpoint kinase inhibitors and selective tumor sensitization

Chk1 inhibitors combined with gemcitabine and/or radiation selectively sensitize tumor cells — particularly pancreatic cancer — while sparing normal cells. Chk1 is a serine/threonine kinase whose catalytic domain (fast subsystem) is regulated by its SQ/TQ domain and C-terminal regulatory region (slow hidden sector). C1–C4 are satisfied: catalytic and regulatory domains are allosterically coupled (C1); regulatory conformational changes (μs – ms) $\gg$ phosphorylation events (ns – μs), with PTM persistence extending to minutes–hours (C2); the regulatory domain has $\sim 2^{10}$ modification states plus continuous conformational degrees of freedom (C3).

The prediction. Chk1's checkpoint signaling is non-Markovian: its response to a DNA damage signal depends on its recent activation history — specifically, on the PTM pattern and conformational state written by previous damage events. A Chk1 inhibitor that binds the regulatory domain disrupts the *memory structure* of the kinase, altering the history-dependence of future checkpoint responses.

Why selectivity emerges. In tumor cells (defective G1, relying on Chk1), the non-Markovian memory is the primary mechanism maintaining genomic stability through repeated replication cycles. Disrupting this memory is catastrophic. In normal cells (intact G1), the Chk1 memory is redundant — the G1 checkpoint provides an independent, Markovian damage response. The selectivity is not just about checkpoint redundancy — it is about the differential role of non-Markovian dynamics.

Why schedule matters. The finding that the *order and timing* of gemcitabine + Chk1 inhibitor administration is critical for efficacy is a direct signature of non-Markovian dynamics. In a Markovian system, only concentrations matter. In a non-Markovian system, the first drug writes information into the enzyme's memory, and the second drug's effect depends on what was written.

3.2 Radiation adaptive response

A small priming dose of radiation (~ 0.01 – 0.1 Gy) makes cells more resistant to a subsequent larger dose (~ 2 Gy), with the effect lasting hours to days. The OI framework identifies this as information backflow from the chromatin hidden sector: dose 1 writes modification marks into histones, which persist and are read by the DDR when dose 2 arrives. The “adaptation” is the DDR network reading the memory of the previous damage event.

3.3 Therapeutic strategies from memory asymmetry

Memory-selective scheduling. The optimal dose interval depends on τ_B of the tumor’s checkpoint kinases. The second dose should arrive when the memory written by the first dose is maximally sensitizing — a timing that differs between tumor cells (deregulated Chk1, altered τ_B) and normal cells (normal τ_B). Current schedules are empirically optimized from a few discrete intervals. The framework says the optimal interval is *calculable* from the measured τ_B and is tumor-specific.

Low-dose memory priming. Instead of a single high dose, give repeated low “priming” doses that write sensitizing memory into the tumor’s checkpoint system, followed by a moderate treatment dose. The predicted protocol: Days 1, 3, 5 — gemcitabine at $\sim 20\%$ MTD (priming phase); Day 7 — gemcitabine at $\sim 50\%$ MTD + Chk1 inhibitor (kill phase). Total drug exposure: $\sim 70\%$ of standard protocol. Predicted toxicity: substantially lower, because normal cells’ Markovian G1 backup is unaffected by priming.

Memory-targeted drugs. Drugs that specifically disrupt the *memory structure* without blocking catalysis: accelerating the regulatory domain’s slowest conformational mode (decreasing τ_B) to erase checkpoint memory without blocking checkpoint activation. Such a drug would be invisible to standard kinase inhibitor screens, which measure catalytic inhibition.

4. Neurodegeneration

4.1 The failure of the amyloid hypothesis

Alzheimer’s disease has been attacked through the amyloid hypothesis for three decades. Drugs that successfully clear $A\beta$ plaques (aducanumab, lecanemab, donanemab) have failed to restore cognition. The central hypothesis appears to target a symptom rather than the mechanism.

4.2 CaMKII as the canonical C1–C4 system

CaMKII — the canonical molecular memory device in neurons — satisfies C1–C4:

- **C1:** The kinase domain is coupled to the regulatory/association domain through the autoinhibitory segment
- **C2:** Autophosphorylation at T286 creates a bistable switch with persistence time of minutes to hours — far longer than individual calcium transients ($\sim$ ms)
- **C3:** CaMKII forms dodecameric holoenzymes with 12 subunits, each independently phosphorylatable — $2^{12} = 4,096$ distinct modification states

CaMKII’s memory function (long-term potentiation) is the *intended* biological use of non-Markovian dynamics. Neurodegeneration involves pathological perturbation of this memory system.

4.3 The OI prediction

Neurodegenerative disease involves pathological alteration of τ_B in synaptic signaling kinases:

- **Normal aging:** Gradual increase in τ_B (slower conformational relaxation due to oxidative damage) $\rightarrow$ excessive memory retention $\rightarrow$ synaptic rigidity $\rightarrow$ reduced plasticity
- **Alzheimer’s:** $A\beta$ oligomers interact with CaMKII and alter its regulatory domain dynamics, shifting $\tau_B \rightarrow$ pathological memory states that drive tau hyperphosphorylation through downstream cascades (including Chk1–CIP2A–PP2A)
- **Parkinson’s:** α -synuclein aggregates alter the conformational dynamics of LRRK2 and other PD-associated kinases

4.4 Therapeutic implication

τ_B -normalizing drugs. Instead of targeting the misfolded proteins themselves (which has largely failed clinically), target the *altered memory timescale* of signaling kinases. A drug that restores τ_B to its normal value could preserve synaptic function without clearing aggregates.

Testable prediction: Measure CaMKII conformational dynamics (by FRET or HDX-MS) in neurons exposed to $A\beta$ oligomers vs. controls. The framework predicts that $A\beta$ shifts τ_B of CaMKII’s regulatory domain. A compound that reverses this τ_B shift should rescue synaptic plasticity (measurable by LTP induction) independently of $A\beta$ clearance.

4.5 PTSD as τ_B pathology with reconsolidation as therapy

PTSD involves abnormally persistent and intrusive memory of traumatic events. From the framework’s perspective, this is a representative case of *failure of τ_B to be appropriately finite* at the relevant memory layer — the trauma-encoded trace persists at a layer where it should have decayed and continues to drive visible-sector dynamics (intrusive re-experiencing, hyperarousal).

The clinical efficacy of the reconsolidation paradigm — recalling a traumatic memory under propranolol blockade reduces its later emotional valence — is directly framework-predicted (§2.4). The mechanism is structural: recall opens the memory to obligatory re-storage, and pharmacologically blocking the noradrenergic re-encoding selectively erases the emotional content while preserving the declarative trace.

Predictions for further development:

- The same approach should generalize to other layers. Targeting memory layers slower than noradrenergic emotional encoding (e.g., the cortical-layer trace itself) would require longer-acting interventions during longer reconsolidation windows.
- The optimal interval between reconsolidation-and-blockade sessions should match the targeted layer’s τ_B . Single-session protocols are likely sub-optimal for memories that have consolidated to slower layers.
- Layer-selective drugs (currently lacking) would substantially improve efficacy. Propranolol works because noradrenergic encoding has a specific molecular substrate; analogous compounds for other encoding modalities are not yet available.

4.6 Direct mechanical access via low-intensity focused ultrasound

Standard pharmacology writes to *chemical* degrees of freedom (receptor occupancy, phosphorylation state). It cannot directly access the conformational and mechanical dynamics that constitute τ_B at the molecular layer. Low-intensity focused ultrasound (LIFU) is the first therapeutic modality that *directly* perturbs these mechanical degrees of freedom: acoustic radiation force acts on mechanosensitive ion channels (Piezo1, TRAAK, TRP family) and on membrane conformational states without requiring receptor binding.

This makes LIFU structurally aligned with the framework in a way pharmacology is not. It is a τ_B -*writer* rather than a catalytic inhibitor.

Clinical evidence consistent with the framework:

- The 2025 Korean trial (Ye et al., J. Neurosurgery) showed cognitive improvement in Alzheimer’s patients from focused ultrasound BBB opening *without* concurrent drug administration. This is unexpected on the amyloid hypothesis but framework-consistent if ultrasound directly normalizes molecular-memory substrate dynamics (e.g., CaMKII regulatory-domain kinetics) independent of amyloid clearance.
- LIFU has shown reversible neuromodulation effects in the anterior limb of the internal capsule (depression target), consistent with direct perturbation of axonal conduction via mechanosensitive K2P channels at the nodes of Ranvier.

Framework-specific predictions for LIFU (extension-level):

- Pulse repetition frequency should match characteristic molecular τ_B values of the target substrate. CaMKII intervention should use minute-scale to hour-scale pulse trains, not continuous sonication.
- Schedule dependence should follow the same τ_B -matched-interval logic as drug scheduling (§3.3, §5.3). Current LIFU protocols use month-scale intervals, which the framework predicts is too slow for molecular-layer targeting.
- Therapeutic window should be biphasic: low intensity for τ_B normalization, high intensity for non-specific mechanical damage. The window between is framework-predicted to be narrow.

These predictions are not currently how LIFU parameters are selected clinically (parameters are largely empirical). Framework-informed trials would systematically scan PRF and intensity against measured molecular τ_B values.

4.7 The brain memory hierarchy (extension)

The five-layer extension introduced in §2.2 — synaptic, circuit, systems, cortical — applies most directly to the brain. Each layer is the hidden sector for the layer above; each layer has its own τ_B inherited from but coarse-grained relative to the layer below.

This hierarchical picture provides natural framework-level interpretations of several established neuroscience phenomena:

- **Long-term potentiation** is the canonical example of cross-layer information transfer. A fast input (Ca^{2+} transient) writes to a fast molecular memory (CaMKII), which writes to synaptic-layer state (AMPA receptor density), which writes to slower structural state (spine enlargement), which eventually writes to the slowest available layer (gene expression and protein synthesis). Each transition is information transfer between hidden-sector layers.
- **Sleep consolidation** is the period during which the systems-layer τ_B becomes accessible. Waking input dominates the cortex; sleep allows the natural relaxation dynamics of the slow hidden sector to proceed. The replay events documented in slow-wave sleep and REM are the framework-required mechanism for inter-layer transfer.
- **Phenomenology of remembering across timescales** — recent memories with rich sensory detail, medium-age memories with preserved event structure but fading sensory content, old memories with semantic gist but reconstructed episodic detail — maps onto the layered τ_B architecture quantitatively.

These are *extensions* of the framework’s primary content (molecular memory and disease pharmacology) rather than independent derivations. They provide a coherent interpretive frame for brain memory phenomena without claiming the framework derives the specific phenomenology of episodic recall, the hippocampal indexing role, or working memory as distinct from long-term storage.

These remain biological and cognitive-science questions outside the framework’s current scope.

5. Antibiotic Resistance

5.1 Persister cells as memory accumulation

Persister cells survive antibiotic treatment through transient phenotypic tolerance, not genetic resistance. The SOS response pathway — regulated by RecA — satisfies C1–C4: RecA’s nucleotide binding (fast) is allosterically coupled to filament formation on ssDNA (slow, seconds to minutes), with filaments extending over hundreds of bases (exponentially large state space).

5.2 The OI prediction

Persister formation is not random switching — it is the accumulation of SOS memory past a threshold. Each antibiotic exposure writes information into the RecA filament state, which persists and modulates subsequent responses.

5.3 Memory-optimized antibiotic scheduling

The interval between doses should be optimized for the bacterial SOS memory timescale (τ_B of RecA filament dynamics), not just the drug’s pharmacokinetic half-life:

- *Dose interval* $< \tau_B$: SOS memory accumulates $\rightarrow$ drives persister formation (counterproductive)
- *Dose interval* $> \tau_B$: Memory decays $\rightarrow$ each dose is independent (no accumulation)
- *Optimal interval* $\approx \tau_B$: Partial memory decay prevents persister threshold crossing while residual memory maintains sensitization

Clinical studies showing pulsed antibiotic regimens outperforming continuous regimens are consistent with this prediction — the mechanism is disruption of bacterial memory, not pharmacokinetic optimization.

6. Immunotherapy and T Cell Exhaustion

6.1 The exhaustion problem

PD-1/PD-L1 checkpoint inhibitors work in only ~ 20 – 30% of patients. The primary barrier is T cell exhaustion — progressive loss of effector function due to chronic antigen stimulation.

6.2 TCR signaling as a C1–C4 system

The TCR signaling cascade ($\text{Lck} \rightarrow \text{ZAP-70} \rightarrow \text{LAT} \rightarrow \text{downstream effectors}$) involves kinases with regulatory domains satisfying C1–C4, with multiple phosphorylation sites creating a combinatorial memory register.

6.3 The OI prediction

T cell exhaustion begins as accumulated non-Markovian memory in the TCR signaling kinases. Each antigen encounter writes PTM/conformational information. In acute infection, this memory resets. In chronic stimulation (persistent tumor antigen), memory accumulates and progressively shifts the signaling dynamics toward exhaustion. The transcriptional changes (TOX upregulation) are *downstream consequences* of the accumulated kinase memory, not the primary cause.

6.4 Therapeutic implication

Memory erasure + checkpoint inhibition. PD-1 inhibitors release the inhibitory brake but do not erase accumulated TCR signaling memory. Combining PD-1 inhibitor with a “memory eraser” — a drug that accelerates conformational relaxation of TCR signaling kinases (decreasing τ_B) — should rescue T cell function more effectively than PD-1 inhibitor alone.

7. Cardiac Pharmacology

7.1 Use-dependent block as non-Markovian channel dynamics

Antiarrhythmic drugs show use-dependent block — efficacy depends on heart rate, not just plasma concentration. Voltage-gated ion channels (hERG, Nav1.5, Cav1.2) satisfy C1–C4: the pore domain (fast gating, $\sim \mu\text{s}$) is coupled to voltage-sensing and regulatory domains with slow inactivation timescales (~ 100 ms to seconds) and multiple inactivation states (C3).

7.2 Therapeutic implication

Heart-rate-adapted dosing. Antiarrhythmic dosing should be adapted to the patient’s *heart rate pattern*, which determines the channel memory state. A patient with persistent tachycardia has channels in a different memory state than a patient with intermittent tachycardia. For drugs with strong use-dependence (flecainide, lidocaine), this could substantially improve the therapeutic window by avoiding pro-arrhythmic effects at high heart rates.

8. Autoimmune Disease

8.1 Disproportionate efficacy of partial JAK inhibition

JAK inhibitors (tofacitinib, baricitinib) show a nonlinear dose-response: 50% reduction in JAK activity produces $> 80\%$ reduction in disease activity. JAK kinases have a pseudokinase domain (JH2) that regulates the kinase domain (JH1), satisfying C1–C4.

8.2 The OI prediction

Partial inhibition disrupts the *memory structure* of JAK signaling without fully blocking transduction. The pathological signaling in autoimmune disease involves accumulated PTM/conformational memory from chronic cytokine stimulation. Partial inhibition erases this memory faster than it blocks acute signaling, producing disproportionate reduction in the chronic (memory-dependent) disease component.

8.3 Therapeutic implication

Memory-selective JAK modulation. Drugs that selectively accelerate JH2 conformational relaxation — erasing accumulated inflammatory memory without blocking acute immune responses — would predict wider therapeutic windows than current JAK inhibitors: disease activity reduced while acute immune competence is preserved.

9. Epigenetics as the Biological Hidden Sector

9.1 The C1–C4 architecture of chromatin

Epigenetic regulation satisfies C1–C4 exactly. The transcriptional machinery (visible sector) is coupled (C1) to the chromatin state (hidden sector), which changes slowly relative to transcription (C2) and has astronomically large capacity (C3: $\sim 2^{28 \times 10^6}$ possible CpG methylation patterns).

The chromatin hidden sector operates at multiple nested timescales:

Layer	Mechanism	τ_B	Biological function
1 (fastest)	Histone acetylation	Minutes–hours	Rapid signal response
2	Histone methylation	Hours–days	Lineage commitment
3	DNA methylation	Cell generations	Cell type memory

Layer	Mechanism	τ_B	Biological function
4	Chromatin compaction	Cell generations	Permanent silencing
5 (slowest)	Germline methylation	Transgenerational	Intergenerational memory

9.2 Cancer as pathological epigenetic memory

Cancer is, in OI language, a disease of the epigenetic hidden sector. The malignant transcriptional program is maintained by aberrant epigenetic memory — stable patterns of methylation and histone modifications that lock the cell into a proliferative program. Disrupting the memory *structure* (the stability of the hidden sector) is predicted to be more selective than disrupting the memory *content* (reactivating specific genes).

9.3 Aging as memory accumulation

The epigenetic clock (Horvath 2013) quantifies progressive accumulation of methylation marks associated with declining function. In OI language, aging is memory accumulation in the slowest epigenetic layers beyond the cell’s ability to maintain homeostasis. Partial reprogramming (transient Yamanaka factor expression) erases recent epigenetic memory while preserving deeper developmental identity — the framework predicts a critical pulse duration matching τ_B of aging-associated marks.

9.4 The substratum-emergent operator distinction at chromatin

The companion paper’s treatment of the Standard Model identifies a *substratum-emergent operator distinction* [SM §8.7]: operators that look identical at the emergent level can have distinct conservation properties depending on whether they act at the substratum (substrate) level or at the emergent (coarse-grained, observable) level. The canonical example: baryon number is exactly conserved at the OI substratum level (a property of the bilinear, block-diagonal lattice dynamics) but is violated at the emergent level through sphaleron processes that have no substratum analog.

This distinction has a direct biological analog at the chromatin layer, which is empirically established by current biomedical literature. The two operators are:

- **Substratum-level memory operator:** the molecular pattern of DNA methylation, histone modifications, and chromatin-state assignments. This is a directly measurable molecular substrate.
- **Emergent-level memory operator:** the gene expression program, the transcriptional phenotype, the cellular function readout. This is what cells display and what clinical phenotypes measure.

The two operators are related by the trace-out, which here is the cell’s signal-transduction and transcription machinery integrating fast molecular fluctuations into slower cellular response. Critically, the operators can dissociate: **substrate-level memory can be preserved while the emergent-level operator is disrupted, and vice versa.**

Three lines of empirical evidence confirm this distinction operates in biology:

- (i) *iPSC reprogramming residuals* (Lin et al. 2024, doi:10.3389/fcell.2024.1306530). When somatic cells are reprogrammed to induced pluripotent stem cells, the cellular phenotype changes (now pluripotent) while substantial residual epigenetic memory of the source cell type persists. The substrate-level operator retains source-cell-type information; the emergent-level operator (cell identity) has flipped. The two operators are dissociated. The framework predicts this should be a structural feature of any reprogramming protocol, not a contingent reprogramming limitation.
- (ii) *Bivalent chromatin states* (Bernstein et al. 2006; reviewed in Voigt-Tee-Reinberg 2013). Embryonic stem cells maintain “bivalent” chromatin at lineage-control genes — simultaneous presence of activating H3K4me3 and repressive H3K27me3 marks. The substrate-level operators (the two mark types) coexist with definite conservation properties at the molecular level; the emergent-level operator (gene expression) is held at zero by their balanced presence. This is a *multi-operator substratum* structure with poised emergent output — exactly the kind of substratum-emergent operator structure characterized in [SM §8.7], realized empirically at the chromatin level.
- (iii) *Multi-timescale chromatin dynamics* (Stasevich et al. 2014; Saxton et al. 2023). Single-molecule and single-cell measurements directly observe distinct dynamics on subsecond and minutes timescales for the same histone modifications on the same nucleosomes. Histone acetylation modulates transcriptional burst frequency rather than burst size (Nicolas-Phillips et al. 2018), a stochastic-emergent-output pattern consistent with the framework’s prediction that emergent operators have specific functional relationships to substratum operators rather than being identical to them.

These three phenomena instantiate a single structural pattern — the substratum-emergent operator distinction from [SM §8.7] — at the chromatin level. The pattern is the same one that organizes the SM’s treatment of strong CP, baryon number, sphalerons, and anomaly matching; at the chromatin layer it organizes methylation maintenance, bivalent poised states, and transcriptional burst dynamics.

9.5 Reader, writer, and eraser pharmacology as the operational substratum-emergent axis

Current epigenetic therapeutics are organized along three axes — drugs targeting **writers** (DNA methyltransferase inhibitors like 5-azacytidine, decitabine), **erasers** (histone deacetylase inhibitors like vorinostat, romidepsin), and **readers** (BET bromodomain inhibitors like JQ1, ZEN-3694, BI 894999). This three-axis organization is the substratum-emergent operator distinction made operational at the drug-development level: the substratum-level memory operator at chromatin is constructed and maintained by readers, writers, and erasers as the molecular machinery, and pharmacology can target any of these three functional roles independently.

This three-axis organization follows structurally from the operator distinction: any therapeutic approach to a substratum-emergent operator pair must target the substratum maintenance machinery (writers and erasers, which establish and remove substrate marks) or the substratum-to-emergent translation machinery (readers, which couple substrate state to downstream function). The three axes correspond to the three distinct functional roles in maintaining a substratum-emergent operator pair.

Combination synergy with scheduling sensitivity. The framework predicts that combinations across the reader/writer/eraser axes should show specific compositional advantages, with scheduling determined by the underlying substratum-level τ_B at each operator. Current clinical evidence:

- *Writer + reader* combination (p300/CBP HAT inhibitor + BET inhibitor) shows synergistic anti-tumor effects in NUT carcinoma (Bauer et al. 2020; Marek et al. 2022, doi:10.1186/s40478-022-01373-0).
- *Eraser + reader* combination (HDAC inhibitor + BET inhibitor) is under clinical investigation in NUT carcinoma and other contexts.
- *Writer + eraser* combination (DNMT inhibitor + HDAC inhibitor) shows synergy preclinically and clinical activity in MDS/AML, with **scheduling as a critical variable** — multiple Phase II trials have addressed sequencing and timing as primary endpoints rather than dose-finding alone (Quintás-Cardama et al. 2011, doi:10.4155/fmc.11.13; Prebet et al. 2014).

The clinical pattern — synergy exists, scheduling matters substantially, global mark changes don't deterministically predict clinical response — is consistent with the framework's prediction that the *timing* of substratum modifications relative to read-back cycles determines emergent response, exactly as the reconsolidation-window paradigm of §2.4 implies. Within the framework, the open scheduling questions in the literature have a structural resolution: optimal combination scheduling should be set by the τ_B of each substratum operator (DNMT effects persist on the methylation timescale, $\tau_B \sim$ days; HDAC effects persist on the acetylation timescale, $\tau_B \sim$ minutes-hours; reader-targeted effects on the chromatin-binding timescale, $\tau_B \sim$ seconds-minutes). Optimal sequencing aligns the fast operator's modification with the slow operator's

read-back window.

This is a substantive prediction: empirical optimization of combination scheduling should converge on τ_B -matched intervals between agents acting on different timescales, not on dose-response or maximum-tolerated-dose optimization alone.

10. Genetic Disorders: Non-Markovian Treatment Management

10.1 The boundary of the framework

Pure loss-of-function genetic disorders — hemophilia (absent Factor VIII/IX), cystic fibrosis (absent/misfolded CFTR), PKU (deficient phenylalanine hydroxylase), Tay-Sachs (absent hexosaminidase A) — are not memory diseases. The core problem is that a protein is missing or non-functional. There is no τ_B to normalize when the protein does not exist. The framework does not claim otherwise.

However, the *management* of genetic disorders — replacement therapy scheduling, immune responses to replacement proteins, gene therapy durability, and compensatory pathway dynamics — involves non-Markovian dynamics at every level. The framework’s contribution is to these surrounding problems.

10.2 Coagulation cascade memory and factor replacement

The clotting cascade (intrinsic and extrinsic pathways converging at Factor X $\rightarrow$ thrombin $\rightarrow$ fibrin) is a multi-kinase signaling cascade with C1–C4 architecture. Thrombin activates Factor V and Factor VIII (positive feedback), while antithrombin and TFPI provide negative feedback on different timescales. The cascade has memory: prior subthreshold activation primes it for faster response to subsequent triggers.

In hemophilia patients receiving replacement factor, the cascade operates in a partially-reconstituted state where the memory dynamics differ from normal. The framework predicts that the *timing* of replacement factor dosing relative to the cascade’s memory state matters — not just the trough factor level. A patient who bleeds and partially activates the cascade before receiving factor concentrate is in a different memory state than a patient receiving prophylactic factor on schedule.

Prediction: Non-Markovian dosing that accounts for the cascade memory state (recent bleeding events, prior subthreshold activation) should reduce breakthrough bleeding rates compared to standard pharmacokinetic-based dosing at equivalent total factor consumption. *Test:* Correlate breakthrough bleeding frequency with the interval between the most recent cascade activation event and the next scheduled prophylactic dose — OI predicts this interval matters; standard PK models predict only trough level matters.

10.3 Inhibitor development as immune memory

Approximately 30% of severe hemophilia A patients develop inhibitory antibodies against replacement Factor VIII — the most serious complication in hemophilia management. This is squarely within the framework’s territory. The immune system’s response to repeated Factor VIII exposure is non-Markovian: each infusion writes information into the B cell and T cell memory compartments through the same TCR/BCR signaling cascades described in §6.

Inhibitor development is not random — it depends on the patient’s exposure history, the timing and intensity of prior infusions, and concurrent immune status. The same architecture applies to immune reactions against enzyme replacement therapy in lysosomal storage diseases (Gaucher, Fabry, Pompe) and against gene therapy vectors (anti-AAV antibodies).

Prediction: Inhibitor risk correlates with the *schedule* of factor exposure (which determines immune memory state), not just cumulative dose. Specifically, initial exposure regimens with intervals tuned to the τ_B of the relevant B cell memory compartment should show lower inhibitor rates than standard intensive prophylaxis schedules. *Test:* Compare inhibitor development rates in patients started on different prophylaxis schedules (daily low-dose vs. twice-weekly standard-dose) at equivalent cumulative factor exposure over the first year.

Immune tolerance induction (ITI) — frequent high-dose factor infusion to overcome inhibitors — works by overwriting pathological immune memory. The framework predicts that ITI success depends on matching the infusion schedule to τ_B of the relevant B cell memory compartment, and that optimized ITI schedules calculated from measured B cell memory kinetics would show higher success rates.

10.4 Gene therapy durability as epigenetic memory

Gene therapies for hemophilia (Hemgenix for Factor IX, Roctavian for Factor VIII) deliver functional gene copies via AAV vectors. The clinical challenge is variable and sometimes declining transgene expression across patients. The framework’s epigenetics analysis (§9) is directly relevant: the transgene’s expression level is determined by the chromatin state at the integration site — which layer of the epigenetic hidden sector it occupies.

If the transgene integrates into a region with high- τ_B silencing marks (stable DNA methylation, compact heterochromatin), expression will be silenced over time. If it integrates into a region with low- τ_B activating marks (histone acetylation), expression will be maintained but variable.

Prediction: Transgene expression durability correlates with the epigenetic τ_B at the integration site — measurable by chromatin accessibility (ATAC-seq) and histone mark stability at the locus. Combining gene therapy with targeted epigenetic modifiers that stabilize activating marks at the transgene locus (increasing

local τ_B of activating marks without affecting other loci) should improve long-term expression. *Test:* In animal models of hemophilia gene therapy, measure expression decline rate alongside integration-site-specific chromatin dynamics. Sites with higher τ_B of activating marks should show slower expression decline.

10.5 Compensatory pathway memory

In many genetic disorders, compensatory pathways develop over time and partially mask the deficiency. In spinal muscular atrophy (SMA), for example, SMN2 partially compensates for lost SMN1 — but the degree of compensation is variable and depends on the patient’s developmental history. In sickle cell disease, fetal hemoglobin (HbF) reactivation provides partial compensation, with levels influenced by the patient’s prior erythropoietic history and epigenetic state.

These compensatory responses are non-Markovian: their current strength depends on the cumulative history of demands placed on them, stored in the epigenetic and signaling memory of the relevant cell populations. The framework predicts that interventions to boost compensatory pathways (HbF inducers in sickle cell, SMN2 upregulators in SMA) would be more effective if timed to the memory state of the compensatory system — analogous to the memory-priming strategy in cancer pharmacology (§3.3).

10.6 The general principle

For genetic disorders, the framework’s contribution is not to the genetic defect itself but to three categories of treatment challenge:

Treatment challenge	C1–C4 system	OI prediction
Replacement therapy scheduling	Coagulation cascade / metabolic pathway	Timing relative to cascade memory state improves efficacy
Immune reactions to therapy	TCR/BCR signaling cascades	Inhibitor risk depends on exposure <i>schedule</i> , not just dose
Gene therapy durability	Chromatin state at integration site	Expression correlates with local epigenetic τ_B
Compensatory pathway optimization	Epigenetic regulation of backup genes	Compensation strength depends on developmental history
Enzyme replacement therapy tolerance	B cell memory compartment	Tolerance induction schedule matches immune τ_B

The summary: genetic disorders are not memory diseases. But the treatments for genetic disorders operate through biological systems that are memory sys-

tems. Optimizing these treatments for the non-Markovian dynamics of the underlying biology is a distinct and testable therapeutic strategy.

10.7 Reader-writer disorders as substratum-emergent operator disruption

A specific class of genetic disorders falls within the framework’s scope of memory diseases rather than the §10.1 surrounding-treatment category. These are disorders caused by loss-of-function mutations in the *readers*, *writers*, or *erasers* of epigenetic memory rather than disorders of the substrate itself. In framework language, these are disorders of the substratum-to-emergent translation machinery, not disorders of the substratum.

Rett syndrome as the paradigm. Rett syndrome is caused by loss-of-function mutations in *MECP2*, a methyl-CpG-binding protein — a *reader* of DNA methylation. The DNA methylation substrate (the molecular memory substratum) is intact in Rett patients; what is lost is the protein that translates substrate state into downstream cellular function. Within the substratum-emergent operator framework of §9.4, Rett syndrome instantiates the dissociation: substratum-level memory operator preserved, emergent-level operator disrupted because the trace-out machinery (*MECP2*) is defective.

The framework predicts that phenotypes from such disorders should be reversible by restoration of the missing reader/writer, because the underlying substrate is preserved and the trace-out is what was broken. This is empirically confirmed for Rett by mouse genetic reactivation studies: Guy et al. (2007, *Science* 315:1143-1147) demonstrated that re-expression of endogenous *Mecp2* in adult mutant mice reverses neurological phenotype, even in animals with significant established pathology. Long-term potentiation in the hippocampus is restored; soma size normalizes; synaptic structures recover.

Generalization. The same structural logic applies to other monogenic disorders affecting epigenetic reader/writer/eraser machinery, even though clinical convention typically treats them as developmental and irreversible:

- *Fragile X syndrome* (FMR1 protein loss, mRNA-binding reader): substrate (RNA target landscape) is preserved; loss is in the reader. Reactivation strategies are now in clinical trials.
- *Angelman syndrome* (UBE3A ubiquitin ligase loss, a writer of degradation marks): paternal allele substrate is preserved; loss is in the writer of one parental imprint. UBE3A reactivation via antisense oligonucleotides is in clinical investigation.
- *Kabuki syndrome* (KMT2D writer disruption, histone H3K4 methyltransferase): substrate landscape is preserved; loss is in the writer. Phenotypic plasticity in mouse models supports the reversibility prediction.
- *ATR-X syndrome* (ATR-X chromatin remodeler disruption): substrate is preserved; loss is in machinery that reads chromatin state. Variable severity correlates with residual ATRX function, consistent with emergent op-

erator output being graded by reader/writer machinery function.

Therapeutic principle. For reader-writer disorders, the structural therapeutic axis is *restoration of the missing reader/writer* rather than substrate modification or symptomatic management. Gene therapy (AAV-based MECP2 replacement in Rett; UBE3A reactivation in Angelman) and direct pharmacological replacement (trofinetide for Rett, FDA-approved 2023) target this axis. The framework predicts that the therapeutic window is wider than for substrate-modifying interventions, and that response correlates with timing of intervention relative to substrate state at each affected locus. The Rett mouse model results — phenotypic reversal in adult animals, not just in juveniles — are consistent with this wide therapeutic window.

Reader-writer disorders therefore are memory disorders in the framework’s sense (substratum preserved, emergent operator disrupted), distinct from other genetic disorders where the defect interacts with memory systems only through treatment (§10.1–§10.6).

11. Distinguishing Predictions

The following predictions follow from the non-Markovian (memory) architecture. Two cautions frame them. First, the relevant comparison is not “Markovian pharmacology” — modern pharmacology already models use-dependence, hysteresis, priming, and chronotherapy — but *whether the framework’s prediction is distinct from standard non-Markovian accounts*; several predictions below (positive autocorrelation of successive events, stretched-exponential single-molecule waiting times, history-dependent antiarrhythmic efficacy) are reproduced by classical dynamic-disorder models and are therefore consistency checks rather than distinguishing tests. Second, these are tests of *memory* (grade-ii); they do not test the framework’s distinctive P-indivisibility (grade-iii, information backflow), which — as §2 and [SM] note — is a substratum-level property. The genuinely distinguishing content is concentrated in the *quantitative* predictions that tie a clinical observable to an independently measurable molecular timescale or quantity: the LIFU pulse-repetition-frequency resonance at $1/\tau_B$ (Prediction 8b), the dependence of DNMT-inhibitor efficacy on methylation *entropy* rather than level (Prediction 16), τ_B -resonant dosing (Prediction 9), and the decoupling predictions 7 and 12. In those cases the discriminating quantity is the measured τ_B (or entropy) value, not the resonance concept itself, which is generic.

11.1 Cancer

Prediction 1: Resistance mutations to regulatory-domain-targeting Chk1 inhibitors cluster in regions that alter the *slowest conformational modes* (τ_B), not k_{cat} or K_m . *Test:* Measure catalytic parameters and conformational dynamics

for resistance mutants. Standard models predict correlated changes; OI predicts uncorrelated changes.

Prediction 2: Checkpoint recovery times are temporally correlated across successive damage events in single cells, decaying on timescale τ_B . *Test:* Time-lapse imaging with Chk1 FRET biosensor under repeated damage. Standard models predict zero autocorrelation; OI predicts positive autocorrelation.

Prediction 3: Cells with prior damage-and-recovery cycles are *more* sensitive to Chk1 inhibition than naive cells at identical current damage levels. *Test:* Compare clonogenic survival with and without prior low-dose cycling. Standard models predict identical survival; OI predicts greater kill in pre-treated cells.

Prediction 4: HDAC inhibitors between radiation fractions specifically abolish the adaptive response by erasing chromatin memory. *Test:* Three-arm experiment comparing inter-fraction vs. pre-treatment HDAC inhibitor. Standard models predict similar radiosensitization; OI predicts greater cell kill from inter-fraction administration.

Prediction 5: Single-molecule Chk1 turnover waiting times follow a stretched-exponential distribution $P(t) \sim \exp(-(t/\tau)^\beta)$ with $\beta < 1$. *Test:* Single-molecule fluorescence assay. Standard models predict $\beta = 1$ (exponential); OI predicts $\beta < 1$ with value calculable from τ_S/τ_B .

11.2 Neurodegeneration

Prediction 6: A β oligomers shift CaMKII's regulatory domain τ_B , measurable by FRET or HDX-MS. *Test:* Compare CaMKII conformational dynamics in A β -exposed vs. control neurons.

Prediction 7: Reversing the τ_B shift rescues LTP independently of A β clearance. *Test:* Treat A β -exposed neurons with a CaMKII conformational modulator; measure LTP induction.

Prediction 8: Normal aging shows gradual increase in CaMKII τ_B , correlating with synaptic rigidity. *Test:* Measure CaMKII regulatory domain dynamics as a function of age in animal models.

Prediction 8a: PTSD reconsolidation efficacy is layer-specific. Propranolol-during-recall blocks emotional re-encoding (noradrenergic substrate); analogous compounds for other encoding modalities should produce dissociable effects on different memory components. *Test:* Compare cognitive vs. emotional vs. sensory components of PTSD memory after layer-targeted reconsolidation interventions.

Prediction 8b: LIFU pulse repetition frequency matched to molecular τ_B produces stronger neuromodulation than empirically-chosen PRF. *Test:* Systematically scan PRF in animal models against measured CaMKII (or other target) τ_B ; predict response peak at PRF matching $1/\tau_B$.

Prediction 8c: LIFU-only intervention (no antibody) in early Alzheimer’s produces cognitive benefit proportional to τ_B normalization, measurable by post-LIFU CaMKII conformational dynamics. *Test:* The 2025 Korean trial design extended with mechanistic assays.

11.3 Antibiotic resistance

Prediction 9: Pulsed antibiotic dosing with interval $\approx \tau_B$ of RecA filament dynamics outperforms continuous dosing and pulsing at intervals $\ll \tau_B$ or $\gg \tau_B$. *Test:* Measure persister fraction under different pulsing intervals.

Prediction 10: RecA filament length/conformation at the time of the second antibiotic pulse predicts survival probability. *Test:* Single-cell imaging of RecA-GFP filaments during pulsed antibiotic treatment.

11.4 Immunotherapy

Prediction 11: Temporal correlations between successive calcium flux events in repeatedly stimulated T cells are positive and decay on τ_B of TCR signaling kinases. *Test:* Time-lapse calcium imaging under repeated anti-CD3 stimulation.

Prediction 12: A kinase regulatory domain “accelerator” (increasing Lck SH2 flexibility) delays functional exhaustion markers (PD-1, Tim-3, LAG-3) in ex vivo exhaustion assays.

11.5 Cardiac pharmacology

Prediction 13: Antiarrhythmic drug efficacy at fixed plasma concentration differs between patients with persistent vs. intermittent tachycardia, beyond what Markovian models predict. *Test:* Correlate drug efficacy with heart rate pattern history, not just current rate.

11.6 Autoimmune disease

Prediction 14: Partial JAK inhibition reverses chronic cytokine-induced gene expression changes more effectively than proportional reduction in acute STAT phosphorylation. *Test:* Compare chronic vs. acute transcriptional effects of partial vs. full JAK inhibition.

11.7 Epigenetics

Prediction 15: Gene expression autocorrelation across cell divisions correlates with τ_B of the dominant epigenetic mark: methylation-regulated genes show stronger memory than histone-regulated genes.

Prediction 16: DNMT inhibitor efficacy correlates with methylation entropy (memory stability), not methylation level.

Prediction 17: Pulsed DNMT inhibitor administration achieves equivalent tumor response at lower toxicity than continuous administration (by allowing normal tissue memory to recover between pulses).

Prediction 18: During iPSC reprogramming, epigenetic marks are erased in order of increasing τ_B : acetylation first, then methylation, then DNA methylation.

Prediction 19: Partial reprogramming (Yamanaka factor pulse) shows a critical duration: methylation age decreases smoothly with pulse length, while cell-type identity remains stable up to a sharp threshold corresponding to τ_B of developmental memory.

11.8 Allosteric pharmacology

Prediction 20: RGS4 allosteric inhibitor selectivity across RGS family members correlates with B-site conformational dynamics (τ_B), not sequence identity. *Test:* Plot IC_{50} against B-site flexibility (from MD simulation) vs. sequence identity.

Prediction 21: Different active-site inhibitors of the same kinase produce different effects on distant regulatory domains because each writes different conformational information into the memory structure.

11.9 Drug resistance

Prediction 22: Resistance mutations far from the active site work by altering the *memory capacity* (C3) or *memory timescale* (C2) of the enzyme, not by blocking catalysis. *Test:* Normal mode analysis of resistance mutants — OI predicts changes in slowest modes, not catalytic modes.

Prediction 23: Drug selectivity among family members (RGS, BTK, Abl) correlates with conformational dynamics, not binding-site sequence. Existing data on BTK inhibitor-specific regulatory effects (Joseph et al. 2020) and RGS4 selectivity (Blazer et al. 2010) already support this prediction.

11.10 Genetic disorder treatment

Prediction 24: In hemophilia patients on prophylaxis, breakthrough bleeding frequency correlates with the interval between the most recent cascade activation event and the next prophylactic dose — not just trough factor level. *Test:* Track bleed timing relative to prior subthreshold activations via thrombin generation assays.

Prediction 25: Inhibitor development rates in hemophilia A correlate with the *schedule* of initial Factor VIII exposure (which determines immune memory state), not just cumulative dose. *Test:* Compare inhibitor rates between daily low-dose vs. twice-weekly standard-dose prophylaxis at equivalent cumulative factor exposure.

Prediction 26: Gene therapy transgene expression durability correlates with epigenetic τ_B at the integration site. *Test:* In animal models, correlate expression decline rate with integration-site chromatin dynamics measured by ATAC-seq and histone mark stability.

12. Existing Evidence

Several predictions are already supported by published data:

Resistance through dynamics, not catalysis (Prediction 22). Taldaev et al. (PNAS, 2021) identified an Abl kinase imatinib-resistance mutation that preserves drug binding affinity but increases conformational dynamics — altering τ_B , not K_d . The BTK T316A mutation (Joseph et al., eLife, 2020) causes resistance by shifting the conformational ensemble, not by blocking drug binding. Deep mutational scanning of Src (Chakraborty et al. 2024) found that 25–45% of resistance mutations for each inhibitor were unique to that inhibitor’s conformational binding mode.

Selectivity through dynamics, not sequence (Prediction 23). RGS4 allosteric inhibitor selectivity does not correlate with B-site sequence conservation (Blazer et al. 2010). Five different BTK active-site inhibitors produce five qualitatively different effects on distant regulatory domains of the *same* protein (Joseph et al. 2020).

History-dependent checkpoint responses (Prediction 2). Single-cell checkpoint tracking (Chao et al., Cell Systems, 2017) found that responses depend on the cell’s exact cell cycle position, requiring “commitment points” modeled as additional Markov states — the OI interpretation is that these are continuous conformational memory, not discrete states.

Non-exponential enzyme kinetics (Prediction 5). Stretched-exponential waiting times have been observed for horseradish peroxidase (Edman & Rigler 2000) and cholesterol oxidase (Lu et al. 1998) — not yet for checkpoint kinases.

13. Discussion

13.1 The unifying principle

Every application follows the same logic: (1) a disease process involves a signaling molecule with C1–C4 architecture; (2) the disease state involves pathological accumulation or alteration of non-Markovian memory; (3) current drugs target catalytic function; (4) the framework identifies memory structure as a pharmacologically distinct target; (5) memory-targeted therapy predicts wider therapeutic windows because the memory asymmetry between disease and normal tissue is more specific than the catalytic asymmetry.

Disease	Memory mechanism	Current target	OI target
Cancer (Chk1)	Checkpoint PTM accumulation	Kinase catalysis	Regulatory domain τ_B
Alzheimer's	CaMKII τ_B shift by $A\beta$	Protein aggregates	τ_B normalization
Parkinson's	LRRK2 τ_B shift by α -syn	Protein aggregates	τ_B normalization
Antibiotic resistance	SOS/RecA filament memory	Bacterial growth	RecA memory timescale
T cell exhaustion	TCR kinase PTM accumulation	PD-1 brake	TCR kinase τ_B
Cardiac arrhythmia	Ion channel slow inactivation	Channel block	Heart-rate-adapted dosing
Autoimmune disease	JAK-STAT memory	JAK catalysis	JH2 conformational relaxation
Cancer (epigenetic)	Aberrant methylation/histone marks	Gene reactivation	Memory stability (τ_B)
Aging	Epigenetic clock accumulation	Symptom management	τ_B -selective erasure
Hemophilia (treatment)	Coagulation cascade memory	PK-based factor dosing	Cascade memory-state dosing
Genetic disorders (immune)	B/T cell memory from replacement therapy	Dose reduction	Exposure schedule matching immune τ_B
Gene therapy durability	Chromatin state at transgene locus	Dose escalation	Epigenetic τ_B stabilization at locus

13.2 Implications for drug discovery

The framework identifies a new class of drug targets — conformational memory timescale — that current screening assays do not measure. A “memory-targeted” drug screen would assay temporal correlations in enzyme activity, not steady-state kinetic parameters. This requires single-molecule or single-cell time-series measurements, which are technically mature but not routinely used in drug discovery.

13.3 Implications for clinical trial design

Two predictions are immediately testable with existing drugs and standard clinical infrastructure: (1) memory-selective scheduling of gemcitabine + Chk1 inhibitor (modified dosing protocol, no new drugs); (2) inter-fraction HDAC inhibitor for radiation adaptive response erasure (standard radiation biology experiment). Both could be evaluated in Phase I/II settings with minimal additional cost.

13.4 Connection to fundamental physics

The mathematical structure underlying these predictions is identical to the theorem that identifies the quantum representation — and the memory sector that makes it nontrivial — with the statistics of embedded observation [1]. The C1–C4 conditions that produce non-Markovian enzyme dynamics are the same conditions that produce, at the cosmological scale, the memory-bearing dynamics quantum mechanics represents. The read-write cycle of a kinase interacting with its regulatory domain is structurally isomorphic to the read-write cycle of an observer interacting with the hidden sector across the cosmological horizon. This connection is not metaphorical — the characterization theorem applies to any system satisfying C1–C4, regardless of scale. The biological instantiations are classical (no quantum coherence is required or invoked), but the mathematical architecture is the same.

The *substratum-emergent operator distinction* developed in the companion SM paper [SM §8.7] has a direct biological analog at the chromatin layer (§9.4). In the SM context, the distinction identifies operators that are exactly conserved at the substratum (bilinear lattice) level but violated at the emergent EFT level through non-perturbative mechanisms — the canonical example is baryon-number conservation. At the chromatin layer, the same structural distinction identifies operators that operate on the molecular substrate (DNA methylation, histone modifications, chromatin state) versus operators that operate on the emergent cellular phenotype (gene expression, cell identity, behavior). The two operators can dissociate, with iPSC reprogramming residuals, bivalent chromatin states, and Rett-class reader-writer disorders all being instances of this dissociation. The same mathematical pattern characterizes fundamental physics and chromatin biology — not by analogy, but as instances of the same structural feature of trace-out dynamics.

The *unified naturalness pattern* developed in [SM §8.8] extends to biology in the form developed in §9.1: biological “fine-tuning” of τ_B/τ_S ratios at memory layers is the persistence requirement (C2) in biological realization — the τ ratio one route — rather than evolutionarily contingent. The same mechanism that dissolves the hierarchy and cosmological-constant naturalness problems in physics — substratum-level structural constraints forcing specific values of emergent parameters — operates at the biological memory layers.

The same structural picture applied to computational analysis methods

is developed in the companion Bioinformatics paper [18], which addresses non-Markovian dynamics in single-cell trajectory inference, gene regulatory network inference, perturbation prediction, multi-omics integration, and cancer methylome classification. The framework’s reader-writer disorder class (§10.7) extends into specific bioinformatic patterns there.

Acknowledgements

During the preparation of this work, the author used Claude Opus 4.6 (Anthropic) and Gemini 3.1 Pro (Google) to assist in drafting, refining argumentation, and surveying the biomedical literature. The author reviewed and edited all content and takes full responsibility for the publication.

References

- [1] A. Maybaum, “The Incompleteness of Observation,” (2026).
- [2] J. Gunawardena, “A linear framework for time-scale separation in nonlinear biochemical systems,” *PLoS ONE* **7**, e36321 (2012).
- [3] K. K. Bhatt et al., “Essential function of Chk1 can be uncoupled from DNA damage checkpoint and replication control,” *PNAS* **105**, 20752 (2008).
- [4] L. A. Parsels et al., “Gemcitabine sensitization by Chk1 inhibition correlates with inhibition of a Rad51 DNA damage response in pancreatic cancer cells,” *Mol. Cancer Ther.* **8**, 45 (2009).
- [5] L. L. Blazer et al., “Reversible, allosteric small-molecule inhibitors of regulator of G protein signaling proteins,” *Mol. Pharmacol.* **78**, 524 (2010).
- [6] B. P. English et al., “Ever-fluctuating single enzyme molecules: Michaelis-Menten equation revisited,” *Nature Chemical Biology* **2**, 87 (2006).
- [7] L. Edman and R. Rigler, “Memory landscapes of single-enzyme molecules,” *PNAS* **97**, 8266 (2000).
- [8] R. E. Joseph et al., “Allosteric communication in BTK,” *eLife* **9**, e60470 (2020).
- [9] A. Taldaev et al., “A kinetic view of imatinib resistance in chronic myeloid leukemia,” *PNAS* **118**, e2106566118 (2021).
- [10] S. Horvath, “DNA methylation age of human tissues and cell types,” *Genome Biol.* **14**, R115 (2013).
- [11] Y. Lu et al., “Reprogramming to recover youthful epigenetic information and restore vision,” *Nature* **588**, 124 (2020).

- [12] H. X. Chao et al., “Evidence that the human cell cycle is a series of uncoupled, memoryless phases,” *Mol. Syst. Biol.* **15**, e8604 (2019).
- [13] K. K. Bhatt et al., “Chk1-S is a splice variant and endogenous inhibitor of Chk1,” *PNAS* **109**, 197 (2012).
- [14] H. P. Lu, L. Xun, X. S. Xie, “Single-molecule enzymatic dynamics,” *Science* **282**, 1877 (1998).
- [15] H. X. Chao, C. E. Poovey, A. A. Privette, G. D. Grant, H. Y. Chao, J. G. Cook, J. E. Purvis, “Orchestration of DNA damage checkpoint dynamics across the human cell cycle,” *Cell Systems* **5**, 445 (2017).
- [16] S. Chakraborty, E. Ahler, J. J. Simon, L. Fang, Z. E. Potter, K. A. Sitko, J. J. Stephany, M. Guttman, D. M. Fowler, D. J. Maly, “Profiling of drug resistance in Src kinase at scale uncovers a regulatory network coupling autoinhibition and catalytic domain dynamics,” *Cell Chemical Biology* **31**, 207 (2024).
- [17] D. Wilsker, E. Petermann, T. Helleday, F. Bunz, “Essential function of Chk1 can be uncoupled from DNA damage checkpoint and replication control,” *PNAS* **105**, 20752 (2008).
- [18] A. Maybaum, “Non-Markovian Dynamics in Computational Biology,” companion paper (2026).